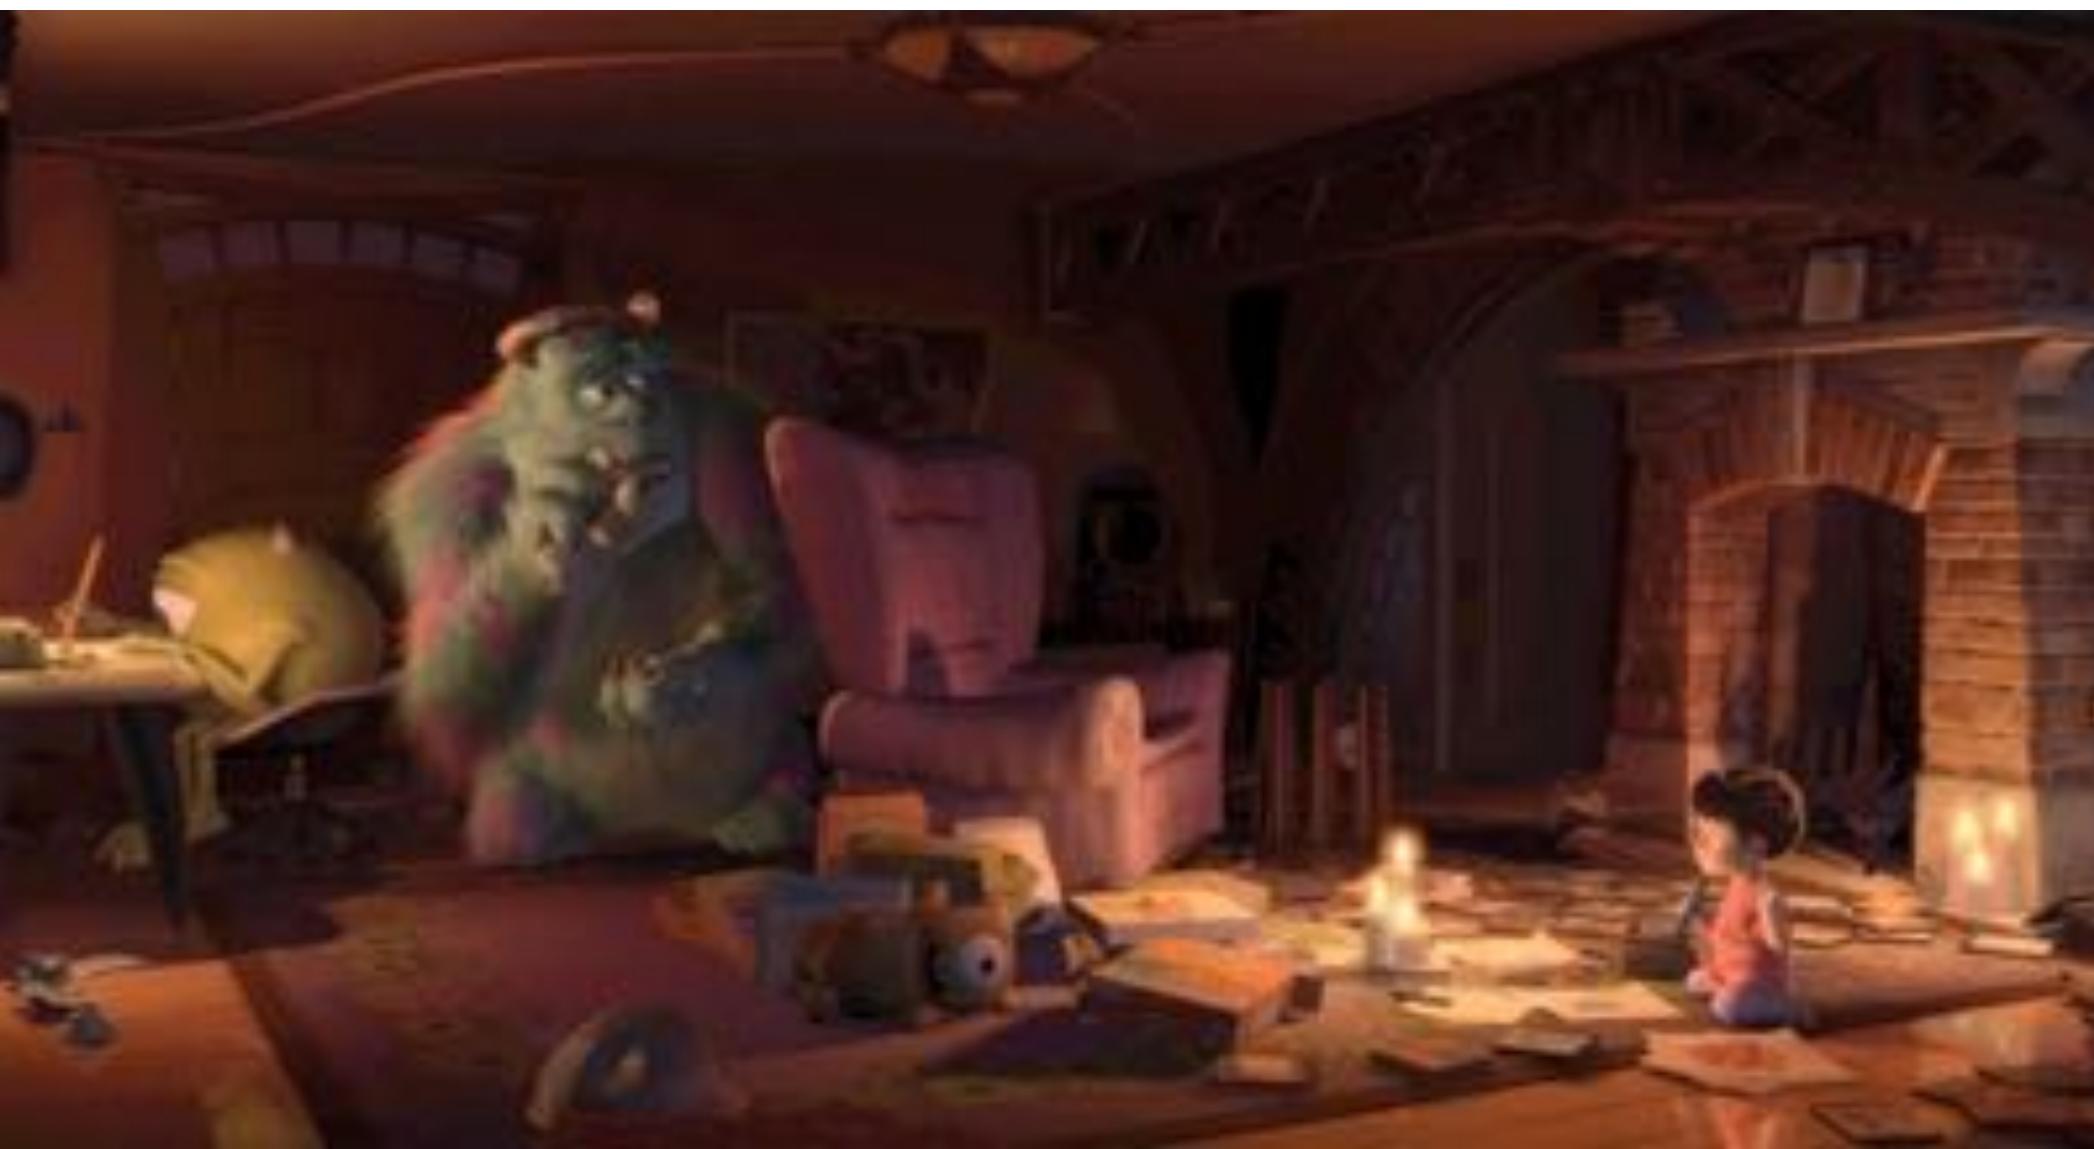




Lighting and Shading

CS 355: Introduction to Graphics and Image Processing



Kinds of Lighting

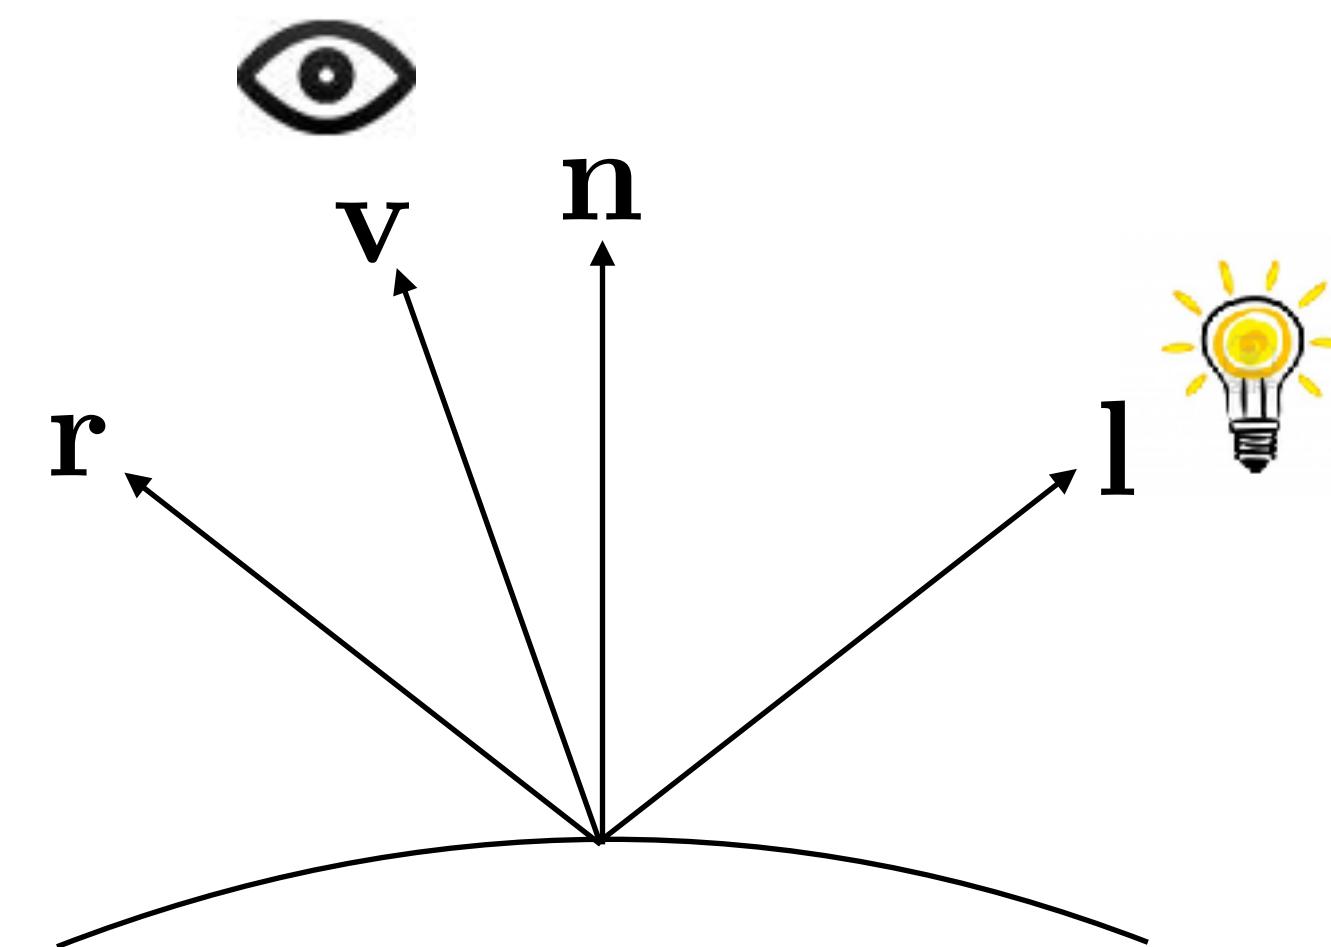
- Direct:
Light falling on an object directly from a light source
- Indirect:  More in CS 455
Light falling on an object after being reflected off (or going through) other objects
- Ambient:
General light bouncing around and scattered enough to be effectively “everywhere”

Light Sources

- Point
 - Area
 - Spot
 - and many other models...
- We'll focus on this for now**
- 

Basic Geometry of Lighting

- The surface **n**ormal
- The **l**ighting direction
(to the light)
- The **v**iewing direction
(to the eye/camera)
- The **r**eflected light direction



Surface Reflectance

- Most objects don't give off light
 - reflect some of the light that falls on them
 - absorb the rest
- The wavelengths reflected give the object its color
- The effect is multiplicative: i.e., "reflects 40% of the green light"
- If we model the light as RGB, we can also model reflectance as RGB
- Reflectance is also sometimes called *albedo*

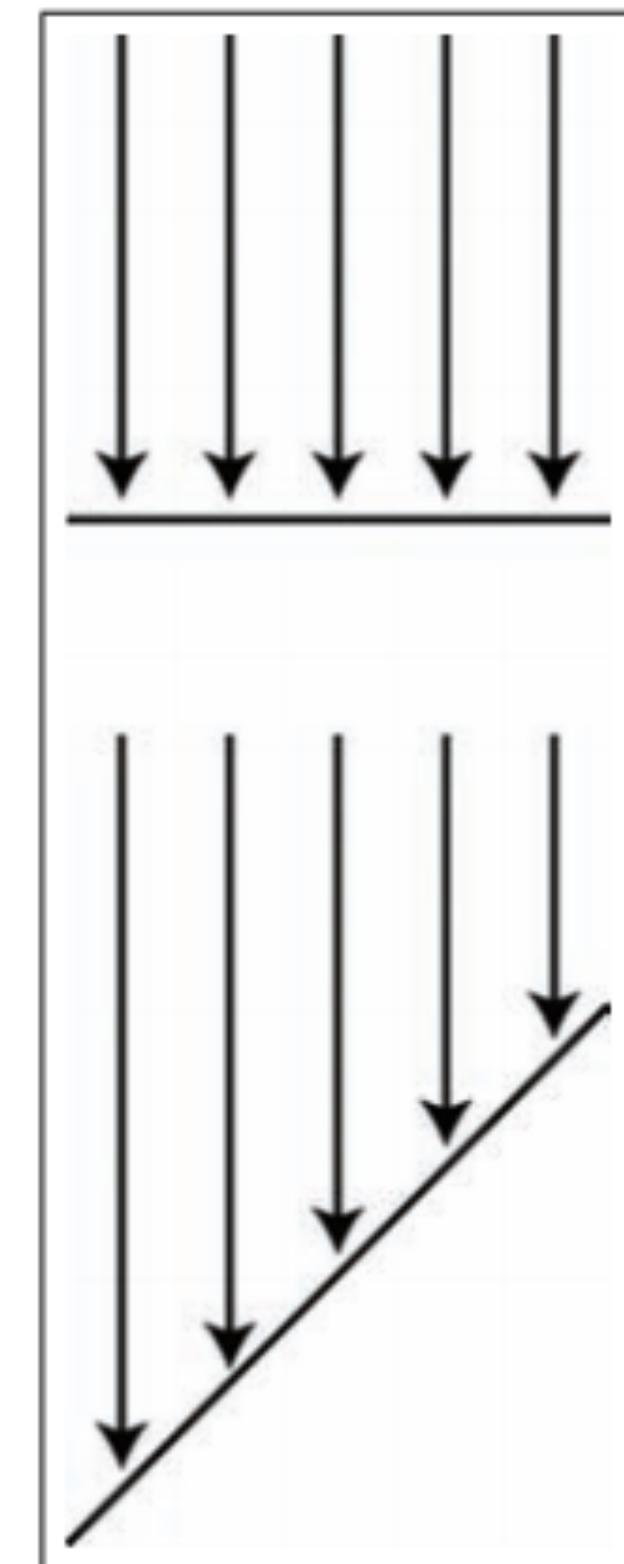
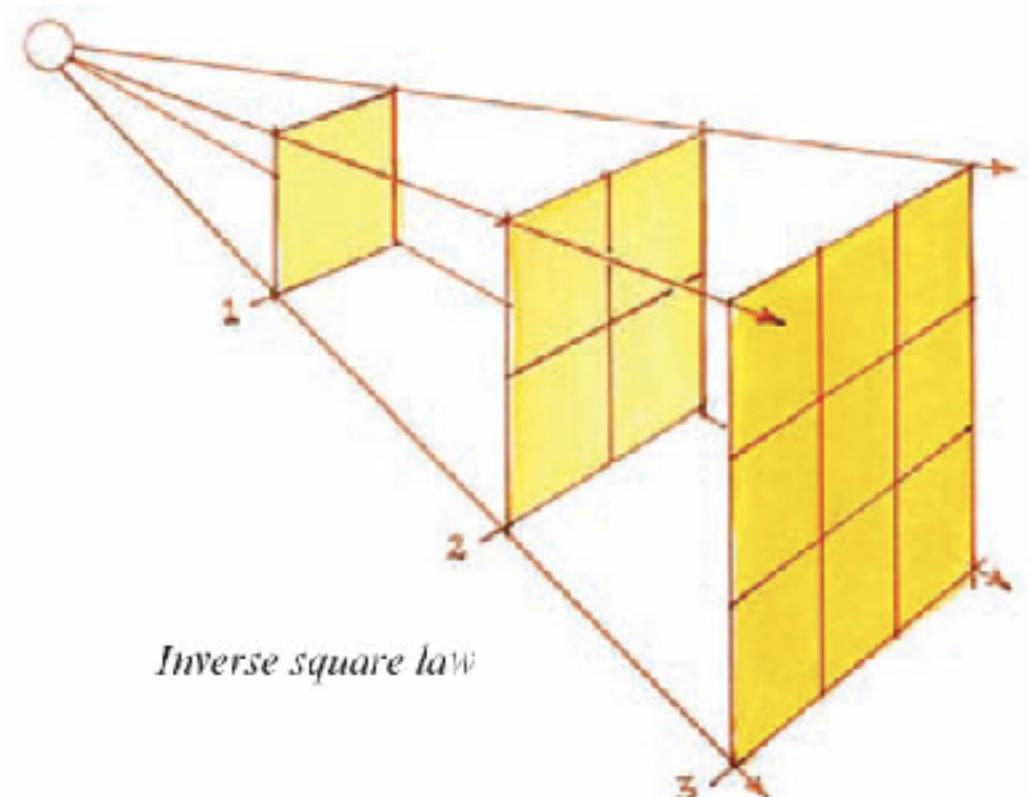
Irradiance

- We sometimes say “the amount of light”
- But it’s really how much light *per unit area*
- This quantity is called the *irradiance*

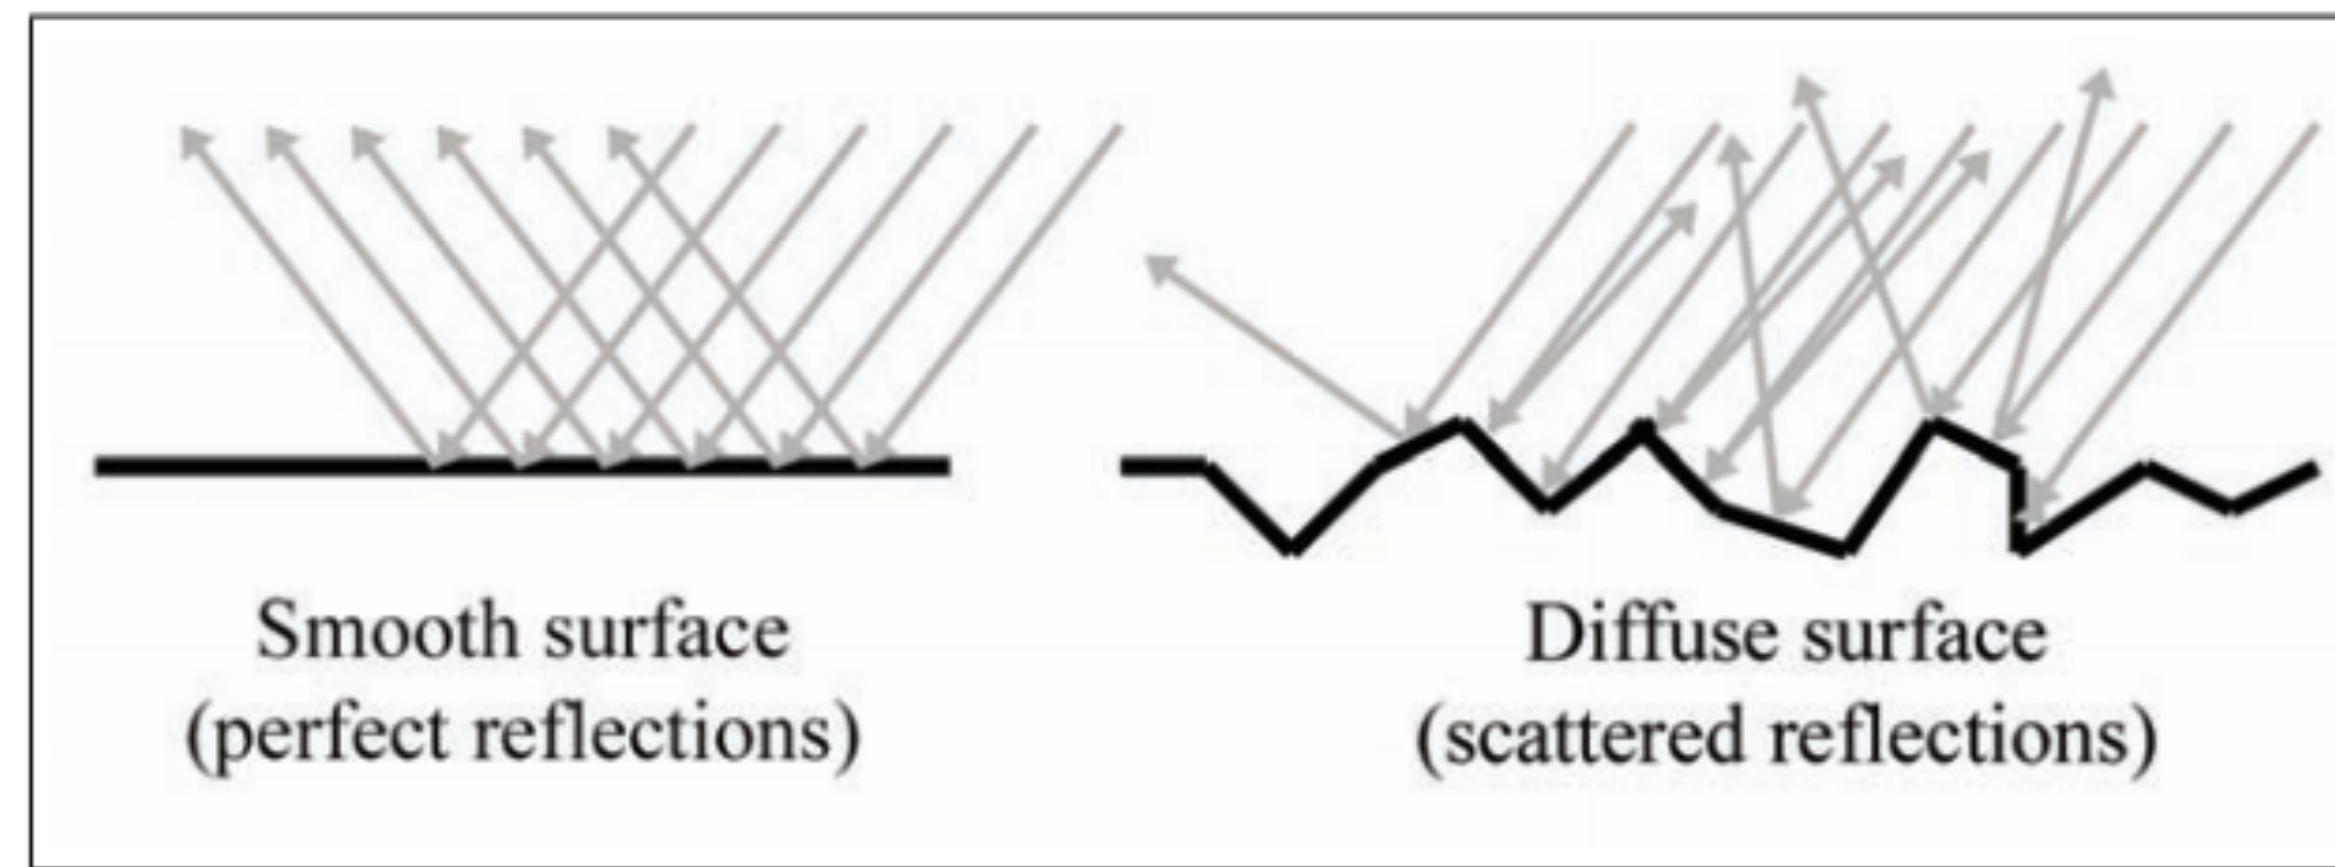


Irradiance

- Two important properties:
 - Irradiance falls off with the square of the distance
 - Irradiance is less when falling on a slanted surface



Specular vs. Diffuse

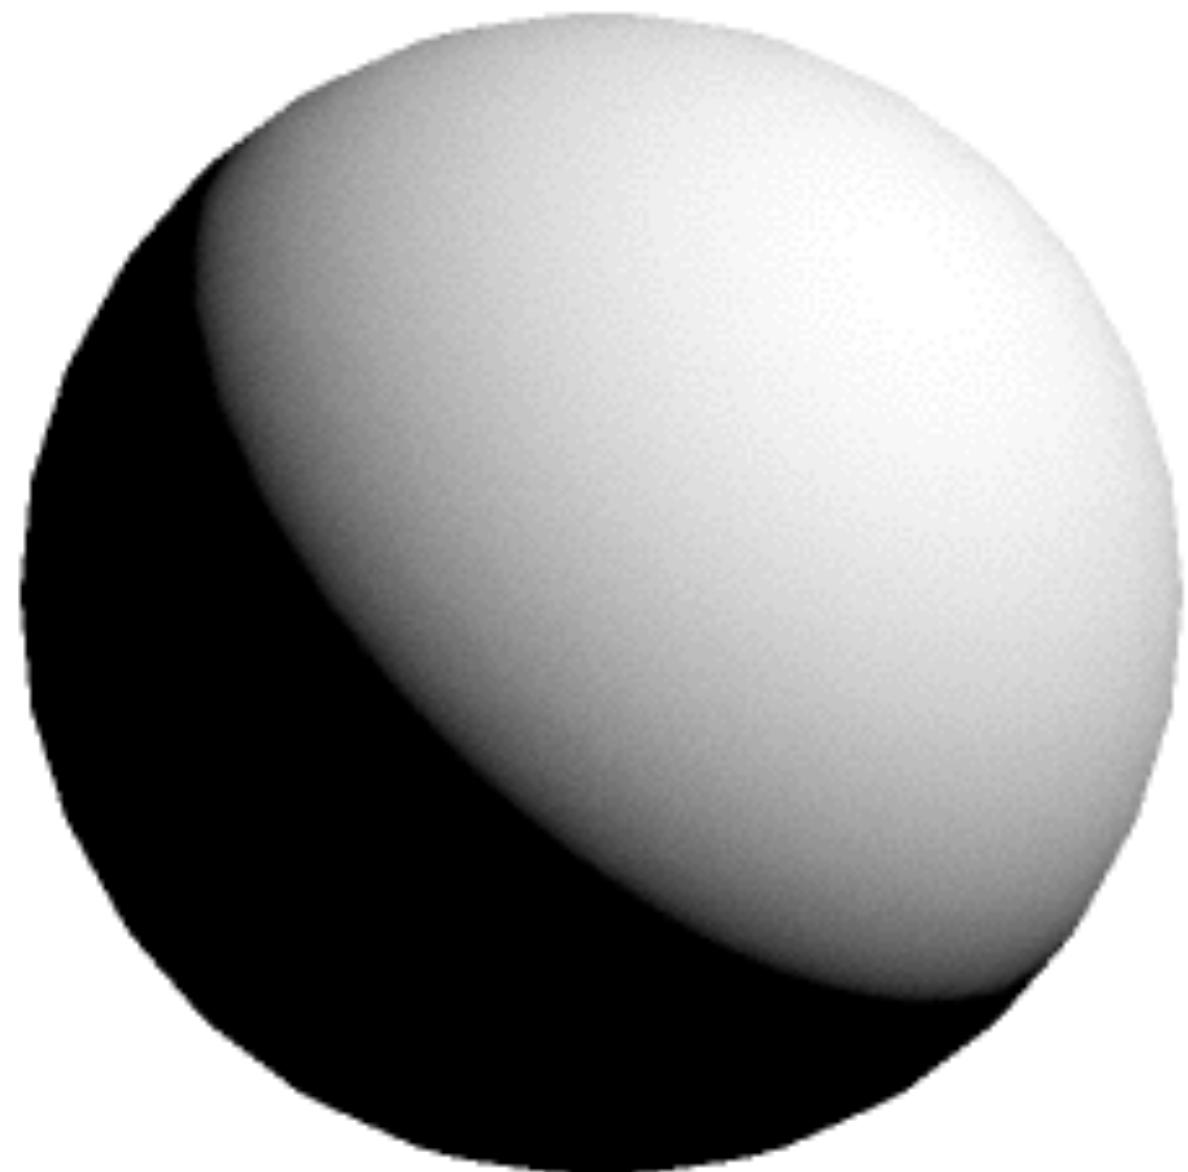


Some light is
reflected perfectly
(specular)

Some light is
scattered
(diffuse)

Diffuse Reflection

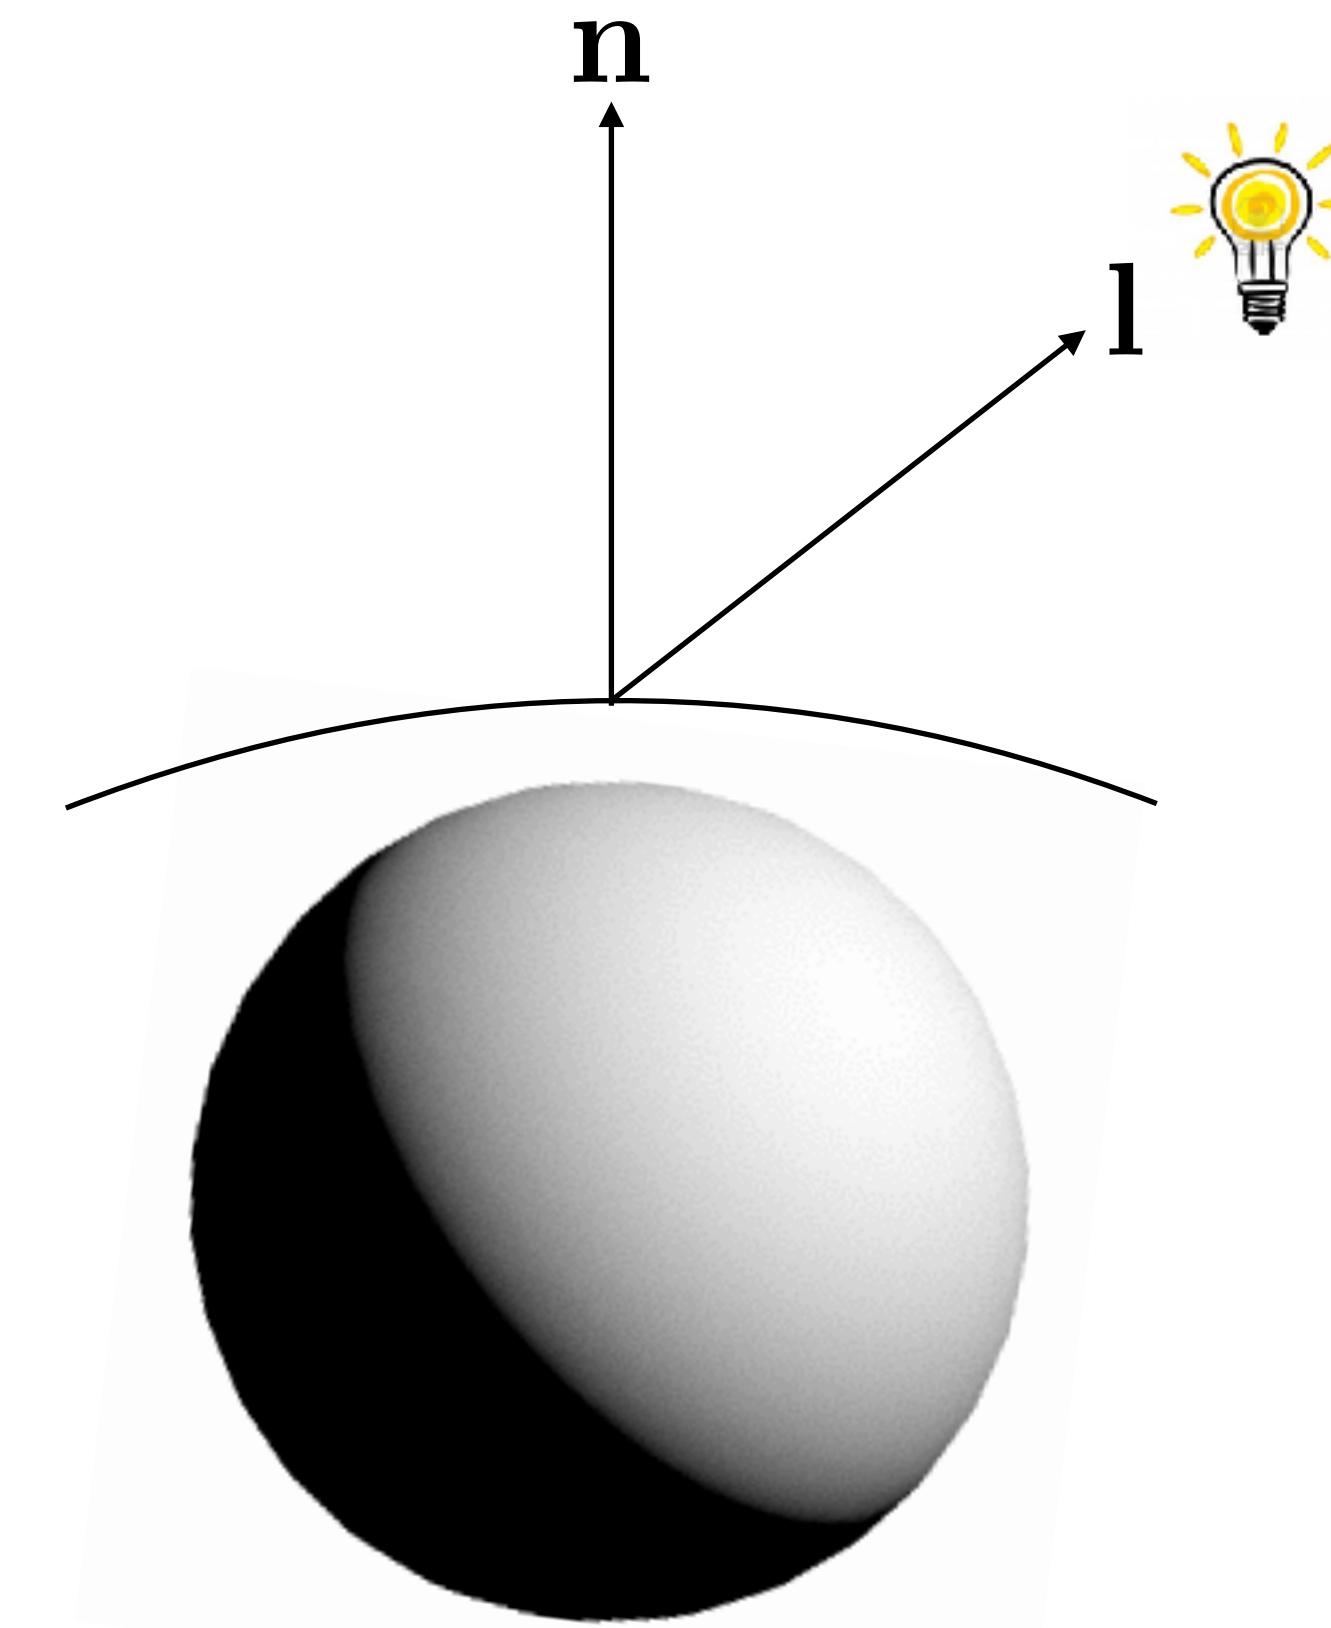
- Light scattered in every direction is called the *diffuse* part of the reflected light
- A perfectly diffuse surface is called *Lambertian*
- Only lighting direction matters
- Viewing direction *does not*



A Simple Diffuse Model

$$\text{pointwise multiply RGB}$$
$$c_{\text{diff}} = (s \otimes m_{\text{diff}})(n \cdot l)$$

diffuse reflected color source intensity material diffuse reflectance

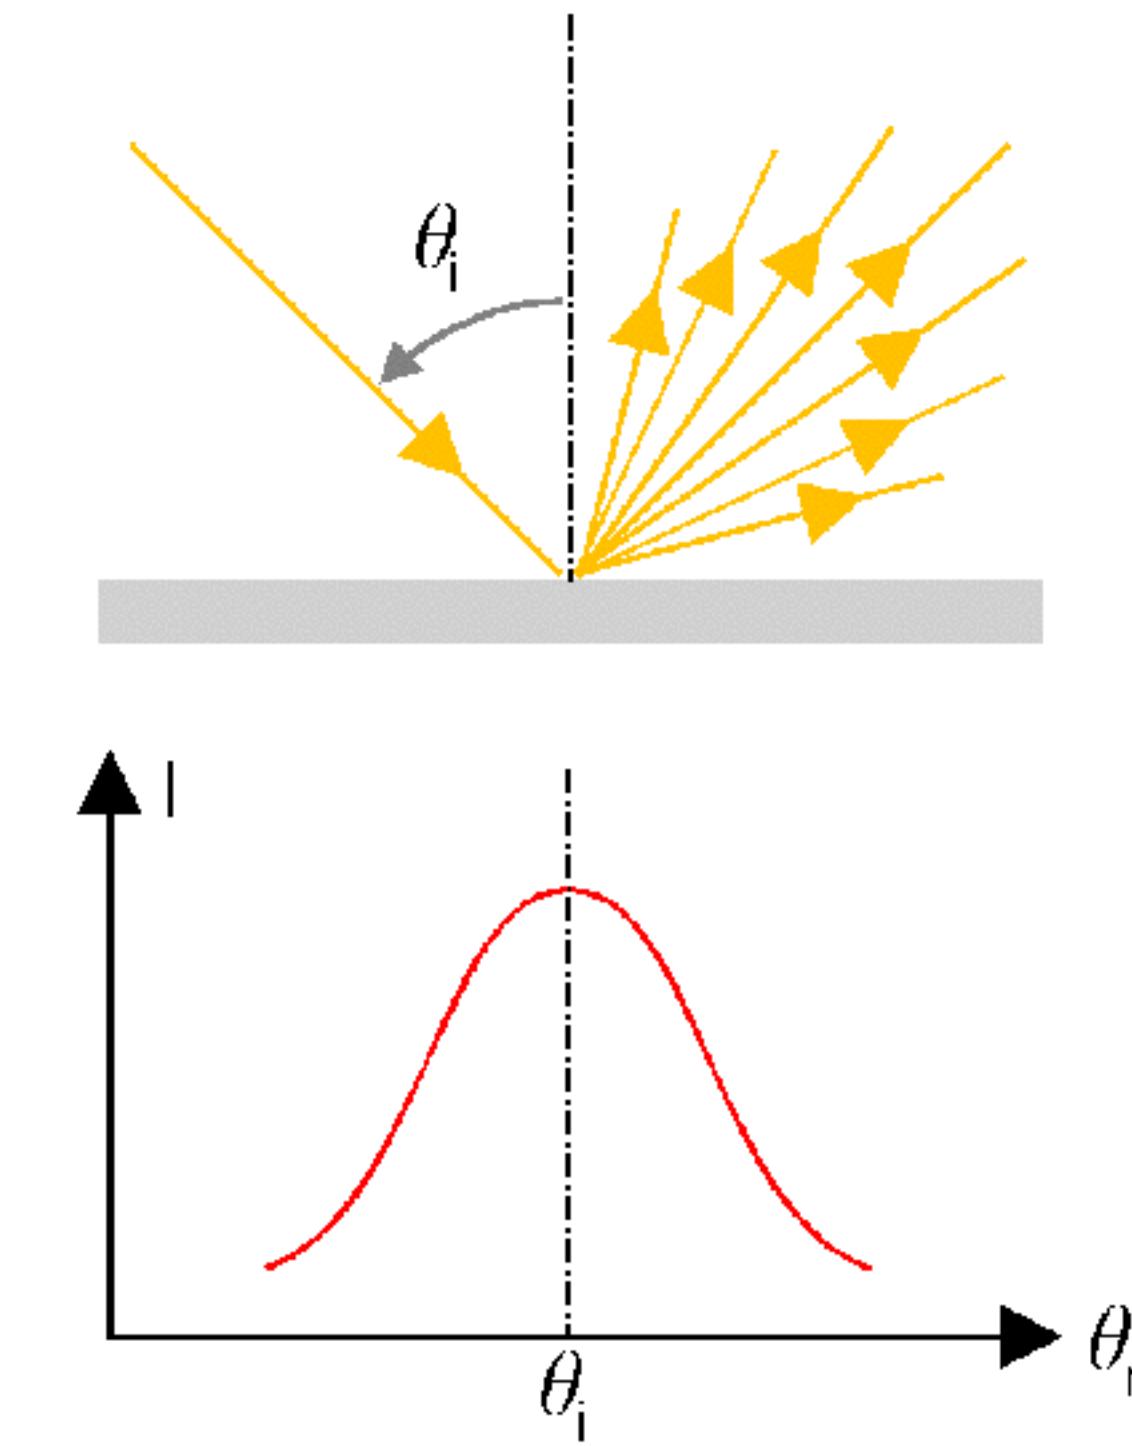
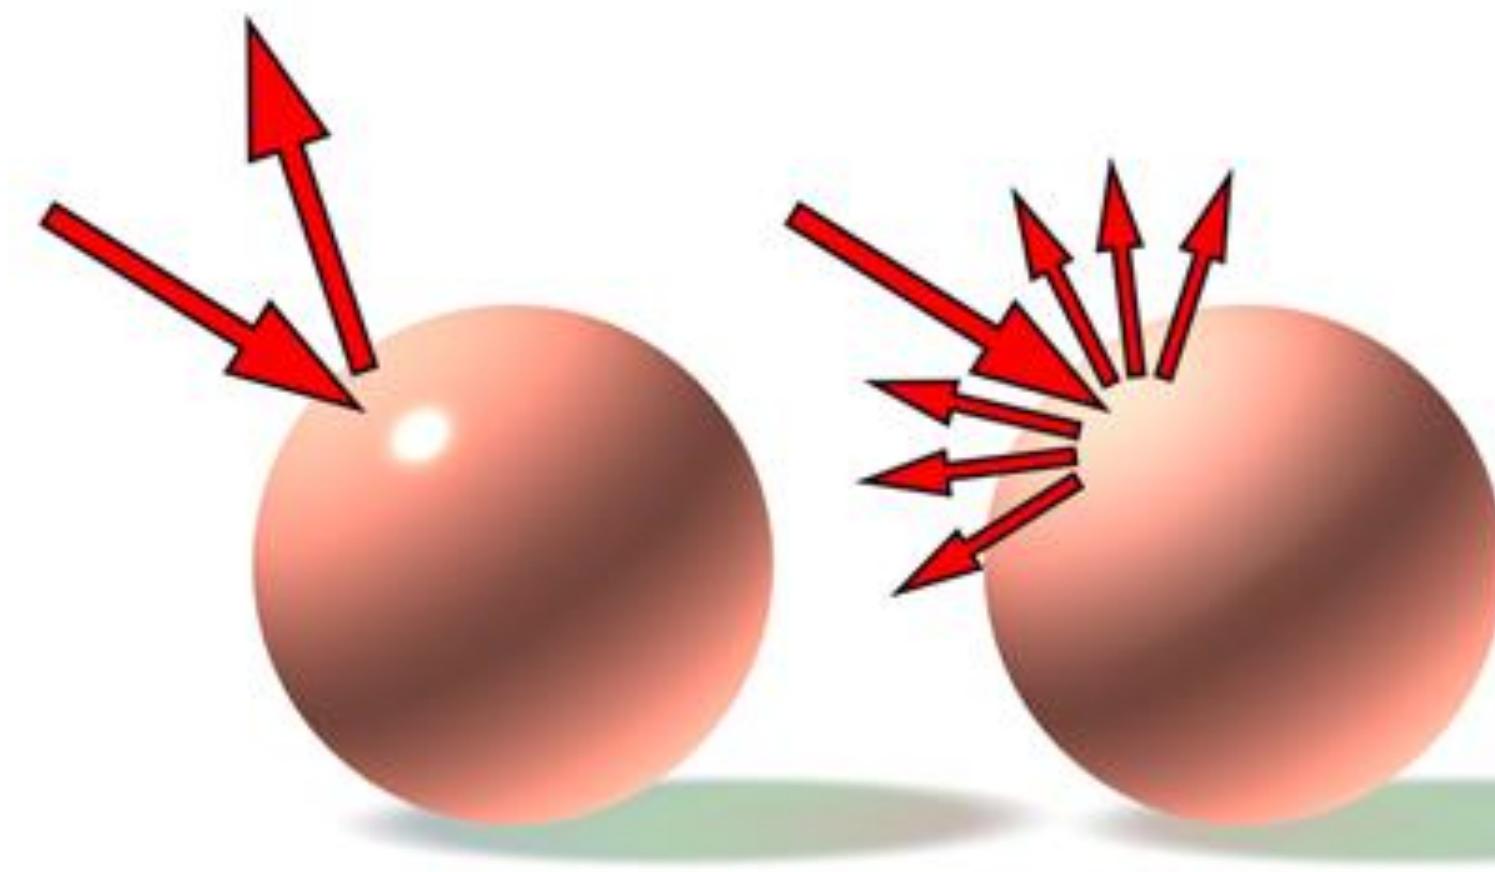


Assumes constant lighting direction and strength

Specular Reflections



Specular Reflections



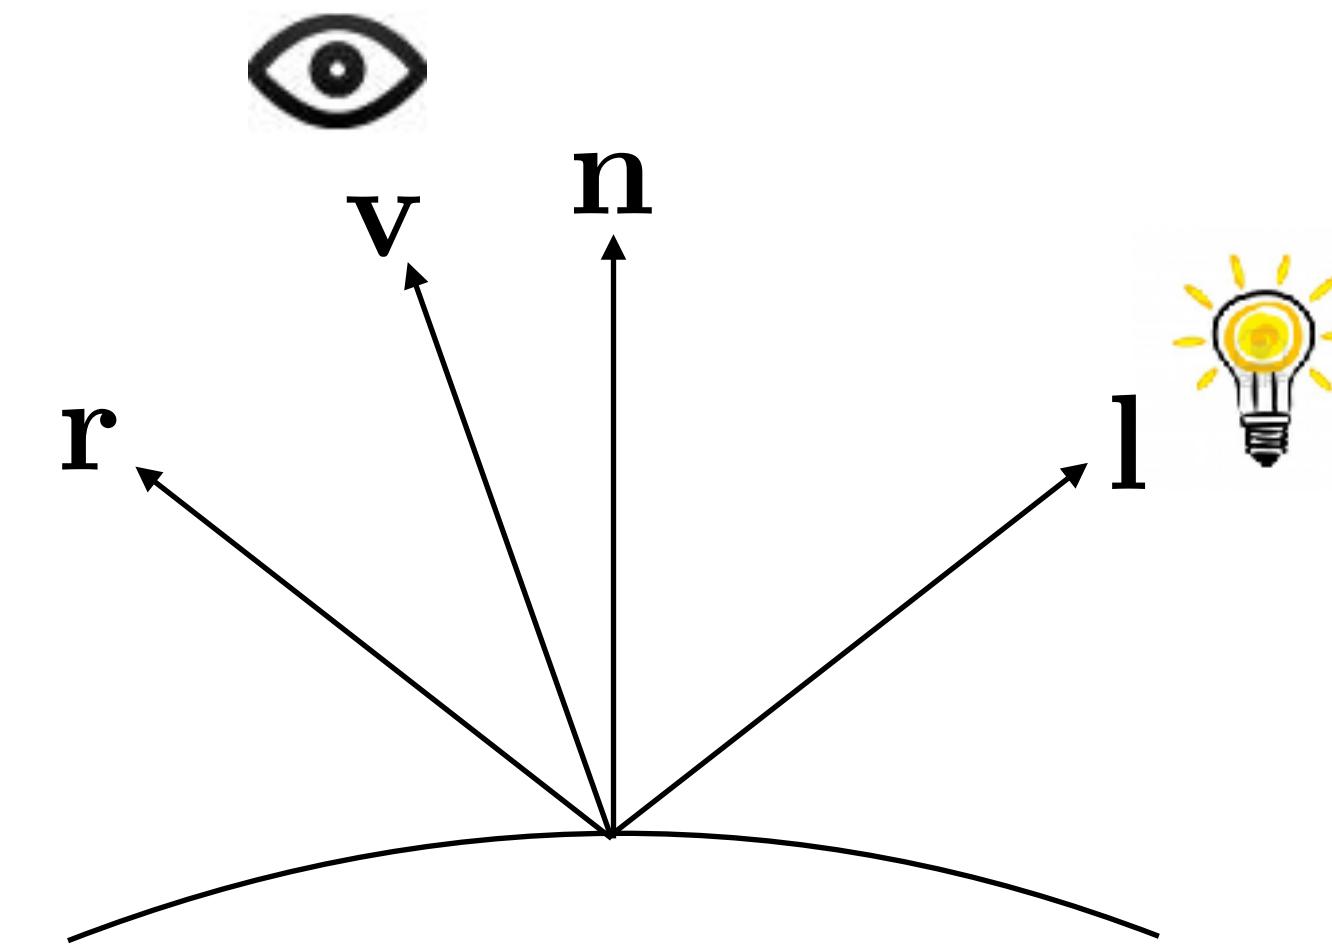
Angle of reflection = Angle of incidence
(but may be blurred)

Specular Reflections

how glossy the surface is

$$c_{\text{spec}} = (s \otimes m_{\text{spec}}) \boxed{(v \cdot r)}^{m_{\text{gls}}}$$

specular reflected color source intensity material specular reflectance



Specular Reflections



m_{gls}

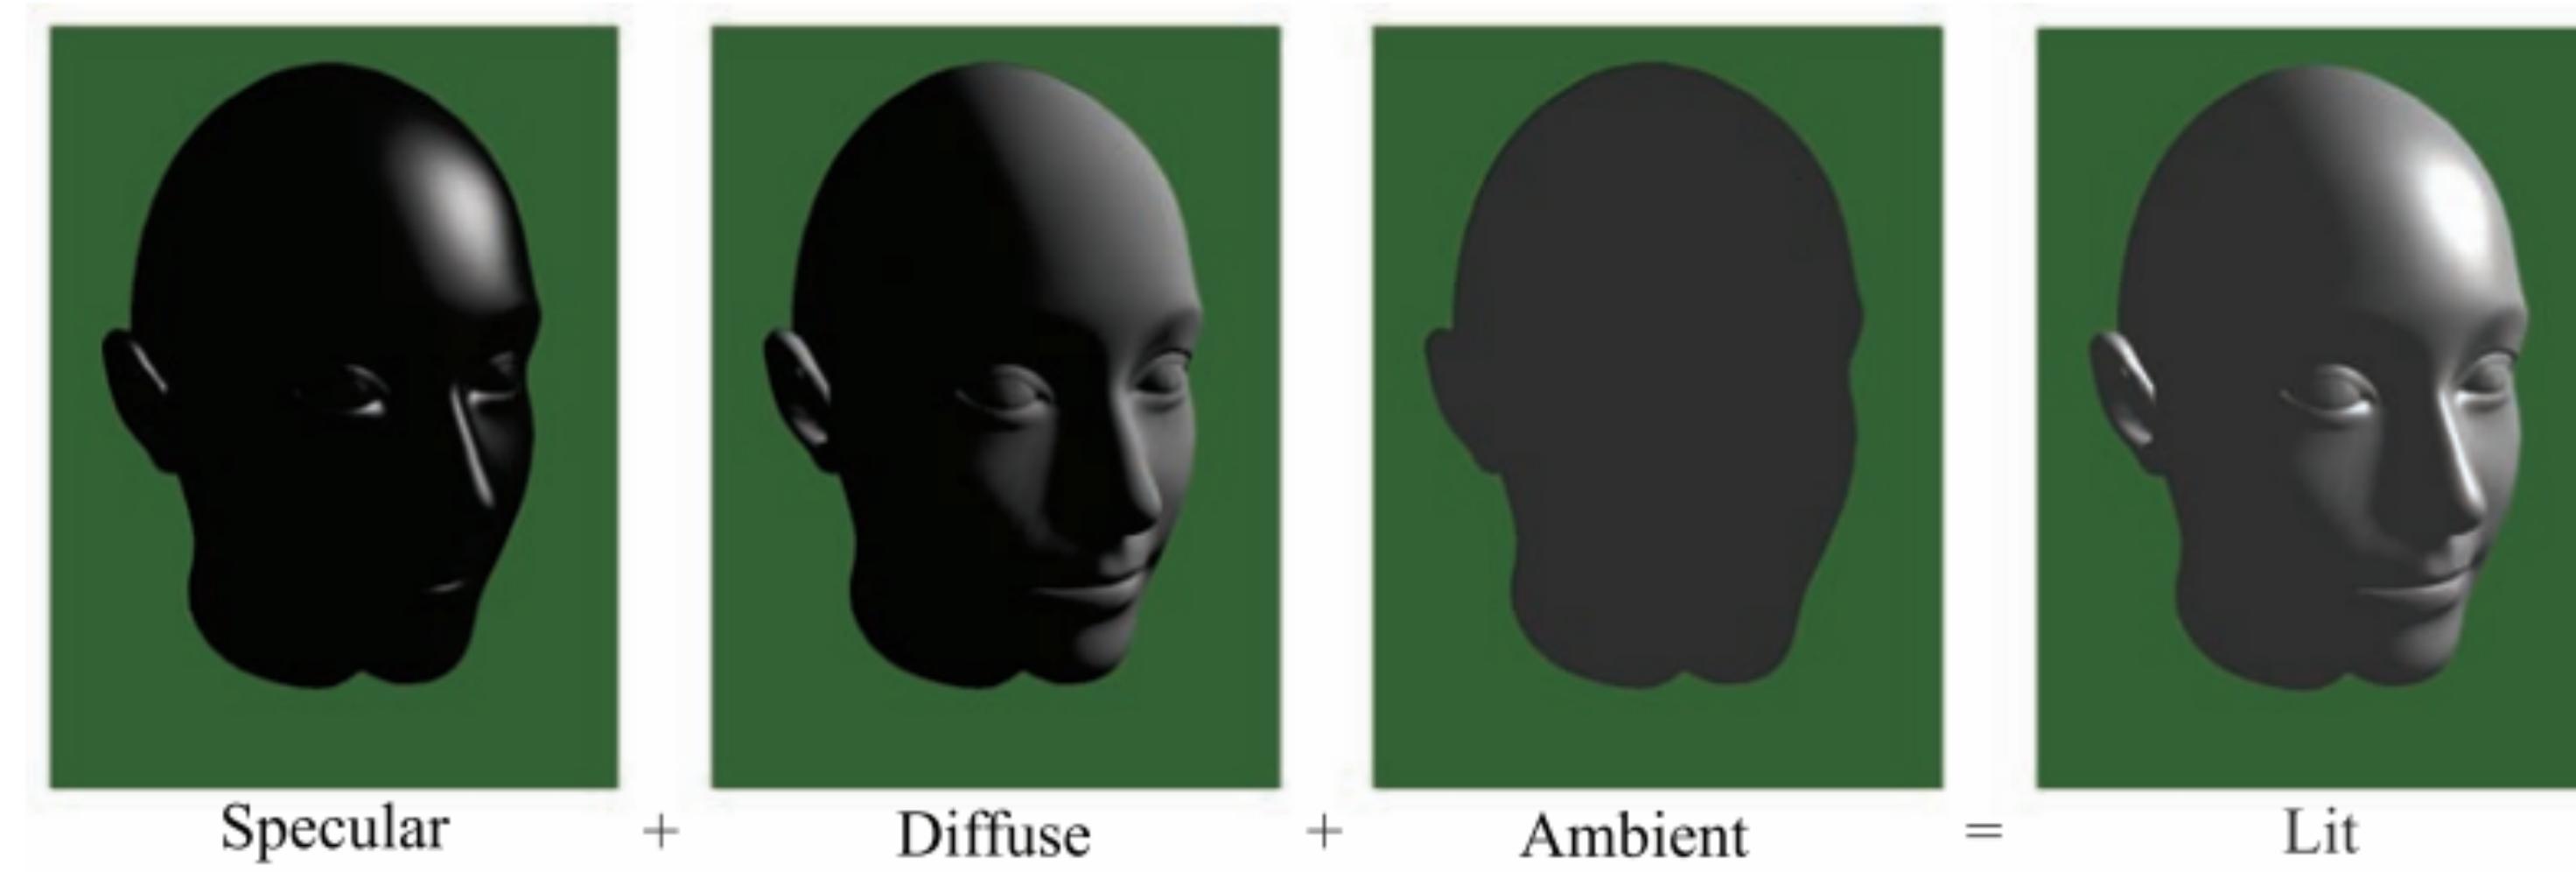
m_{spec}

Ambient Reflection

- Ambient light is “all around”, so directions don’t matter
- Just the product of the ambient light and the surface reflectance

$$\mathbf{c}_{\text{amb}} = \mathbf{s}_{\text{amb}} \otimes \mathbf{m}_{\text{amb}}$$

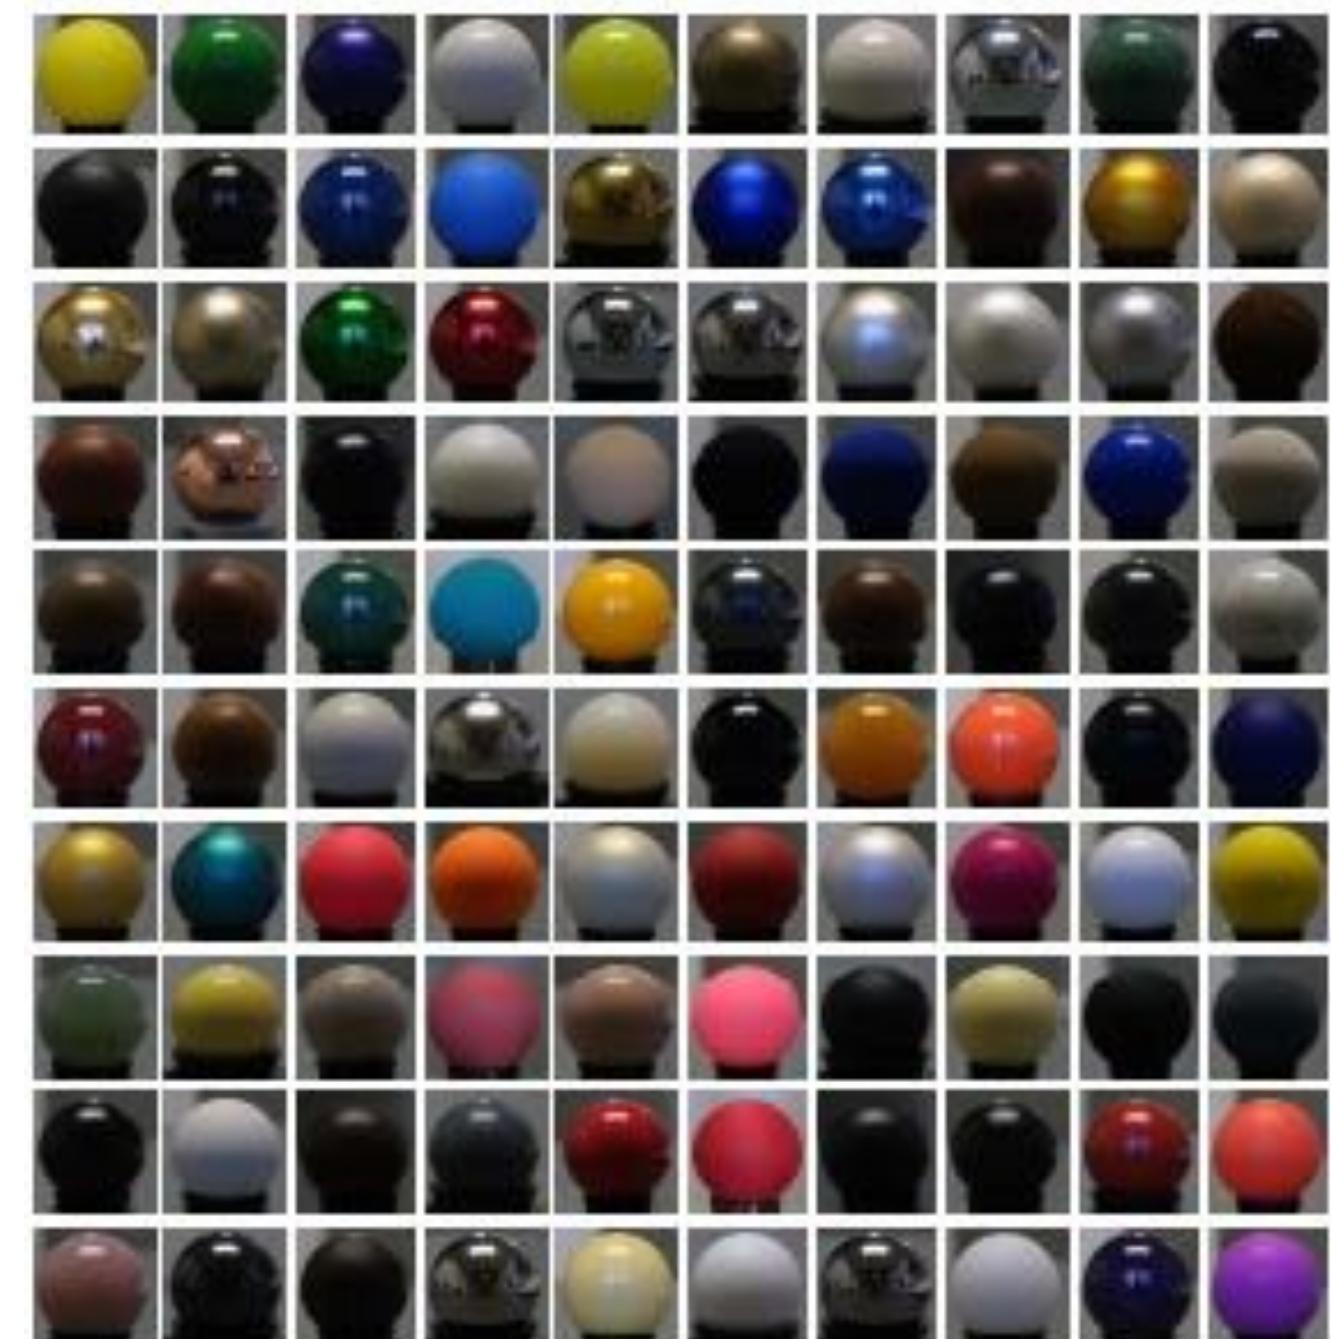
All Together Now...



This is called the *Phong* model
(the *Blinn* model is similar with slightly different specular)

BRDFs

- The Phong model is only an approximation
- Real reflections are not a simple mix of pure diffuse and pure specular
- Function of both incoming direction and outgoing direction
- Reflectance isn't constant across the surface

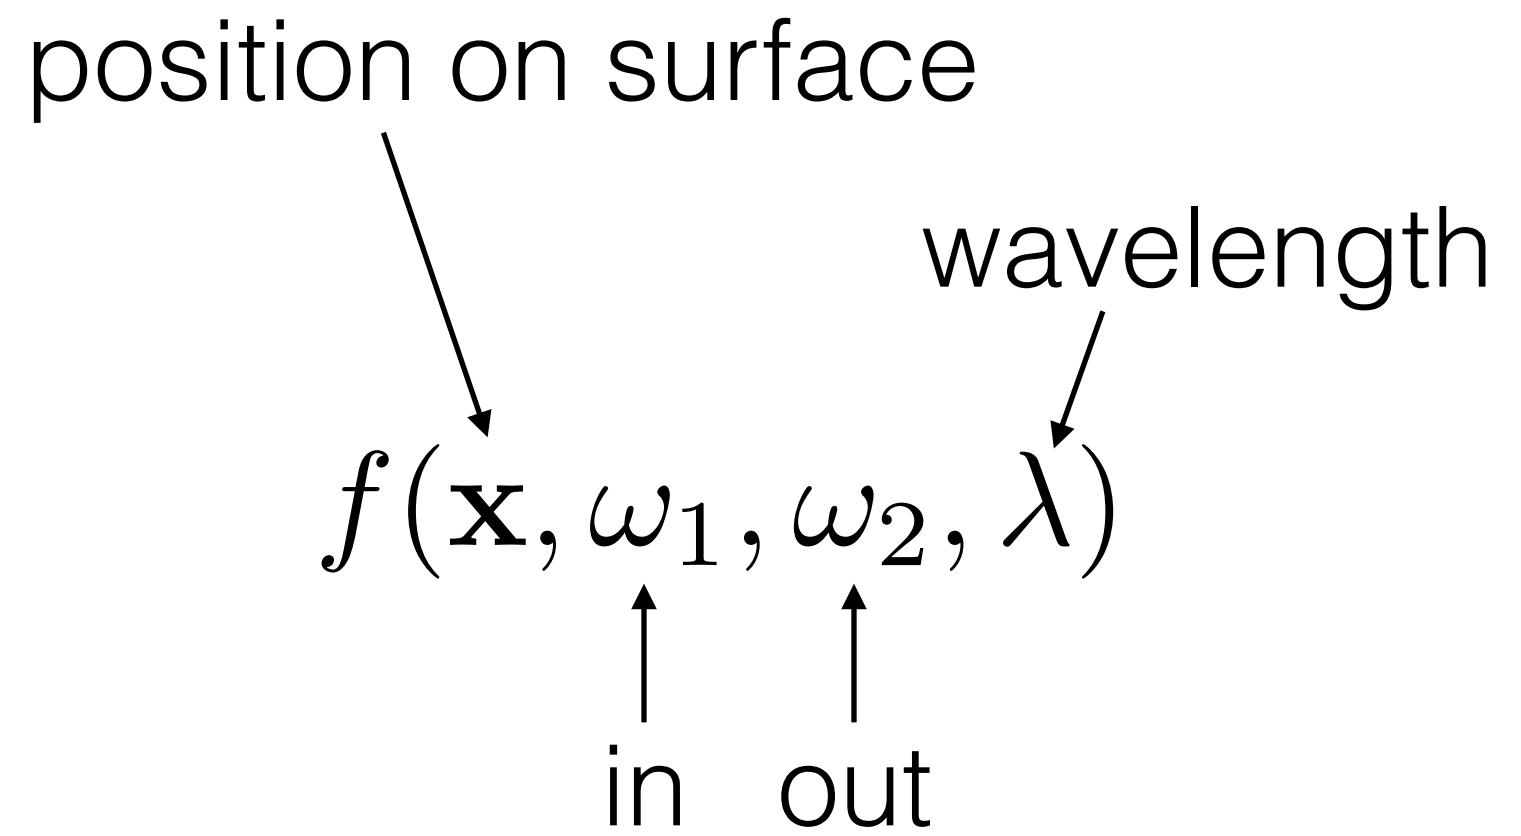


BRDFs

- The Phong model is only an approximation
- Not a simple mix of pure diffuse and pure specular
- Reflectance isn't constant across the surface
- Function of both incoming direction and outgoing direction

$$f(\mathbf{x}, \omega_1, \omega_2, \lambda)$$

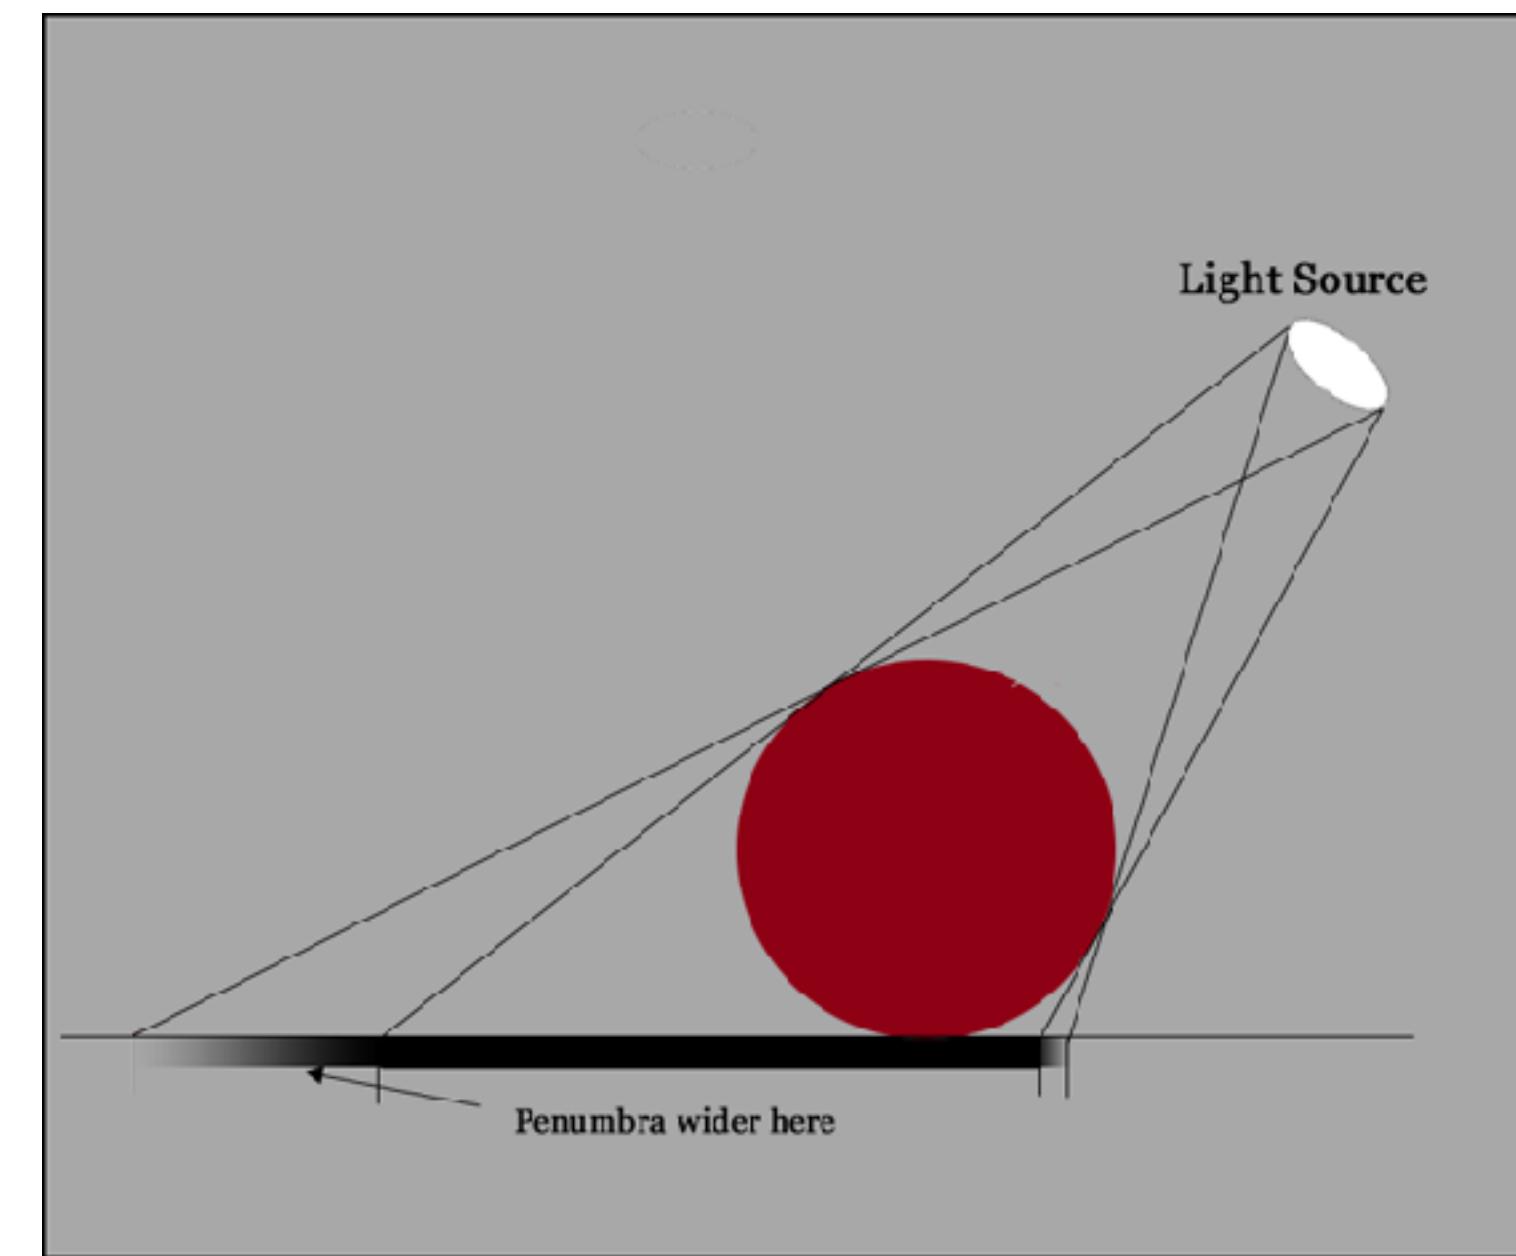
position on surface
wavelength
in out



Bidirectional
Reflectance
Distribution
Function

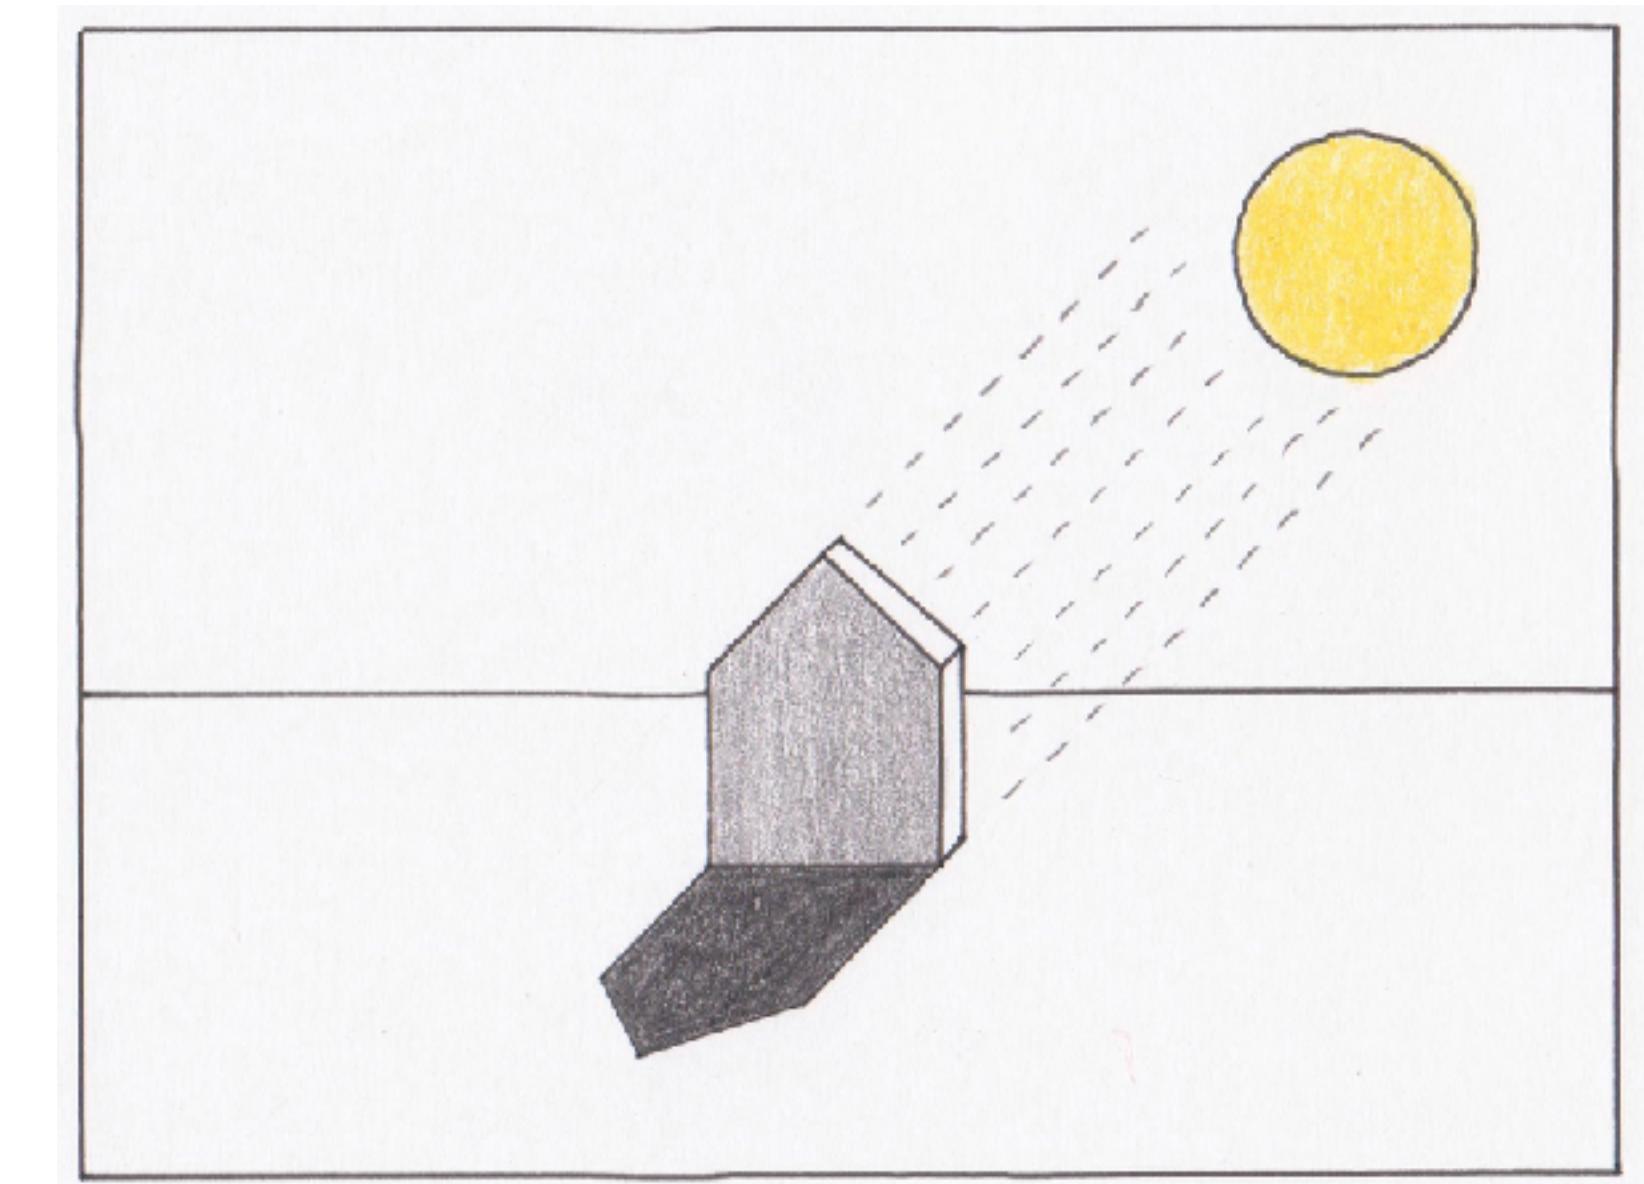
What About Shadows?

- Point lights cast hard shadows
- Area lights cast softer shadows
 - Umbra = area in full shadow
 - Penumbra = area in partial shadow



Simple Shadows

- For point lights, shadows are pretty simple
- *Do a visibility test from the point of view of the light!*
- Z-buffering can also be used for distance-based falloff



Coming up...

- More on lighting and shading



Lighting and Shading (cont'd)

CS 355: Introduction to Graphics and Image Processing

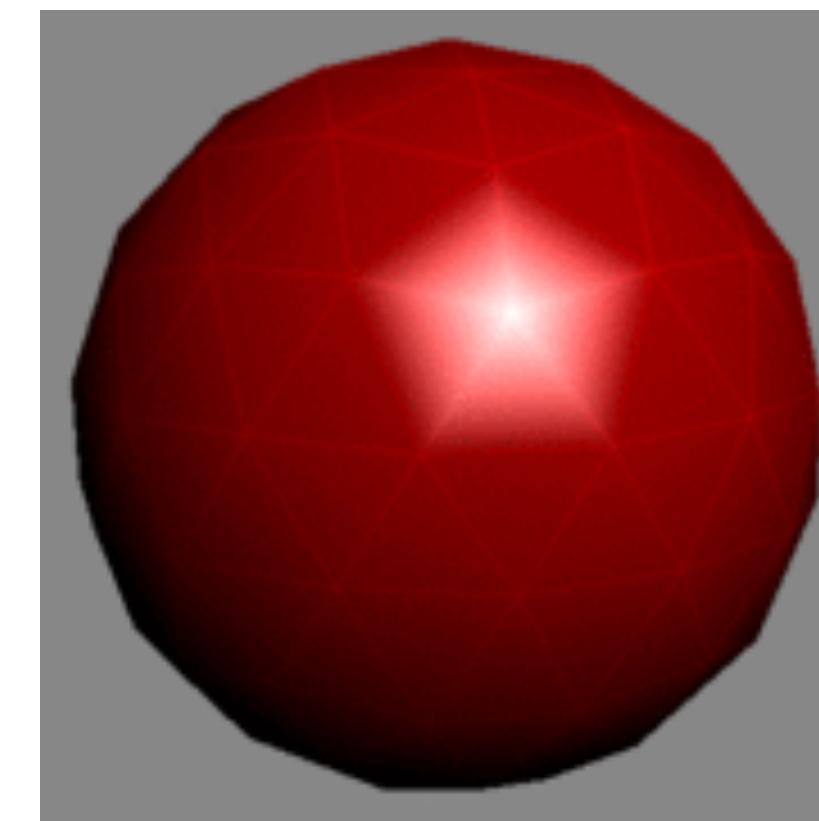
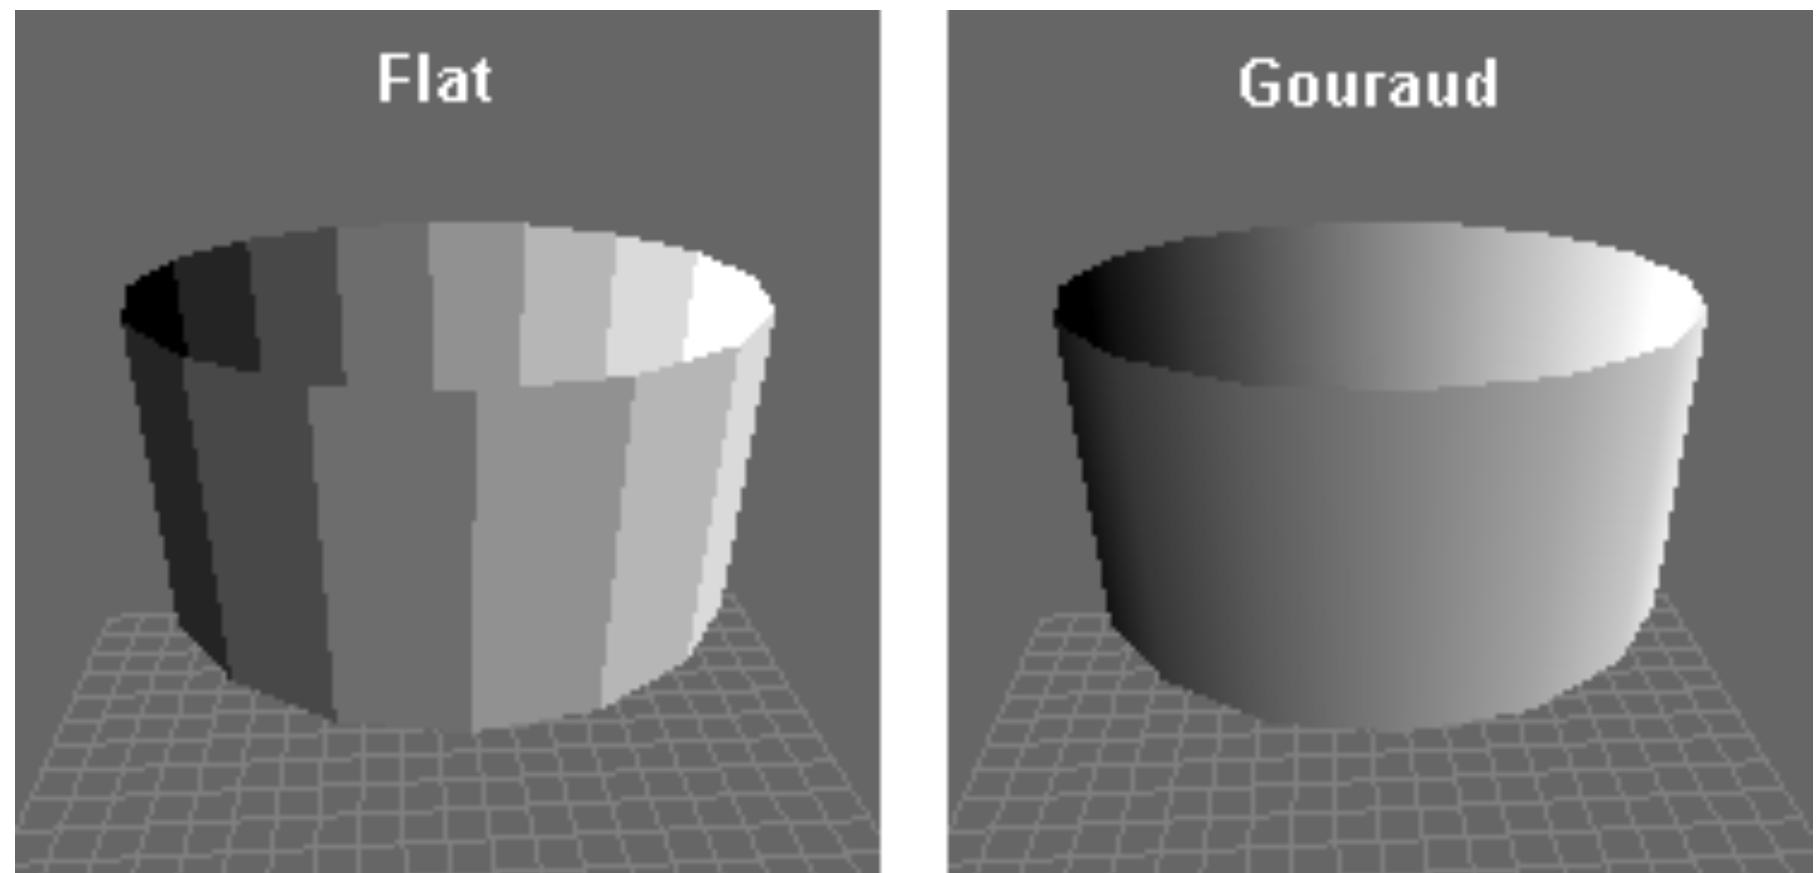
Flat Shading

- The entire polygon has the same normal, so it's all colored the same
- Leads to “flat shading”



Gouraud Shading

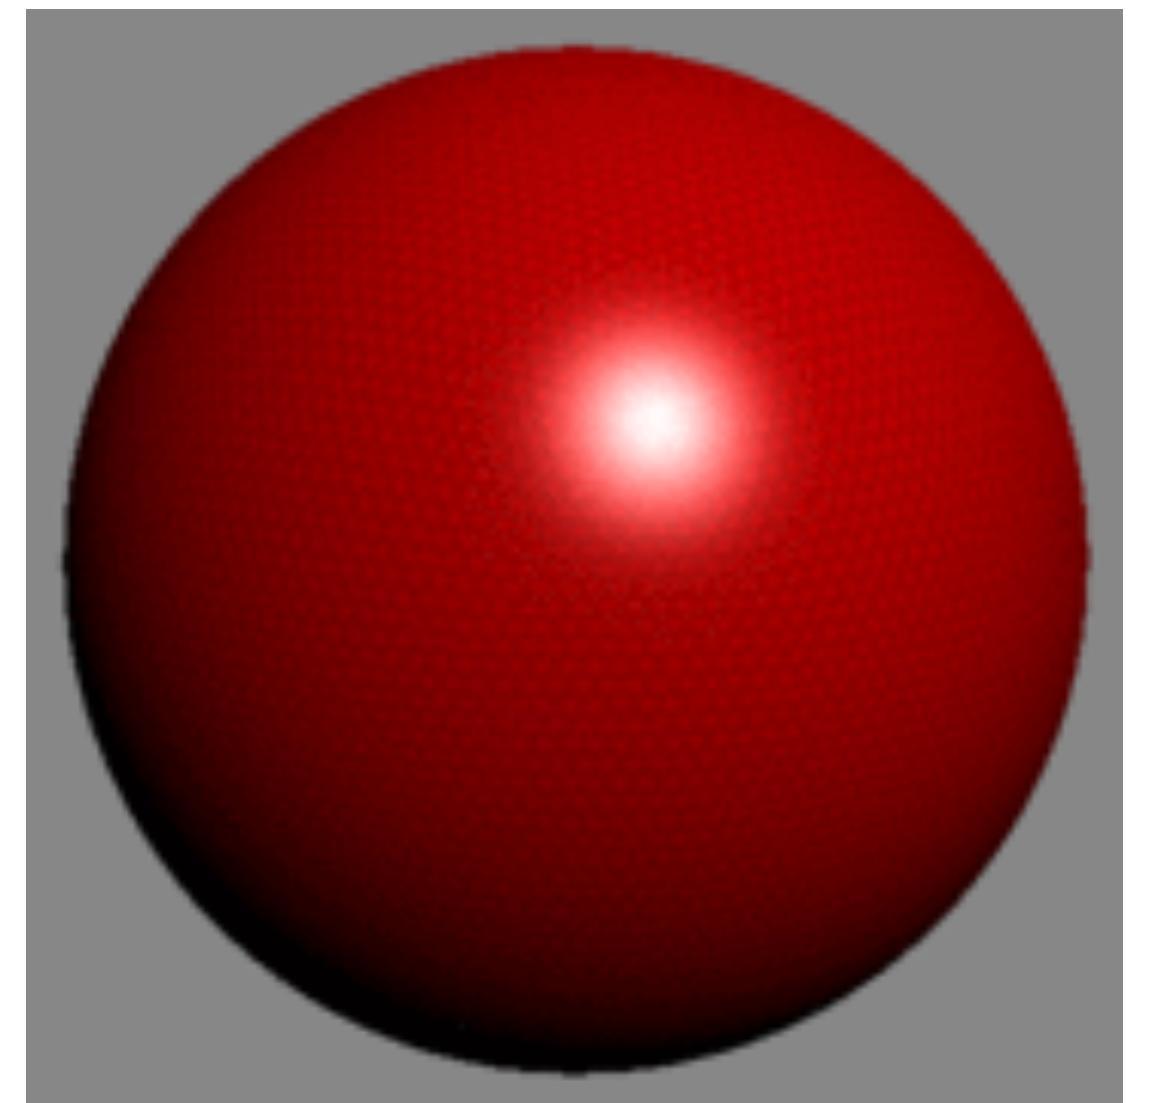
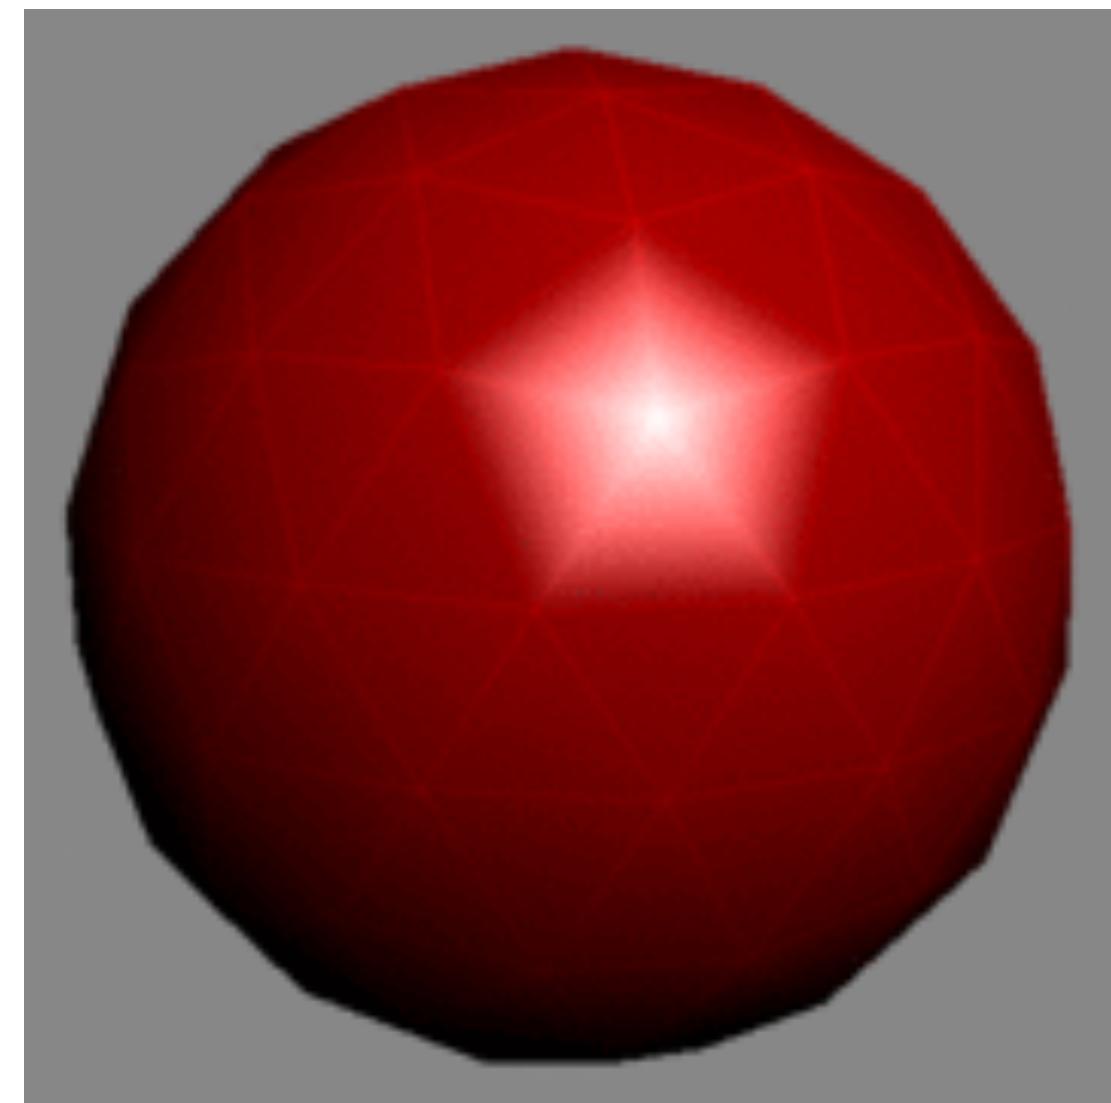
- Simple idea:
 - Compute shading at edges/vertices and interpolate
 - Works well for diffuse
 - Doesn't work as well for specular



Henri Gouraud, 1971

Phong Shading

- Key idea:
 - Instead of interpolating the shading, *interpolate the normals*
 - Compute shading on a per-pixel basis using interpolated normal



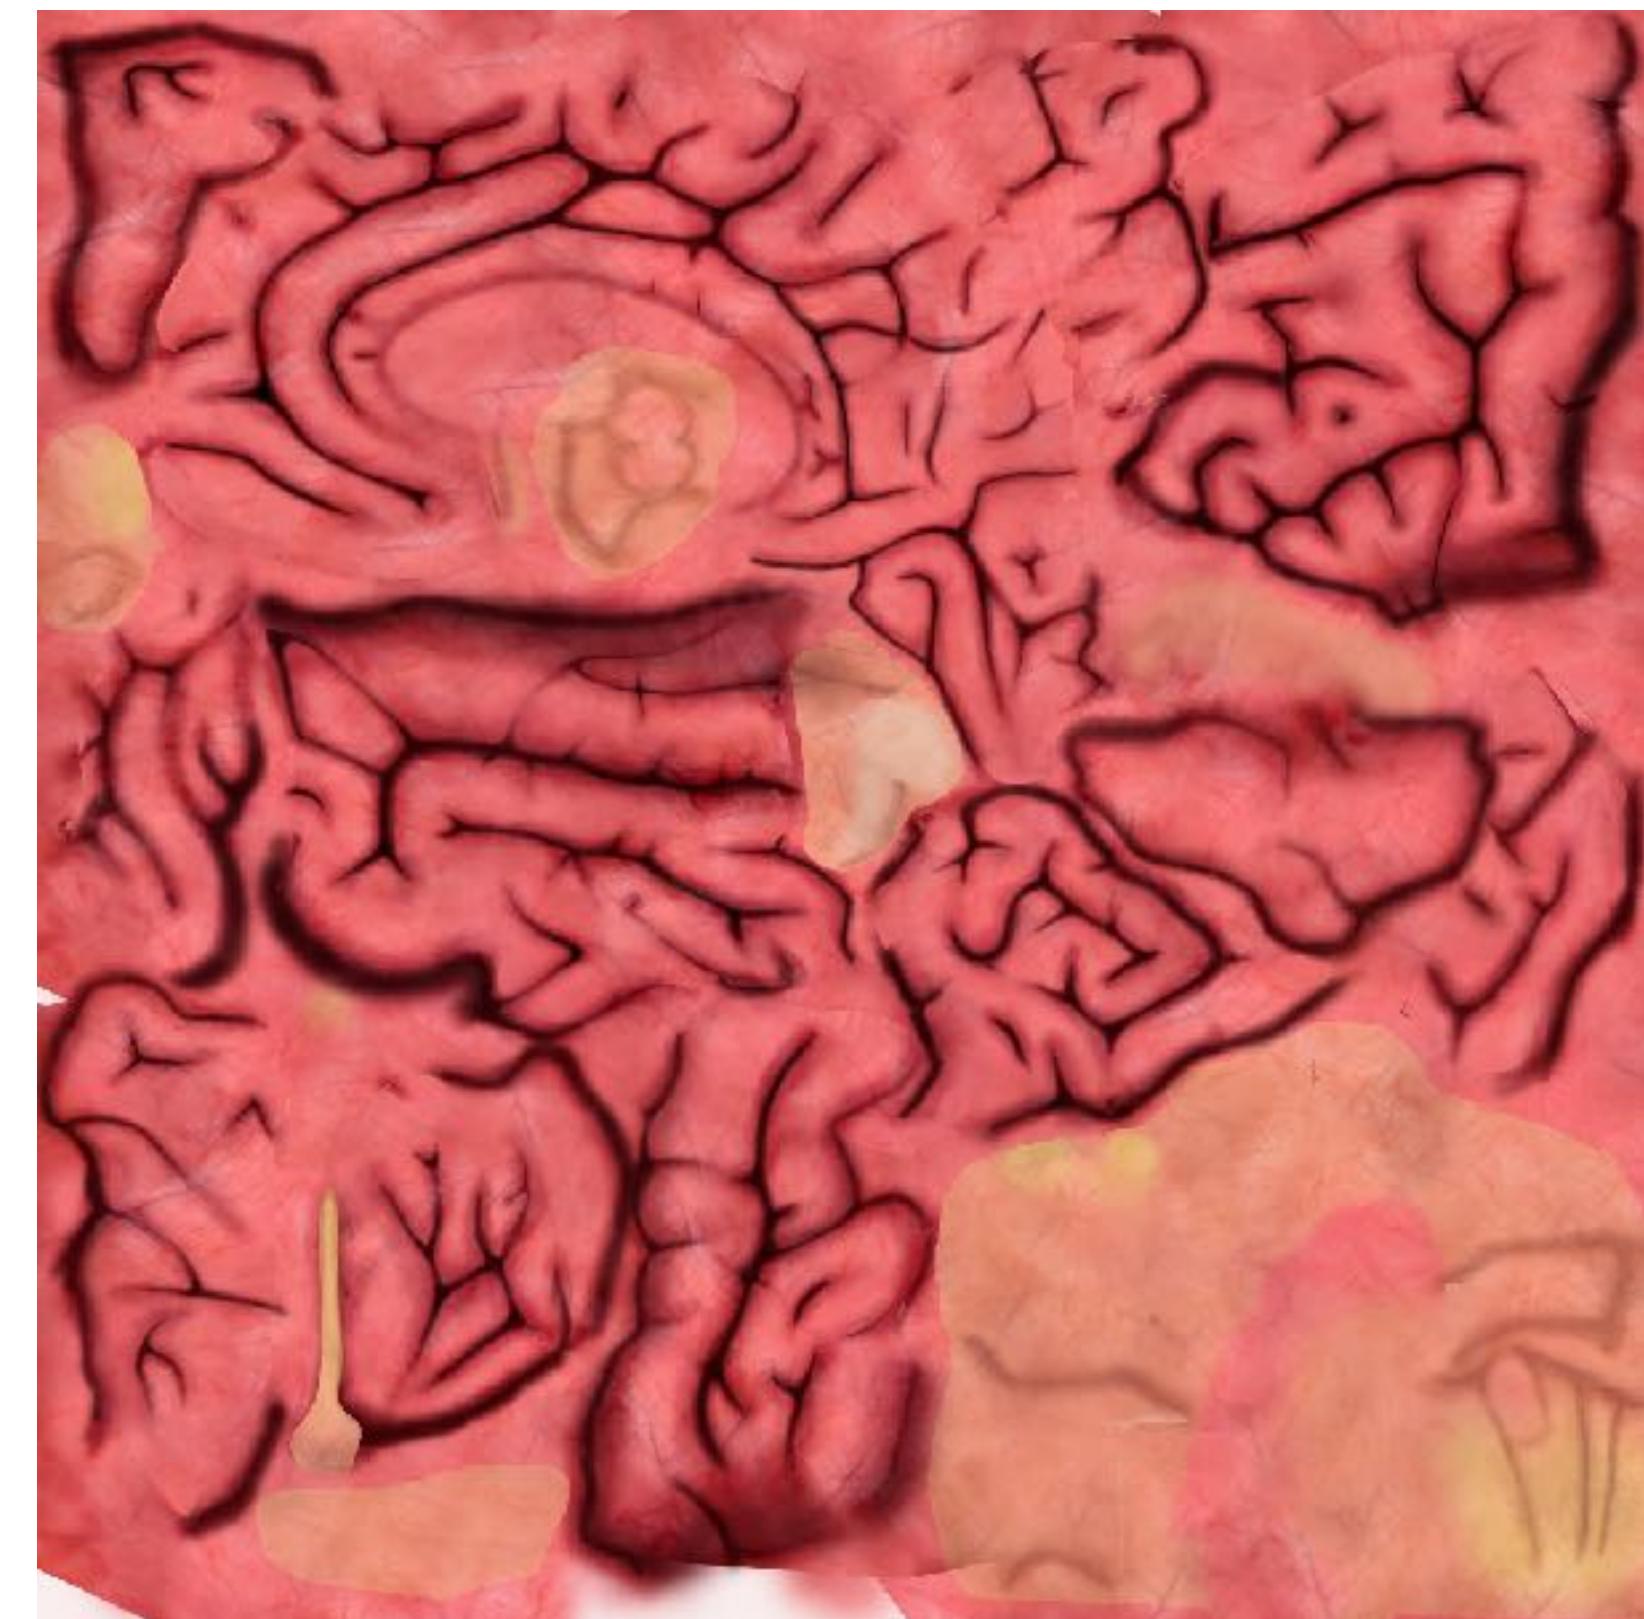
Bui Tan Phong, 1973

Texture Maps



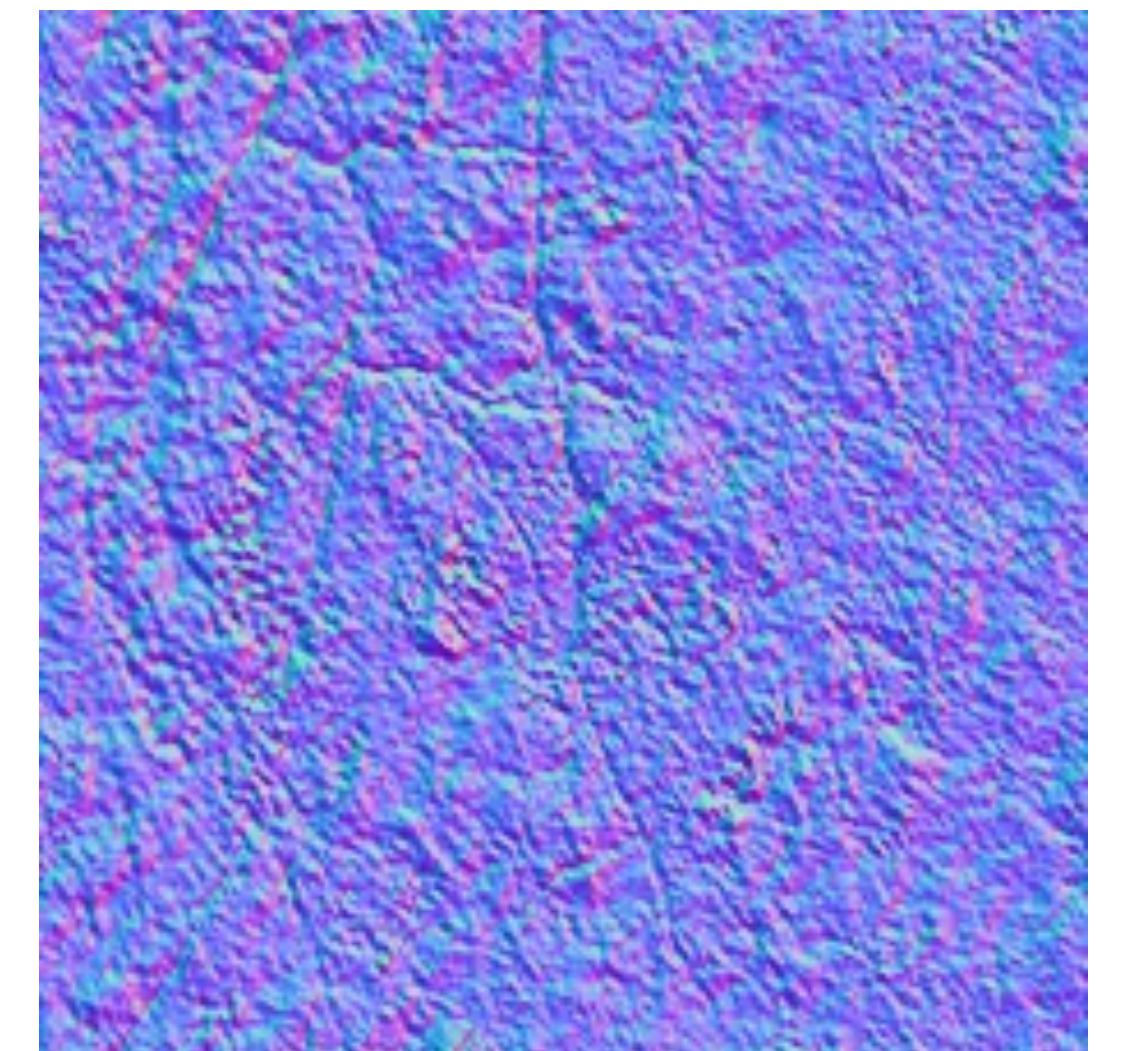
Texture Maps

- Reflectance across a surface varies with position
- Map each vertex (x,y,z) to a point (u,v) in an image
- Interpolate across the polygon to interpolate the corresponding (u,v) positions in the image



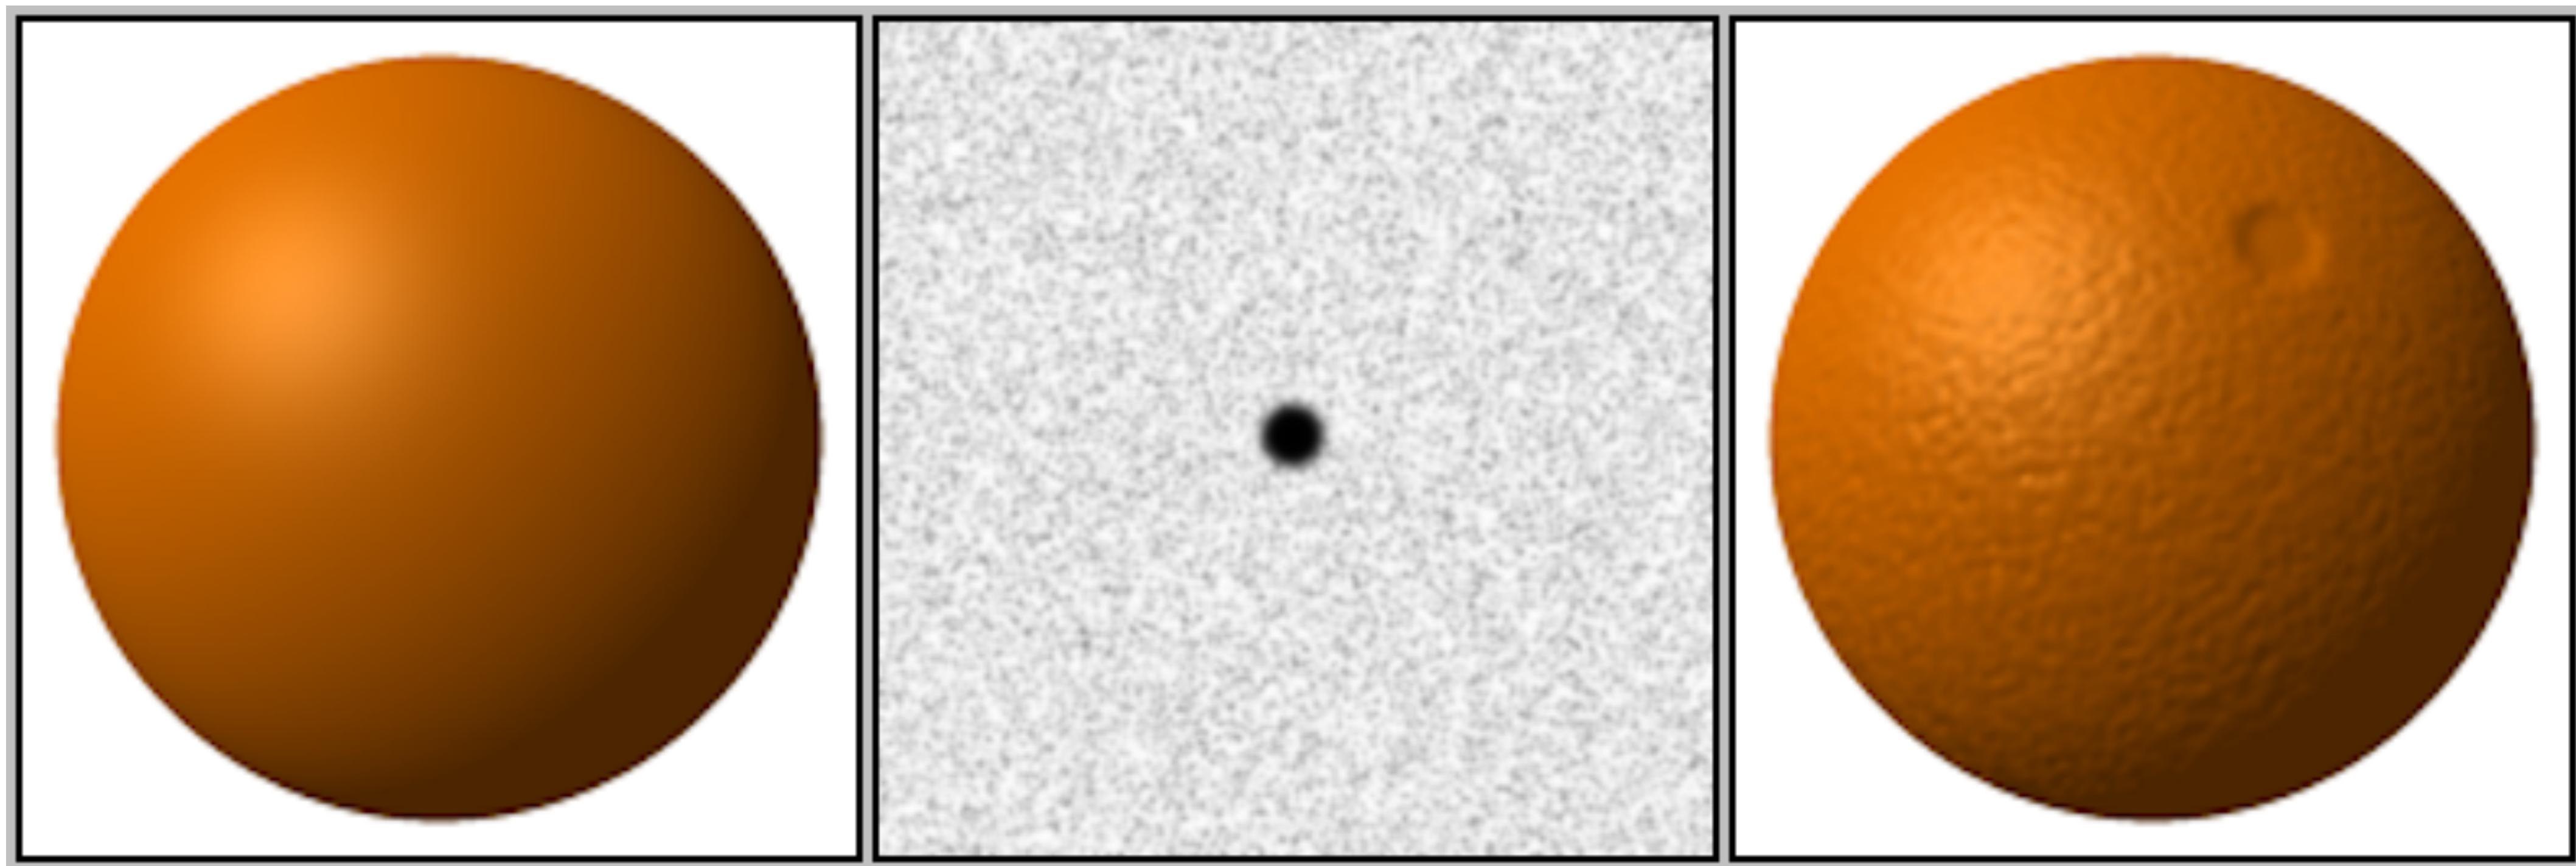
Bump Maps

- Texture maps are only reflectance
- They don't respond to changes in the lighting ("painted on")
- Idea: use a *normal map*

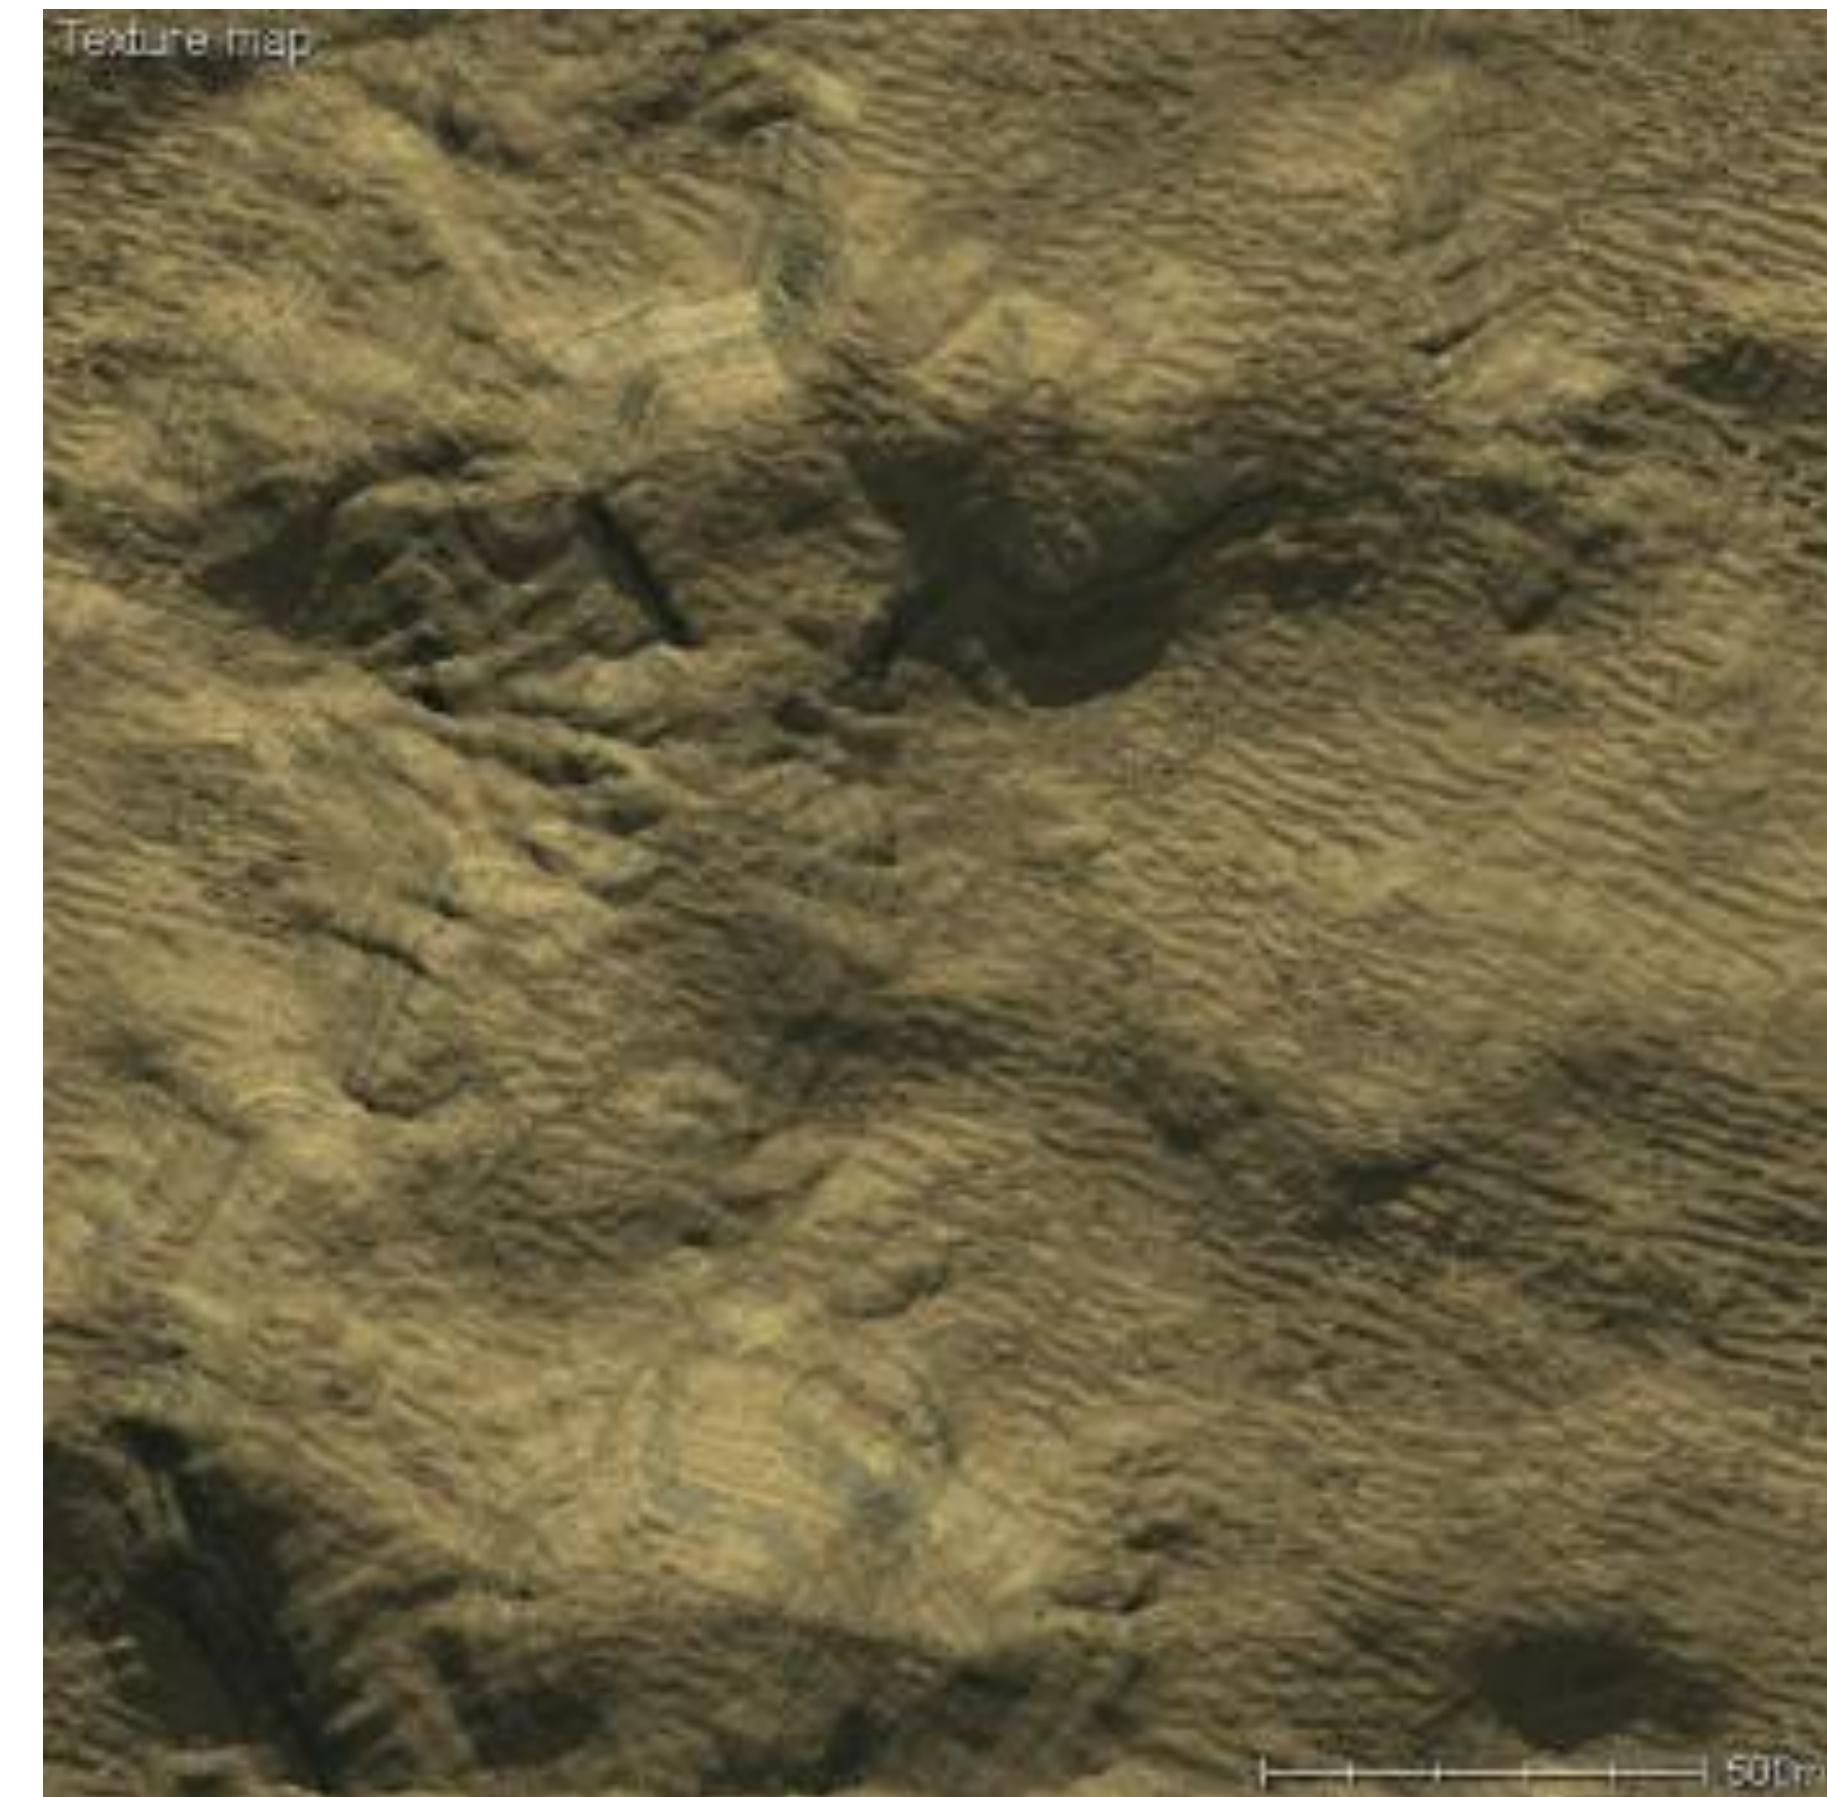
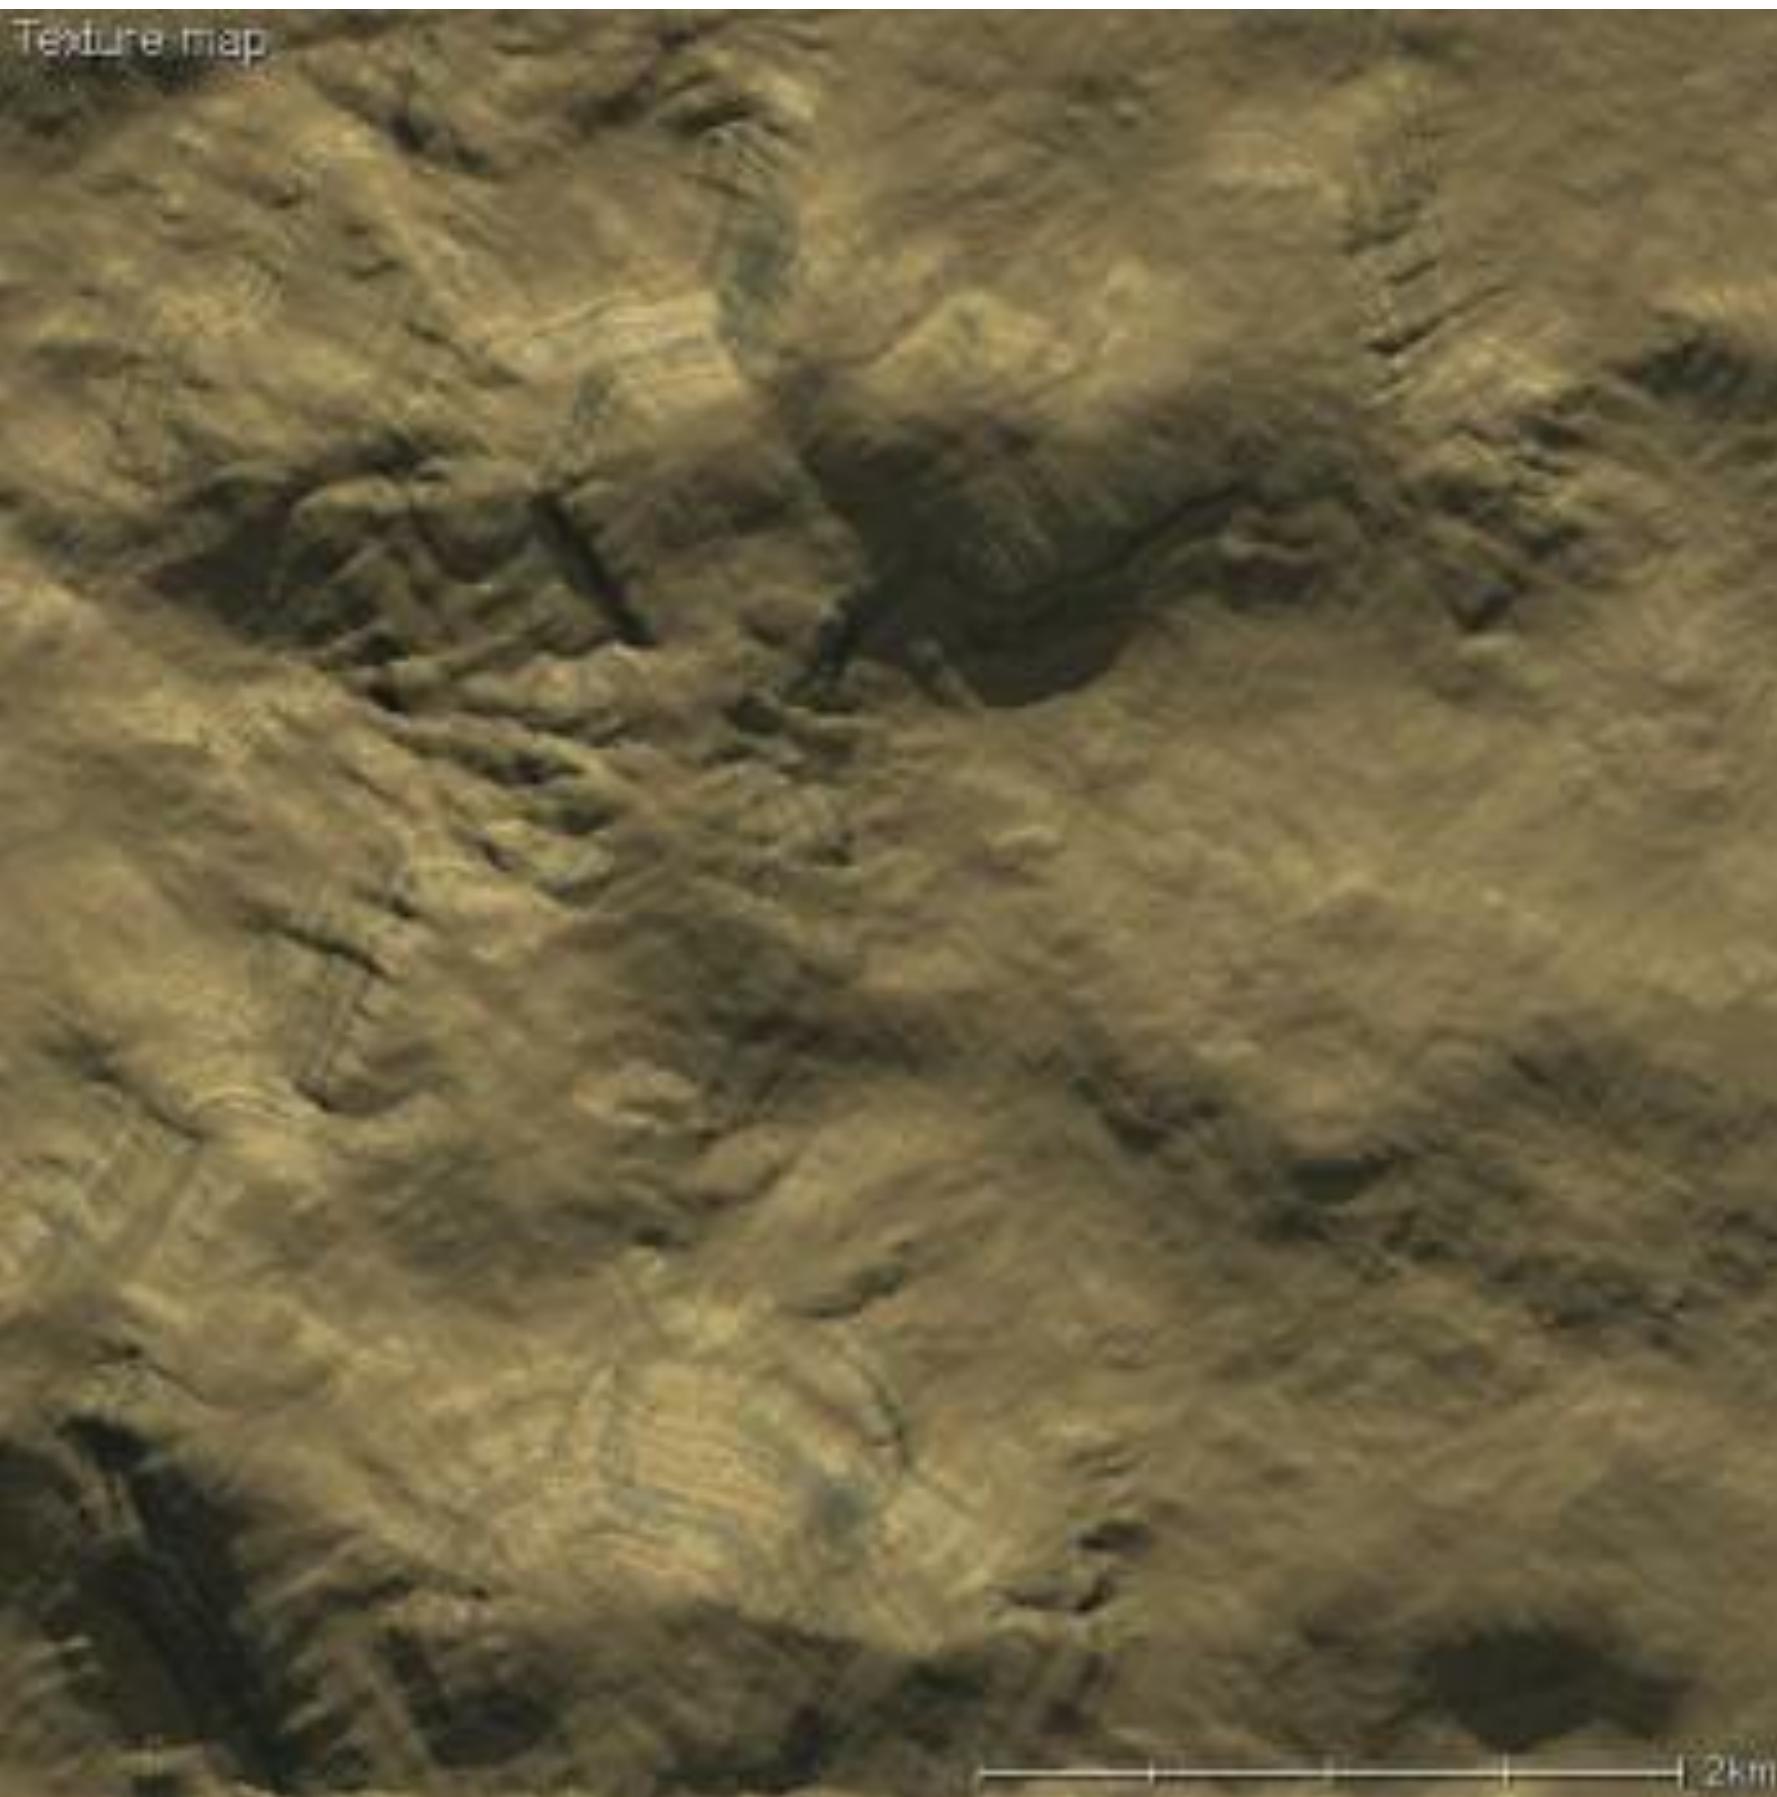


Jim Blinn, 1978

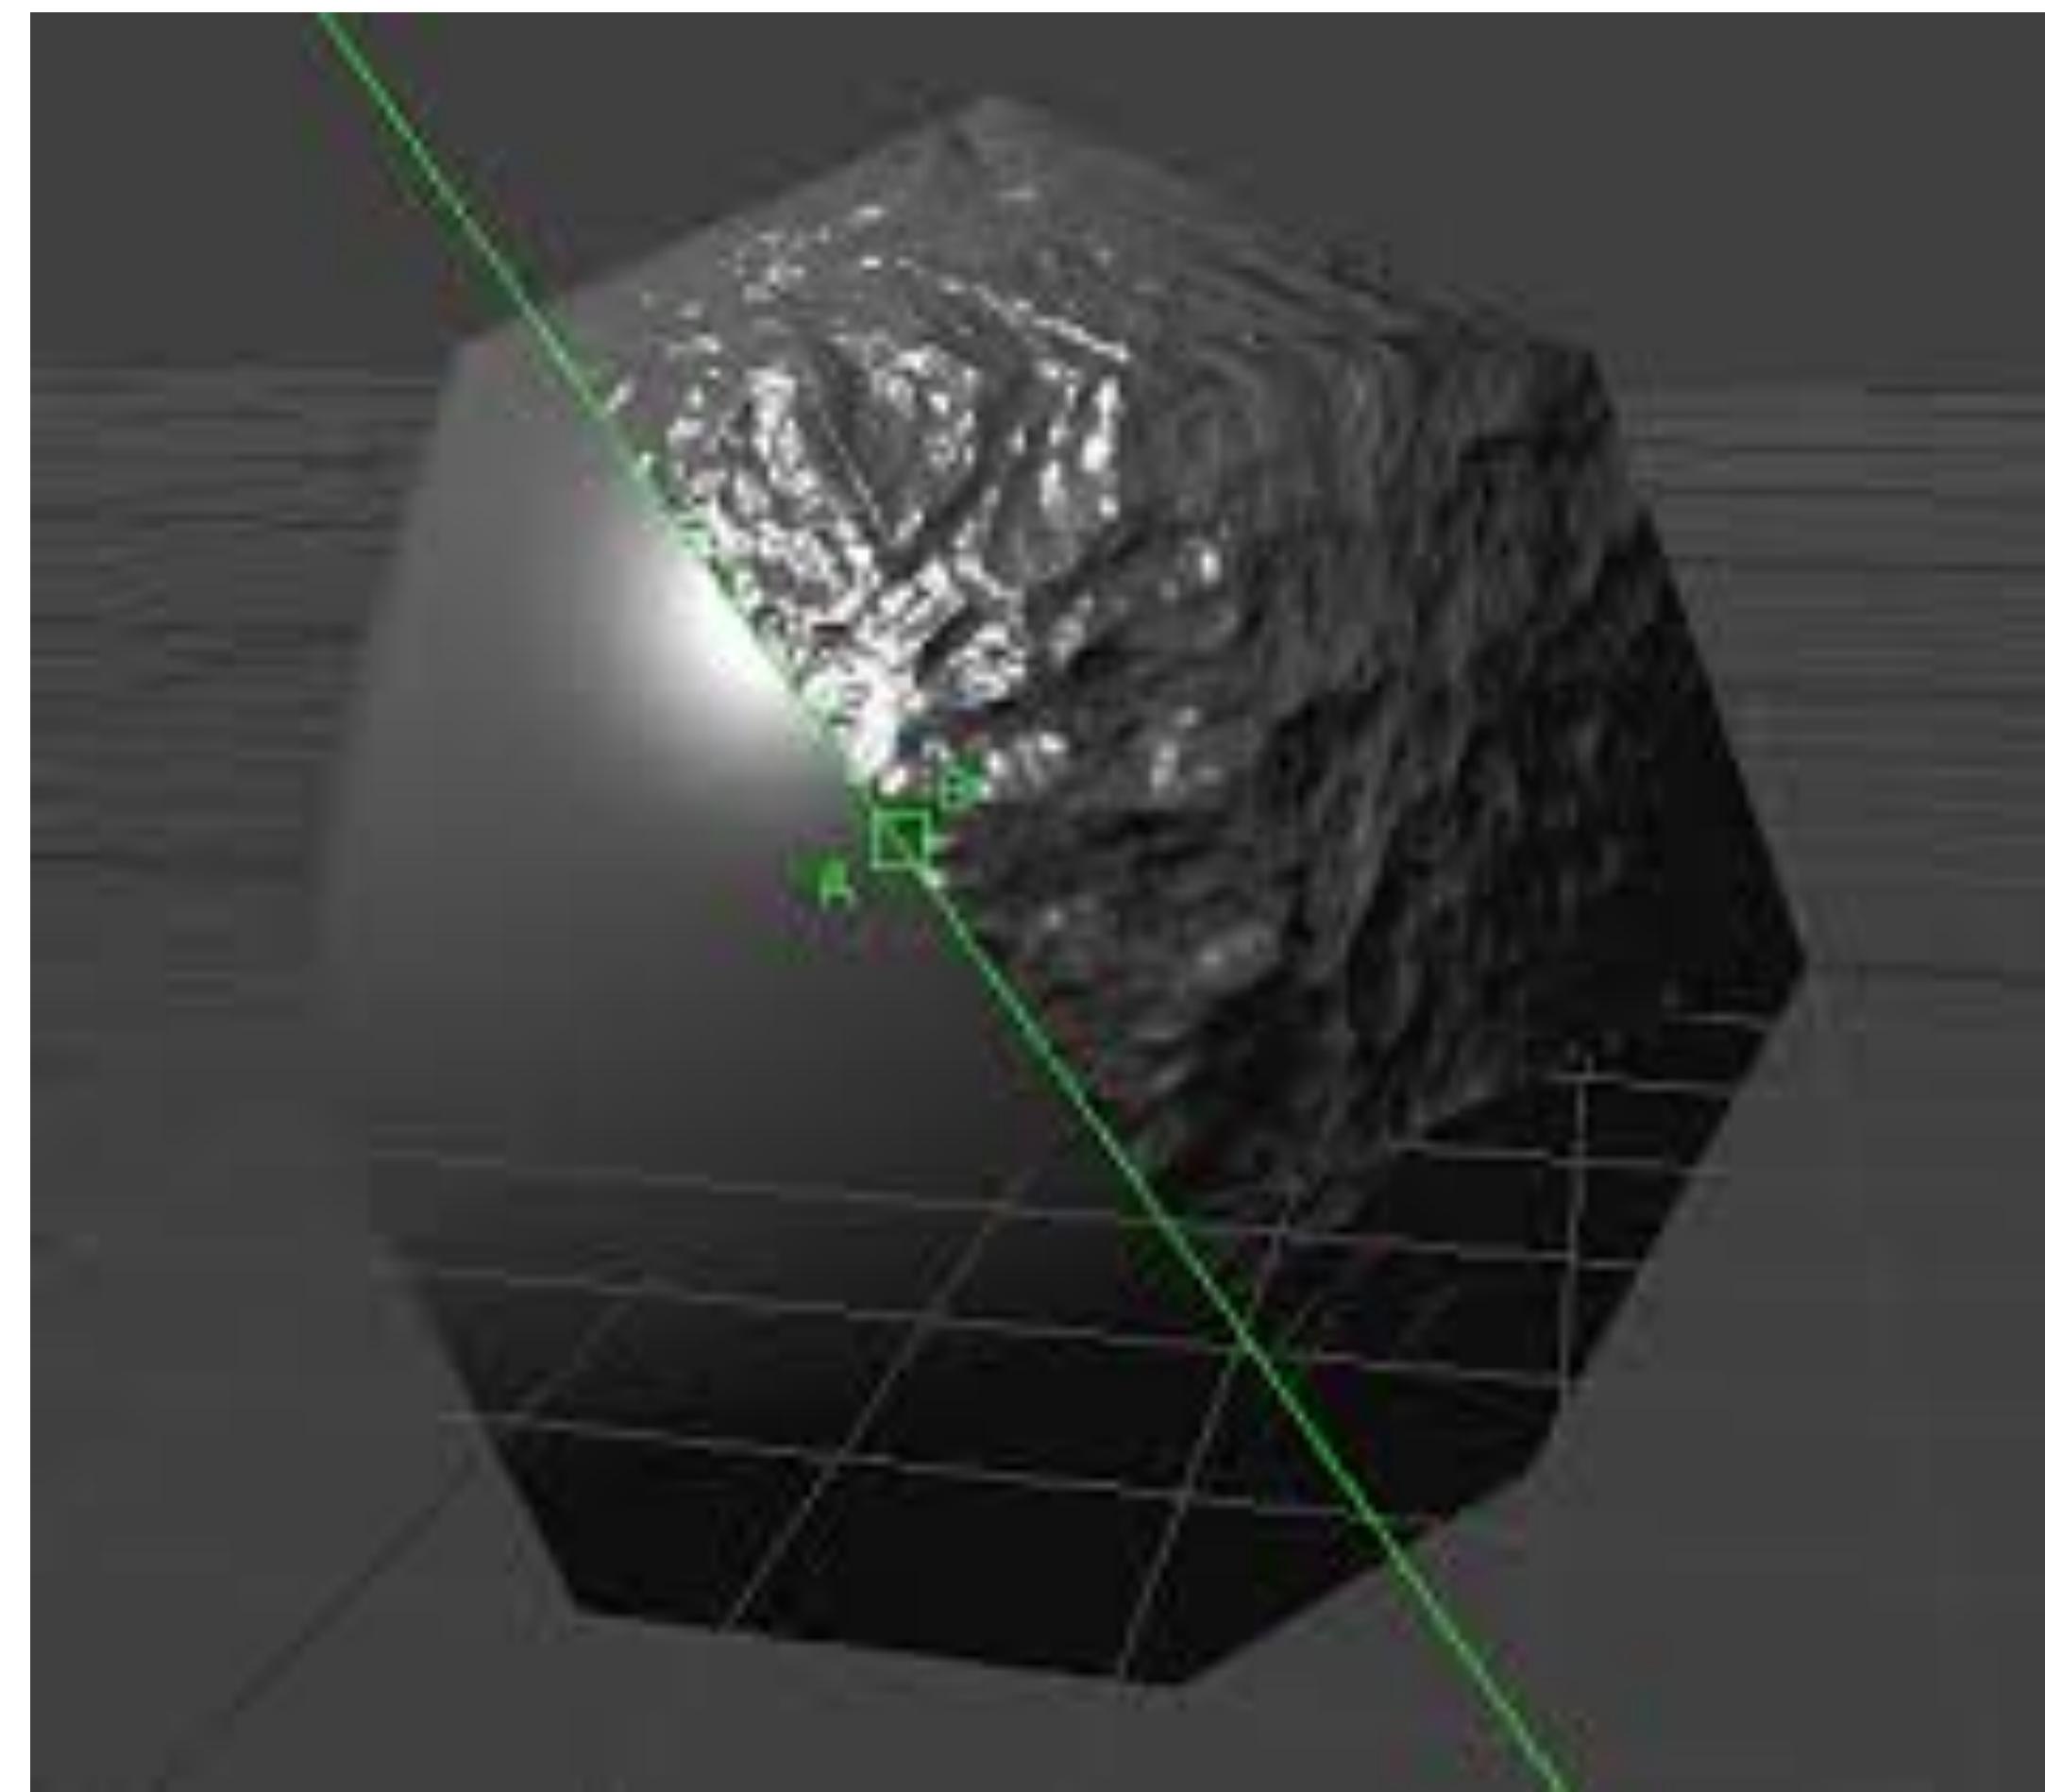
Bump Maps



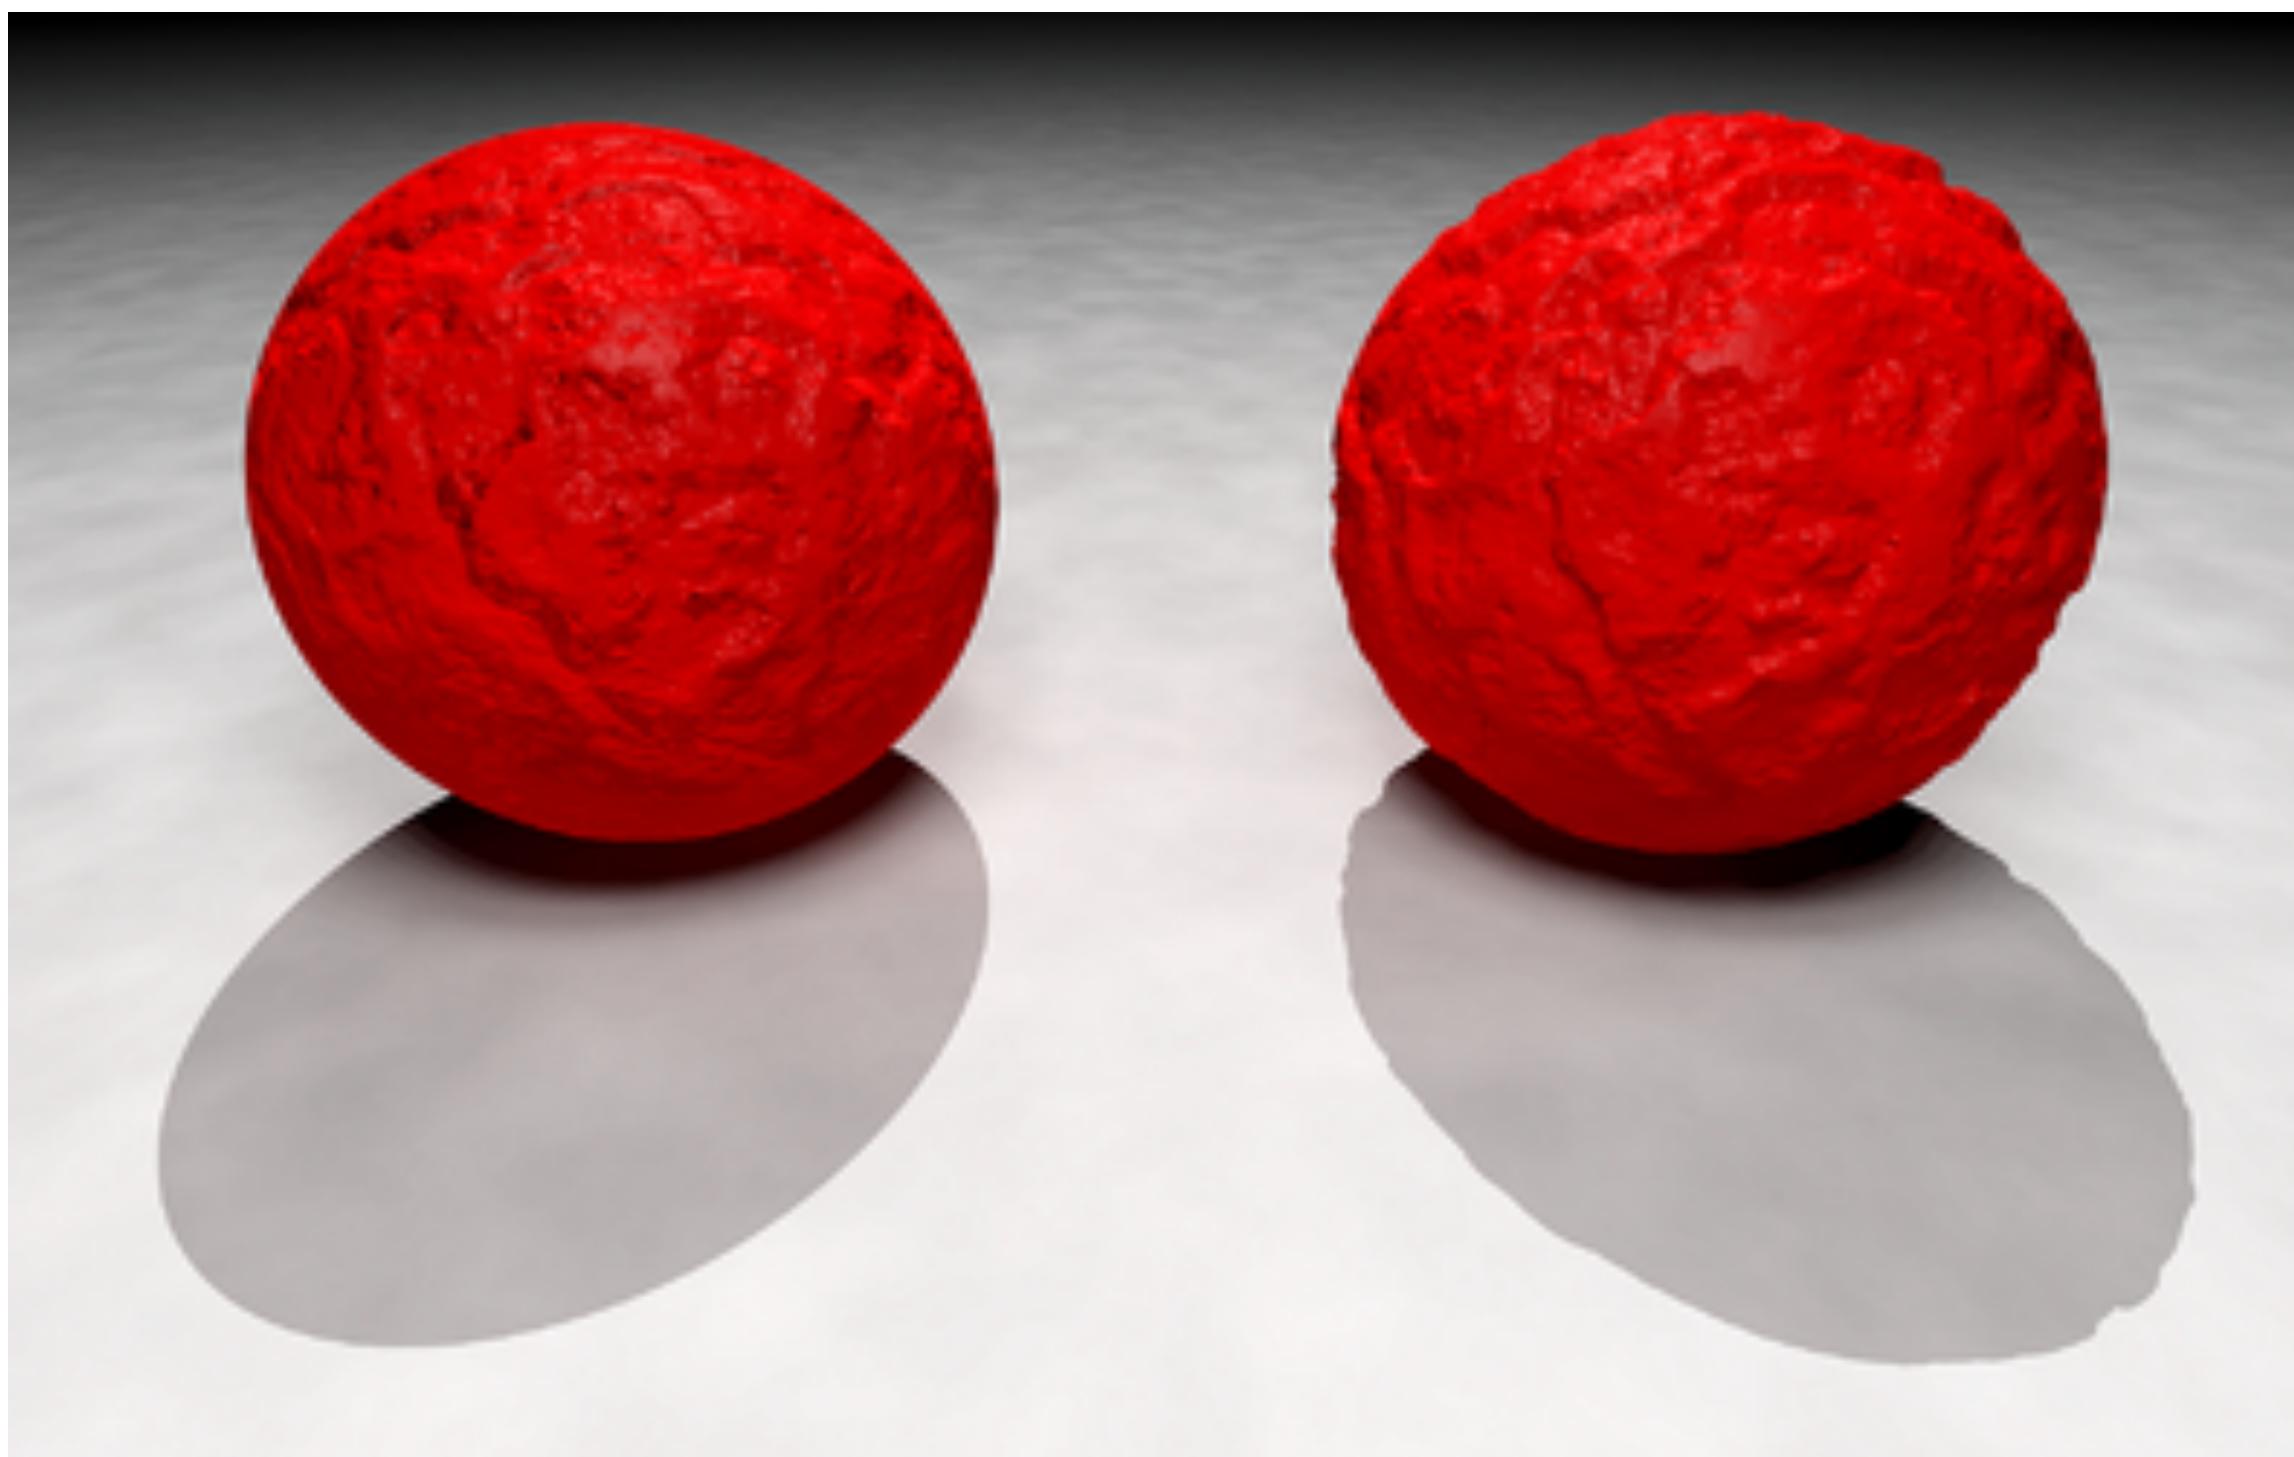
Bump Mapping



Bump Mapping



Bump Mapping



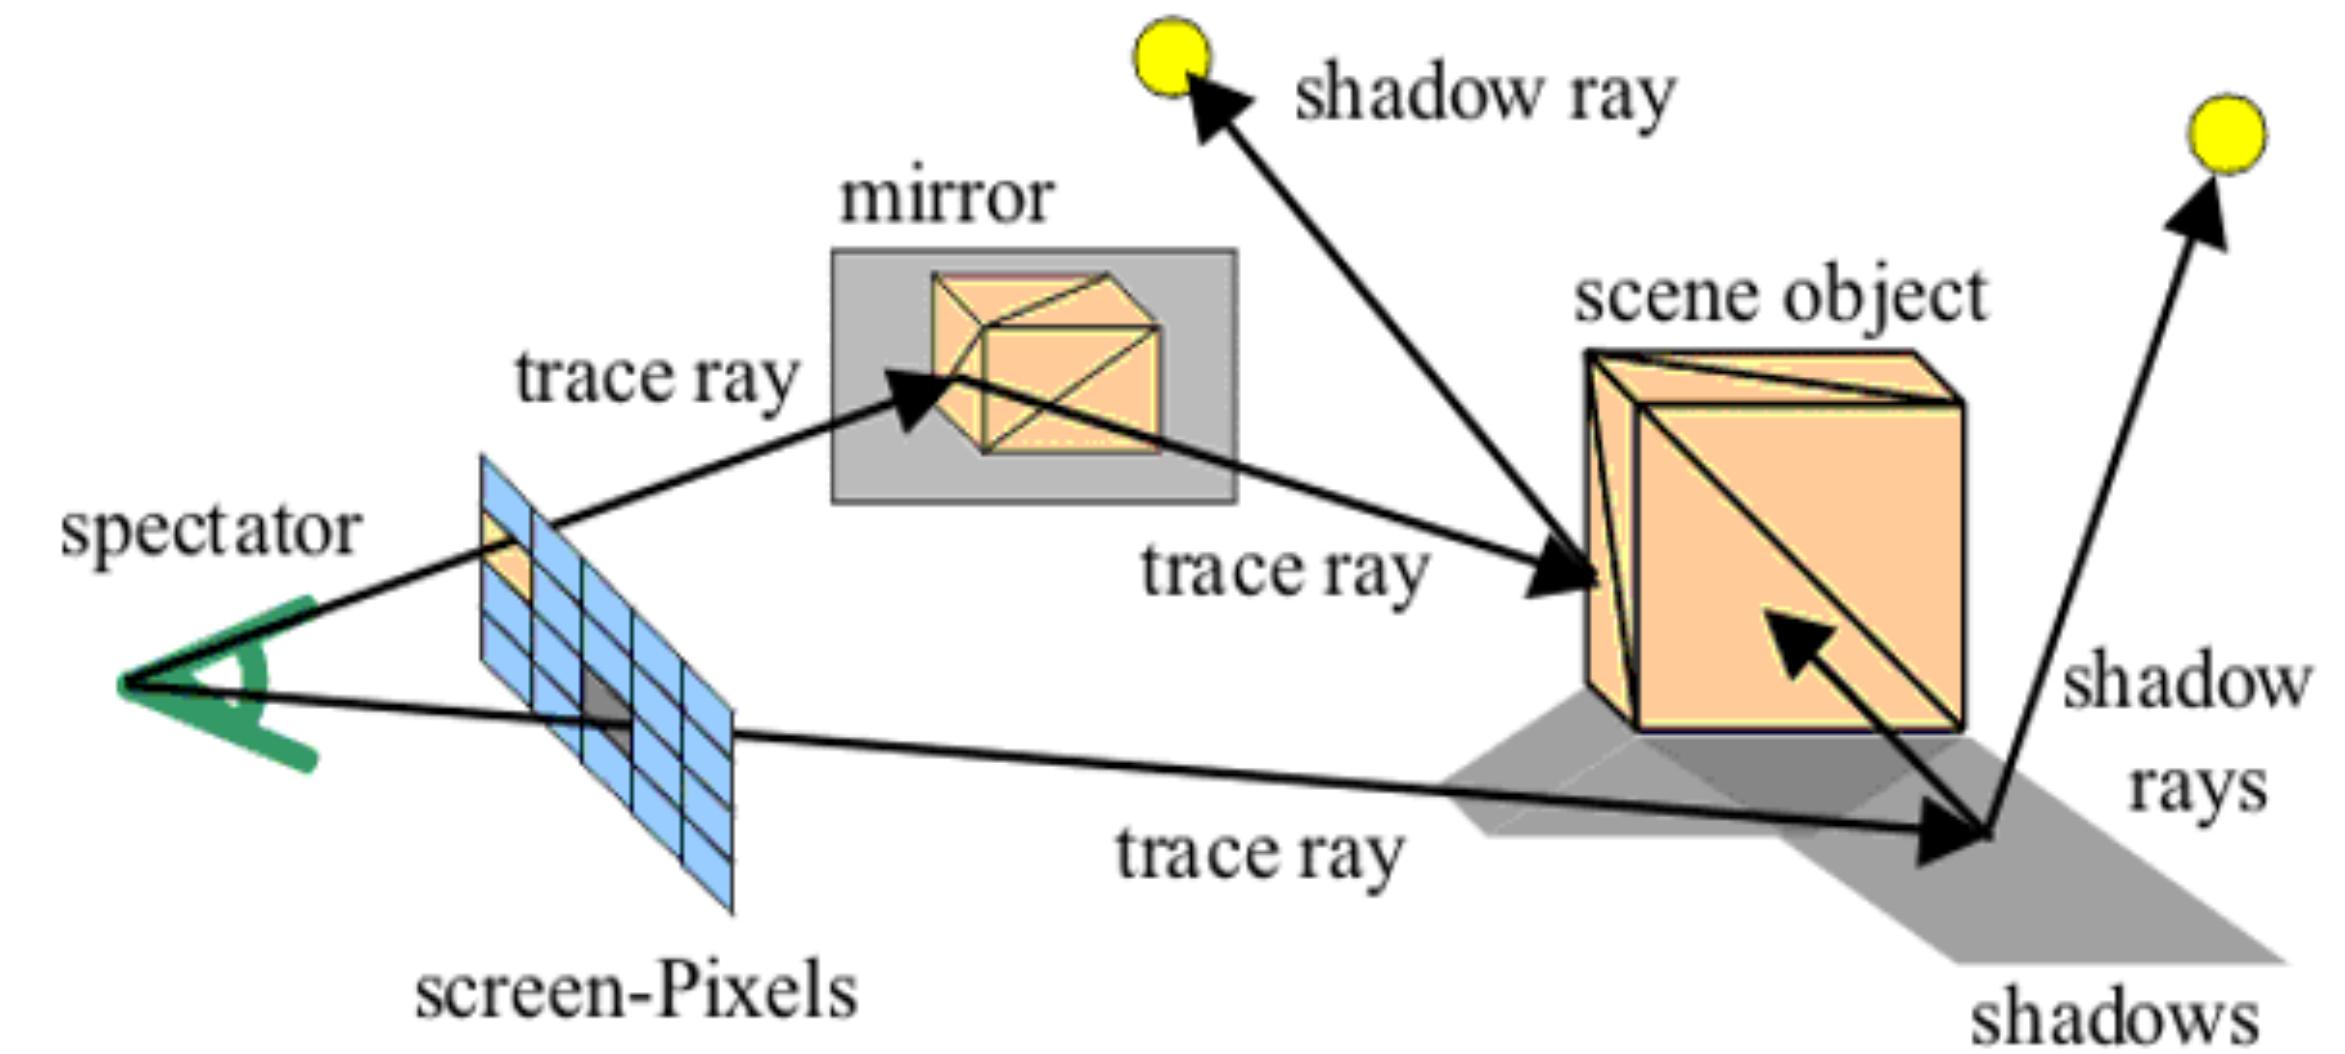
Ray Tracing

- What about things like
 - Reflections?
 - Refraction through transparent materials?



Ray Tracing

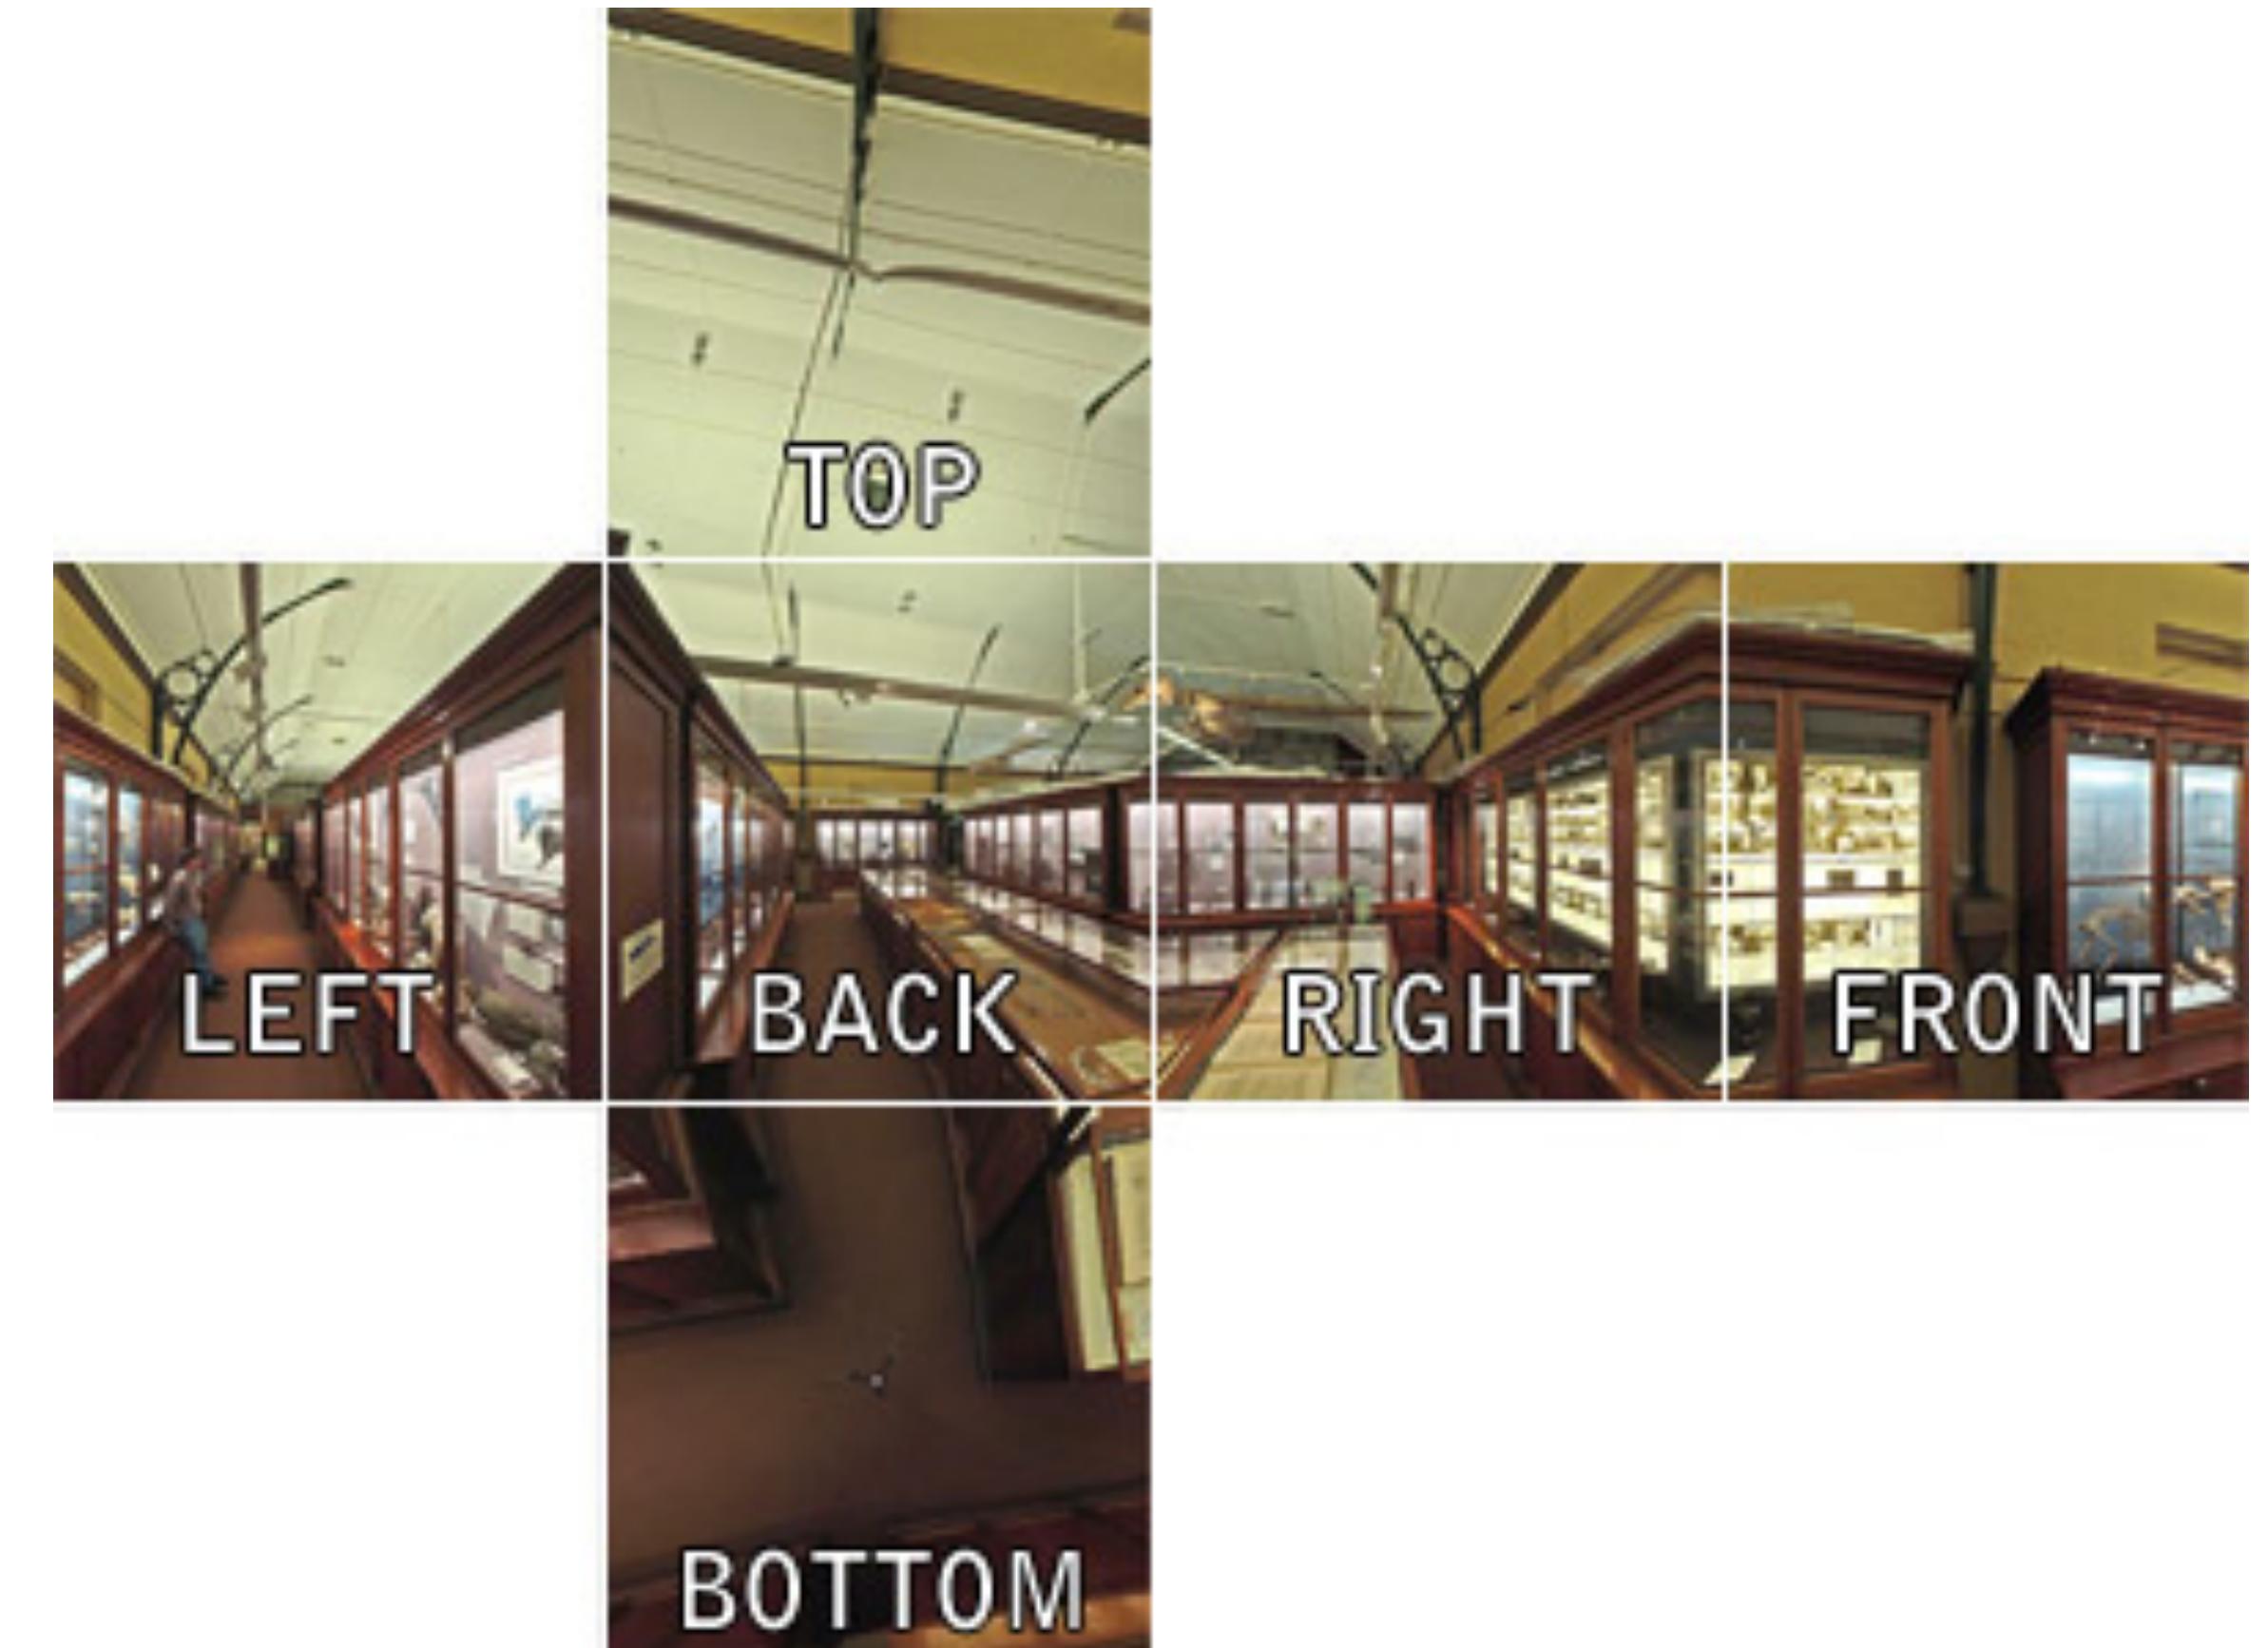
- Shoot a ray out through the pixel on the imaging plane
- When it hits something, ***recursively*** shoot more rays:
 - Towards light source(s) — direct lighting
 - Reflection ray (if applicable)
 - Refraction ray (if applicable)



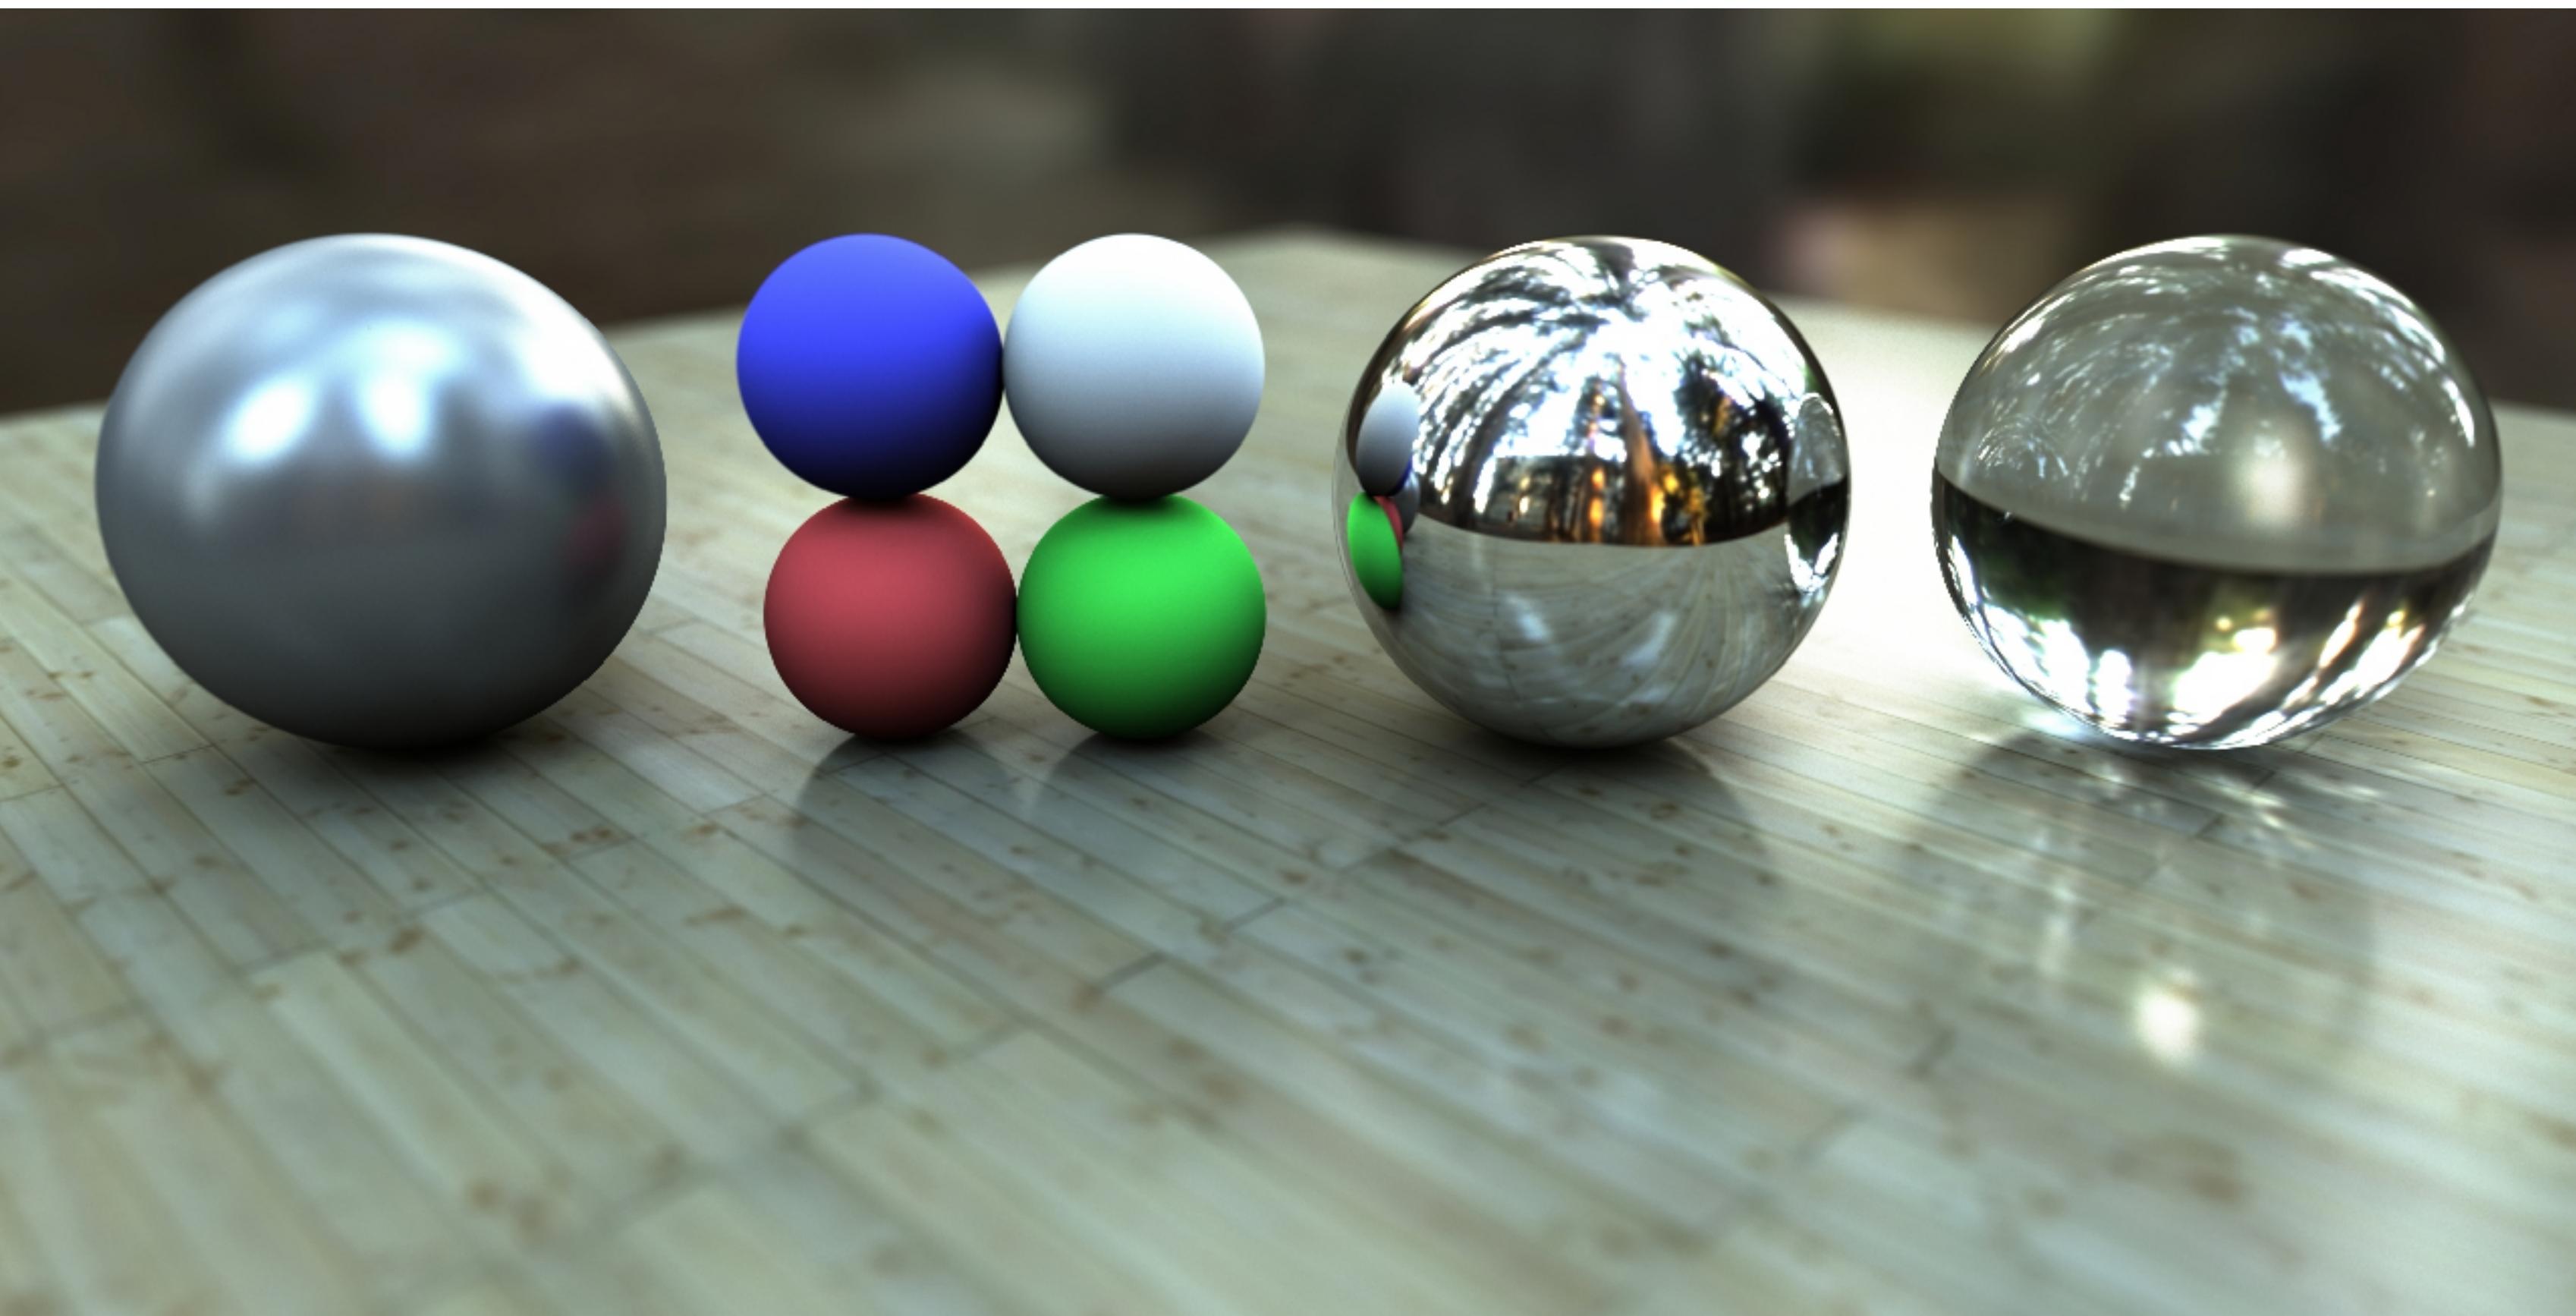
Turner Whitted, 1980

Environment Maps

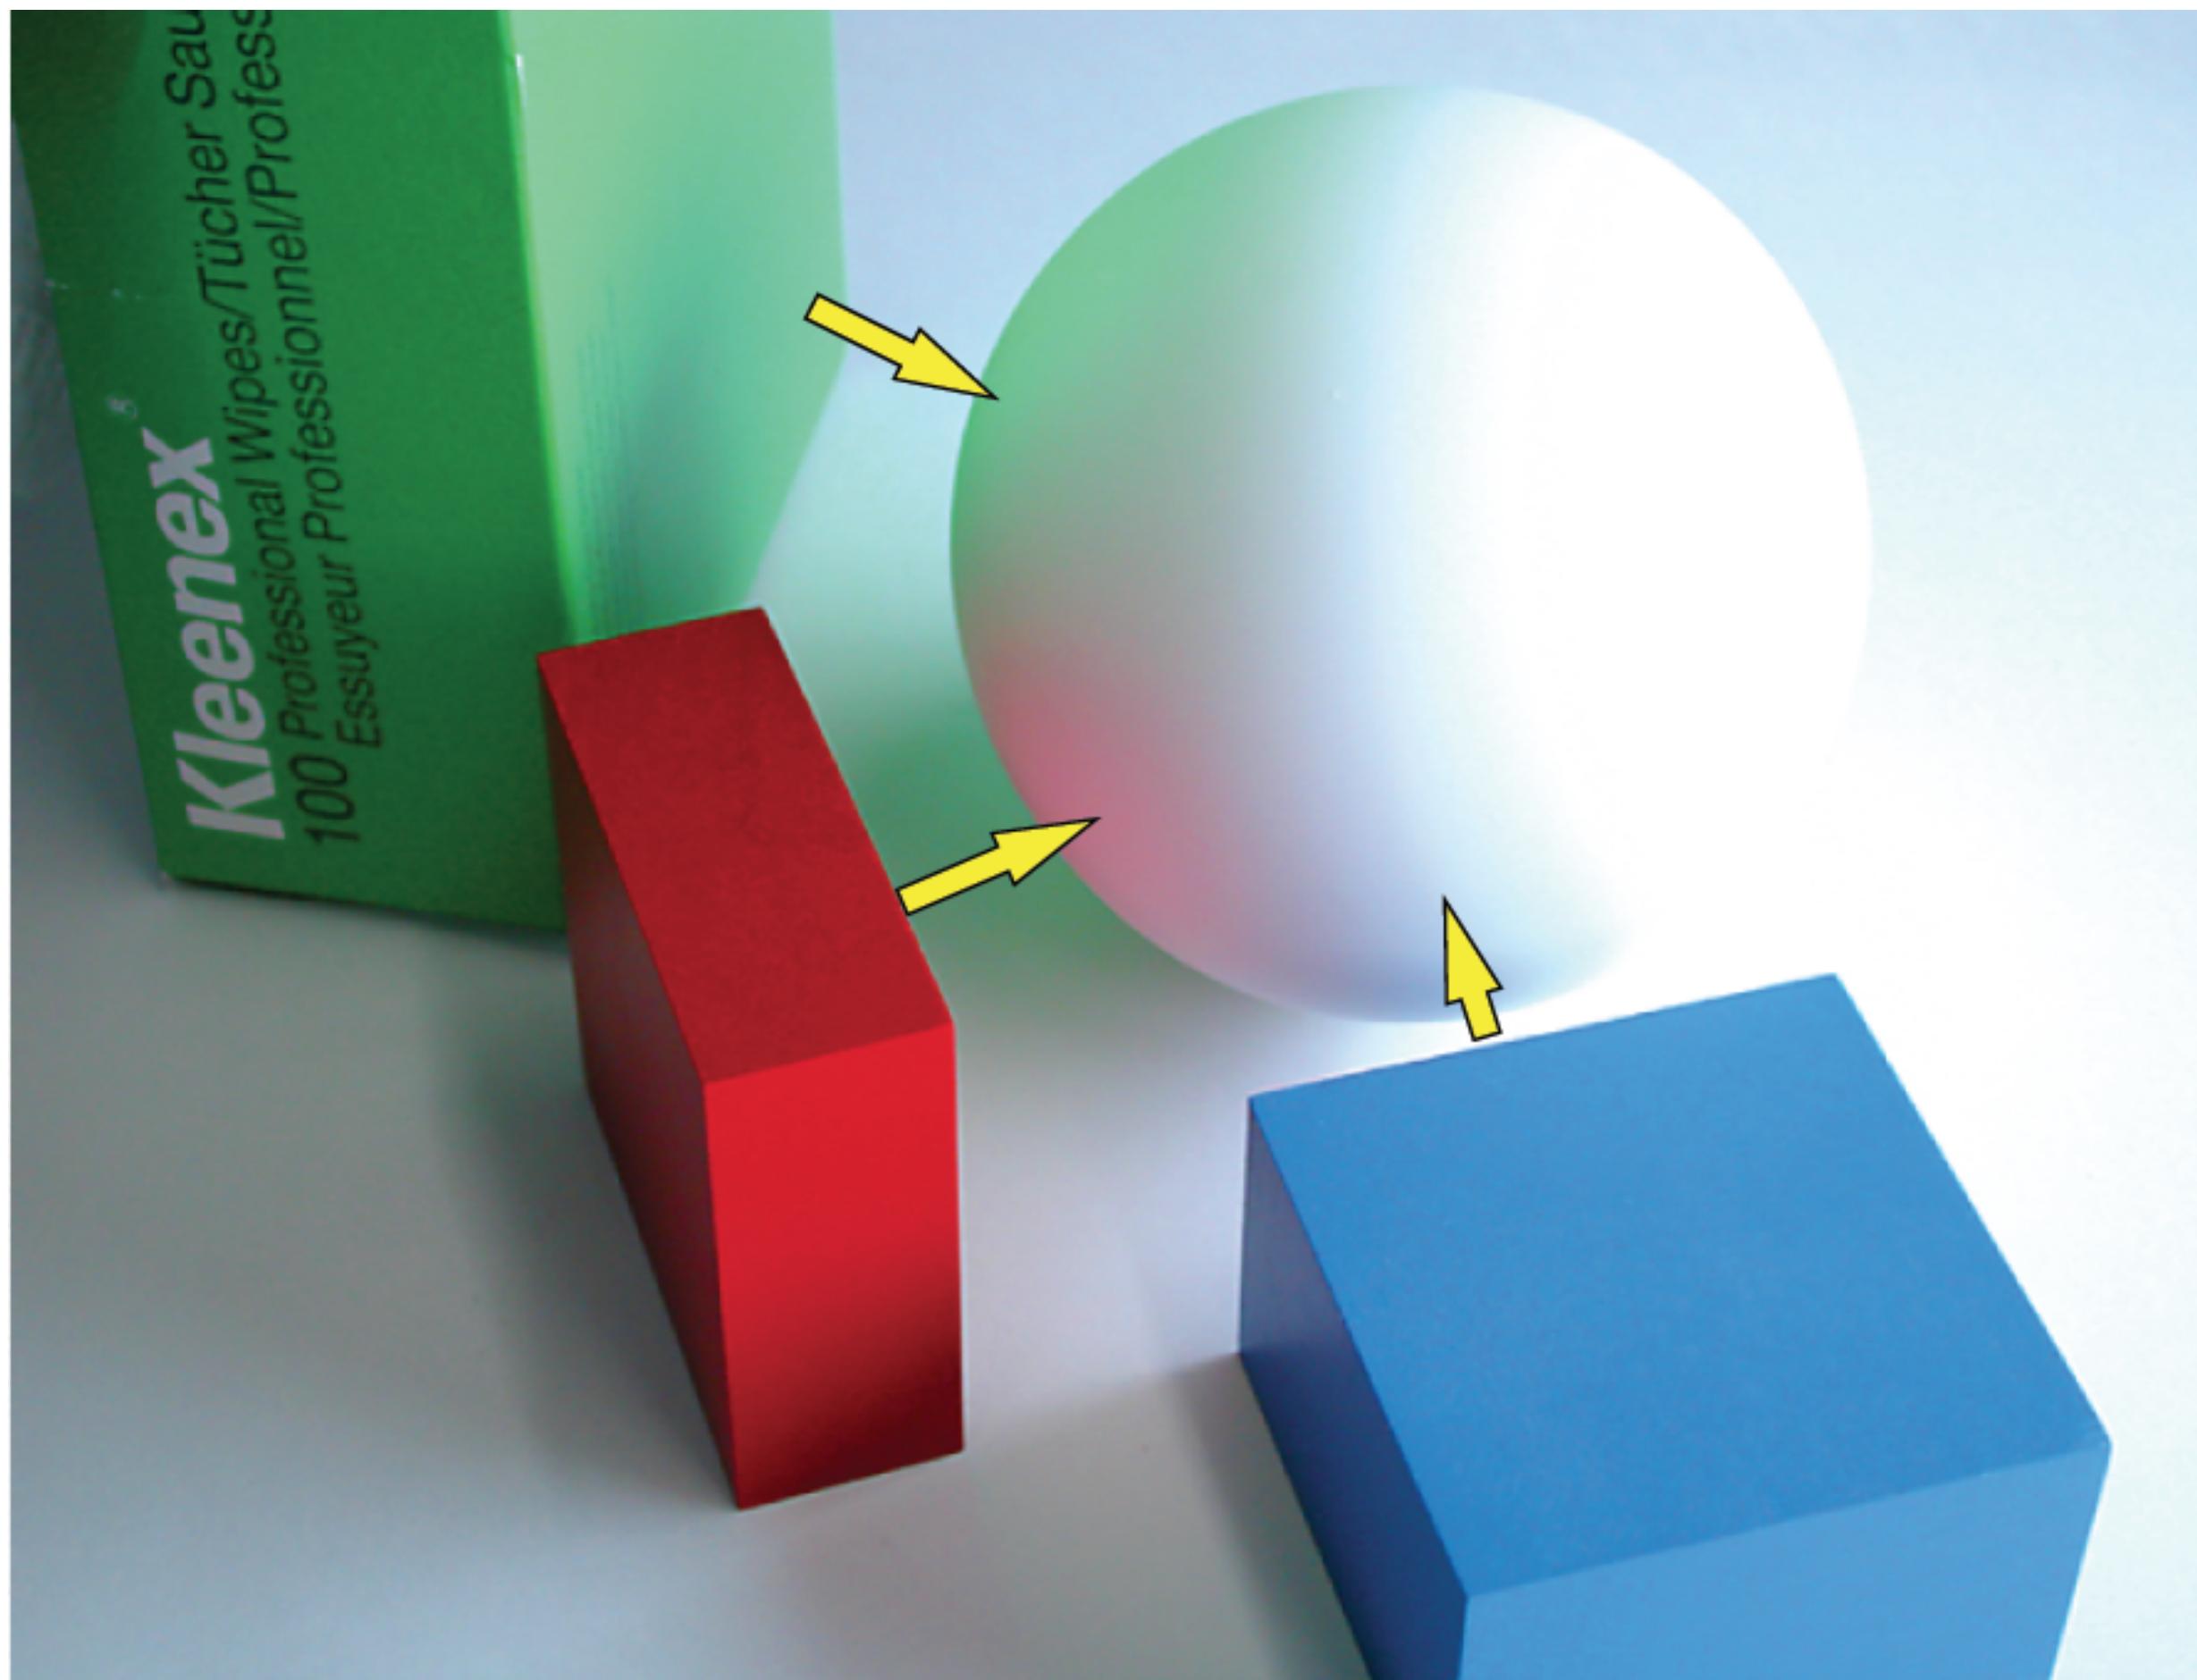
- What about rays that exit the scene?
- Can wrap the entire scene in an *environment map*
- Square, spherical, etc.



Environment Maps

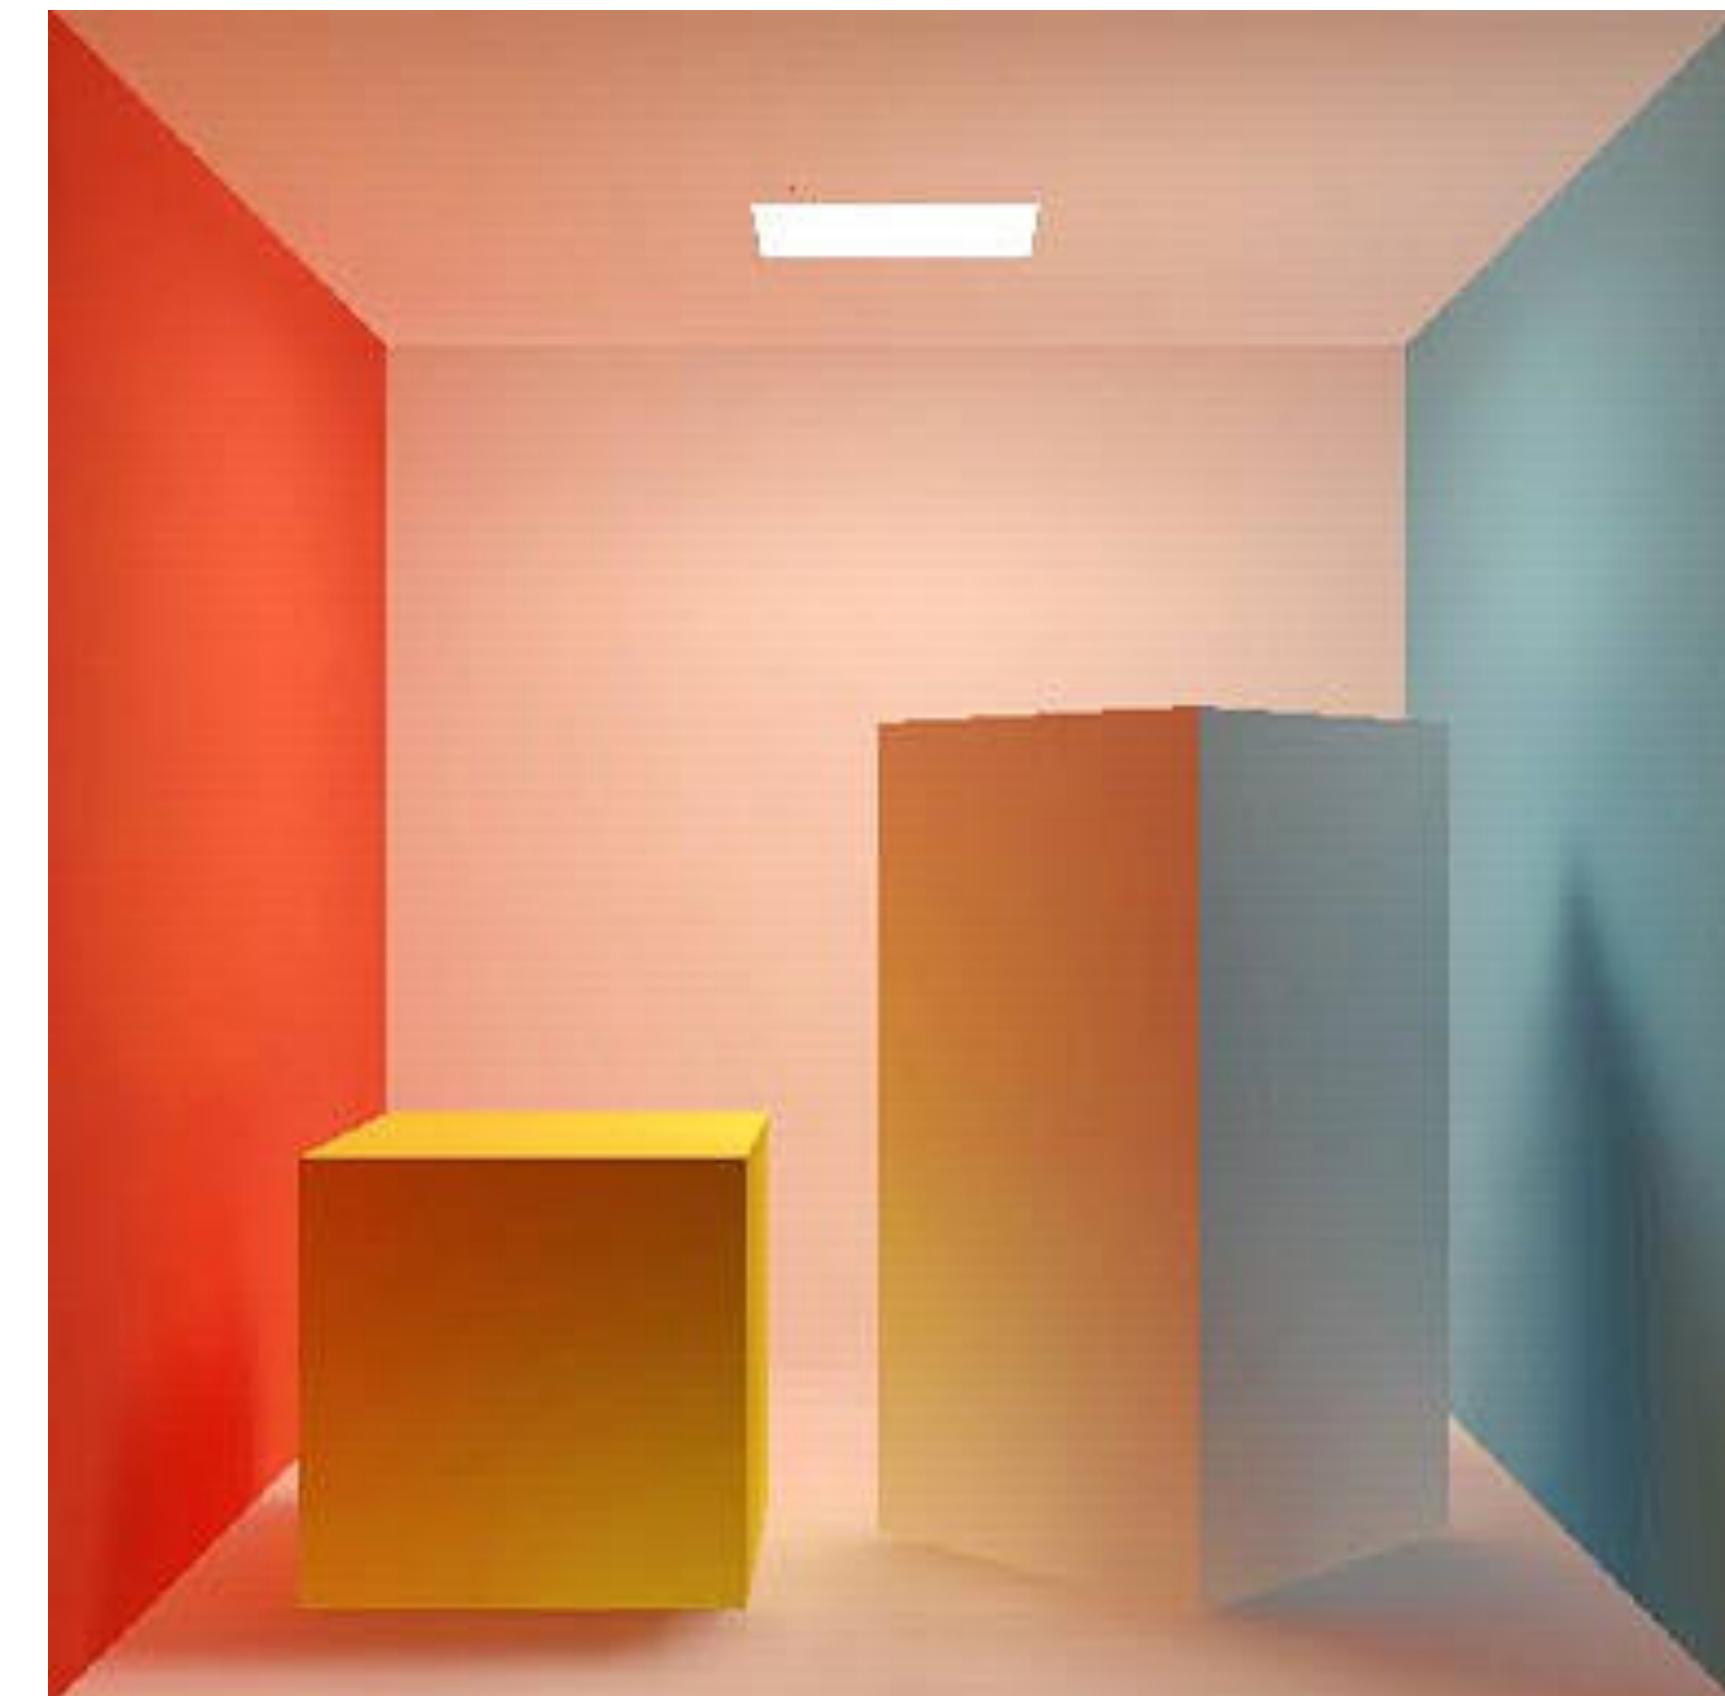


Diffuse Interreflections



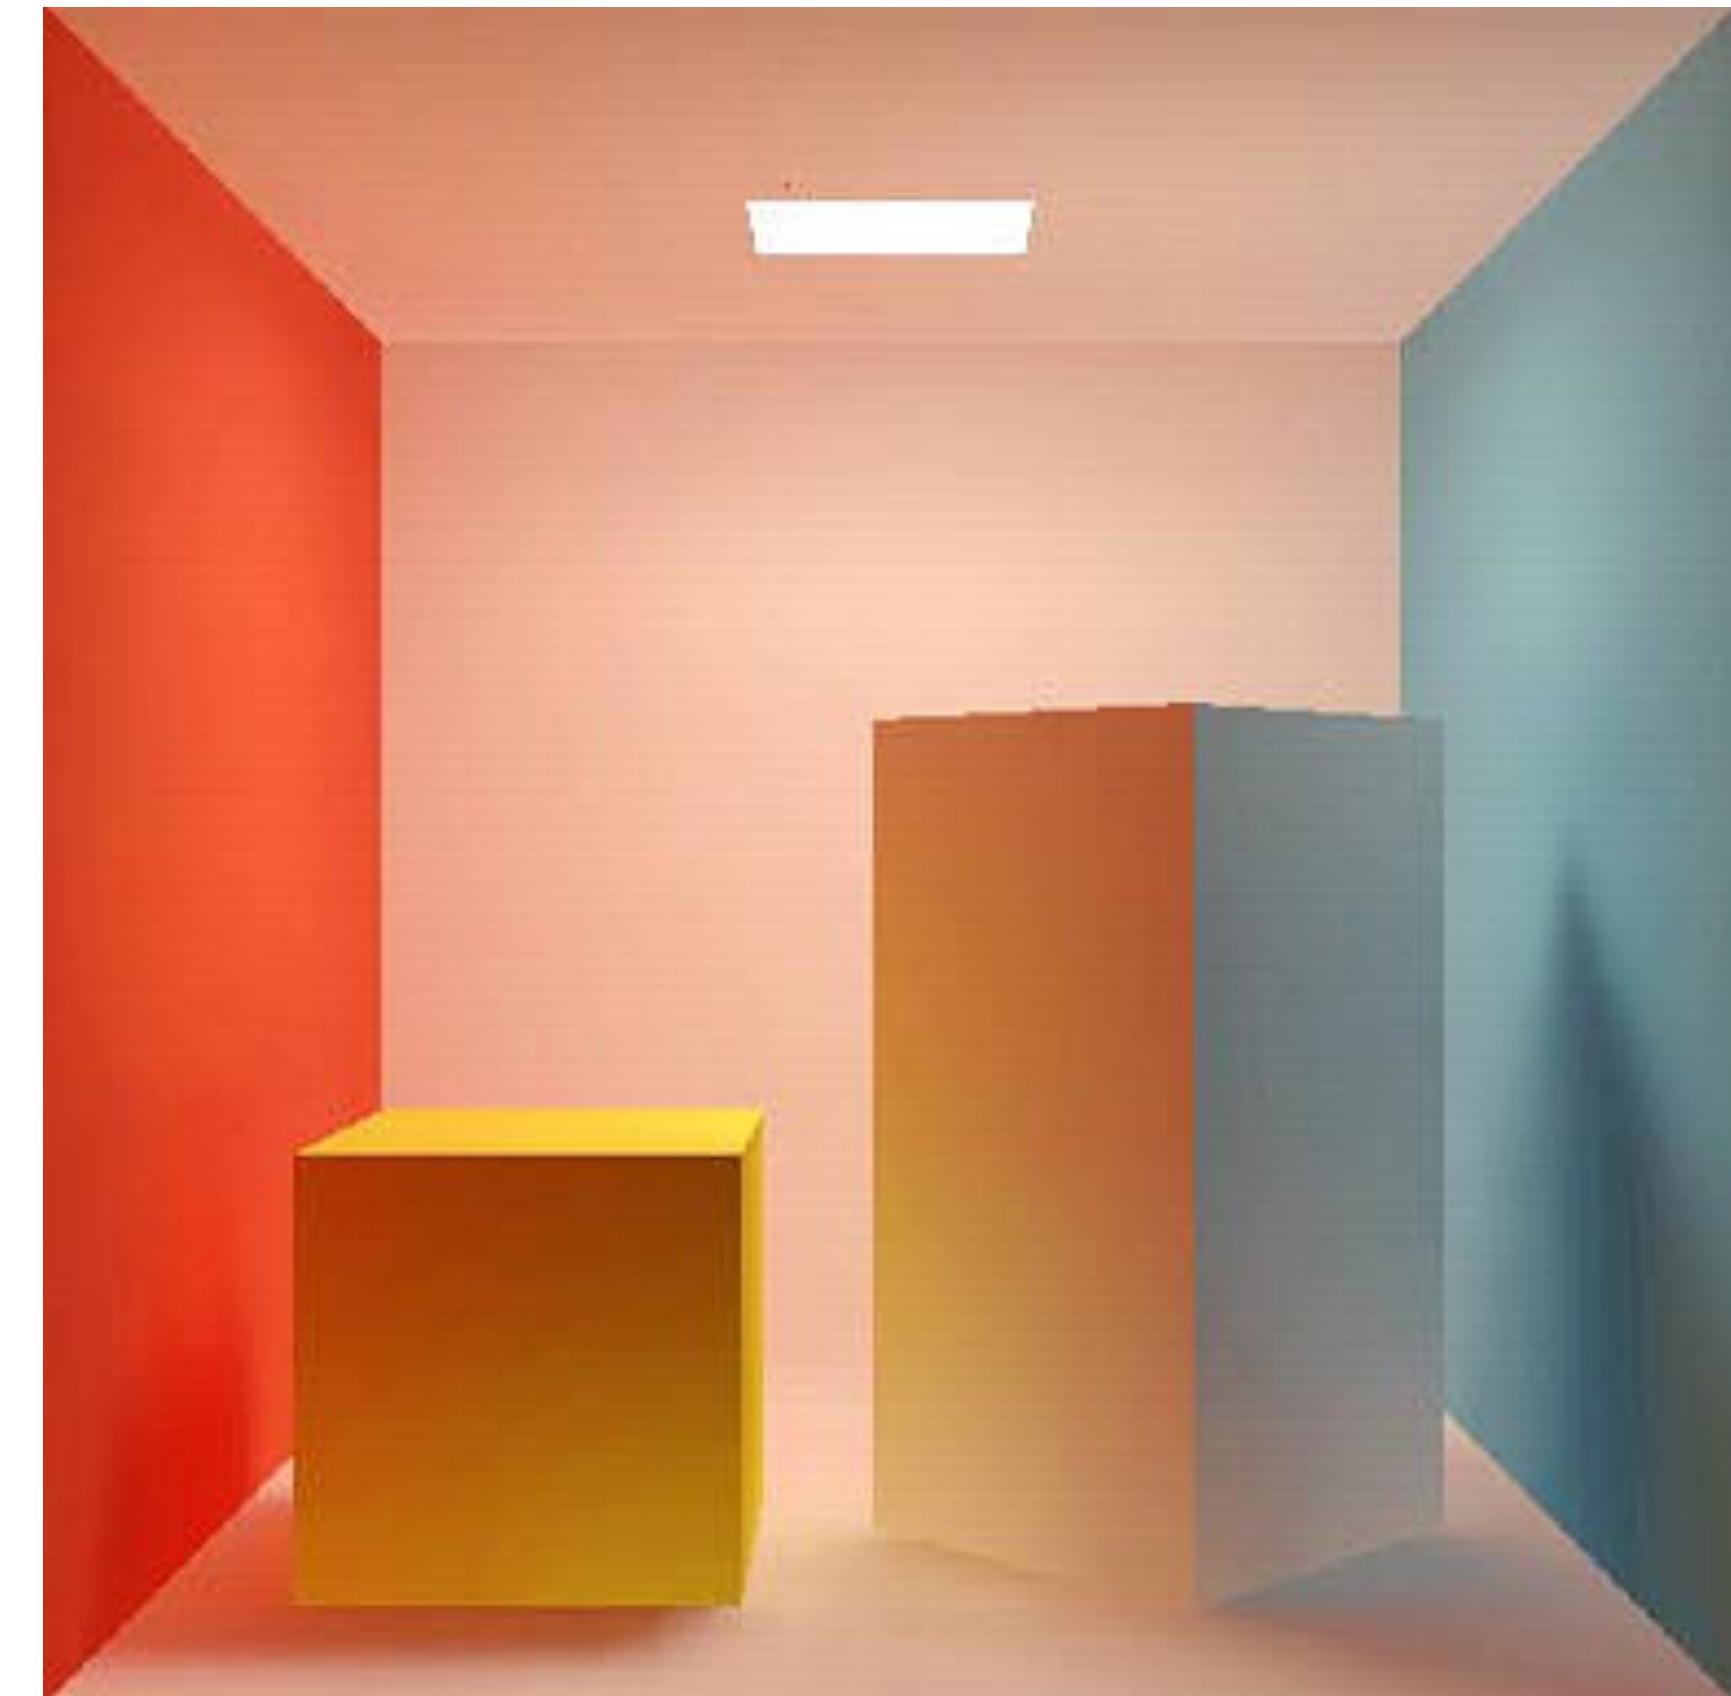
Radiosity

- Math involves integrals over *all angles* of light coming *into* and *off* of a surface
- Approximated using *importance sampling*
- Effectively traces lots and lots of rays
 - compute intensive



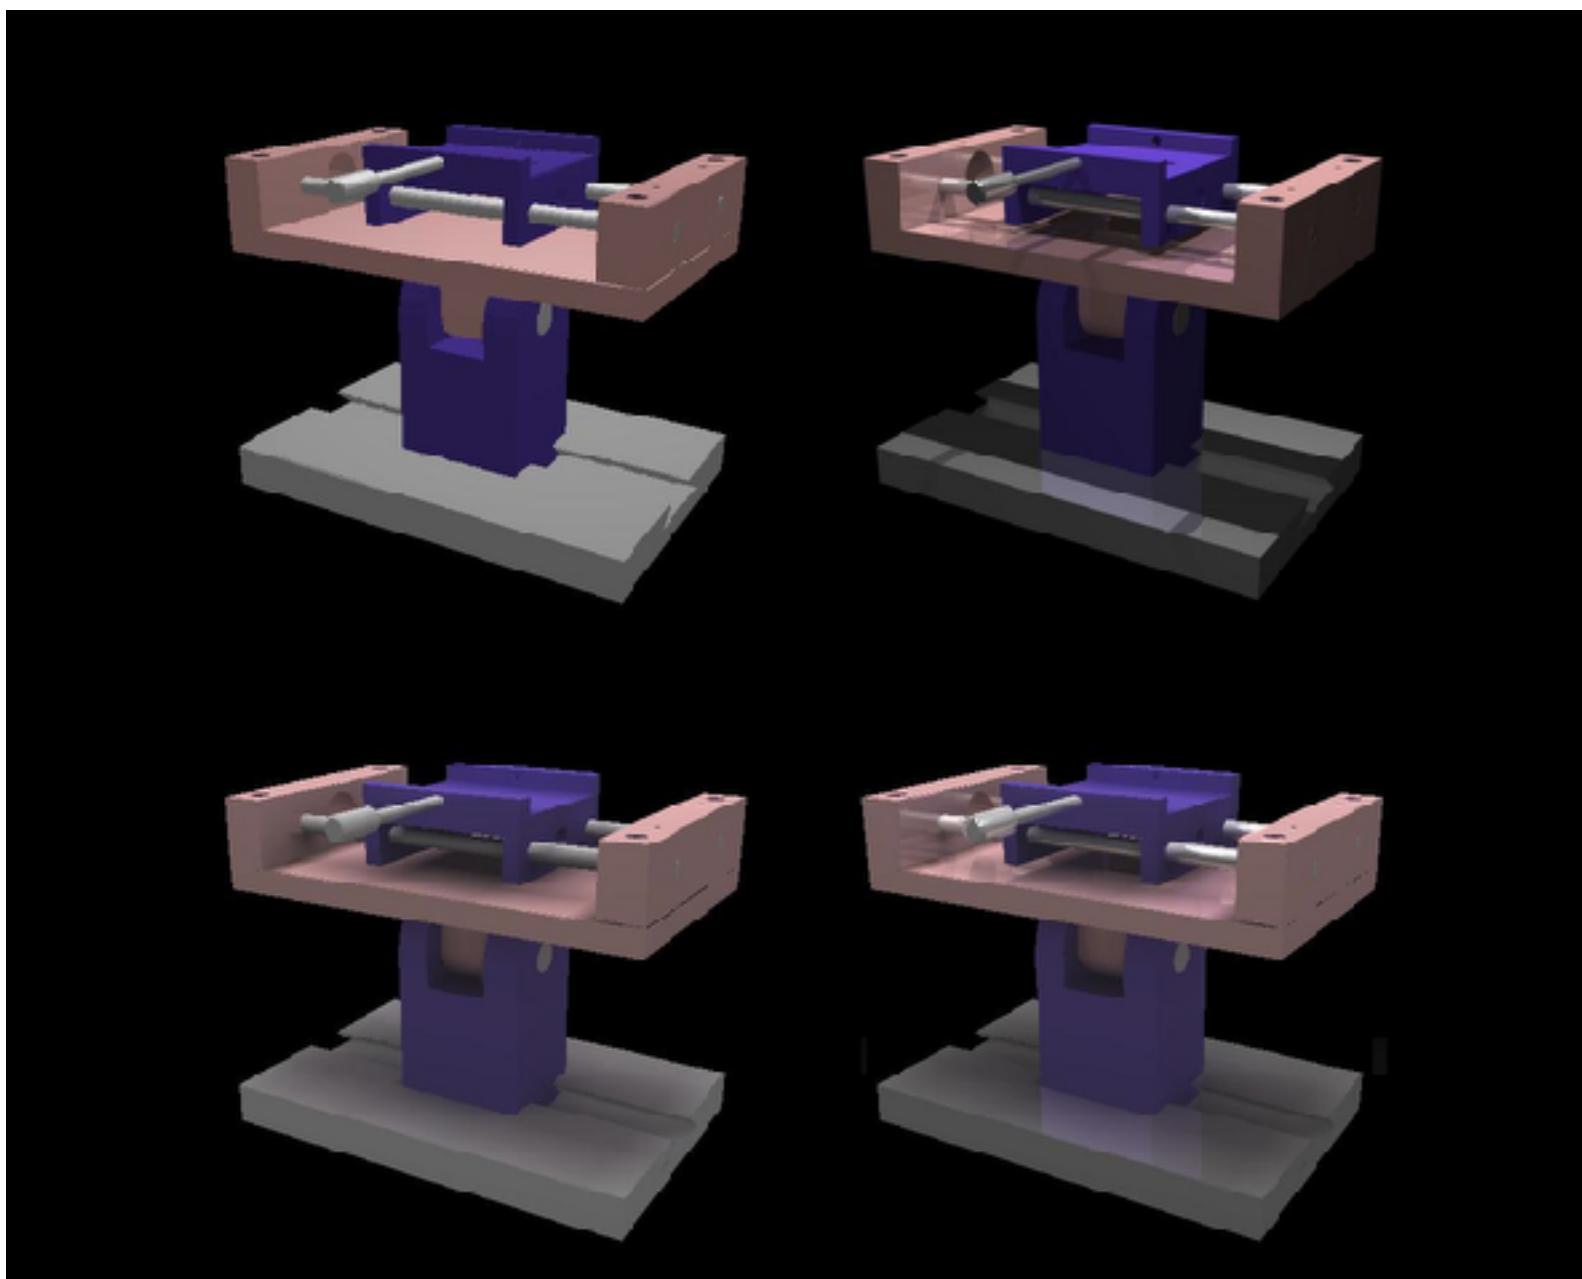
Radiosity

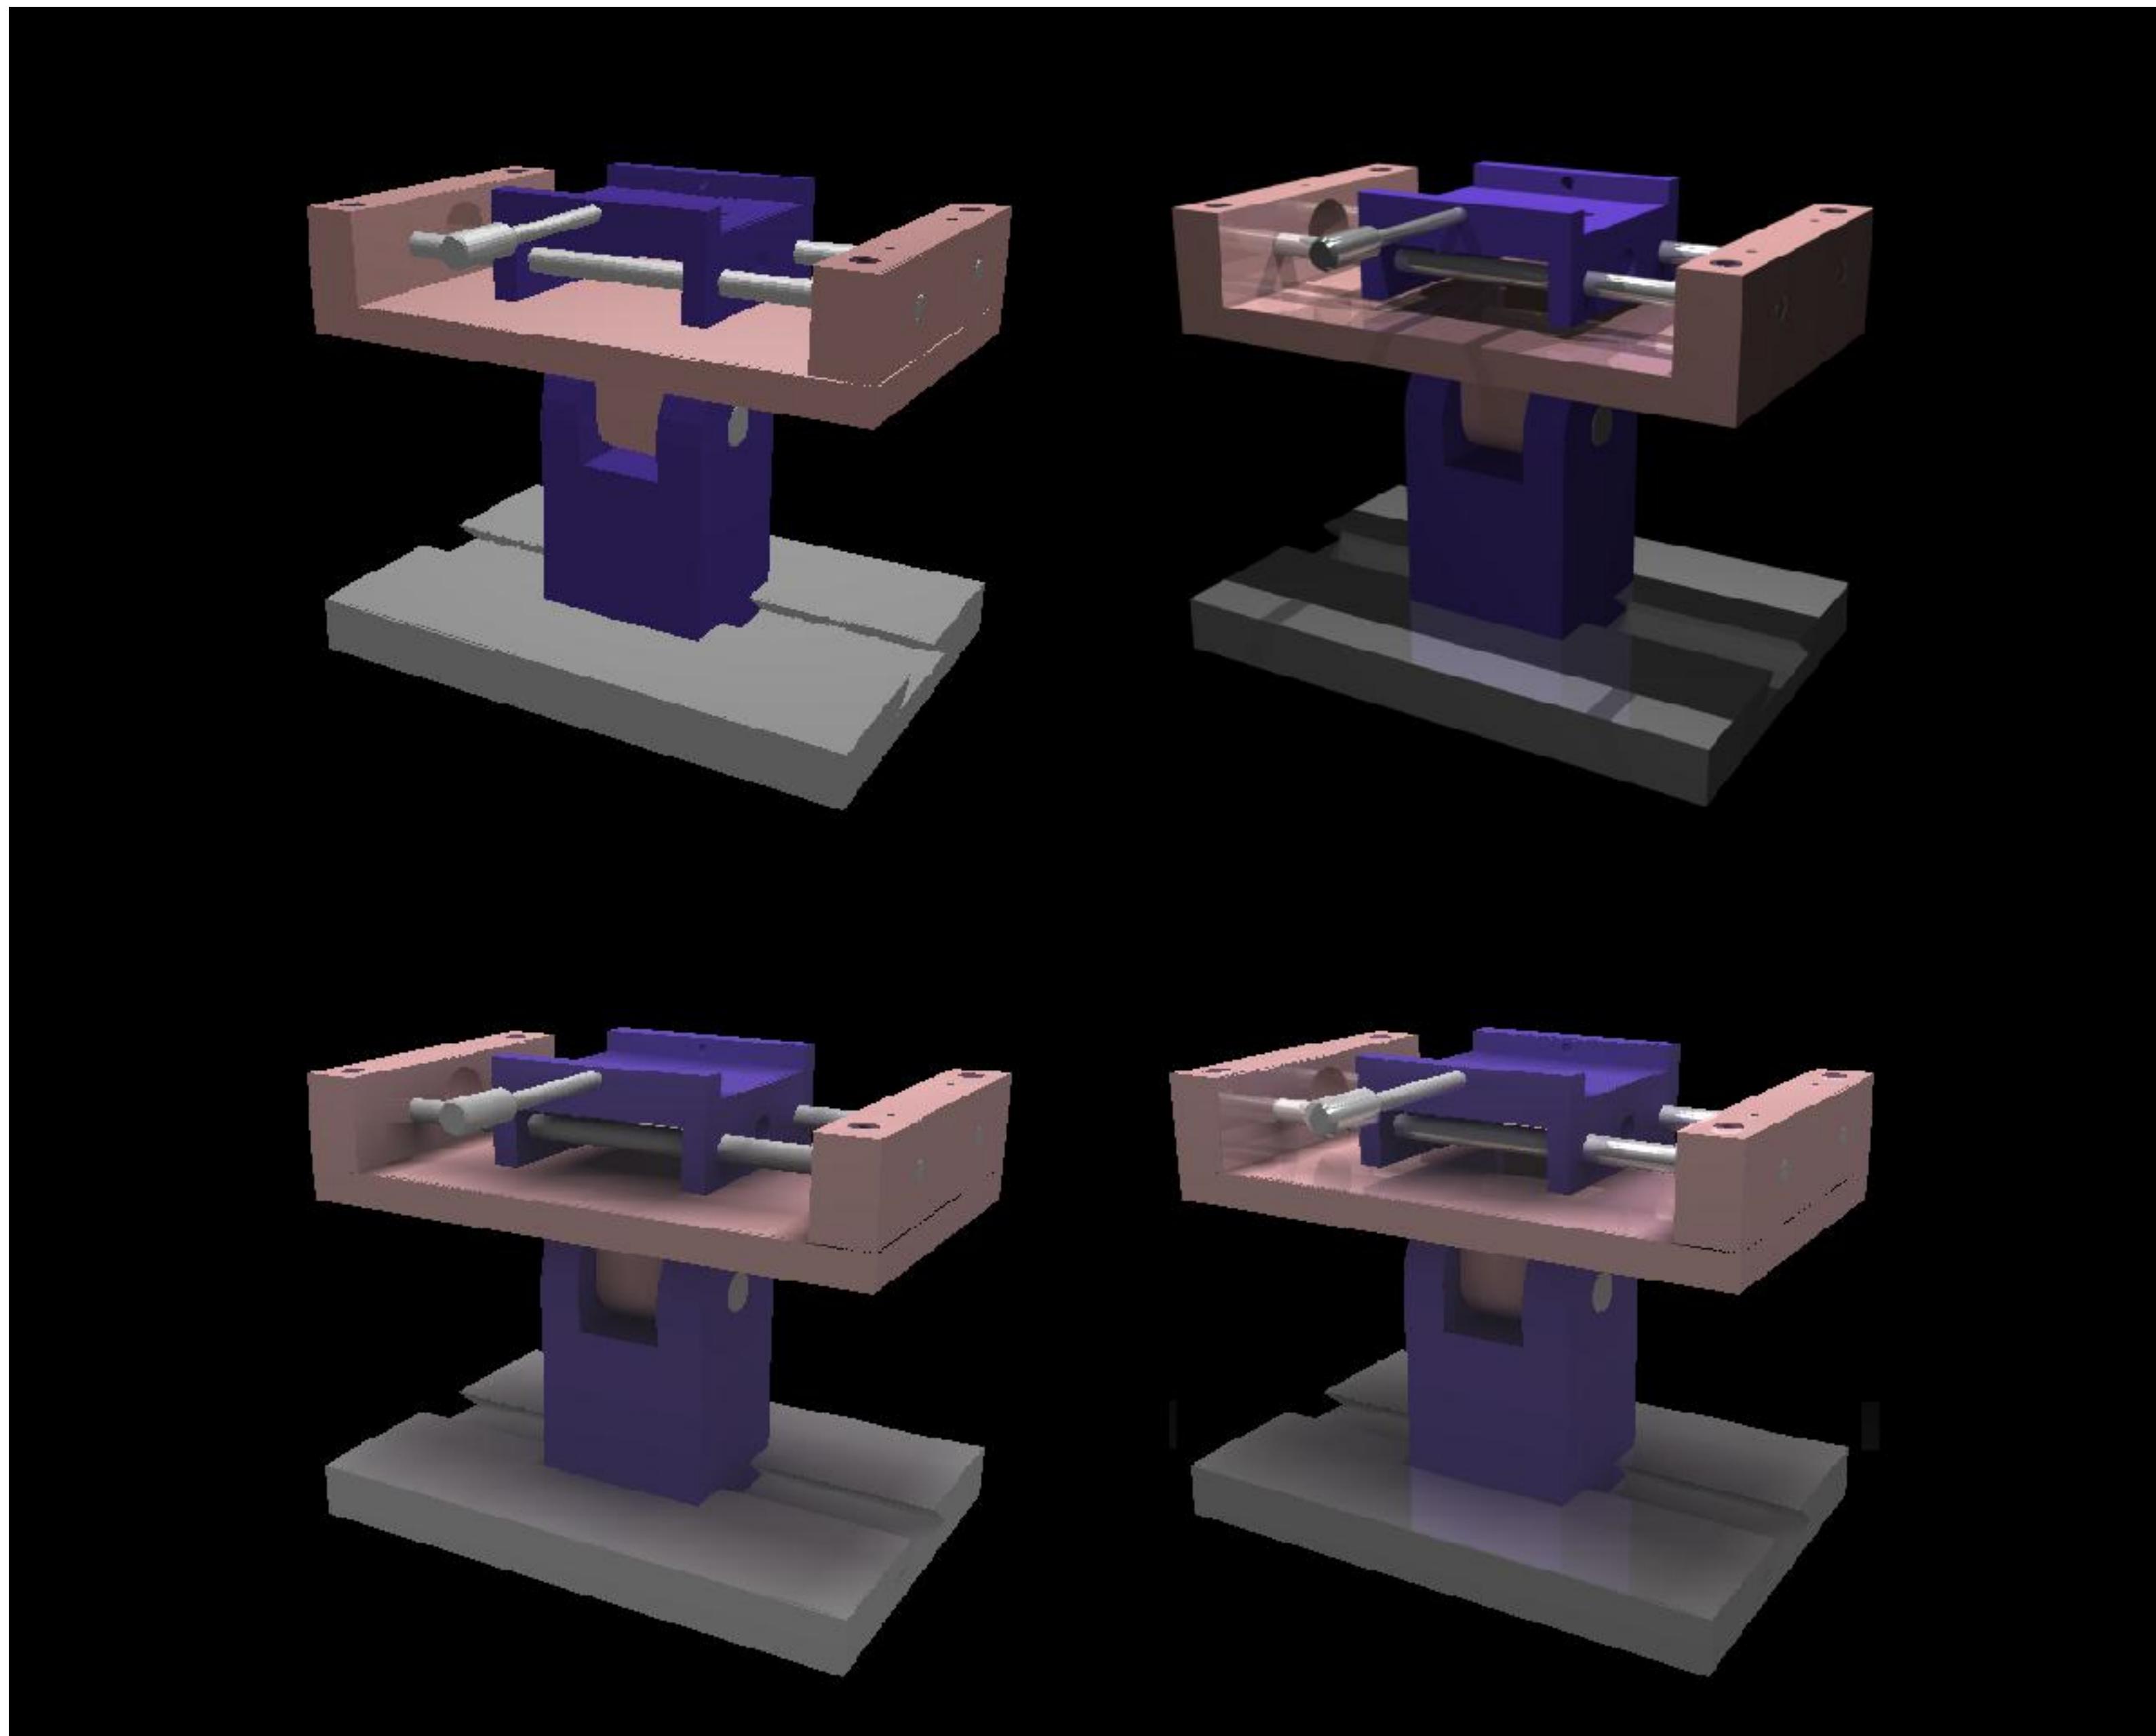
- Because radiosity involves diffuse light, it is independent of the viewpoint
- Can precompute ahead of time if objects and light sources *do not change*
- But if something other than the camera moves...



Combining Approaches

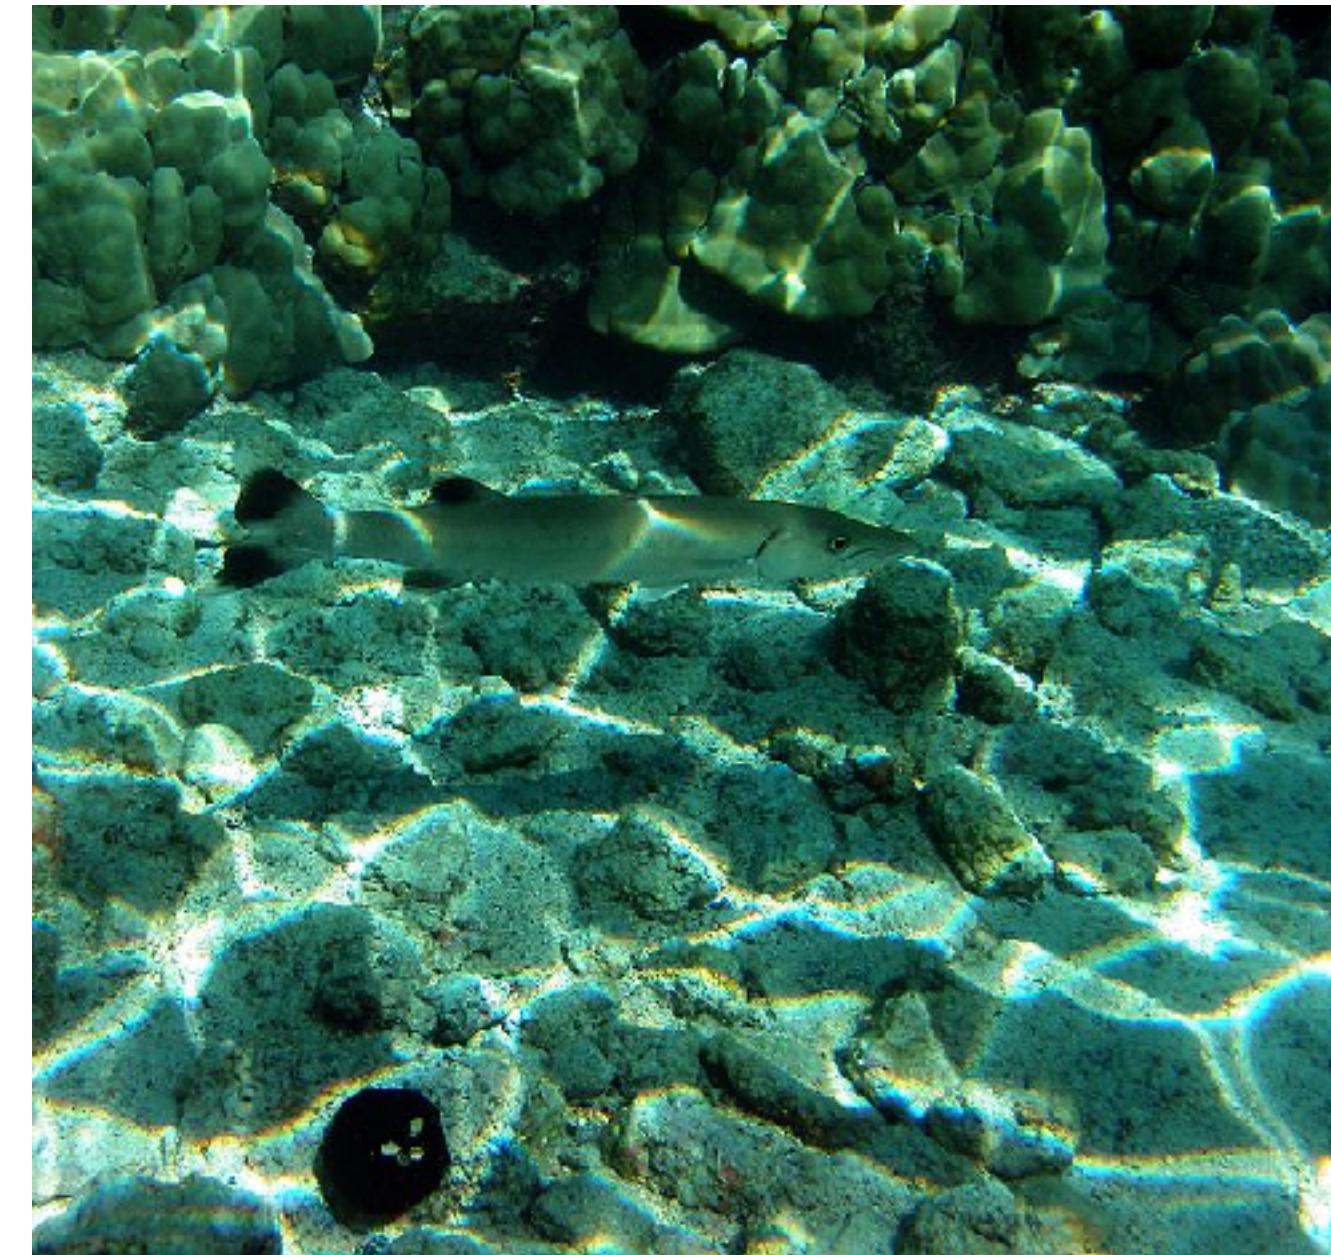
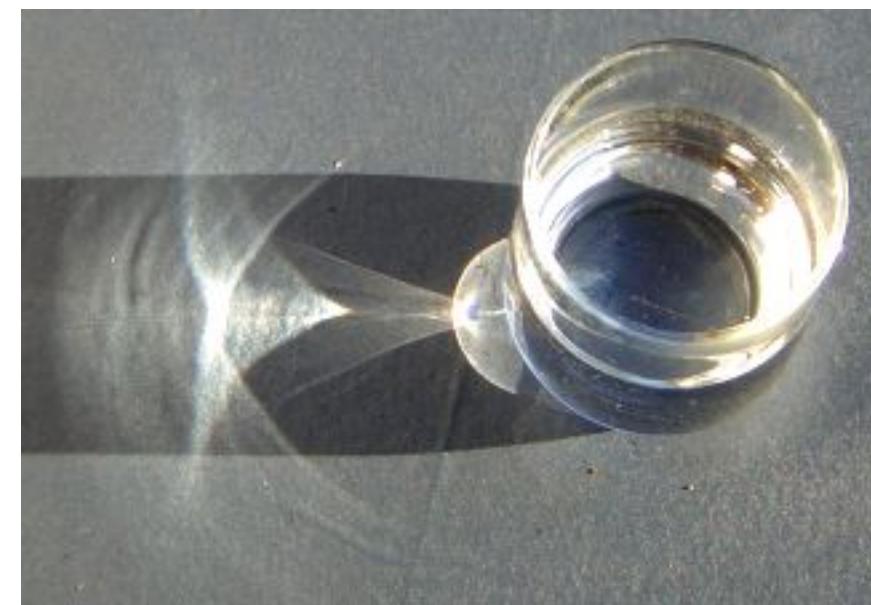
- Ray tracing is good for specular reflections, refraction, etc.
- Radiosity is good for diffuse interreflections
- Can combine methods to get both





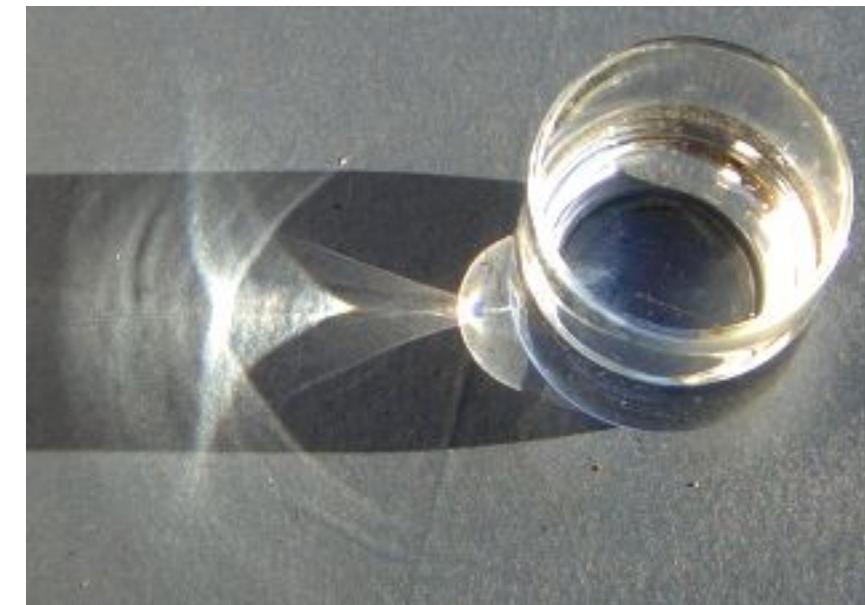
Forward Rendering

- Some phenomena can't be modeled well by going "backward" from the camera
 - ray tracing
 - radiosity
- Some require "forward" rendering of the light from the light source
- Example: "caustics"



Forward Rendering

- Forward or hybrid methods:
 - Path tracing
 - Metropolis light transport
 - Photon mapping
- All try to *approximate* the “rendering equation”
(Jim Kajiya, 1986)



More in CS 455...

- Ray tracing and variations
- Radiosity
- Other lighting models
- Curves and surface modeling
(we'll talk a bit about this coming up)
- Animation (movement)
- Physical-based modeling
- And much more...

The Rest of the Semester...

- Physical cameras revisited
- Interpolation
- Image warping
- A bit more geometry...
- Curves and surfaces
- Frequency-domain processing

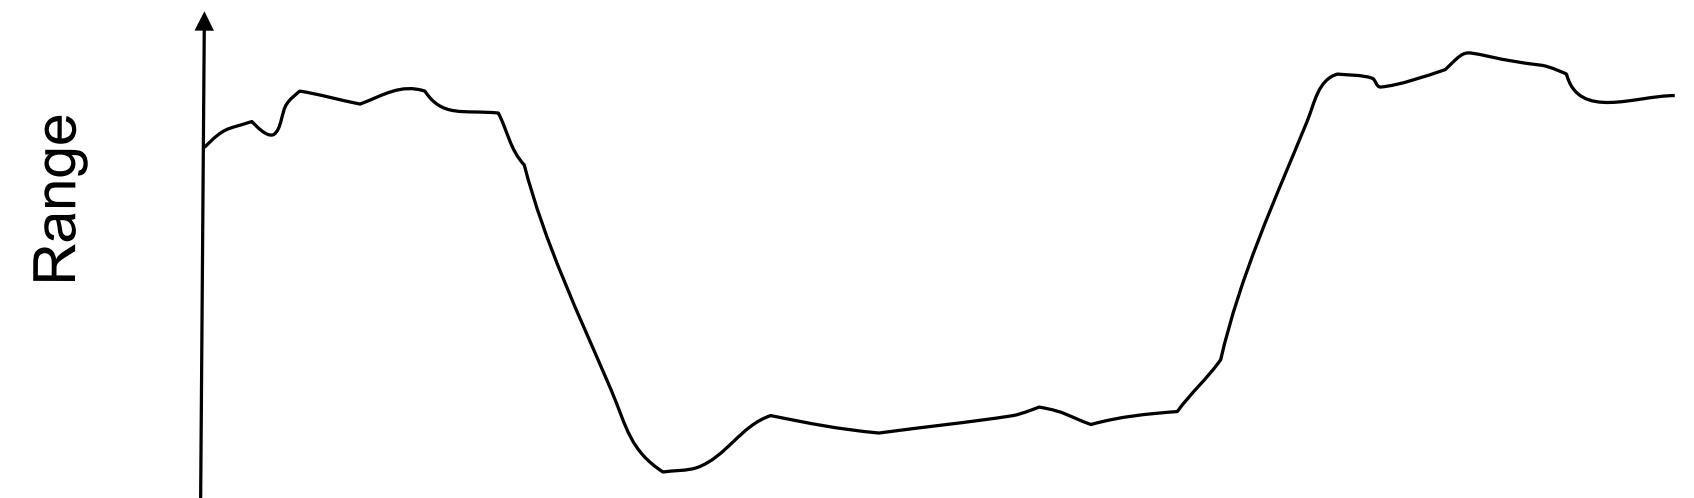


Cameras and Signal Acquisition

CS 355: Introduction to Graphics and Image Processing

Signals as Functions

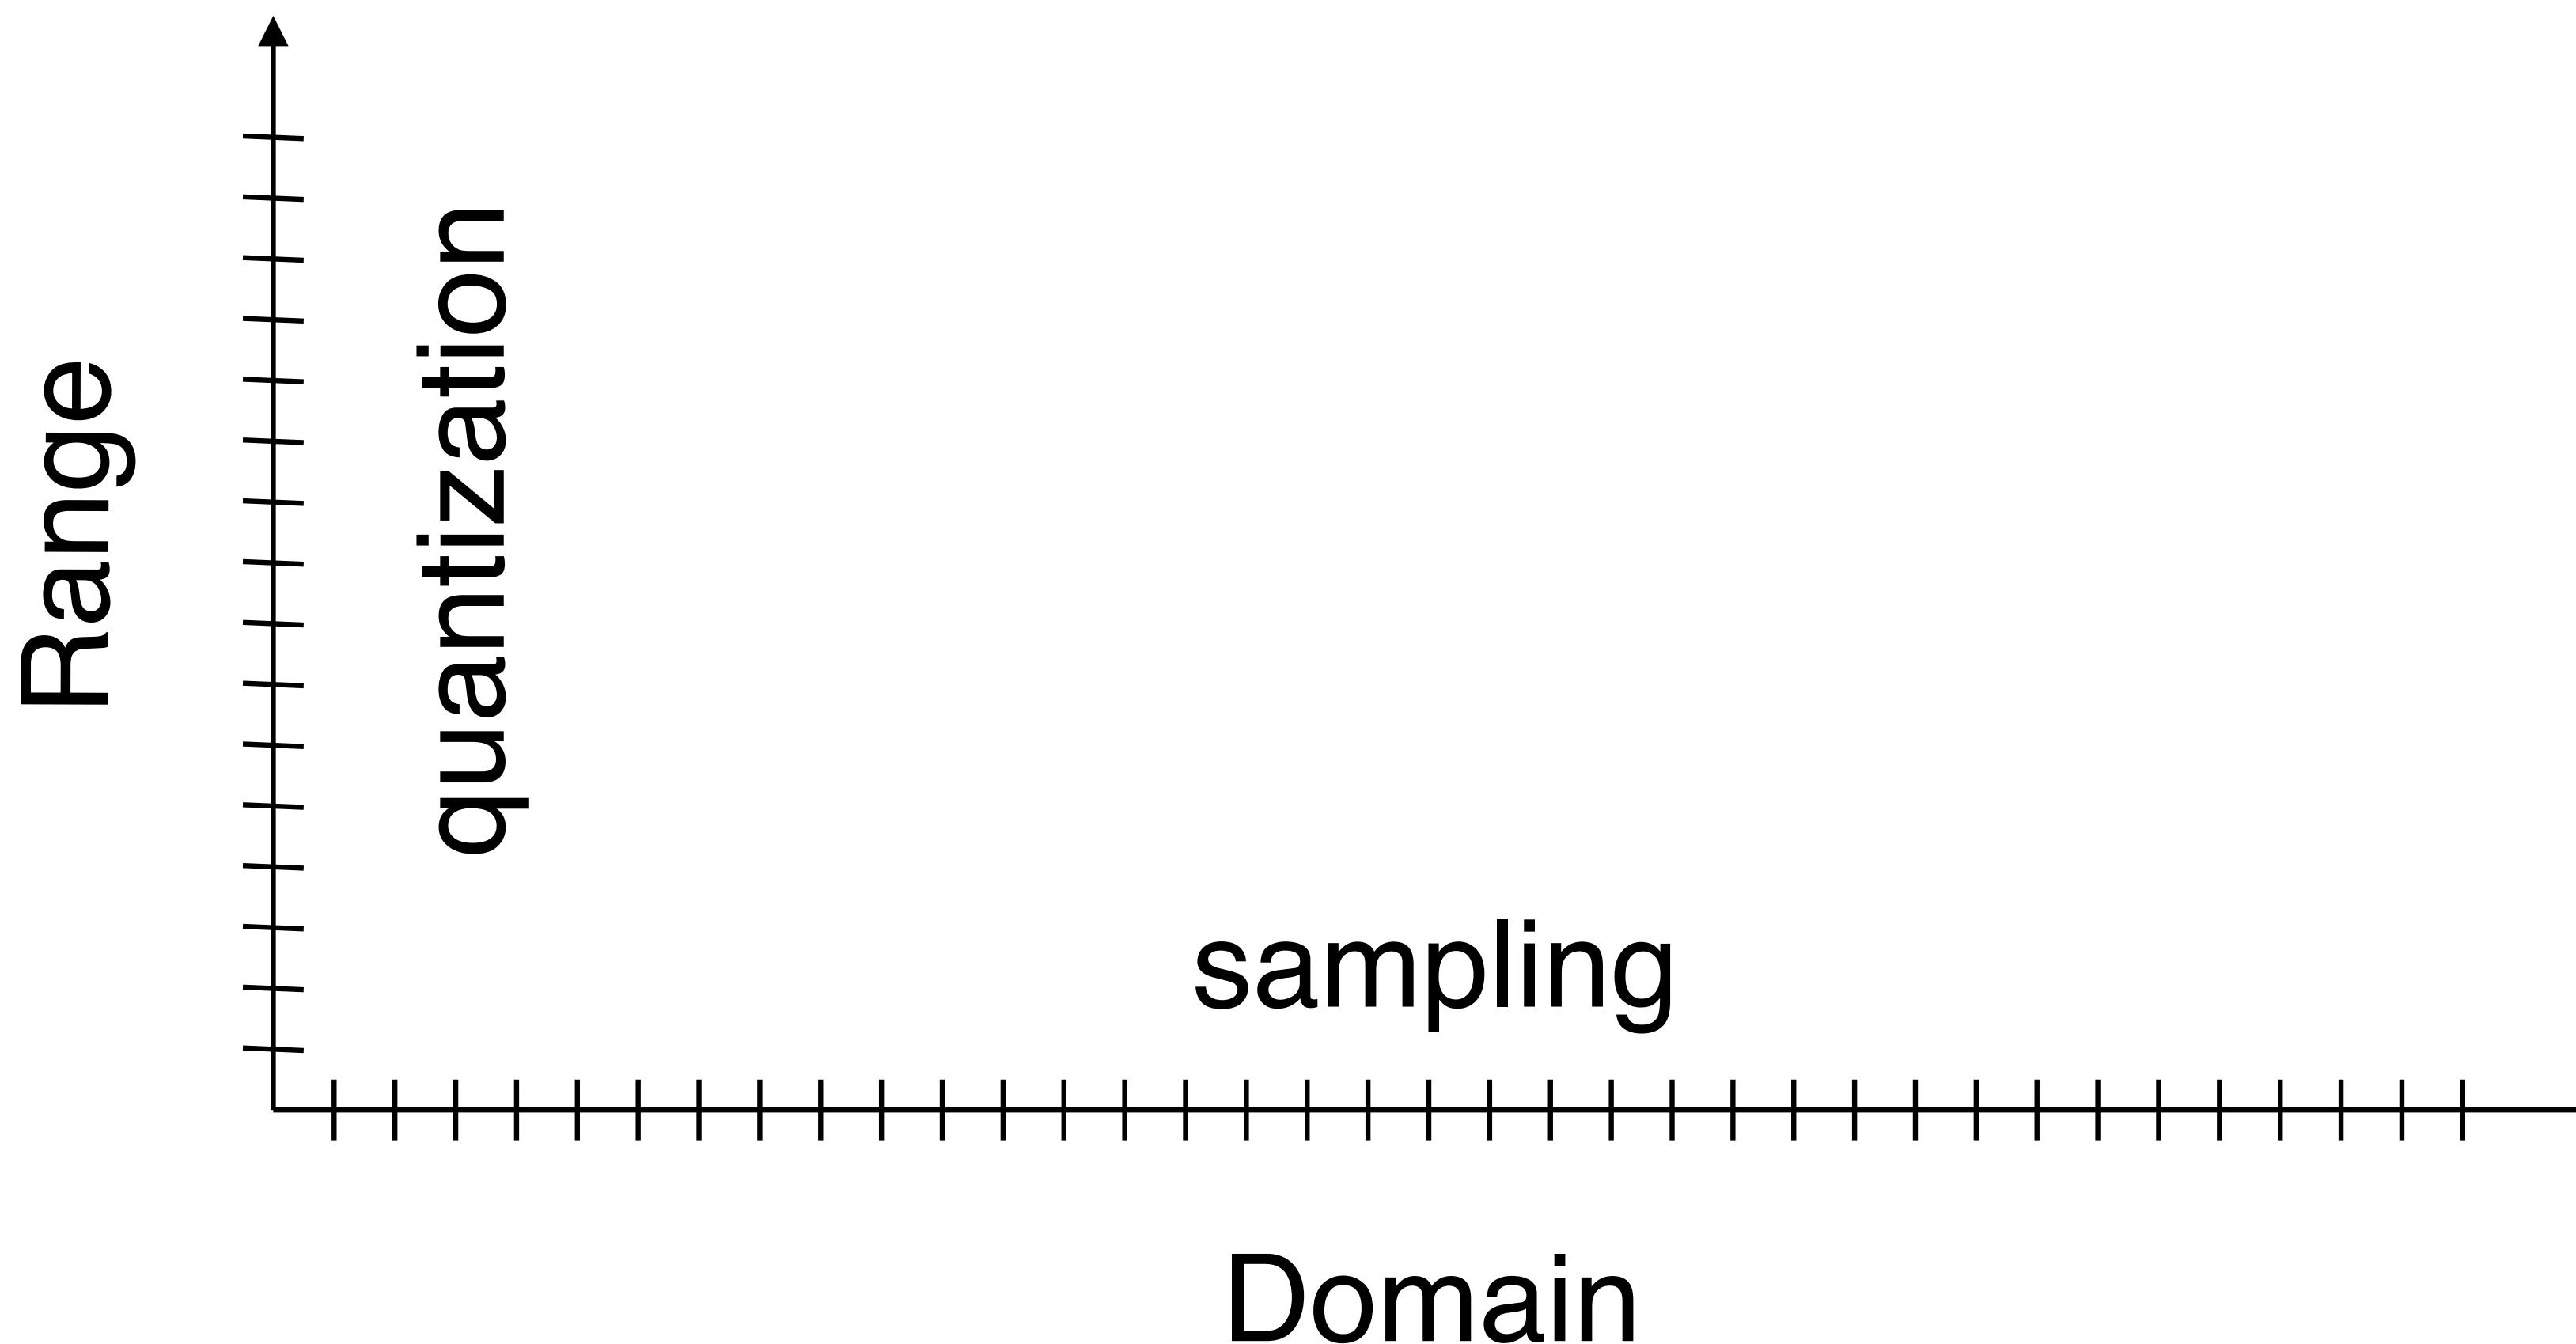
- Digital signals can be thought of as *sampled functions*
- Domains:
 - Time (audio)
 - Space (images)
 - Both (video)
- Ranges:
 - Changing air pressure (audio)
 - Visible light (photographs, video)
 - Other properties
(X-rays, MRI, range images, etc.)



$$f(t)$$

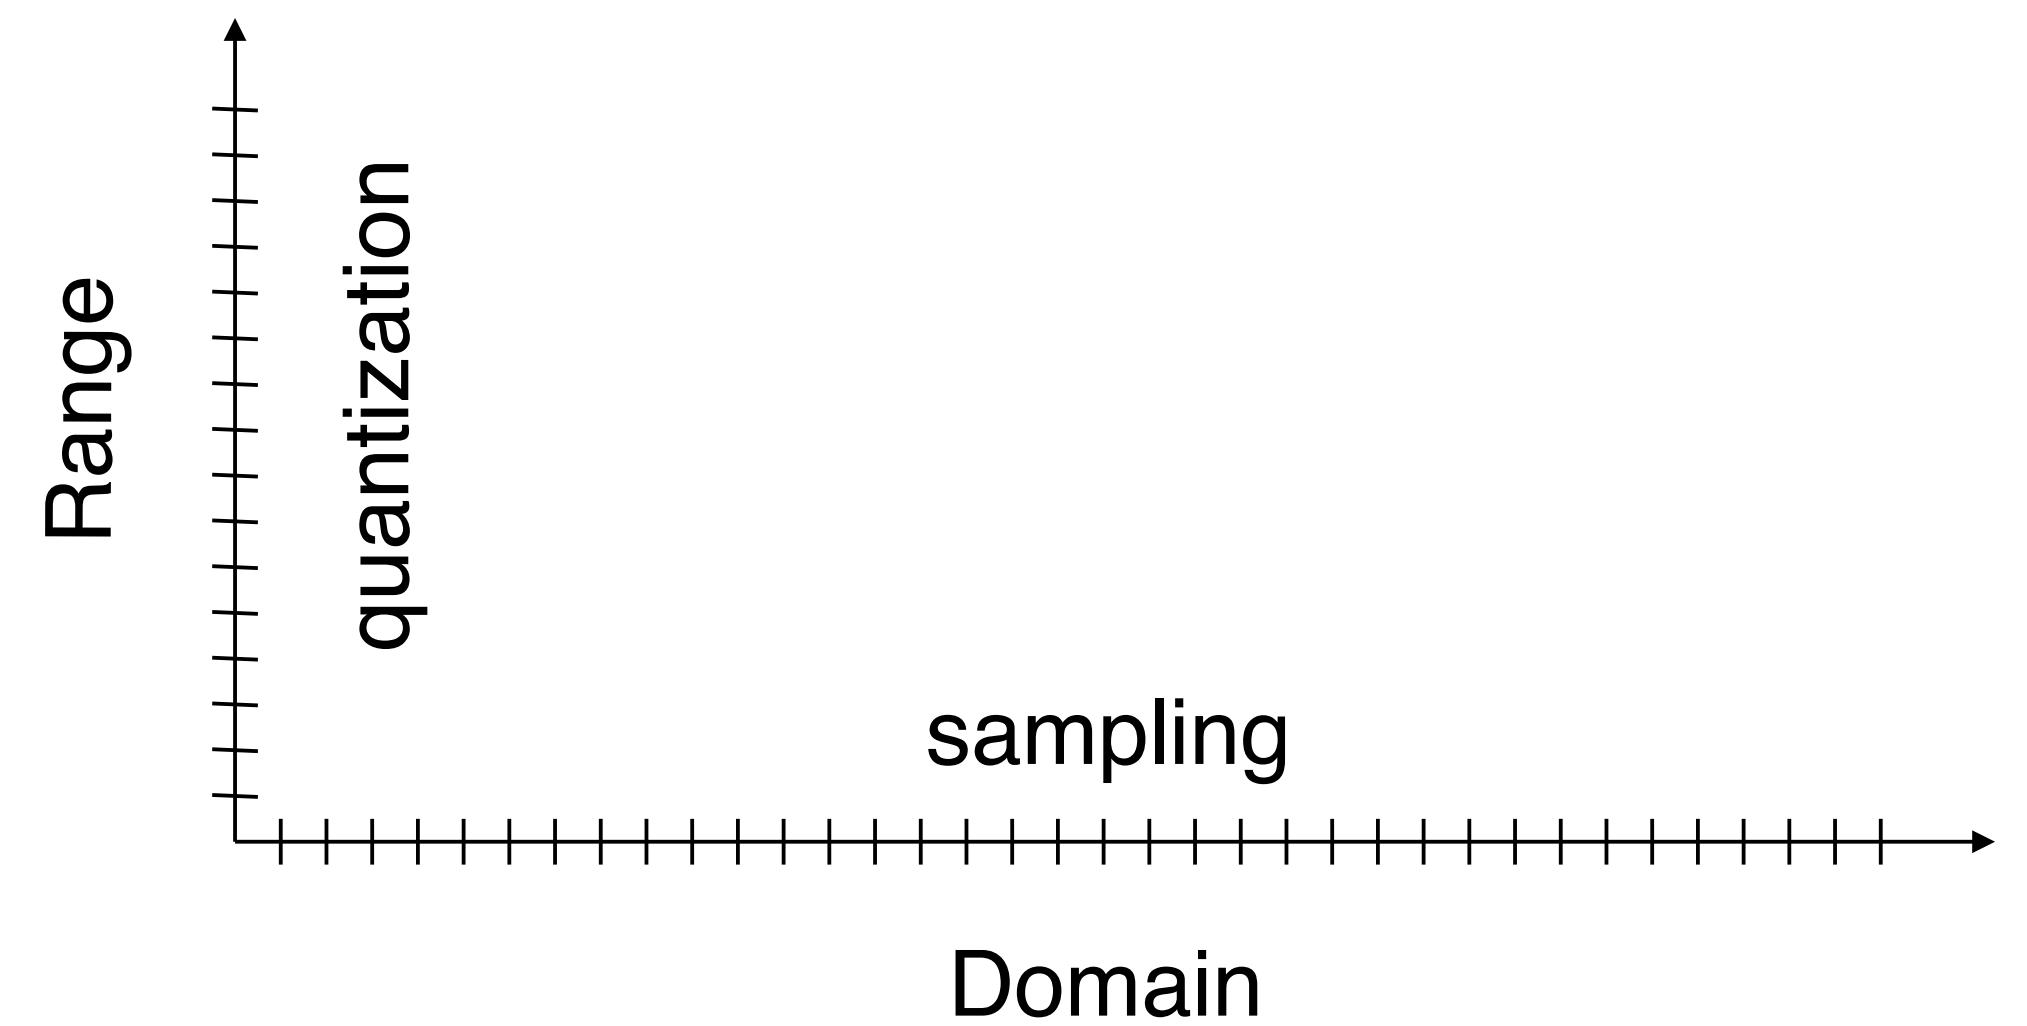
$$I(x, y)$$

Sampling vs. Quantization



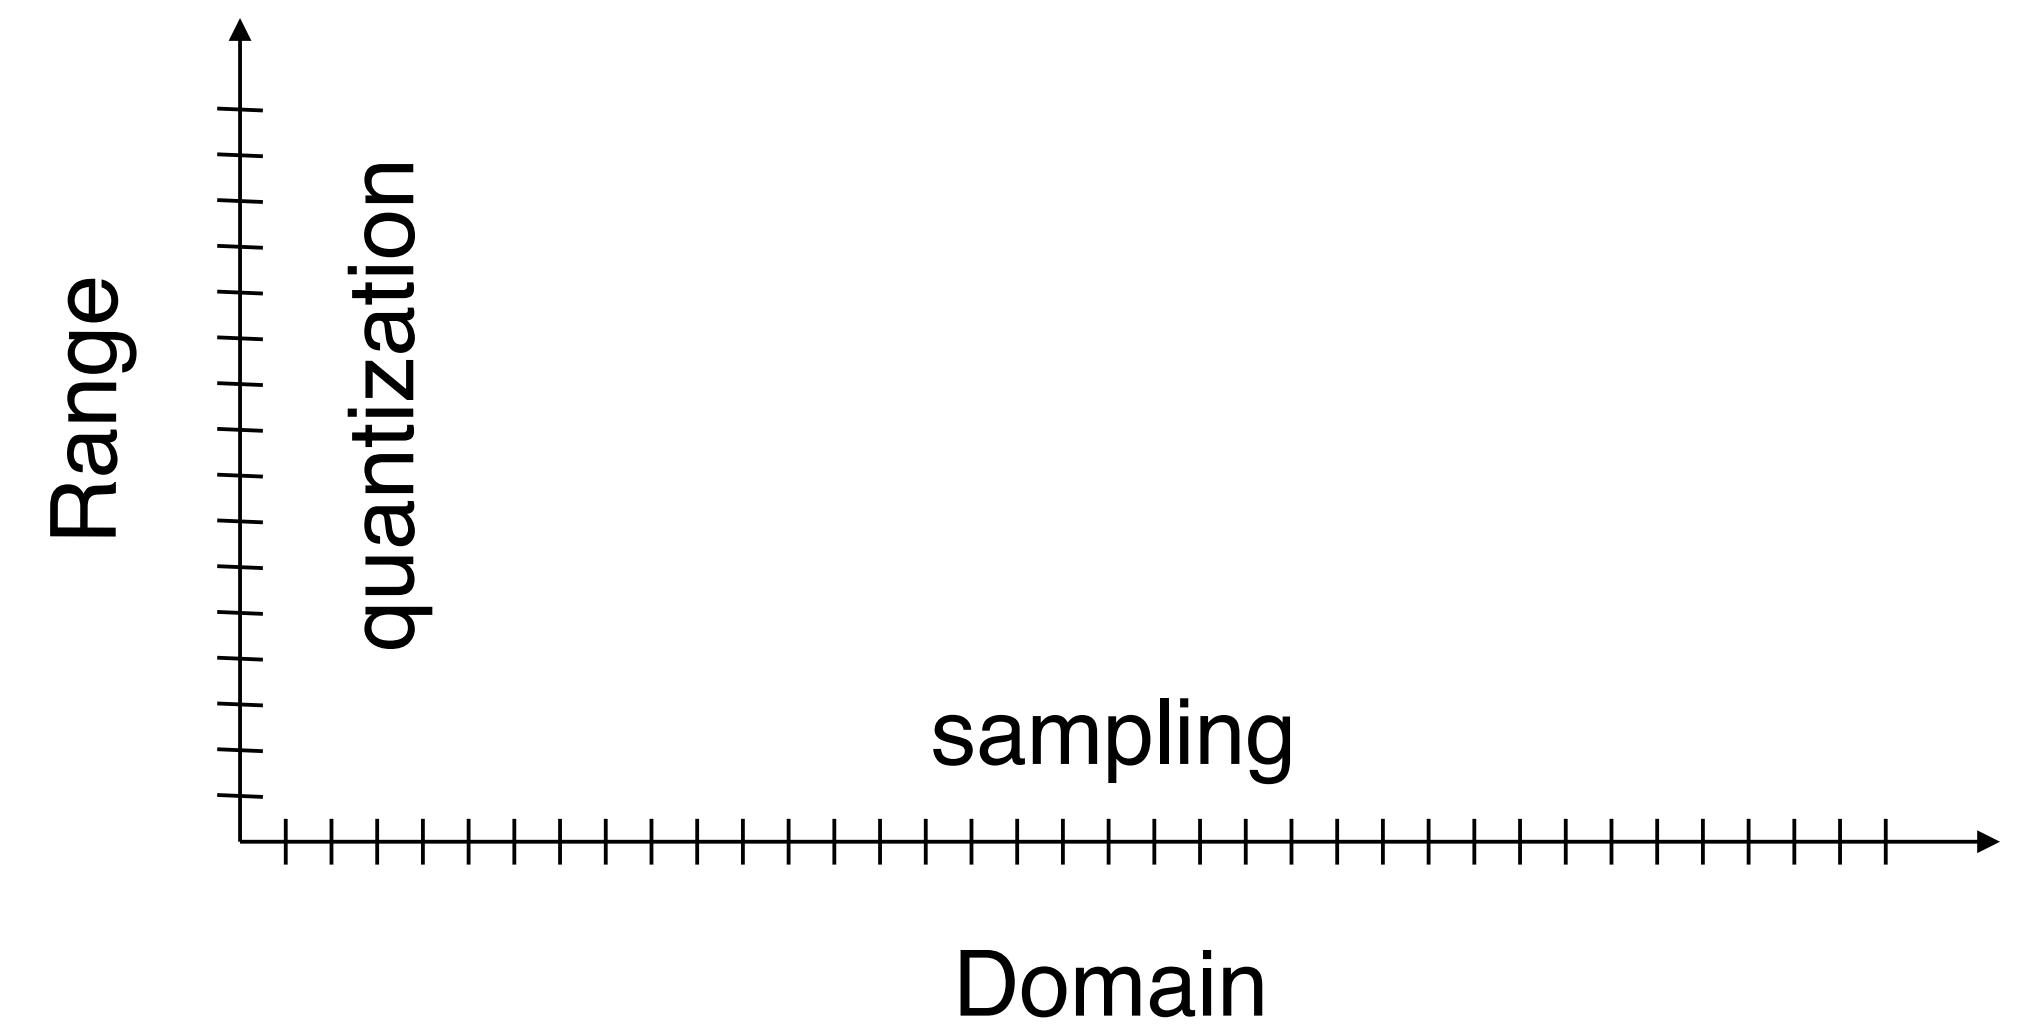
Sampling

- Samples per unit length, area, etc.
- Often expressed as
 - rate
 - spacing
 - density



Quantization

- *Levels of precision* in each sample
- Usually
 - number of levels
 - number of bits



Sampling vs. Quantization

- 600 dots per inch
- black and white images
- 256-level grey
- 8-bit grey
- 30 frames per second
- 24-bit color
- 44.1 KHz audio
- 16-bit audio

Storage

- We usually store digital signals (including images) as arrays
 - Audio: 1-D domain, 1-D array of values (PCM)

90	50	8	...	42
----	----	---	-----	----

- Images: 2-D domain, 2-D array

88	86	8	...	9
91	92	10	...	7
87	91	9	...	8
:	:	:	..	:
90	89	11	...	8

Storage

- In memory: usually just arrays
 - Again, be careful of (x,y) vs. row-major ordering
- On disk: may be something else entirely
 - Tiled storage (think virtual memory)
 - Hierarchical/Interlaced
 - Compressed
- Headers: EXIF, compression settings, etc.

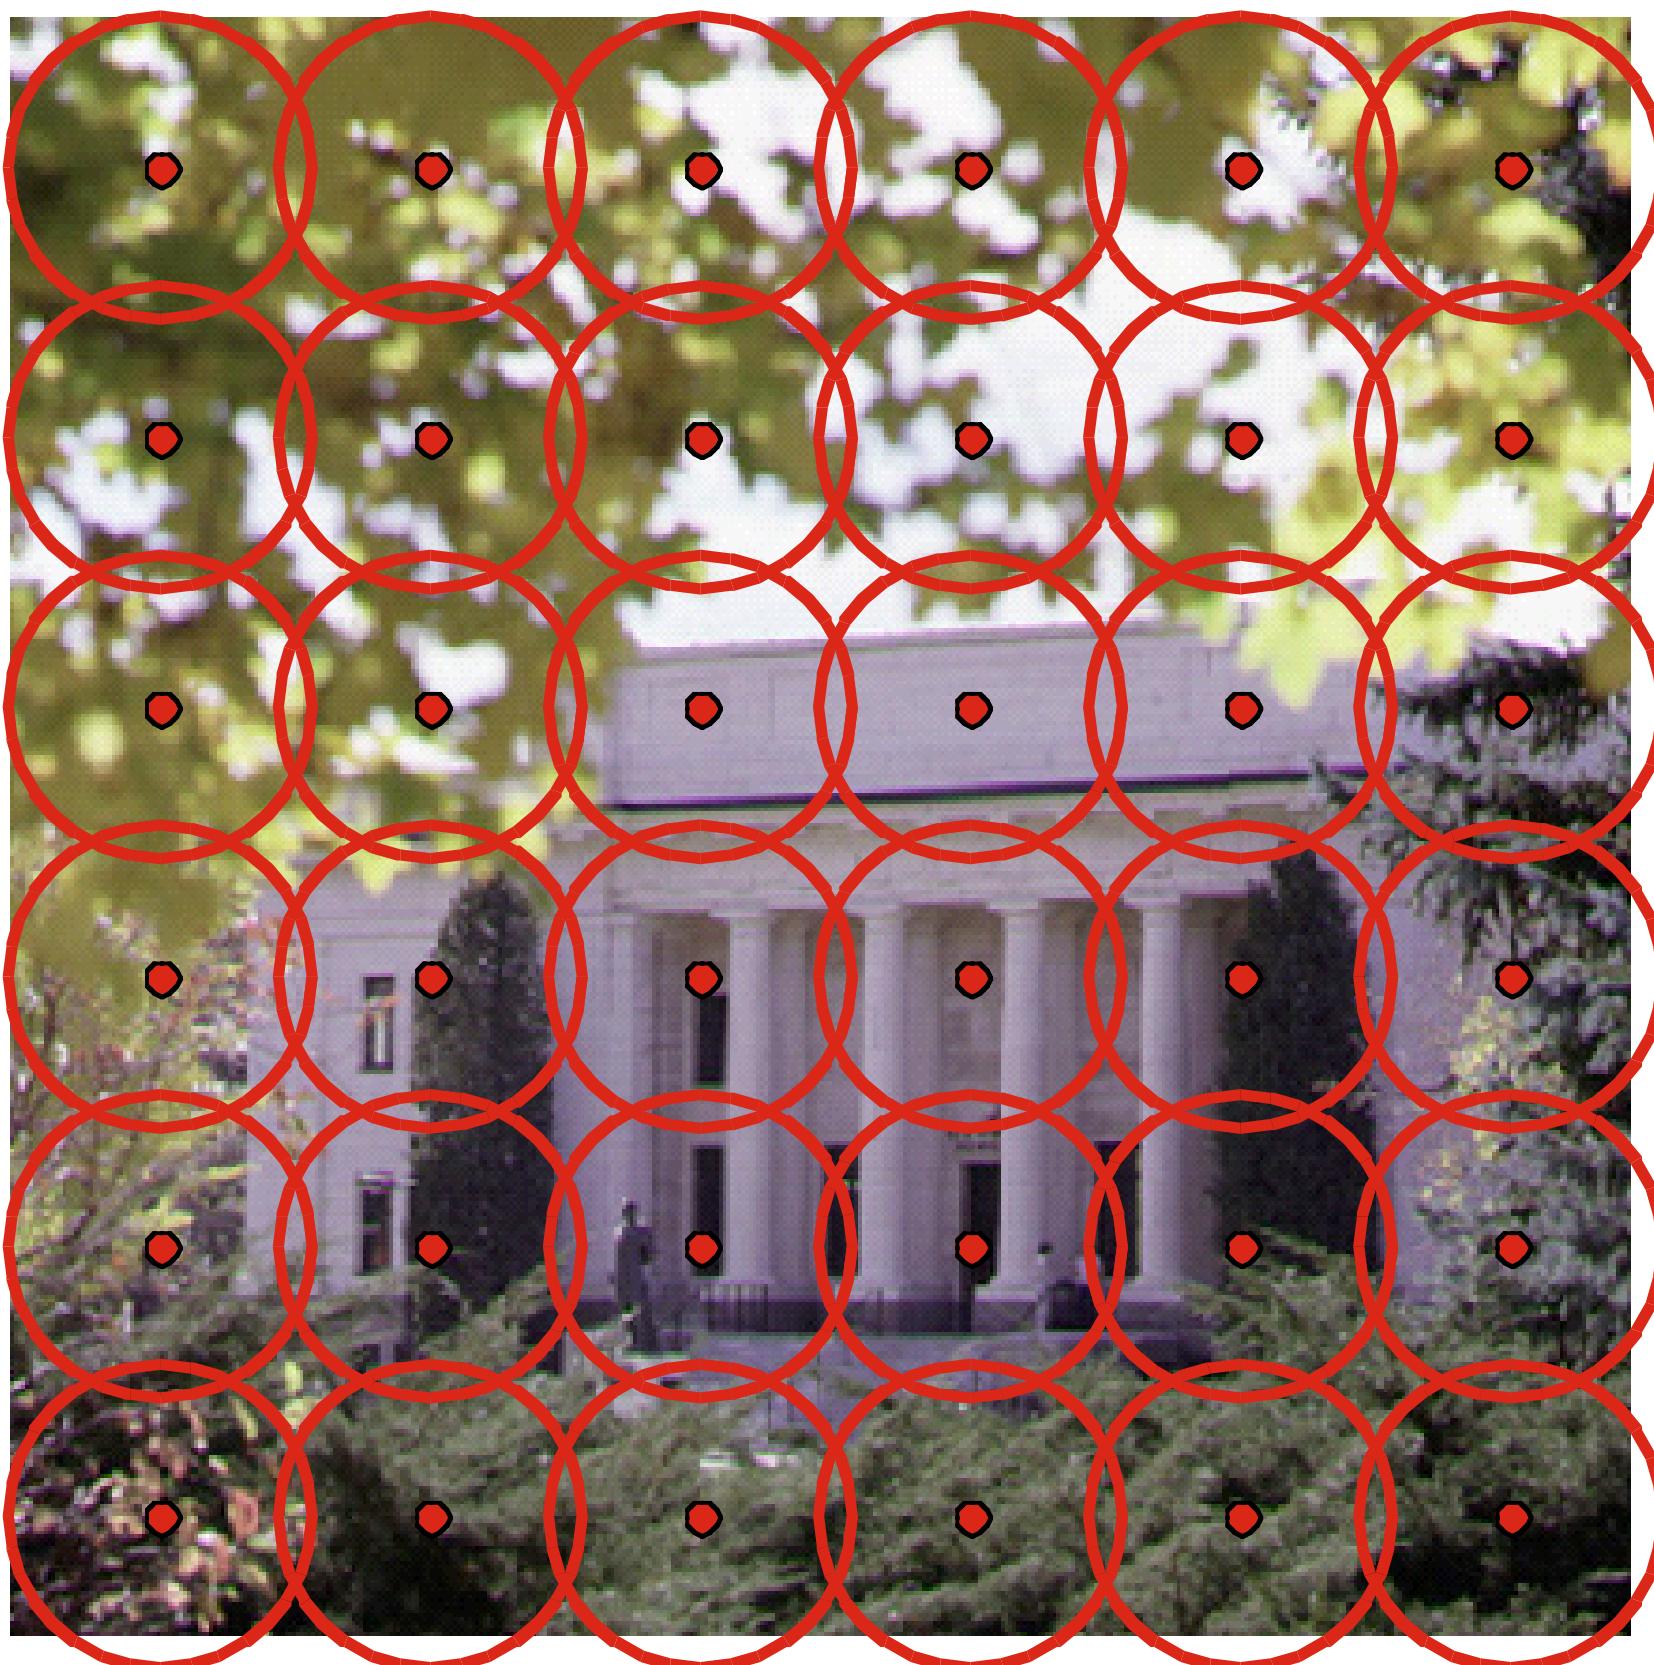
How do we get images?

Acquisition Devices

- Aperture 
- Scanning
- Sensor 
- Quantizer
- Output storage medium

Apertures

- Pixels aren't point samples
- Total light over an *area* of the visible scene
- Controlled by the camera's iris (photographers: F-stop)
- Also caused by physical sensor area (pixel's area on the camera's sensor)

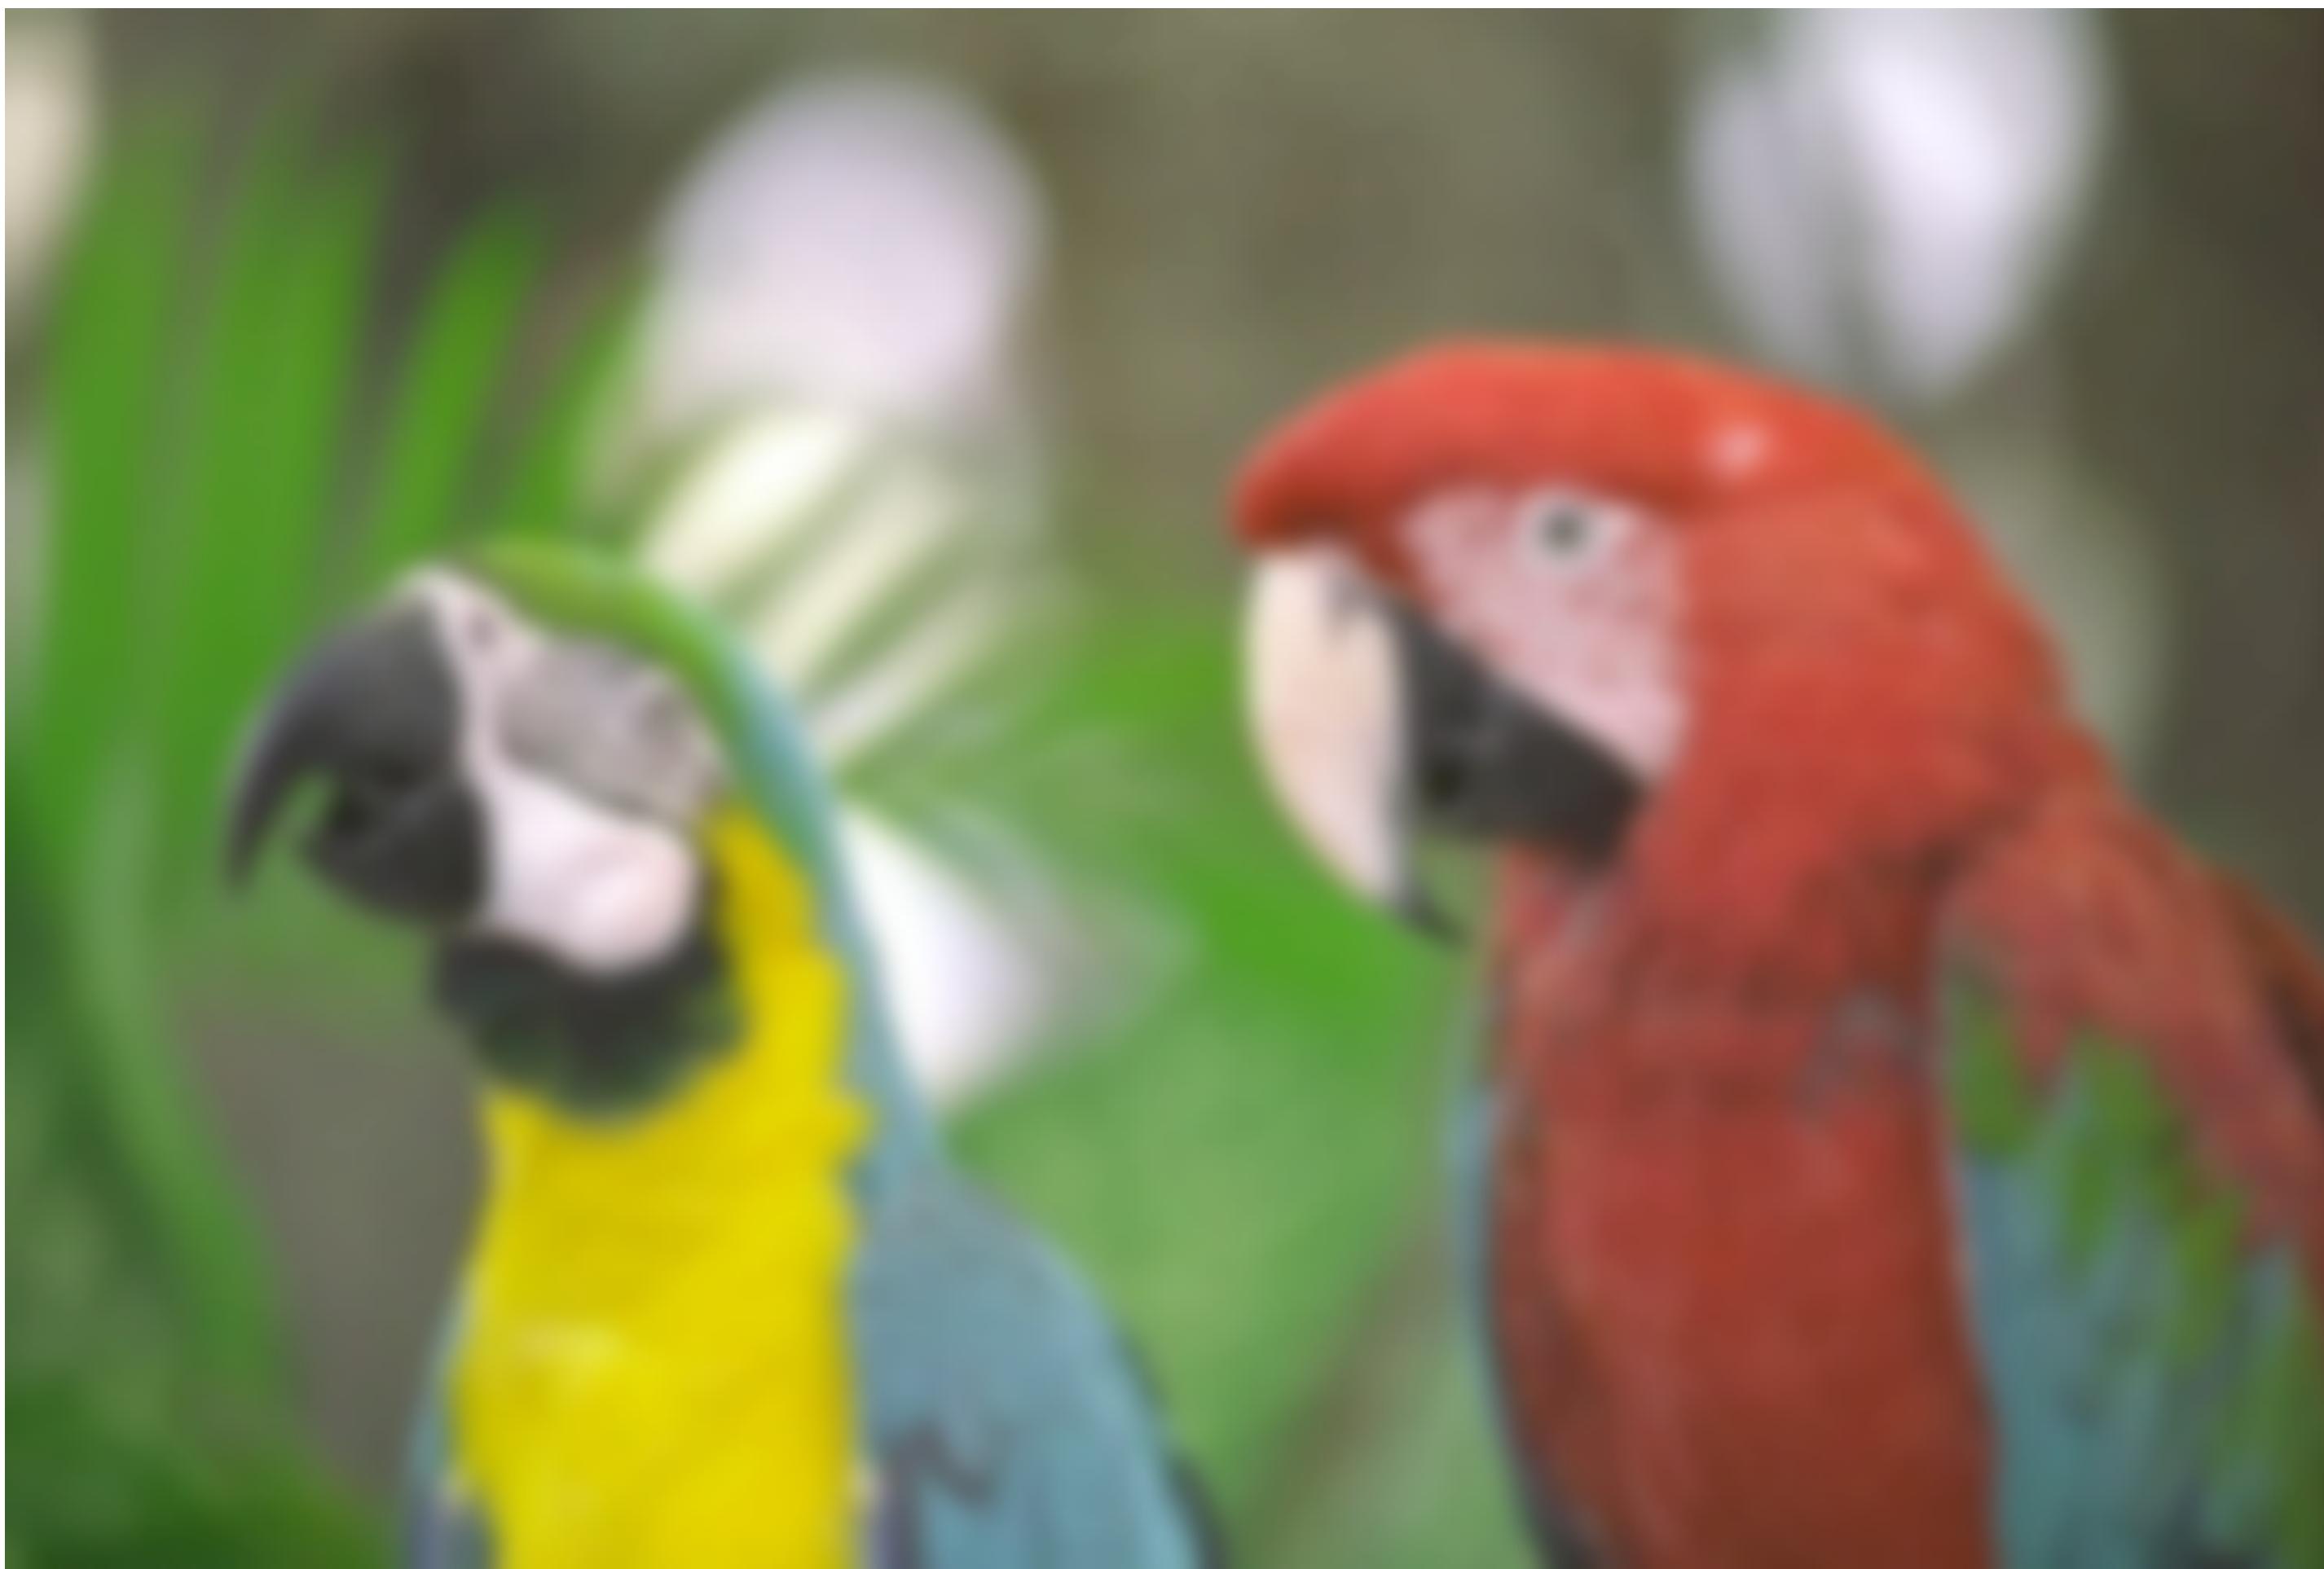


Apertures

- Sampling and size of aperture determine resolution
 - Smaller apertures = better resolution
 - Larger apertures = worse resolution
- Lenses allow a physically larger aperture to act as an effectively smaller one

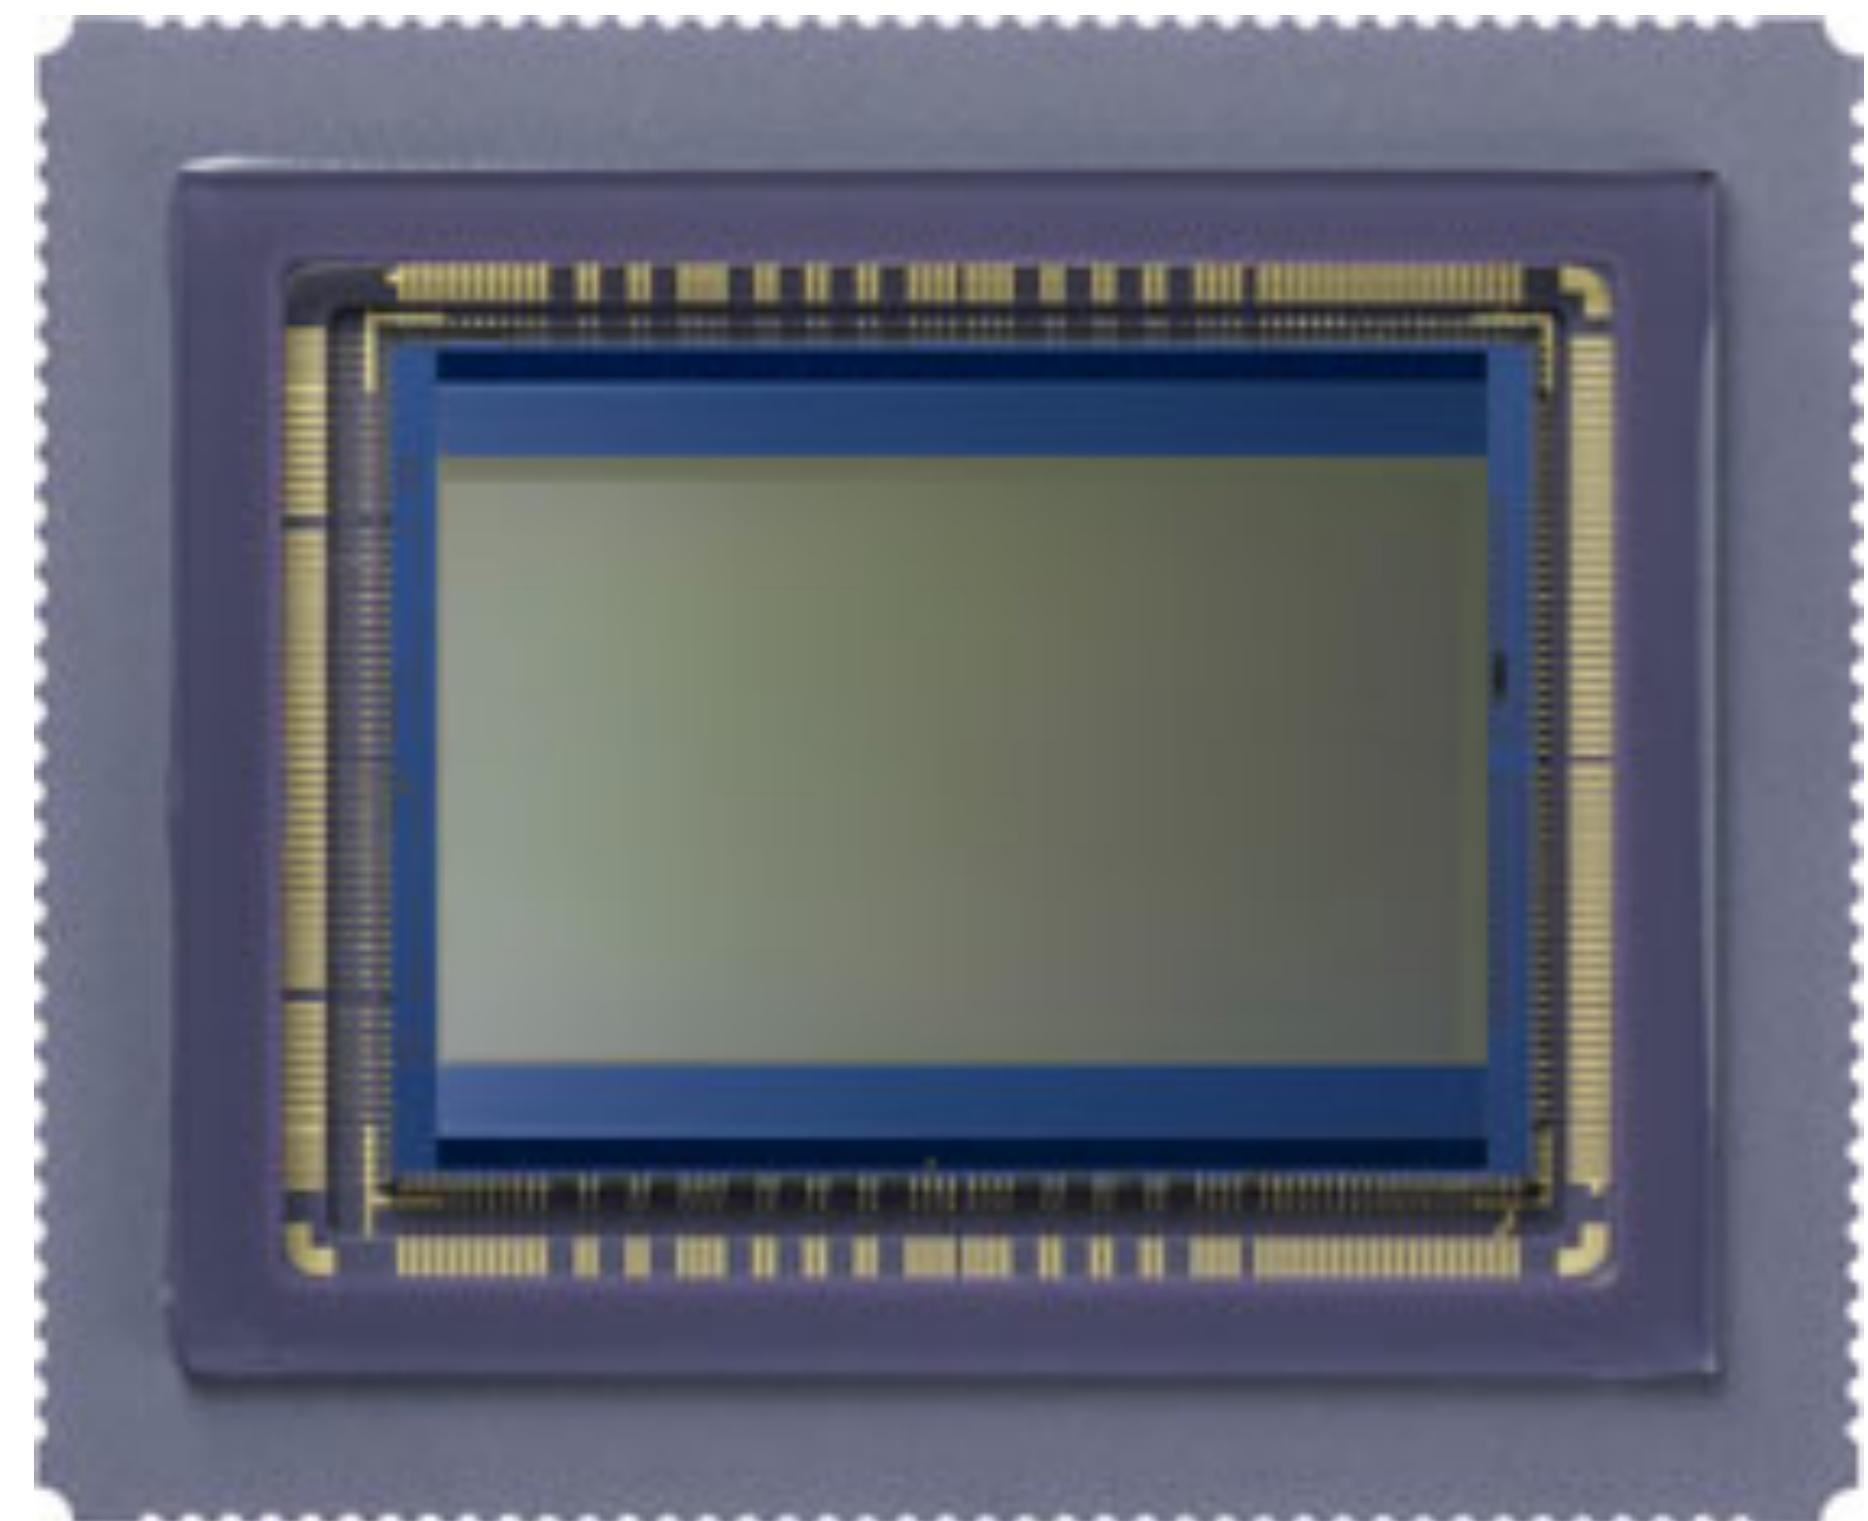


Resolution



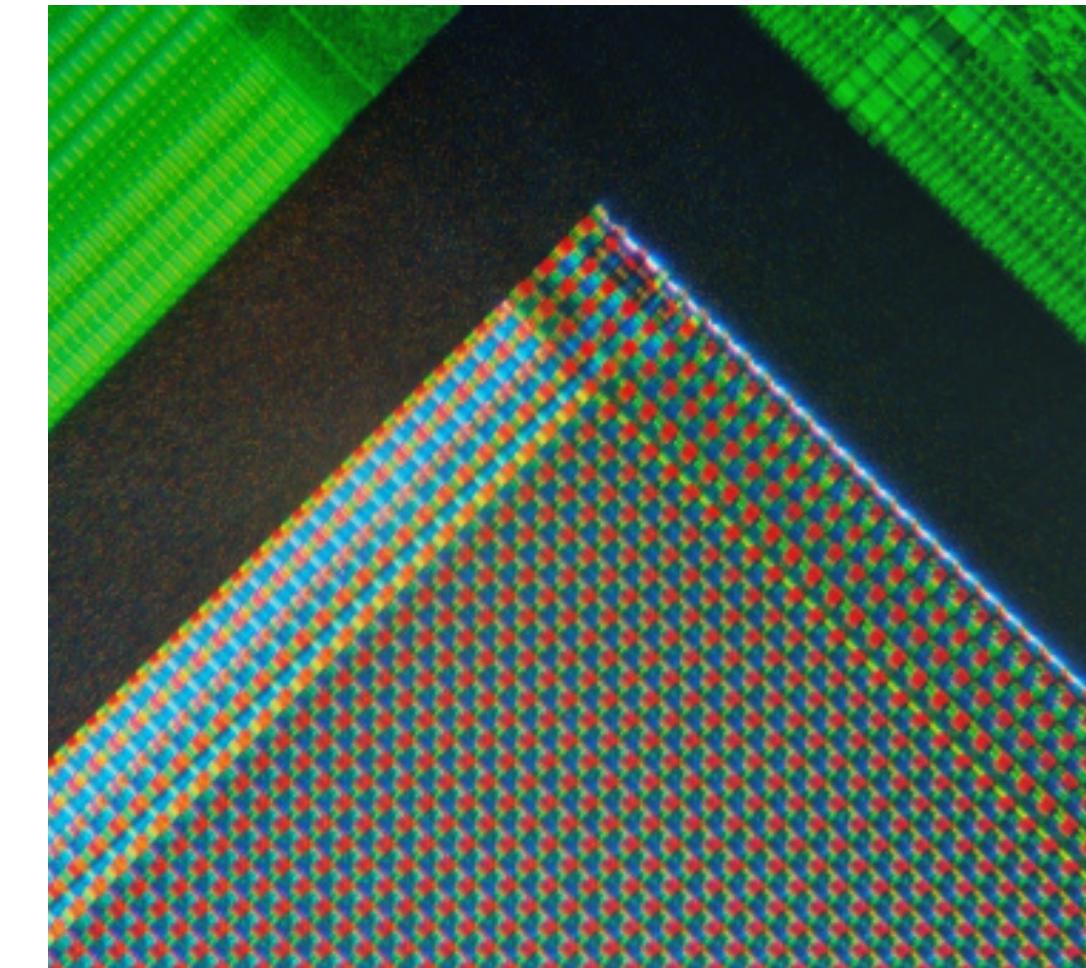
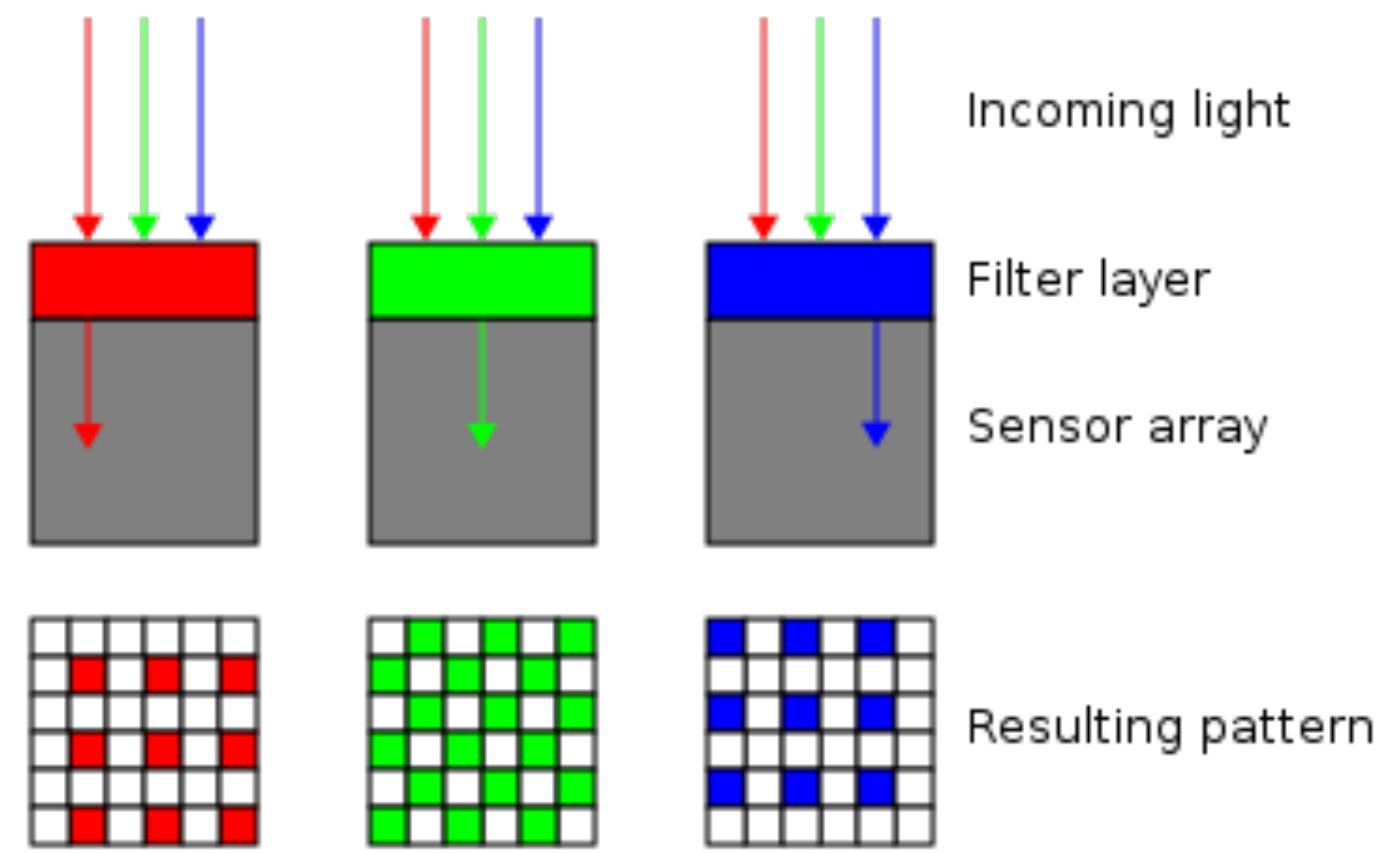
Sensors

- Converts light (photons) to chemical and/or electrical response
- Examples
 - Silver halide crystals (film)
 - Photoreceptors in our eyes (rods, cones)
 - Charge-coupled device (CCD)
 - CMOS arrays



Bayer Patterns

- Most commercial-grade cameras sample only one color per pixel
- Small colored filter over each sensor element
- 16 megapixels =
 - 8 megapixels green
 - 4 megapixels red
 - 4 megapixels blue
- You get interpolated combination



Noise

- Unavoidable random fluctuations from “correct” value
- Can usually be modeled as a statistical distribution with mean at the “correct” value
- A measured sample will vary from that mean according to the distribution std. dev.

μ

σ

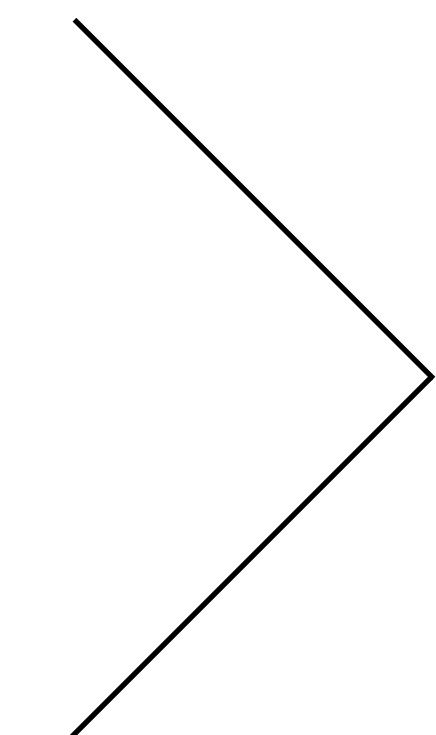
Signal-To-Noise Ratio

- Measure of how “noise free” a signal is

$$\text{SNR} = \frac{\mu}{\sigma}$$

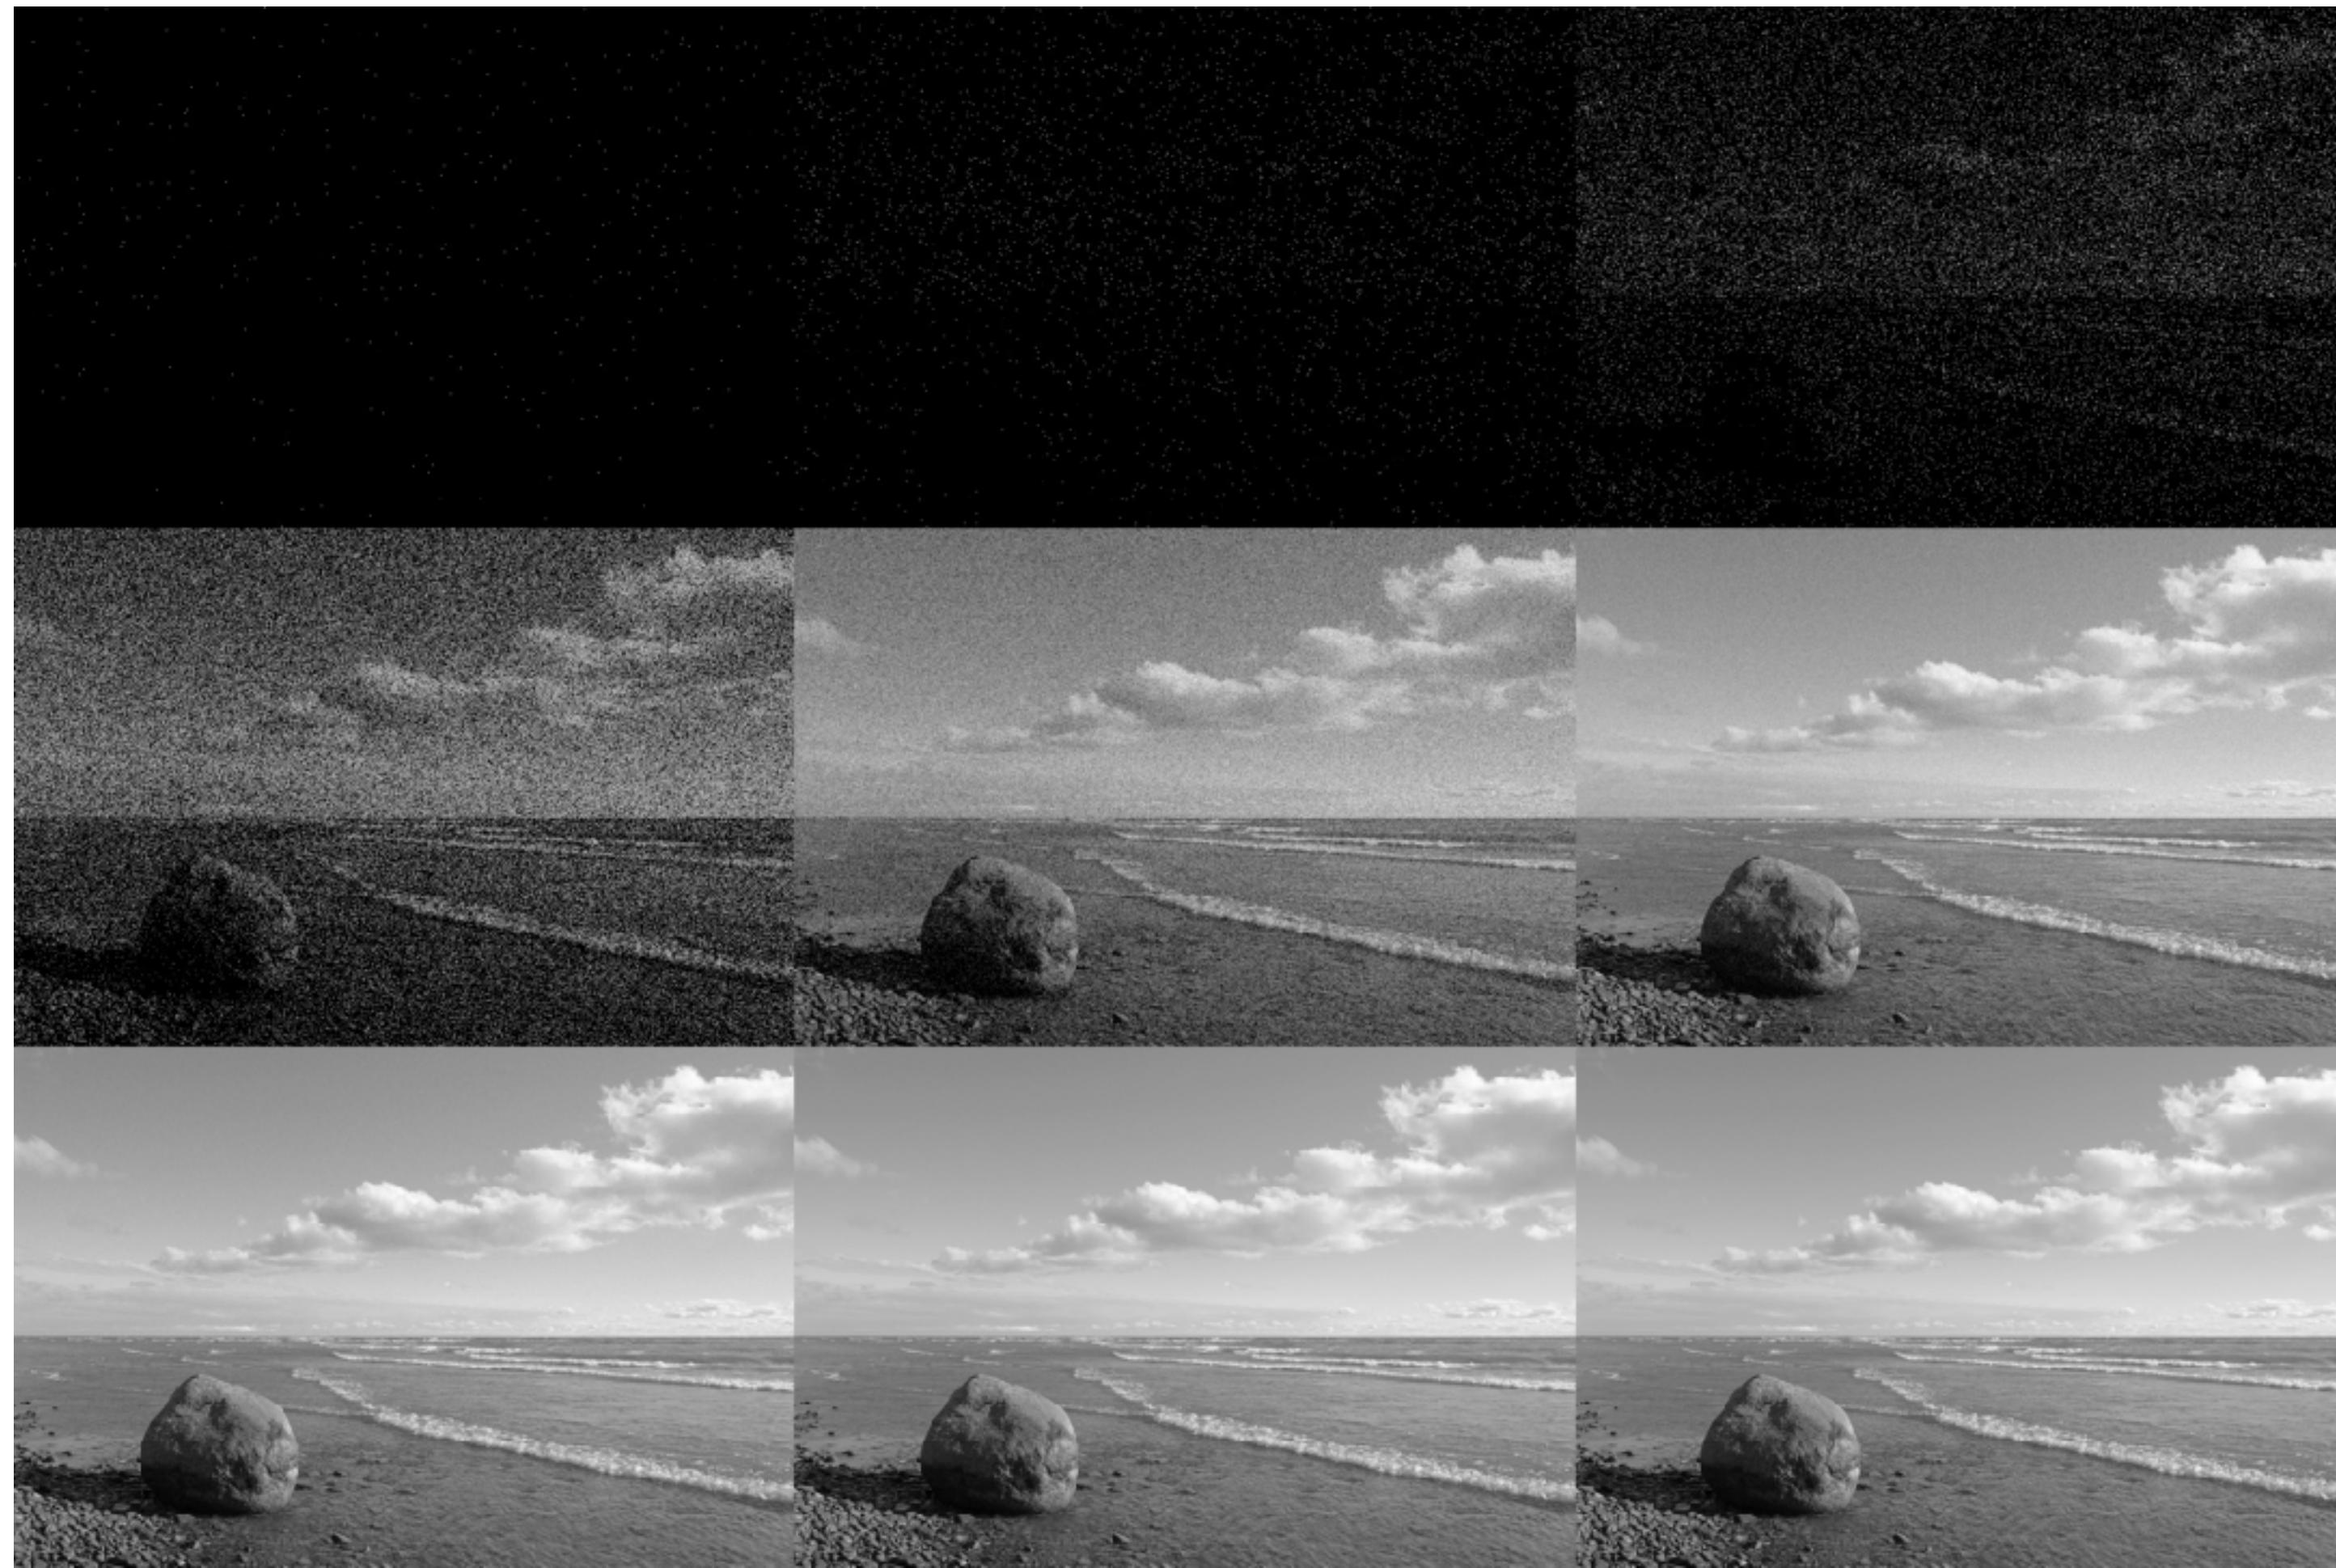
Sources of Noise

- Quantum nature of light
- Sensor inhomogeneity
- Electrical fluctuations
- “Background” noise



May not be
random

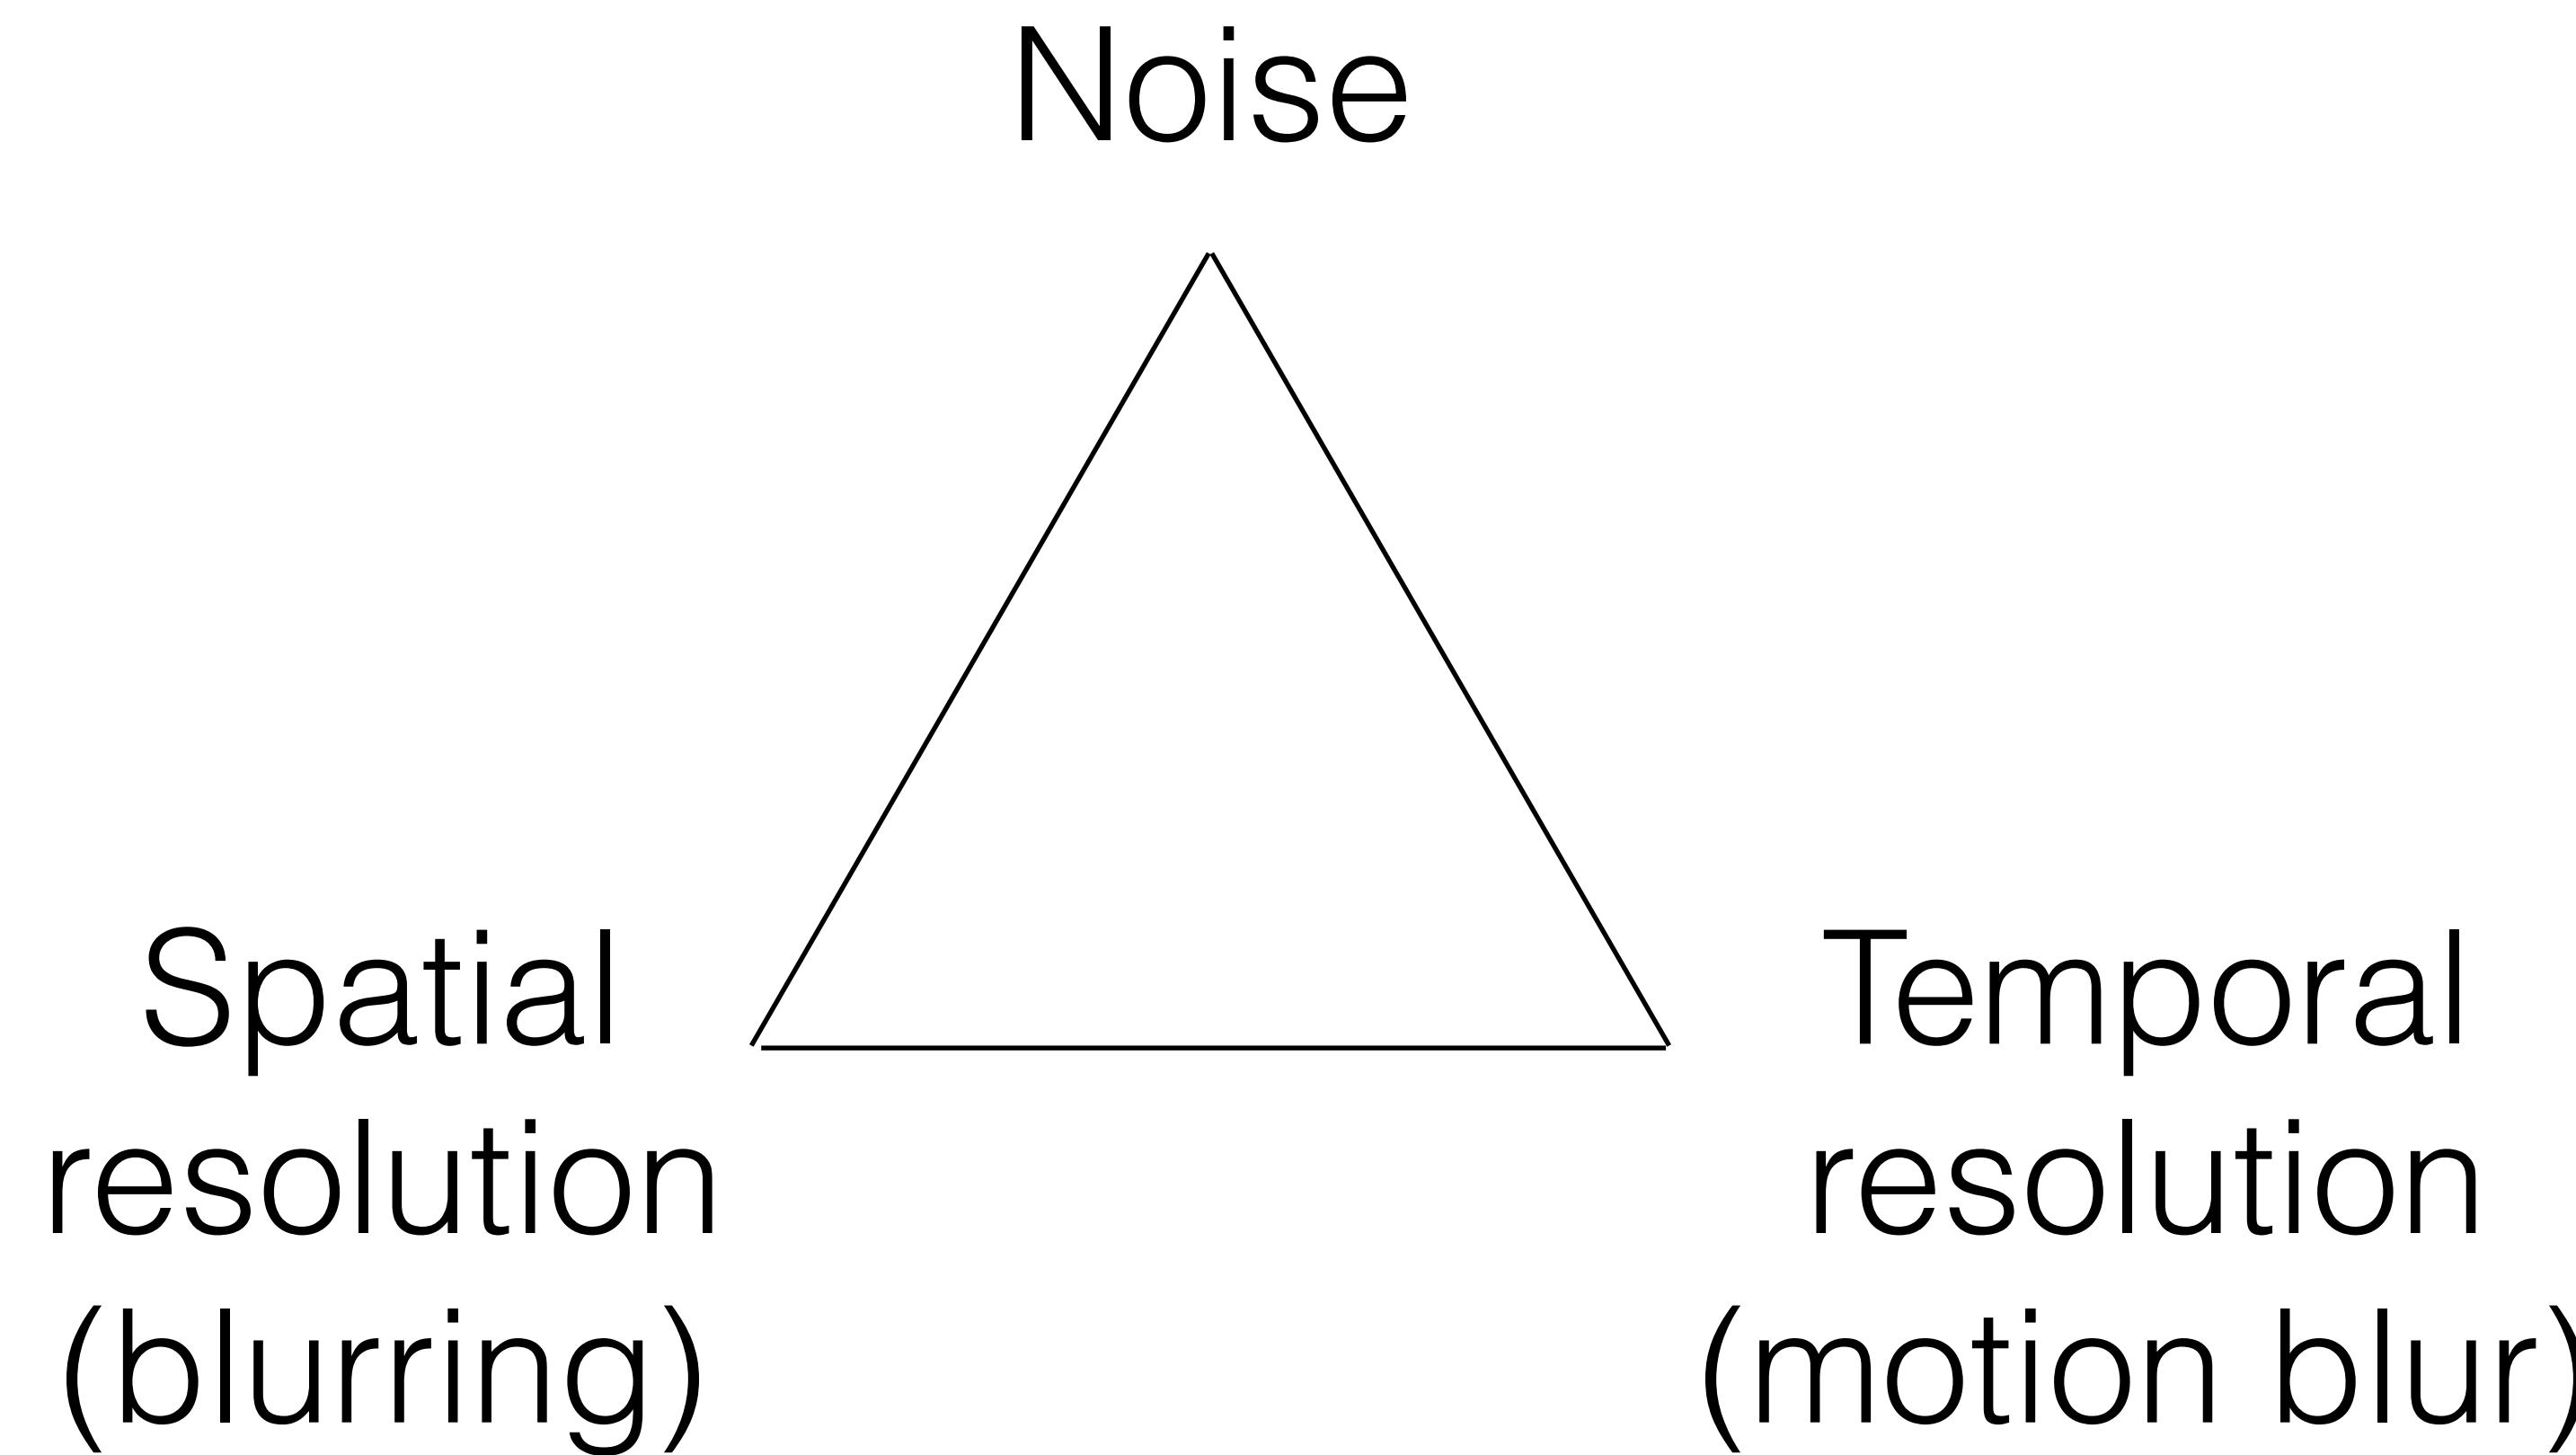
Shot Noise



Reducing Shot Noise

- The only way to reduce quantum noise is to *collect more light*
 - Turn up the source
 - Larger aperture
 - Collect for longer
- What are the tradeoffs?

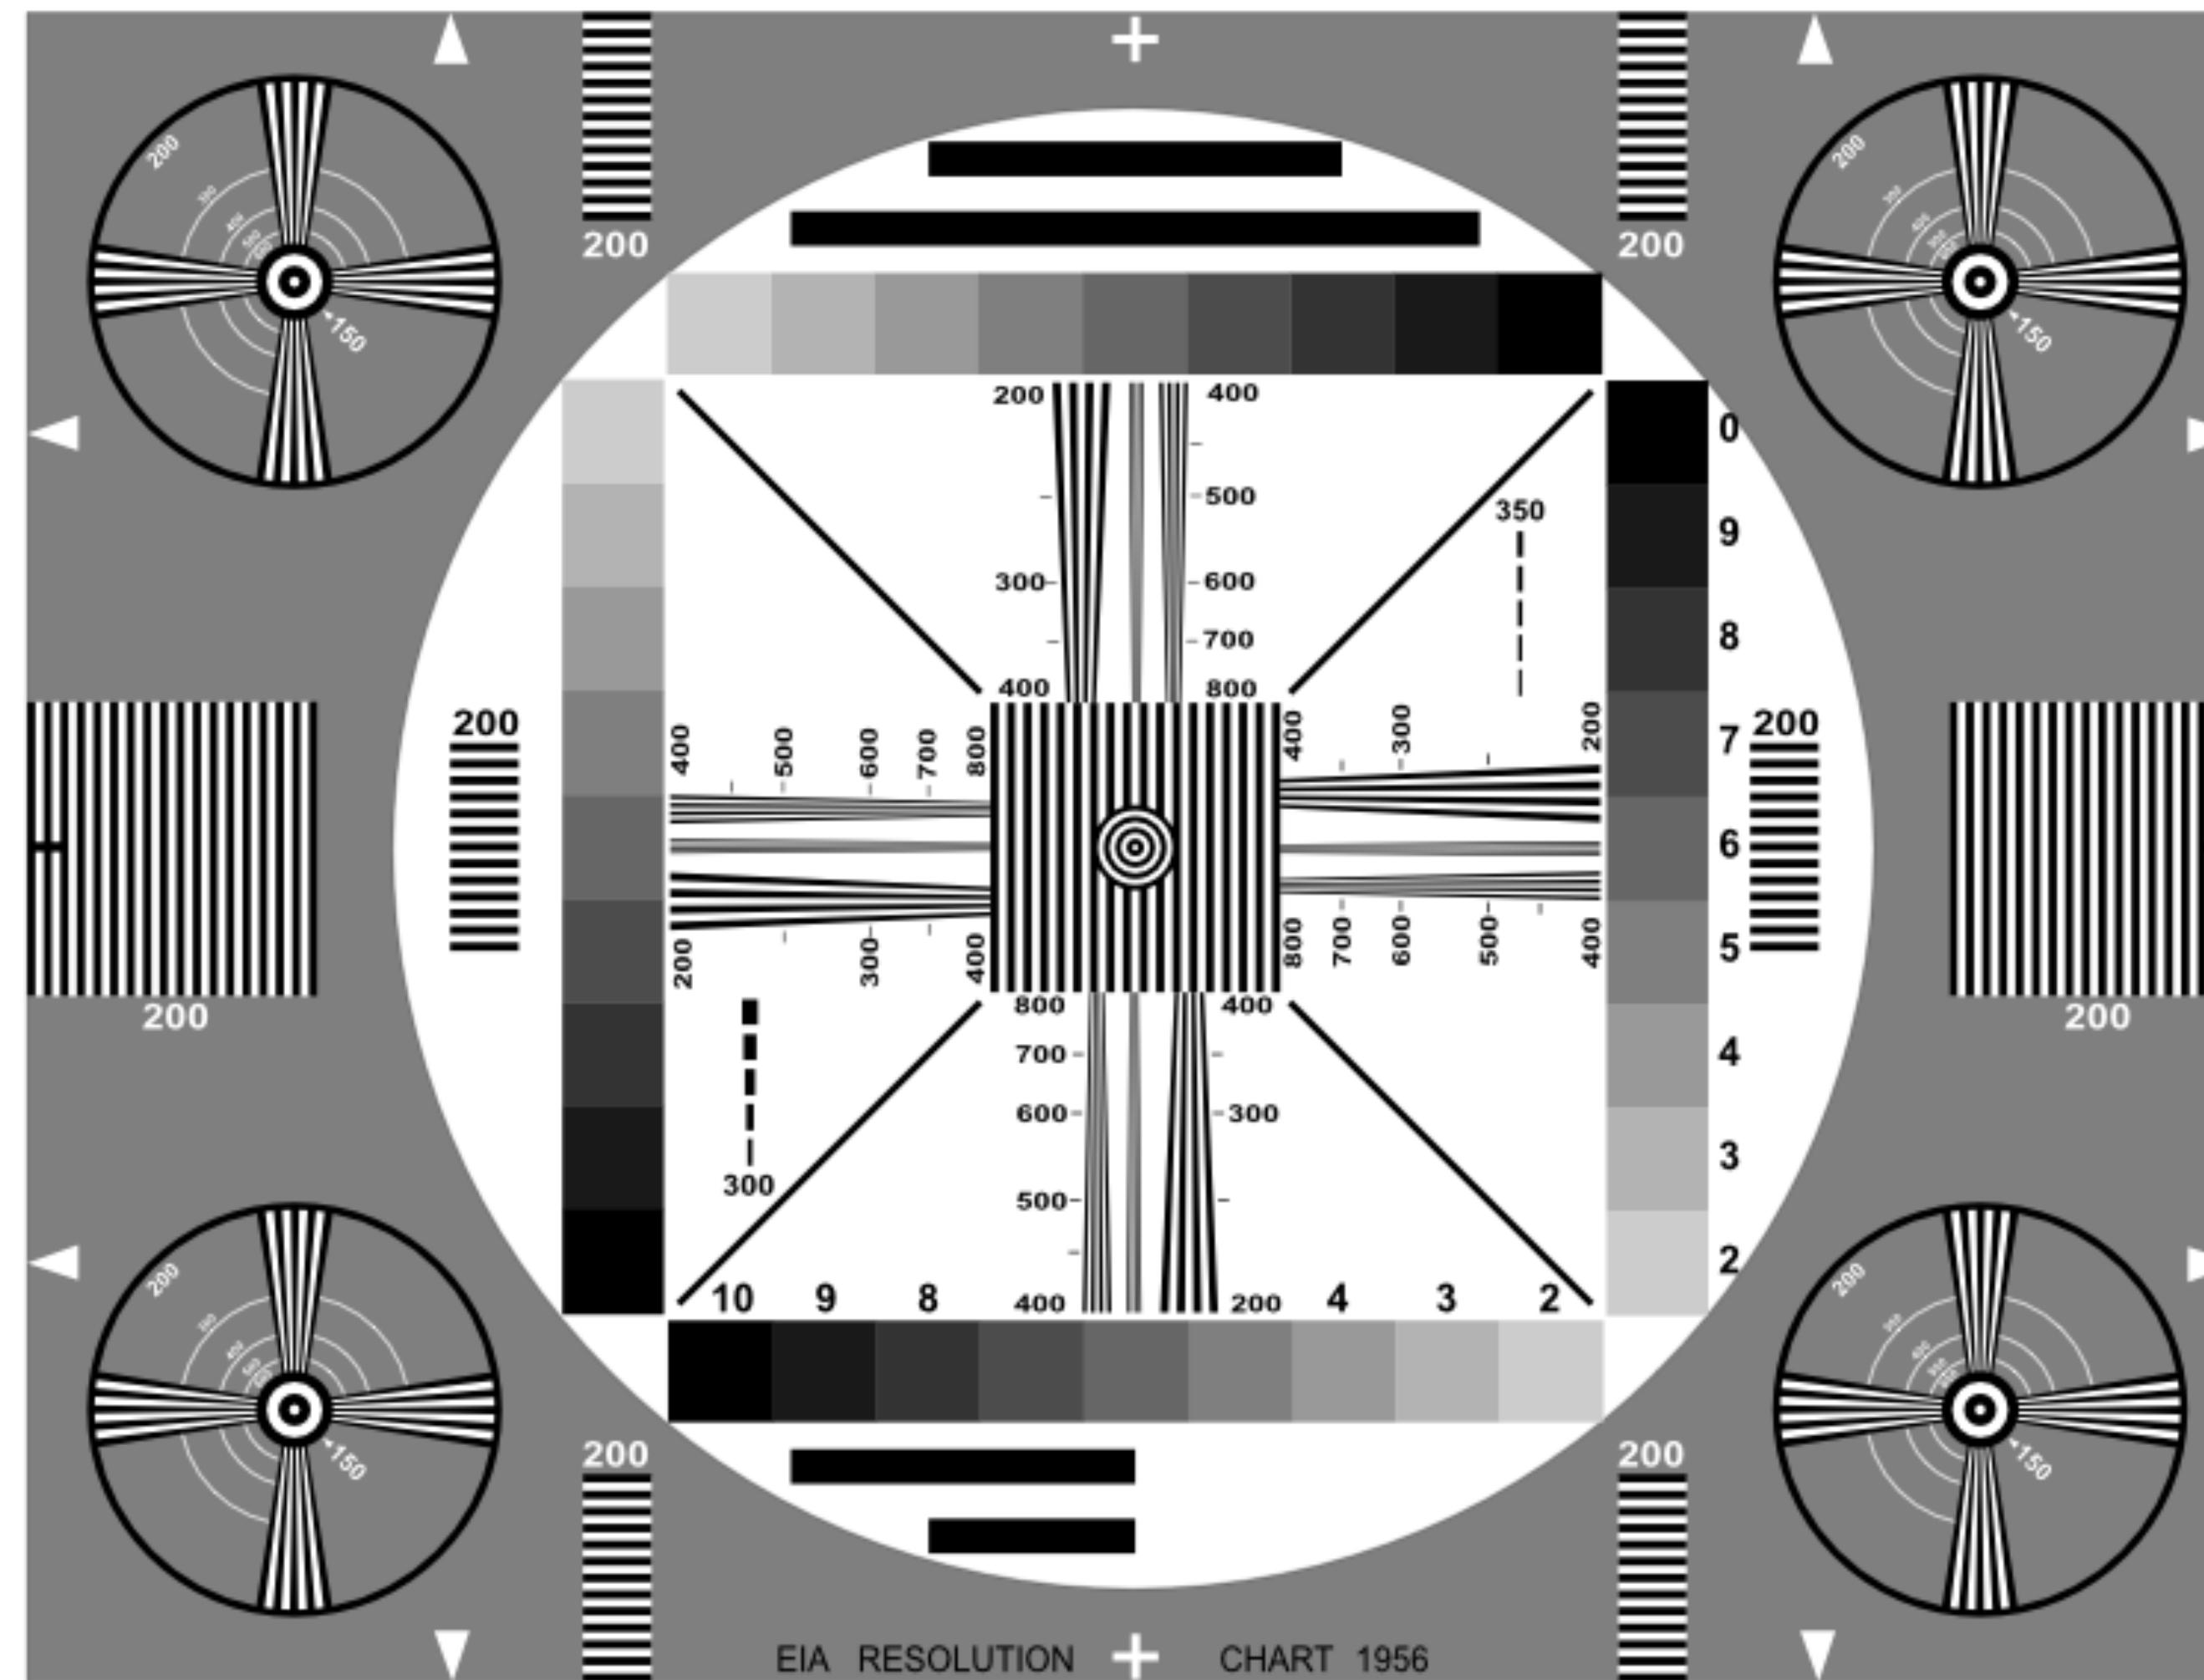
Fundamental Tradeoff



Measuring Resolution

- One way is to use alternating black/white lines with fixed spacing
 - Increase the density until you can't see the separate lines
 - Gradually blurs to grey
 - Stop when half the original contrast
 - Units: line pairs per millimeter

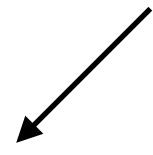
Measuring Resolution



Sampling Revisited

- How much sampling is enough?
 - *Shannon Sampling Theorem:* twice the highest frequency in the signal (in theory)
 - *Nyquist rate*
- What happens if you sample above this?
 - Avoids dangers of theoretical limits
 - Better for intermediate processing
- What happens if you don't sample enough?
 - Aliasing (false low-frequencies components appear)
 - In images this causes *Moiré patterns*
- *Insufficient sampling during acquisition introduces flaws that cannot be corrected in later processing*

Different kind of aliasing
than jaggies, but related



Moiré Patterns



Camera Problems

- Noise
- Spatial blur
- Motion blur
- Bayer sampling artifacts
- Lens distortion
- Chromatic aberration
- Brightness
- Contrast
- Color balance
- Tone mapping
(color responses)

Coming up...

- Interpolation of discrete samples



Interpolation

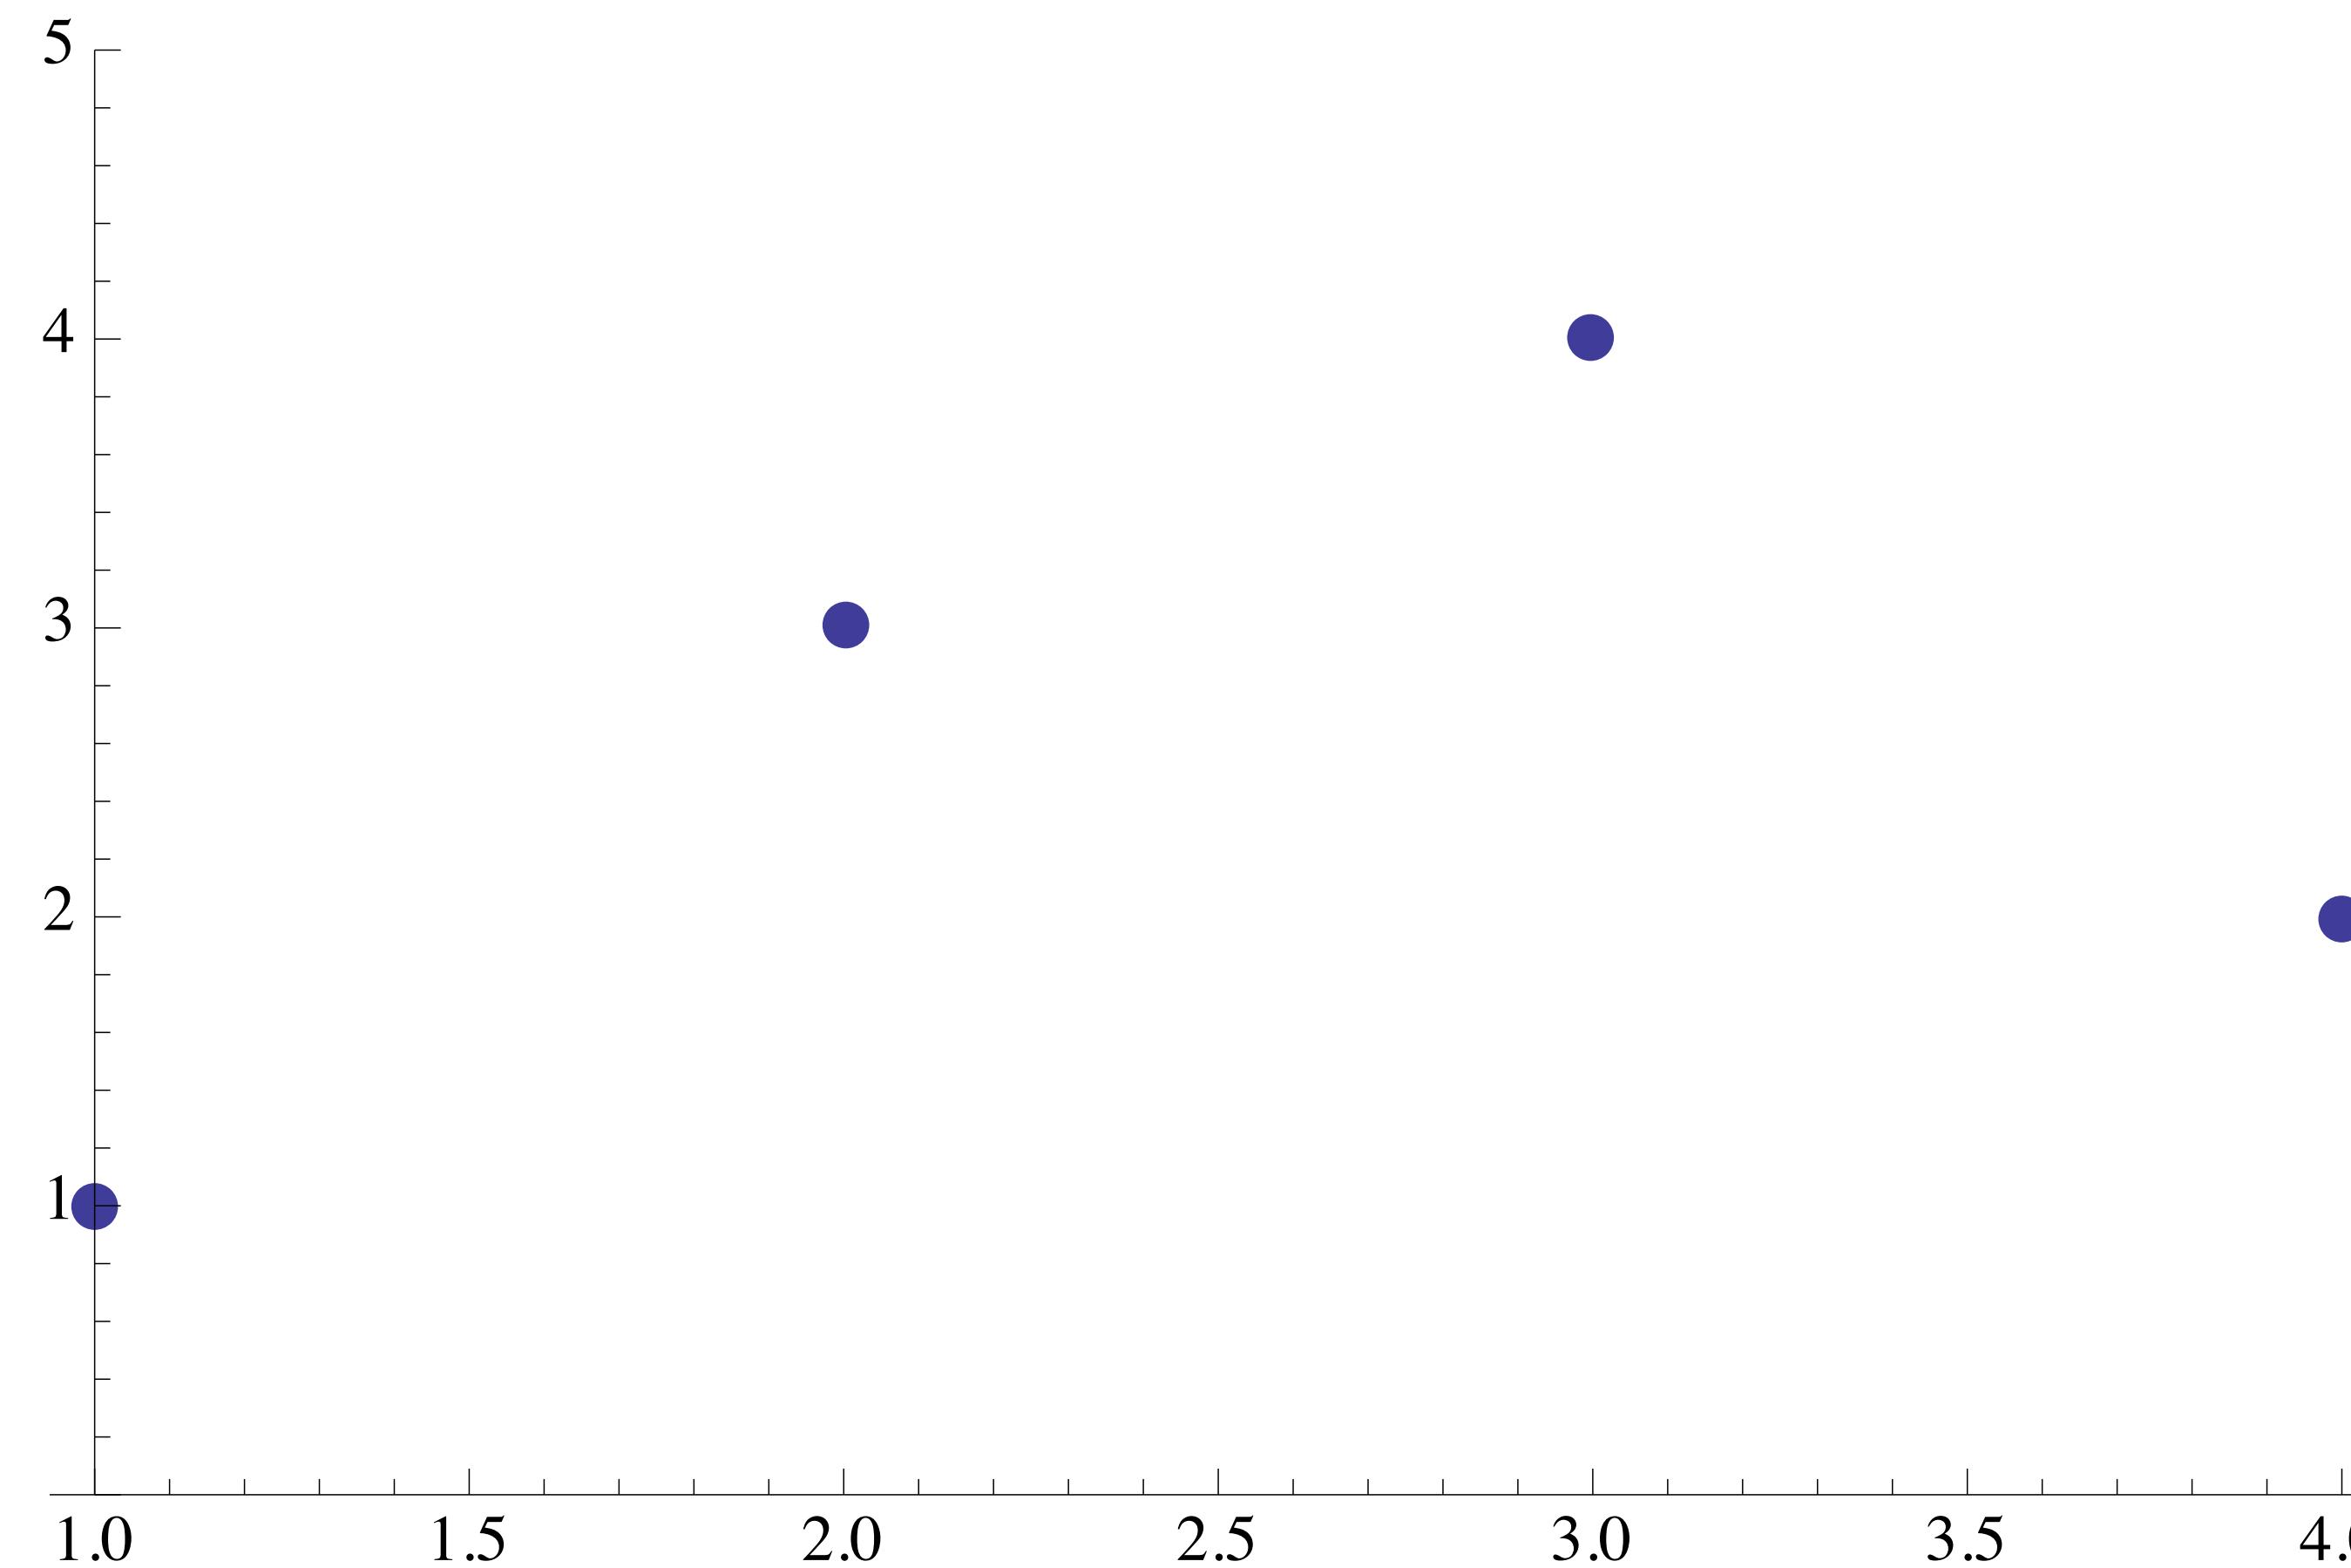
CS 355: Introduction to Graphics and Image Processing

Interpolation

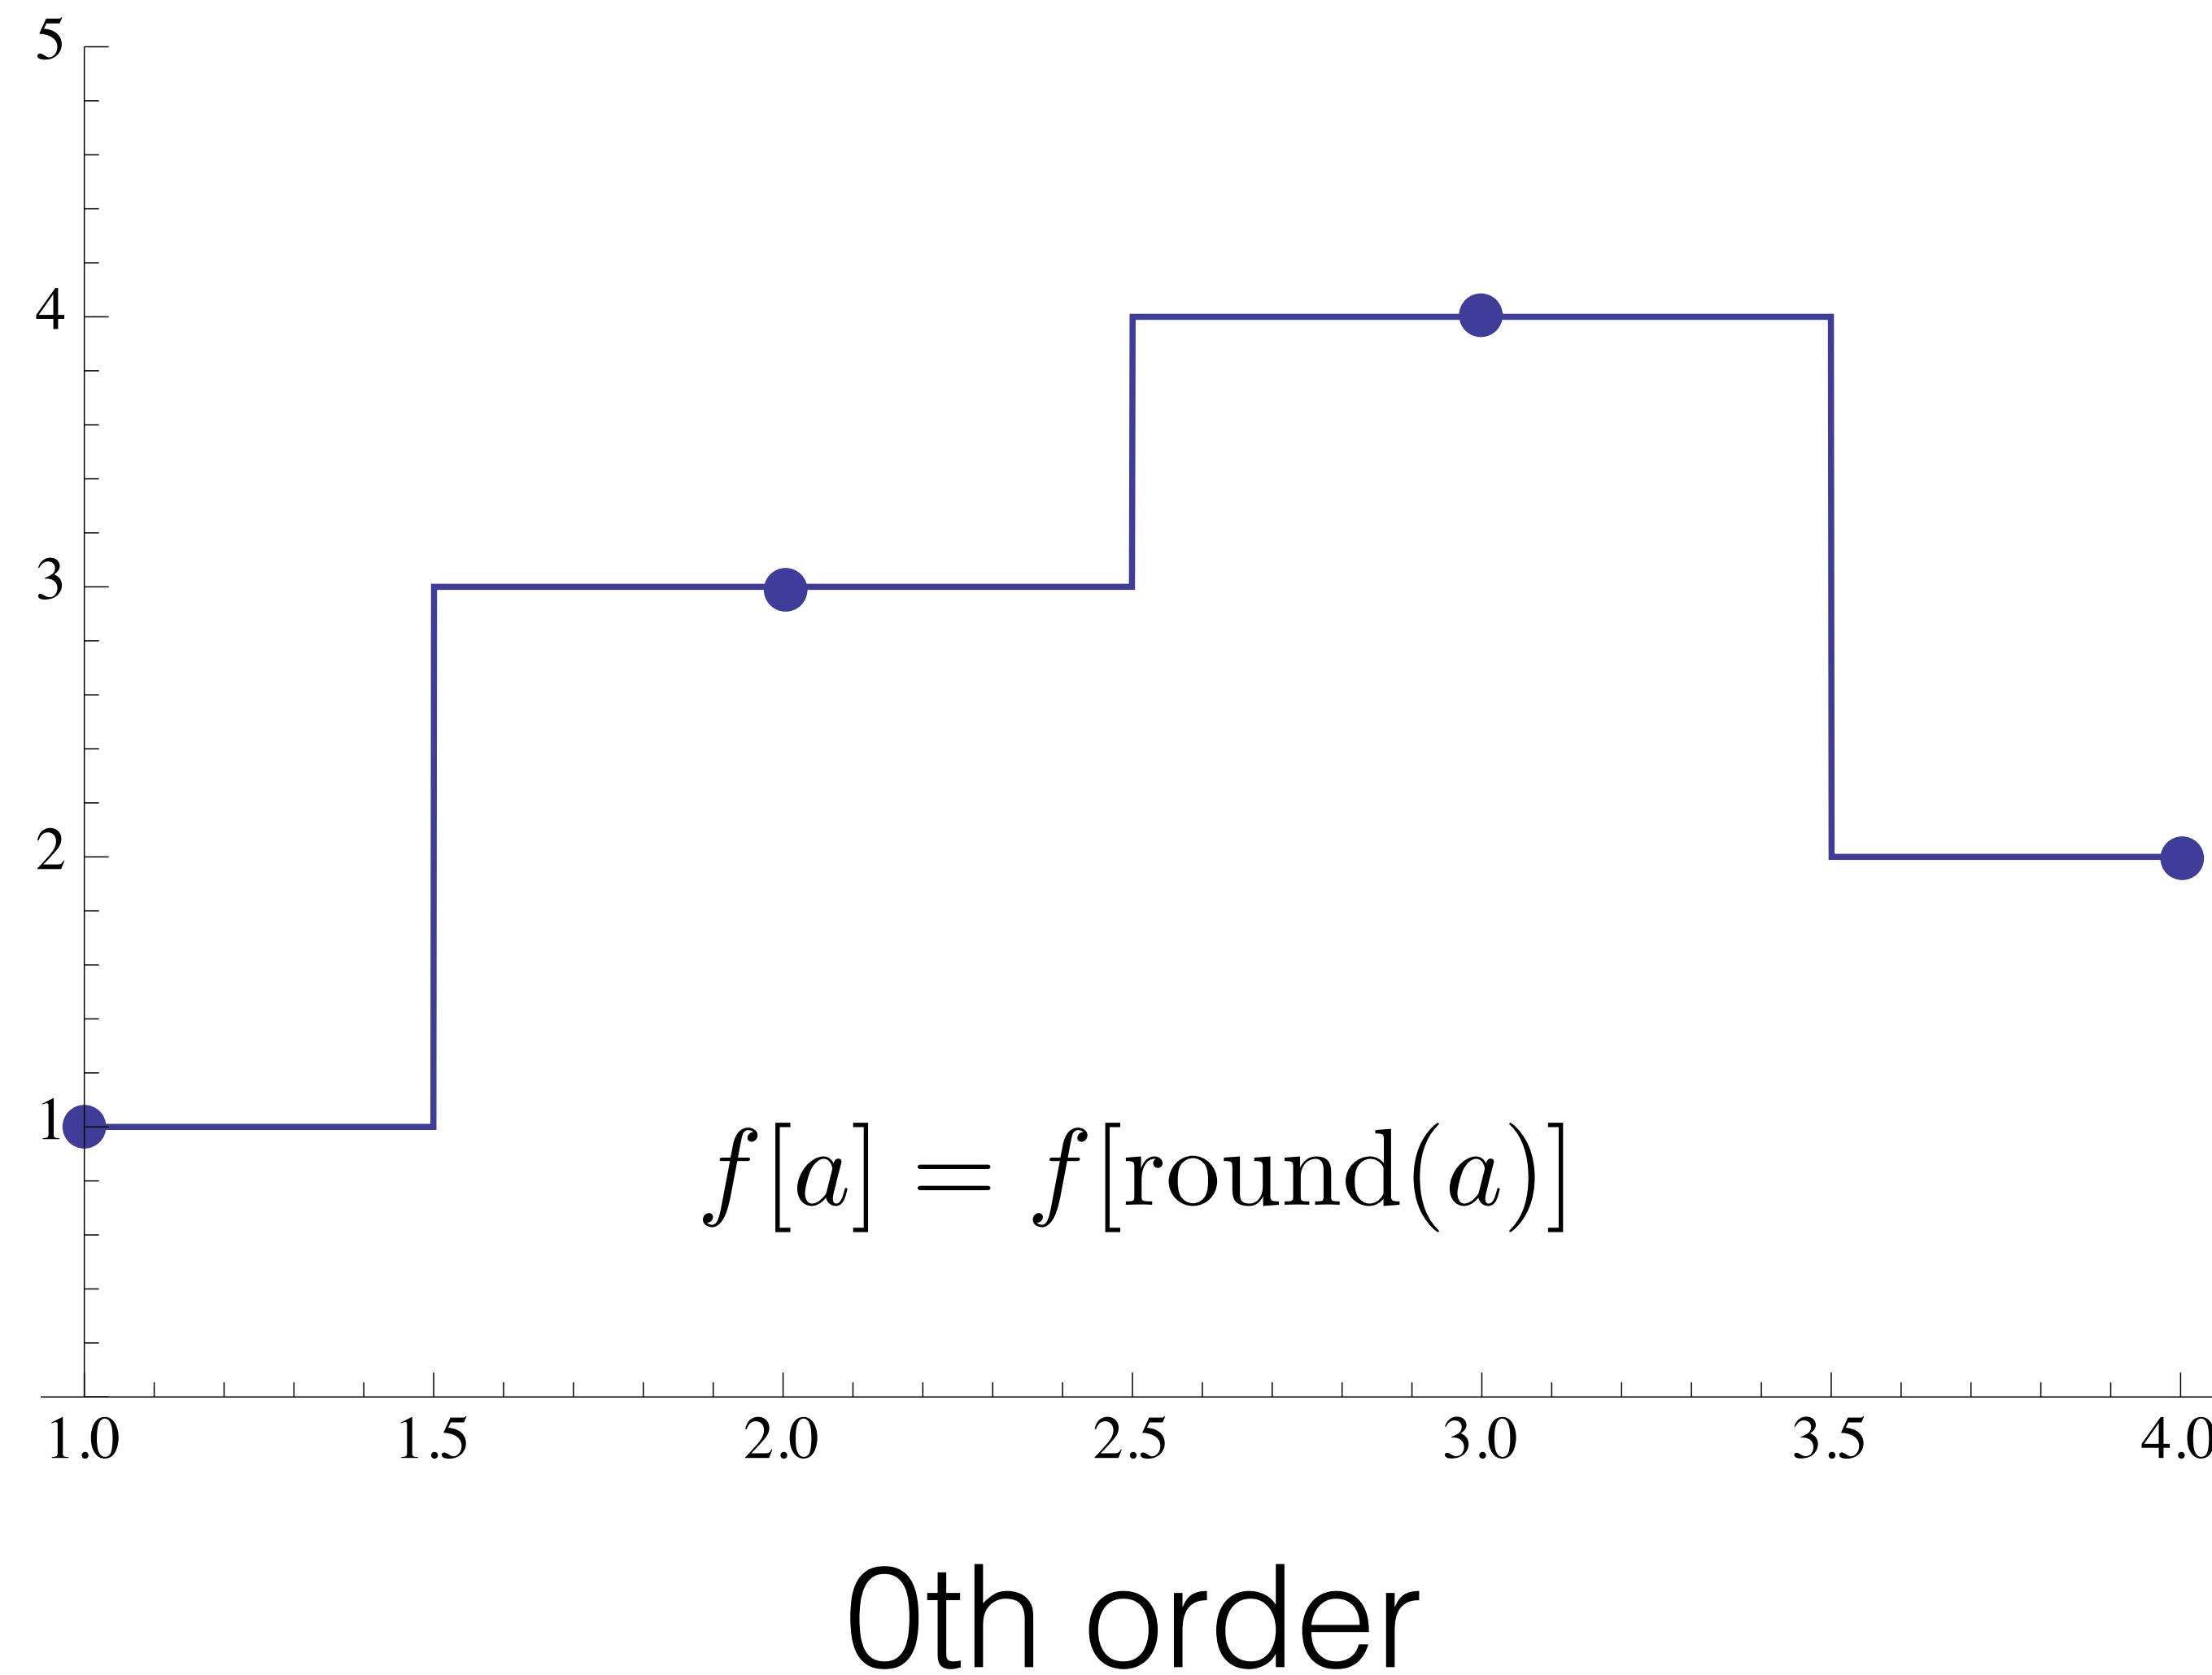
- Used to estimate a function between known sampled values
- Digital signals to analog
 - Audio playback
 - Image display
- Spatial image manipulation
 - Changing size
 - Rotation
 - Warping
- Generate smooth geometric models between defined points
(we'll come back to this later)



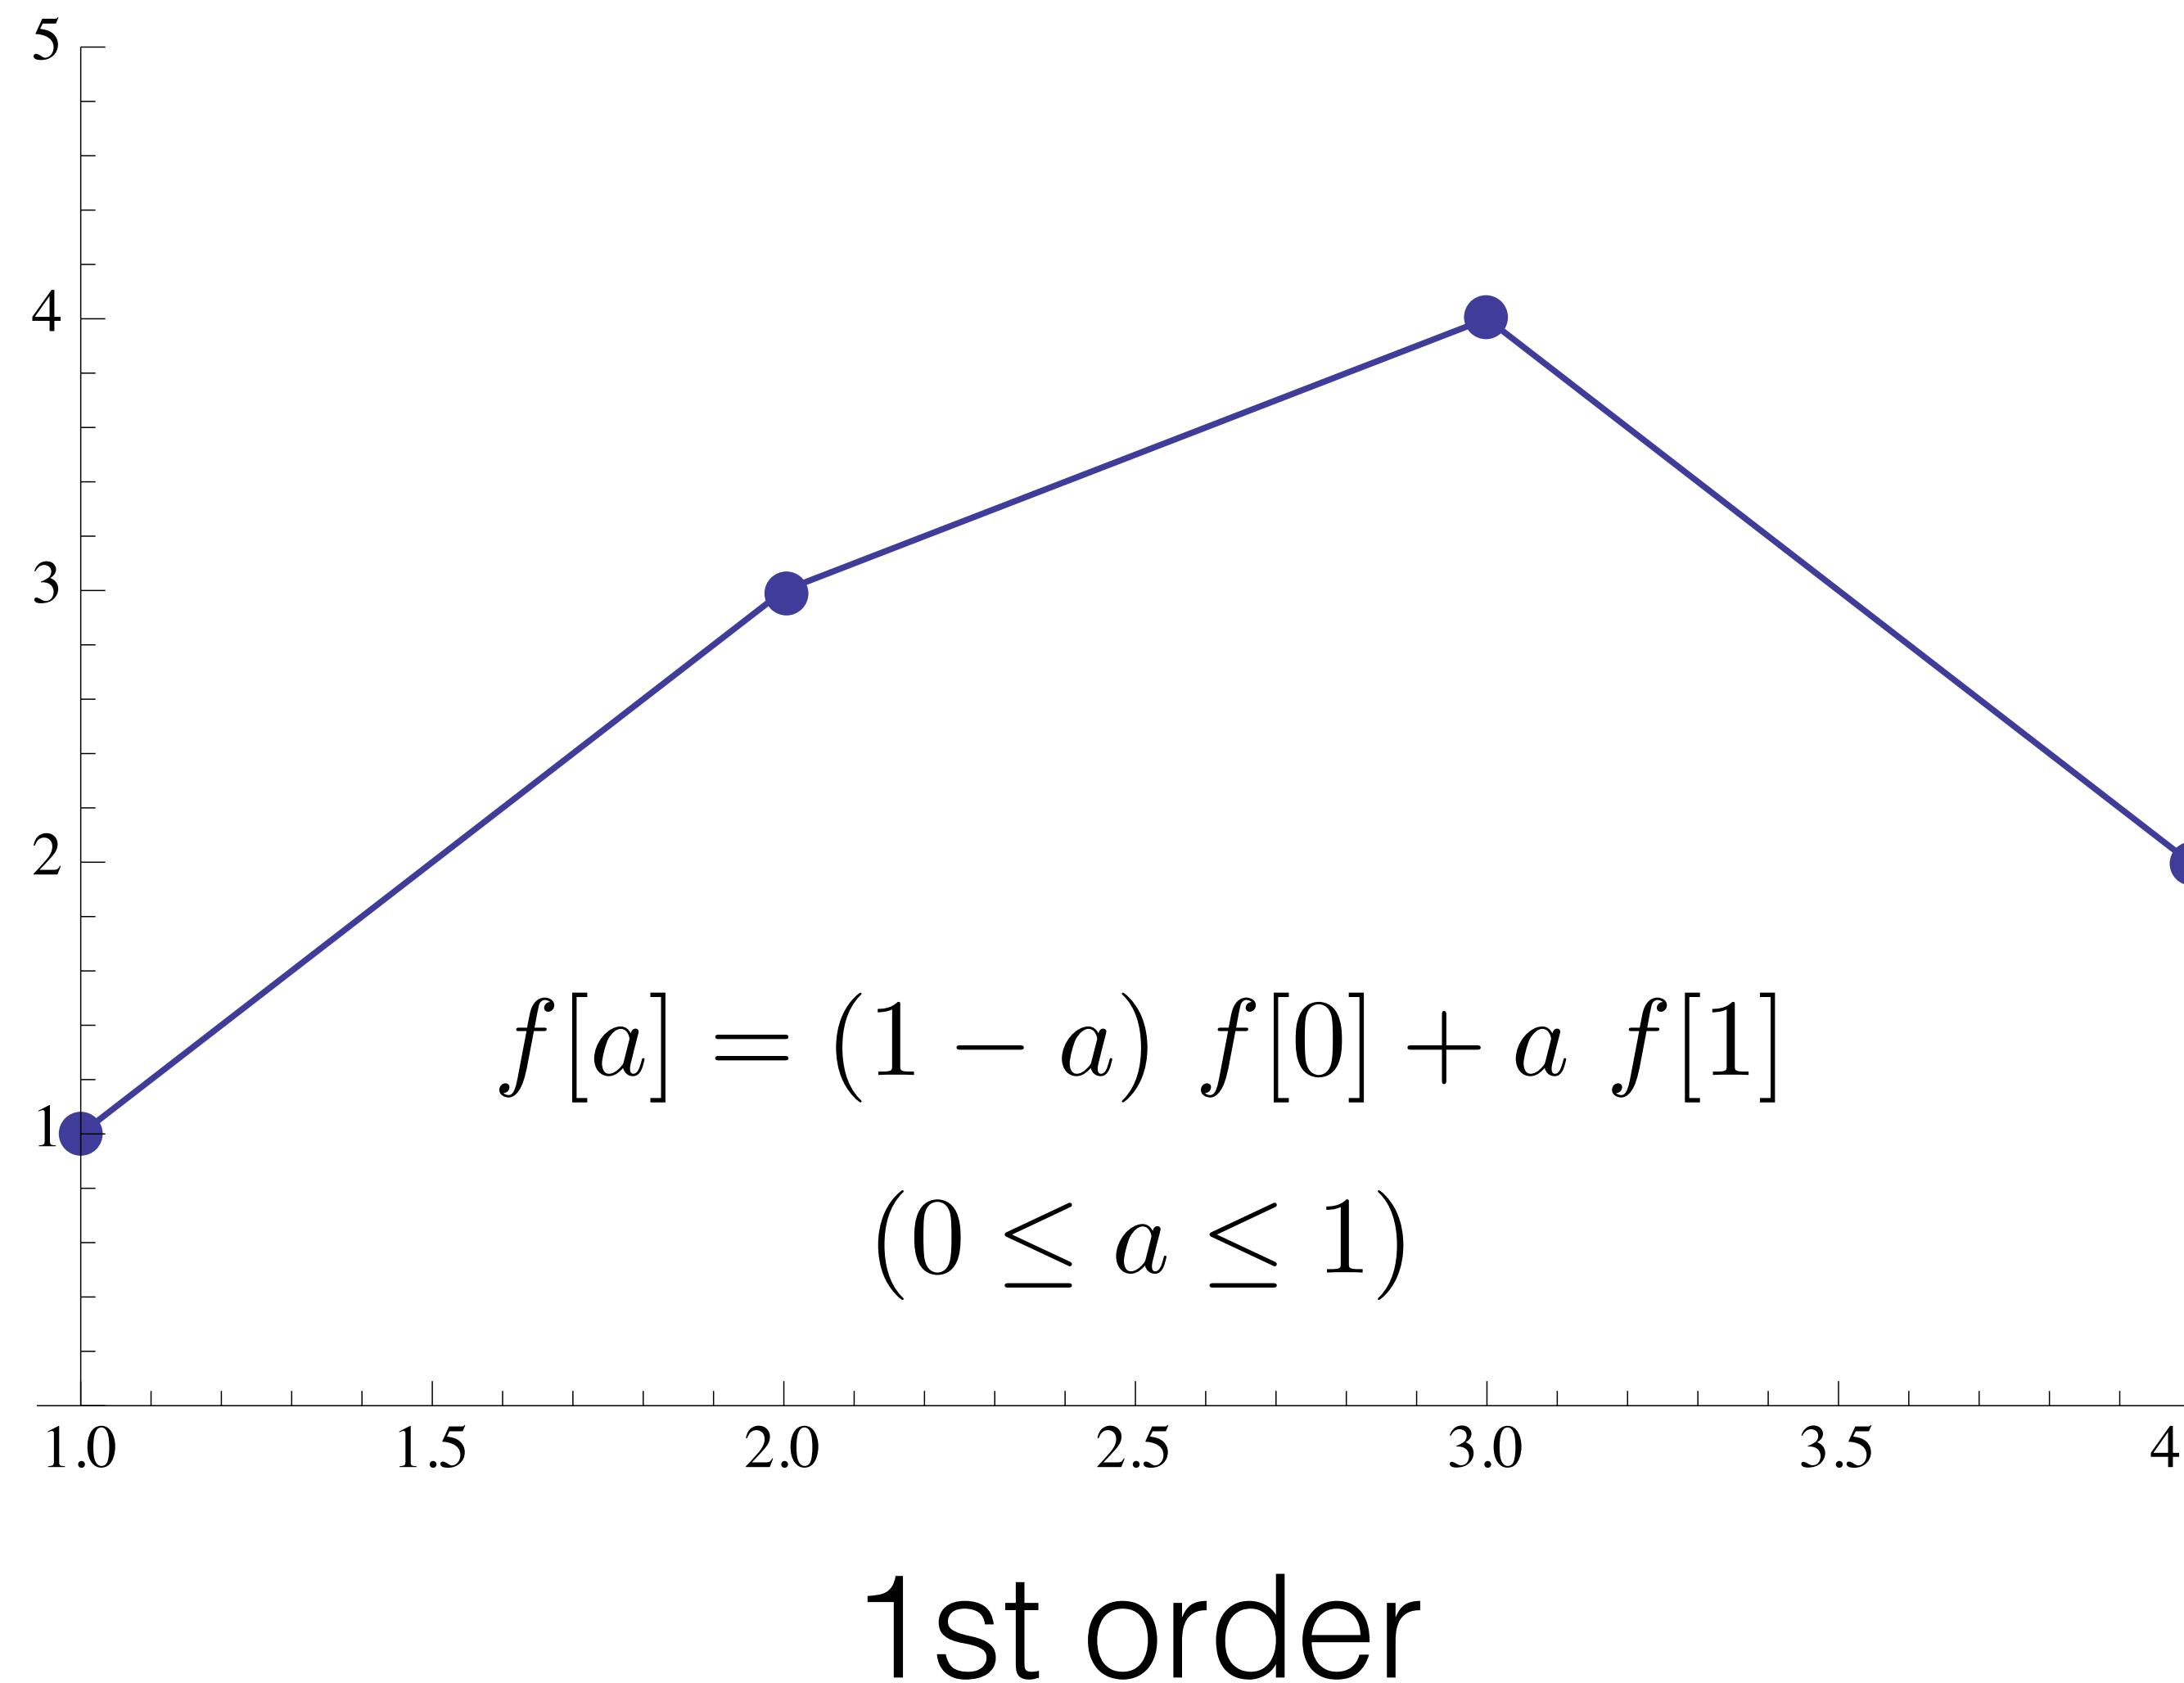
Interpolation



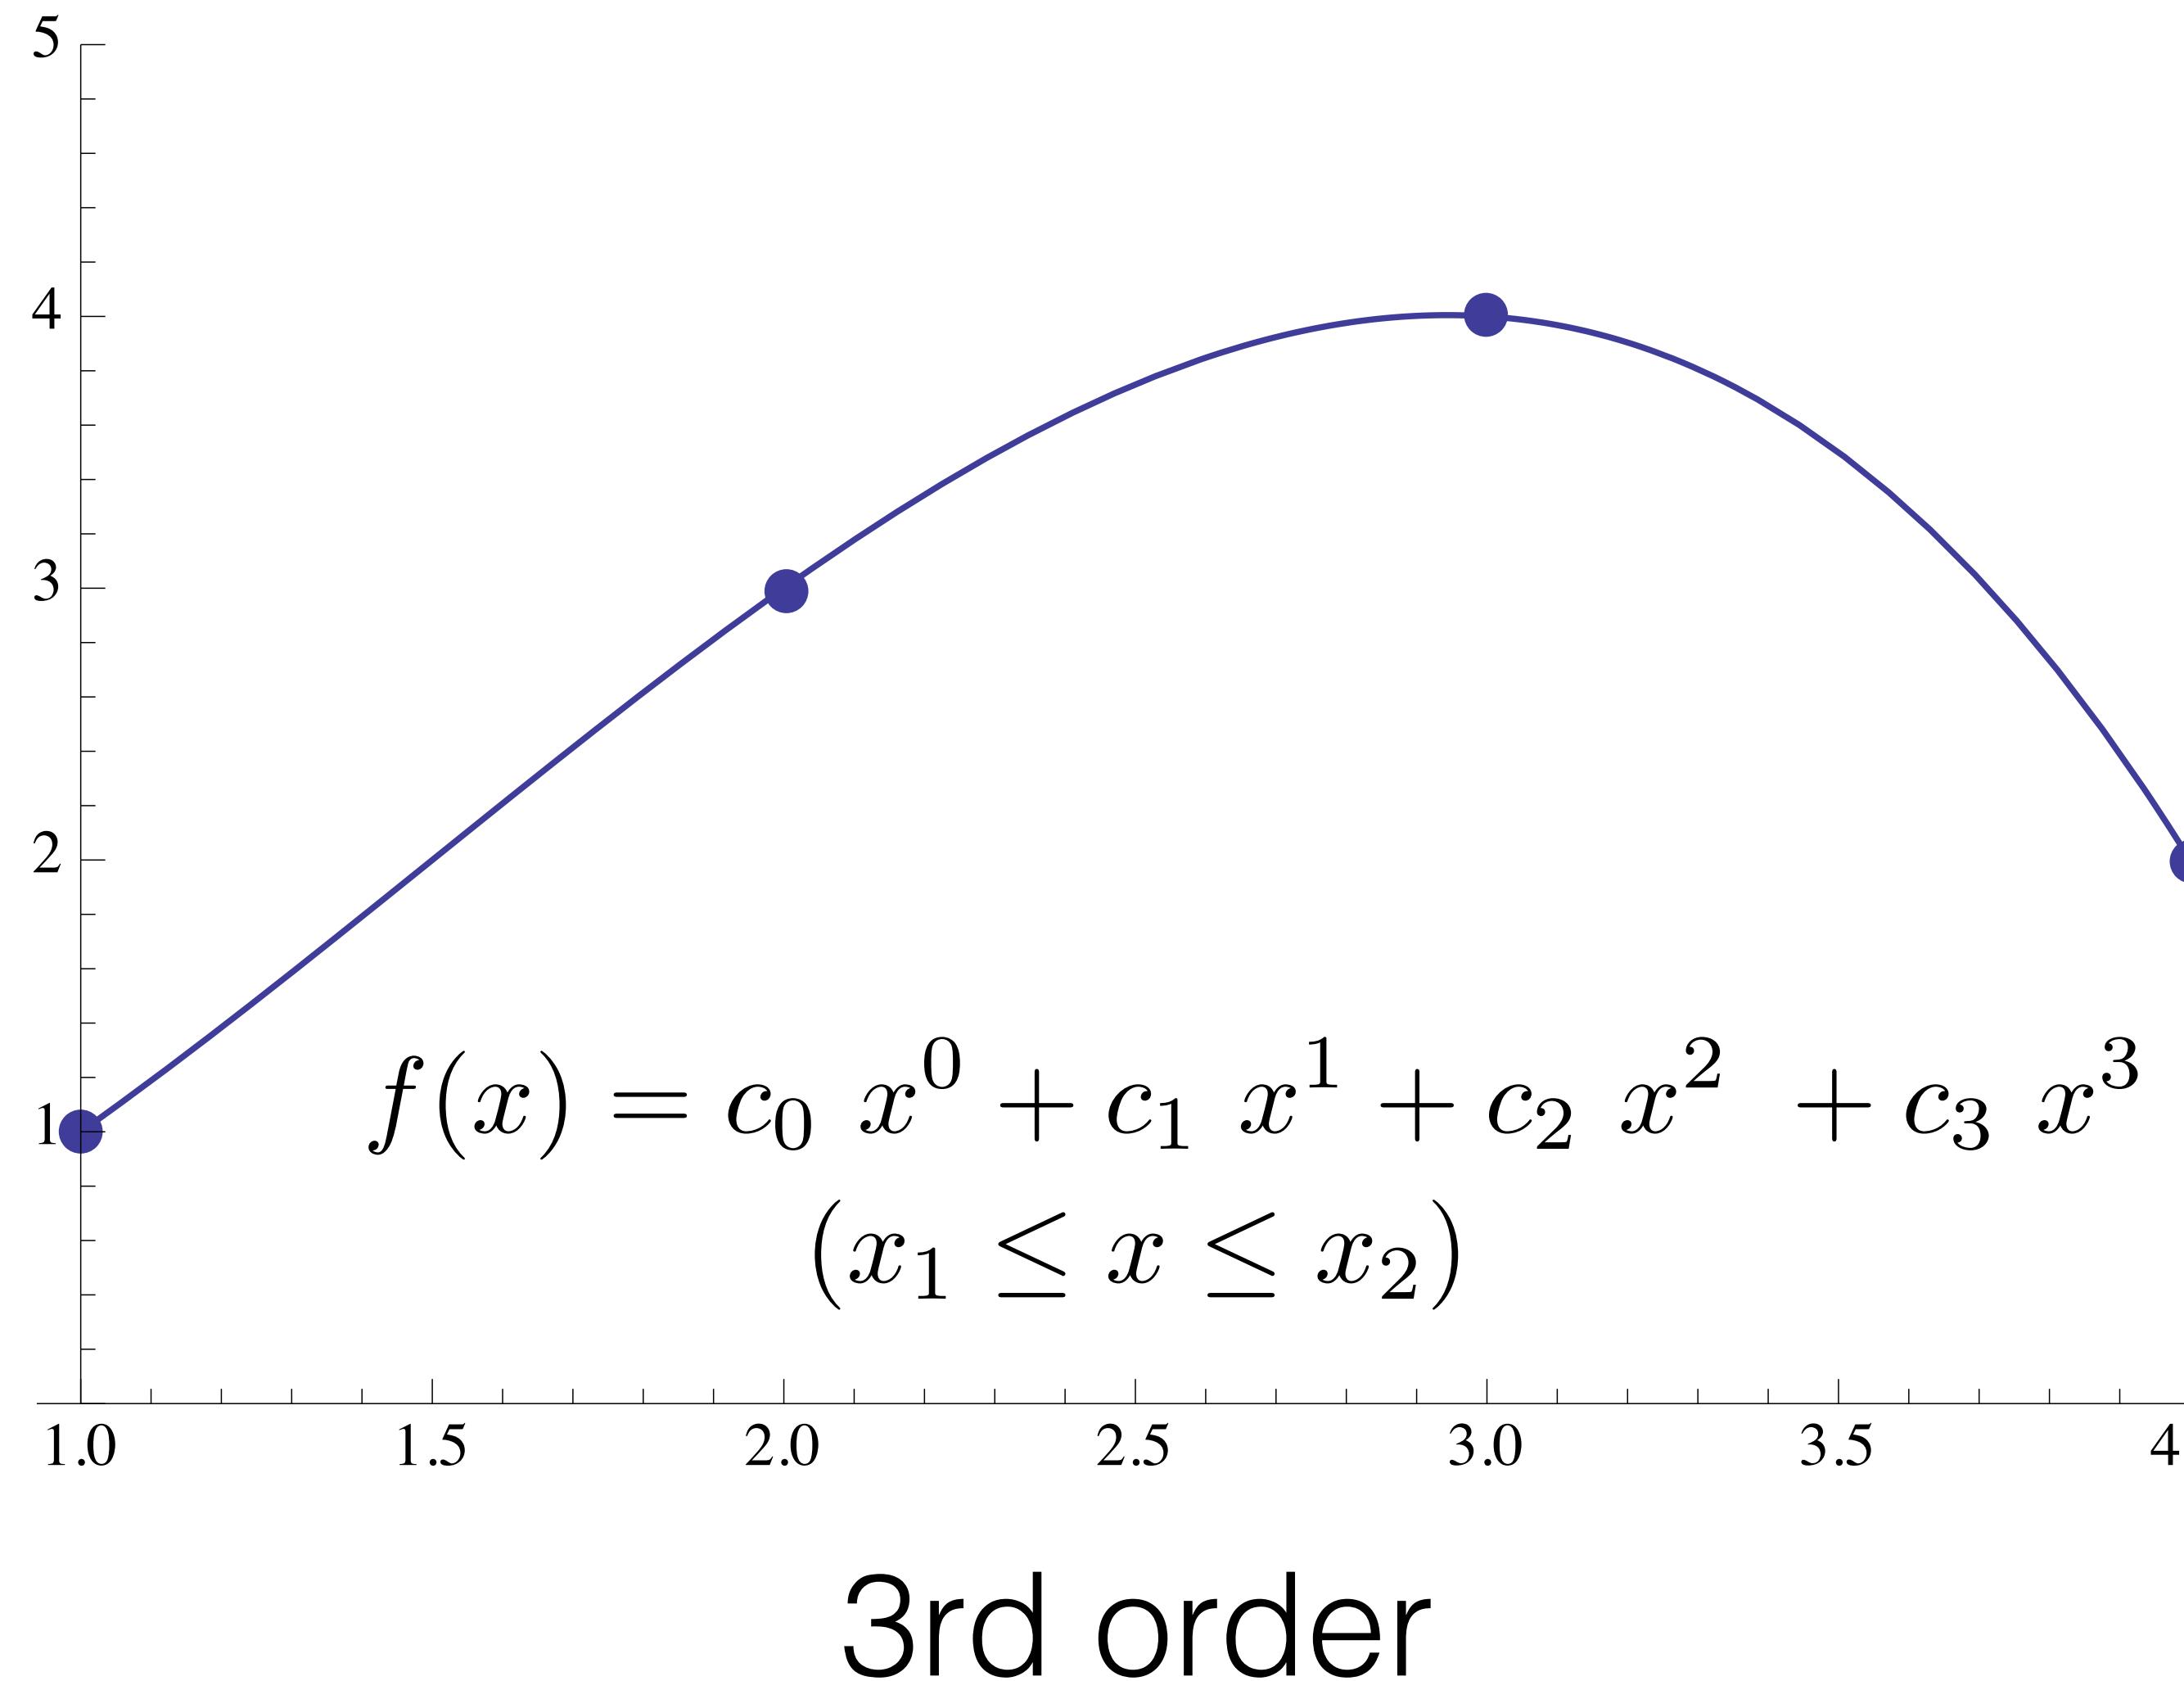
Nearest Neighbor



Linear Interpolation



Cubic Interpolation



Cubic Interpolation

$$f(x) = ax^3 + bx^2 + cx + d \quad (x_1 \leq x \leq x_2)$$

$$f(x_0) = a x_0^3 + b x_0^2 + c x_0 + d$$

$$f(x_1) = a x_1^3 + b x_1^2 + c x_1 + d$$

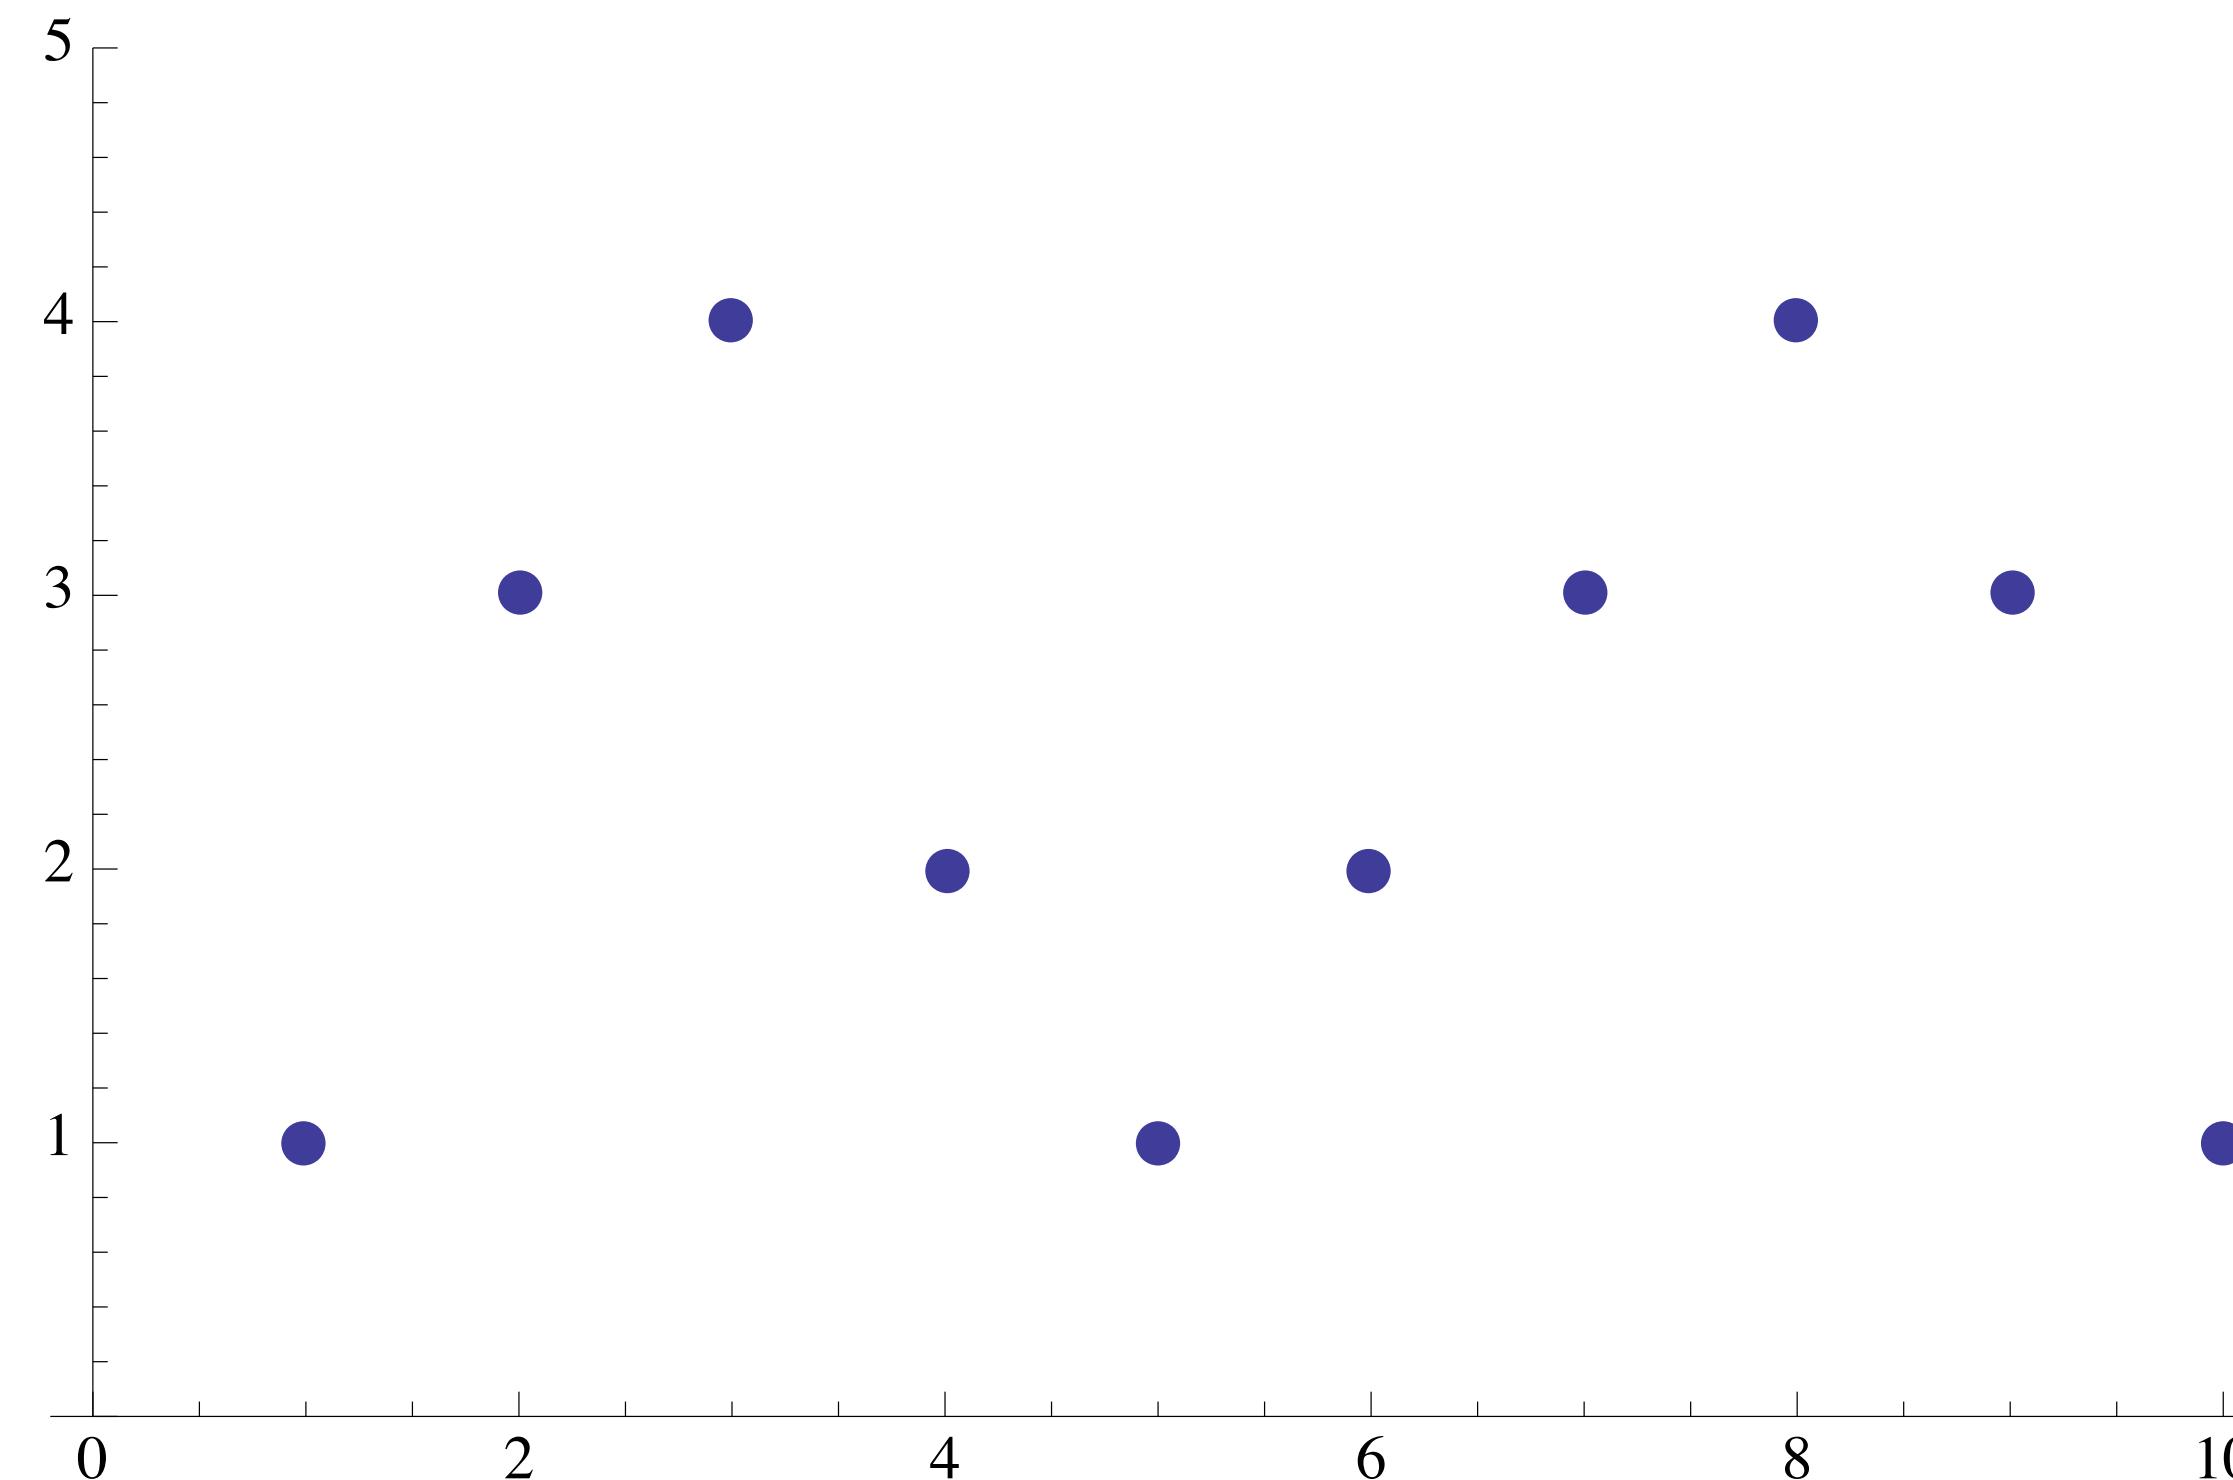
$$f(x_2) = a x_2^3 + b x_2^2 + c x_2 + d$$

$$f(x_3) = a x_3^3 + b x_3^2 + c x_3 + d$$

$$\begin{bmatrix} x_0^0 & x_0^1 & x_0^2 & x_0^3 \\ x_1^0 & x_1^1 & x_1^2 & x_1^3 \\ x_2^0 & x_2^1 & x_2^2 & x_2^3 \\ x_3^0 & x_3^1 & x_3^2 & x_3^3 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \end{bmatrix} = \begin{bmatrix} f(x_0) \\ f(x_1) \\ f(x_2) \\ f(x_3) \end{bmatrix}$$

Solve for
a, b, c, d
and plug
into function

Longer Sequences

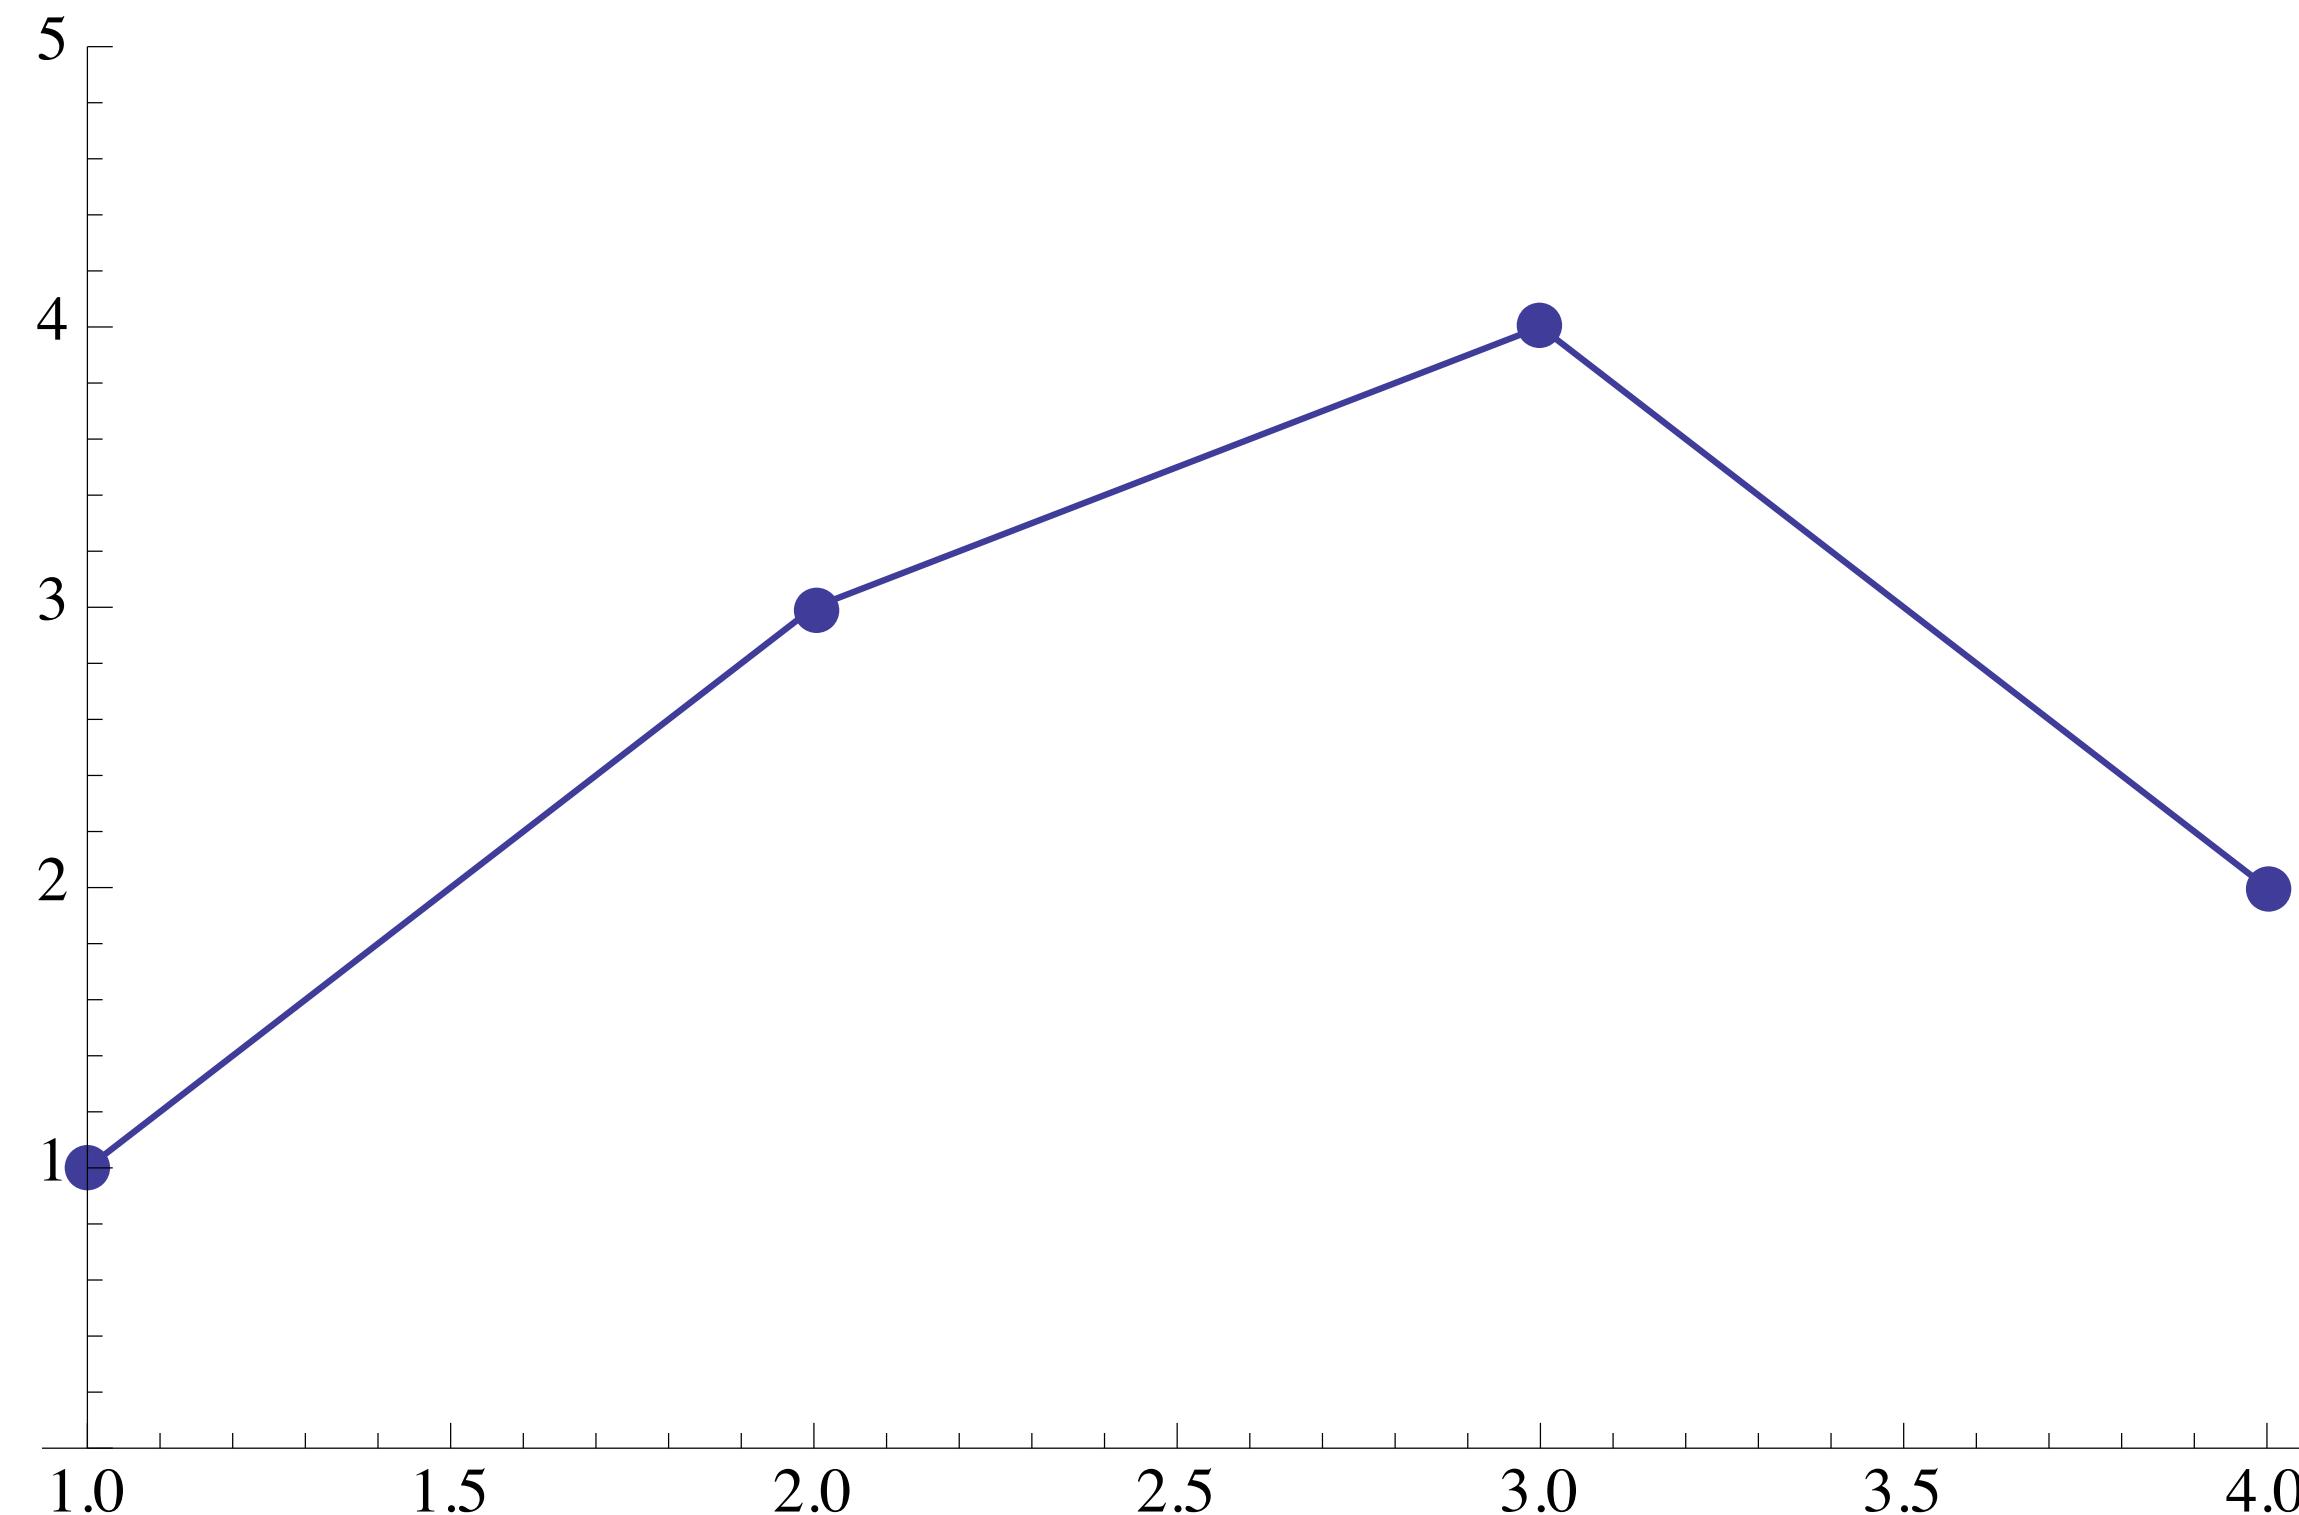


0th - nearest neighbor

1st - use one sample on both sides

3rd - use two samples on both sides

Continuity



0th order: function is continuous

1st order: slope is continuous

nth order: nth derivative is continuous

2D Images

- Extend ideas of single-value interpolation
 - Nearest neighbor
 - Linear
 - Cubic
 - ...



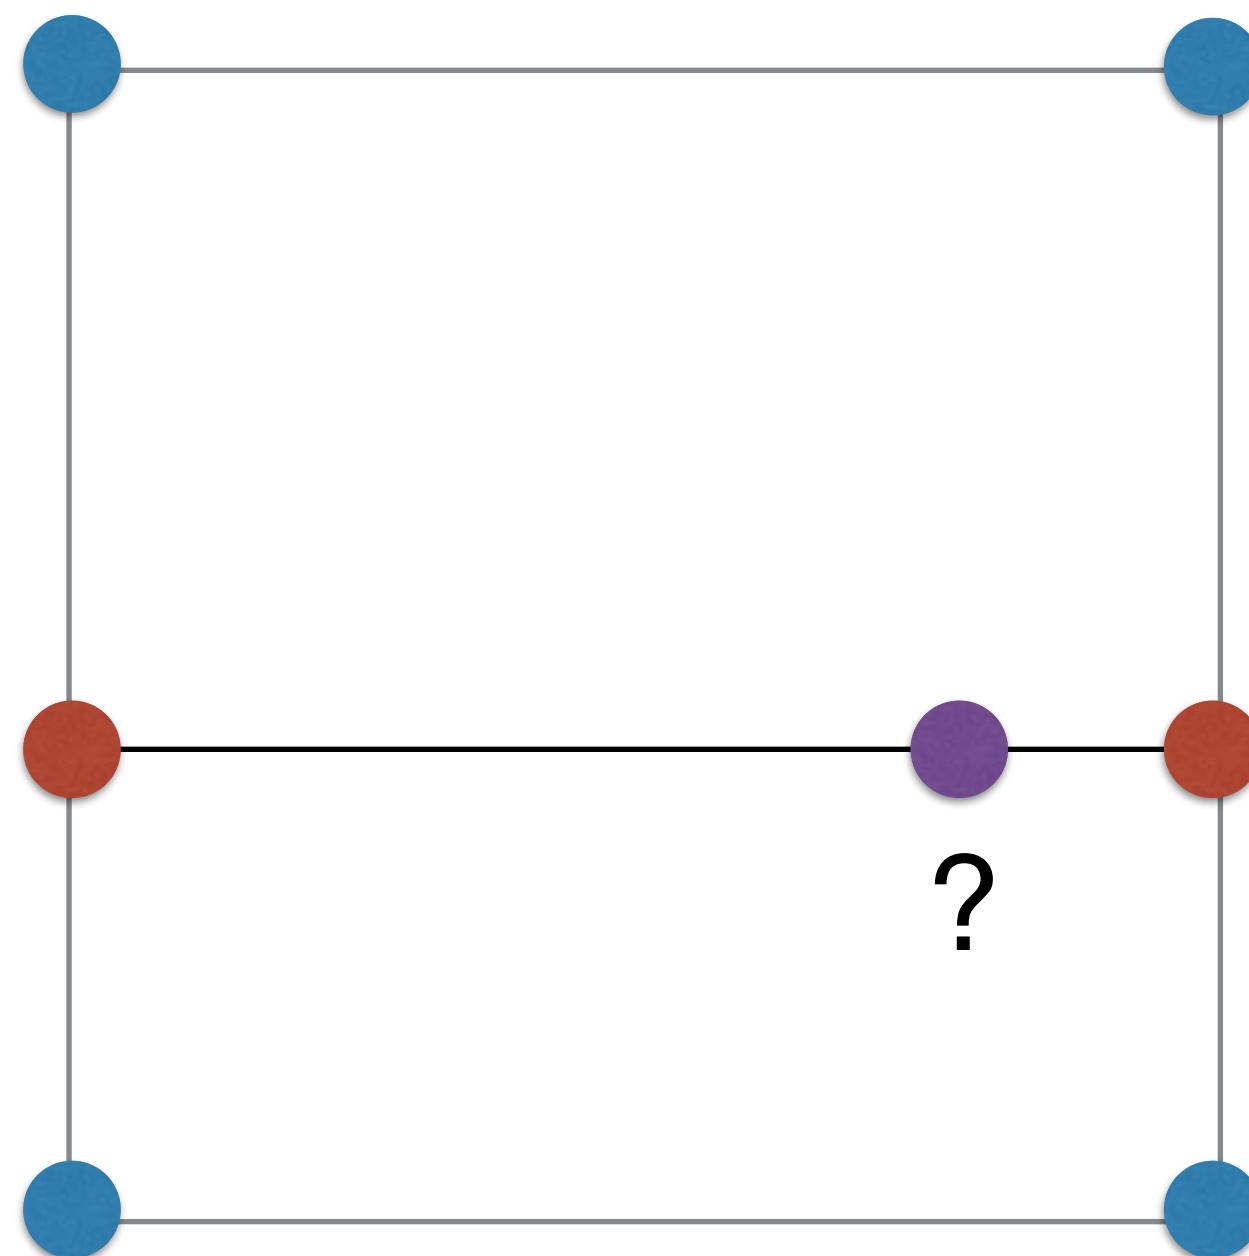
Nearest Neighbor

- Same idea as in 1-d:
Round off to nearest pixel
- Also called “pixel replication”
- Big blocky pixels



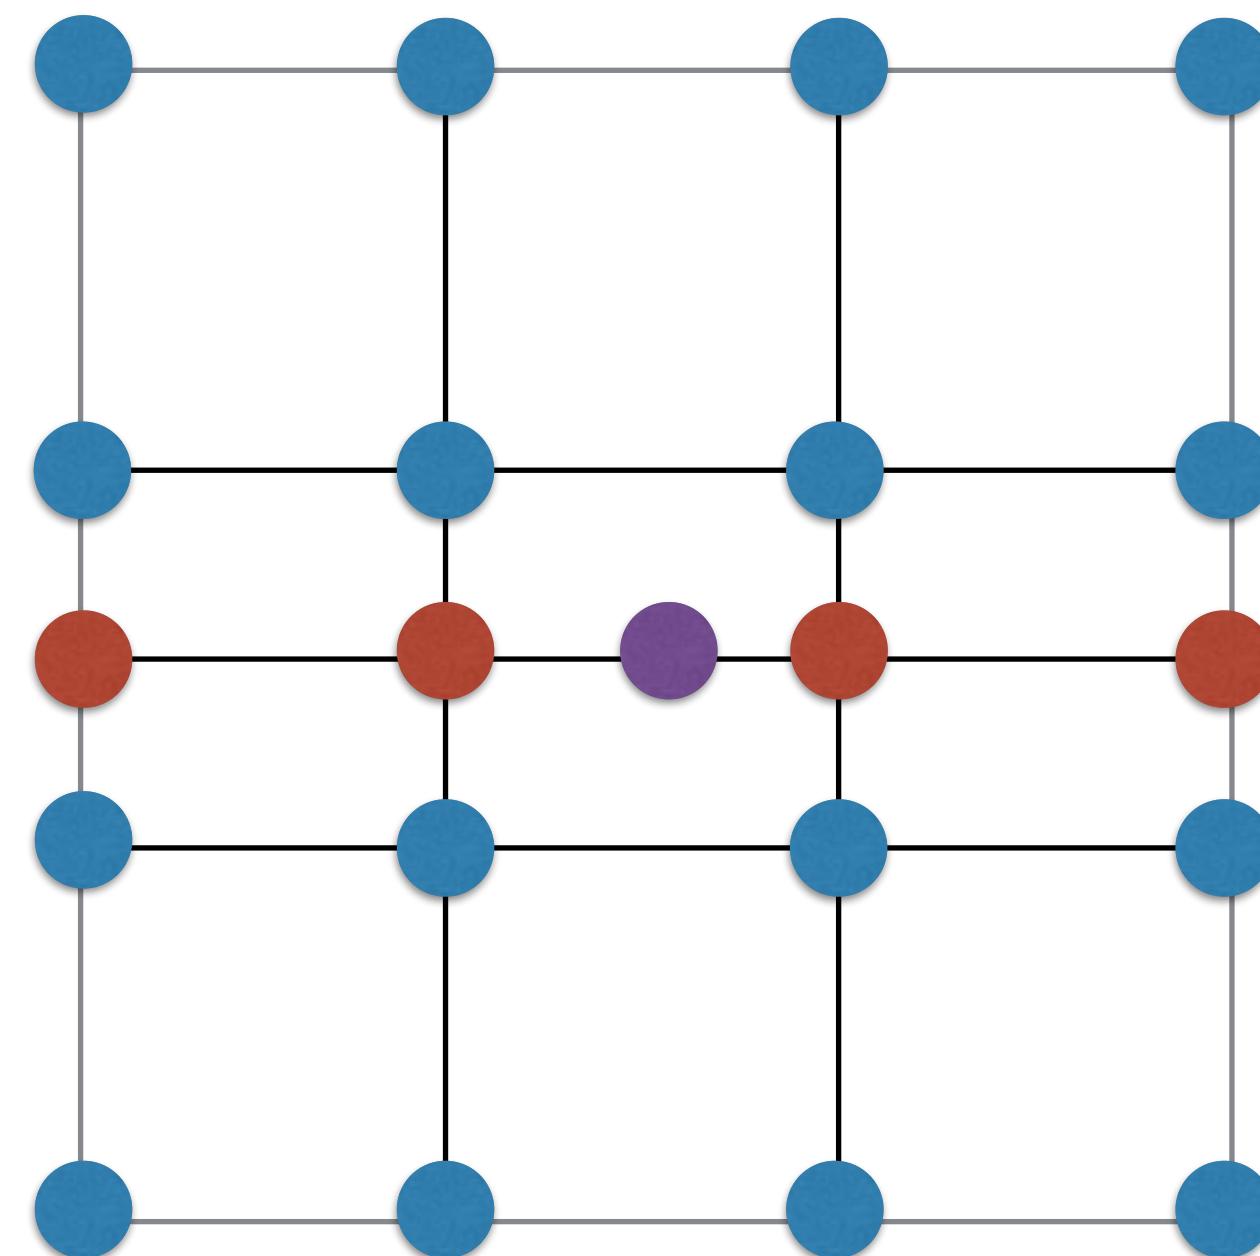
Bilinear

- Very common approach:
 - Interpolate vertically
 - Interpolate horizontally
- For linear interpolation, this is called bilinear



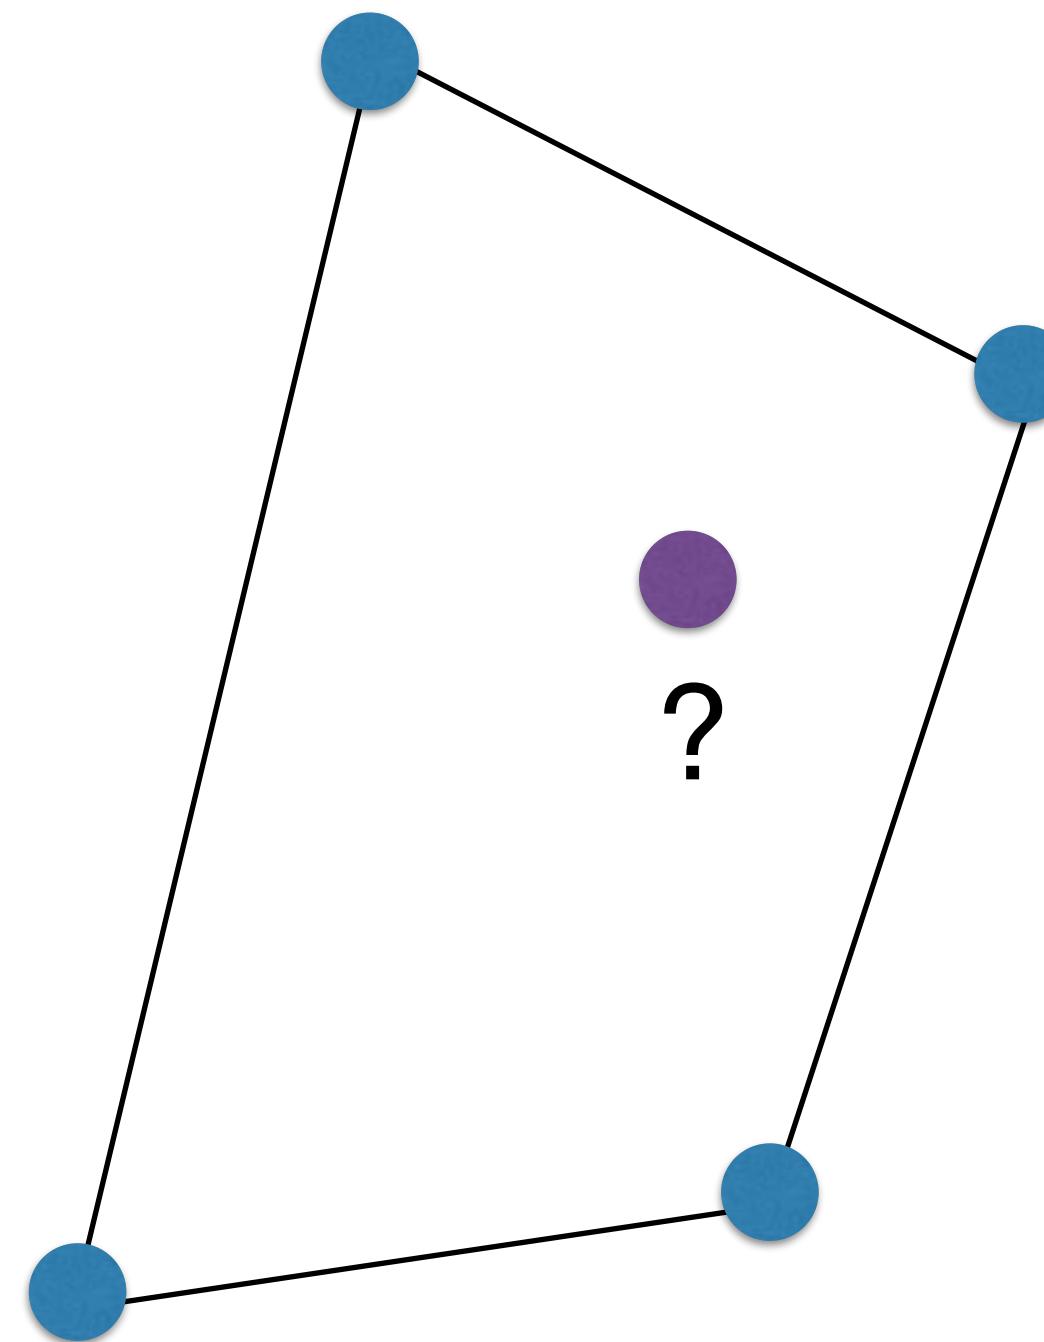
Bicubic

- Same idea as bilinear but using a 4×4 grid of neighbors
- Tends to produce sharper results
- More computationally intensive



Generalizing Bilinear

- Corner points don't have to lie on a square or rectangle
- Can be any quadrilateral



Generalized Bilinear

- General form:

$$f(x, y) = ax + by + cxy + d$$

- Use similar strategy as with curve fitting:

$$f(x_1, y_1) = ax_1 + by_1 + cx_1y_1 + d$$

$$f(x_2, y_2) = ax_2 + by_2 + cx_2y_2 + d$$

$$f(x_3, y_3) = ax_3 + by_3 + cx_3y_3 + d$$

$$f(x_4, y_4) = ax_4 + by_4 + cx_4y_4 + d$$

Solve for a, b, c, d and
plug into function

Coming up...

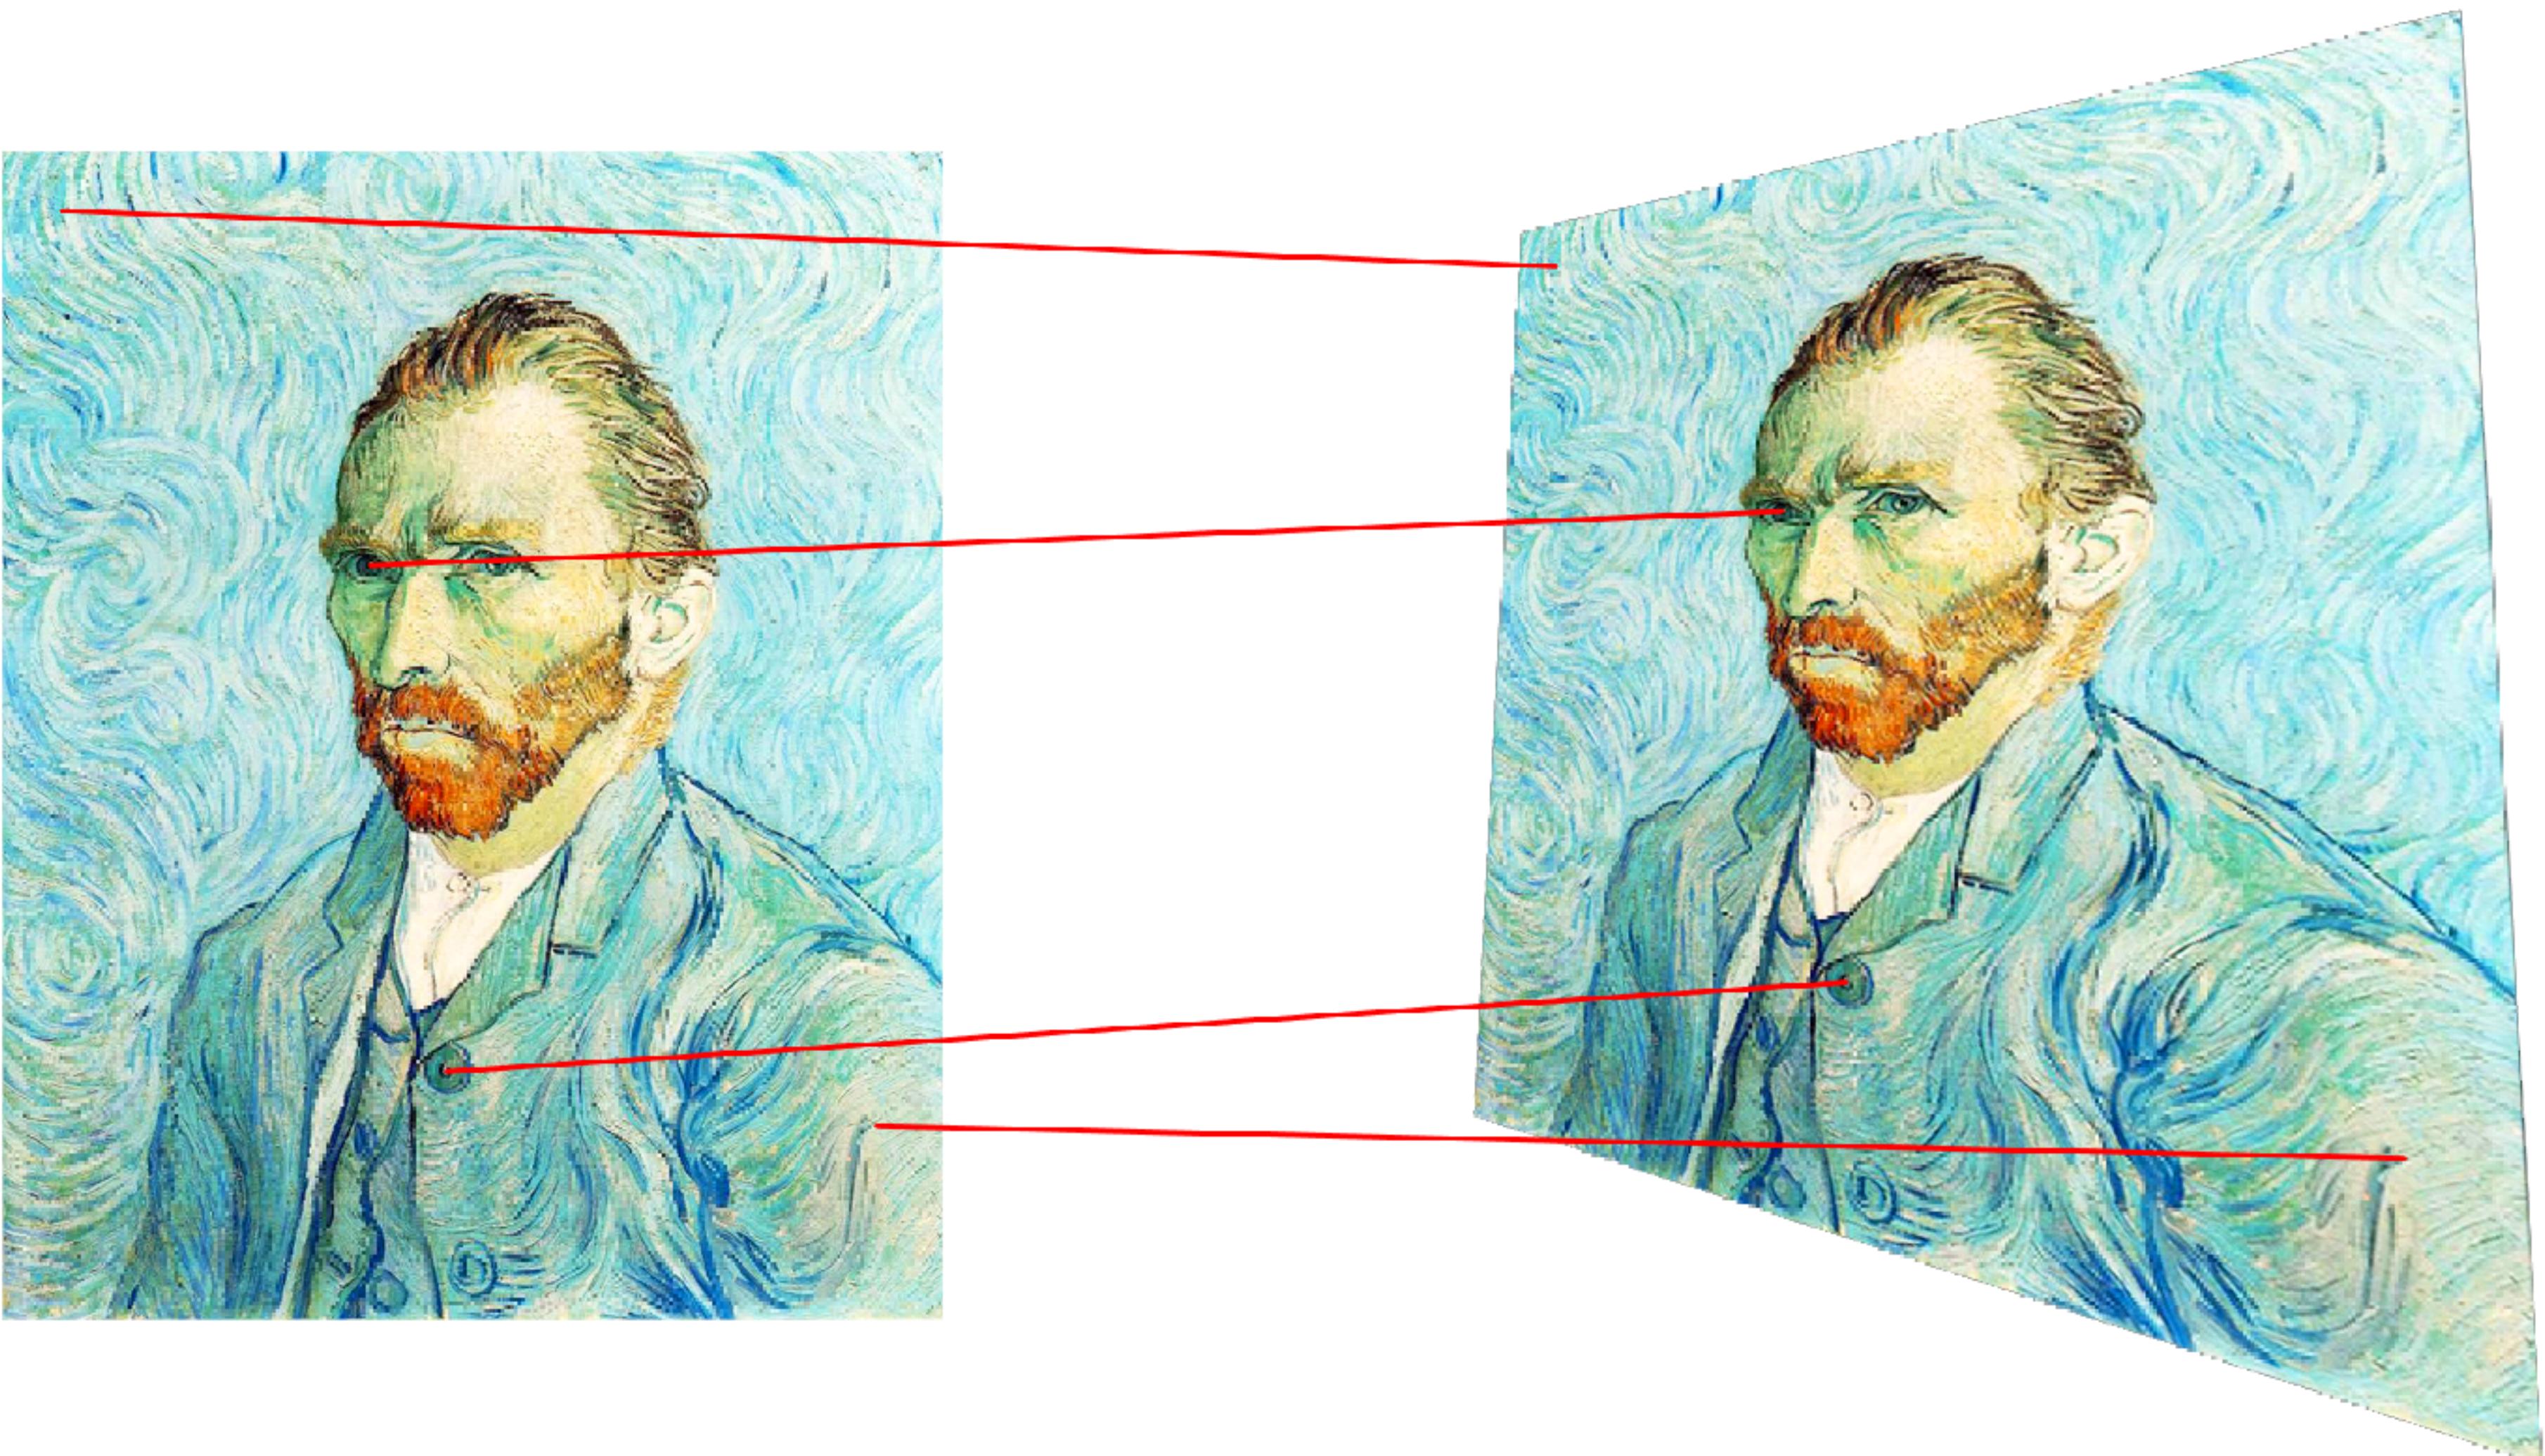
- Image warping
- Texture mapping



Image Warping (cont'd)

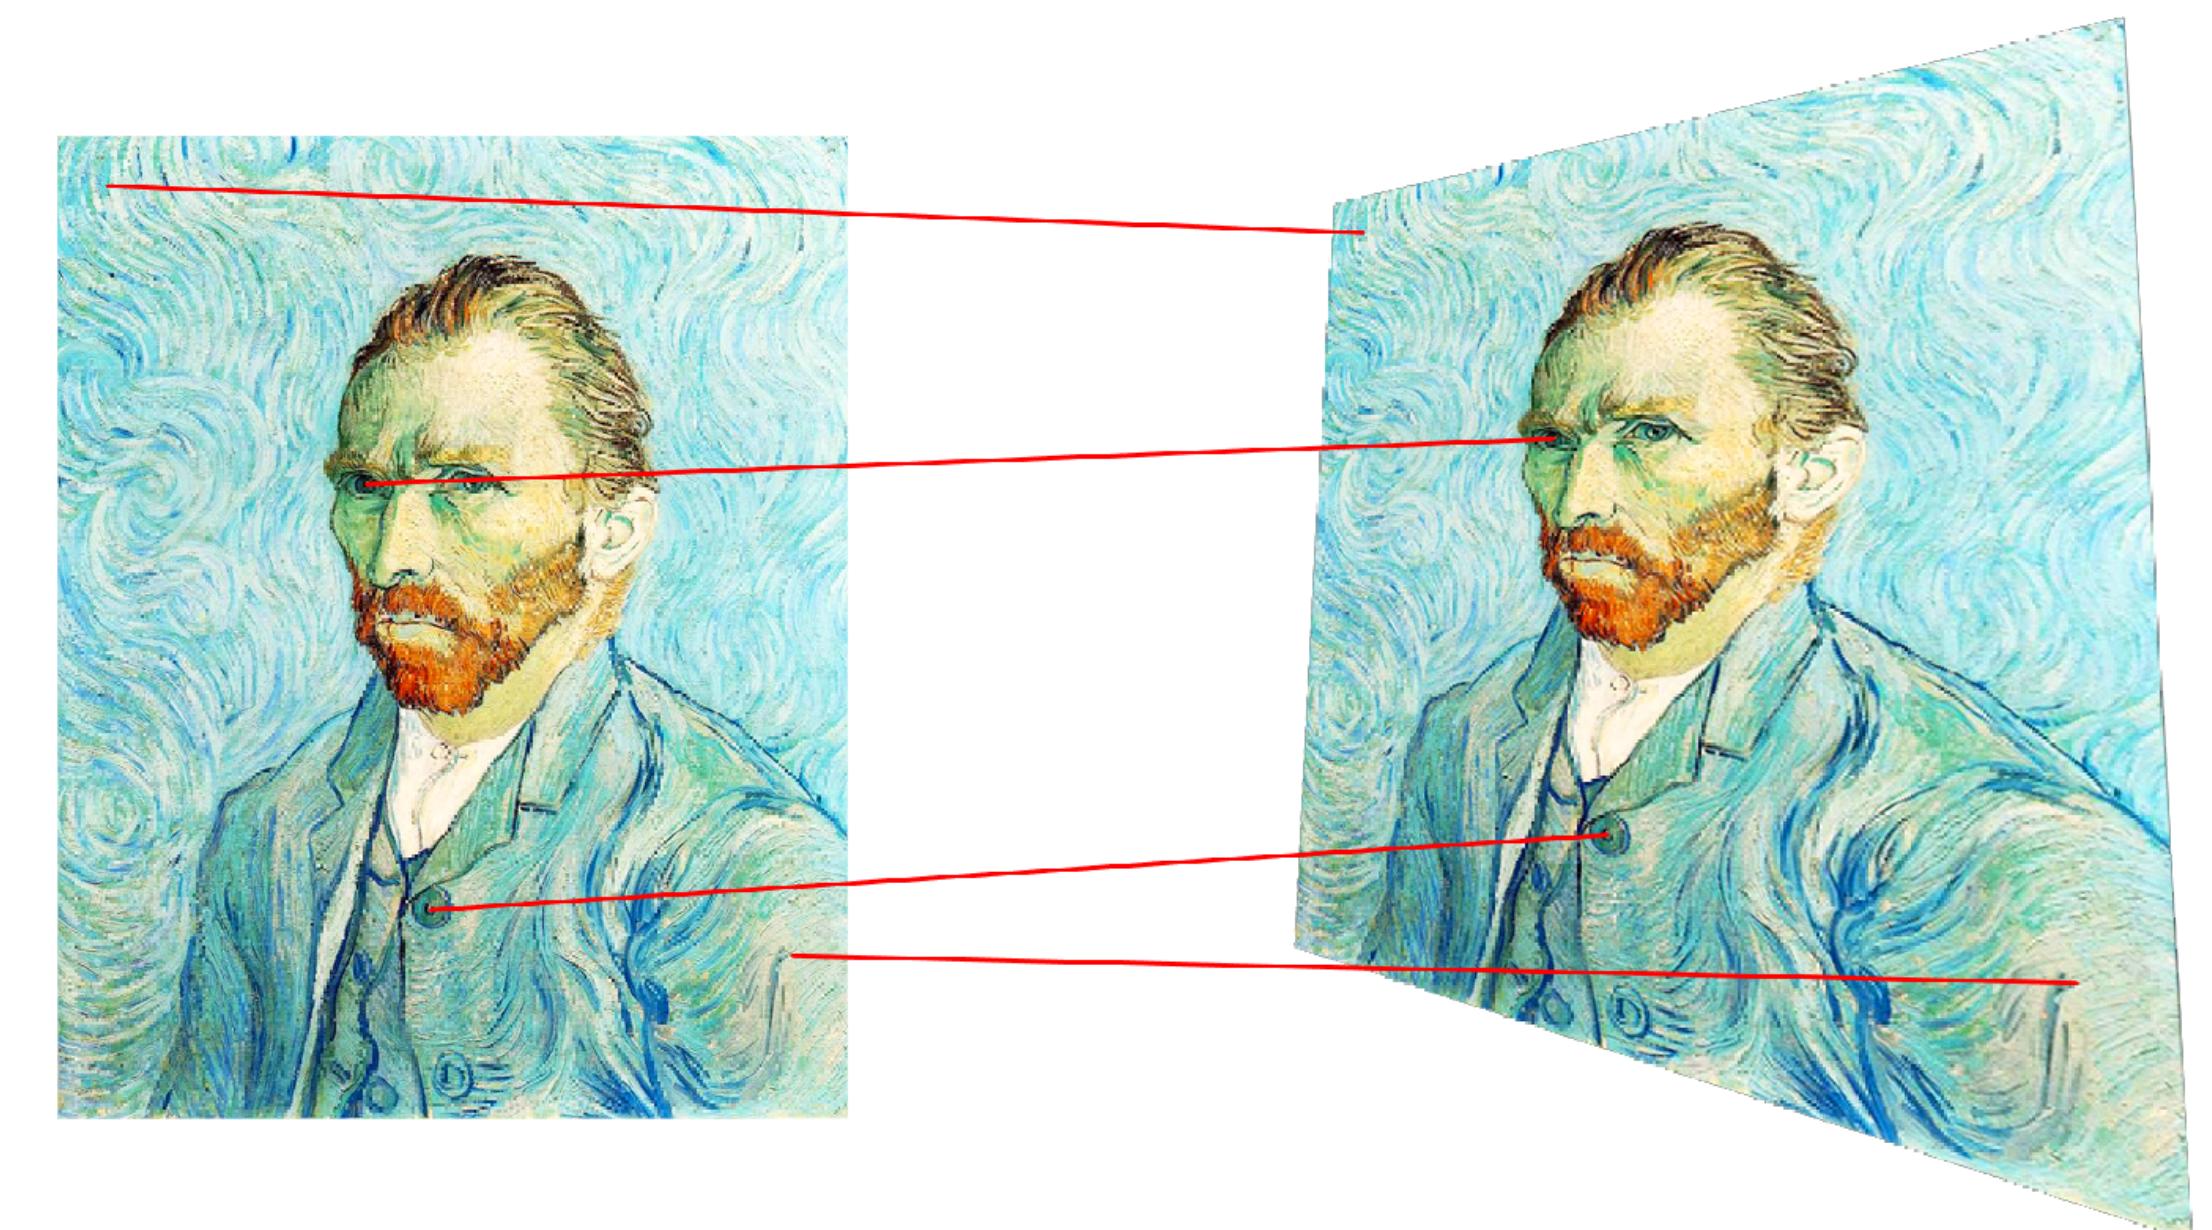
CS 355: Introduction to Graphics and Image Processing

Perspective Warping



Perspective Warping

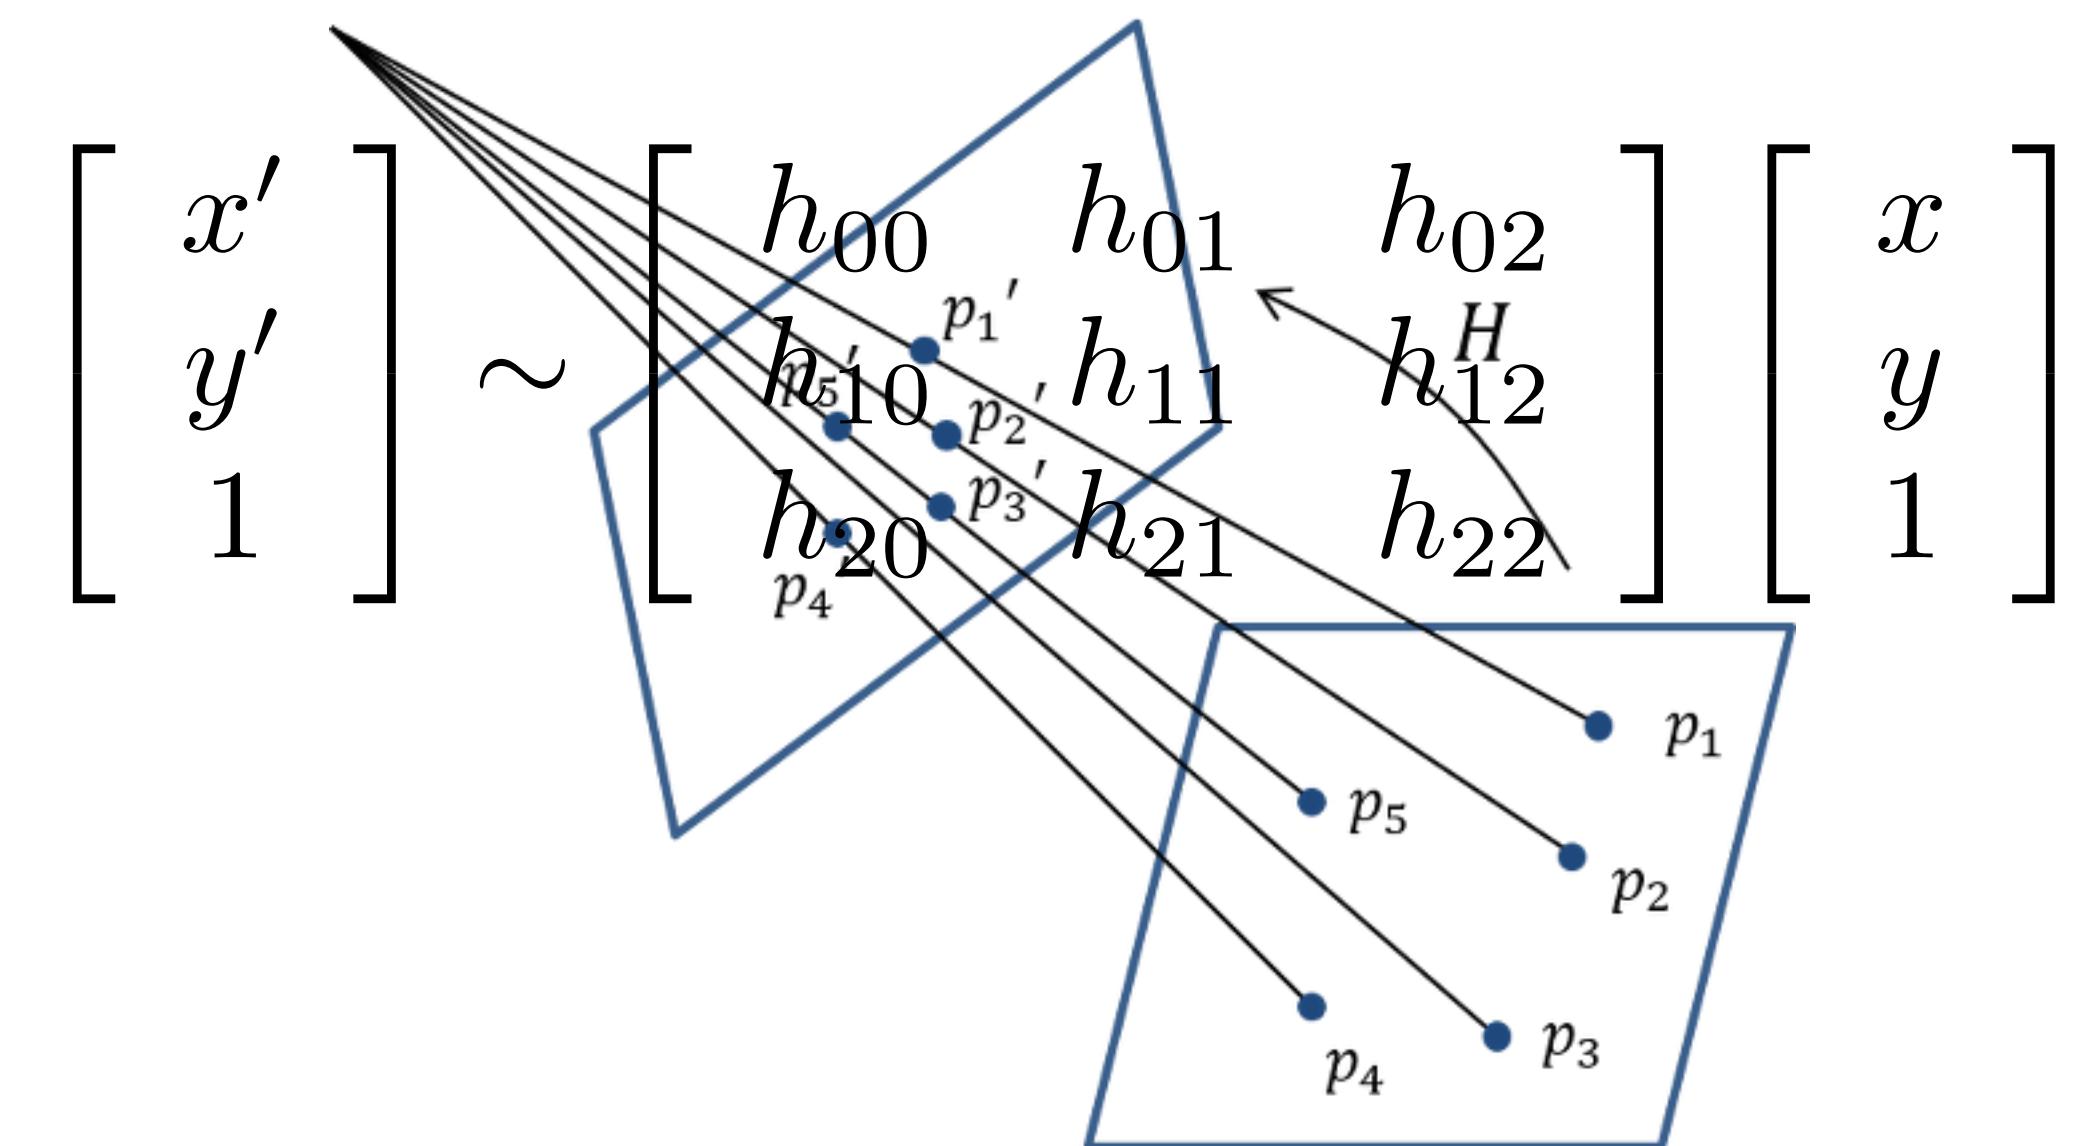
- You've already seen how to project 3D geometry to create an image
- Special case: if the points all lie on a plane in 3D, the transformation can be done with a single 3×3 matrix
- Perspective projection of any 3D plane to any other 3D plane



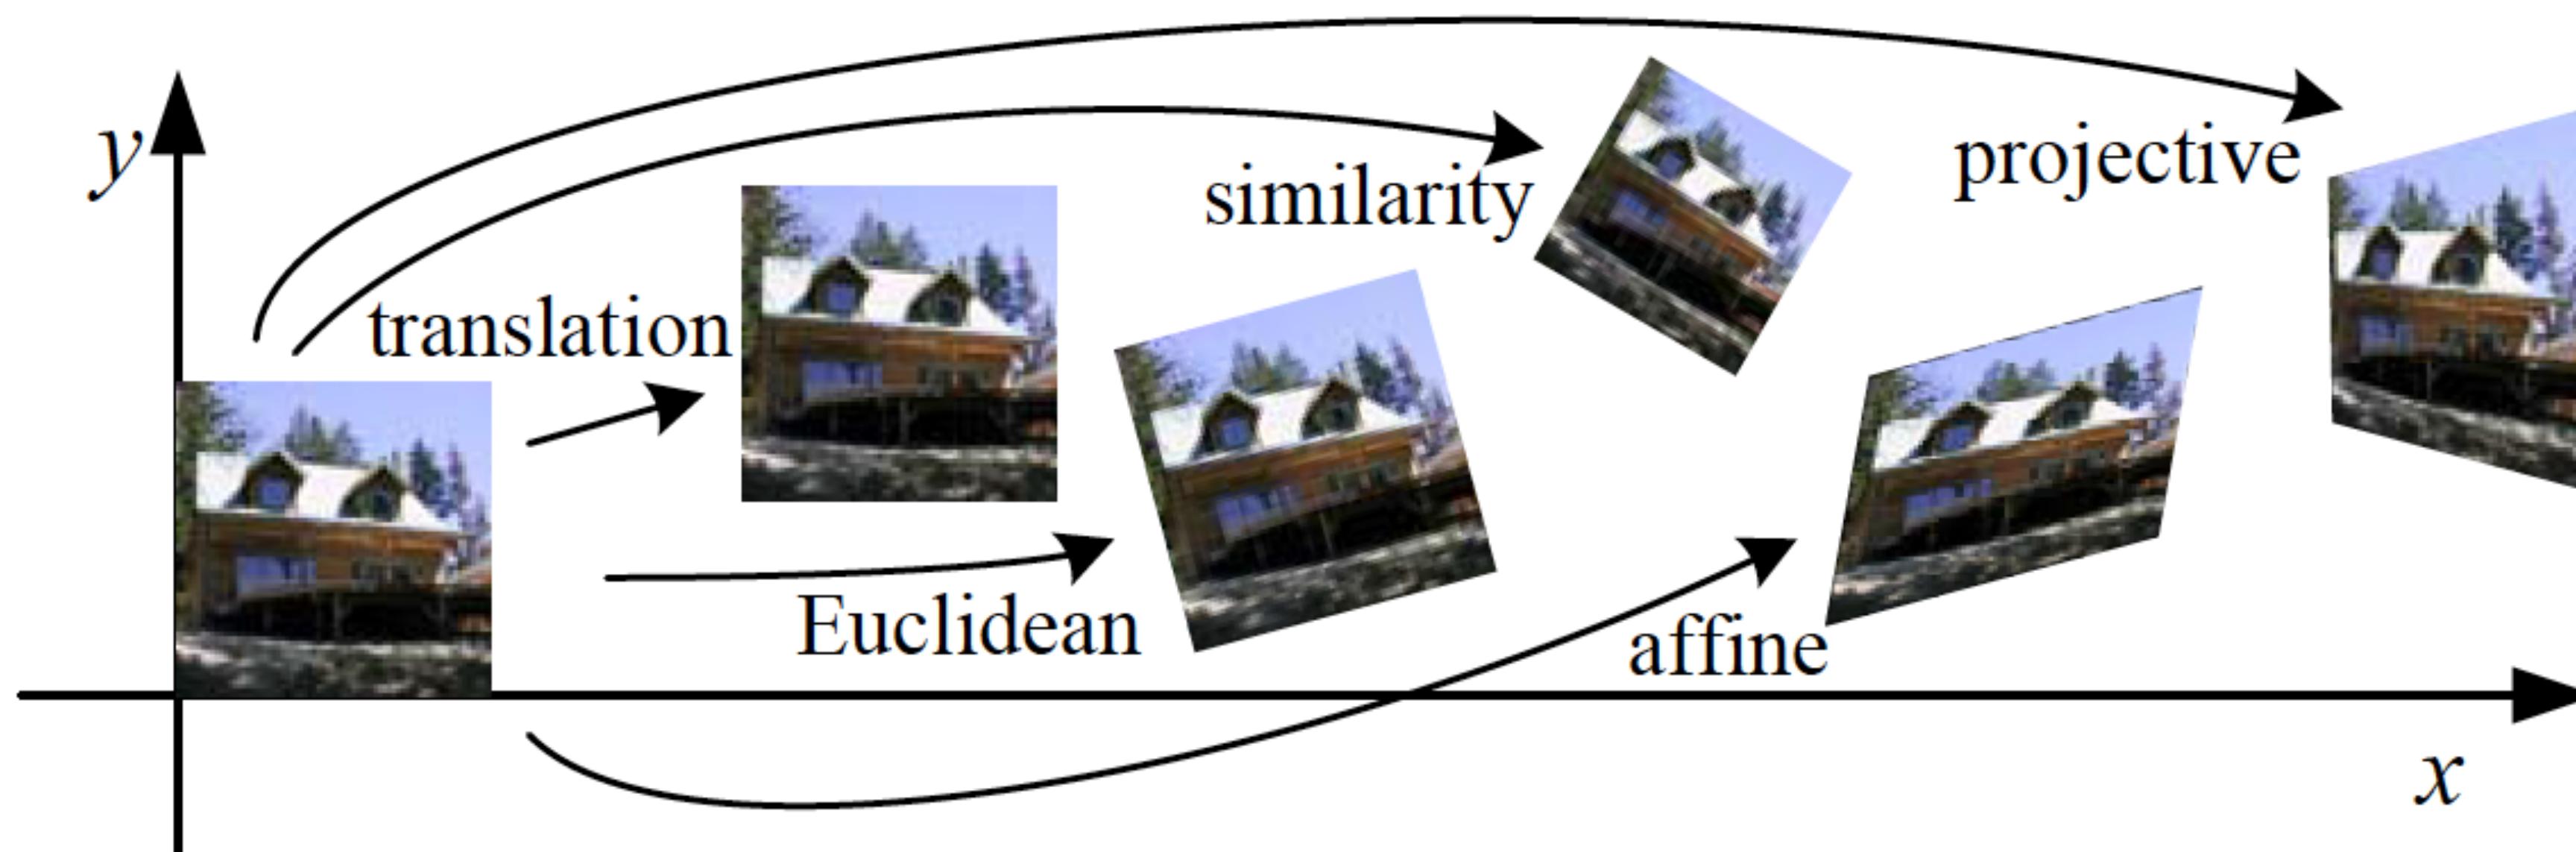
Homographies

- Perspective projection from a plane to another plane is a homography
- Same 3×3 matrix transformation you've seen before, but the bottom row doesn't have to be $[0 \ 0 \ 1]$
- In practice, it's invertible!

$$\mathbf{p}' = \mathbf{H} \ \mathbf{p}$$



Classes of Transformations



Homographies

- Non-zero multiple of a homography does the same transformation
- That means there are really only 8 free parameters, not 9

$$\mathbf{p}' = \mathbf{H} \mathbf{p}$$

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \sim \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Multiple Pictures of the Same Plane

- If two cameras take pictures of a plane from different viewpoints,
 - The projection of the plane to each camera is a homography
 - The transformation of points from one picture to the other is also a homography!

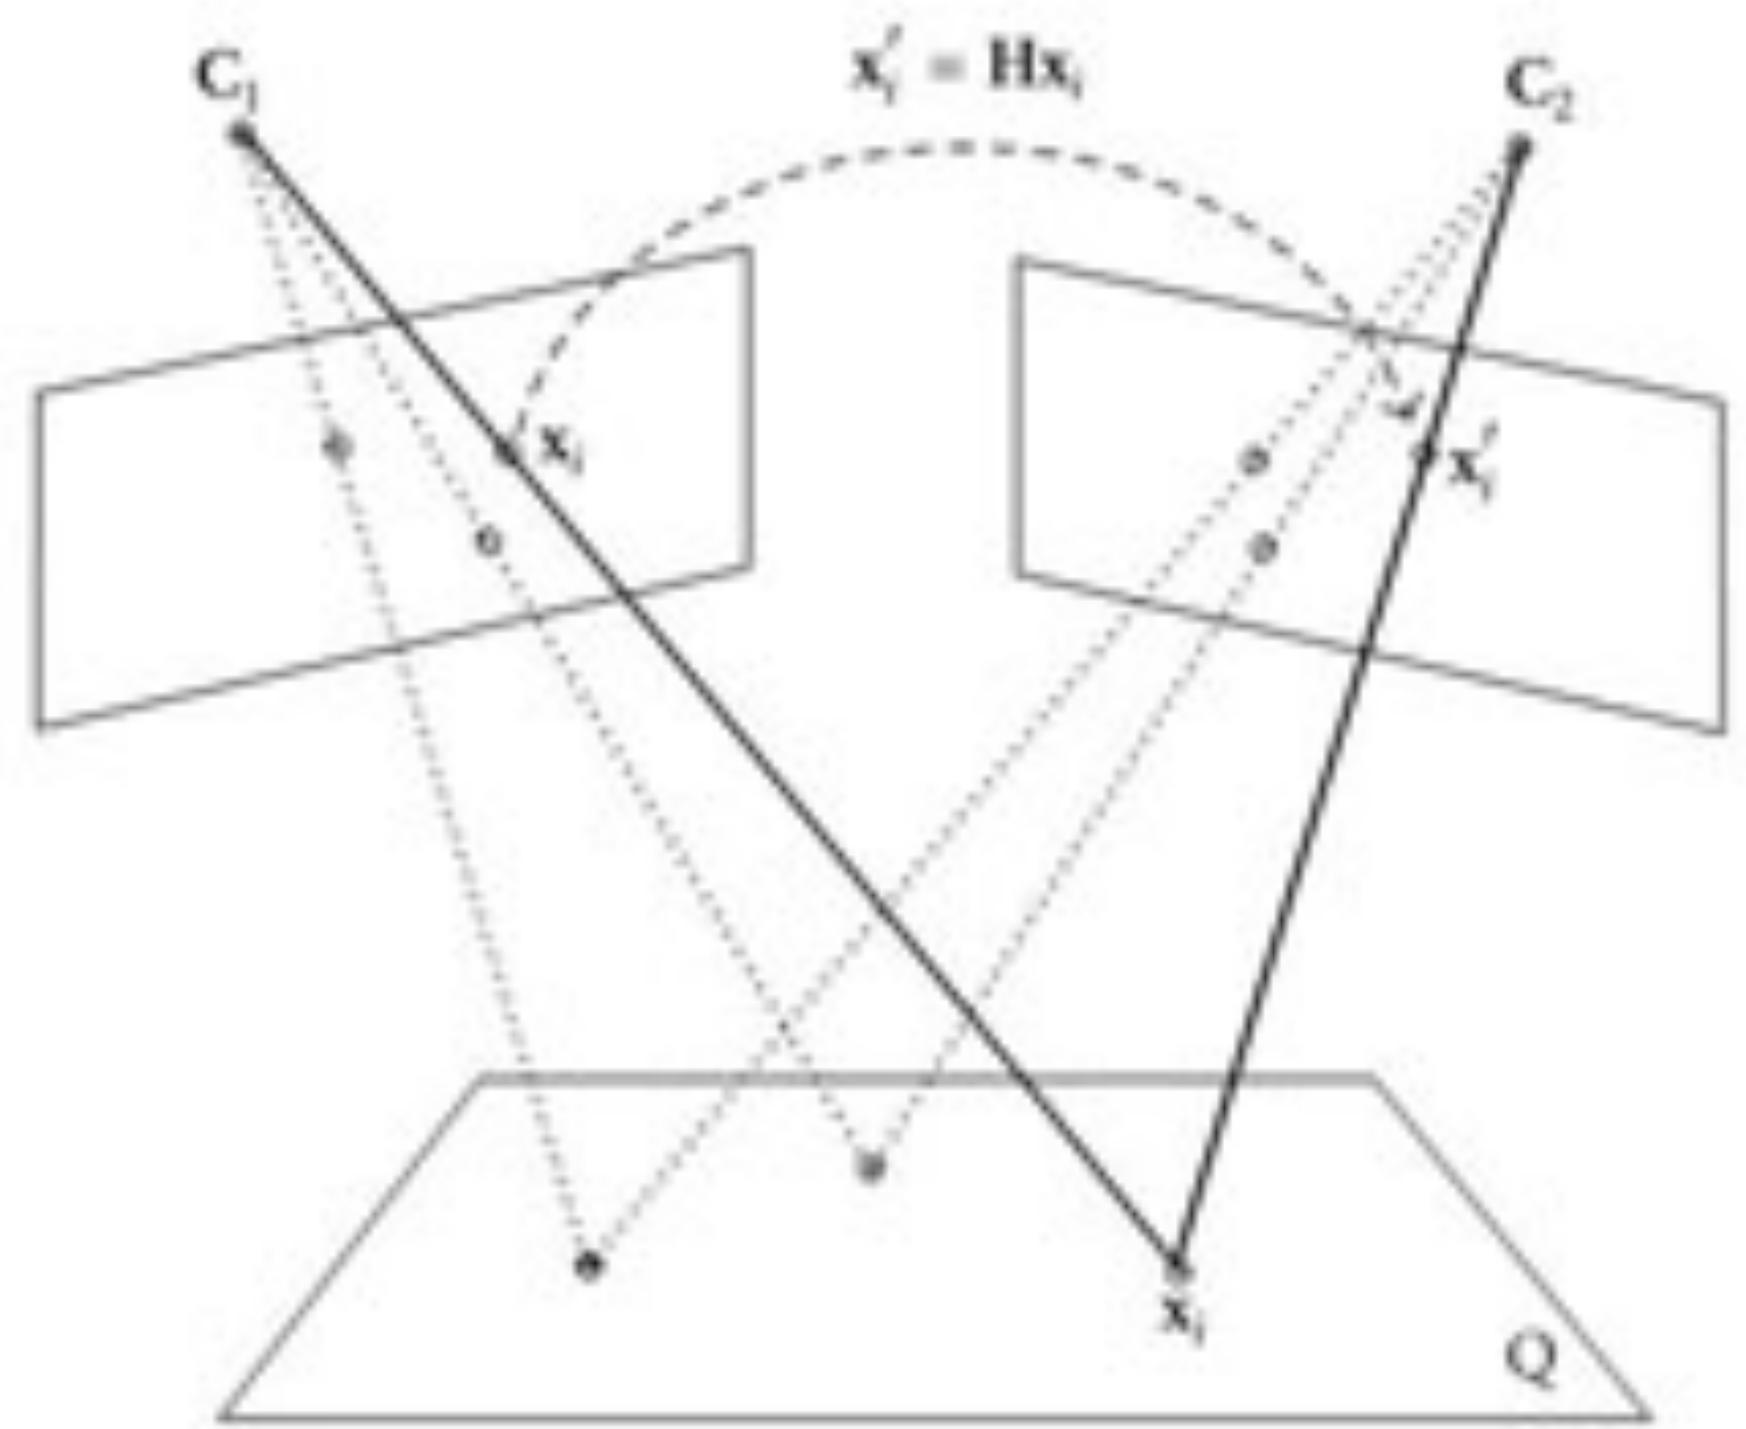


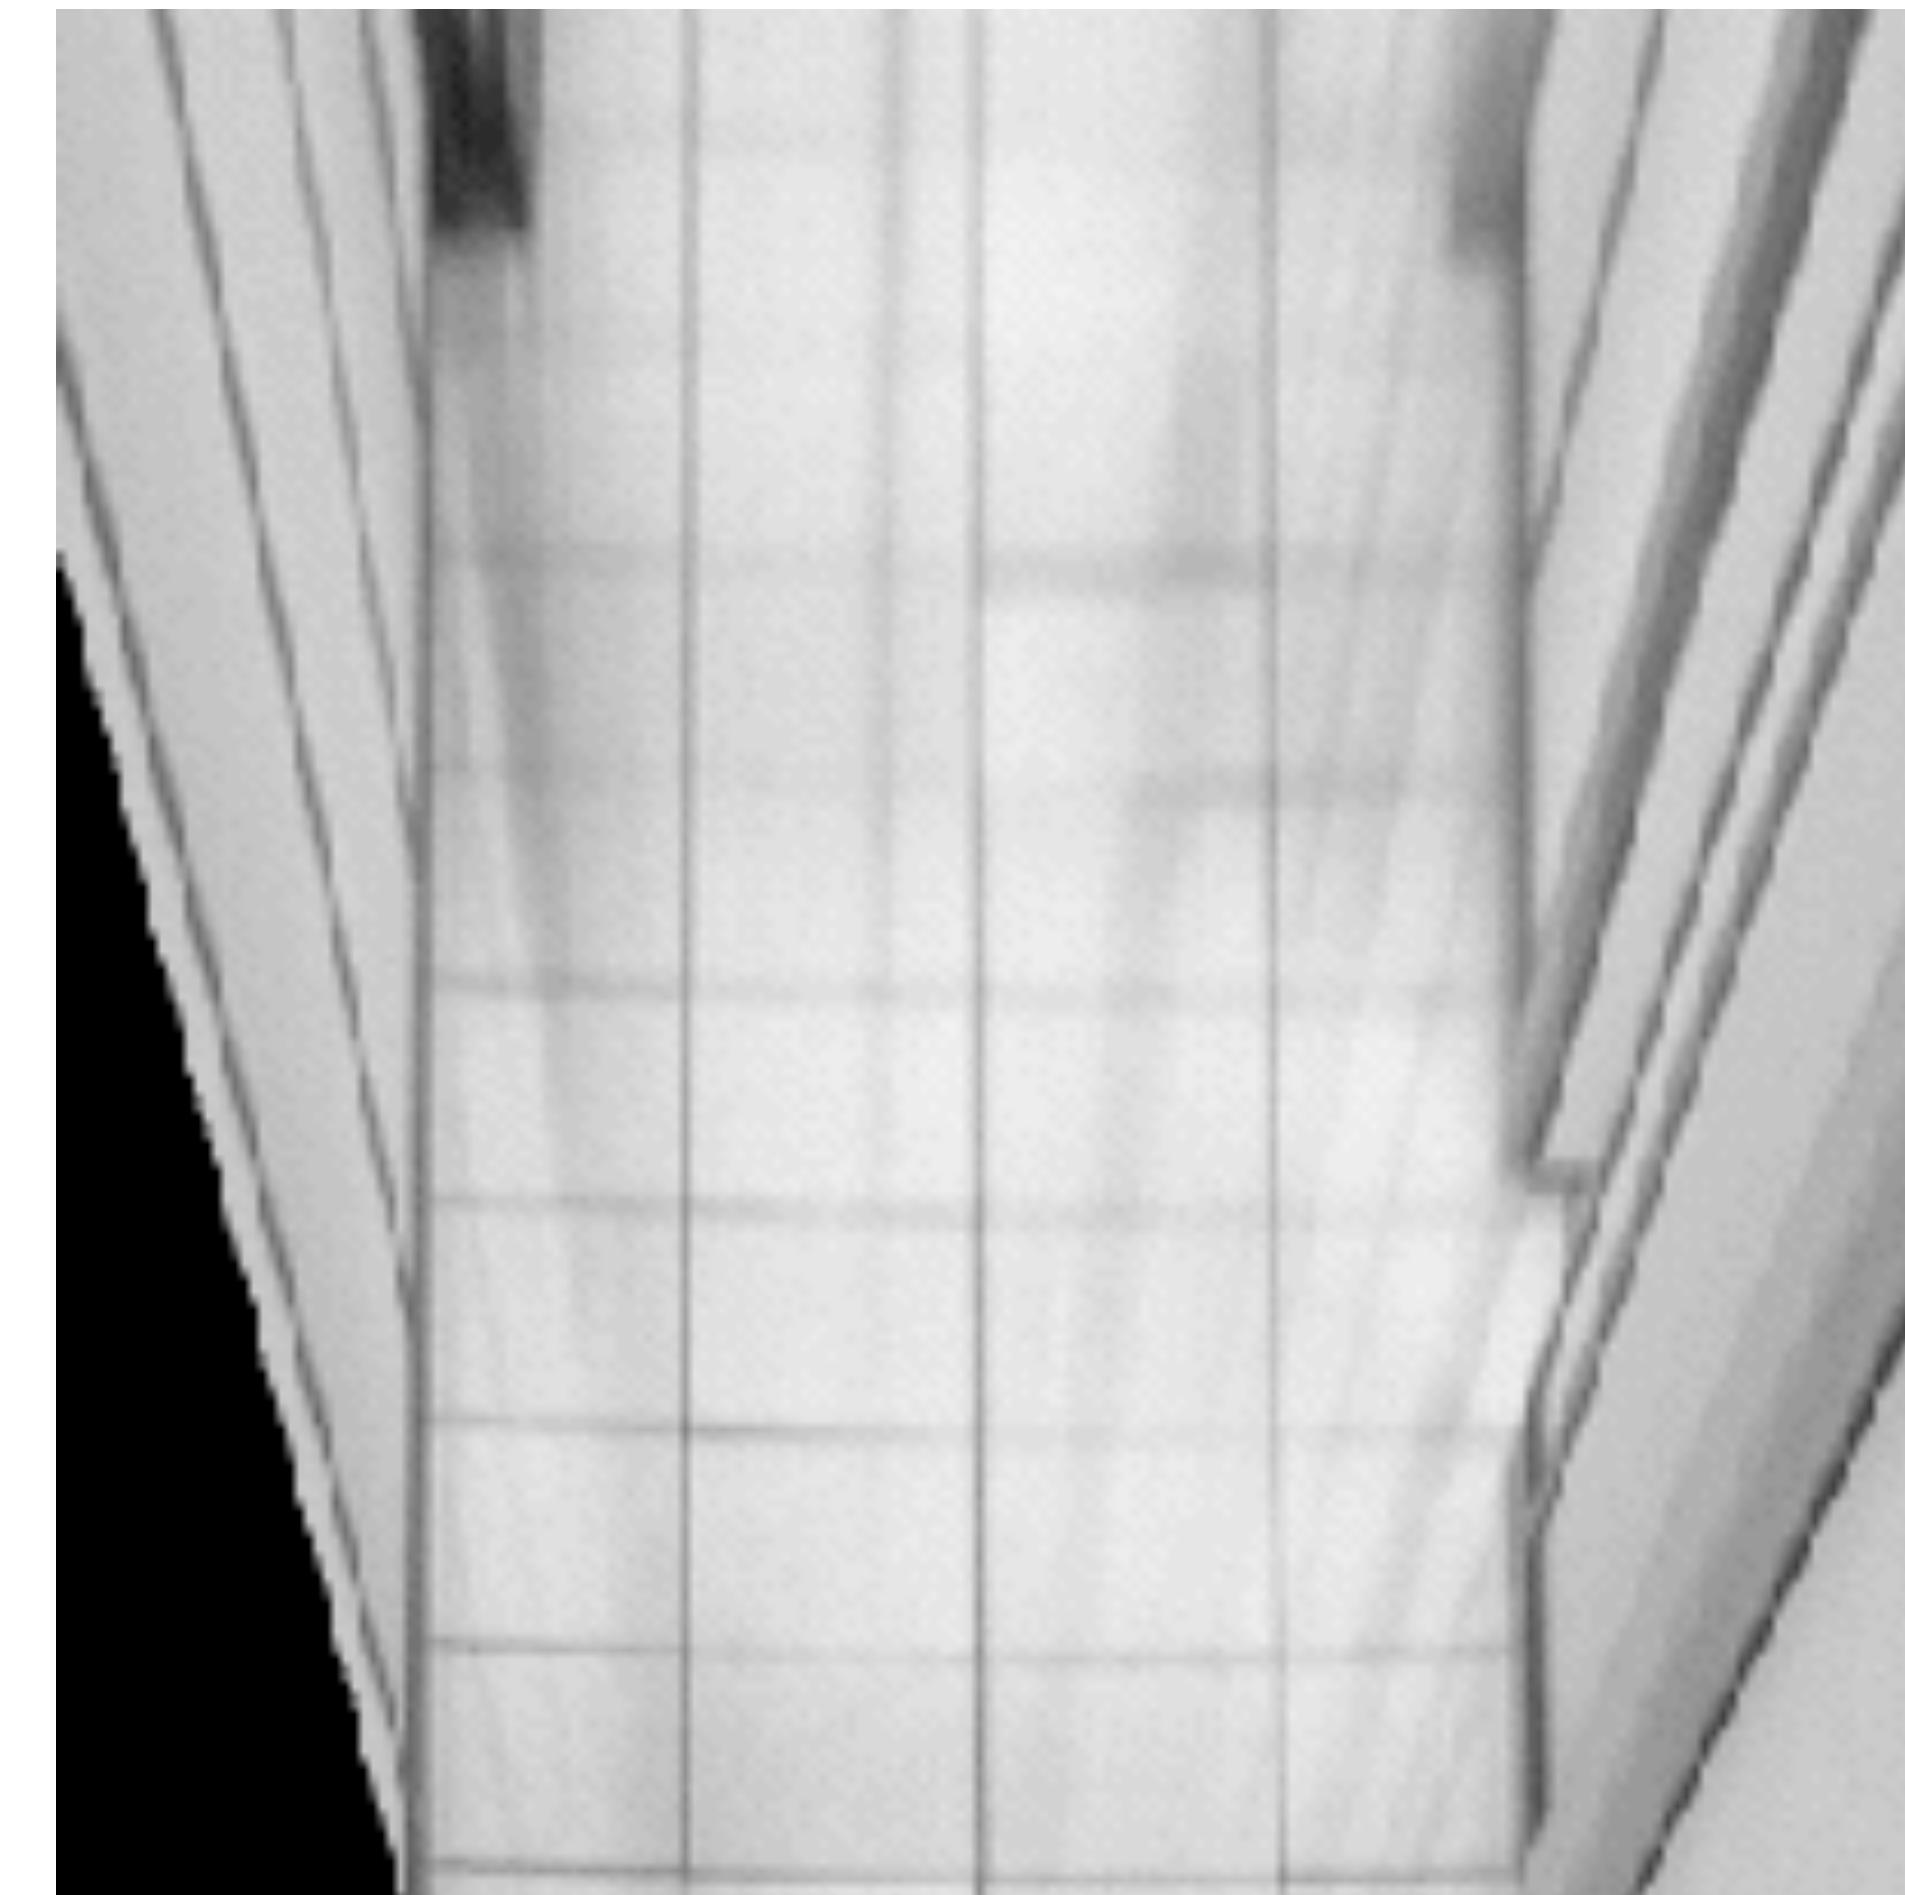
Image Stitching



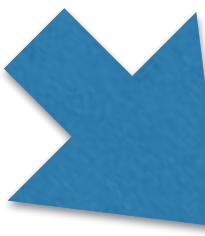
“Straight On” View



“Birds-Eye” View

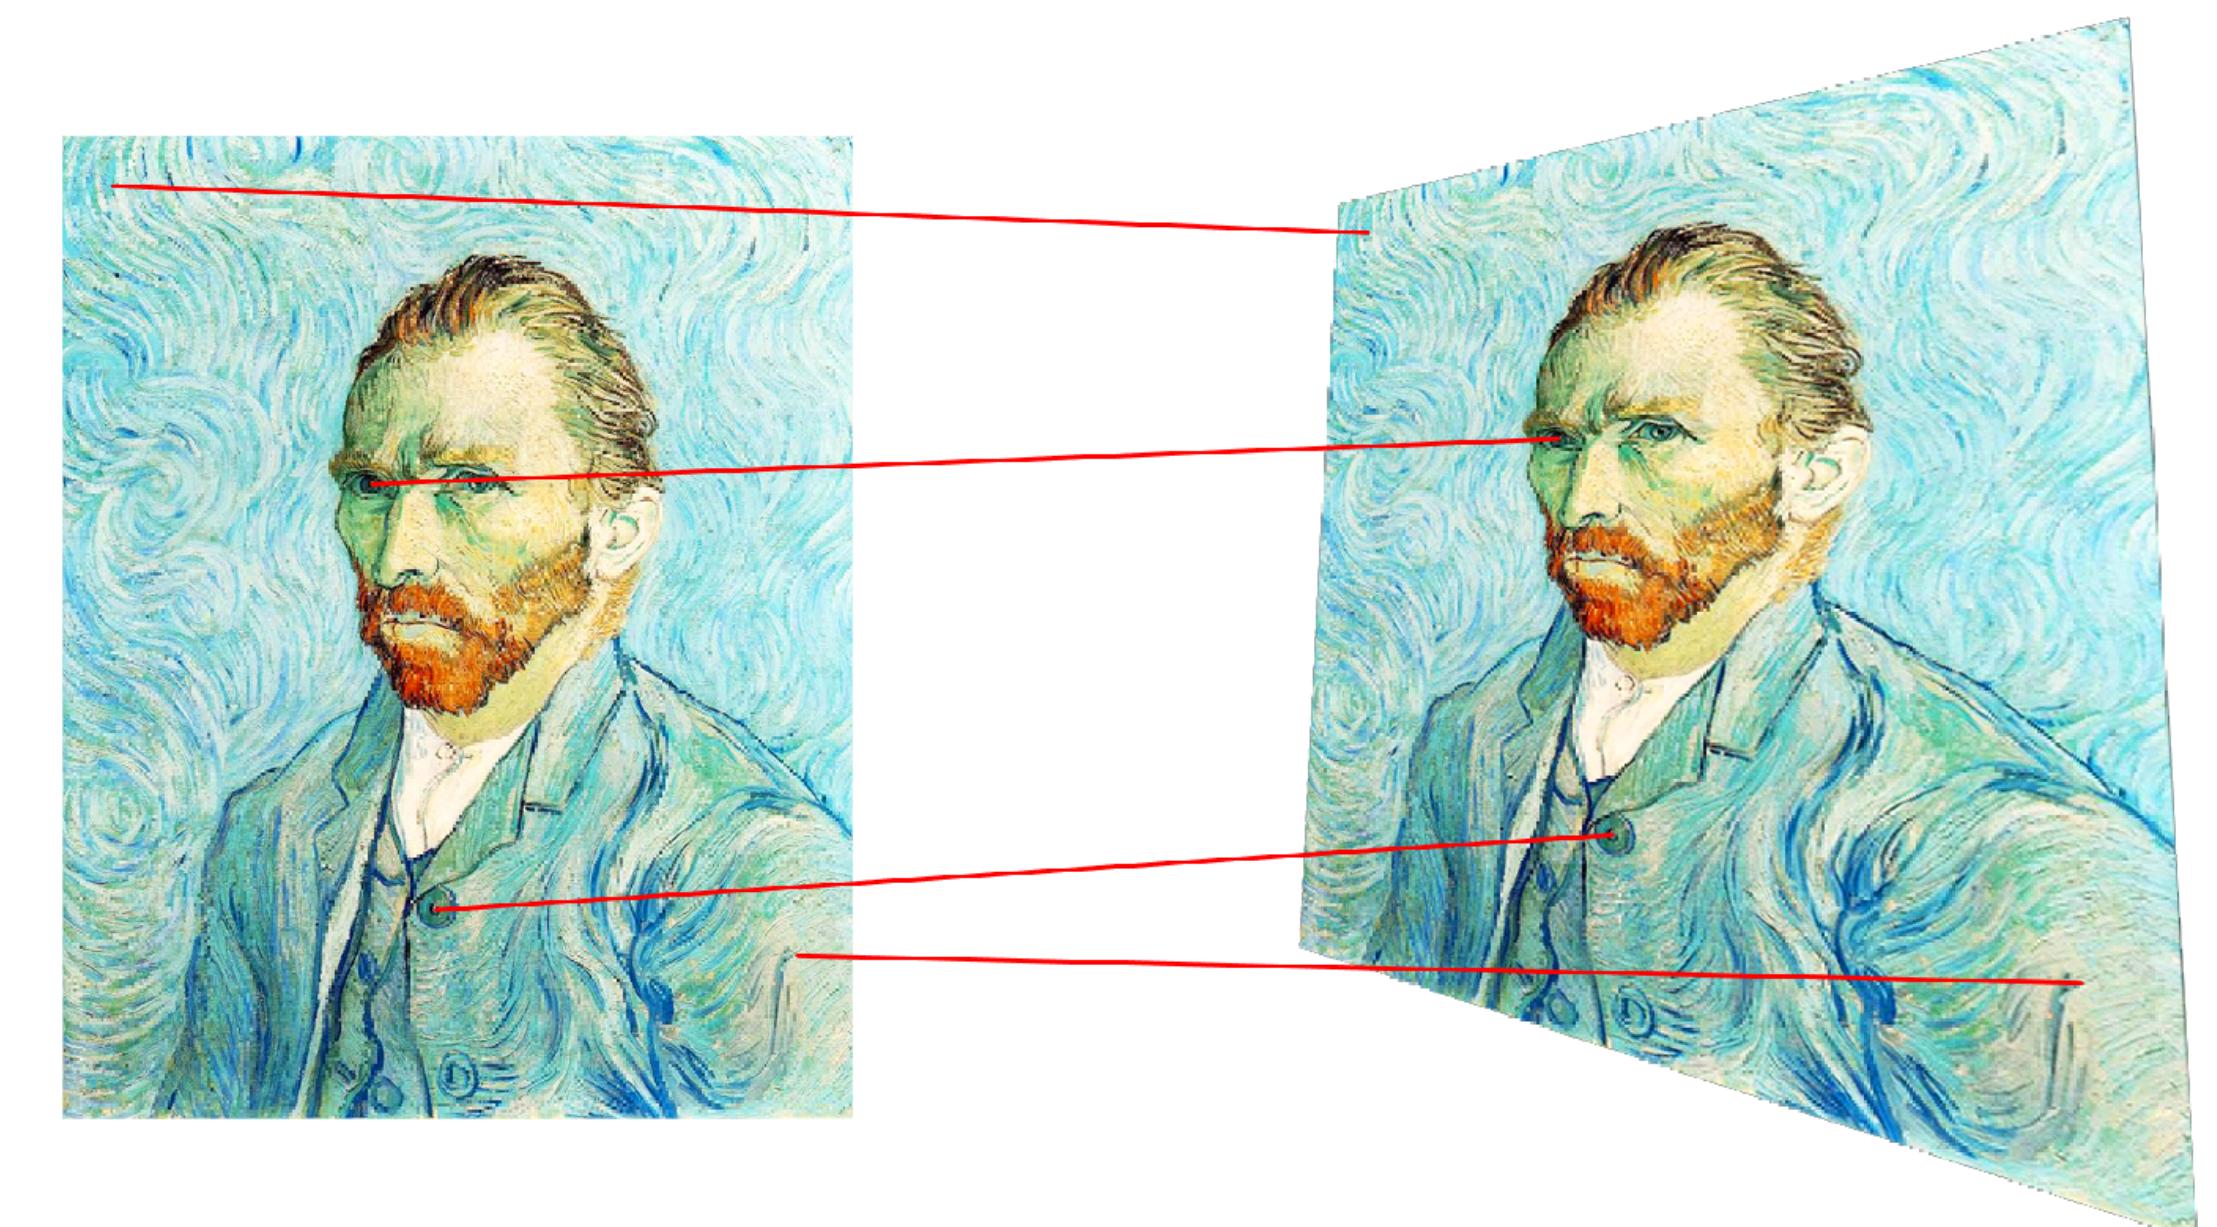


Or Going The Other Way

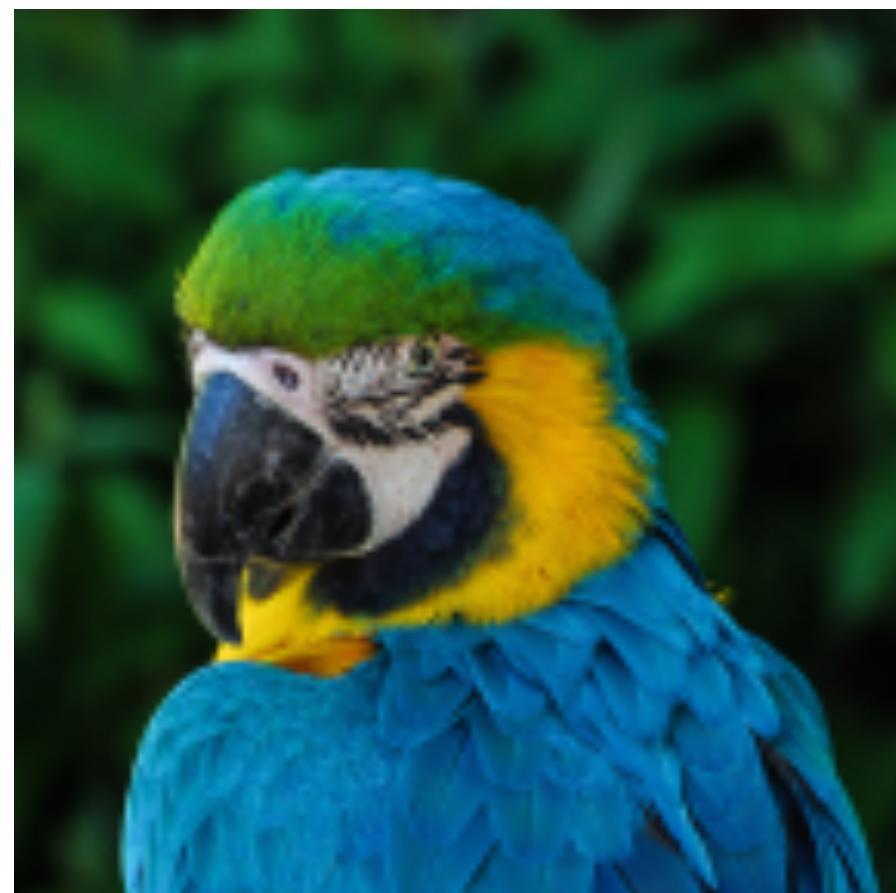


Computing Homographies

- Eight free parameters
- A pair of matching points provides two constraints (one for x, one for y)
- Four points provides eight equations to solve for eight unknowns
- Call the *4-point algorithm*



Lab 9



Lab 9



Lab 9

- We give you an implementation of the 4-point algorithm:
 - You give it matching points
 - It gives you back a homography
- You implement:
 - Backward warping
 - Bilinear interpolation
 - Point-in-frame geometric tests
(more on this later)



Coming up...

- Back to more geometry:
 - Representing lines, curves, surfaces
 - Geometric tests



Lines, Curves, and Surfaces

CS 355: Introduction to Graphics and Image Processing

Lines



Lines extend infinitely in both directions

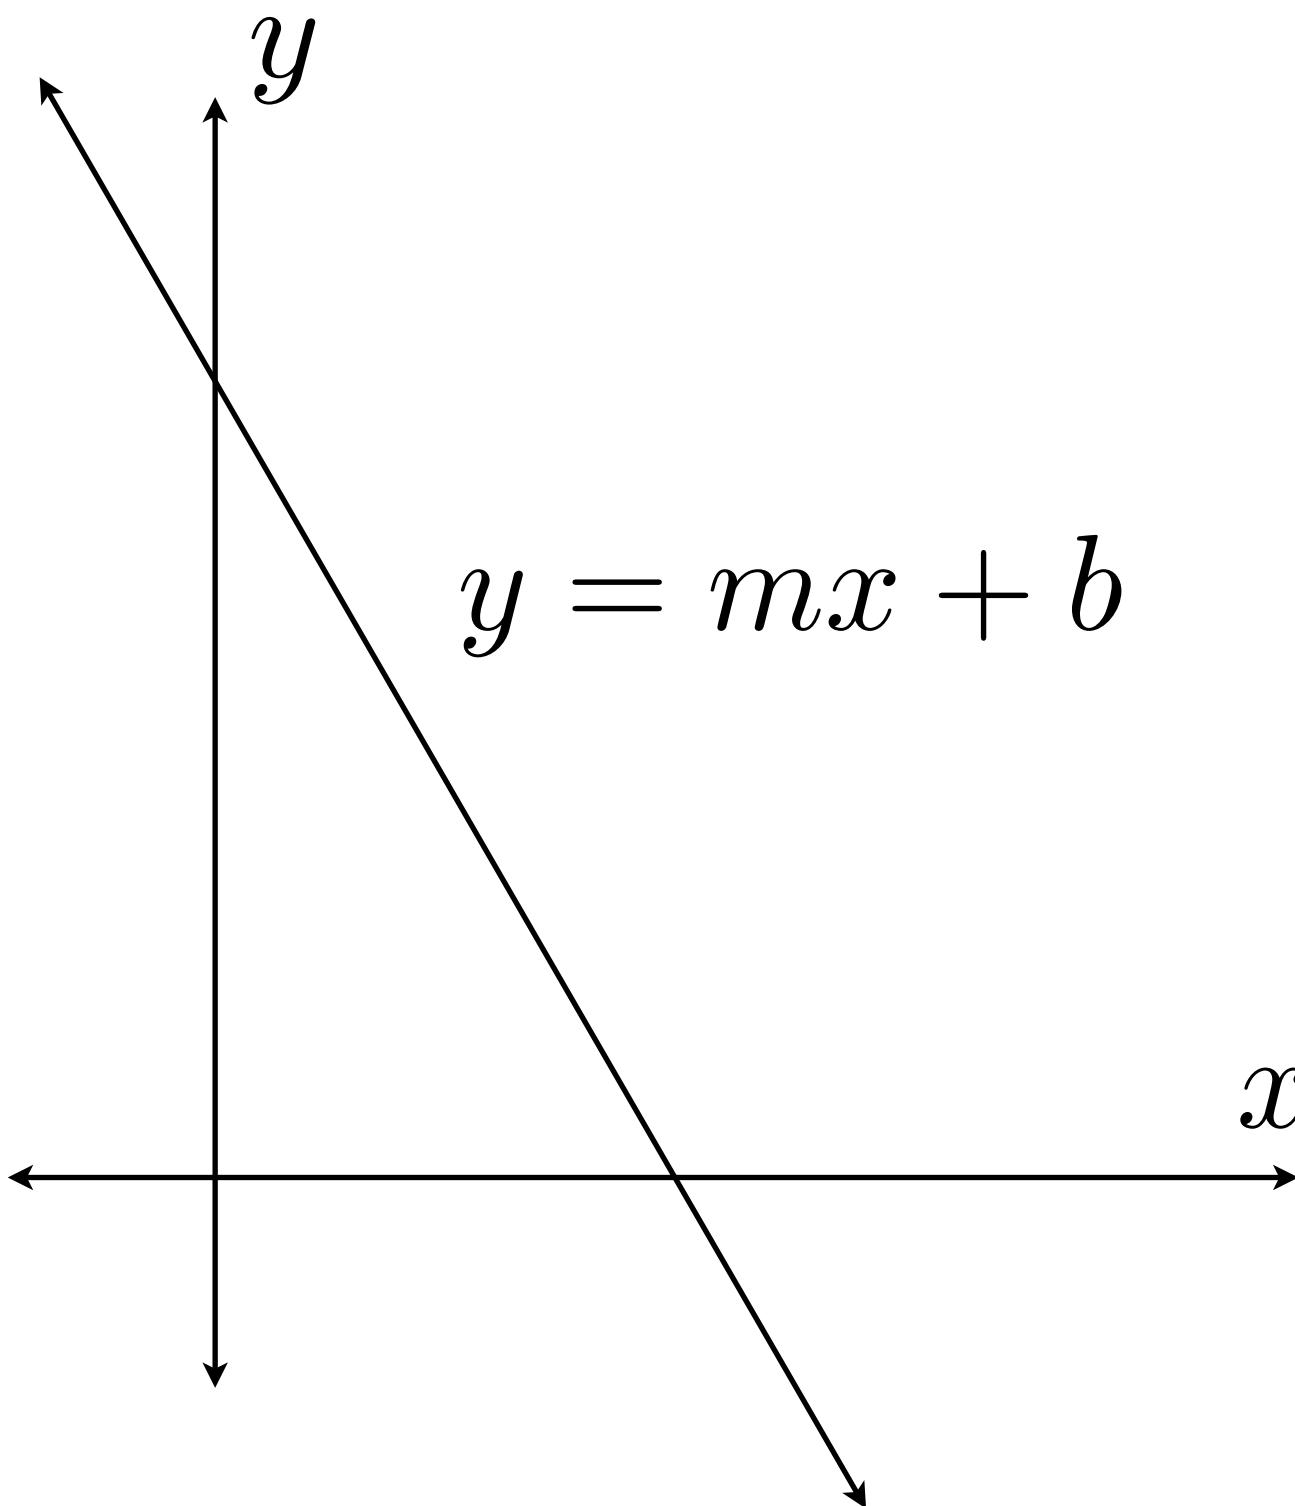


Line segments are finite



Rays extend infinitely in one direction

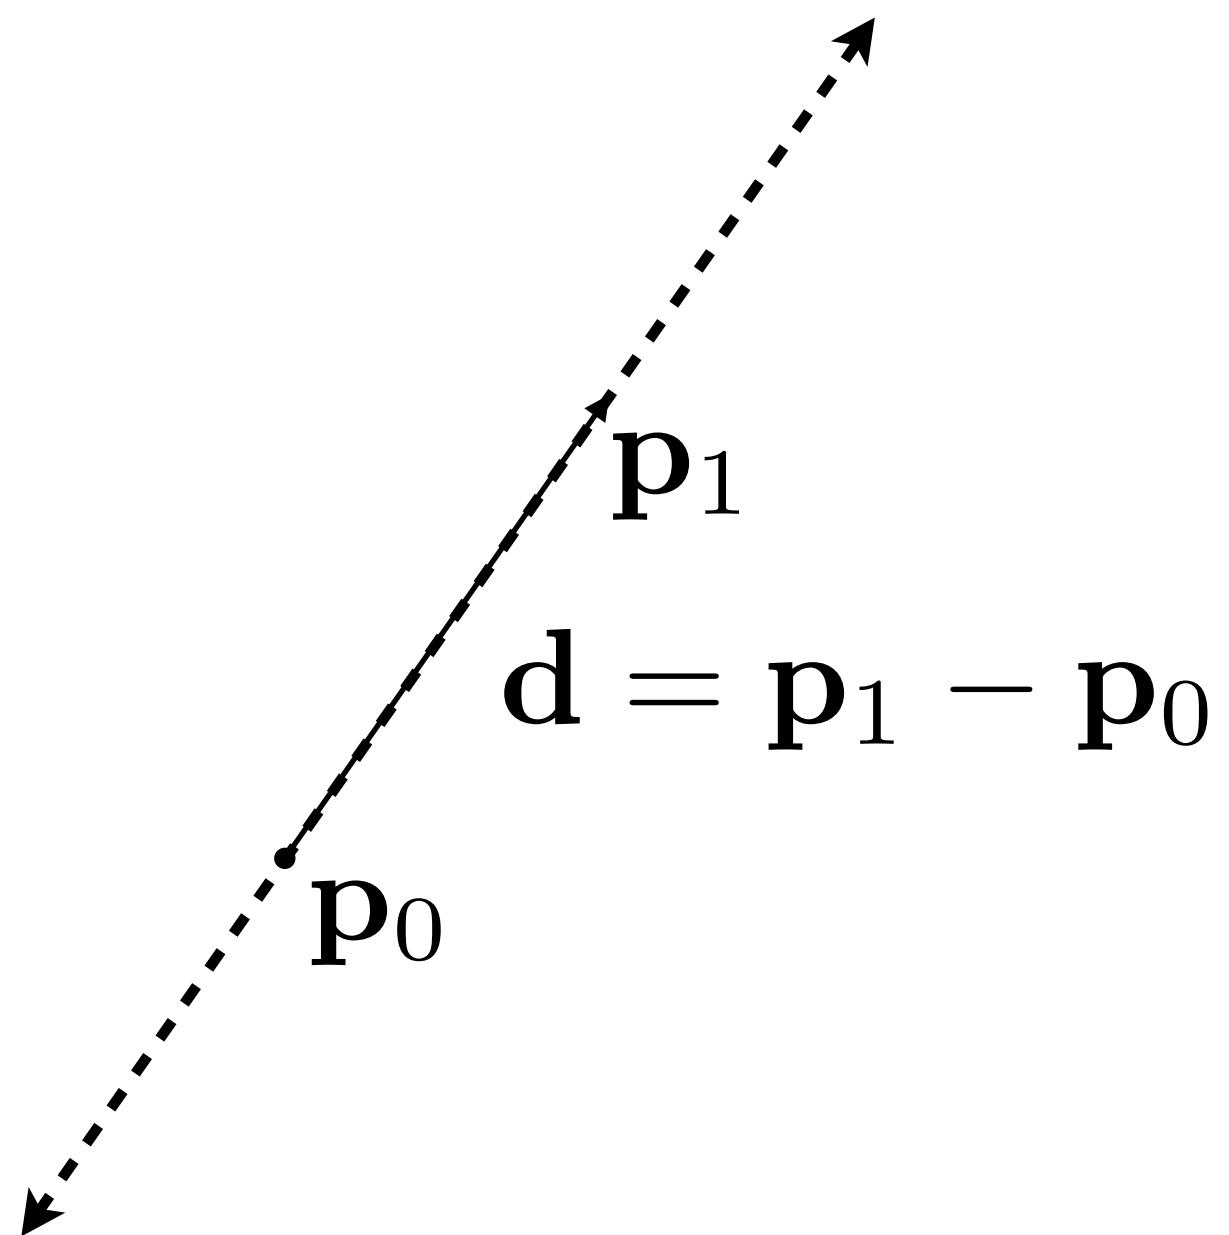
Slope / Intercept



Doesn't work for vertical lines

Hard to extend to 3D

Parametric Representation



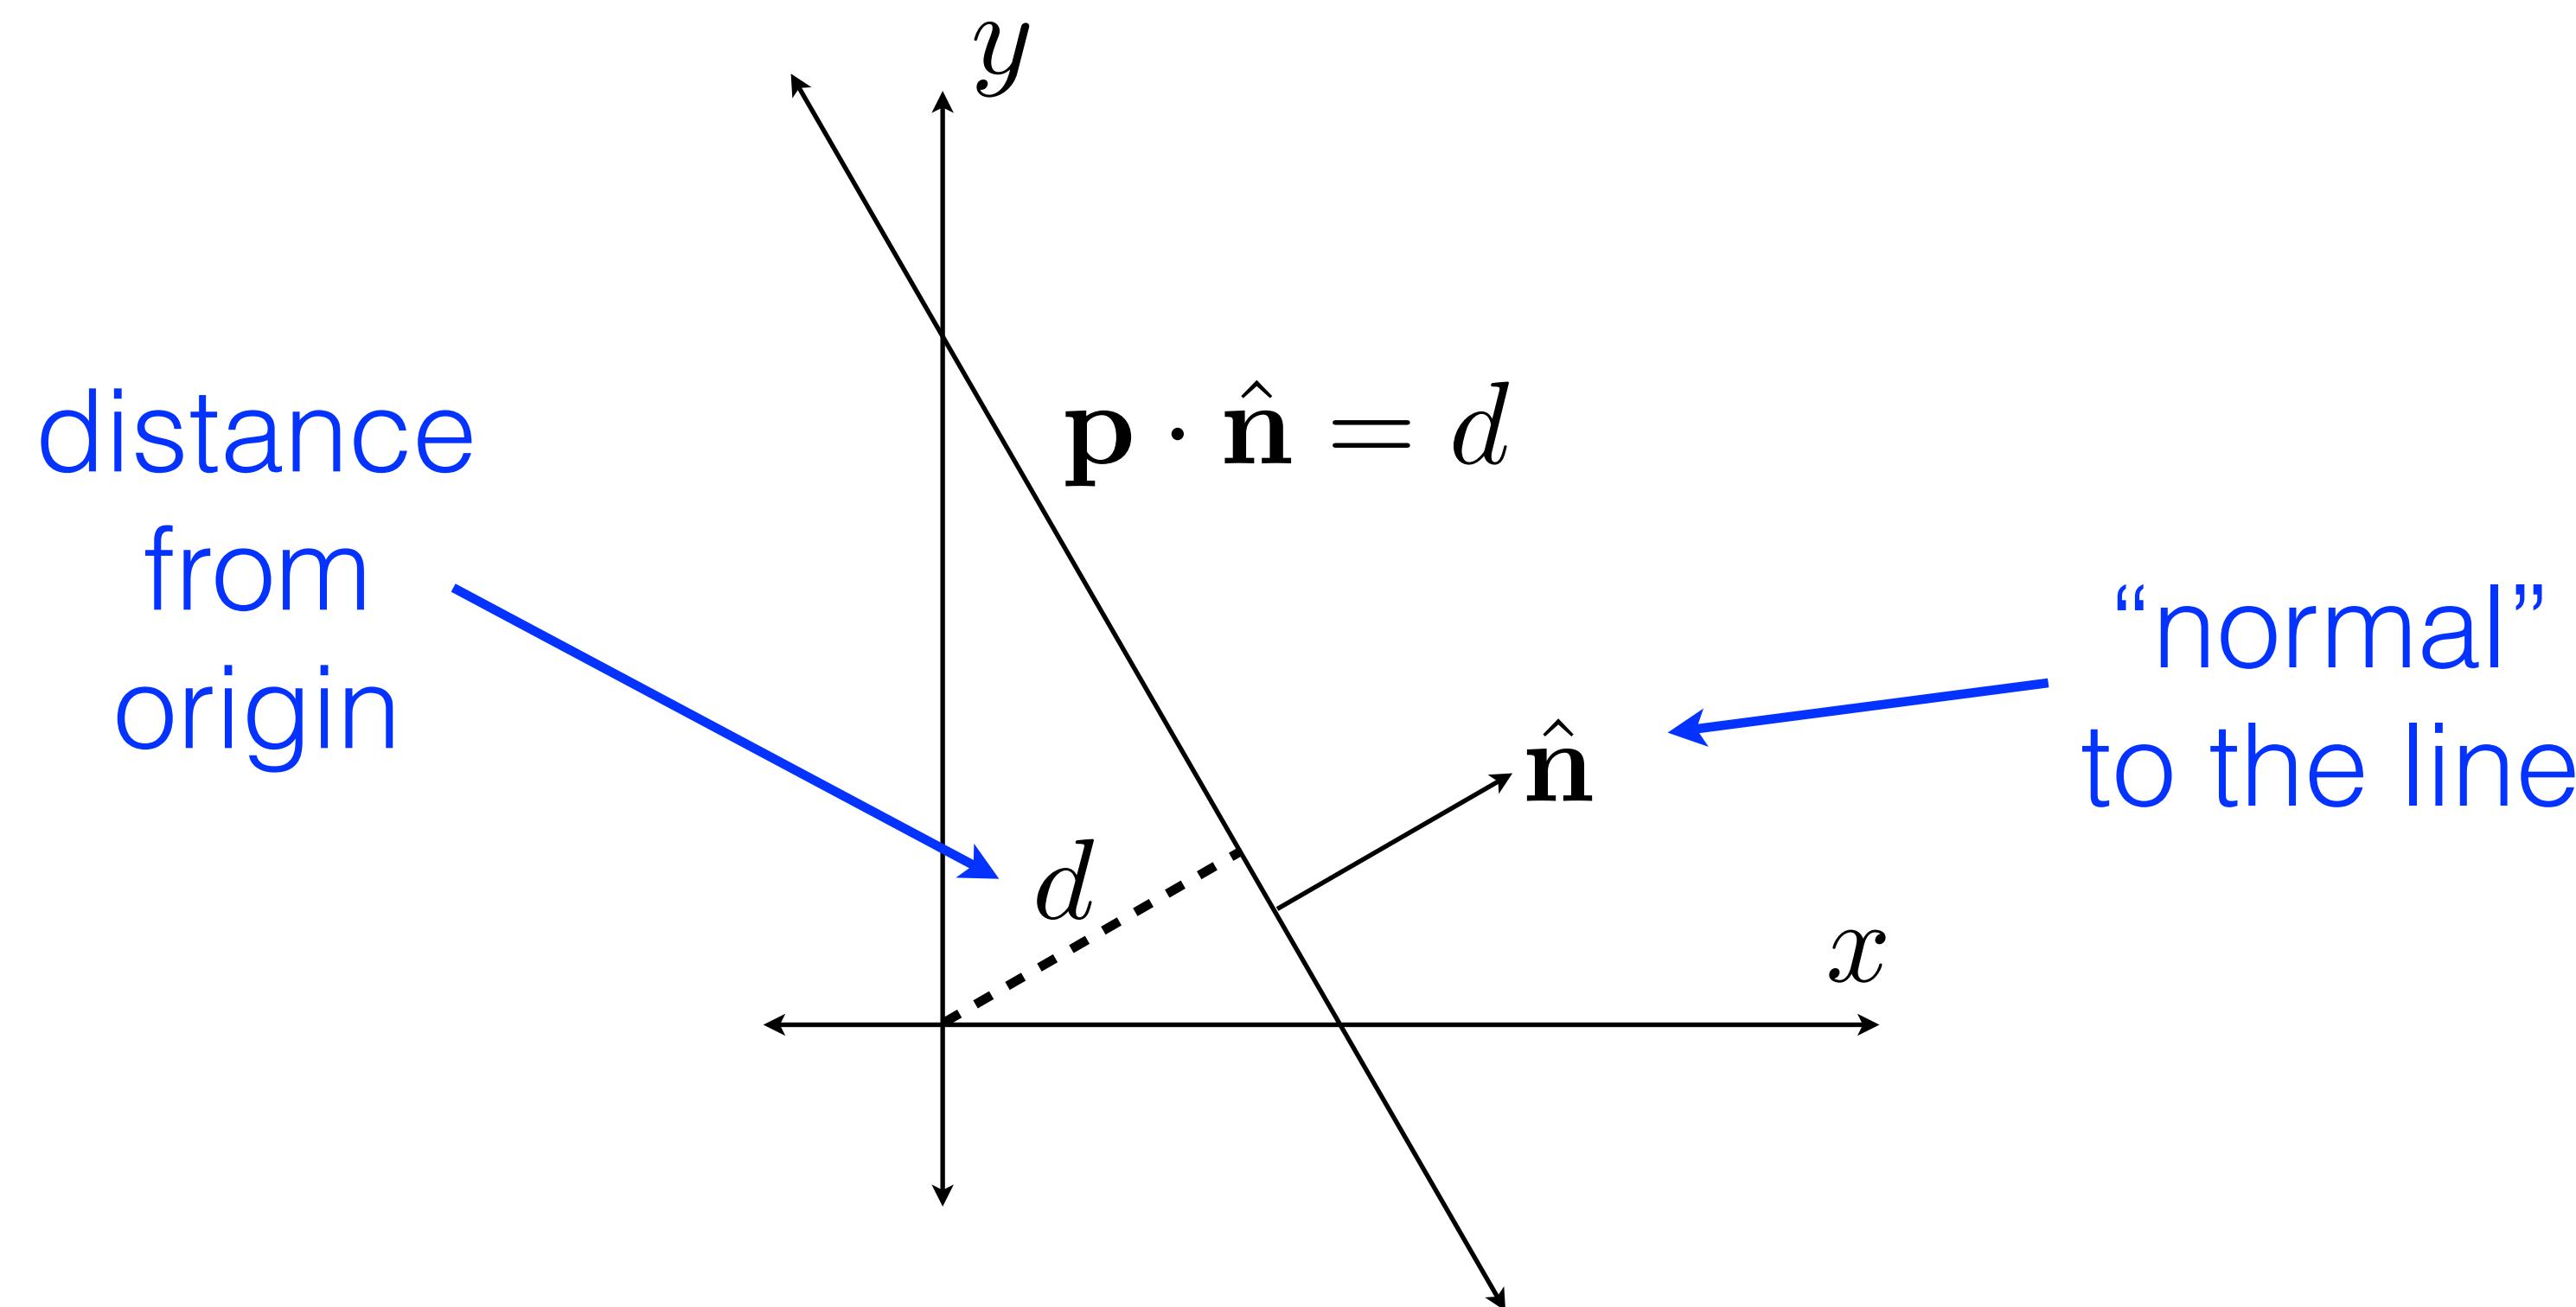
$$\mathbf{p}_0 + t \mathbf{d}$$

Line: $-\infty < t < \infty$

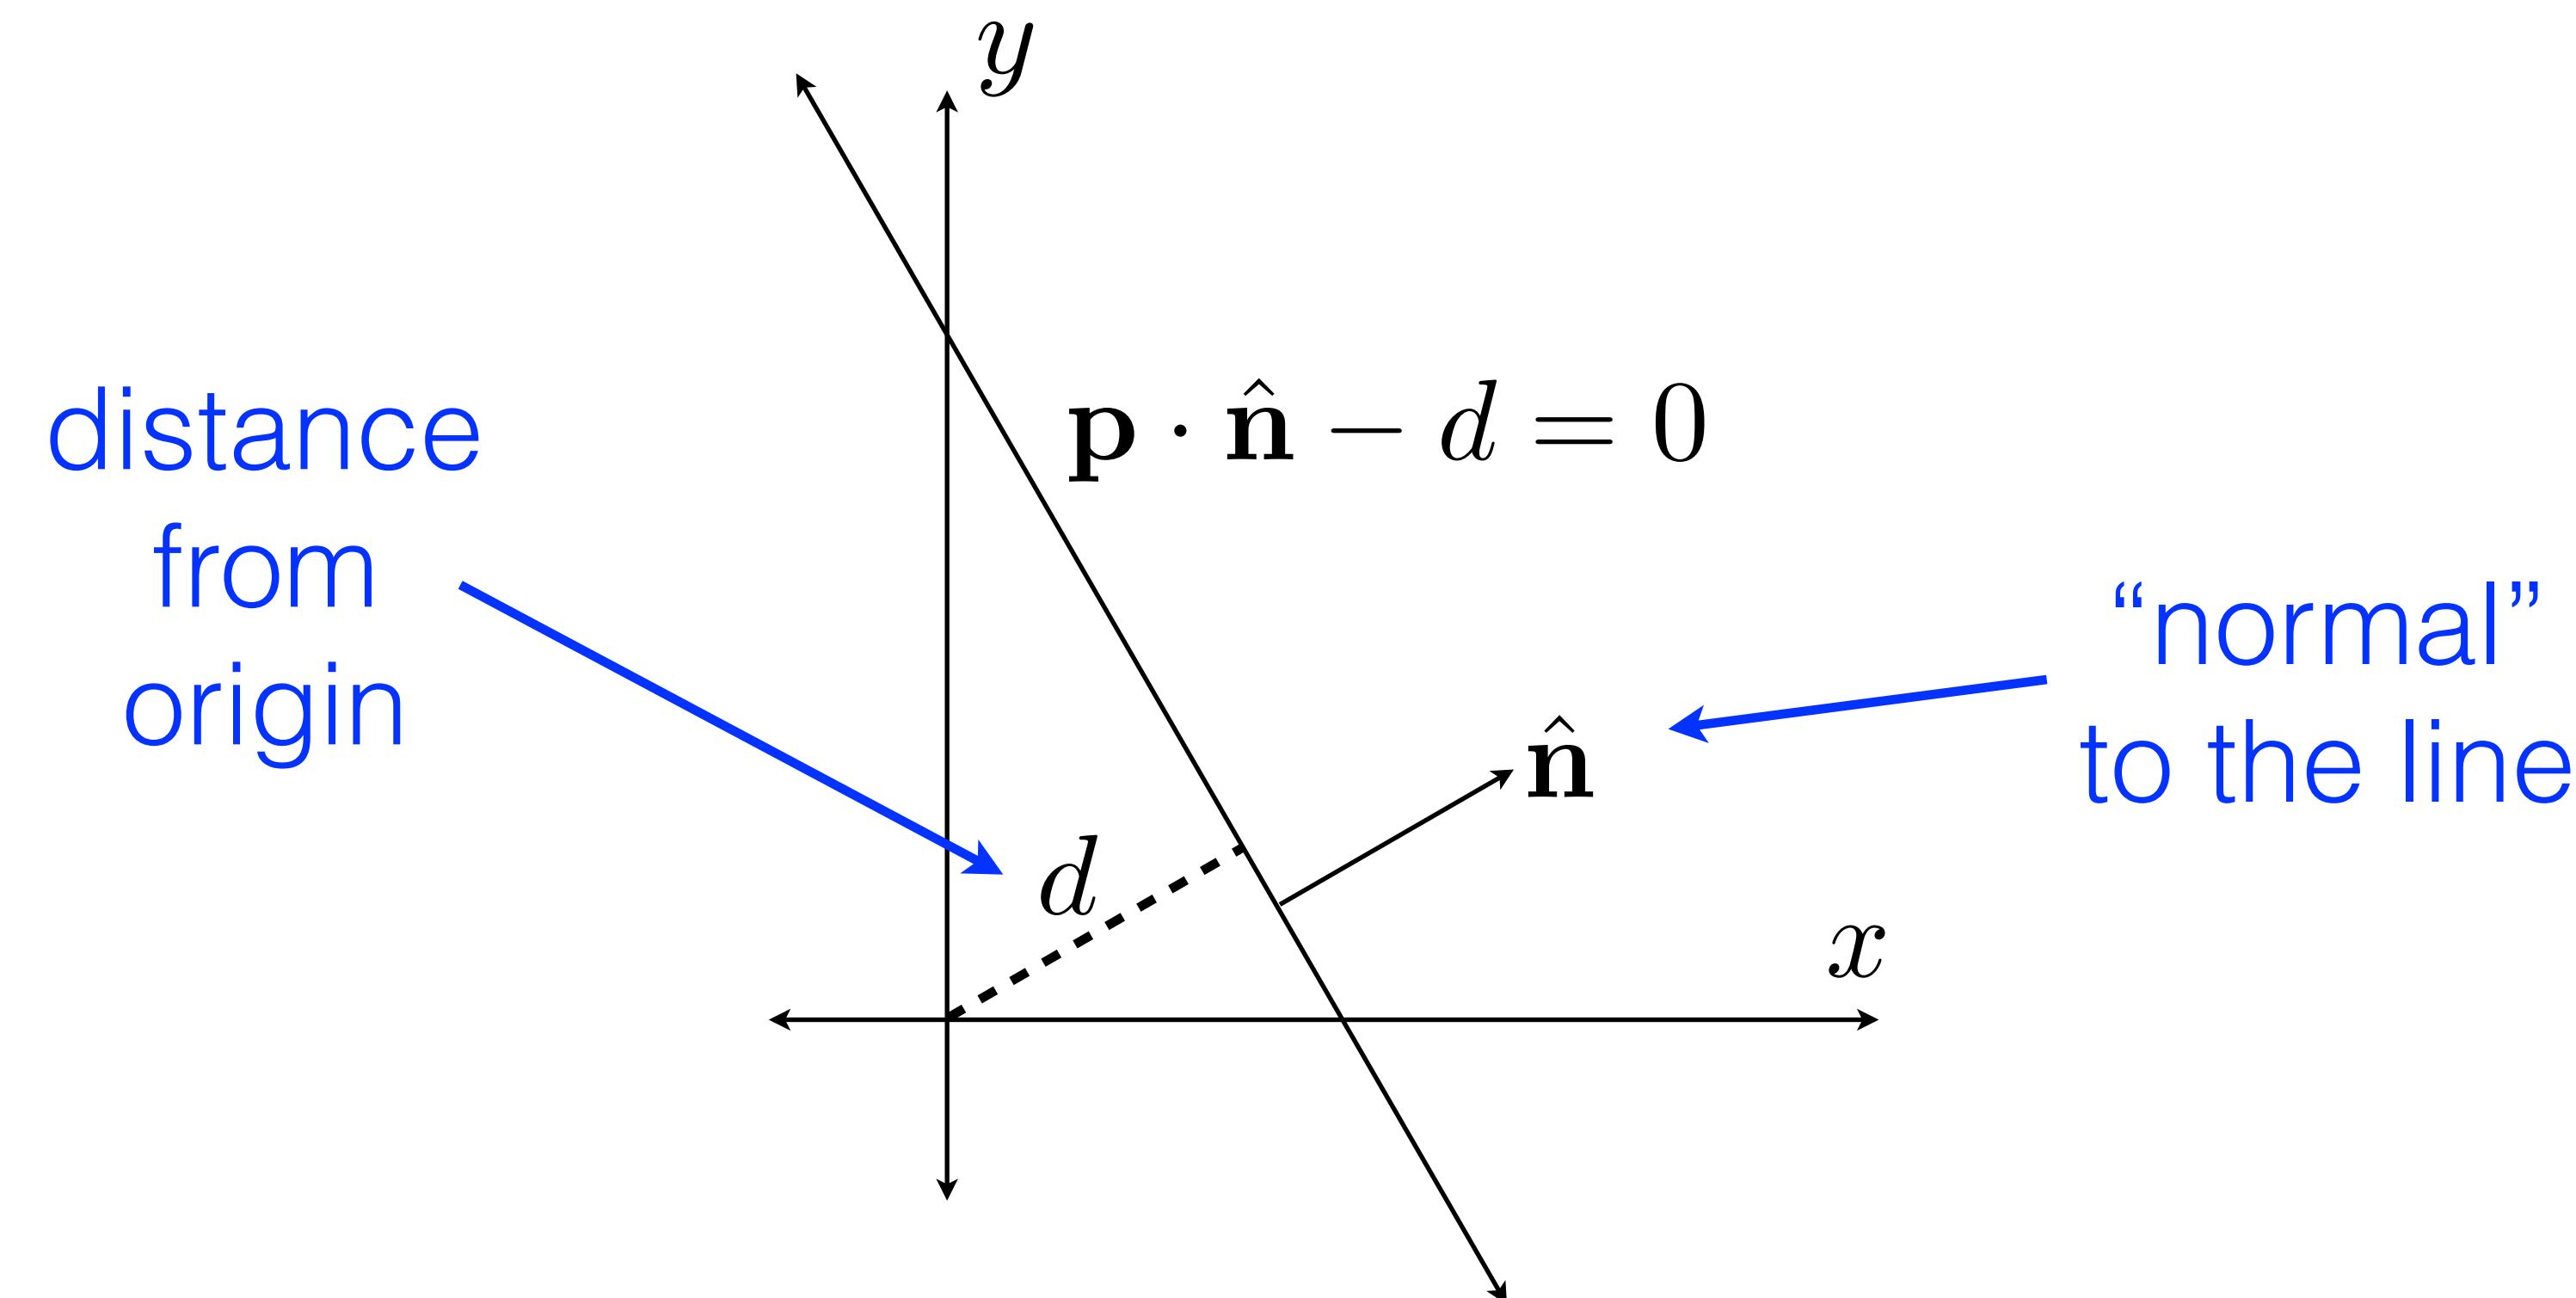
Line segment: $0 \leq t \leq 1$

Ray: $0 \leq t < \infty$

Normal + Distance



Implicit Representation



Aside: Lines and Distance

You've seen this before:

$$ax + by + c = 0$$

↑ ↑ ↑
normal negative
distance

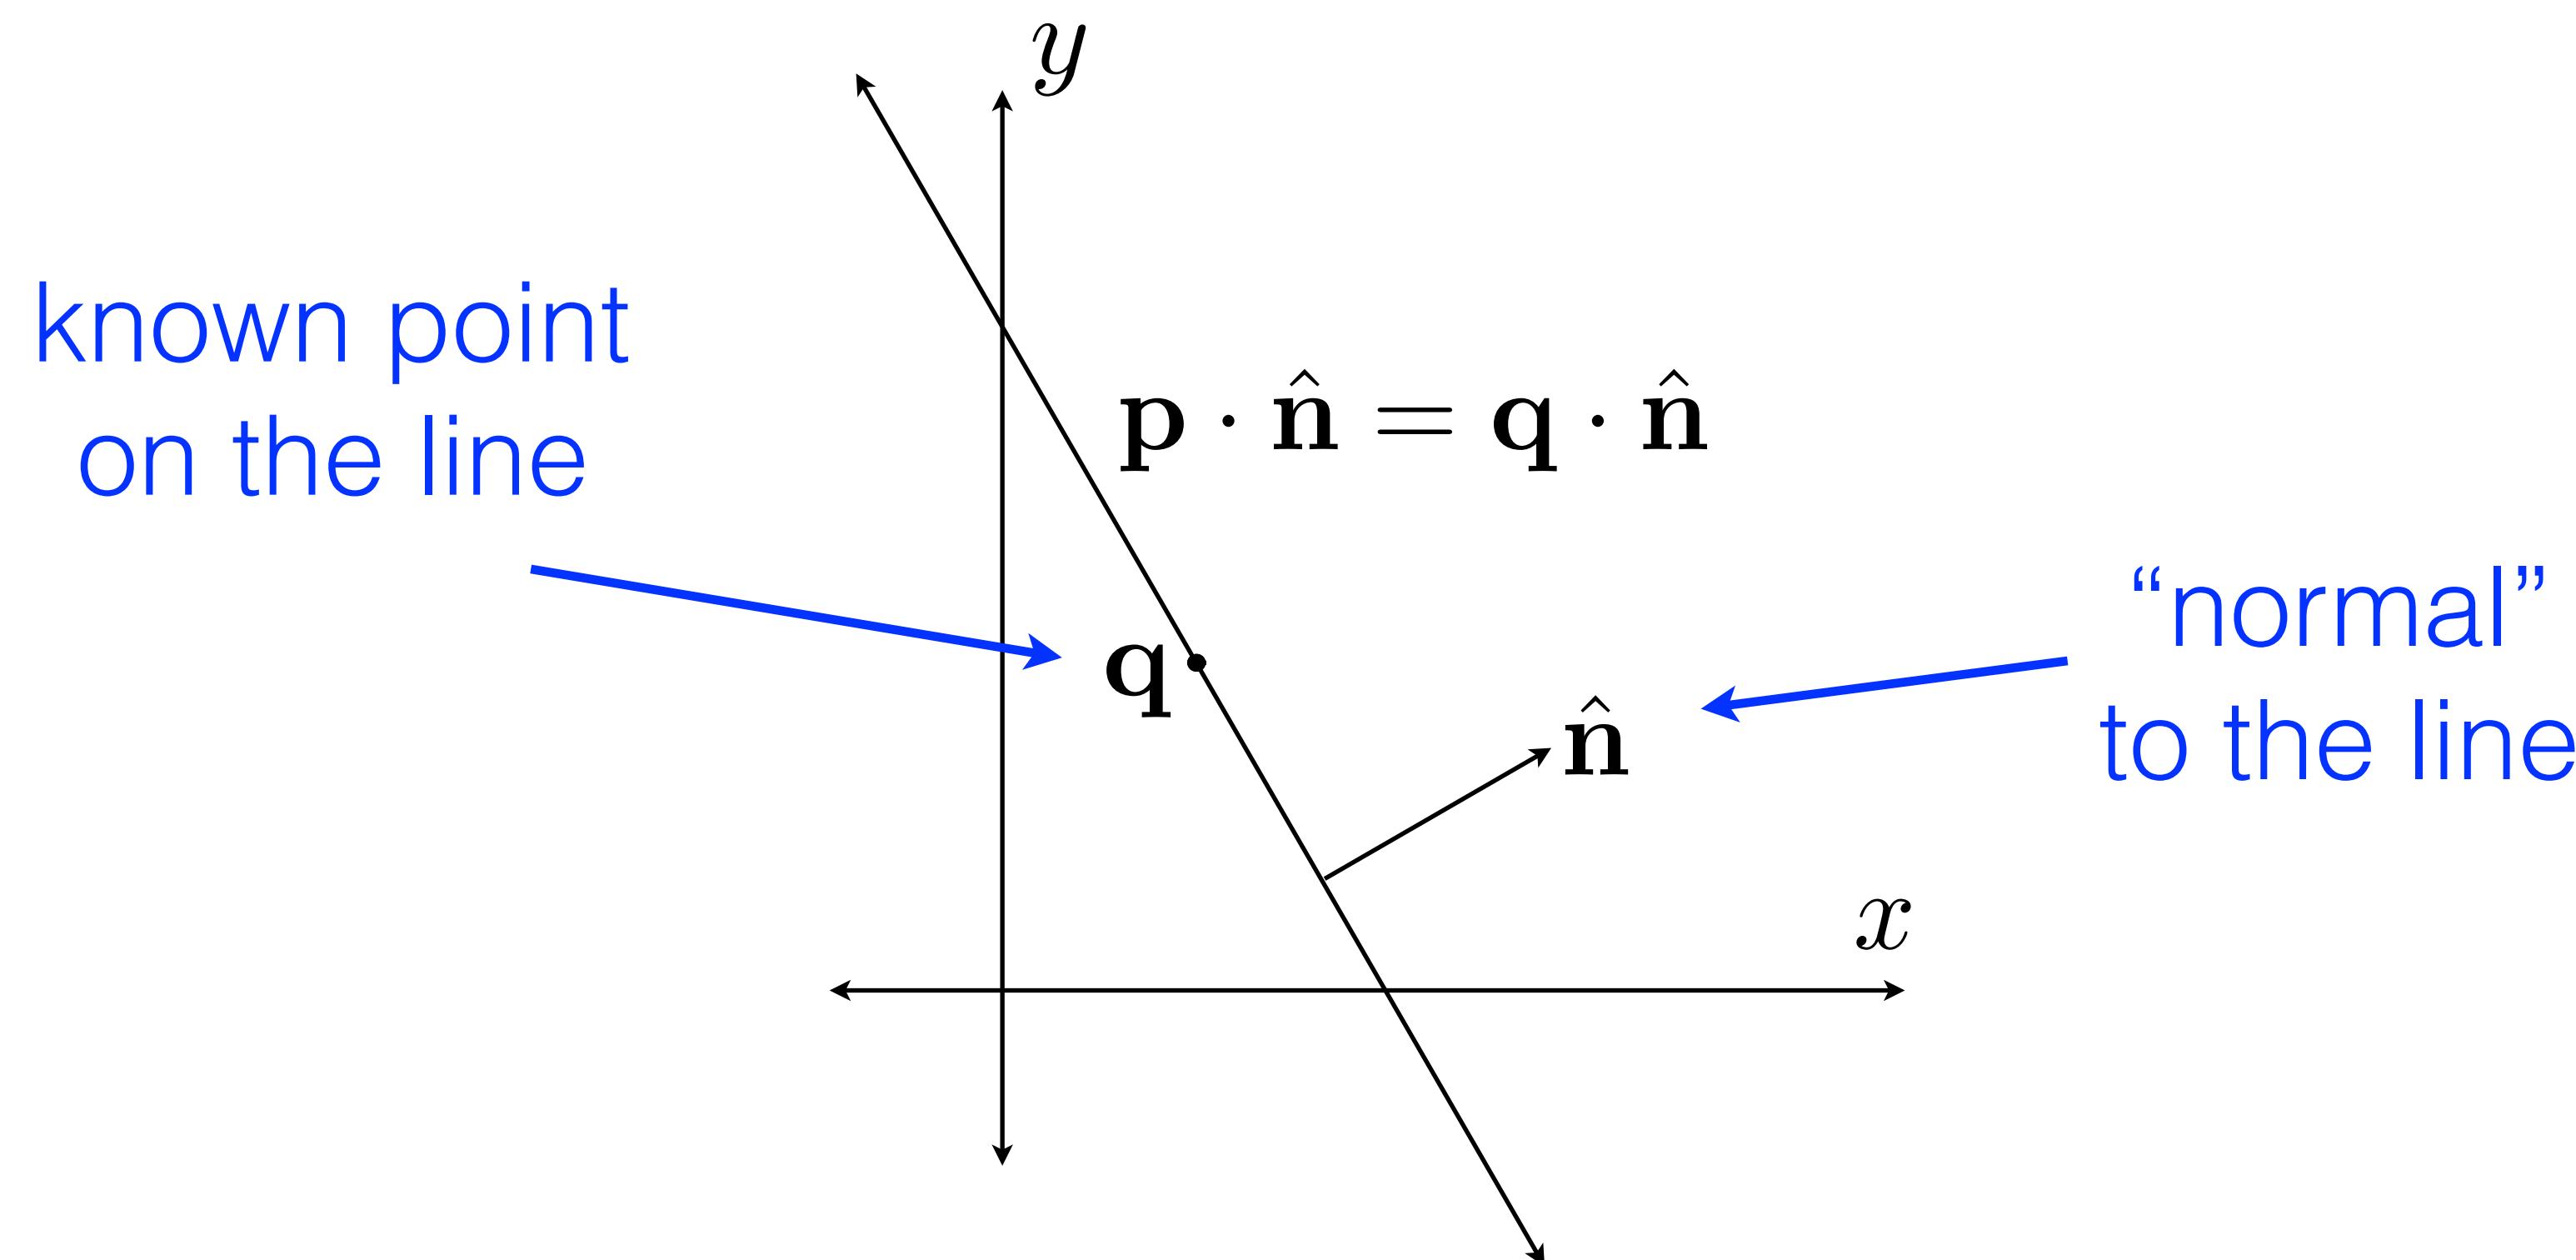
Normal and distance:

$$\mathbf{p} \cdot \hat{\mathbf{n}} = d$$
$$\mathbf{p} = \begin{bmatrix} x \\ y \end{bmatrix} \quad \hat{\mathbf{n}} = \begin{bmatrix} a \\ b \end{bmatrix}$$

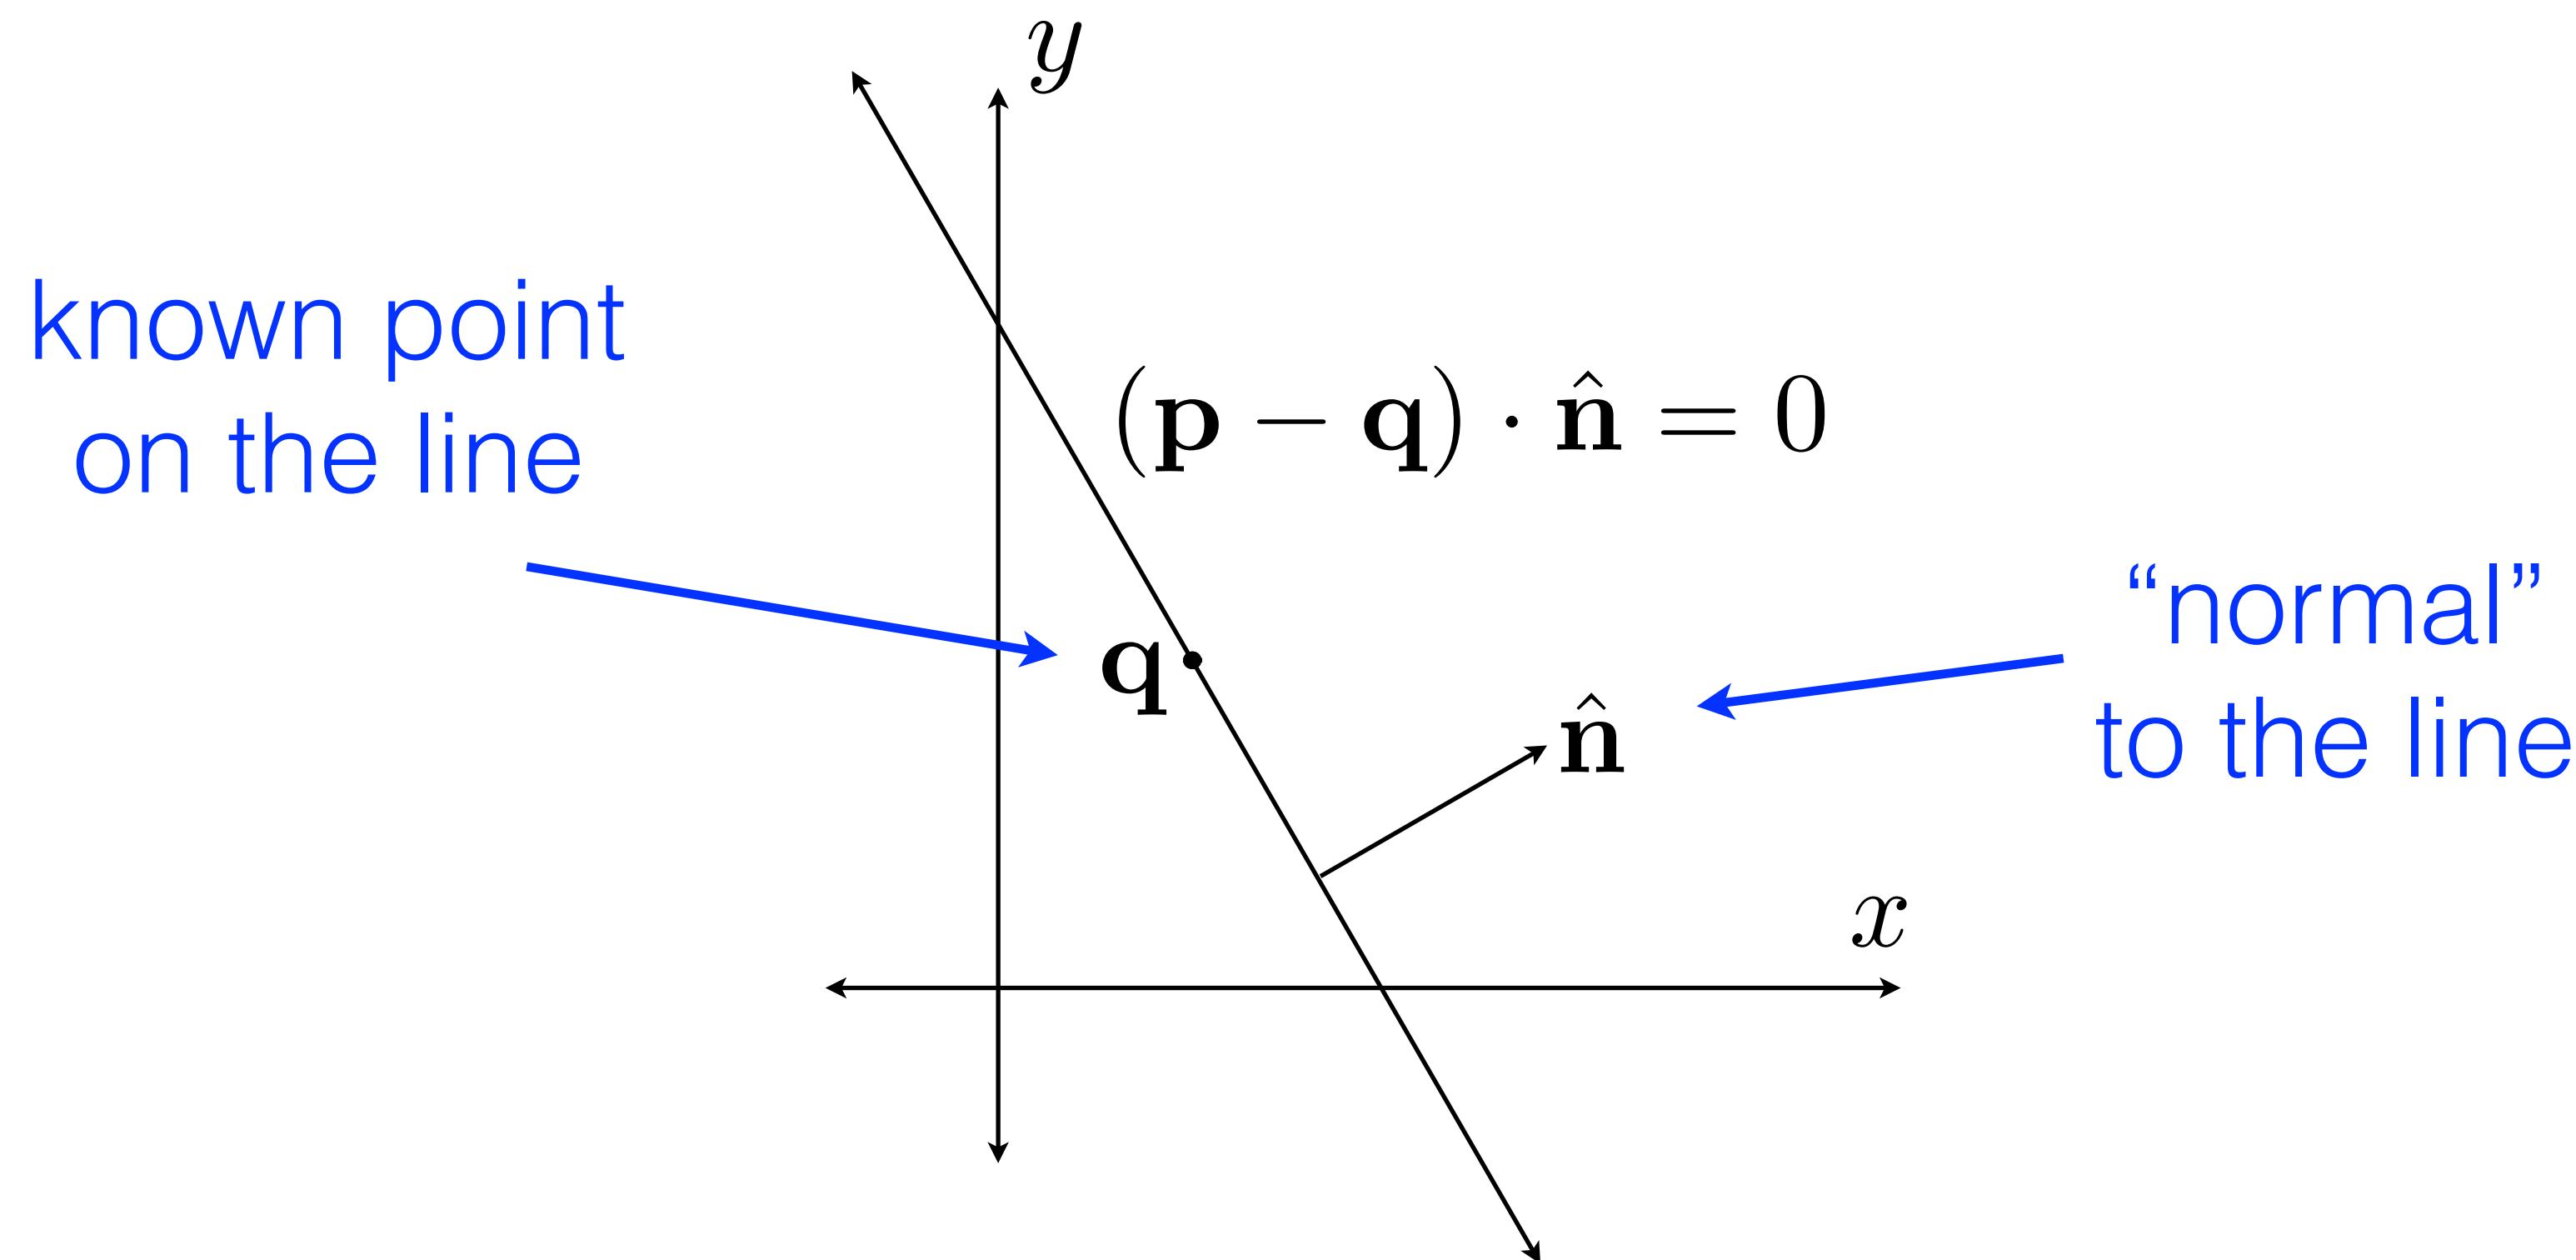
$$d = -c$$

This is really what you've seen years ago,
just in linear algebra form

Normal + Point



Implicit Representation



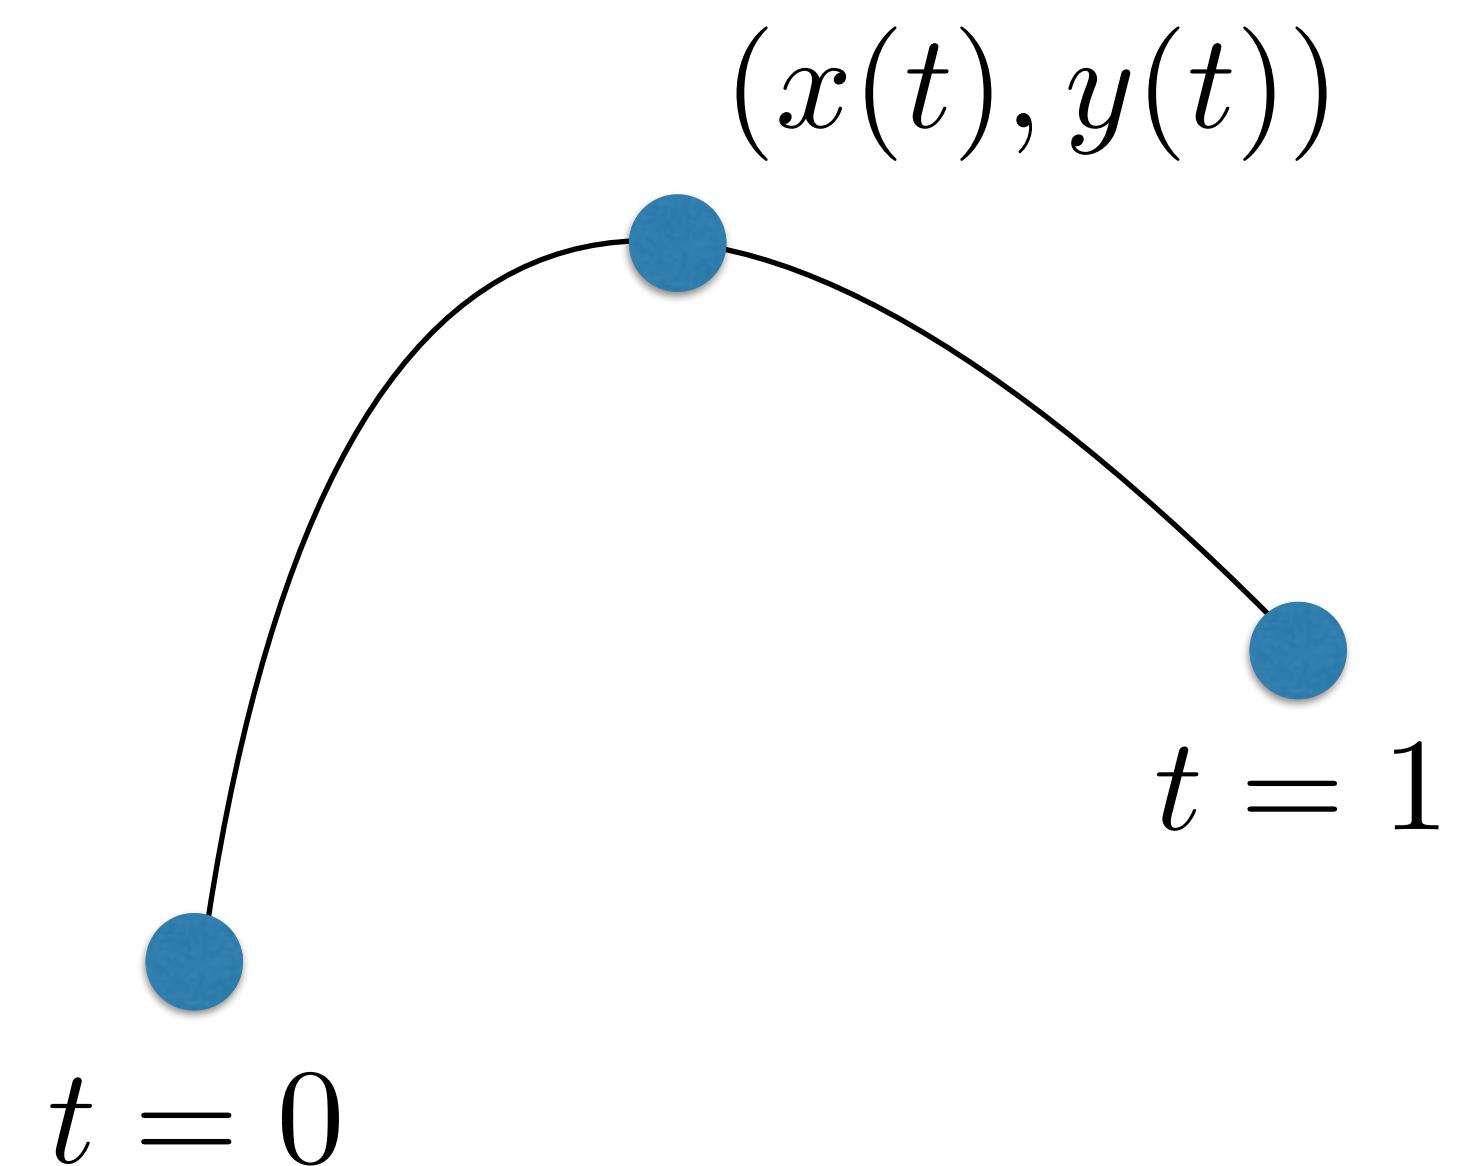
Representing Shapes

- Parametric
(sweeps out the shape as a function of some parameters)
 $\mathbf{p}(t)$
- Implicit
(meets some test)
 $f(\mathbf{p}) = 0$
- Others...

Can usually convert between representations

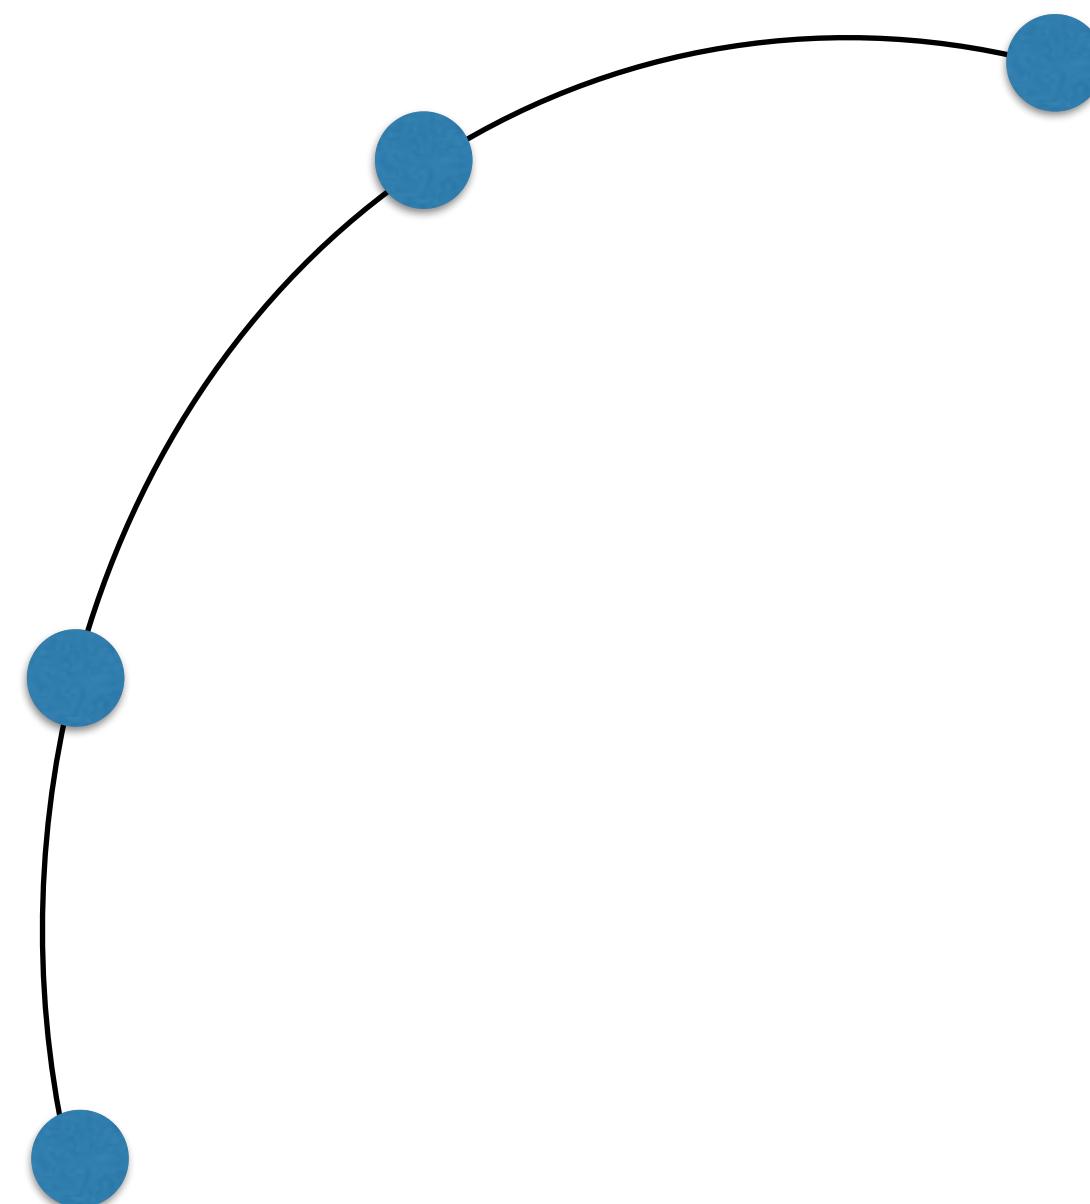
2D Curves

- Parametric curves
 - Parameter t traces the curve
 - One function for each dimension
 - Can extend to a space curve in any 3D or higher



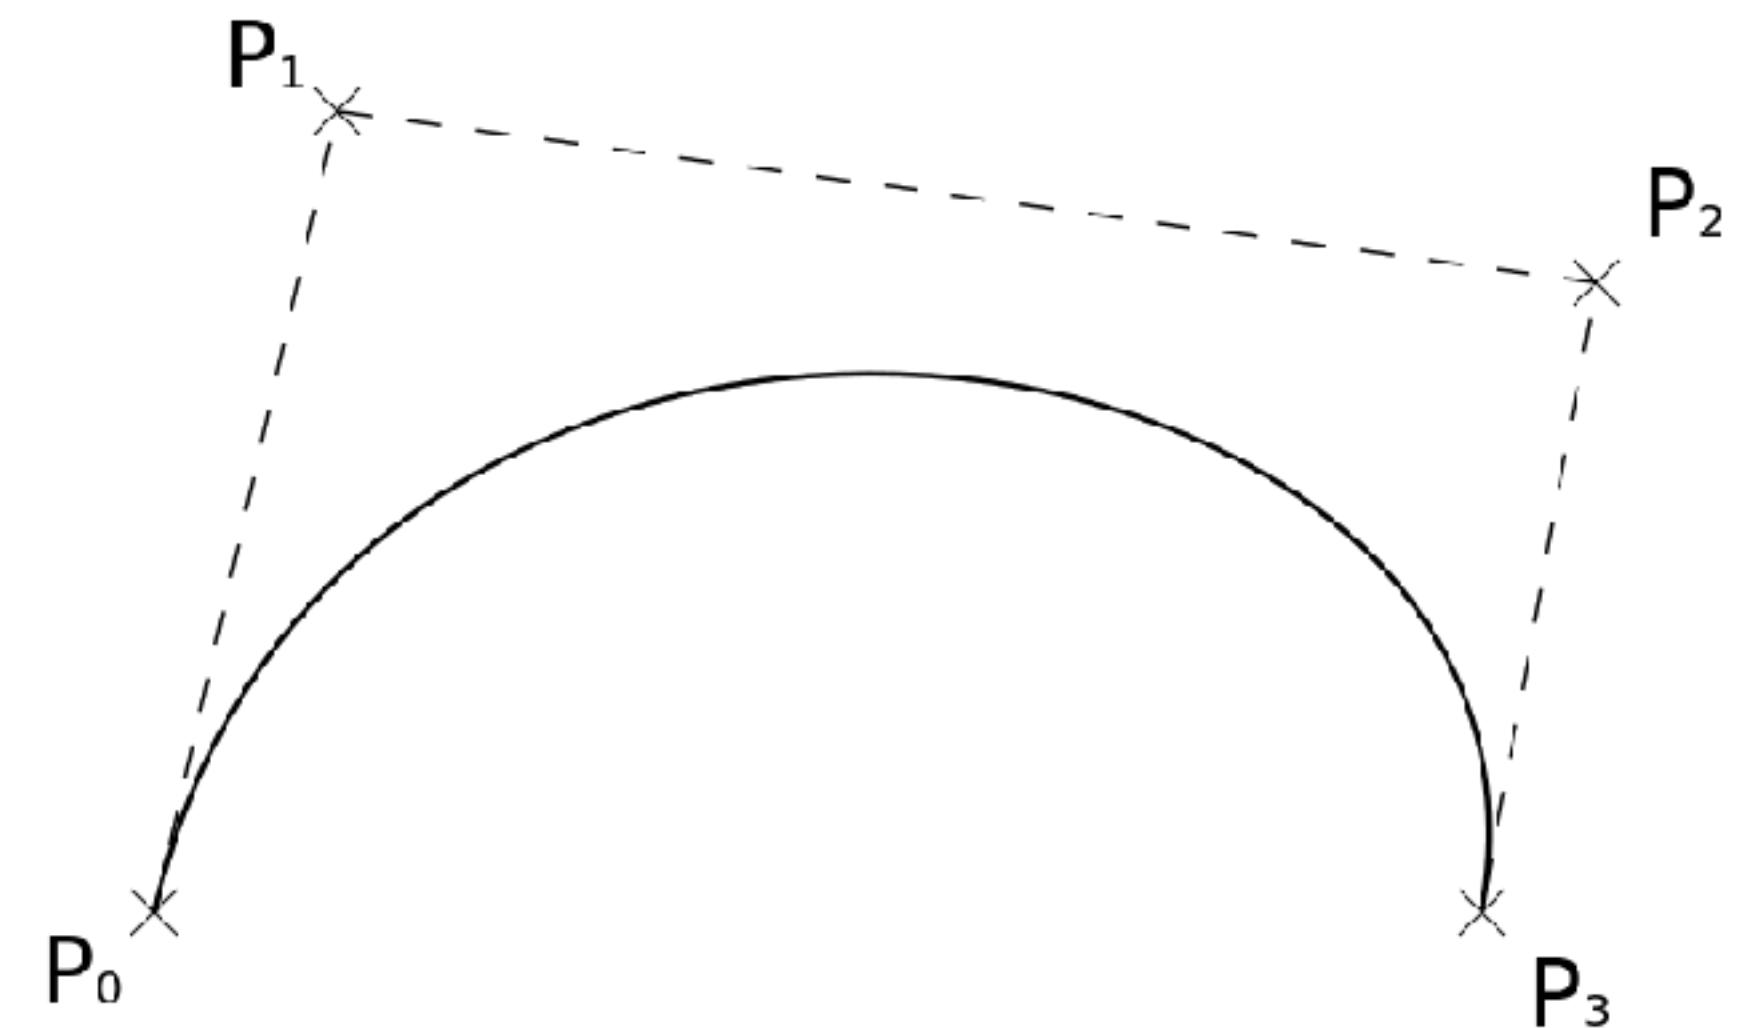
Cubic Curves

- Lots of types
 - Exact-fit curves (splines, etc.)
 - Bezier
 - B-splines
 - ...



Bezier Curves

- Specify *end points* and a set of *control points* in between
- Curve doesn't (usually) pass through the control points
- Control points determine tangents at certain points on the curve
- Common in 2D drawing programs

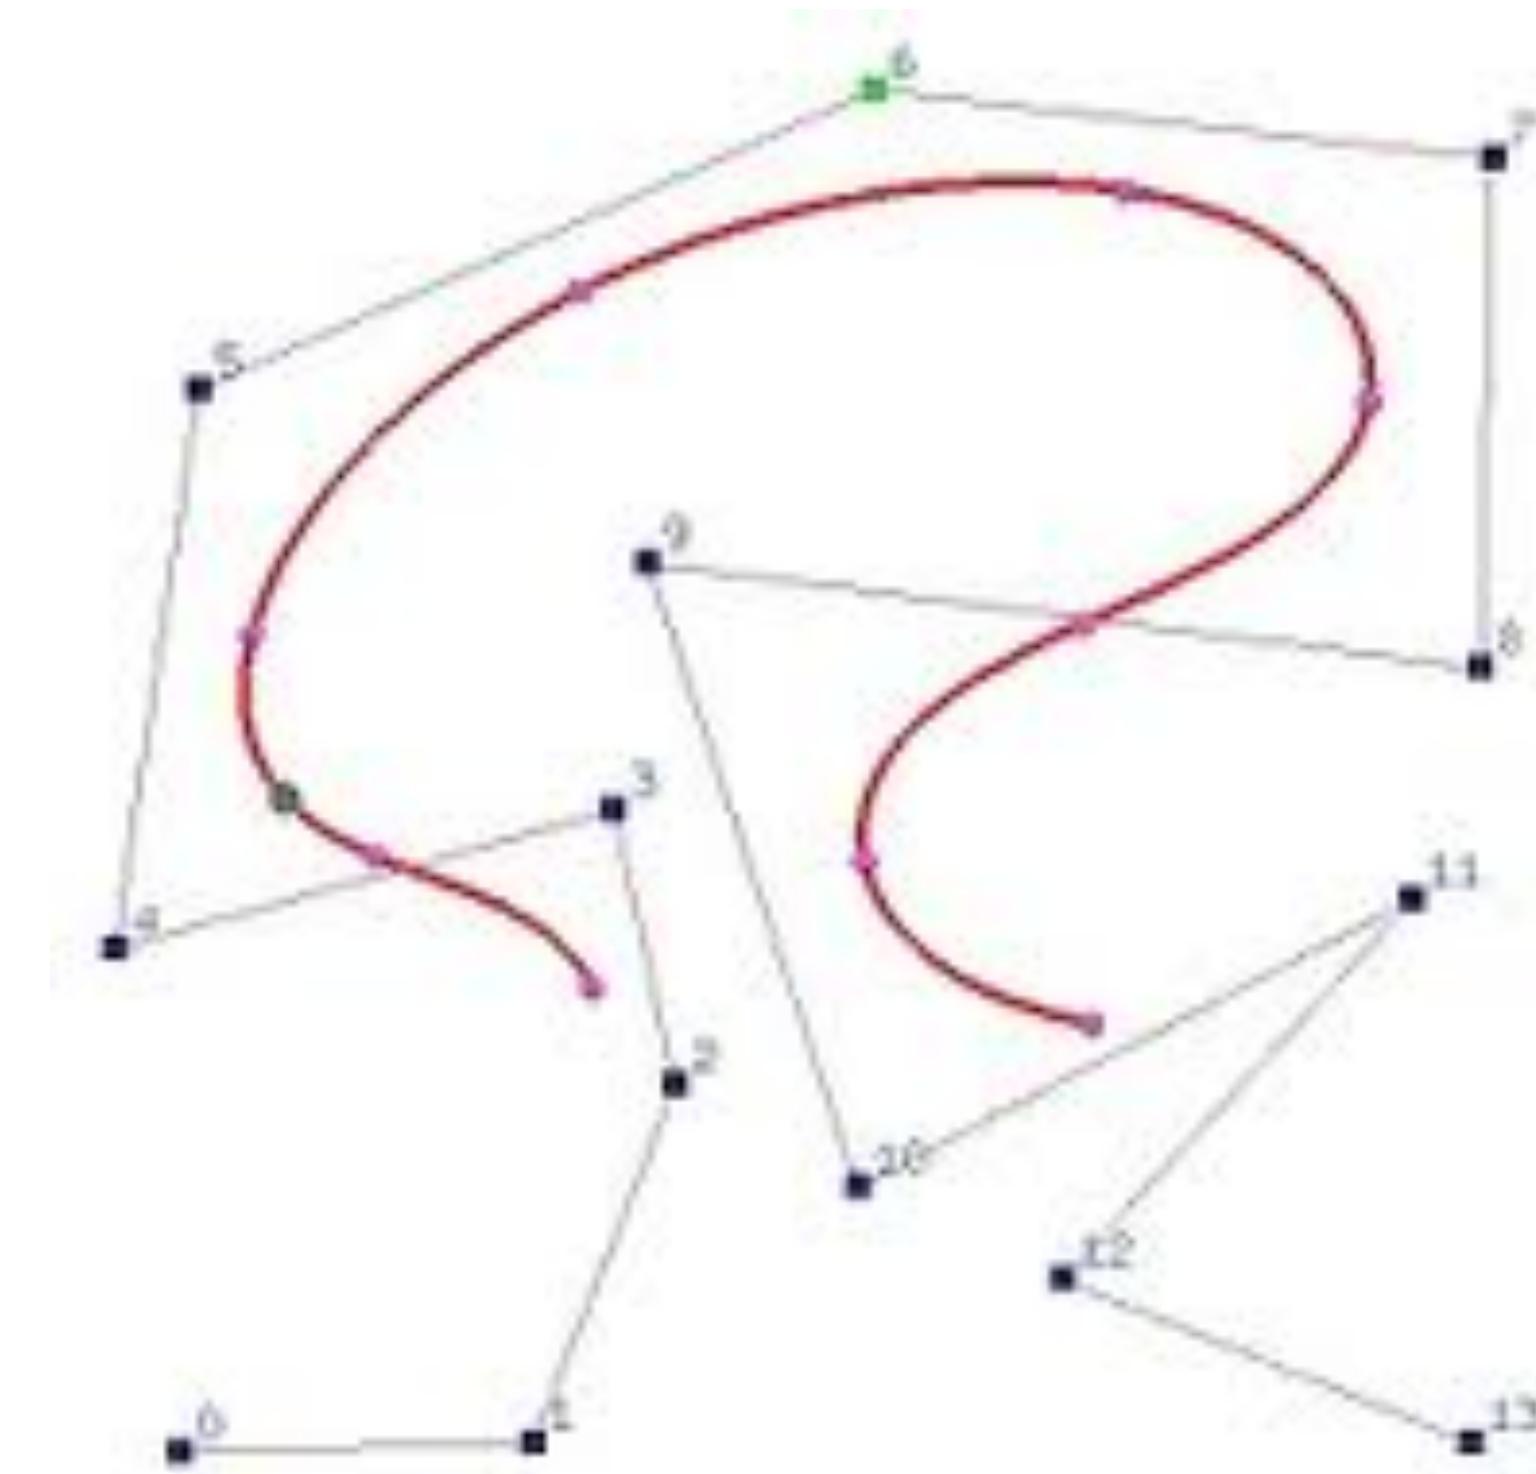


Complicated Curves

- Use higher-degree polynomials
 - Can make an n th-order Bezier curve out of two end points and $n-1$ control points
 - Cumbersome and slow for long curves
 - Non-local control!
- Piecewise cubic segments
 - Can get 0th-order continuity by sharing endpoints
 - Can get 1st-order continuity by aligning control points

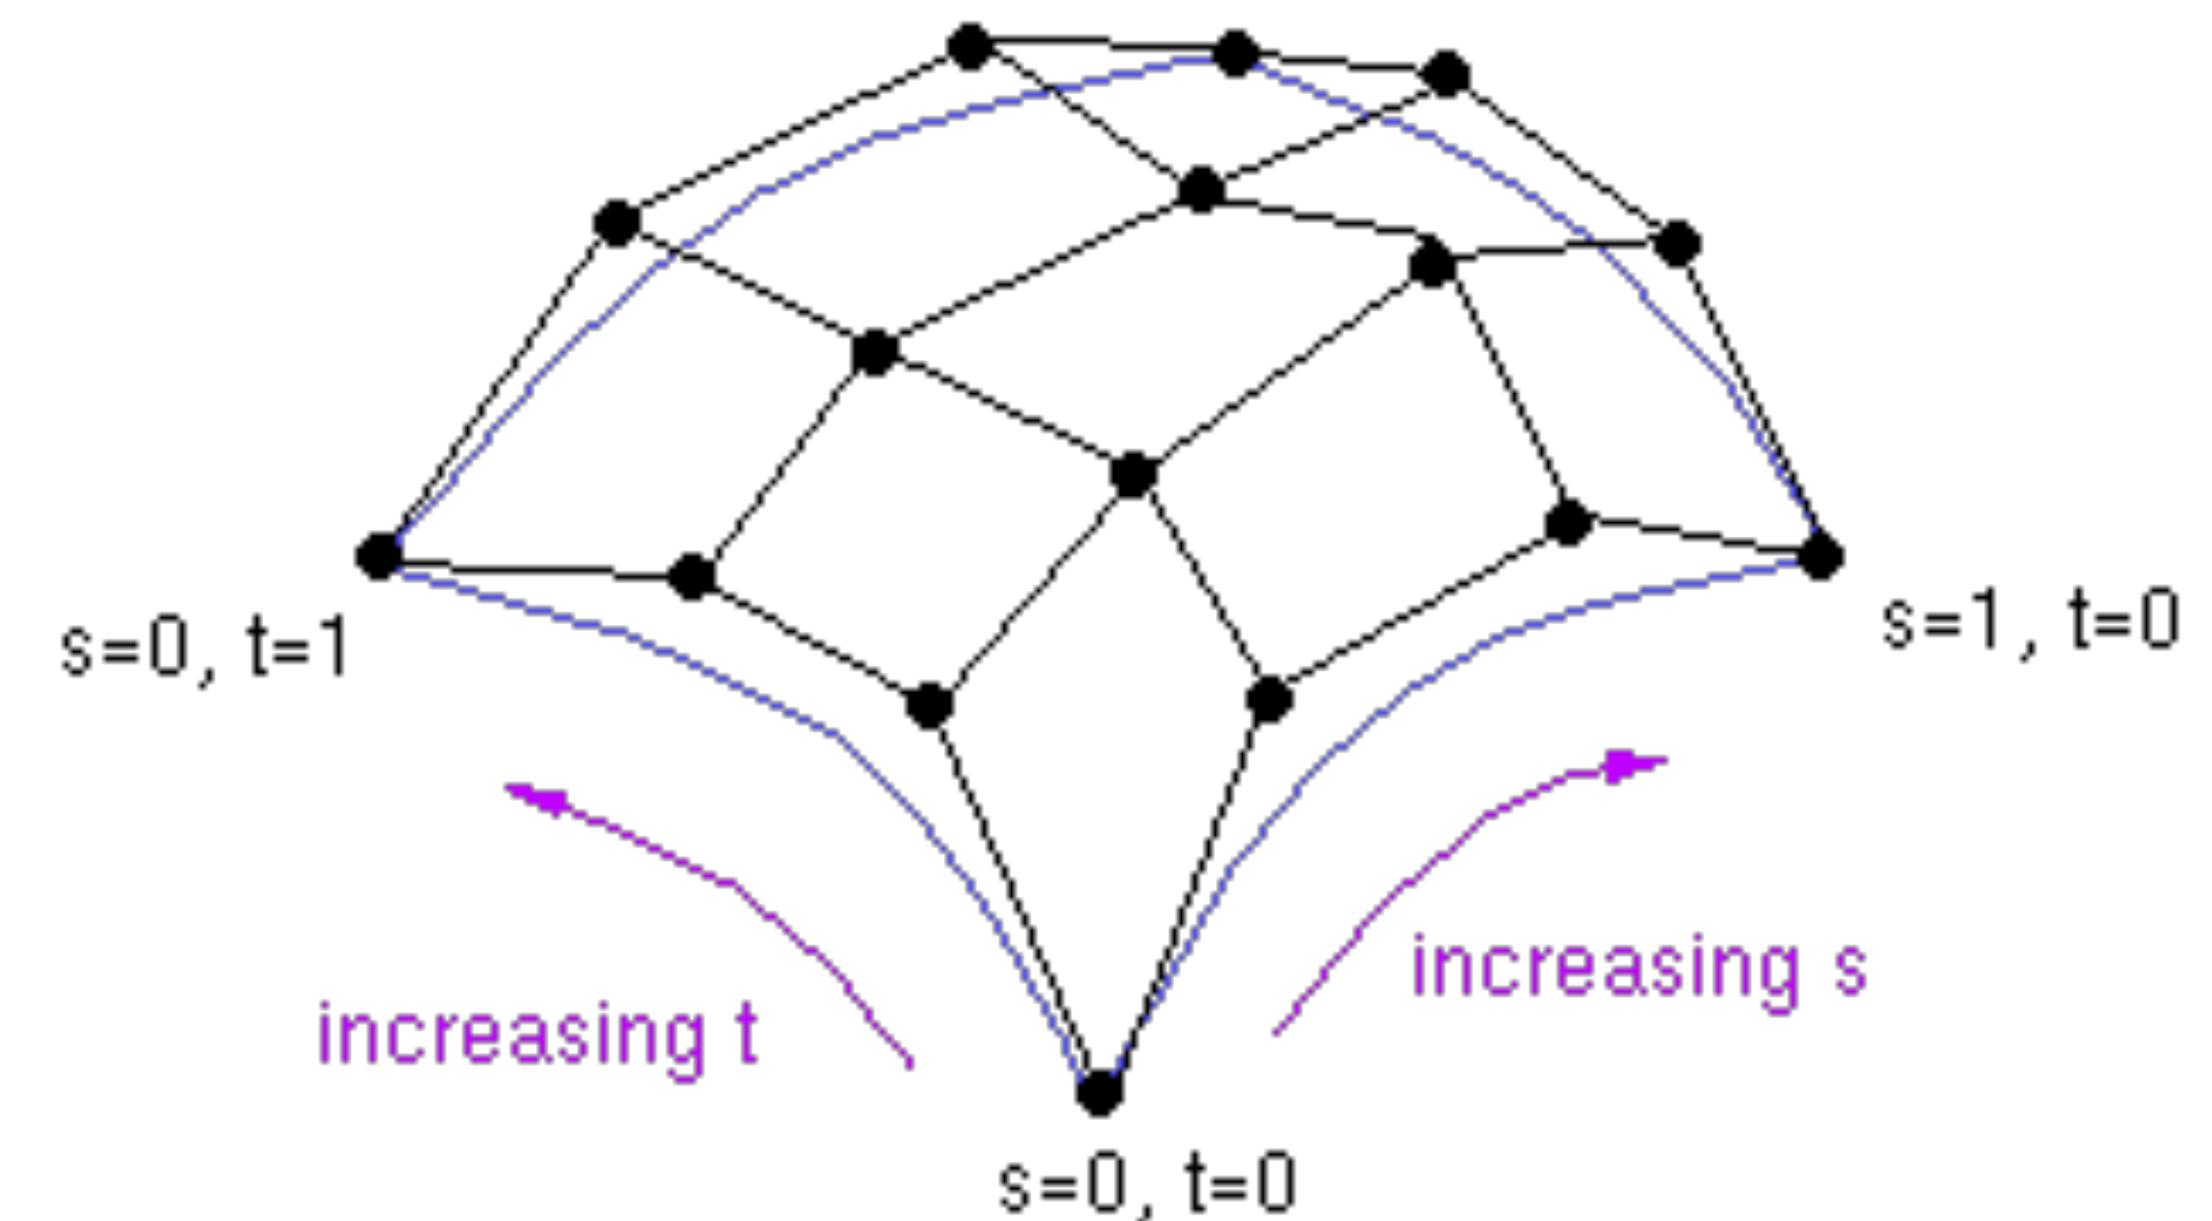
B-Splines

- More common to use *basis splines* (B-splines)
- Each local portion of the curve is a *weighted blend* of nth-order basis functions (usually cubic)



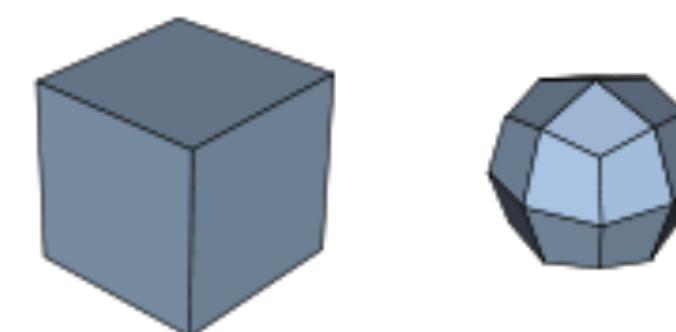
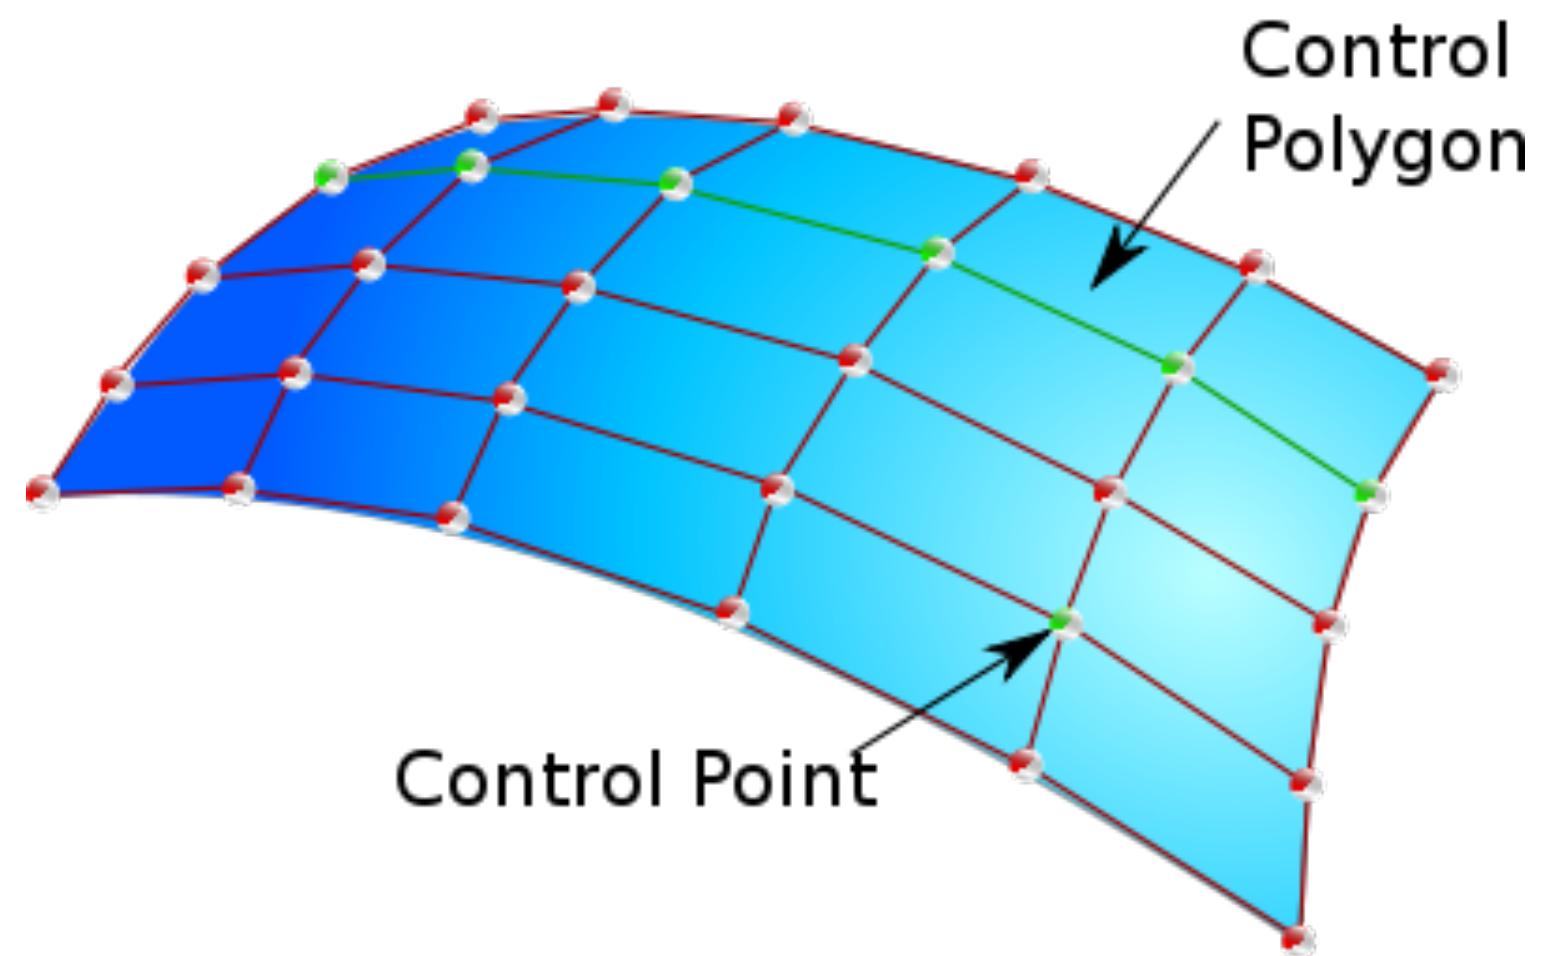
3D Surfaces

- Extend idea of Bezier or B-spline curves to a mesh or grid of control points
- Two parameters across surface: s, t
- Similar in principle to bicubic interpolation



3D Surfaces

- Bezier
- B-spline
- Non-Uniform Rational B-splines (NURBs)
- Catmull-Rom and other spline variations
- Subdivision surfaces



See CS 455 for more...

Coming up...

- Geometric tests

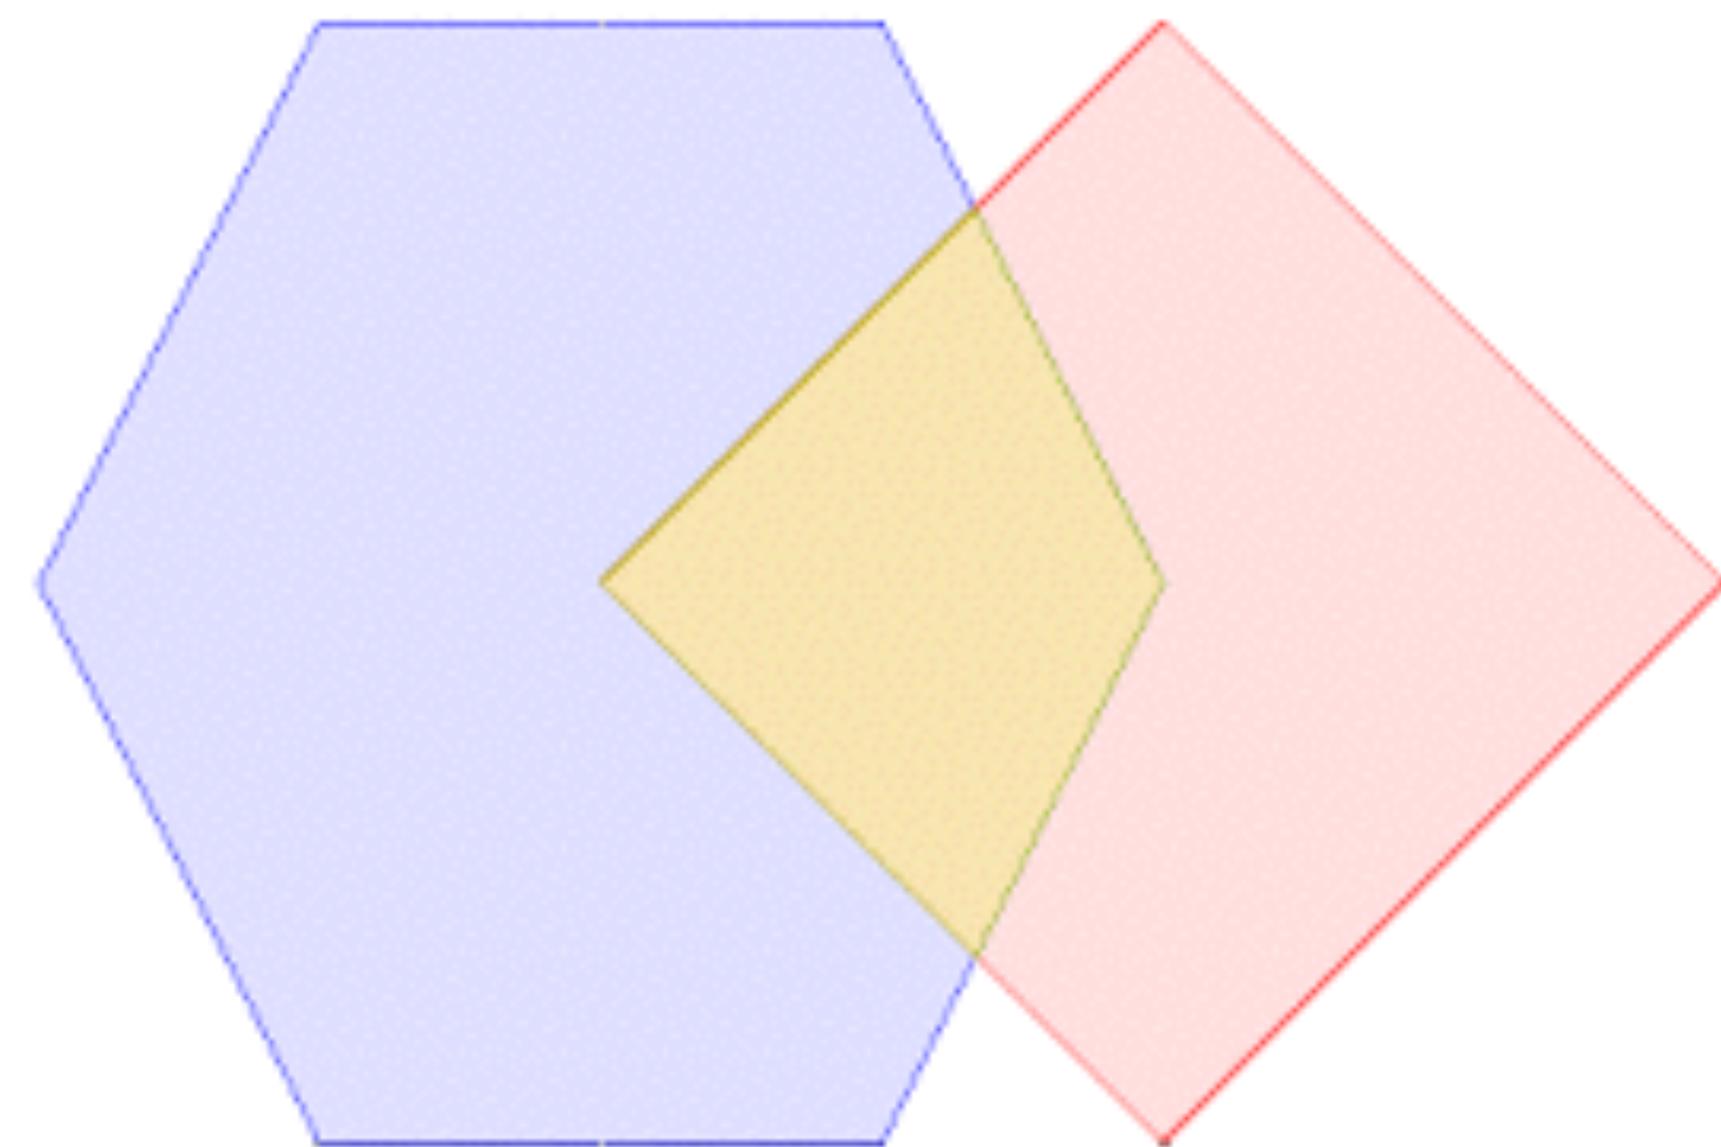


Geometric Tests

CS 355: Introduction to Graphics and Image Processing

Geometric Tests

- Lots of things in graphics involve making geometric tests:
 - Interactive selection
 - Intersections
 - Collisions (e.g., games)
 - Ray intersections
 - Nearest points
 - ...



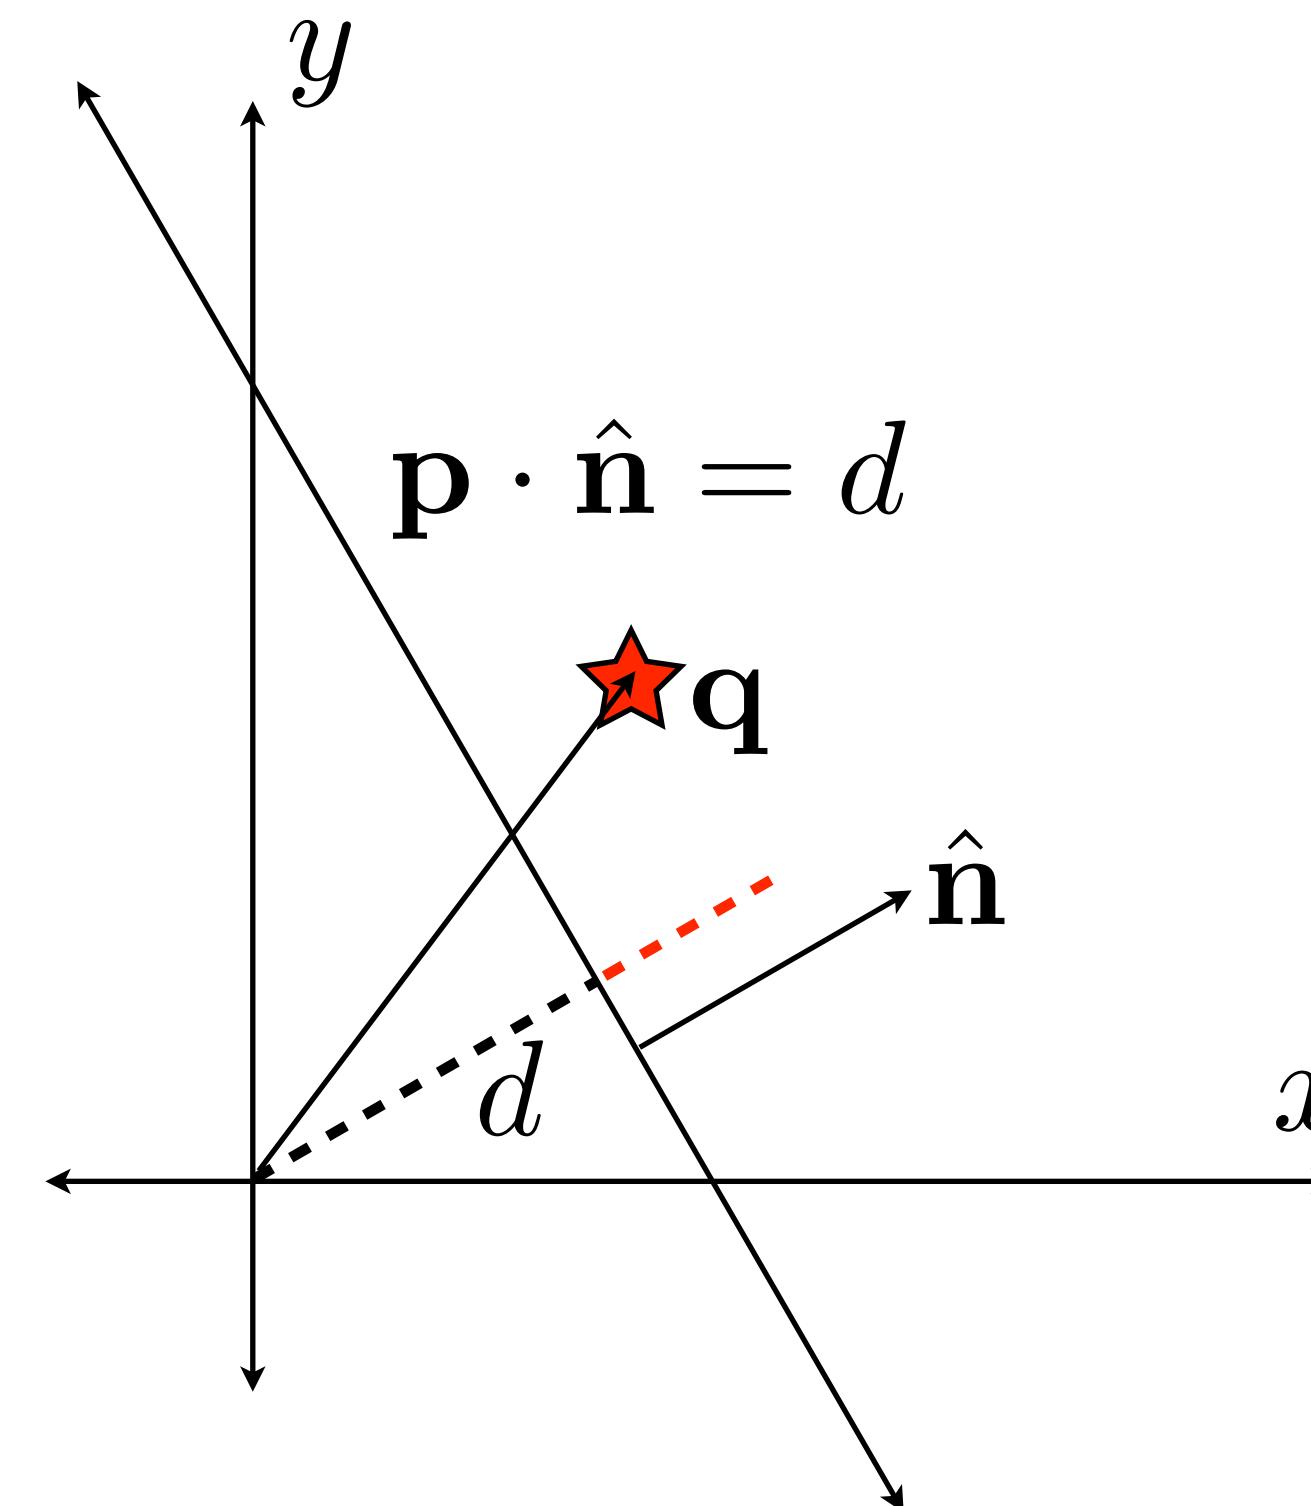
Distance to a Line

- Points \mathbf{p} on the line \mathbf{L} satisfy this constraint:

$$\mathbf{p} \cdot \hat{\mathbf{n}} - d = 0$$

- Distance from point \mathbf{q} to the line \mathbf{L} :

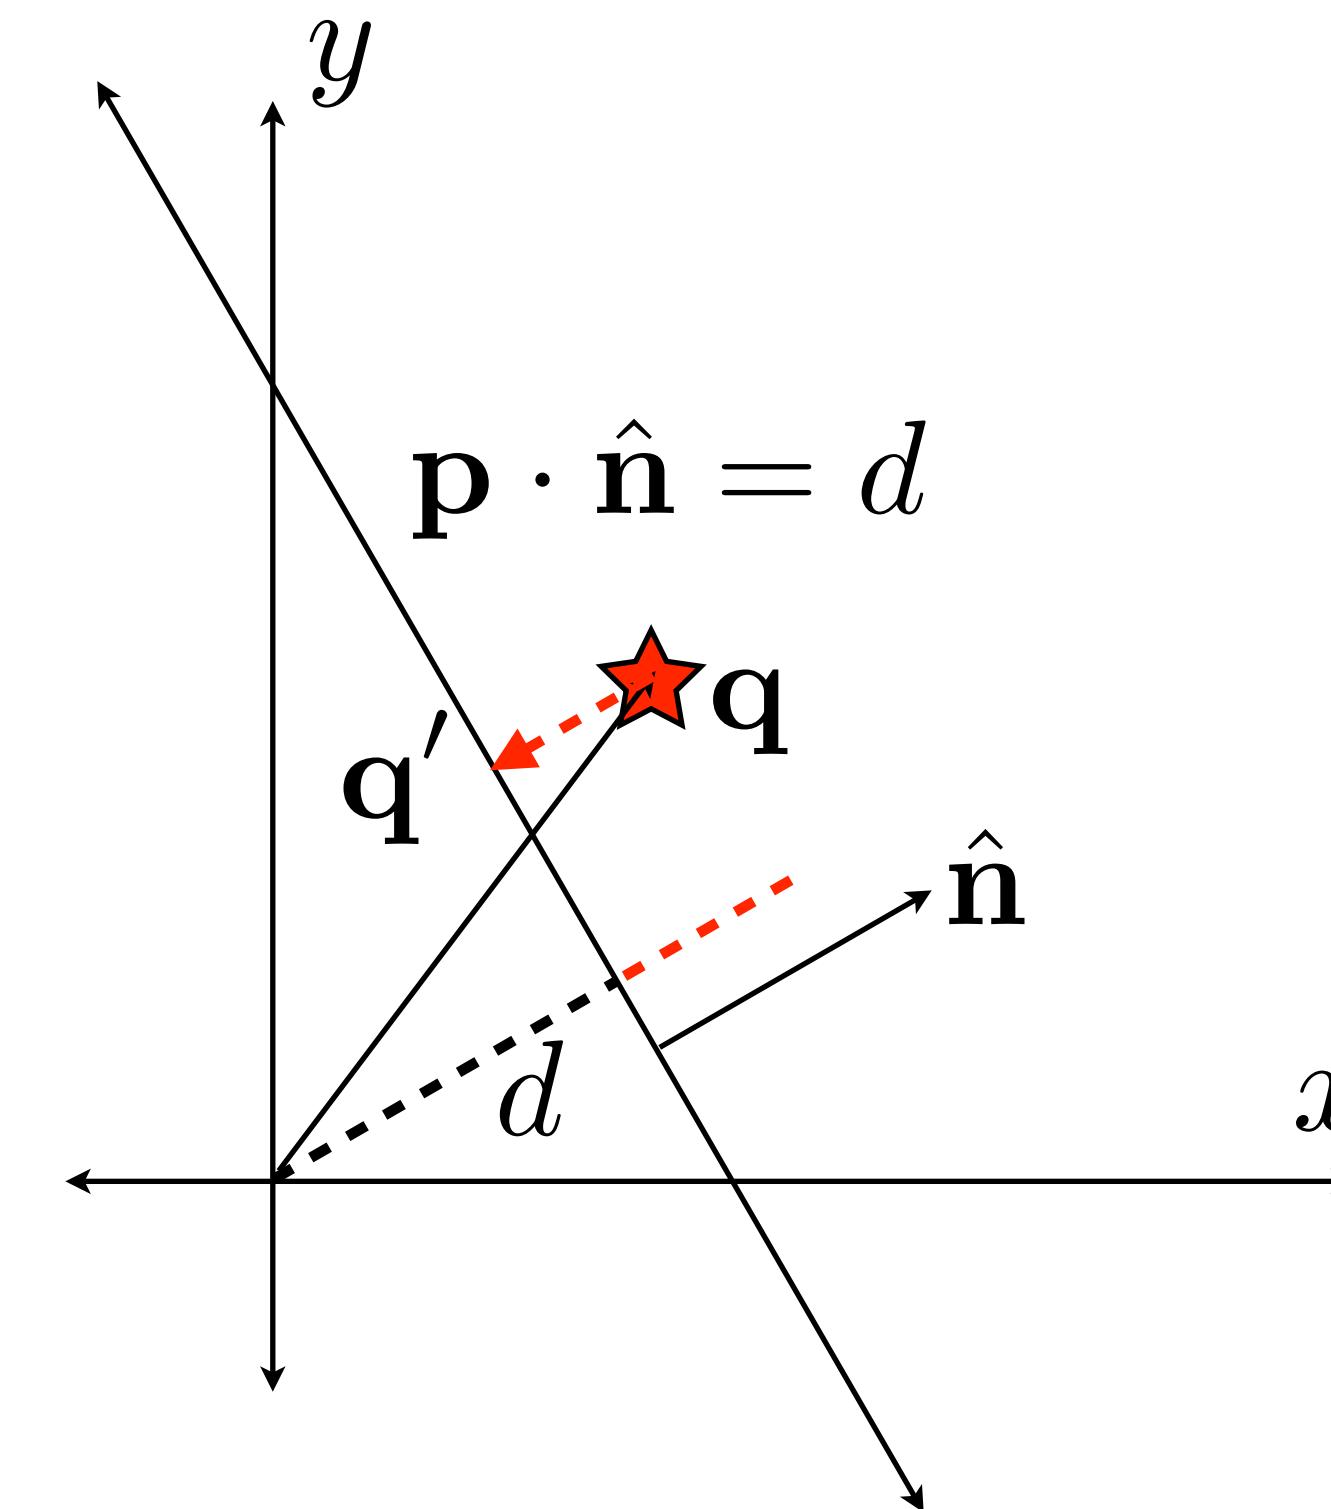
$$|\mathbf{q} \cdot \hat{\mathbf{n}} - d|$$



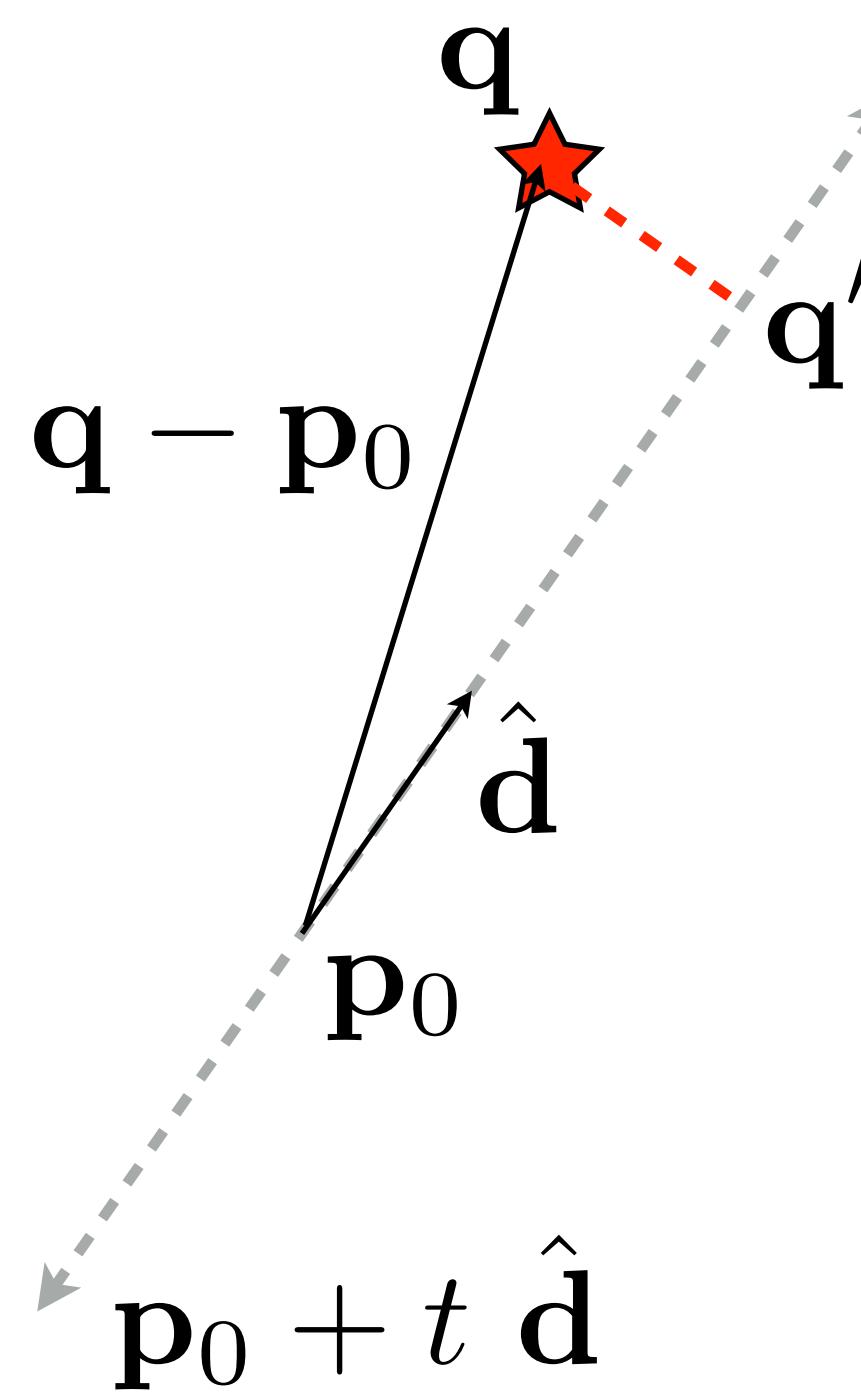
Closest Point to a Line

- To get the closest point on the line \mathbf{L} to point \mathbf{q} , go back along the normal direction:

$$\mathbf{q}' = \mathbf{q} - (\mathbf{q} \cdot \hat{\mathbf{n}} - d) \hat{\mathbf{n}}$$



Closest Point to a Line



At what value of t is
the line closest to q ?

$$t = (\mathbf{q} - \mathbf{p}_0) \cdot \hat{\mathbf{d}}$$

$$\mathbf{q}' = \mathbf{p}_0 + ((\mathbf{q} - \mathbf{p}_0) \cdot \hat{\mathbf{d}}) \hat{\mathbf{d}}$$

$$\text{distance} = \|\mathbf{q} - \mathbf{q}'\|$$

Works for higher dimensions as well!

Point-In-Shape Tests

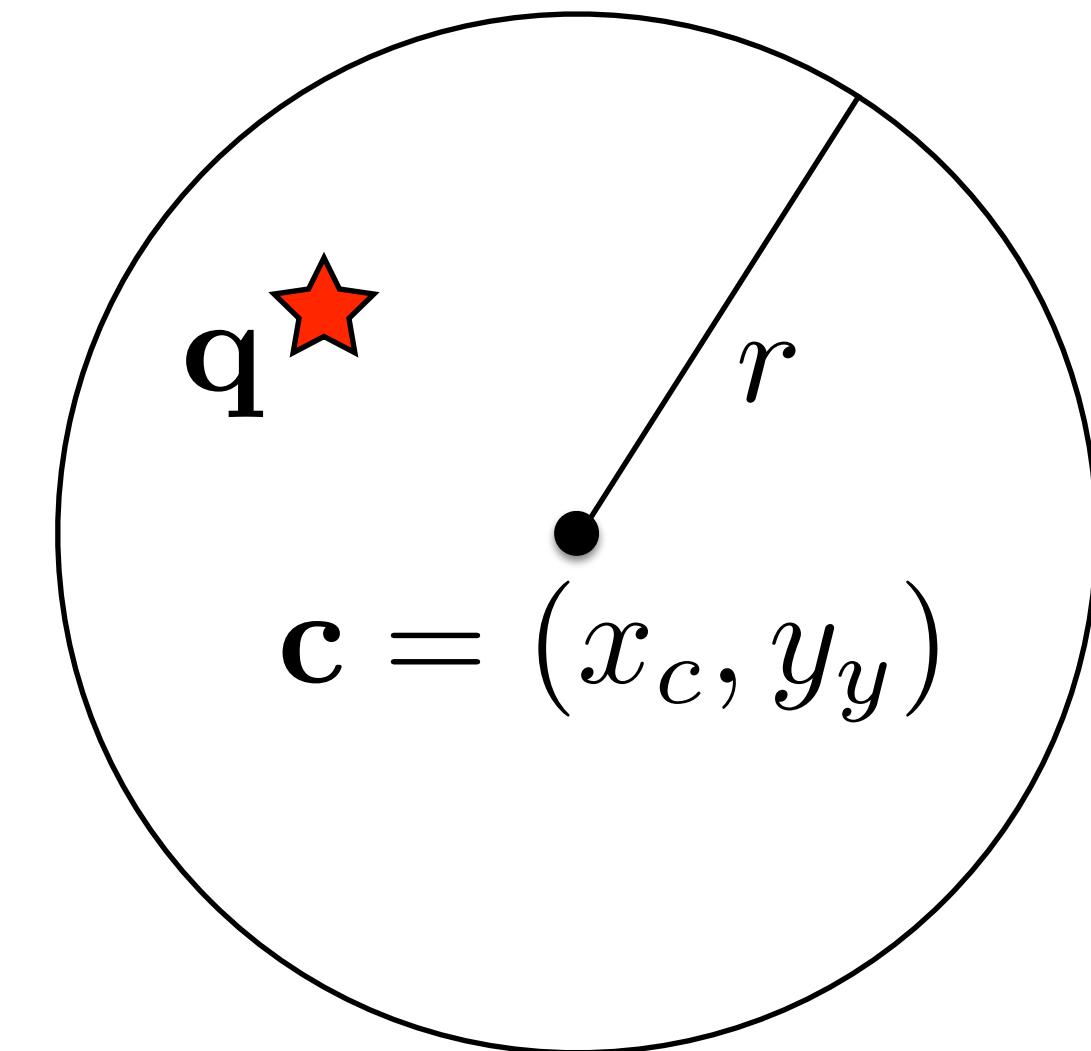
- Useful for selection tests in graphical interaction
- “Did the user click on...”



Point in Circle

Point on the circle:

$$\|\mathbf{q} - \mathbf{c}\| = r$$

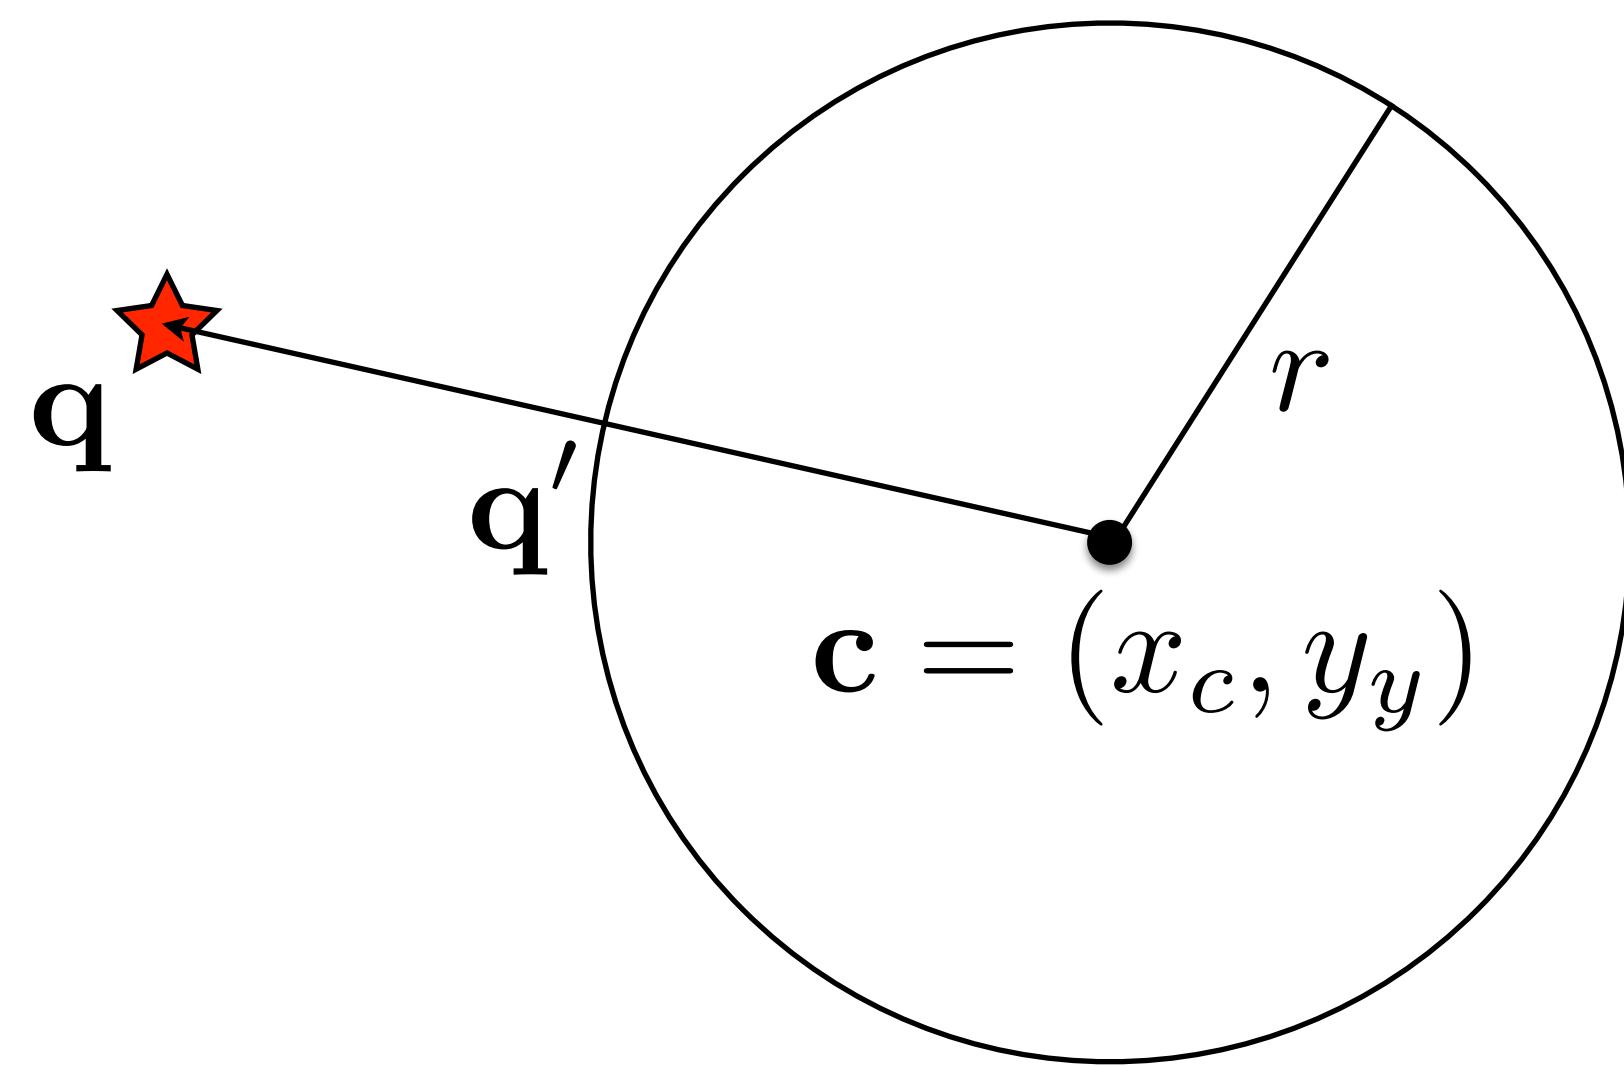


Point in the circle:

$$\|\mathbf{q} - \mathbf{c}\| \leq r$$

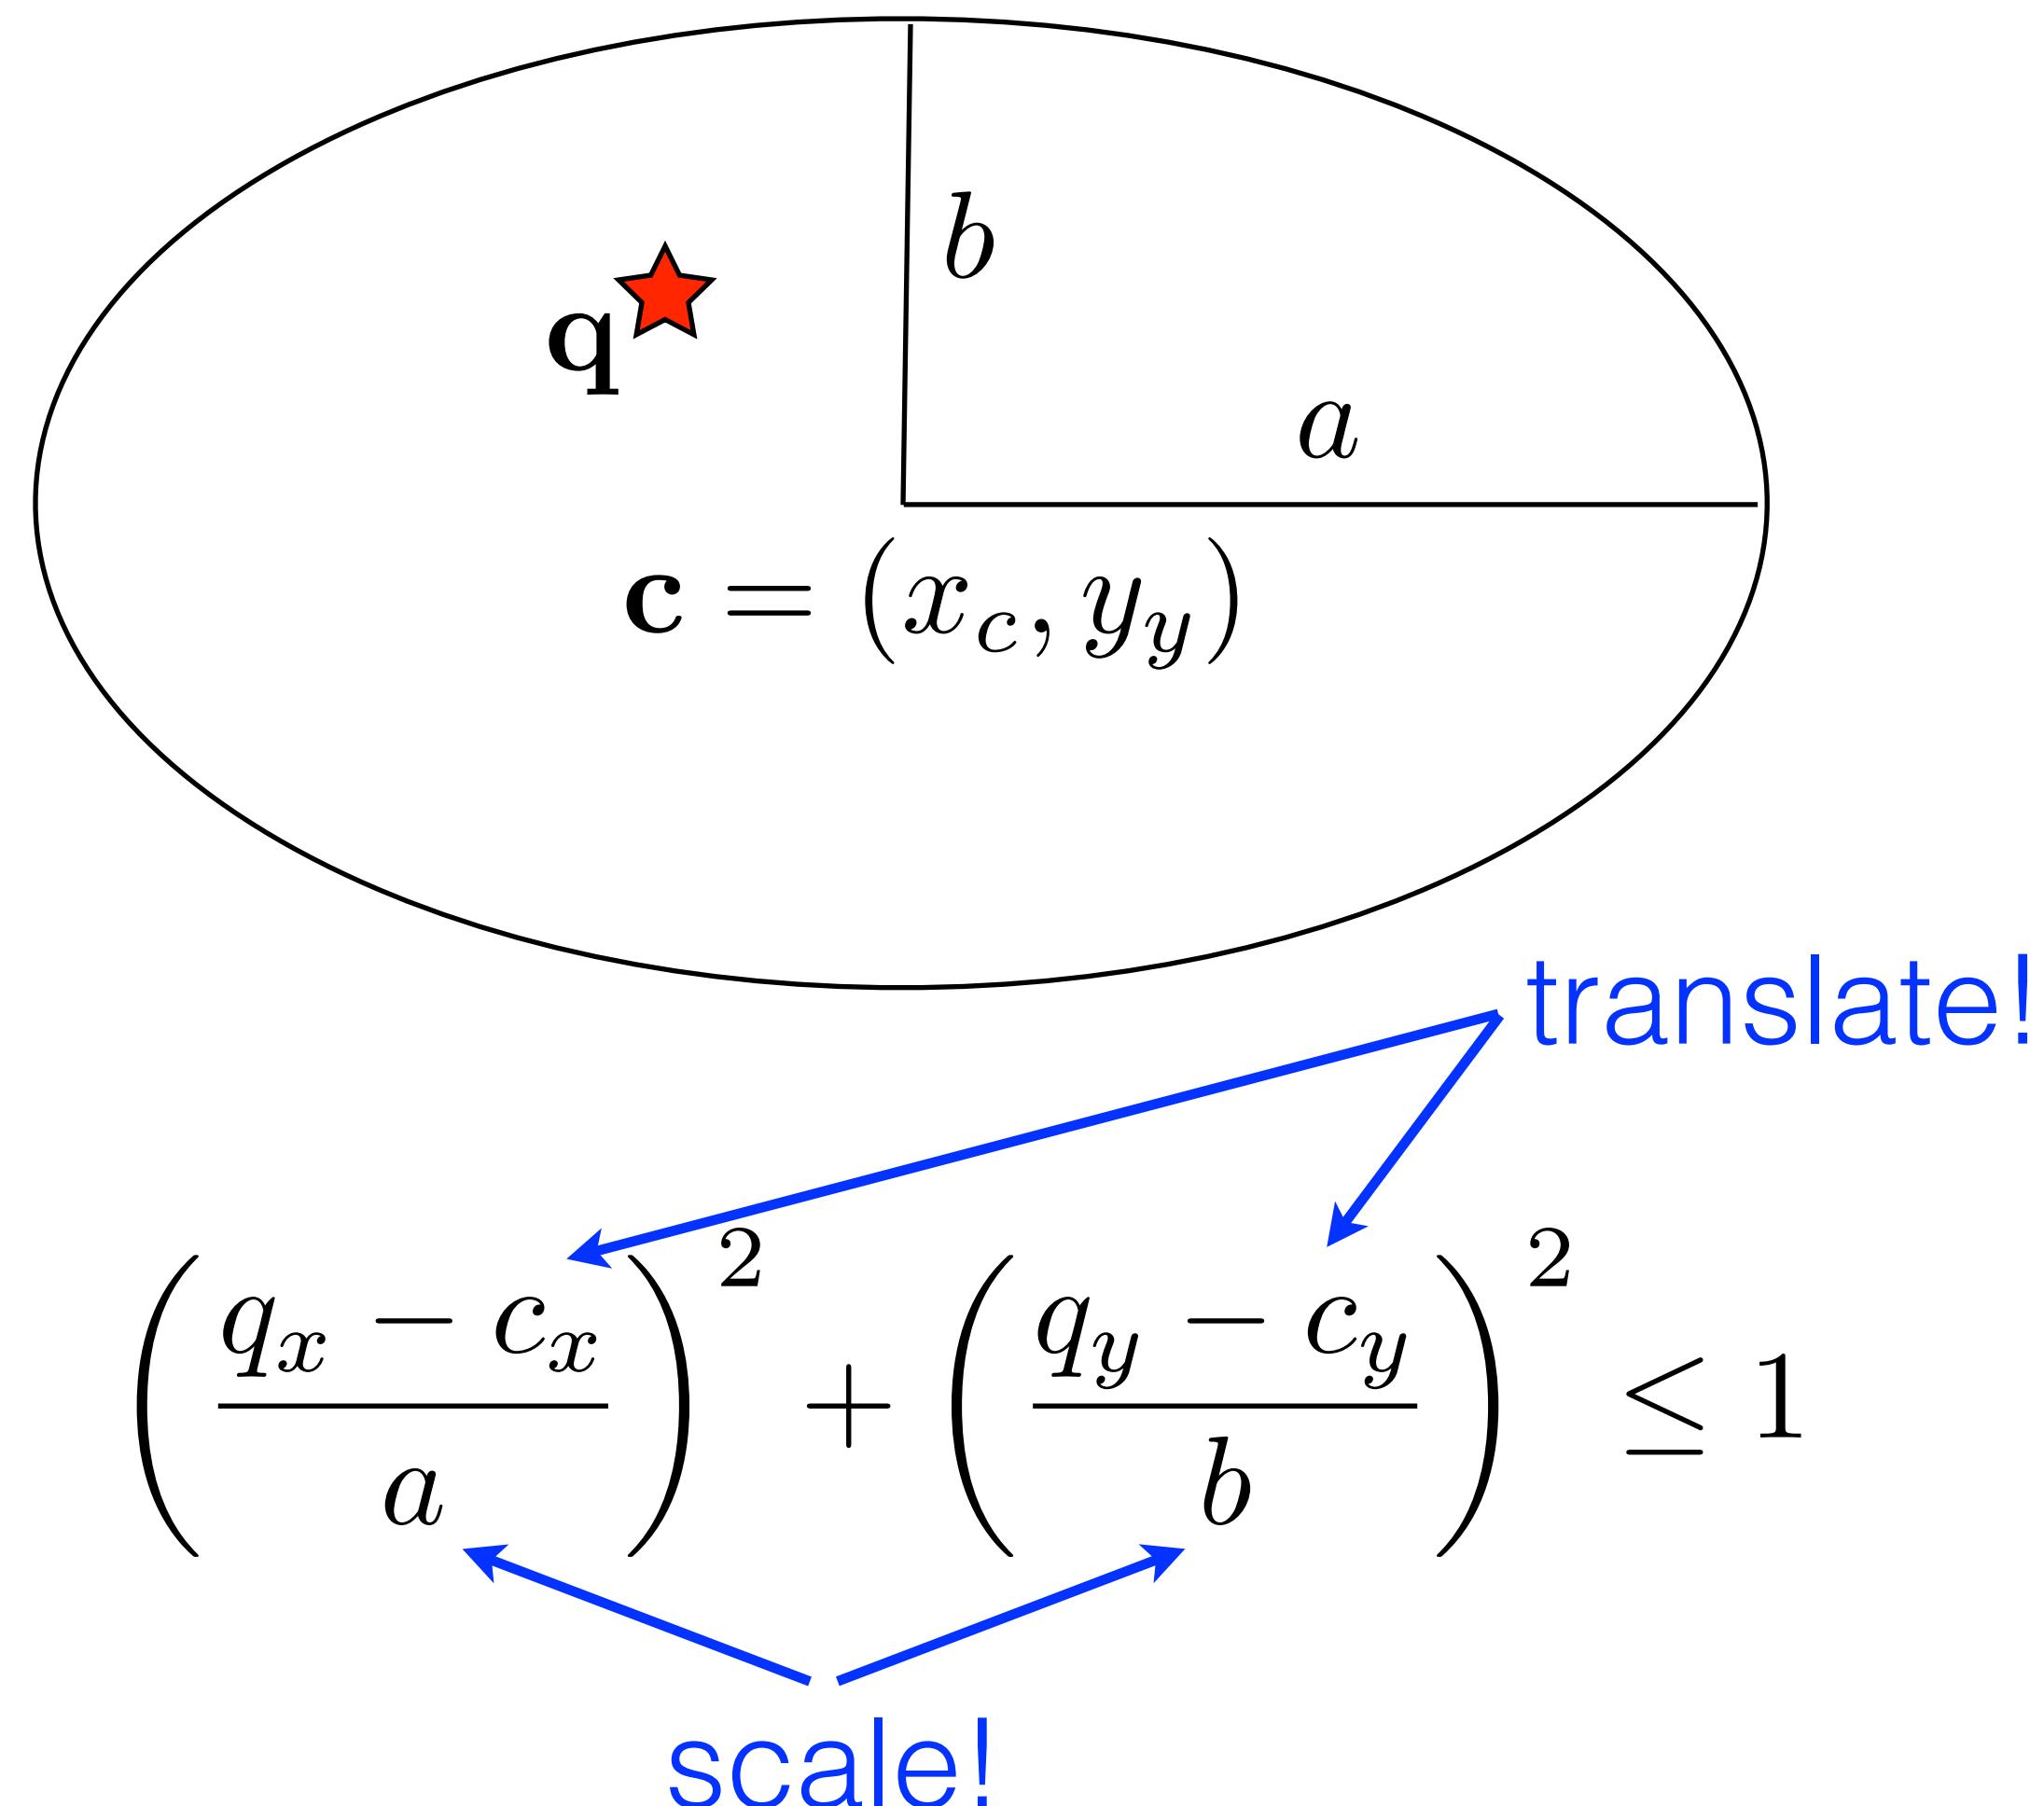
$$(q_x - c_x)^2 + (q_y - c_y)^2 \leq r^2$$

Closest Point on Circle



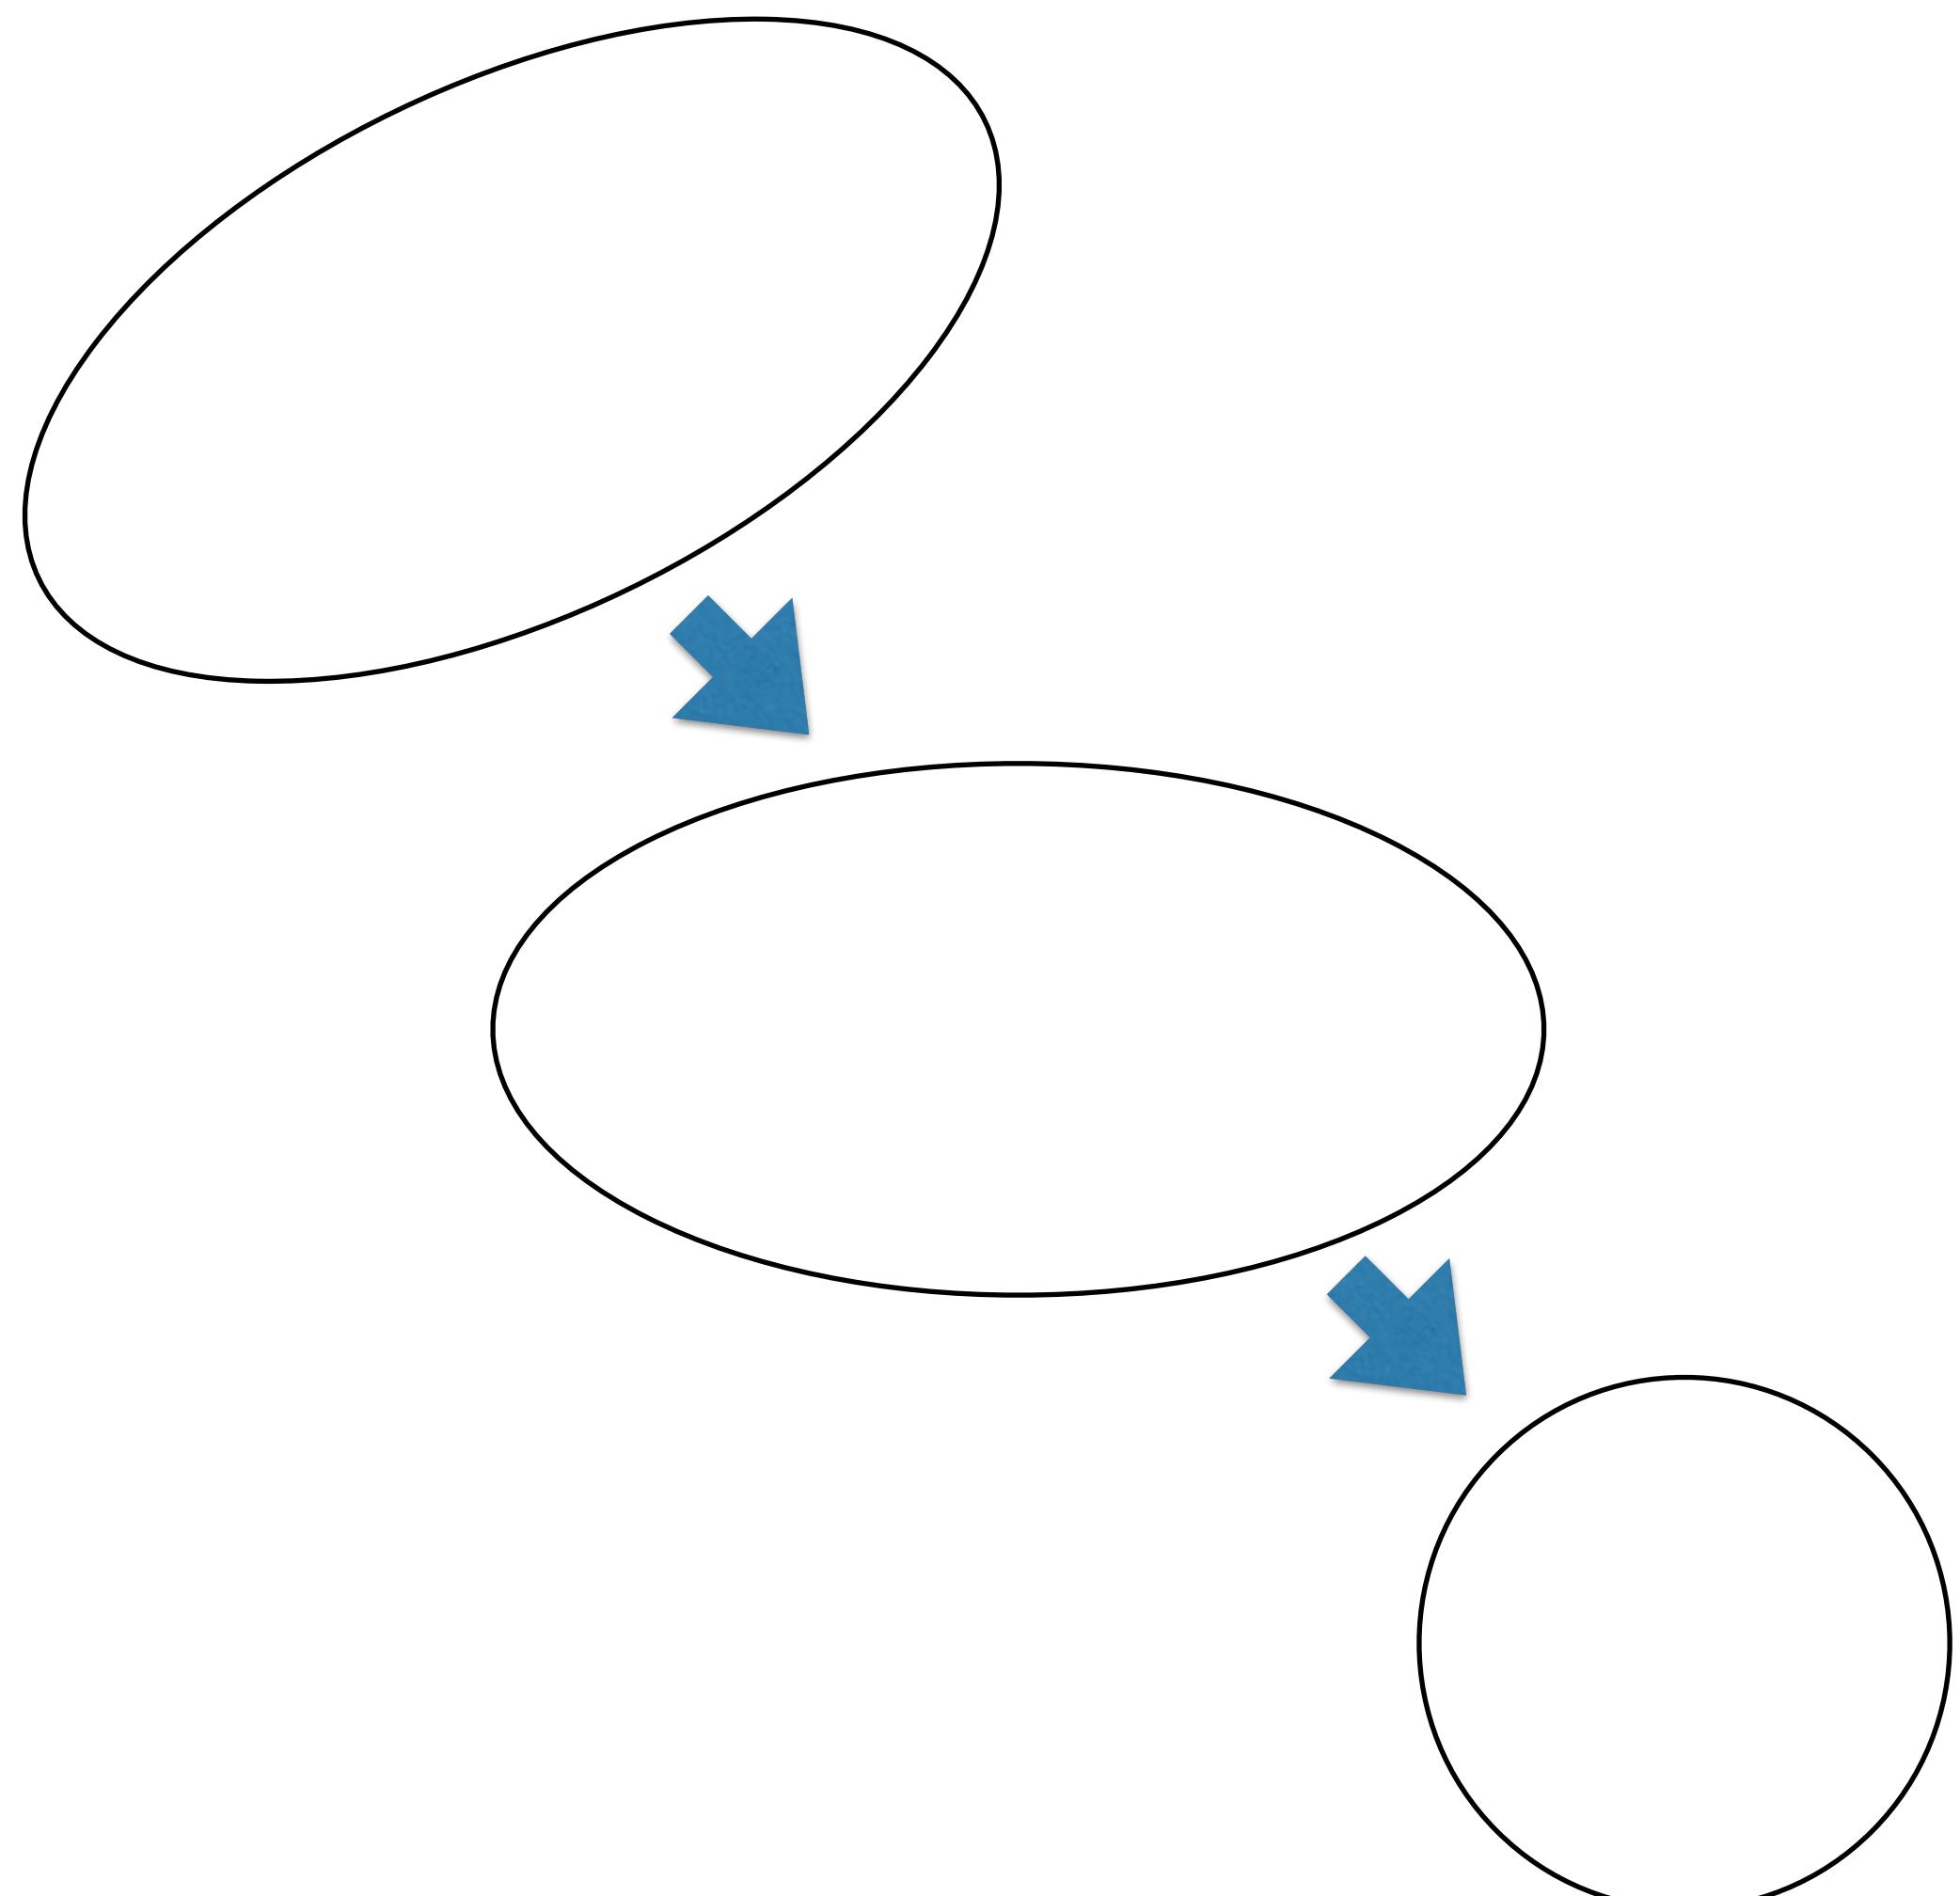
$$\mathbf{q}' = \mathbf{c} + r \frac{\mathbf{q} - \mathbf{c}}{\|\mathbf{q} - \mathbf{c}\|}$$

Point in Ellipse (Oval)



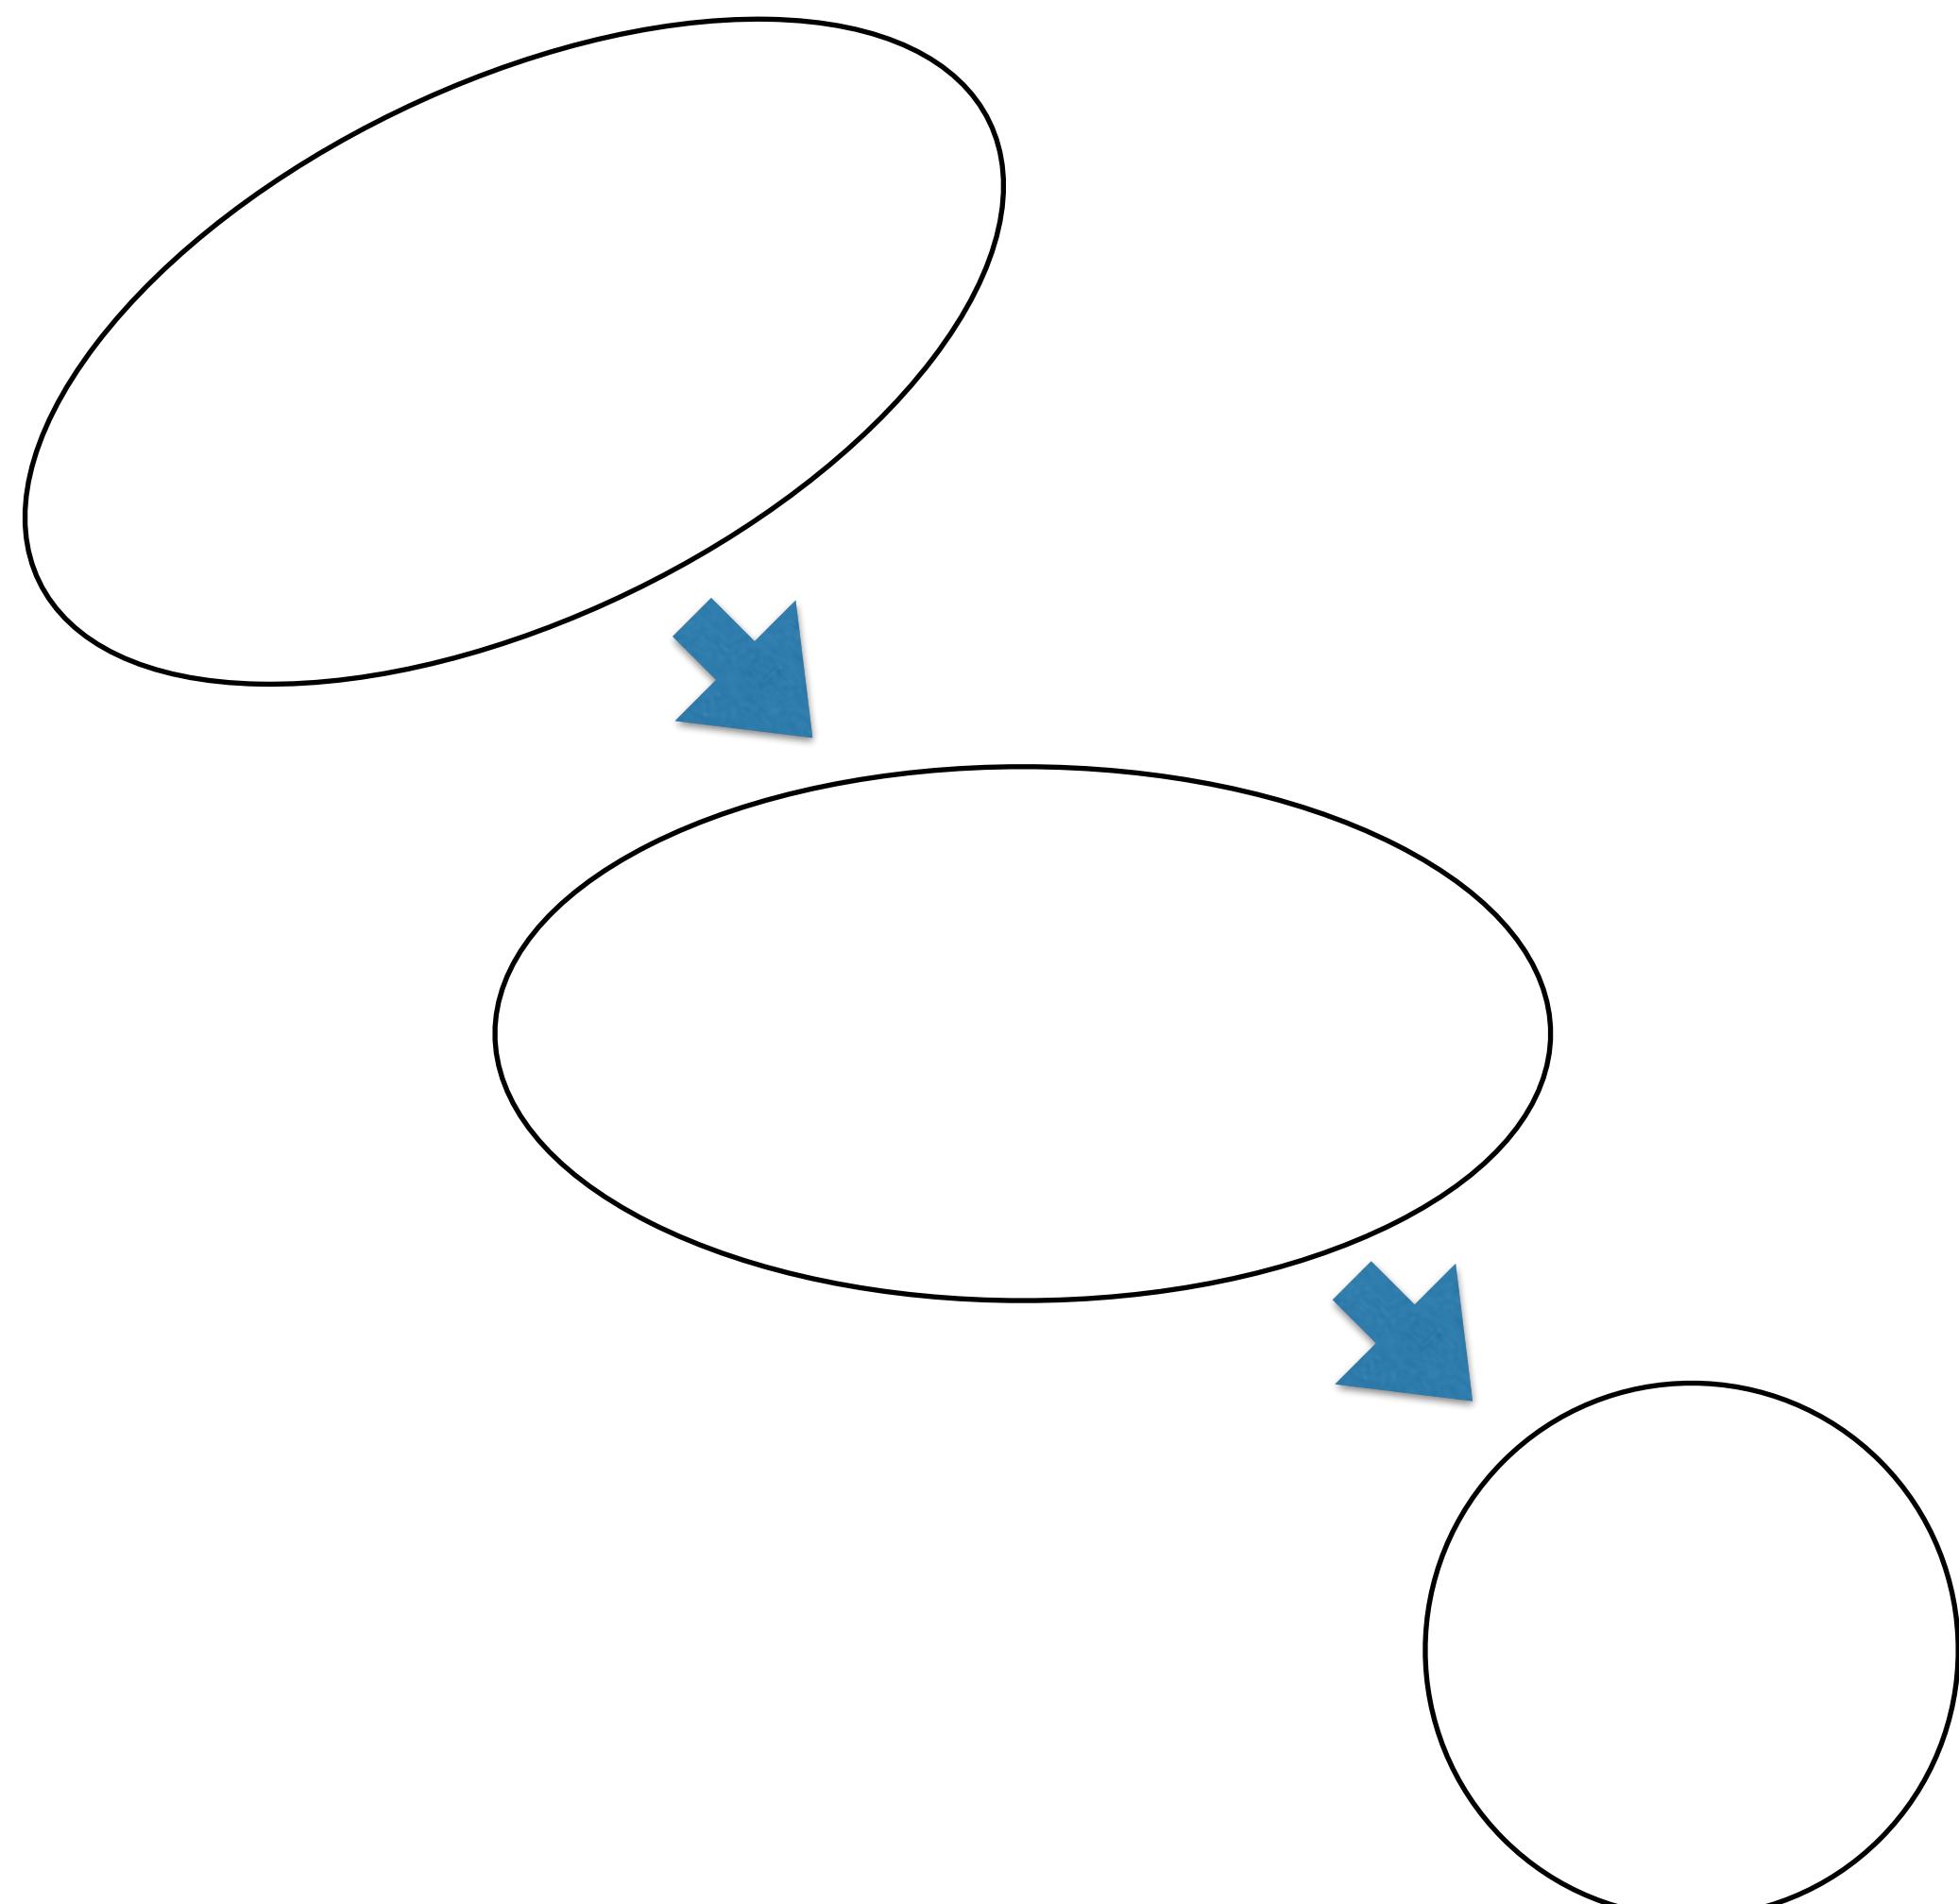
Transformations and Tests

- What seem like complicated tests become easier after transformation
- Center the shape
- Rotate to standard alignment
- Scale if necessary
- Simple tests



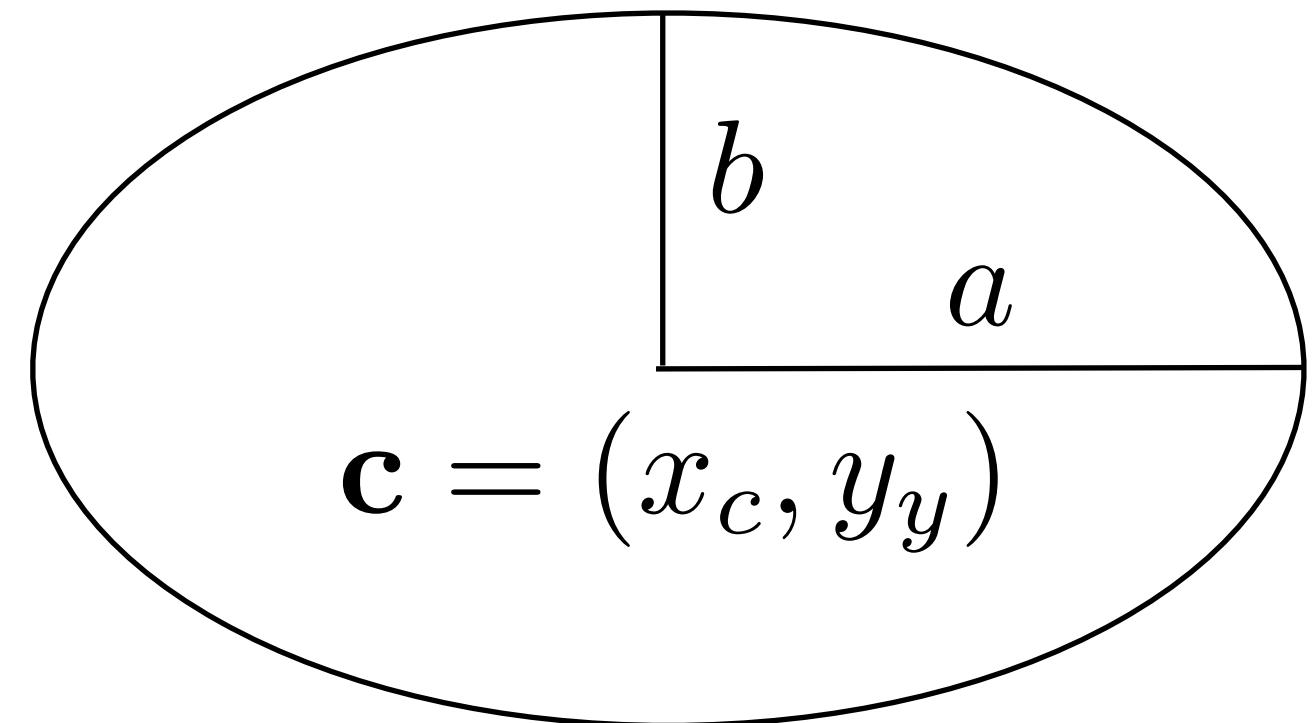
Transformations and Tests

- **Object-to-world** for rendering
- **World-to-object** for selection
(geometric tests are often easiest
the object's own intrinsic space)



Bounding Boxes

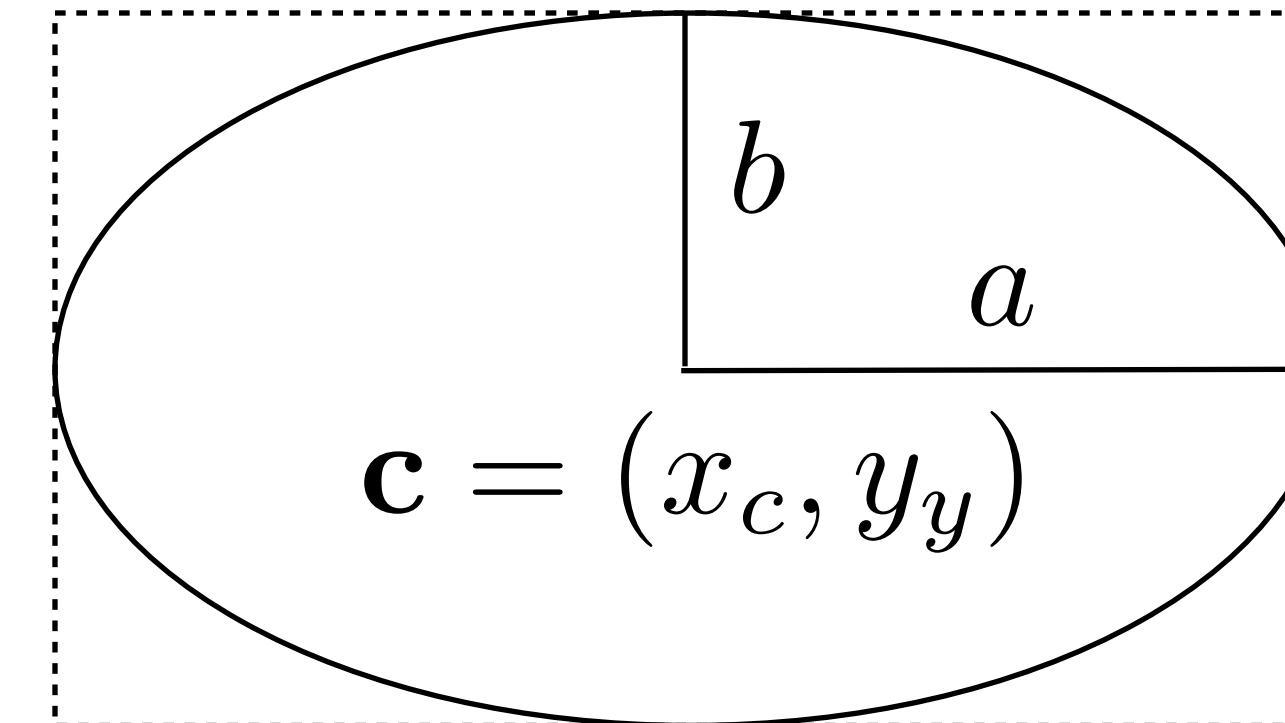
q★



Hard test:

$$\left(\frac{q_x - c_x}{a}\right)^2 + \left(\frac{q_y - c_y}{b}\right)^2 \leq 1$$

q★

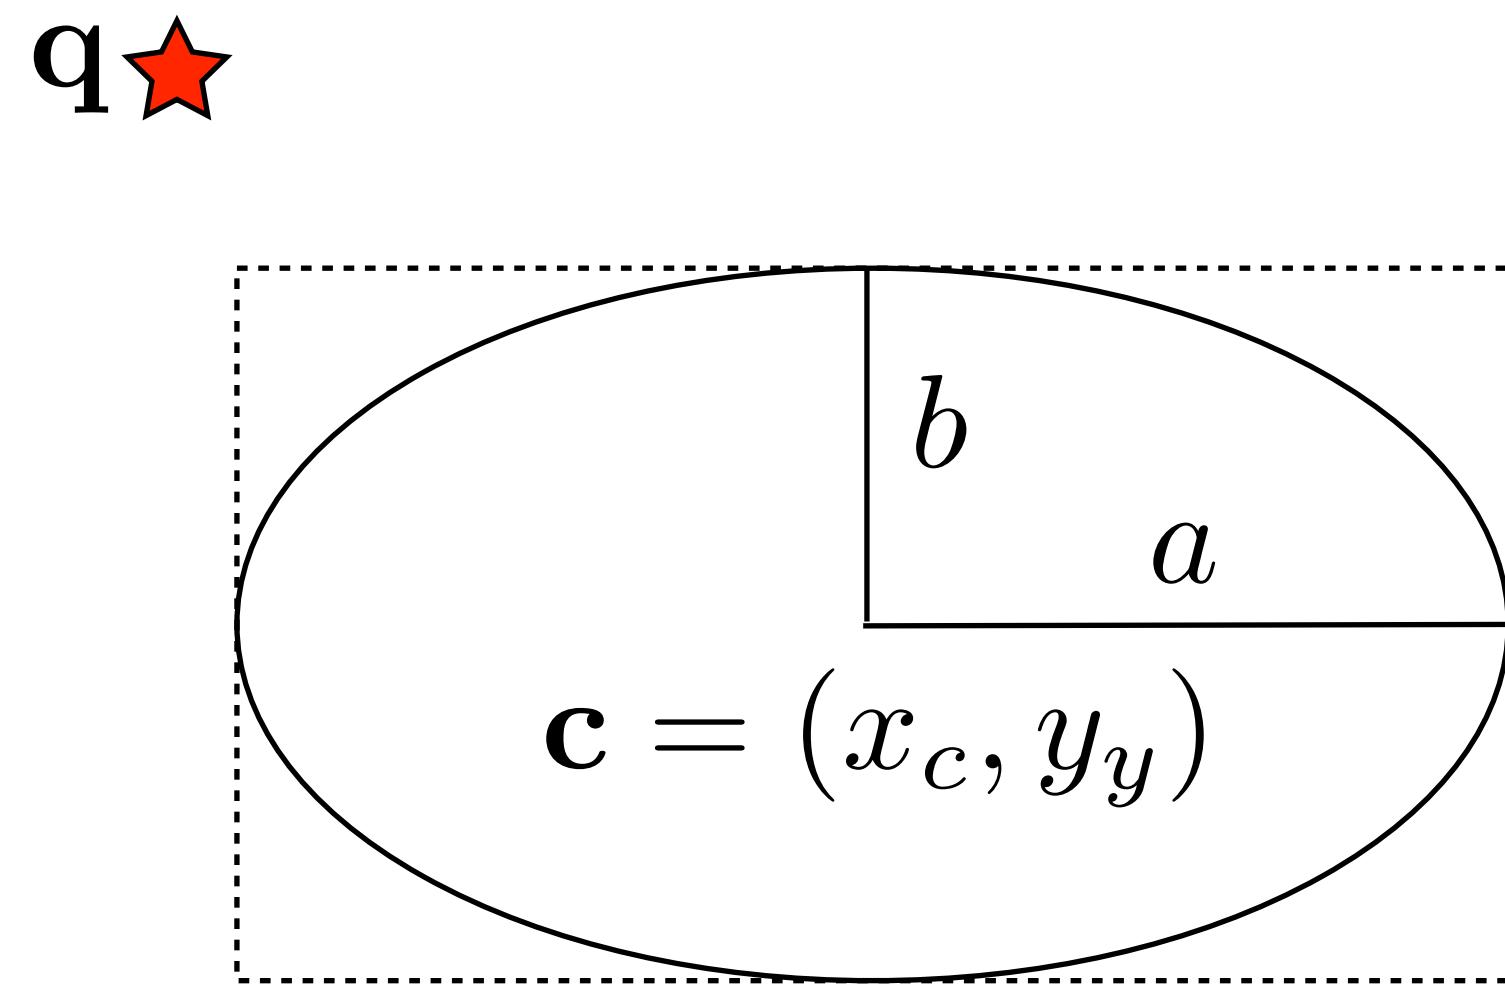


Easy test:

$$\begin{aligned}|q_x - c_x| &\leq a \\ |q_y - c_y| &\leq b\end{aligned}$$

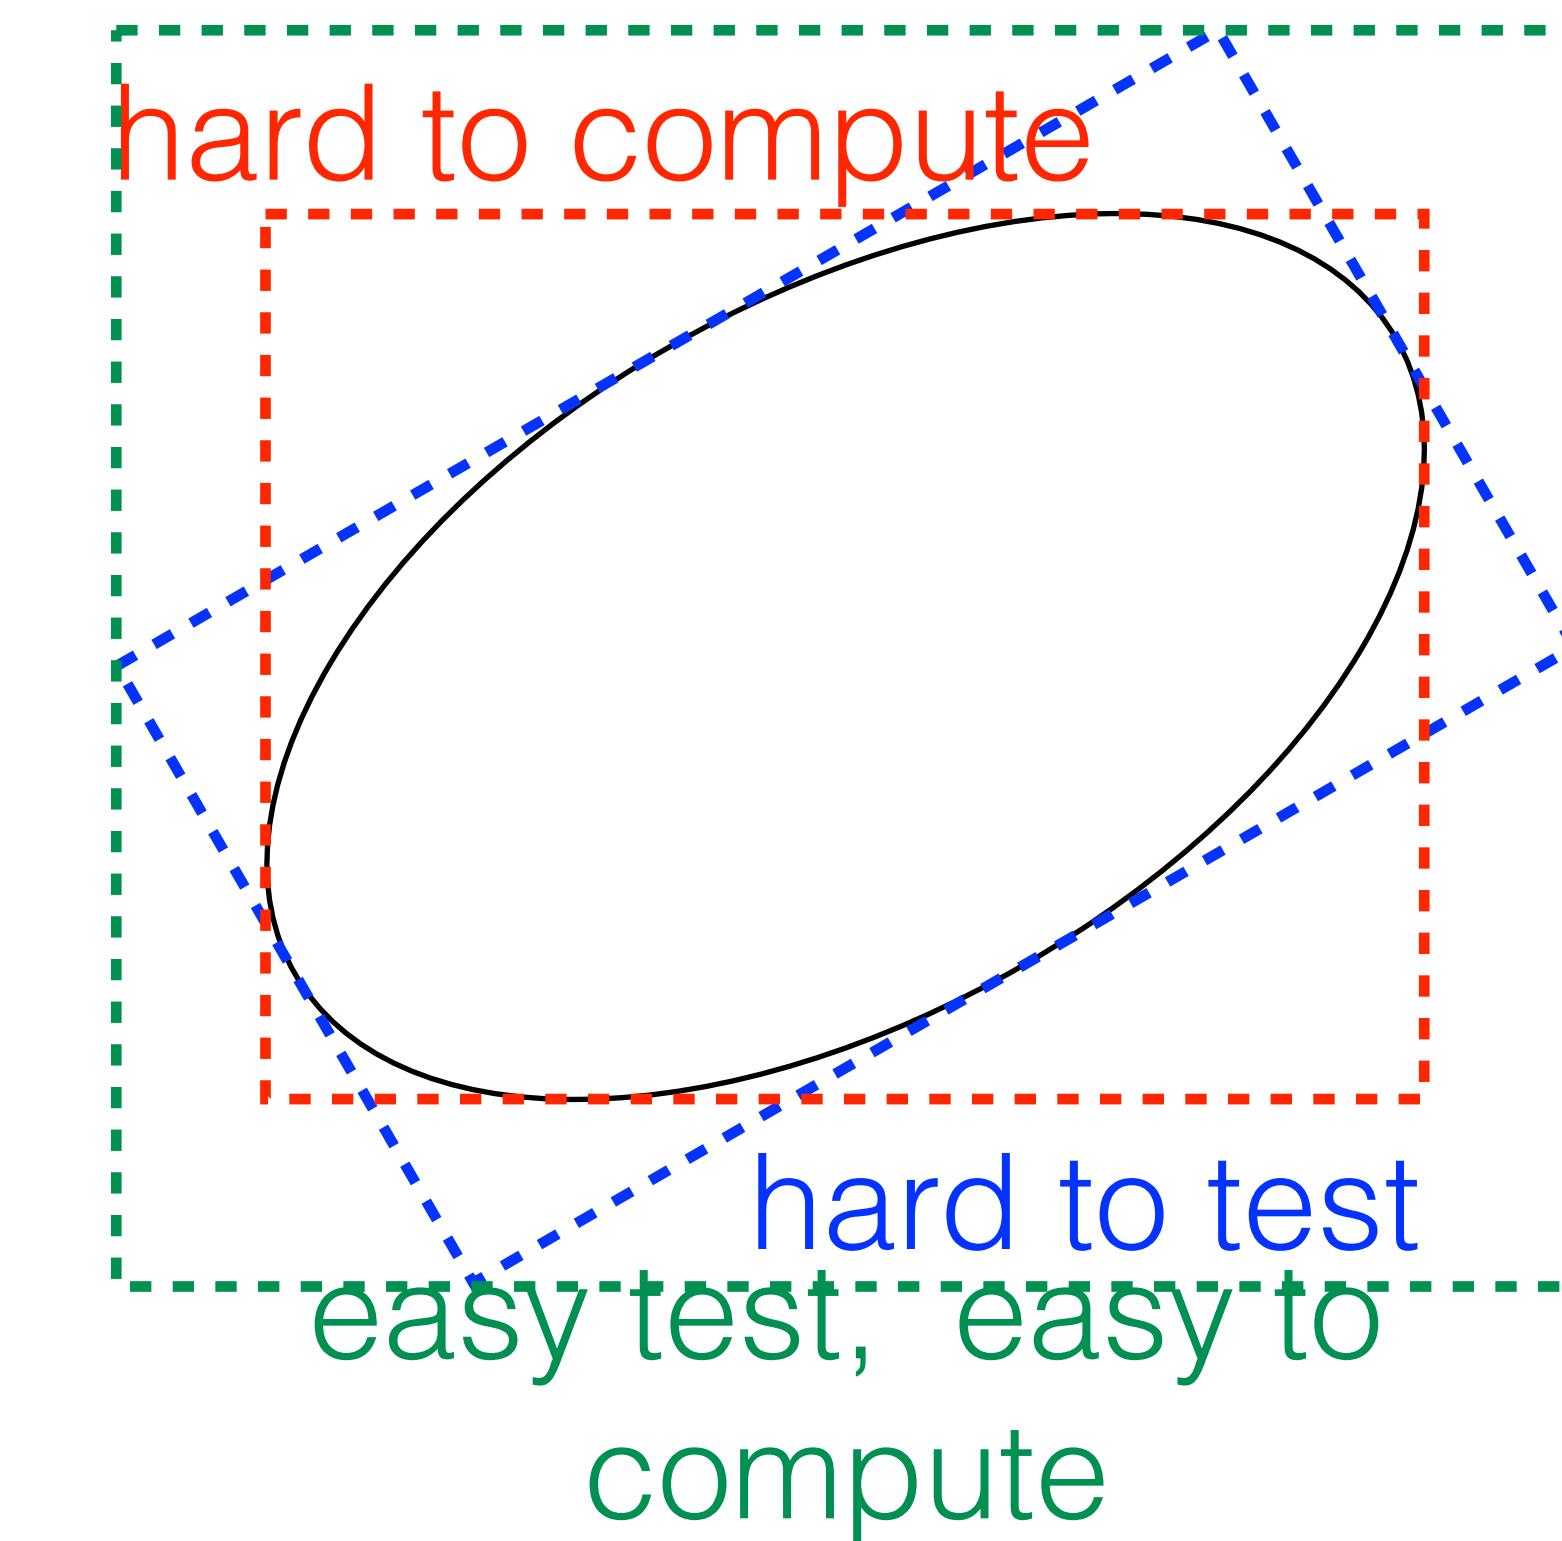
Bounding Boxes

- Idea: use bounding box tests as a “quick reject”
- If passes, then spend time on more complex tests
- The more complex the test, the bigger the win



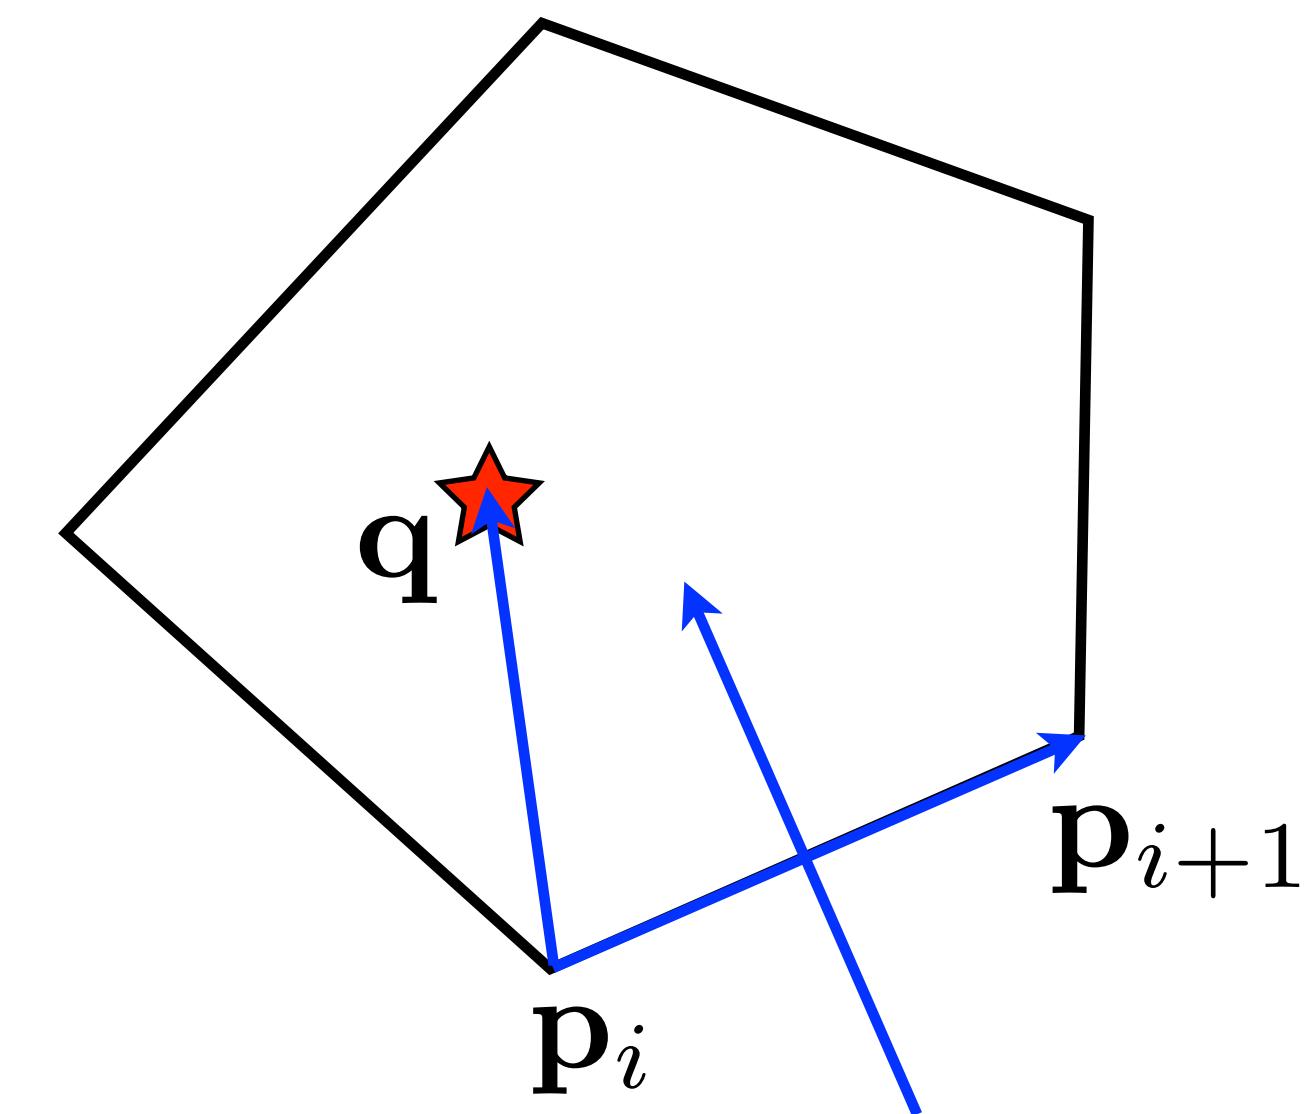
Bounding Boxes

- Boxes don't necessarily have to be "tight" to be useful
- Looser boxes may be easier to compute and test
- "Axis-Aligned Bounding Box" (AABB)



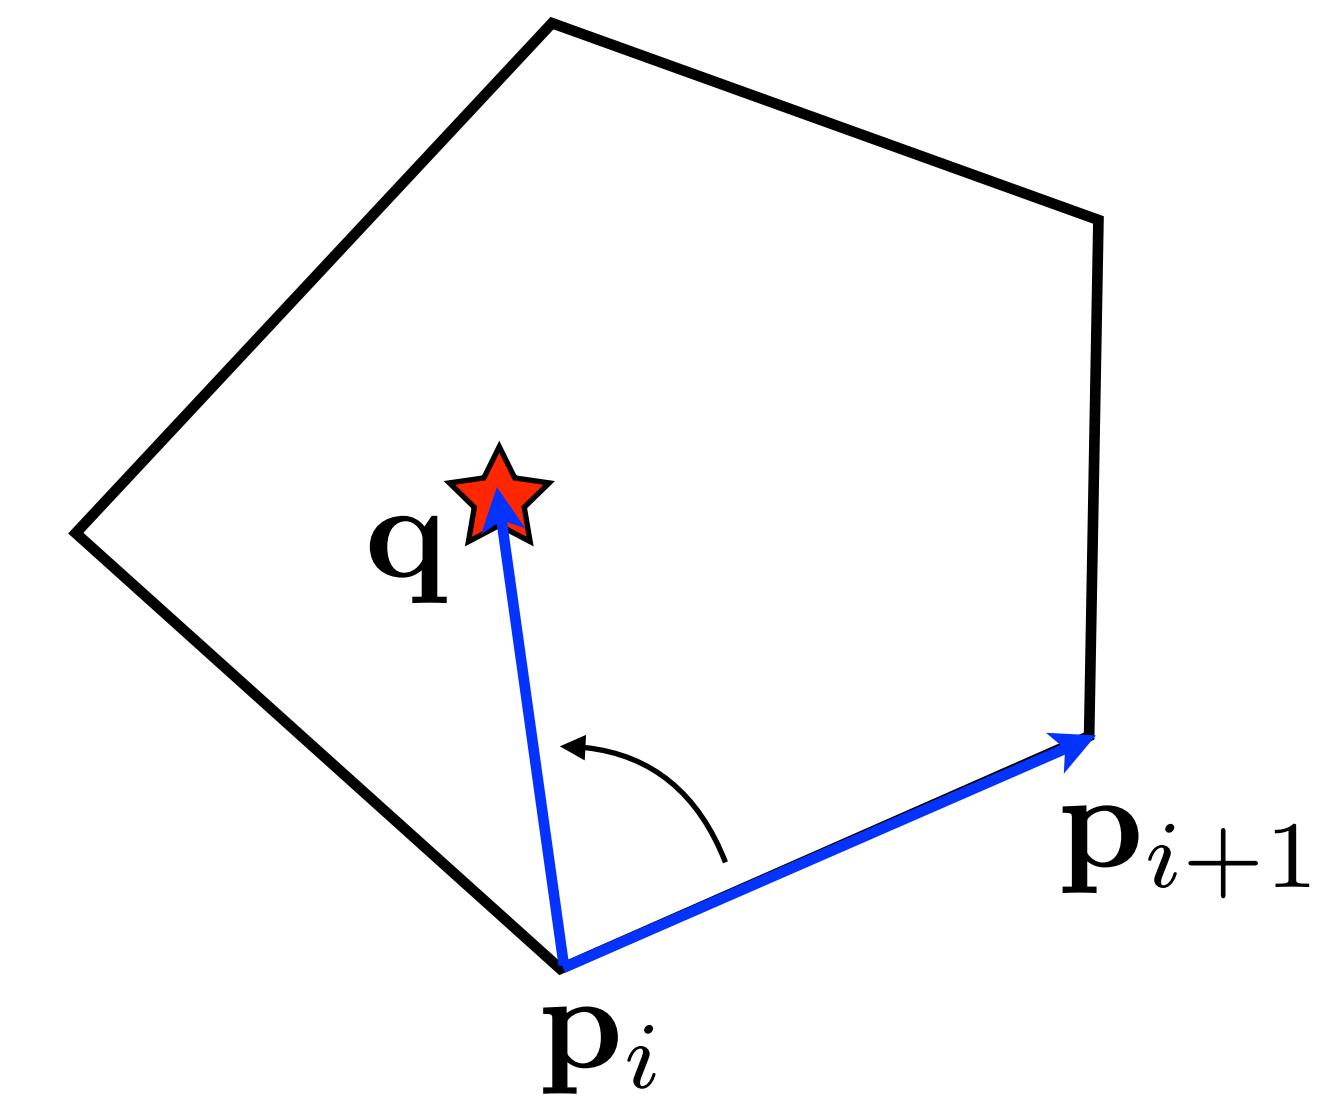
Convex Polygons

- In 2D: for all edges, the point is on the same side of the edge
- Walk around the polygon (in order) and test
$$(\mathbf{q} - \mathbf{p}_i) \cdot (\mathbf{p}_{i+1} - \mathbf{p}_i)_{\perp} > 0$$
- Be consistent with ordering and perpendiculars



Convex Polygons

- Can also do using cross-products
- Turn all 2-D points (x,y) to 3-D ones $(x,y,0)$
- Test the z component of this cross-product:
$$(\mathbf{p}_{i+1} - \mathbf{p}_i) \times (\mathbf{q} - \mathbf{p}_i)$$
- If all are positive, the point is in the shape

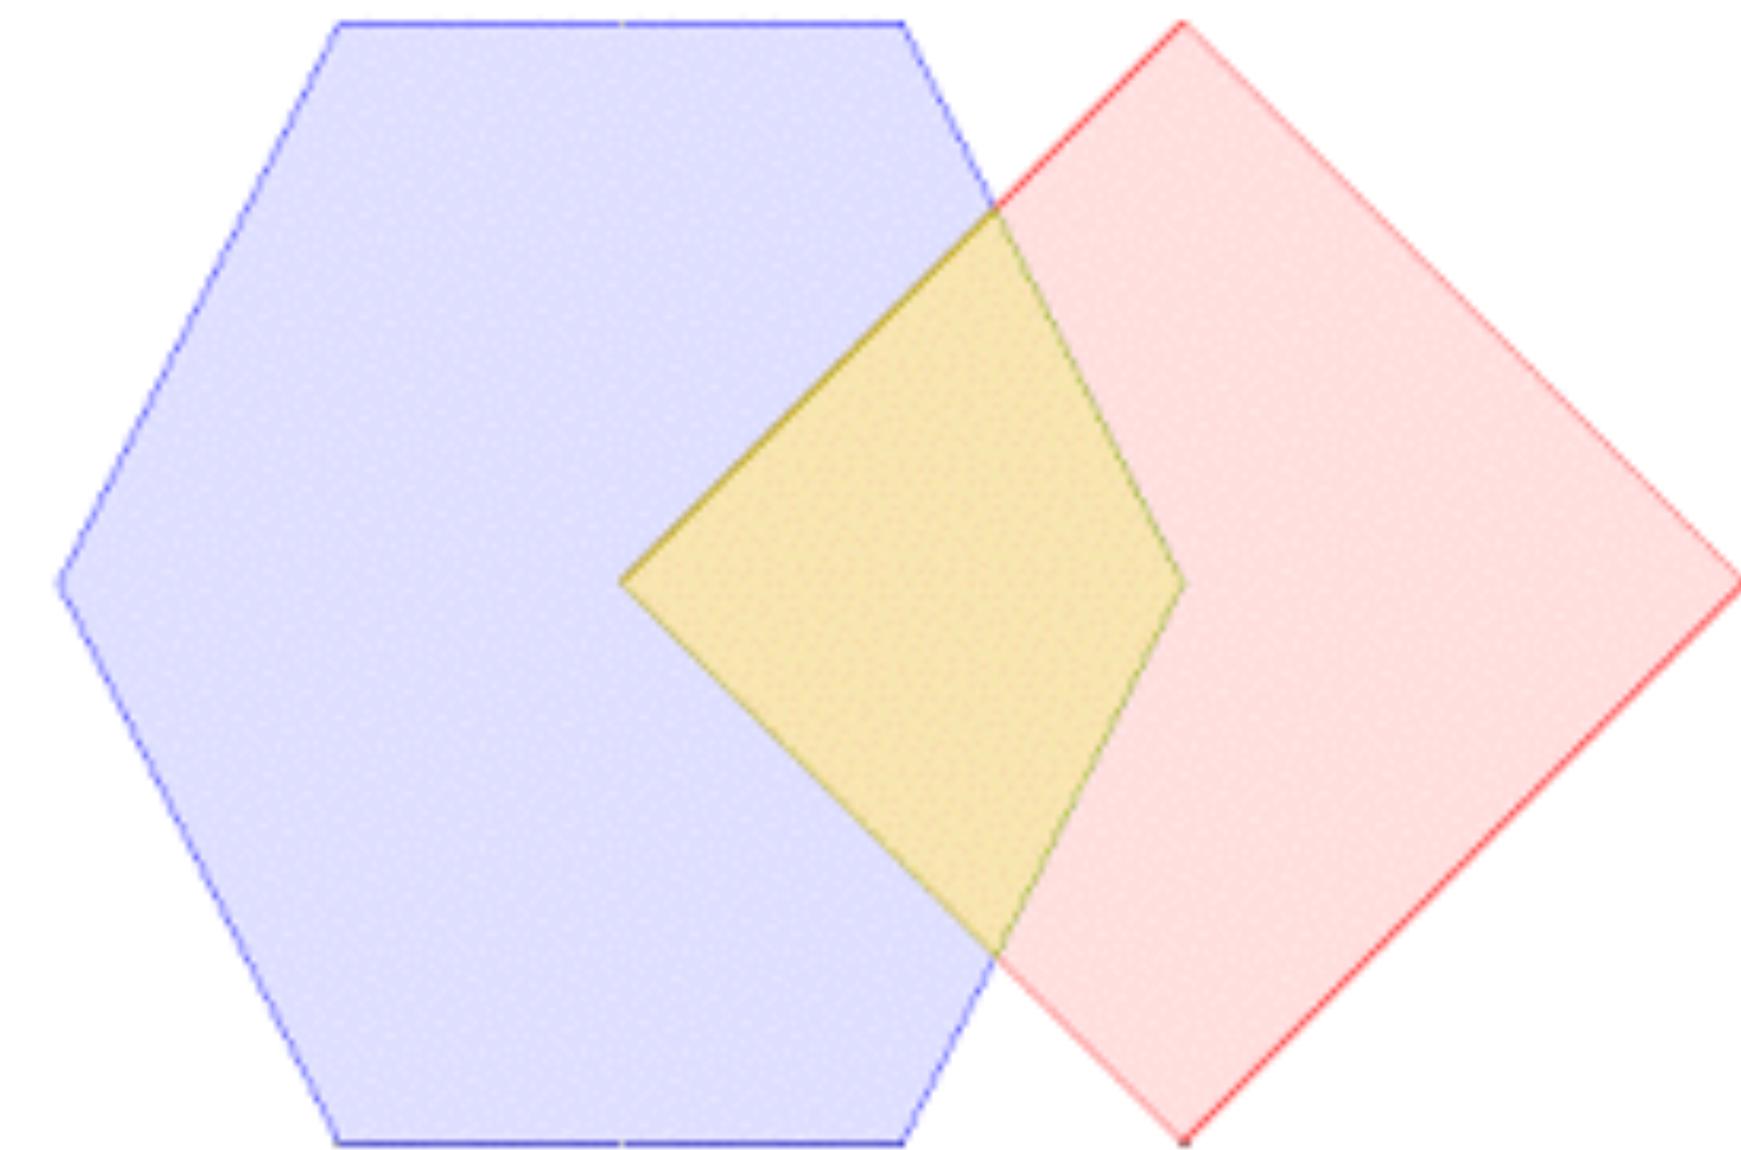


Polygon Bounding Boxes

- Bounding boxes for polygons are really easy $\min(x_i)$
- Just min / max tests over all of the vertices' x and y coordinates $\min(y_i)$ $\max(x_i)$ $\max(y_i)$

Intersecting Polygons

- To test to see if two convex polygons intersect, see if any of the vertices of one are within the other, and vice versa



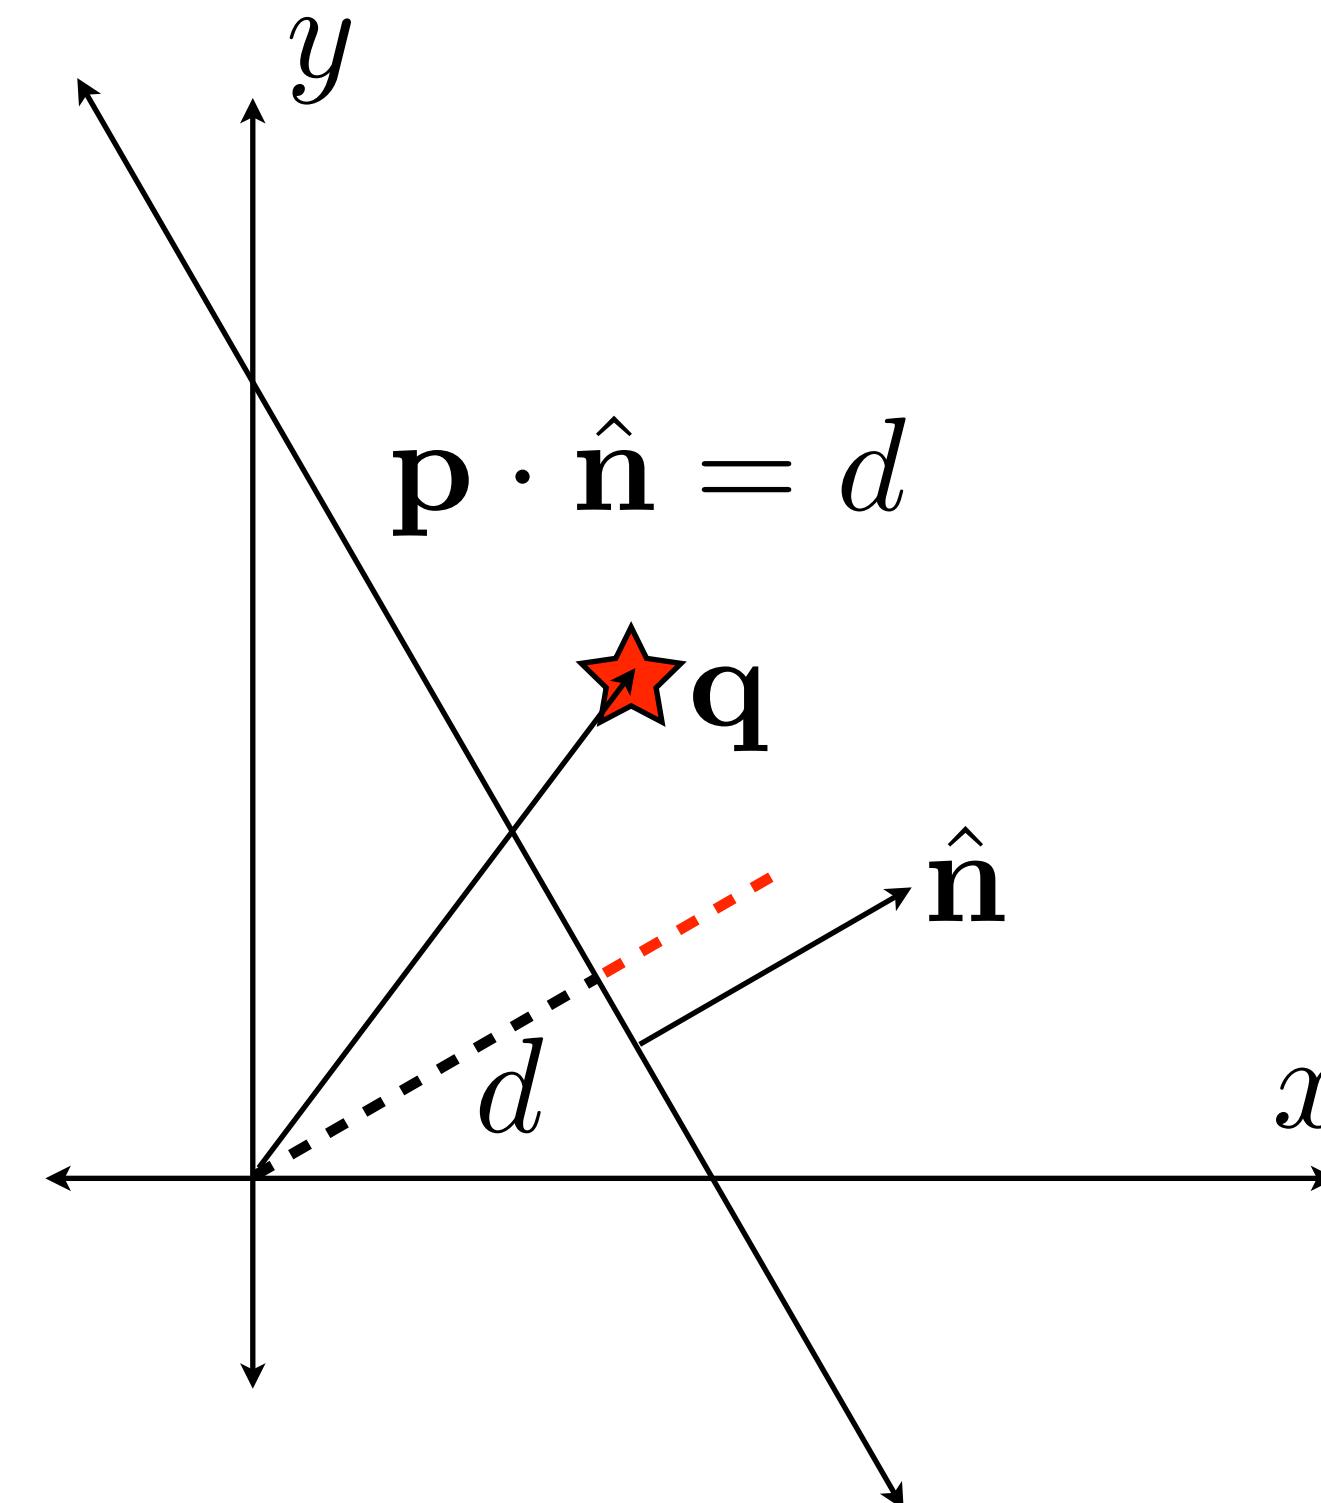
Which Side of a Line?

- Points \mathbf{p} on the line \mathbf{L} satisfy this constraint:

$$\mathbf{p} \cdot \hat{\mathbf{n}} - d = 0$$

- Use to test which side of a line a point is on by looking at the sign of

$$\mathbf{q} \cdot \hat{\mathbf{n}} - d$$



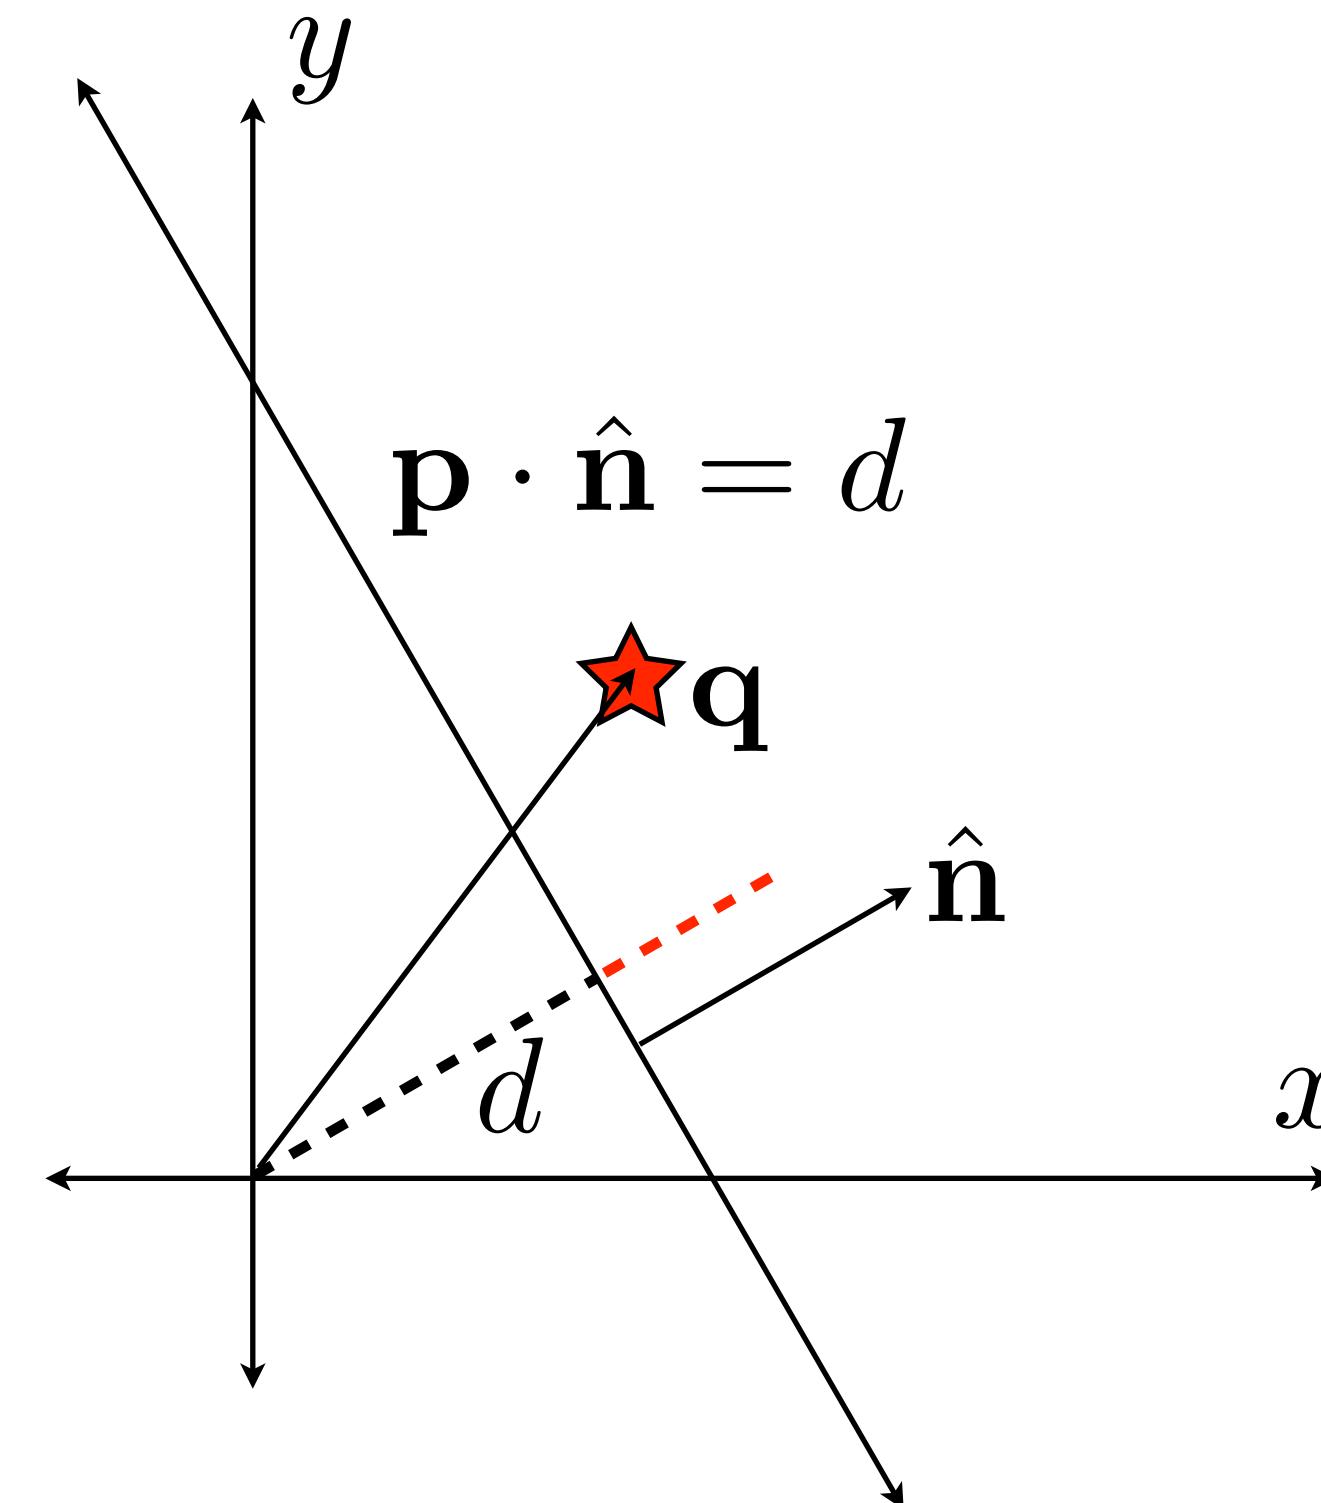
Which Side of a Plane?

- A 2-D plane in 3-D space can be represented the same way:

$$\mathbf{p} \cdot \hat{\mathbf{n}} - d = 0$$

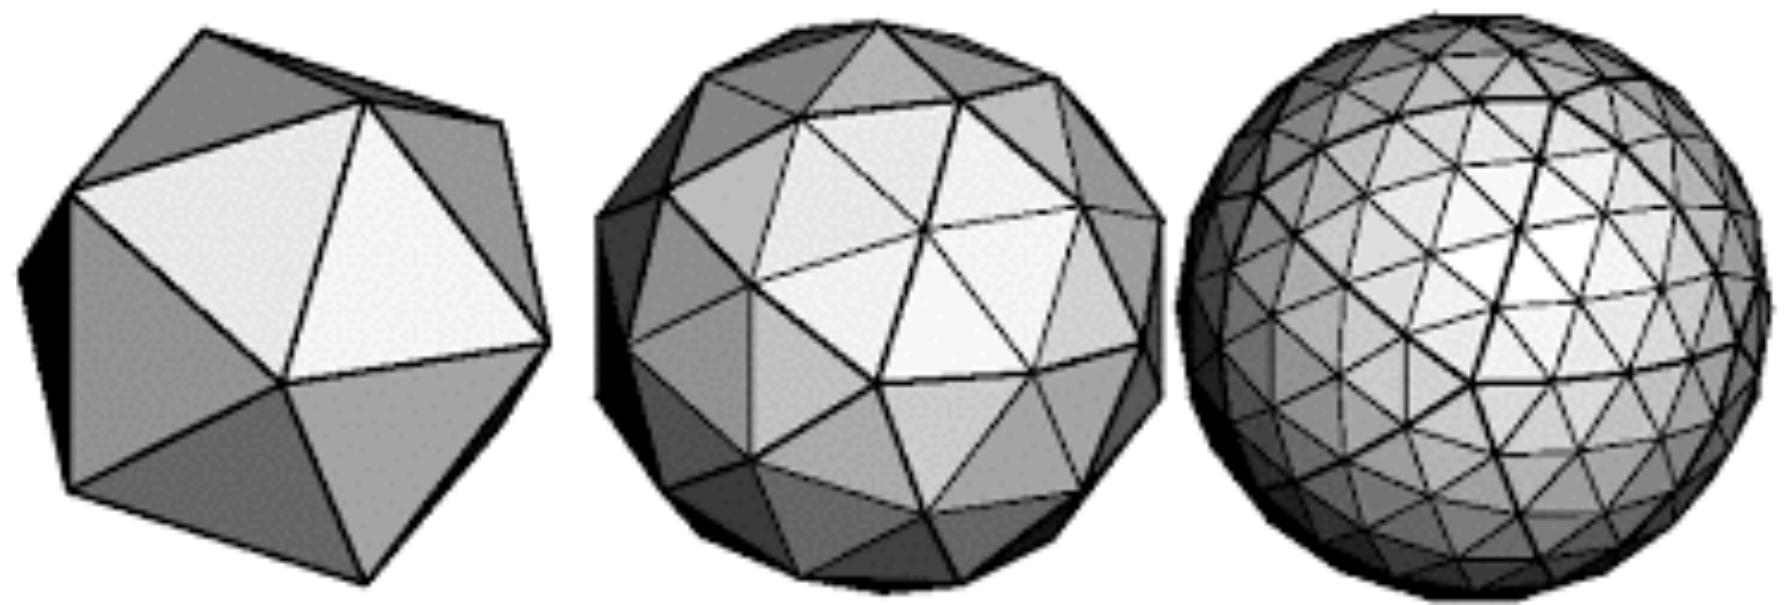
- Use to test which side of a plane a point is on by looking at the sign of

$$\mathbf{q} \cdot \hat{\mathbf{n}} - d$$



Point Inside Polygonal Model

- A point is inside a convex polygonal model if it is on the “inside” side of all of the faces
- Can test for intersection of two convex shapes by seeing if a vertex of one is inside the other, and vice versa



Lab 9

- Don't perform backward mapping for every point in the image
- Compute a bounding box around the four corners of the frame (min/max on the x, y coordinates)
- Test to see if each point is inside the quadrilateral of the frame (point in convex polygon test)



Lab 9

- We give you an implementation of the 4-point algorithm:
 - You give it matching points (match corners of the target frame to corners of the image)
 - It gives you back a homography
- You implement:
 - Point-in-frame geometric tests
 - Iterate over bounding box
 - Point-in-quadrilateral tests
 - Backward warping
 - Bilinear interpolation



Coming up...

- Frequency-domain processing



Introduction to the Frequency Domain

CS 355: Introduction to Graphics and Image Processing

Let's revisit the idea of
transformations...

Basis Sets

- Remember that a basis set is a minimal set of vectors that span a space of vectors
- That means any vector in the space can be represented by a unique sum of the basis vectors
- **All vectors are represented with respect to some basis set**

$$\{\mathbf{e}_i\}$$

$$\mathbf{v} = \sum_i a_i \mathbf{e}_i$$

Change of Basis

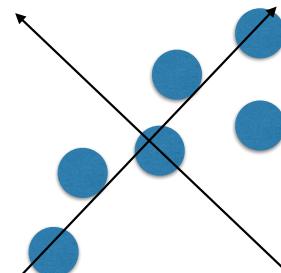
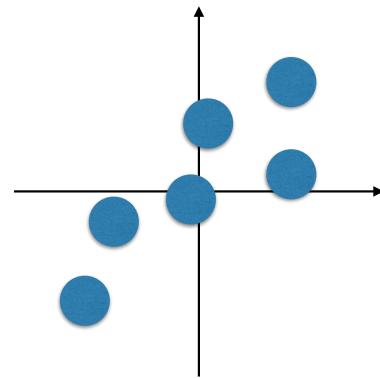
- Can change from
 - representation in terms of one basis set to
 - representation in terms of another basis set
- If basis vectors are orthonormal, this is just simple dot products
(you've seen this already)

$$\mathbf{v} = \sum_i a_i \mathbf{e}_i$$

$$a_i = \mathbf{v} \cdot \mathbf{e}_i$$

Change of Basis

- For many problems, *analyzing points and vectors may be easier in a different coordinate system*
- Key is often to find the right coordinate system
 - Can be based on the problem
 - Can adapt to the specific data

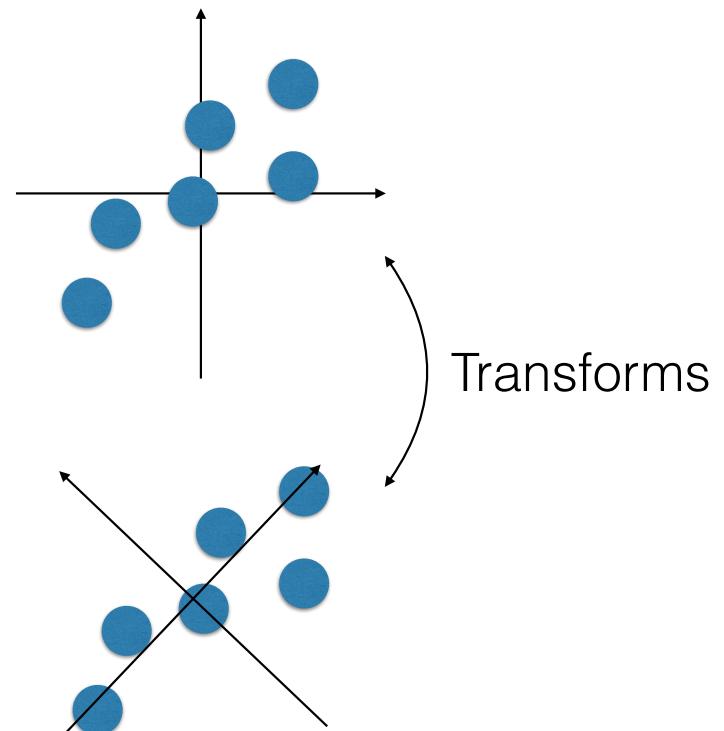


Data as Vectors

Lots of things can be thought of as points/vectors in some space

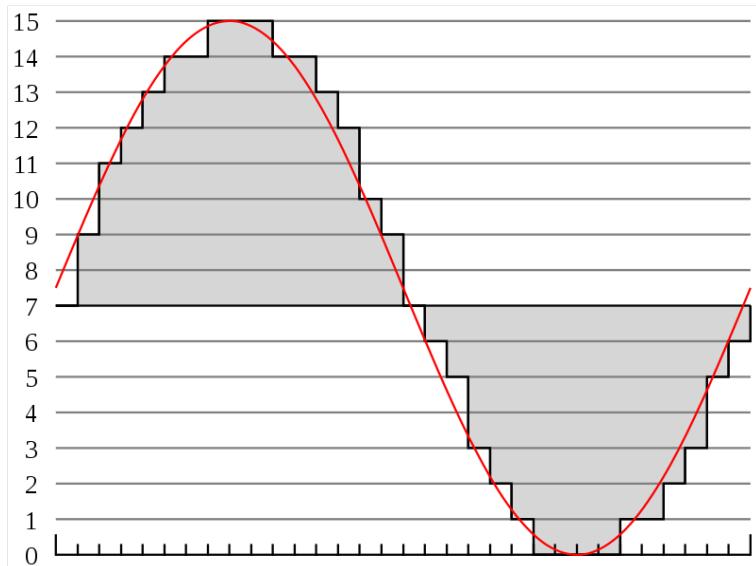
Data Transforms

- Common pattern in data analysis:
 - Represent data as vectors
 - Convert to a different coordinate system
 - Analyze (or change!) while in that coordinate system
 - Convert back (if needed)



Digital Audio

- Raw digital audio is stored as a series of sampled values (Pulse Code Modulation)
- One second of (one channel of) music on a CD is 44,100 samples
- Is this any different from a vector in a 44,100 dimensional space?
- Can we represent it differently?



Fourier Analysis

- Sampled sines and cosines of different frequencies form an orthonormal basis set
- Can decompose any waveform into a weighted sum of sines and cosines of different frequencies
- Great for analysis, manipulation, etc.

$$c[u] = \sum_k f[k] \cos(2\pi u t)$$

$$s[u] = \sum_k f[k] \sin(2\pi u t)$$

(OK, there's a bit more to it, but this is the basic idea.)

Let's hear it...

Key Ideas from Today

- Transformations used for geometry extend to higher-dimensional data
- Can represent any point/vector as a weighted sum of basis vectors
- Choose the right set of basis vectors for your problem:
 - If the basis is orthonormal, get the weights by projecting (dot product)
 - Get back the original using a weighted sum of the basis vectors
- Sines and cosines are orthonormal basis set for 1-D functions like sound
- Can represent sound in the **time domain** or the **frequency domain**
- Can be used for analysis, but can also use it to manipulate the sound!

Coming up...

- A bit of review on complex numbers and sinusoids
- The Fourier Transform
- The Fast Fourier Transform
- How to feed data into FFT code and how to interpret the results
- Filtering in the frequency domain
- Sampling revisited
- Image filtering



The Fourier Transform

CS 355: Introduction to Graphics and Image Processing

First, a bit about
sines and cosines...

Sinusoids

- Useful to think of sinusoids in terms of three properties:
 - their *frequency* (how often they repeat)
 - their *amplitude* (height of the peaks)
 - their *phase* (shifting left or right)
- A sine wave is just a cosine wave with phase $\pi/2$

$$A \cos(2\pi ut + p)$$

Amplitude

Frequency

Phase (shift)

The diagram illustrates the components of a sinusoidal wave equation. The equation is $A \cos(2\pi ut + p)$. Three arrows point from labels to specific parts of the equation: an arrow from 'Amplitude' points to the term A ; an arrow from 'Frequency' points to the term $2\pi u$; and an arrow from 'Phase (shift)' points to the term p .

And a bit about transforms...

Functions as Vectors

- Inner (dot) product between two vectors is the summation of the point-wise product
- Can't we do the same thing with functions?
- For continuous functions, the summation just becomes an integral
- Functions satisfy all of the mathematical requirements for “vectors”
- Can we transform *functions*?

$$\mathbf{u} \cdot \mathbf{v} = \sum_j \mathbf{u}[j] \mathbf{v}[j]$$

$$f(t) \cdot g(t) = \int_{-\infty}^{\infty} f(t) g(t) dt$$

Sinusoidal Basis Functions

- One set of orthonormal basis functions is the set of sines and cosines of different frequencies
- Let's use these as the basis functions for a transform

$$c(u) = \int_{-\infty}^{\infty} f(t) \cos(2\pi ut) dt$$

$$s(u) = \int_{-\infty}^{\infty} f(t) \sin(2\pi ut) dt$$

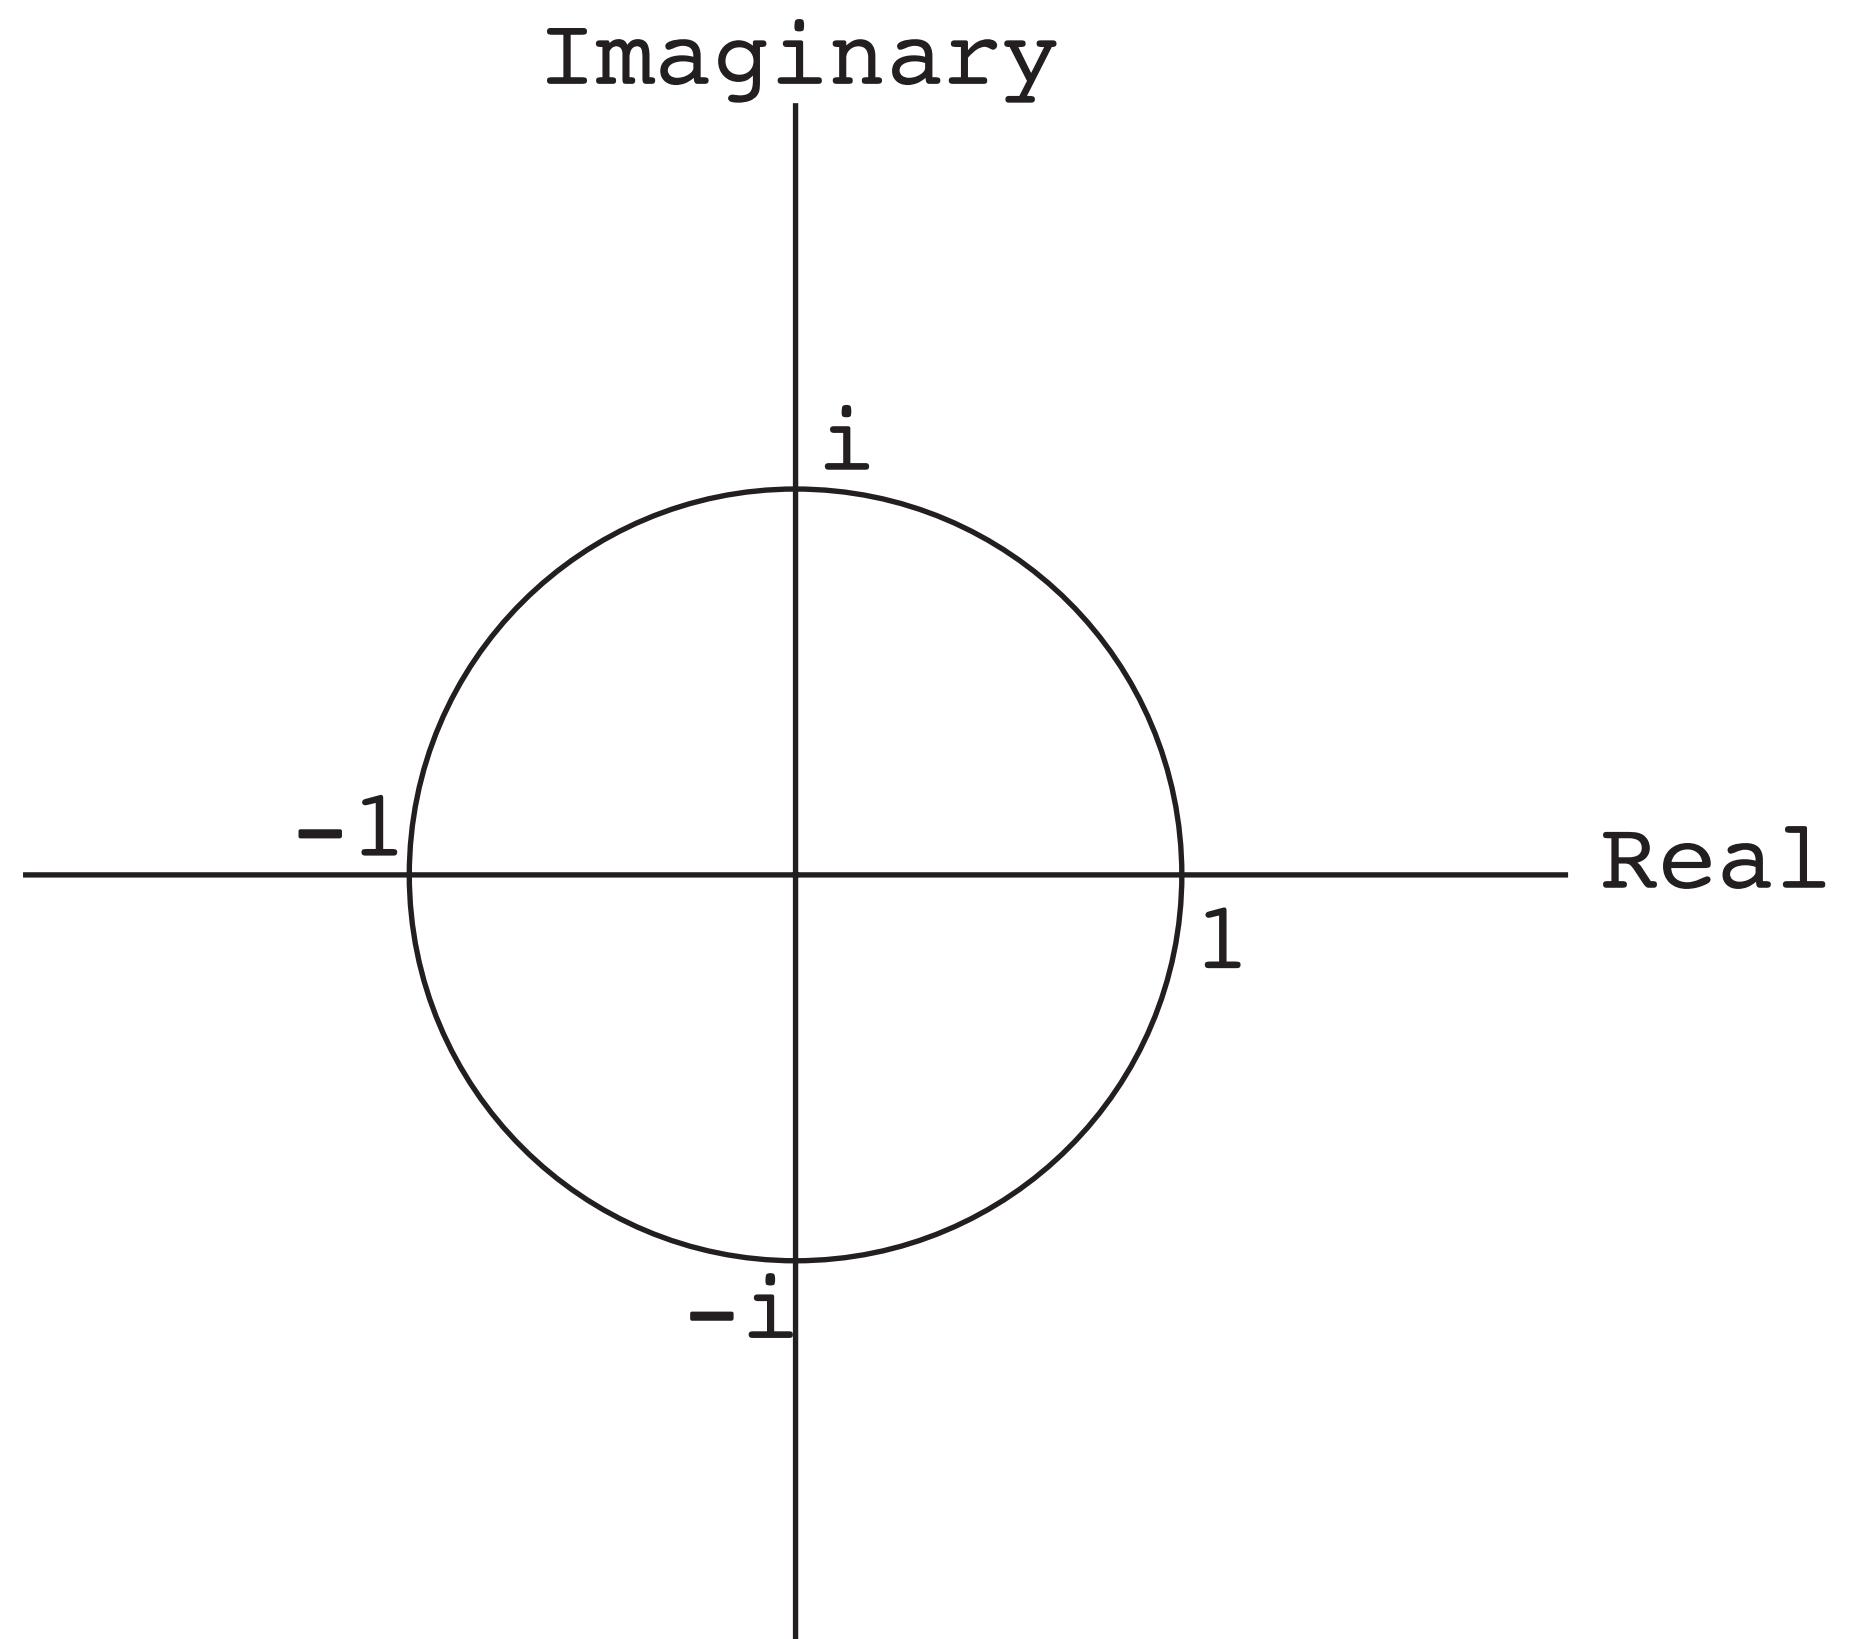
A brief detour...

Complex Numbers

- A complex number is the sum of a real number and an imaginary number:

$$a + bi$$

- Think of as a point on the *complex plane*



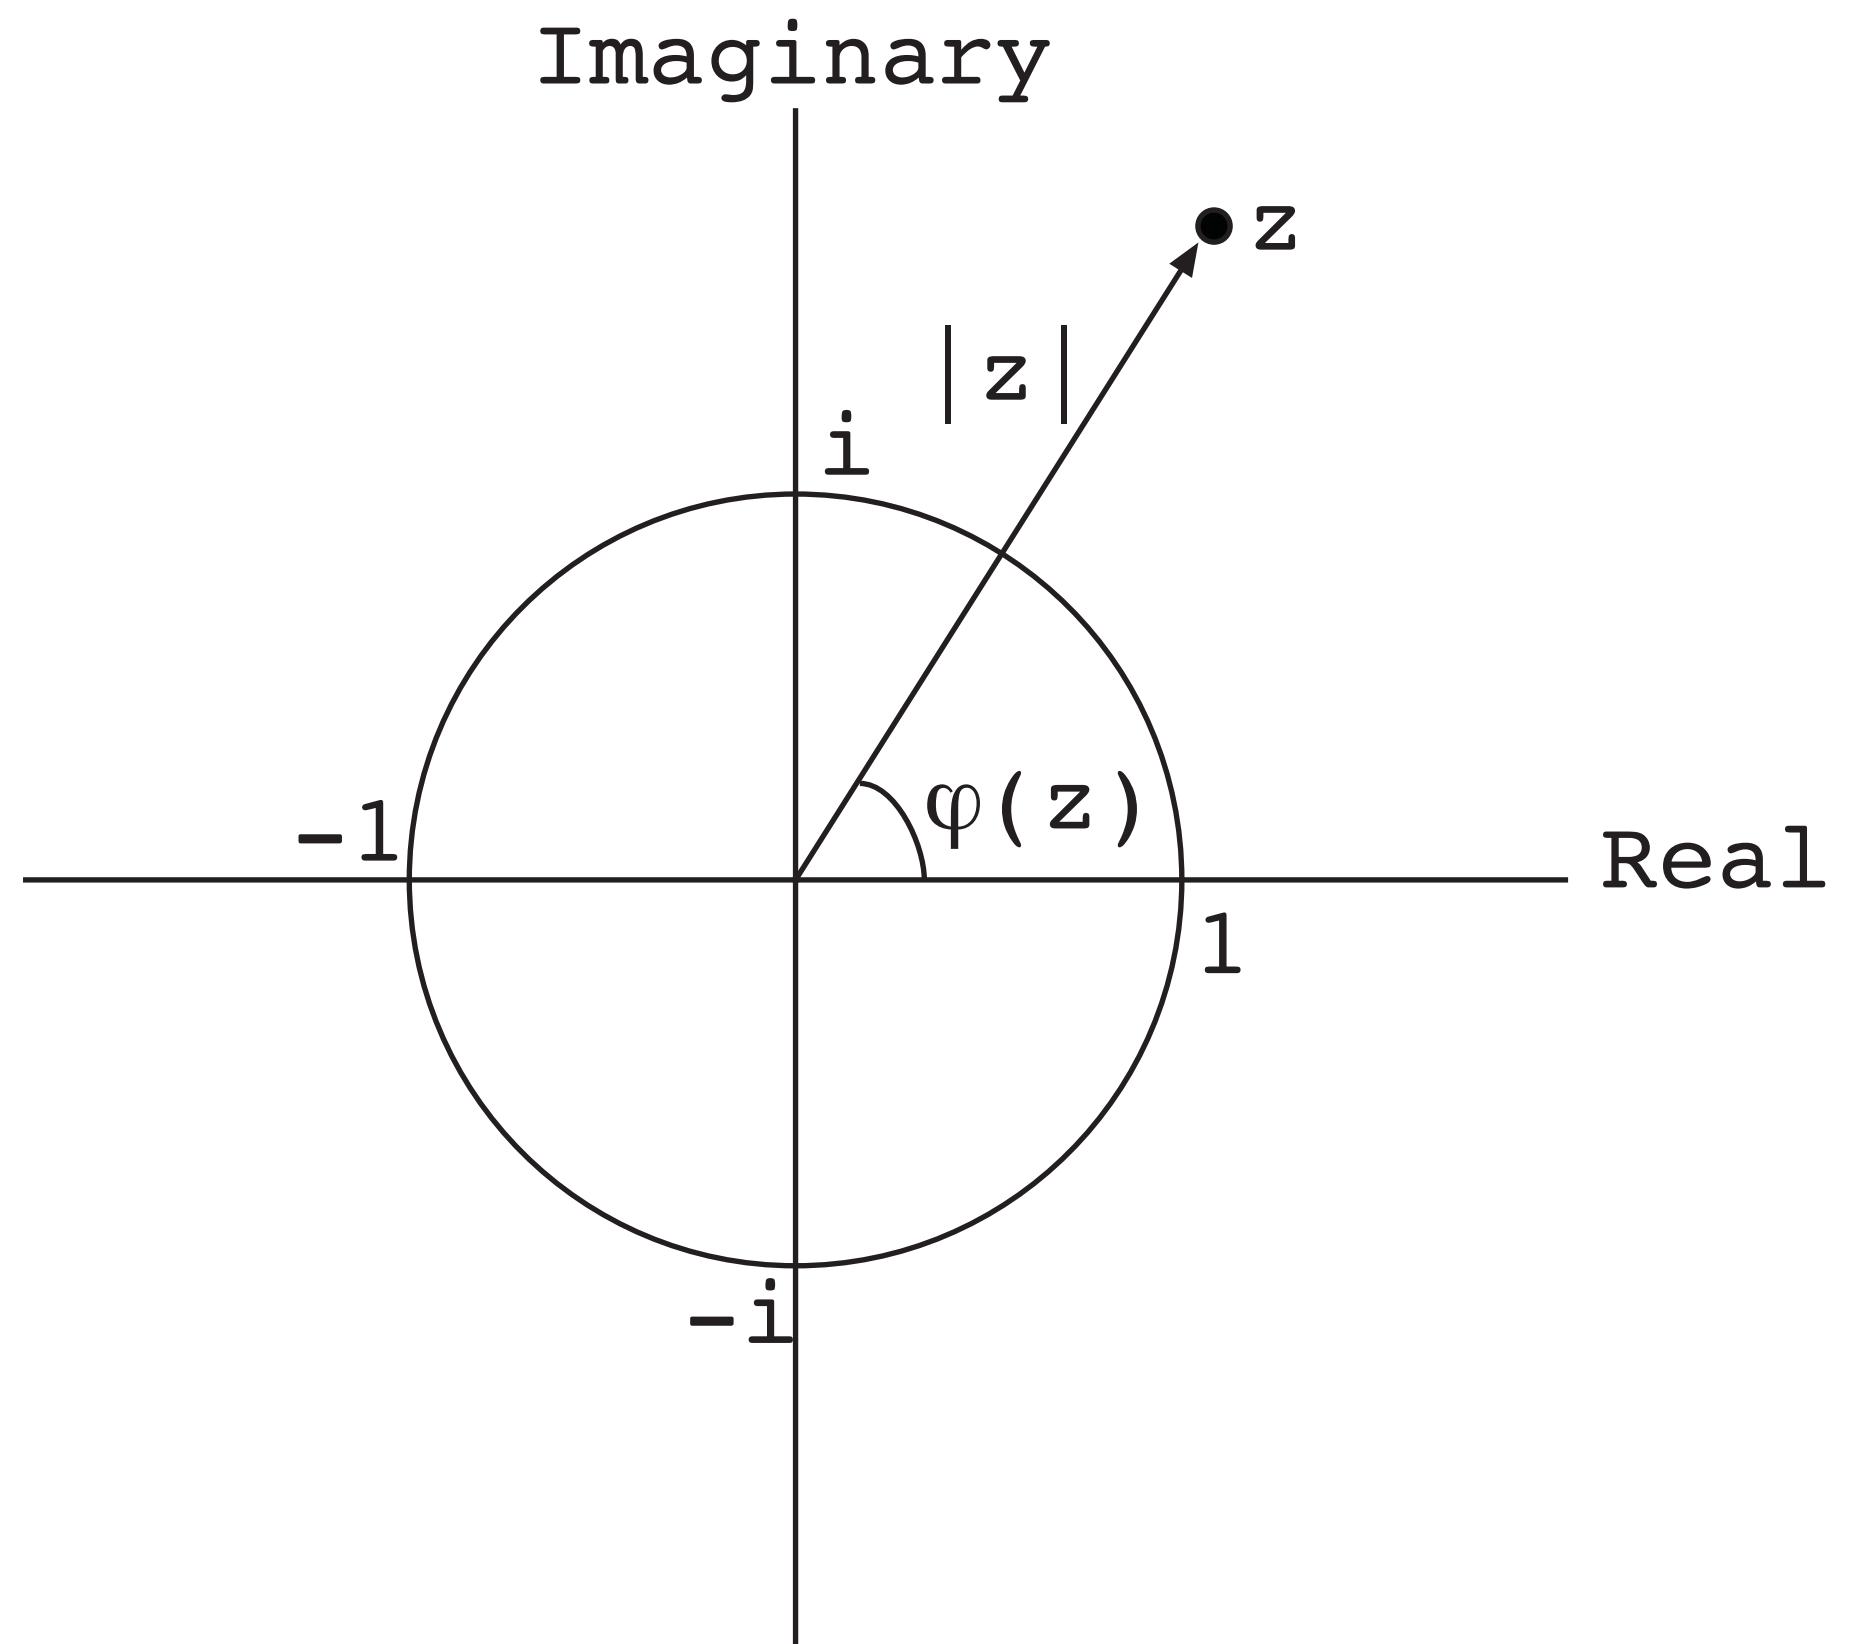
Complex Numbers

- Often useful to think about complex numbers in polar coordinates
- The magnitude is

$$\|a + bi\| = \sqrt{a^2 + b^2}$$

- And the phase (angle) is

$$\phi(a + bi) = \tan^{-1} \left(\frac{b}{a} \right)$$



So why do we care about
complex numbers?

Adding Sinusoids

- When you add some amount of a cosine to another amount of a sine of the same frequency, you get a sinusoid of the same frequency

$$a \cos(2\pi ut) + b \sin(2\pi ut)$$

- If you encode the amount of cosine and the amount of sine as a complex number...

$$z = a + bi$$

- The amplitude of the sinusoid is the magnitude of the complex number

$$\|z\| \cos(2\pi ut + \phi(z))$$

- The phase of the sinusoid is the phase of the complex number

Two Ways to Make Sinusoids

- We thus have two ways to make *the same sinusoid*:
 - Mix a cosine and a sine with specific weights
 - Start with a cosine and
 - Stretch it by the square root of the sum of the squares of the sine and cosine weights
 - Shift it by the arctangent of the ratio of the sine weight to the cosine weight

Sinusoidal Weights as Complex Numbers

- Useful to encode the coefficients from projecting onto sinusoidal basis functions as a single complex number

$$c(u) = \int_{-\infty}^{\infty} f(t) \cos(2\pi ut) dt$$

$$s(u) = \int_{-\infty}^{\infty} f(t) \sin(2\pi ut) dt$$

$$F(u) = c(u) - i s(u)$$

* Note the minus sign here when combining the two, this comes from doing linear algebra with complex quantities

The Fourier Transform

- An encode as a function $f(t)$ as a weighted sum of sines and cosines where the weights are given by a complex-valued function $F(u)$
- Can think of as an operator
- Written as

$$F(u) = \mathcal{F}(f(t))$$

$$c(u) = \int_{-\infty}^{\infty} f(t) \cos(2\pi ut) dt$$

$$s(u) = \int_{-\infty}^{\infty} f(t) \sin(2\pi ut) dt$$

$$F(u) = c(u) - i s(u)$$

The Inverse Fourier Transform

- Can invert the transformation to get $f(t)$ from $F(u)$ by simply adding the sines and cosines back up with the respective weights

- Can also think of as an operator

- Written as

$$f(t) = \mathcal{F}^{-1}(F(u))$$

$$F(u) = a(u) + i b(u)$$

$$f(t) = \int_{-\infty}^{\infty} a(u) \cos(2\pi ut) du$$

$$+ \int_{-\infty}^{\infty} b(u) \sin(2\pi ut) du$$

The Discrete Fourier Transform

- What about a finite-length, sampled signal?
- The Fourier Transform assumes an infinite-length signal (impossible to work with on computers)
- What should we assume about the signal before and after we record it?
- Since we're decomposing it into a sum of periodic functions,
let's assume it's one period of an infinite periodic function

Period Functions

- For a periodic signal with period N units, all of the underlying frequencies must also repeat over the period N
- So, each component frequency must be a multiple of the frequency of the periodic signal itself:
- There are no more than N components for a signal with period N samples!

$$\frac{0}{N}, \frac{1}{N}, \frac{2}{N}, \frac{3}{N}, \dots, \frac{N-1}{N}$$

The Discrete Fourier Transform

- Same idea as the Fourier Transform
- For a finite with N samples:
 - N discrete frequencies u/N
 - Sum over the N discrete samples
- What would the code look like?
- What is the complexity?

$$c[u] = \frac{1}{N} \sum_{t=0}^{N-1} f[t] \cos(2\pi ut/N)$$

$$s[u] = \frac{1}{N} \sum_{t=0}^{N-1} f[t] \sin(2\pi ut/N)$$

$$F[u] = c[u] - i s[u]$$

The Fast Fourier Transform

- The Fast Fourier Transform does exactly the equivalent mathematically
- Divide-and-conquer algorithm provides greater efficiency

DFT $O(N^2)$

FFT $O(N \log N)$

Using the FFT

- Pass in an array of length N
- Returns back an array of length N of type complex
 - The real part of each number is how much of a cosine of that frequency there is
 - The imaginary part of each number is how much of a sine of that frequency there is

Interpreting the Complex Numbers

- Can think of in Cartesian form:
 - The real part of each number is how much of a cosine of that frequency
 - The imaginary part of each number is how much of a sine of that frequency
- Or in polar form:
 - The magnitude of each number is how much of that frequency there signal
 - The phase of each number is the relative shift (from a cosine) of each sinusoid

Negative Frequencies

- Because of their repetitive nature
 - cosines are symmetric
 - sines are antisymmetric
 - discrete ones repeat modulo N

$$\cos(-2\pi u/N) = \cos(2\pi u/N)$$

$$\sin(-2\pi u/N) = -\sin(2\pi u/N)$$

$$N - u \equiv -u \pmod{N}$$

Storage of the DFT Results

Index	0	1	2	...	N/2	...	N - 2	N - 1
Freq.	0	1	2	...	$\pm N/2$...	-2	-1

Because of symmetry, the last half of the array is a mirrored (negated) copy of the first half

$$\cos(-2\pi u/N) = \cos(2\pi u/N)$$
$$\sin(-2\pi u/N) = -\sin(2\pi u/N)$$

Implications

- Because of symmetry
 - Last half of the real part is a mirrored copy of the first half
 - Last half of the imaginary part is a negated mirrored copy
- There are really only N basis functions, as there should be, not $2N$
- Only really computing frequencies up to $N/2$
 - Twice the number of samples are there are frequencies
 - “Sample at twice the highest frequency in the signal...”
 - This is the same as Shannon’s sampling theorem!

Useful Python Functions

- `np.fft.fft` - forward FFT (complex array in, complex array out)
- `np.fft.ifft` - inverse FFT (complex array in, complex array out)
- `np.absolute` - returns the magnitude of a complex number
(or the absolute value of a real one)
- Normal mathematical operators (+, -, *, /) work on complex numbers as well as integer, floating point ones
- Tip: a real number is a complex one with imaginary part equal to zero

Why Are We Doing This Again?

- The Fourier Transform of a signal lets you analyze the mix of frequencies in it
- Can manipulate the transformed signal and then transform it back! (that's the topic for the next class)

Coming up...

- Examples and properties
- Filtering in the frequency domain
- The Convolution Theorem (wait, what? convolution?)
- 2-D FFT and image filtering in the frequency domain

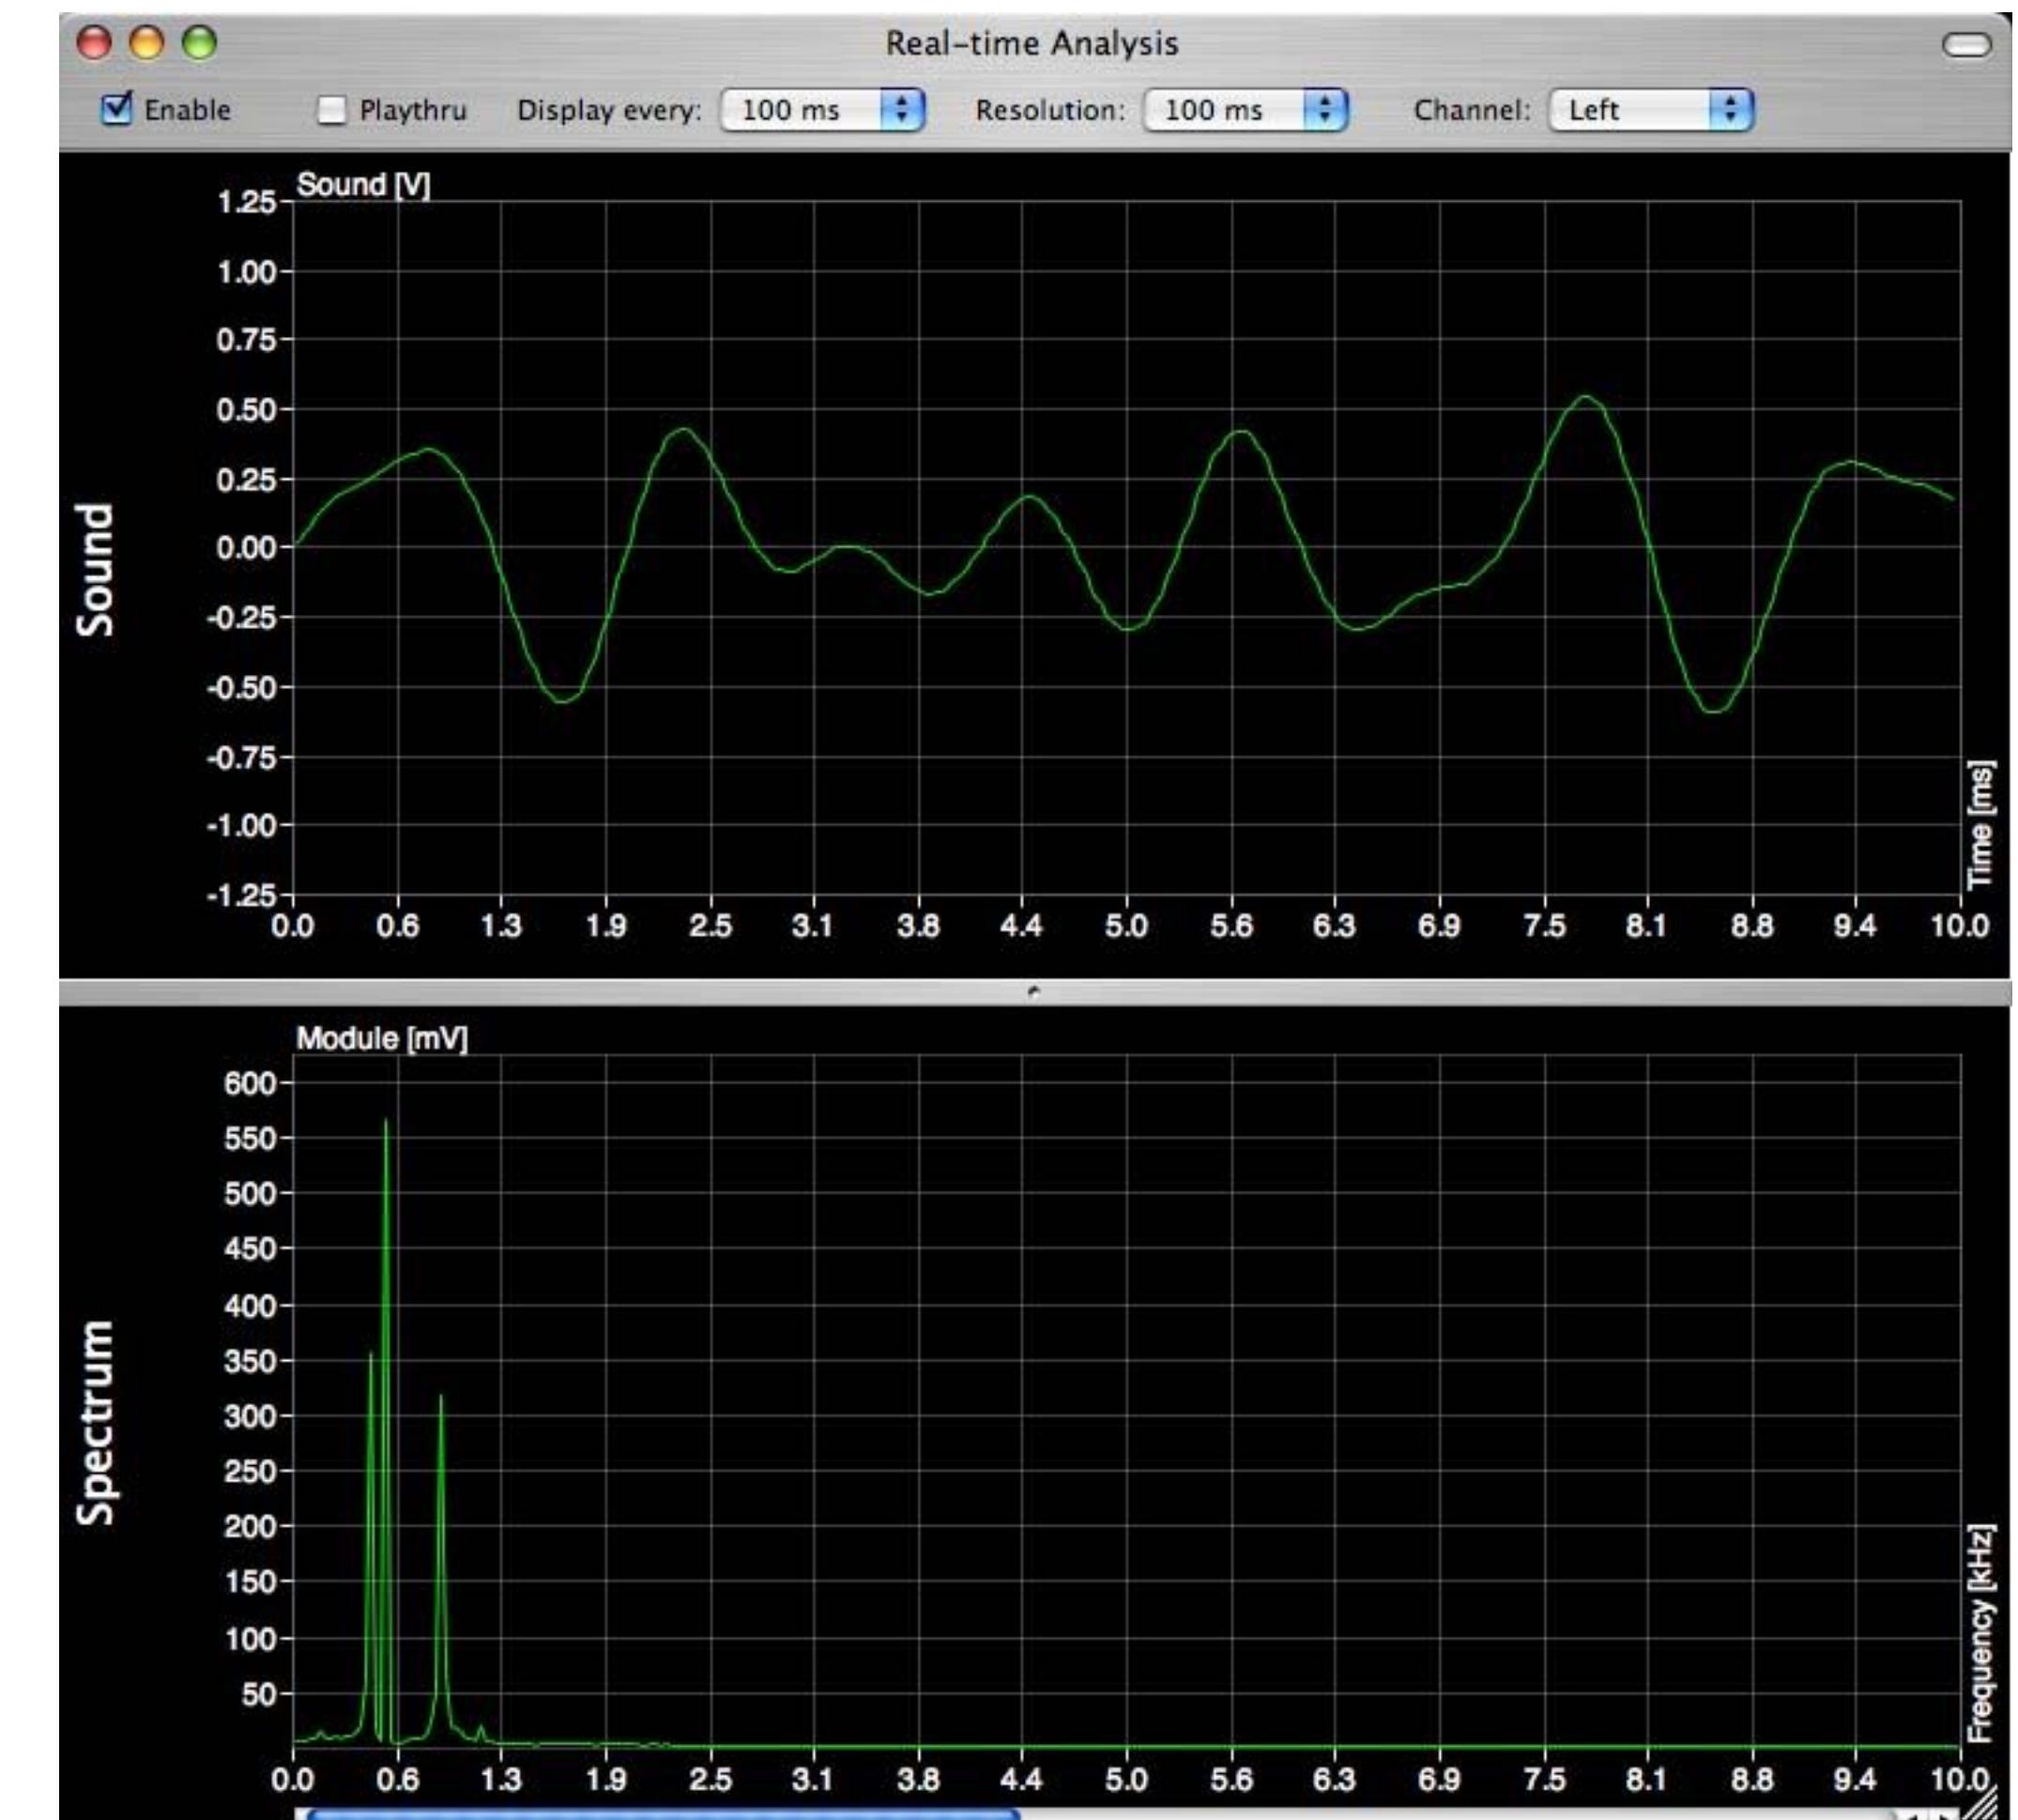


Frequency-Domain Processing

CS 355: Introduction to Graphics and Image Processing

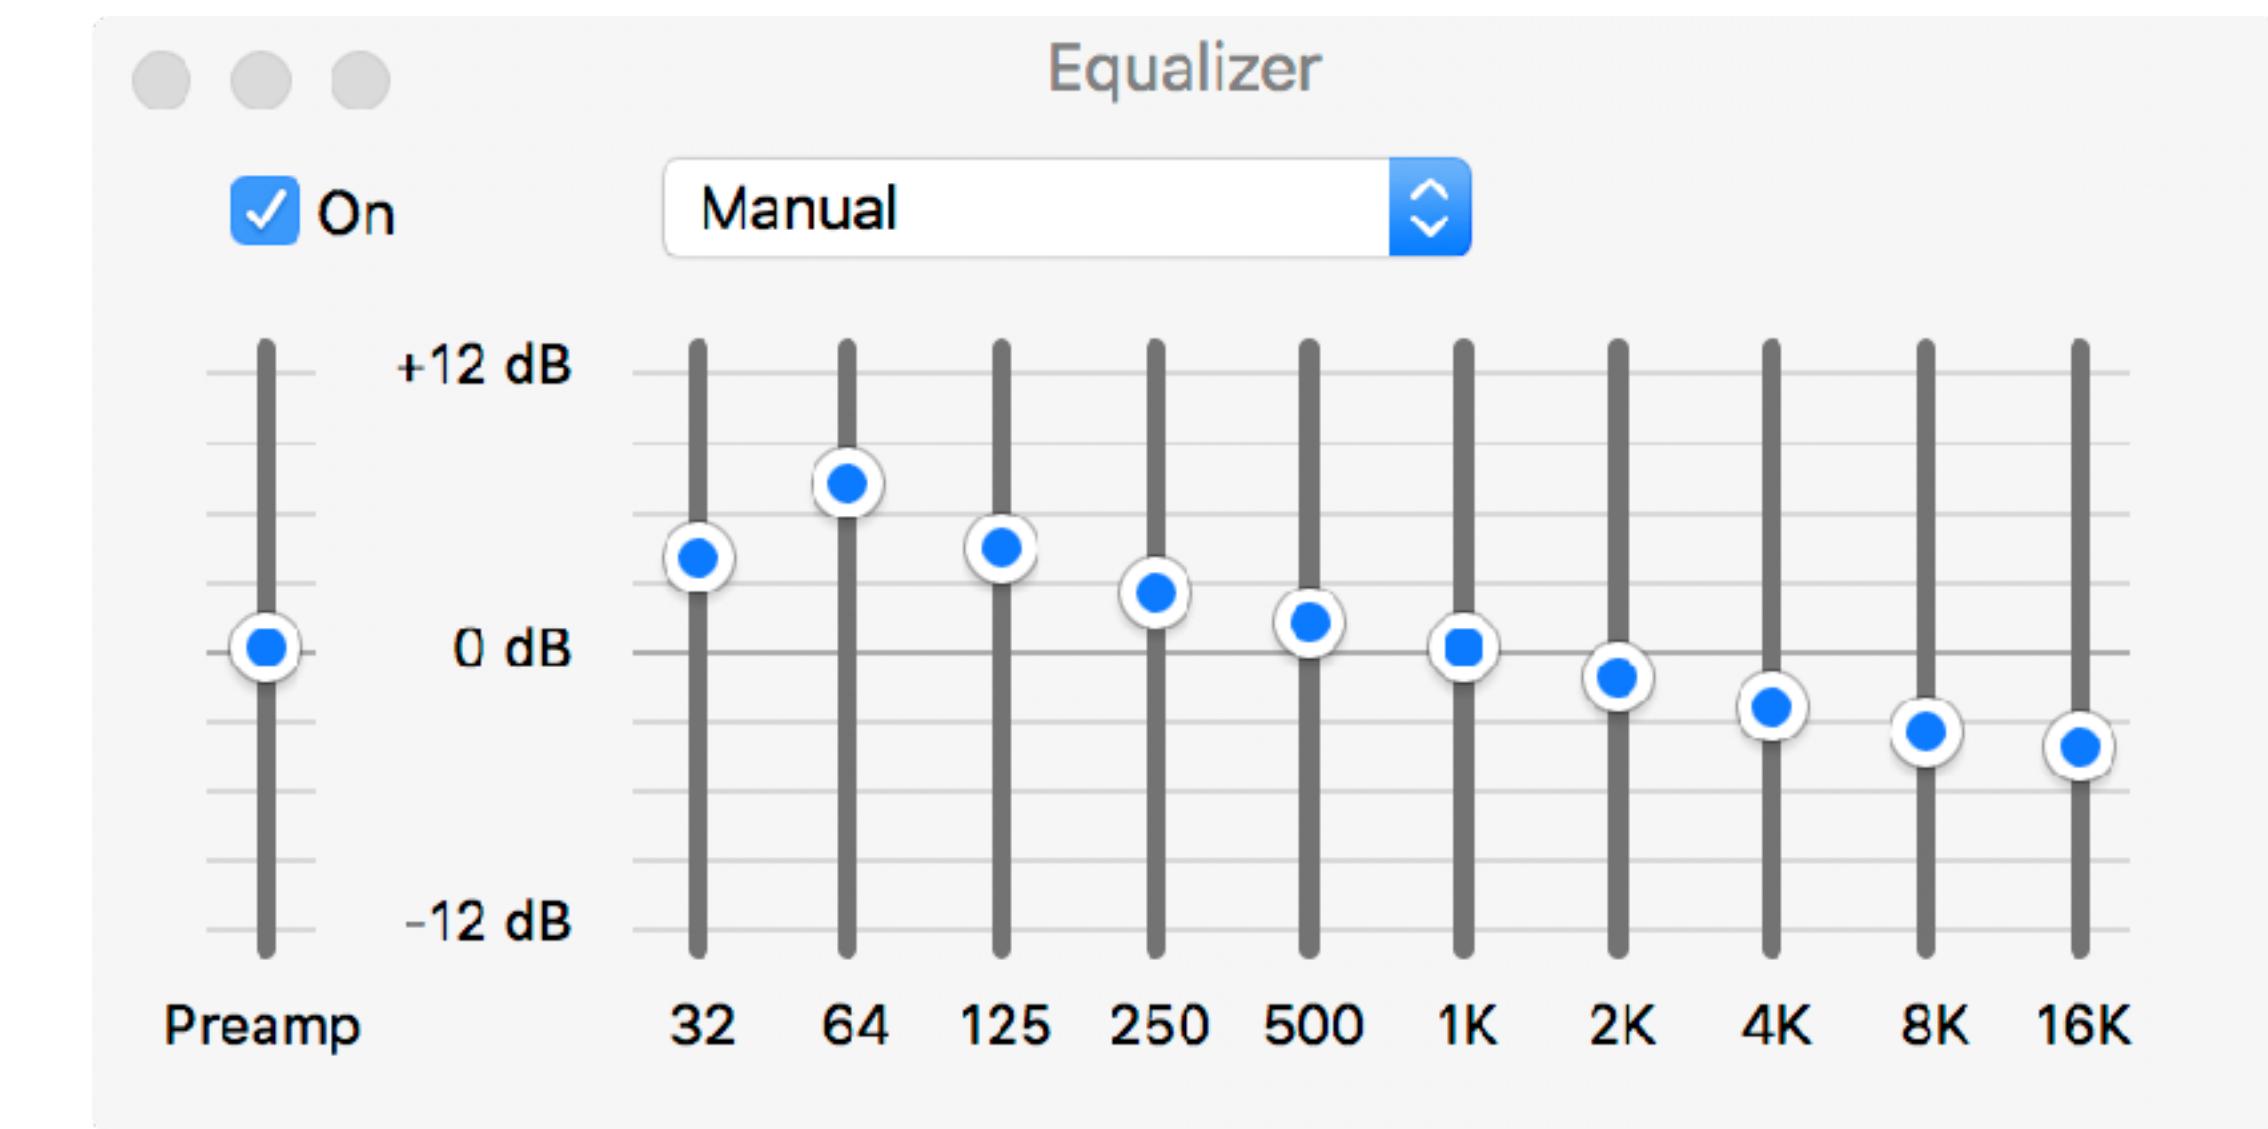
Time Domain vs. Frequency Domain

- For sampled functions of time, we talk about manipulating these samples in the **time domain**
- Earlier in the semester, we learned ways to manipulate images in the **spatial domain**
- Let's talk about manipulating them in the **frequency domain**

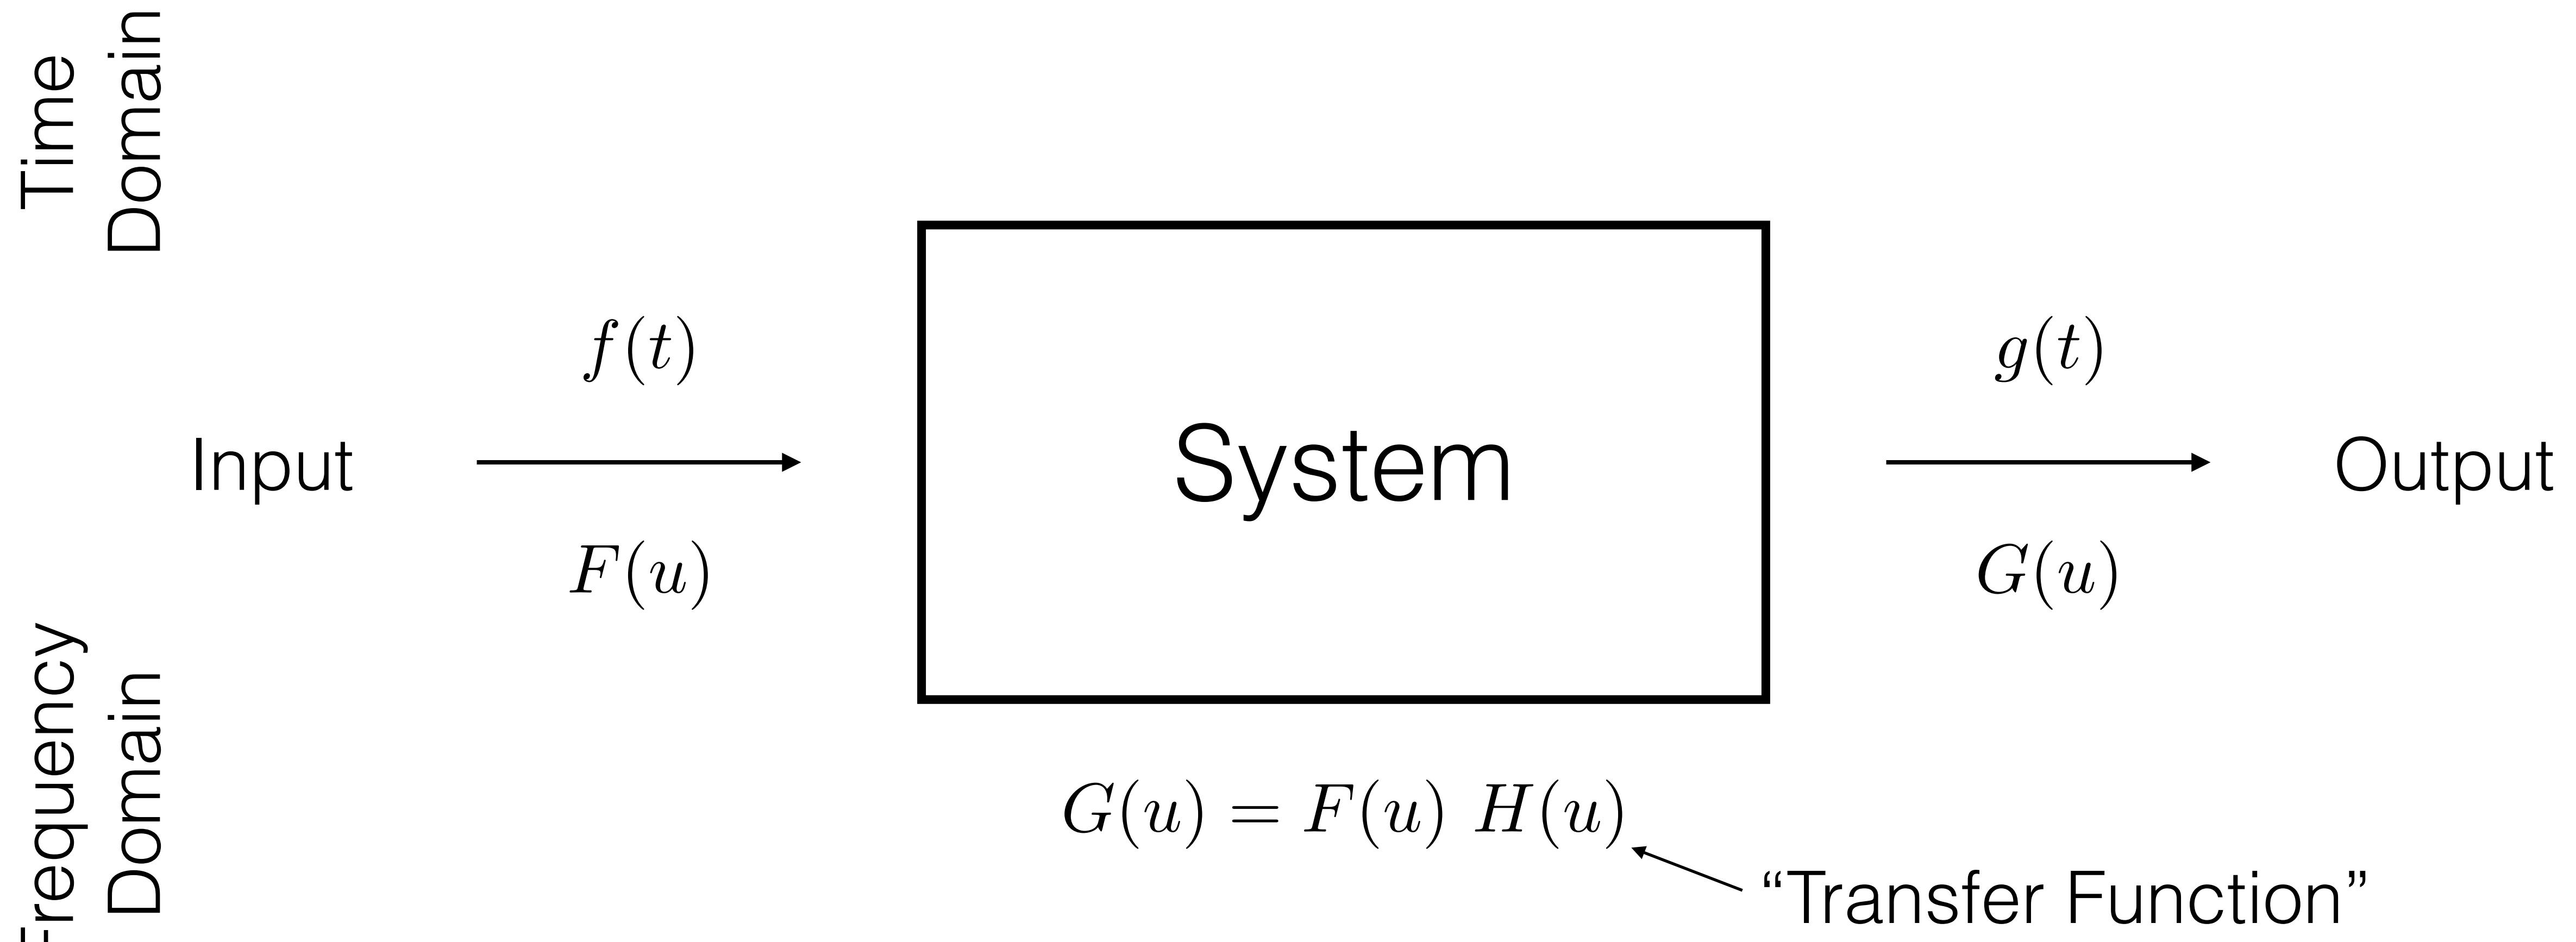


Frequency-Domain Processing

- The simplest way to process signals in the frequency domain is to
 - Selectively amplify (boost) certain frequencies
 - Selectively attenuate (weaken) other frequencies



Common Notation



The Transfer Function

- Intuition: how different frequencies “transfer” from input to output
- Remember:
 - Input $F(u)$ is complex
 - Output $G(u)$ is also complex
- If transfer function $H(u)$ is **real-valued**, it multiplies both the real (cosine) part and the imaginary (sine) part equally
 - Magnitude/amplitude is scaled (amplification/attenuation)
 - Phase is left unchanged

$$G(u) = F(u) H(u)$$

The Transfer Function

- If $H(u)$ is **complex-valued**, it multiplies the real (cosine) part and the imaginary (sine) part differently
 - The magnitudes multiply (amplification/attenuation)
 - The phases add (shifts different frequencies potentially differently)

$$G(u) = F(u) H(u)$$

The Transfer Function

$$G(u) = F(u) \ H(u)$$

Modulation Transfer Function

$$|H(u)|$$

Selectively amplifies/attenuates

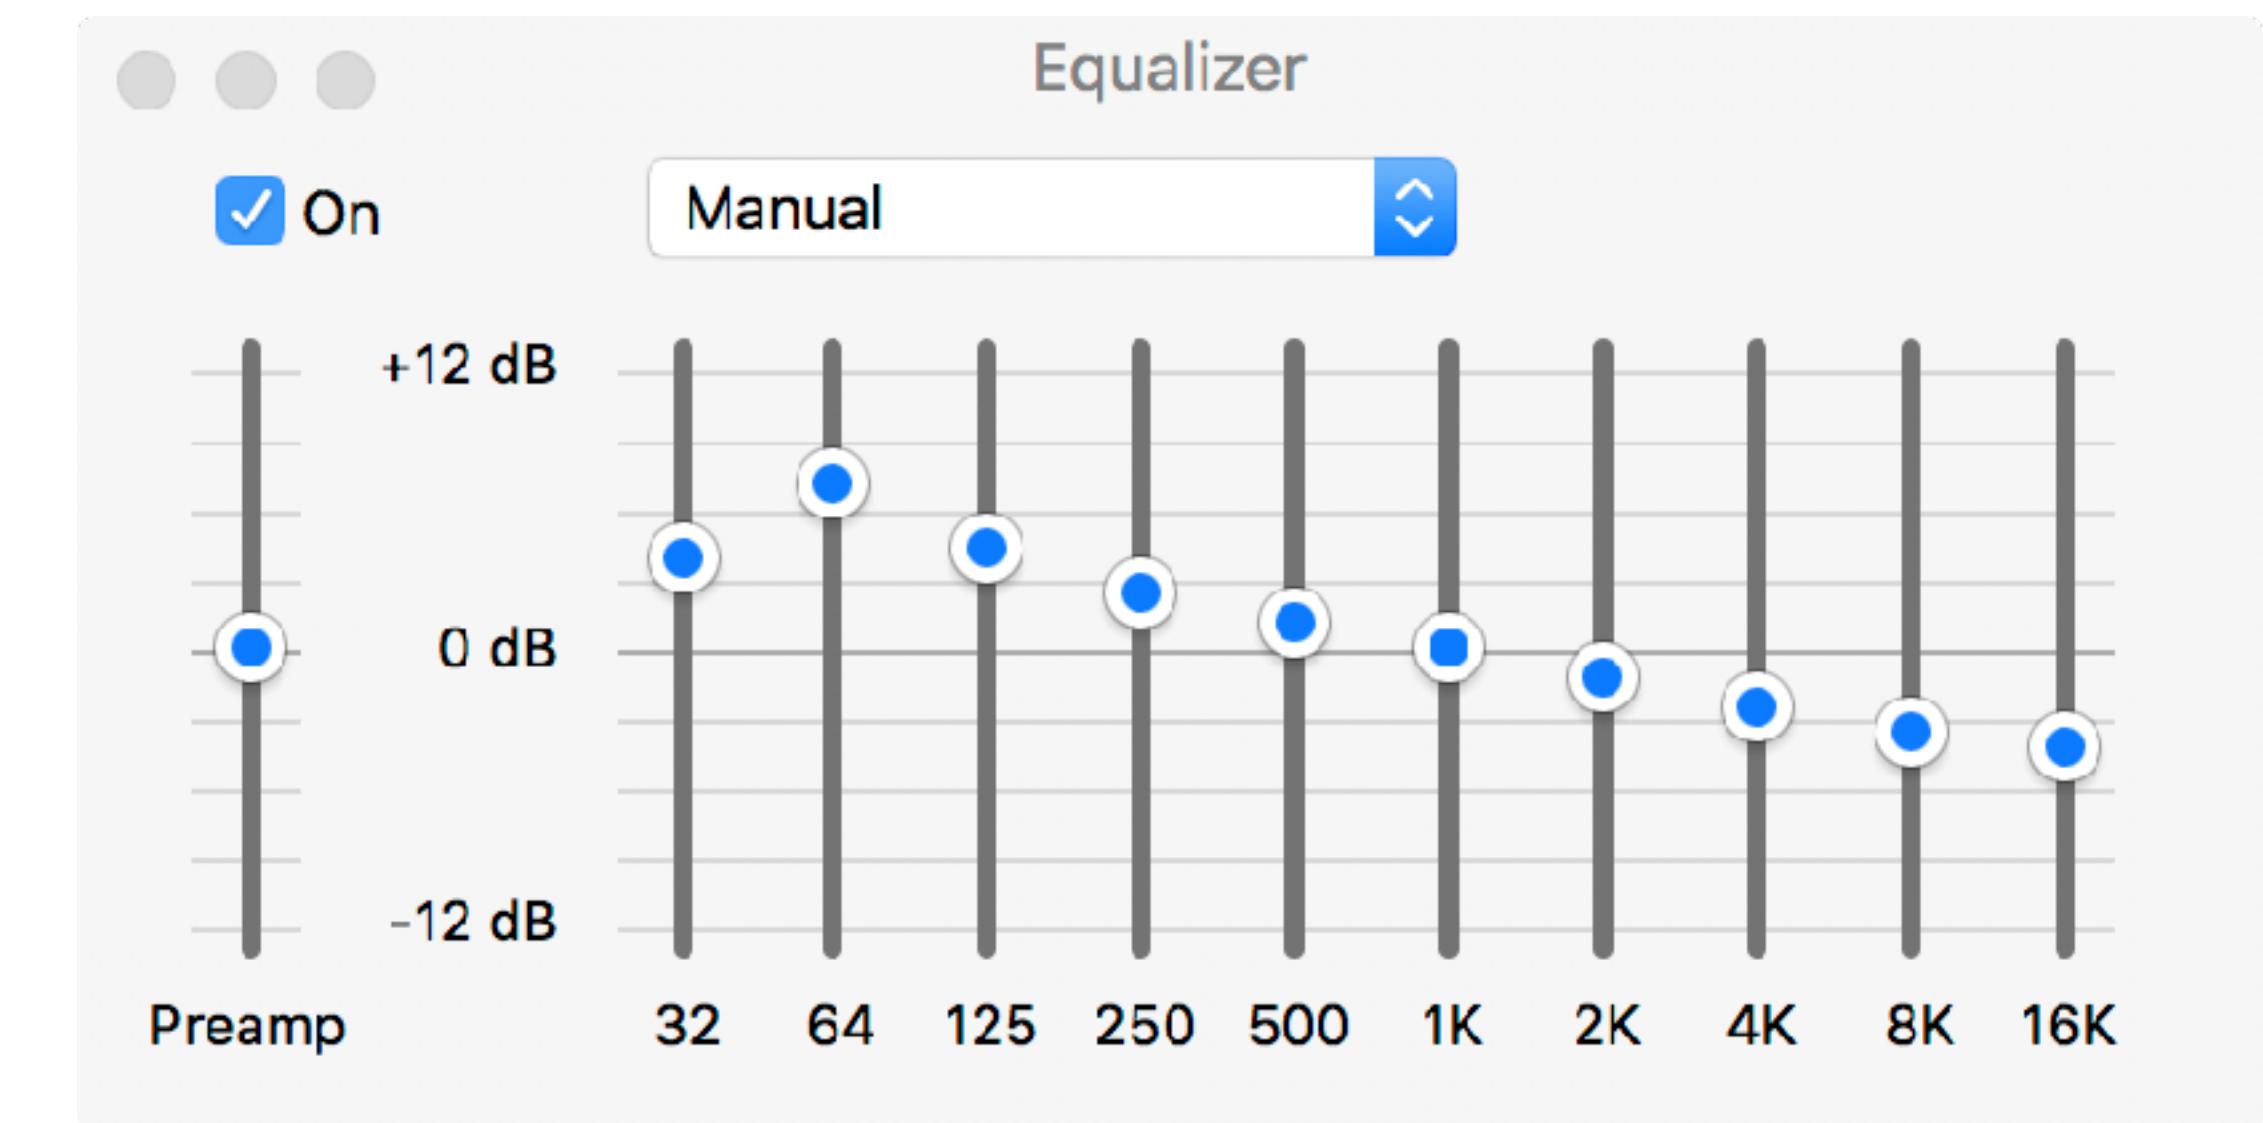
Phase Transfer Function

$$\phi(H(u))$$

Selectively shifts different frequencies

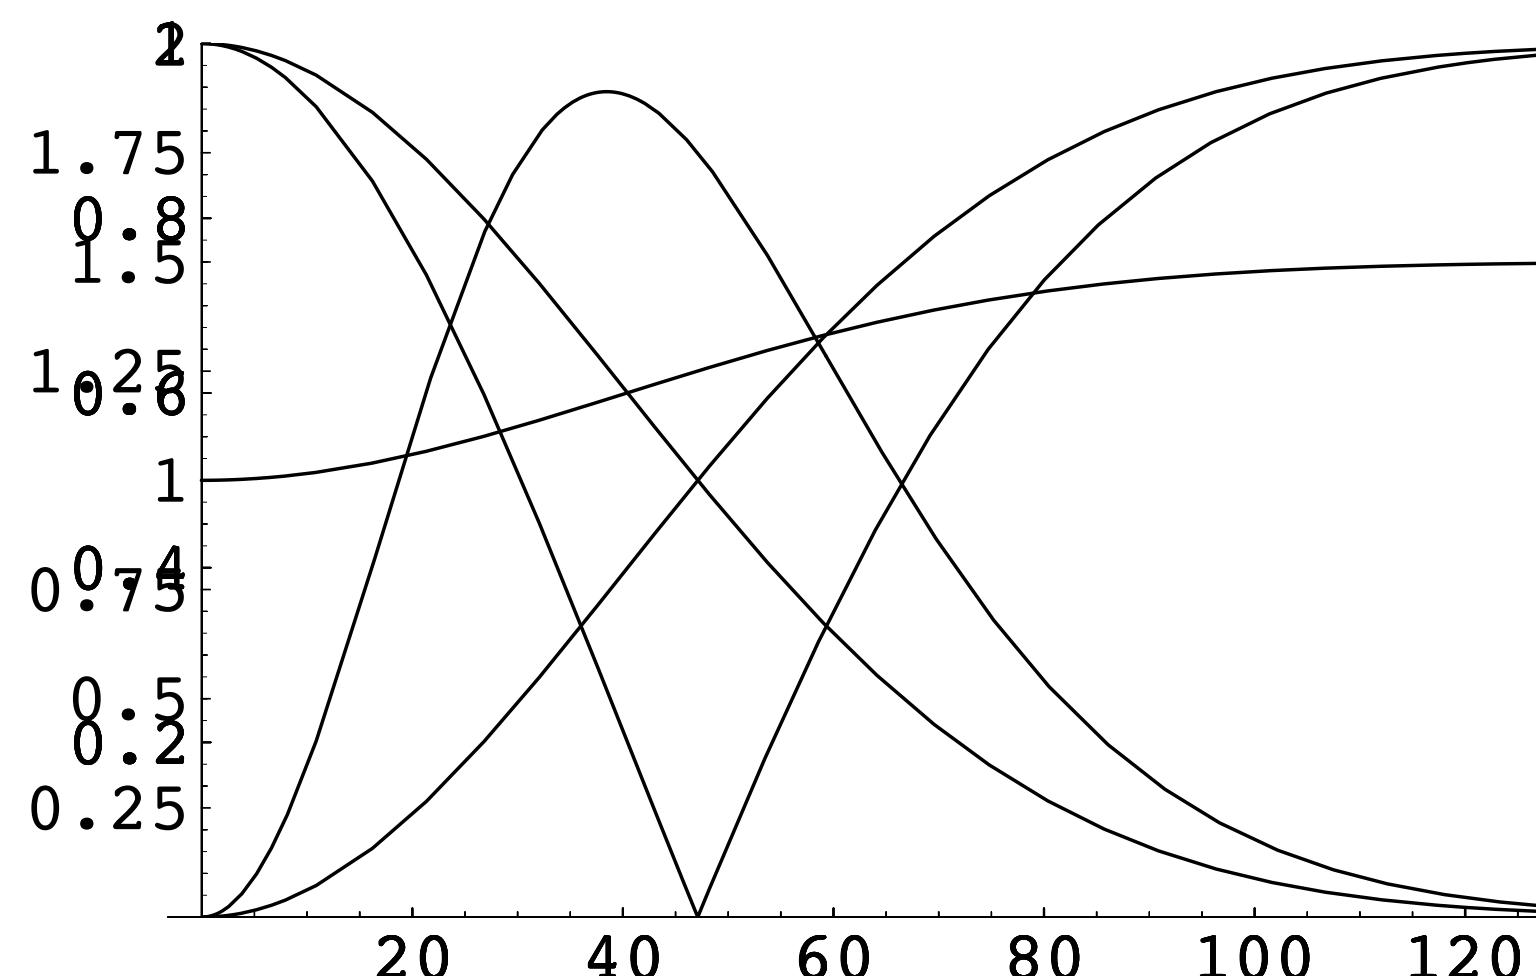
Filtering

- Designing and applying a transfer function to a signal in the frequency domain it is called **filtering**
- You can design $H(u)$ to be whatever you want it to be!



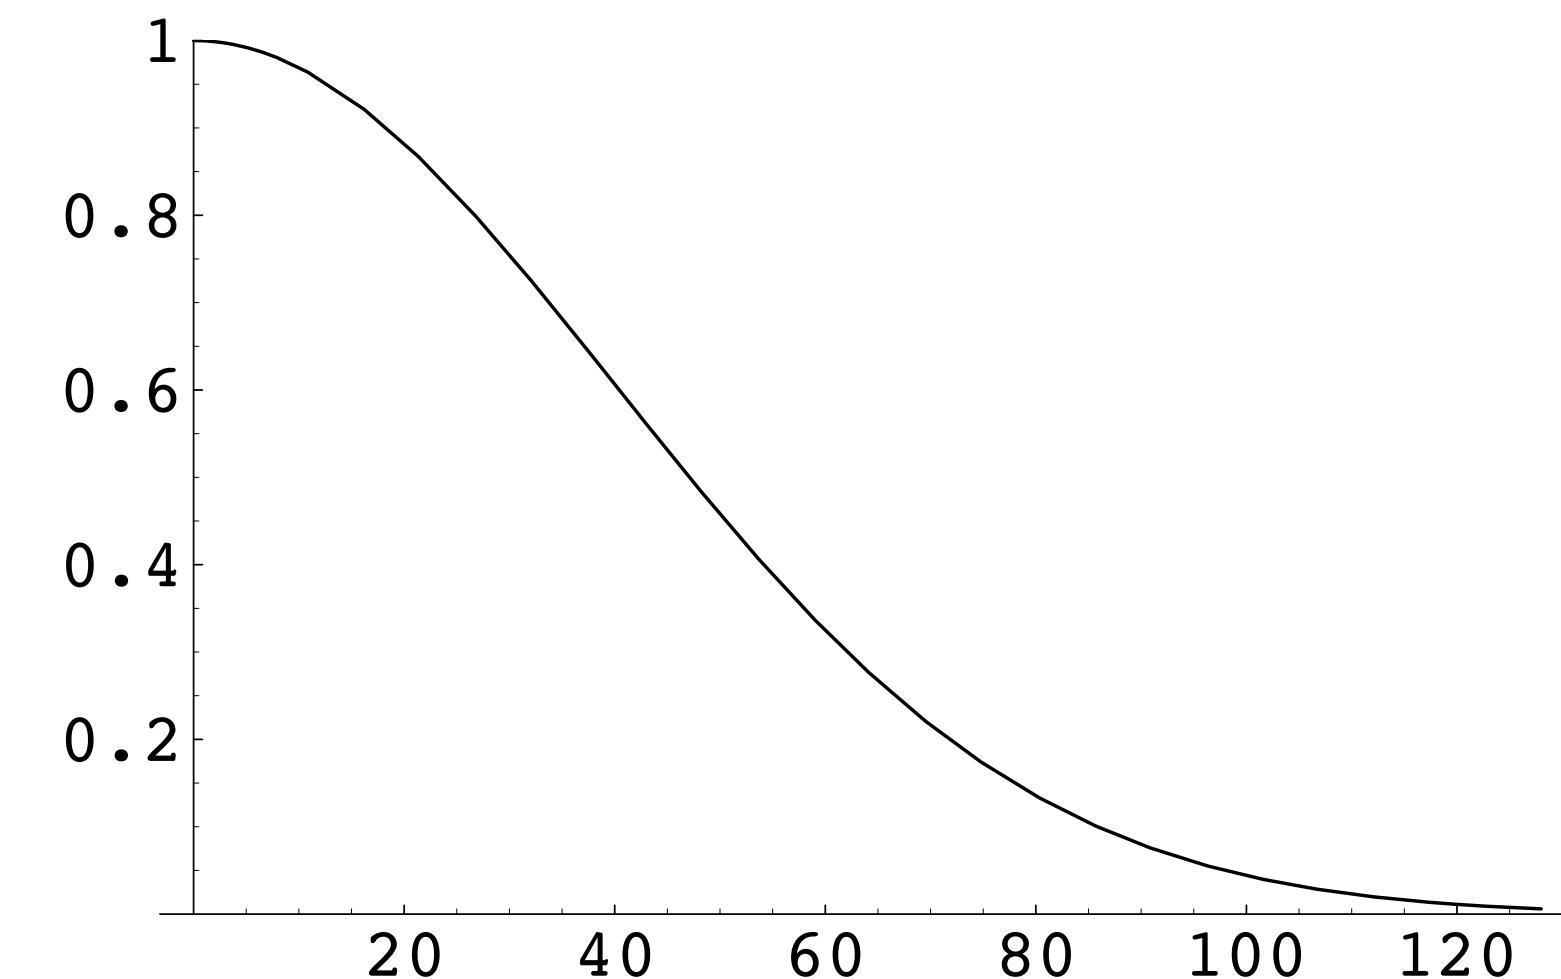
Types of Filters

- | | |
|---------------|---|
| Low-Pass | Keeps low frequencies, attenuates higher ones |
| High-Pass | Keeps high frequencies attenuates lower ones |
| High-Boost | Keeps lower frequencies, boosts higher frequencies |
| Band-Pass | Attenuates frequencies above or below a target band |
| Band-Suppress | Attenuates frequencies in a target band |



Low-Pass Filters

- Lets **low** frequencies **pass** through
- Attenuates higher frequencies
- Examples:
 - Sound through air
 - Sound through other materials
 - RF signals through air
 - Electronic signals through wire

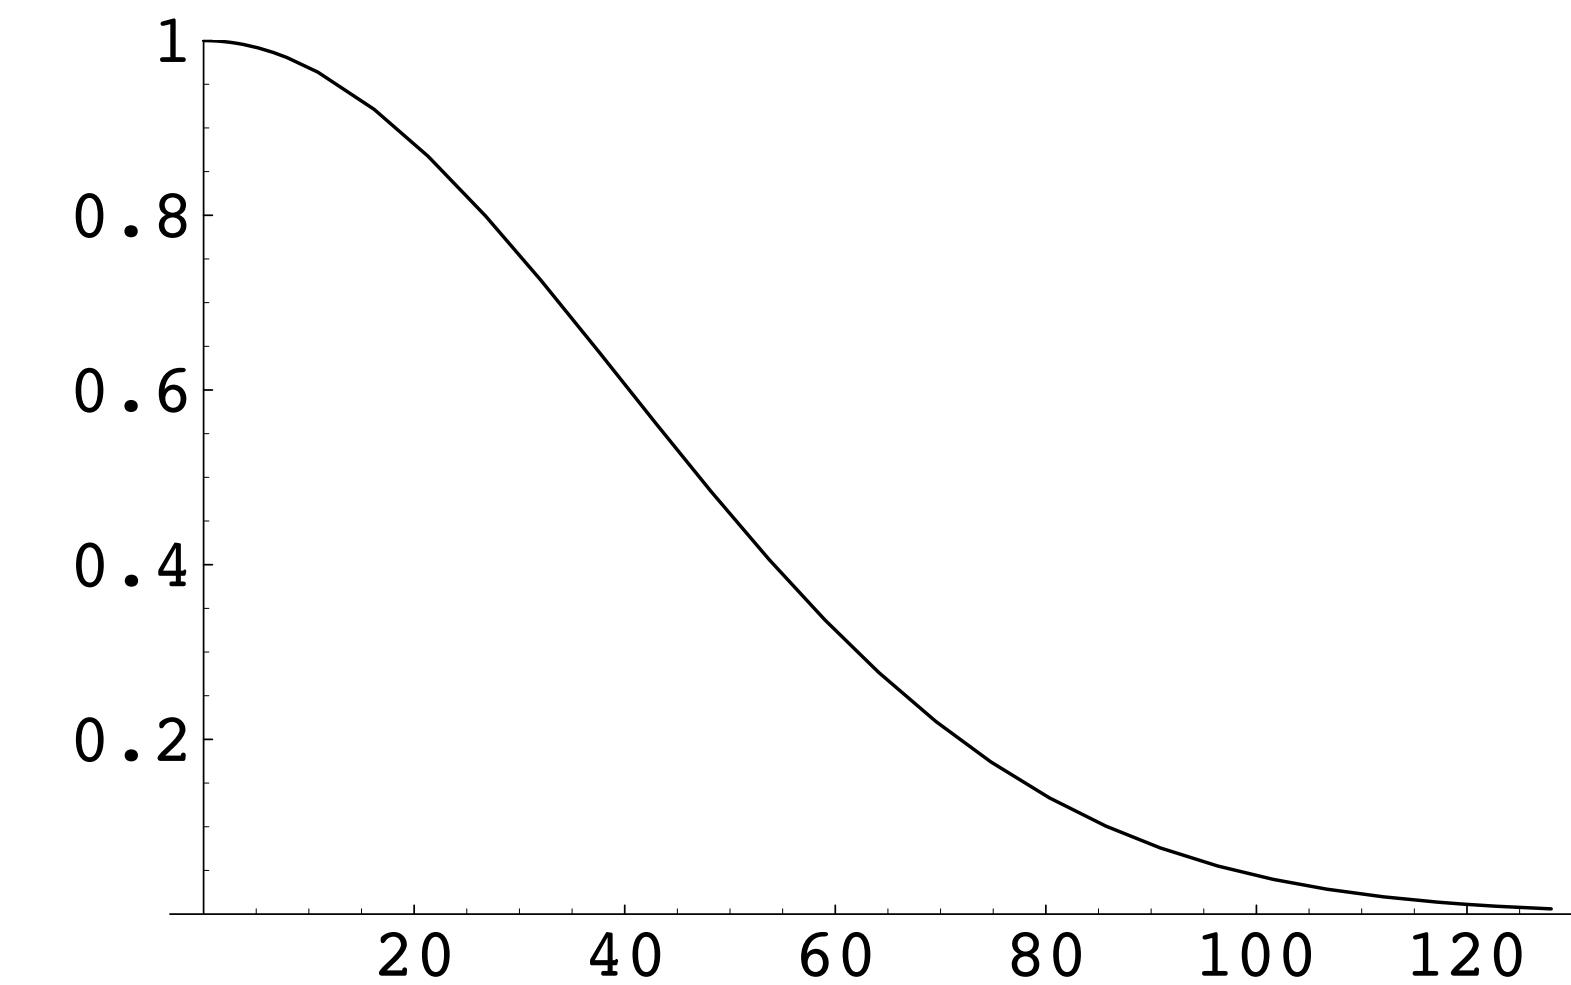


Example: Gaussian Low-Pass Filter

- One common form of a low-pass filter is the Gaussian

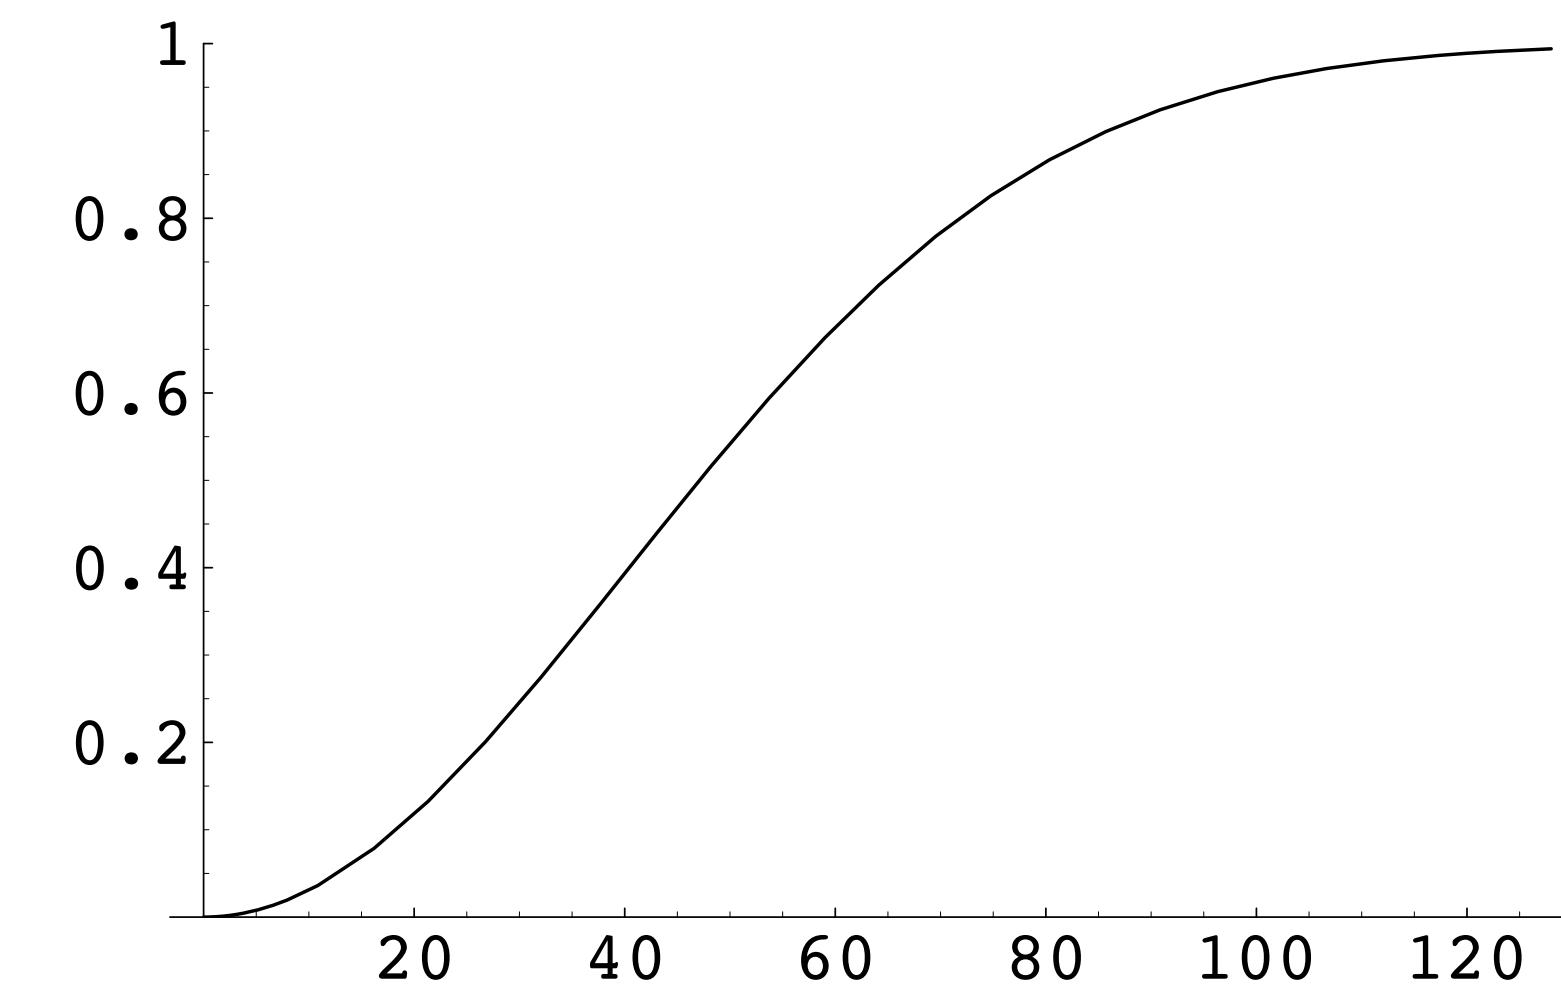
$$H(u) = e^{-\frac{1}{2}u^2/u_c^2}$$

Cutoff frequency



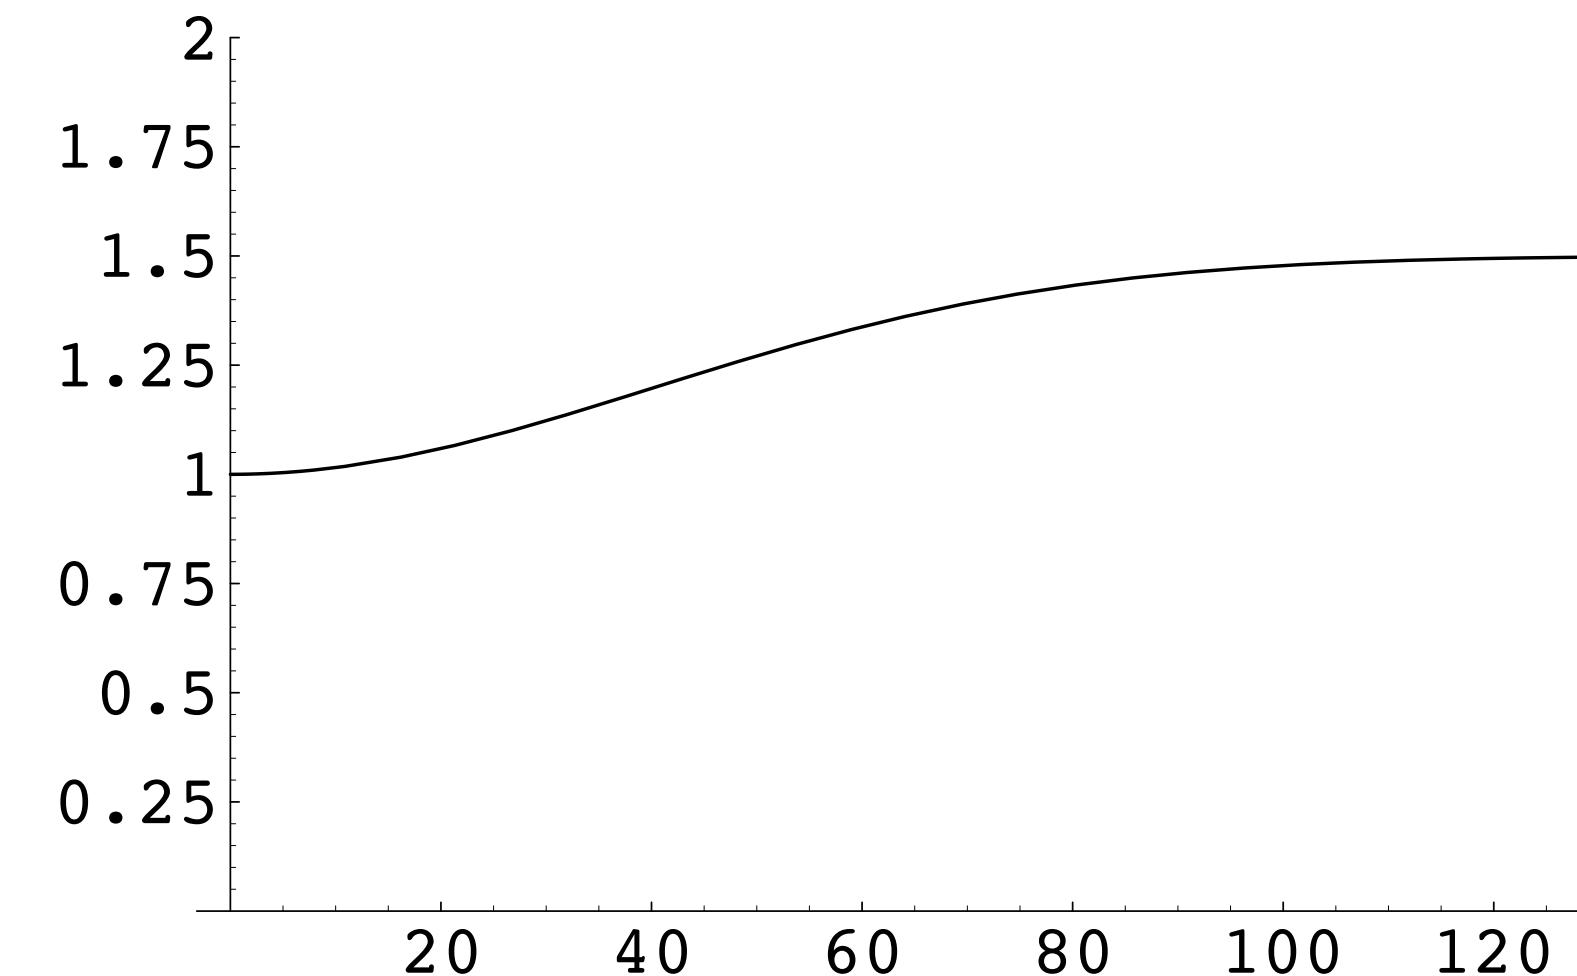
High-Pass Filters

- Lets **high** frequencies **pass** through
- Attenuates lower frequencies
- Not very common in nature but useful in electronics or signal processing
- Often 1.0 minus a low-pass filter



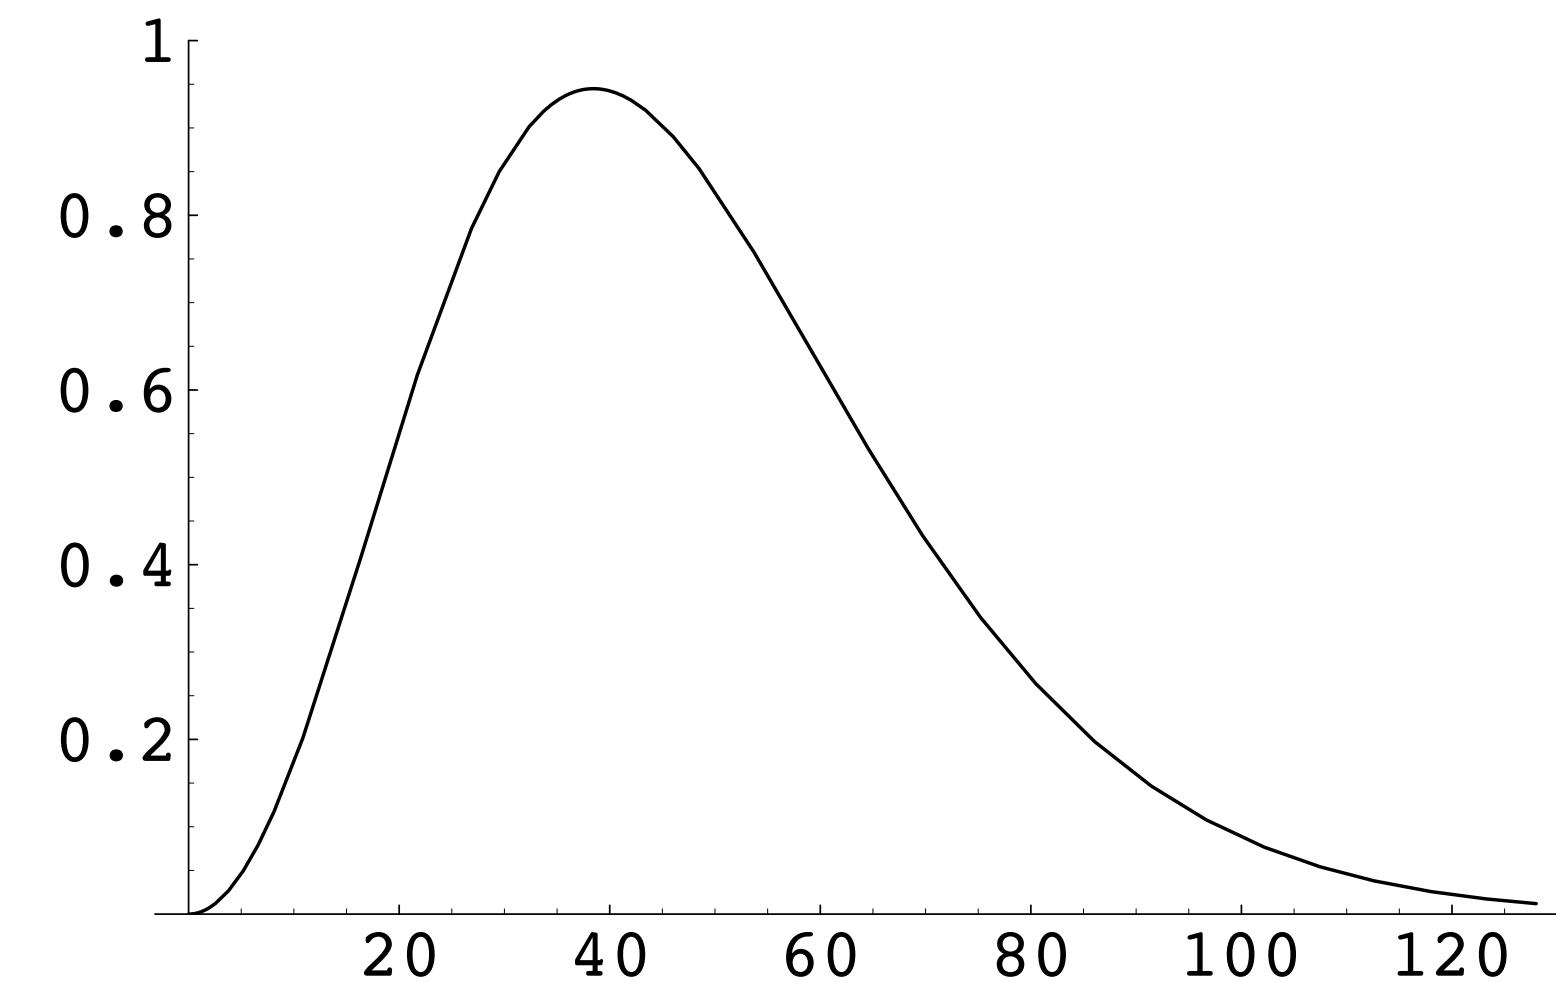
High-Boost Filters

- Boosts high frequencies
- Useful for “undo-ing” low-pass filters
- Can boost any other part of the frequency domain as well
- Often 1.0 plus a high-pass filter



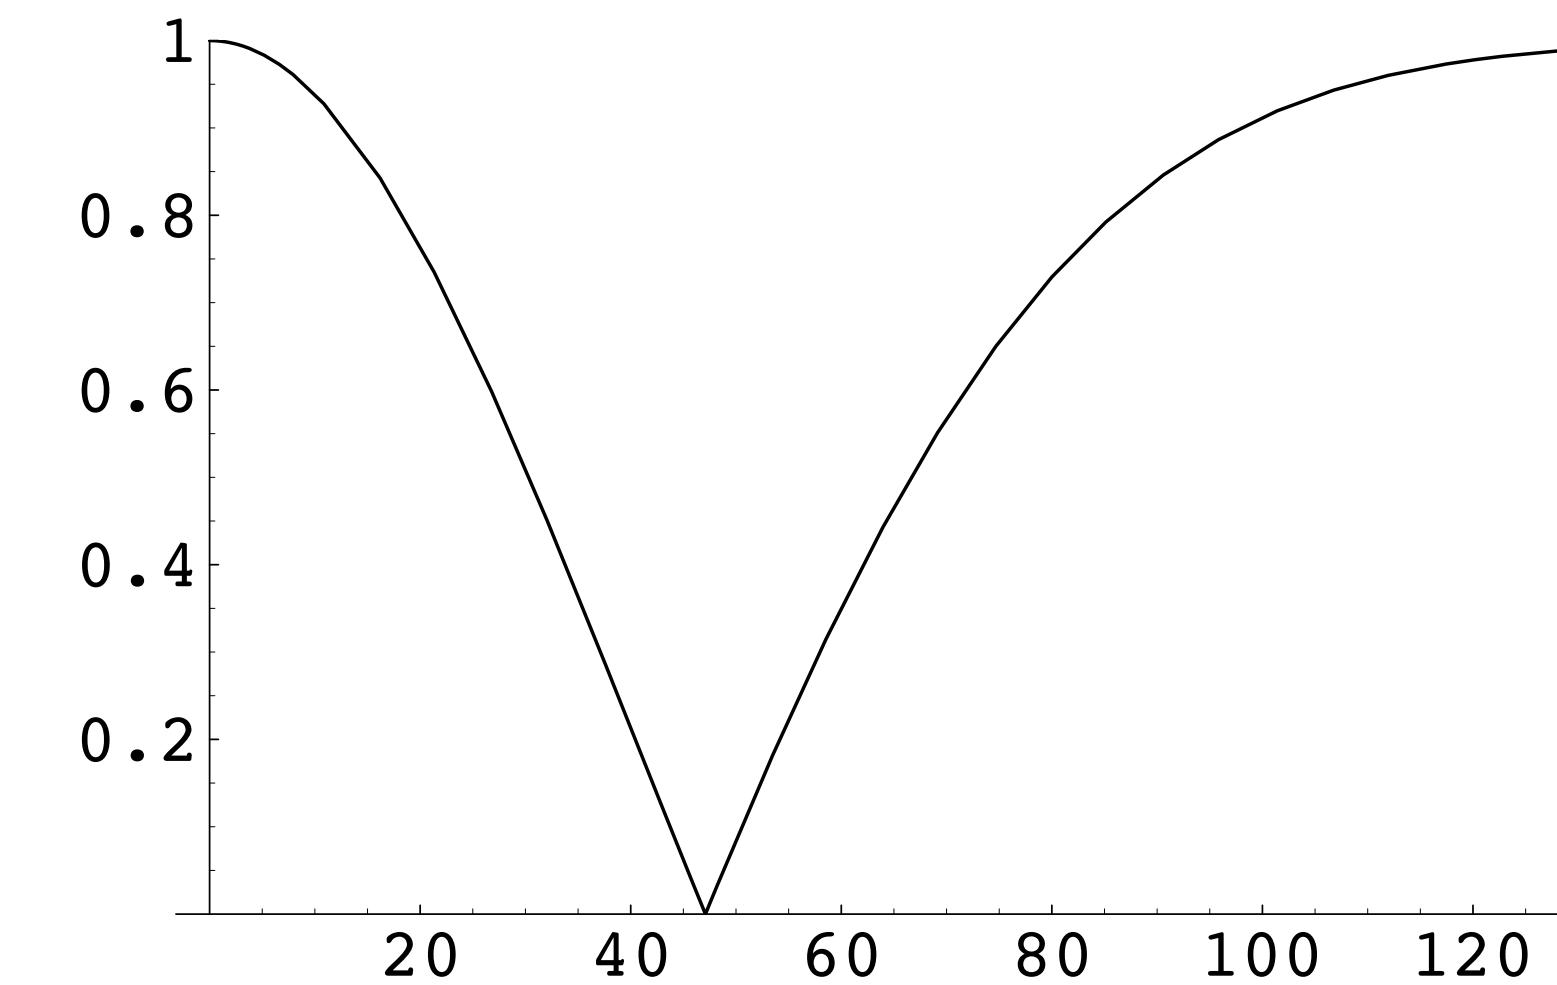
Band-Pass Filters

- Lets frequencies in some target **band pass** through
- Attenuates higher/lower frequencies
- Eg., difference of two low-pass filters



Band-Suppresses Filters

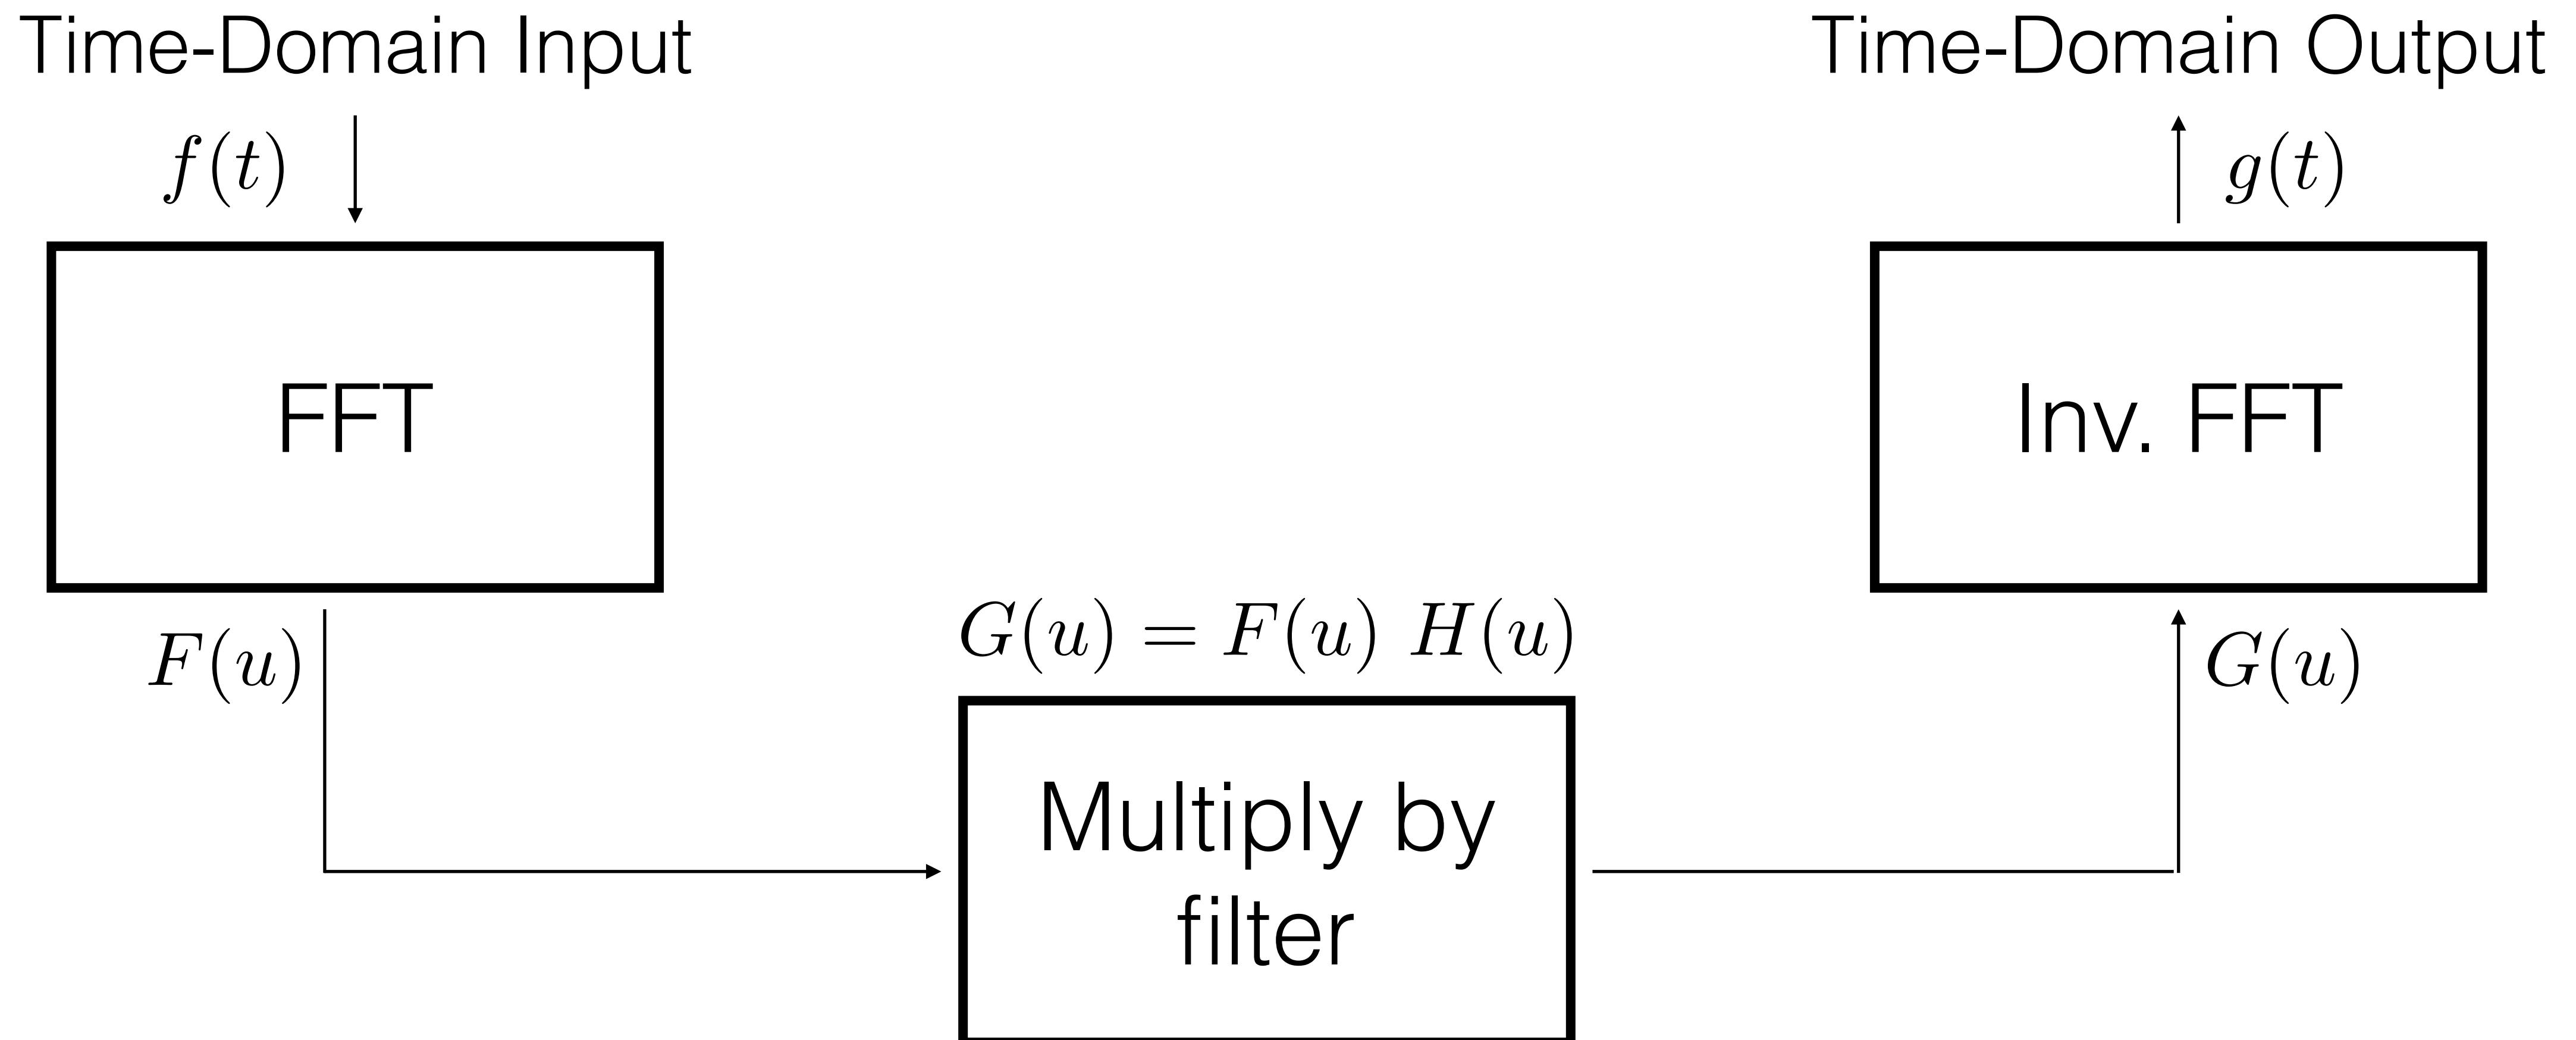
- **Suppresses** frequencies in a target **band**
- Preserves higher/lower frequencies



Phase Filtering

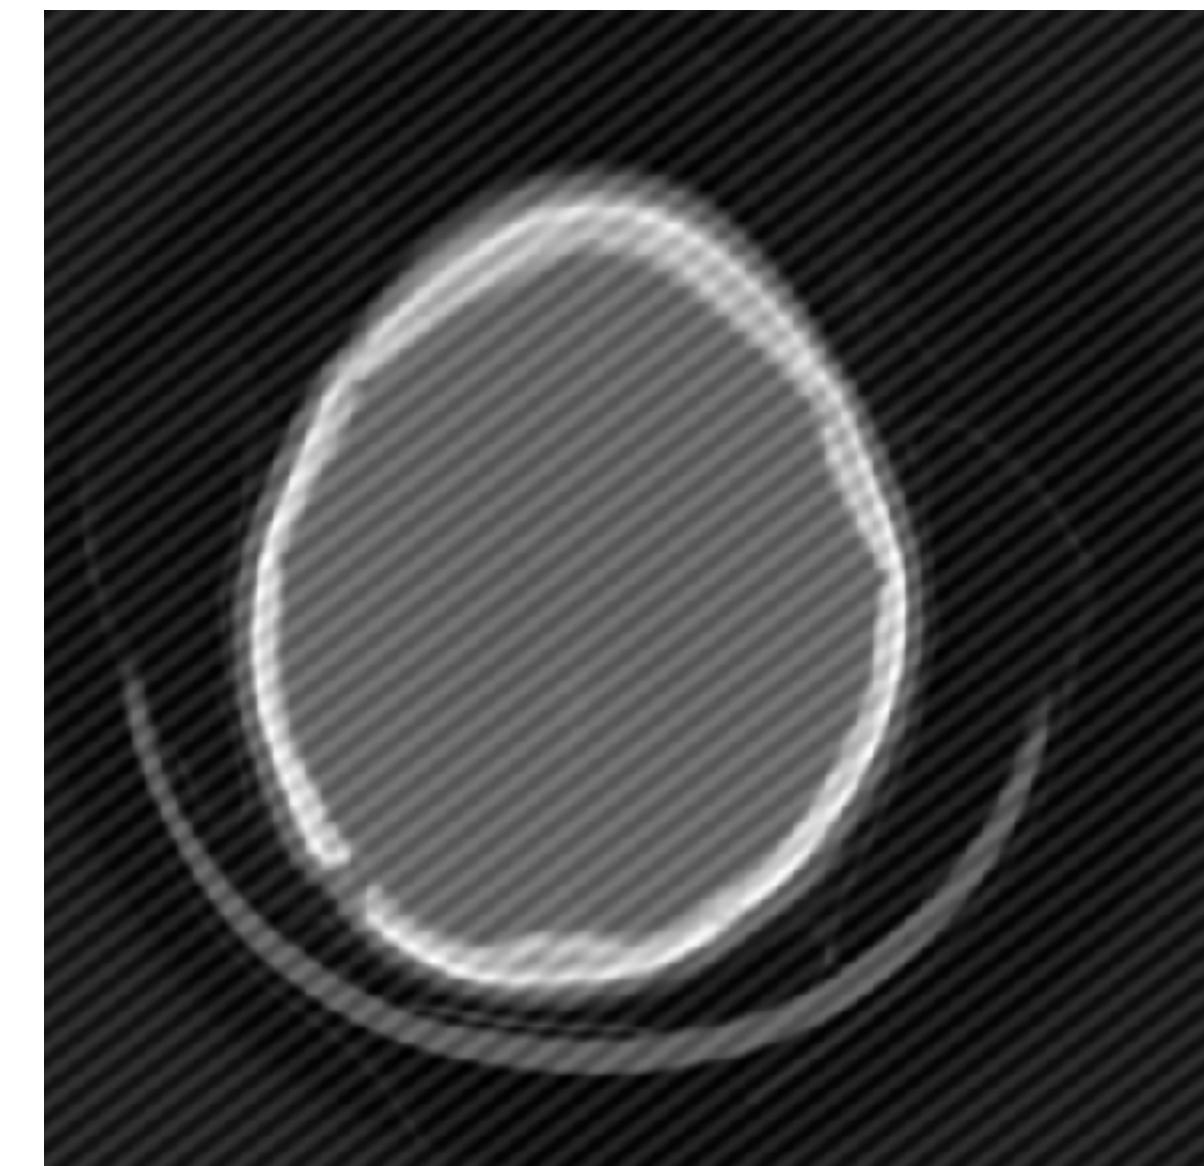
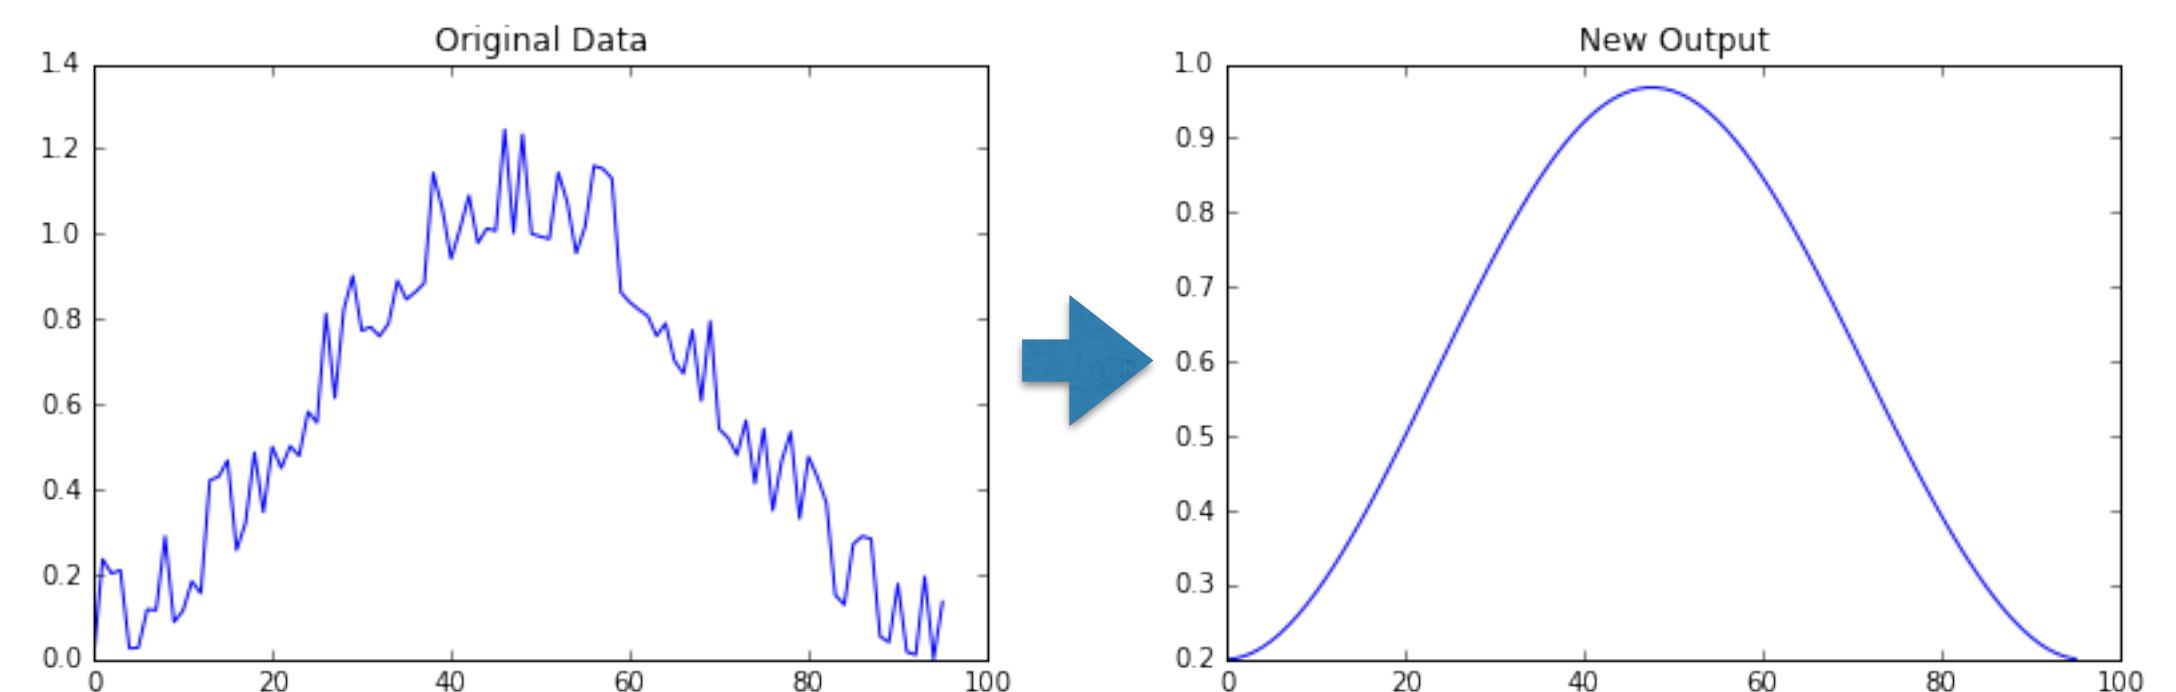
- Time delay
- Example: transmission of electronic signals through wires
 - Different frequencies pass at different speeds through the same media
 - Can become noticeable over long distances
 - Solution: periodically gather, re-sync the phases, retransmit

Implementing Filtering Digitally



Lab 10

- Filtering 1-D signals
 - Synthetic example (noisy function)
 - Audio example
- Filtering 2-D images
 - Smoothing/sharpening
 - Isolating and removing a single-frequency interference pattern



Coming up...

- 2-D FFT and image filtering in the frequency domain
- And then we're done!



2D Fourier Transform

CS 355: Introduction to Graphics and Image Processing

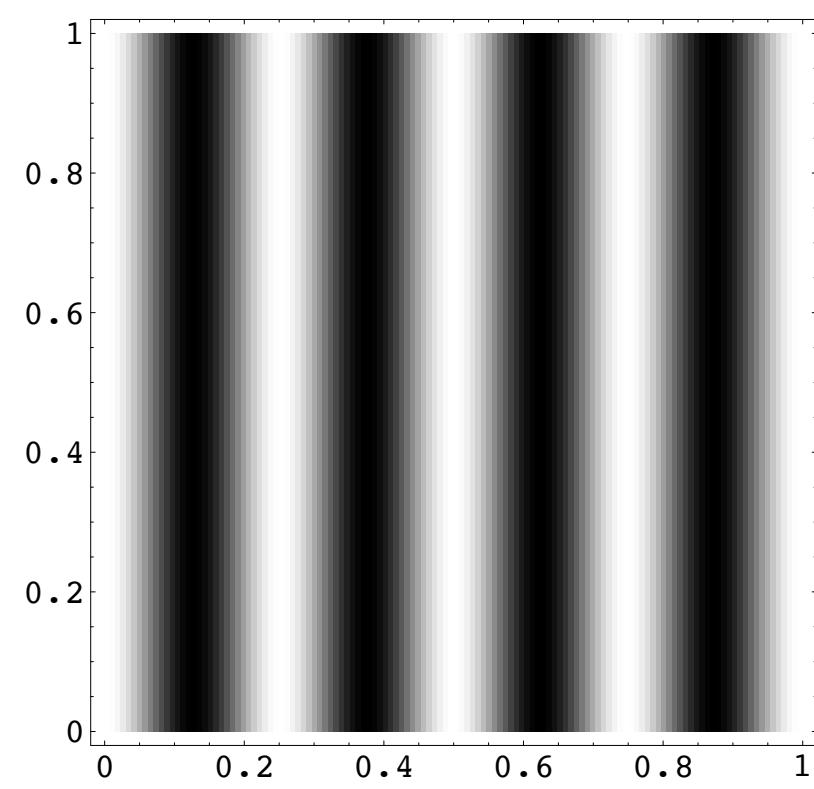
The 2D Discrete Fourier Transform

- Uses basis functions of two variables
 - u = frequency in the x direction
 - v = frequency in the y direction

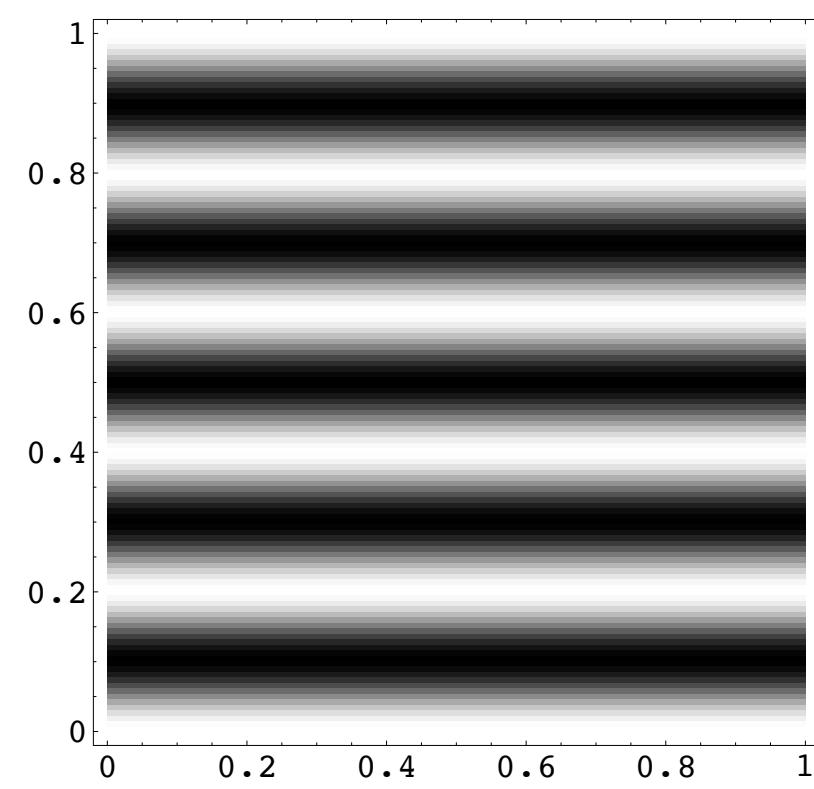
For an $N \times M$ image:

$$\cos(2\pi(ux/M + vy/N))$$

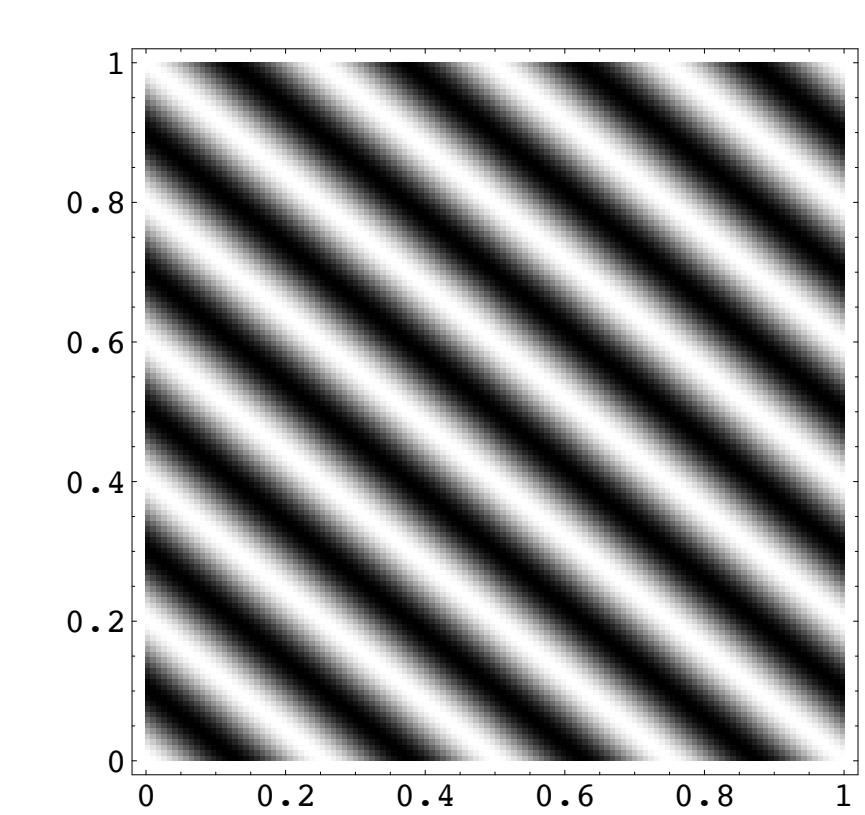
$$\sin(2\pi(ux/M + vy/N))$$



$u = 4, v = 0$



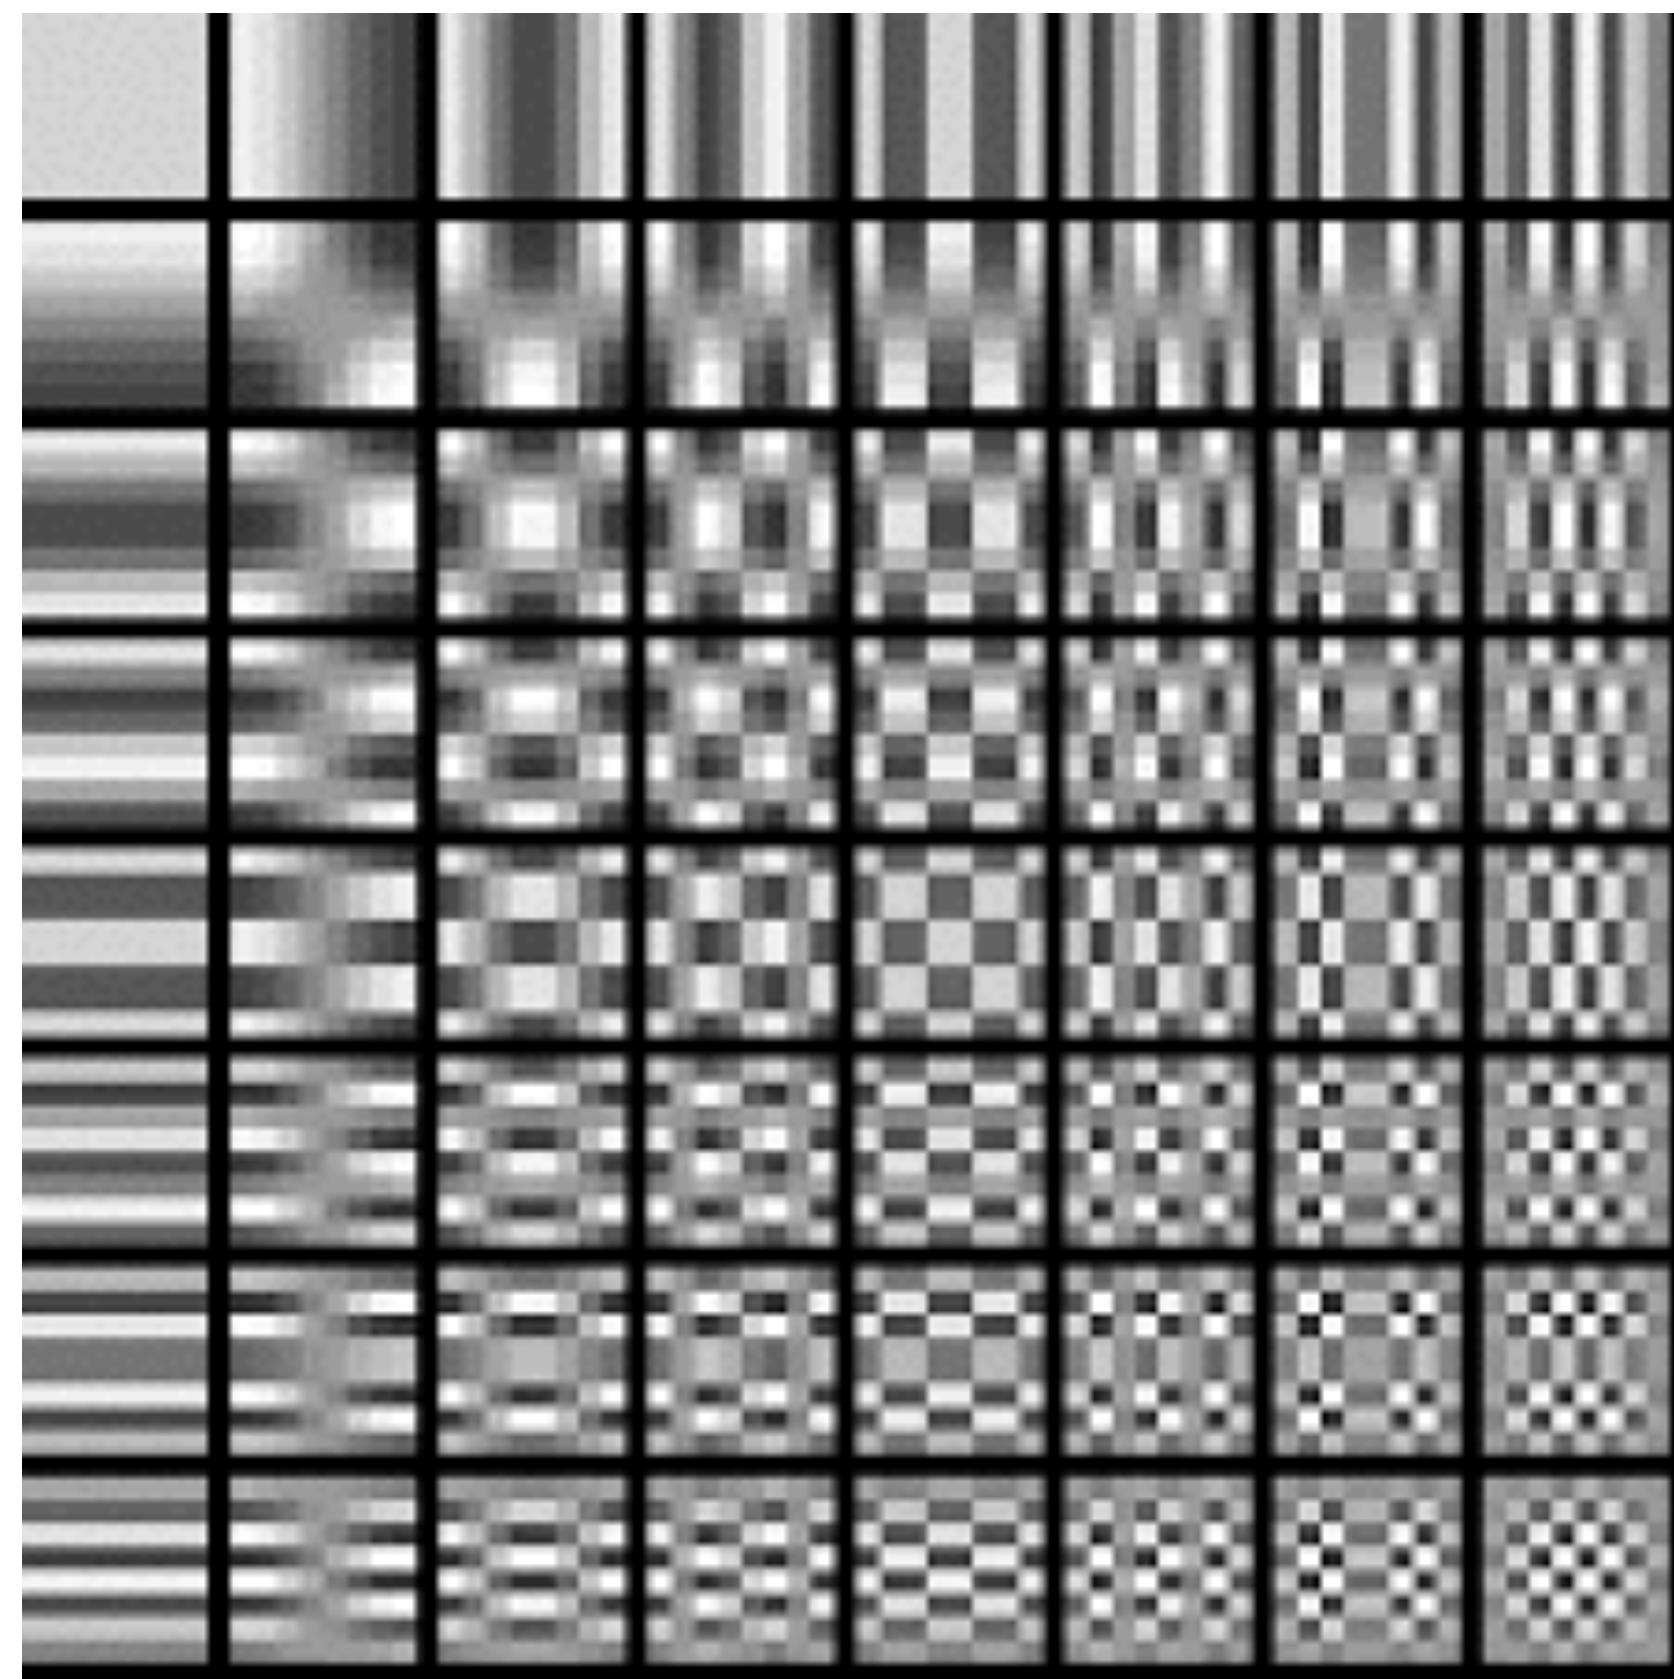
$u = 0, v = 5$



$u = 4, v = 5$

The 2D Discrete Fourier Transform

- Example: 64 basis images can be combined in a weighted sum to form **any** 8×8 image



* Technically, these are for the DCT, a close cousin to the DFT

Spatial Frequencies

- We call these different image components “spatial frequencies”
- Analogous to 1-D frequencies like audio
- Intuition:
 - low frequencies are slow, gradual changes
 - high frequencies are rapid changes (e.g., textures, edges)

The 2D Discrete Fourier Transform

$$\begin{aligned} F[u, v] &= \frac{1}{NM} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f[x, y] \cos(2\pi(ux/M + vy/N)) dx dy \\ &+ i \frac{1}{NM} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f[x, y] \sin(2\pi(ux/M + vy/N)) dx dy \end{aligned}$$

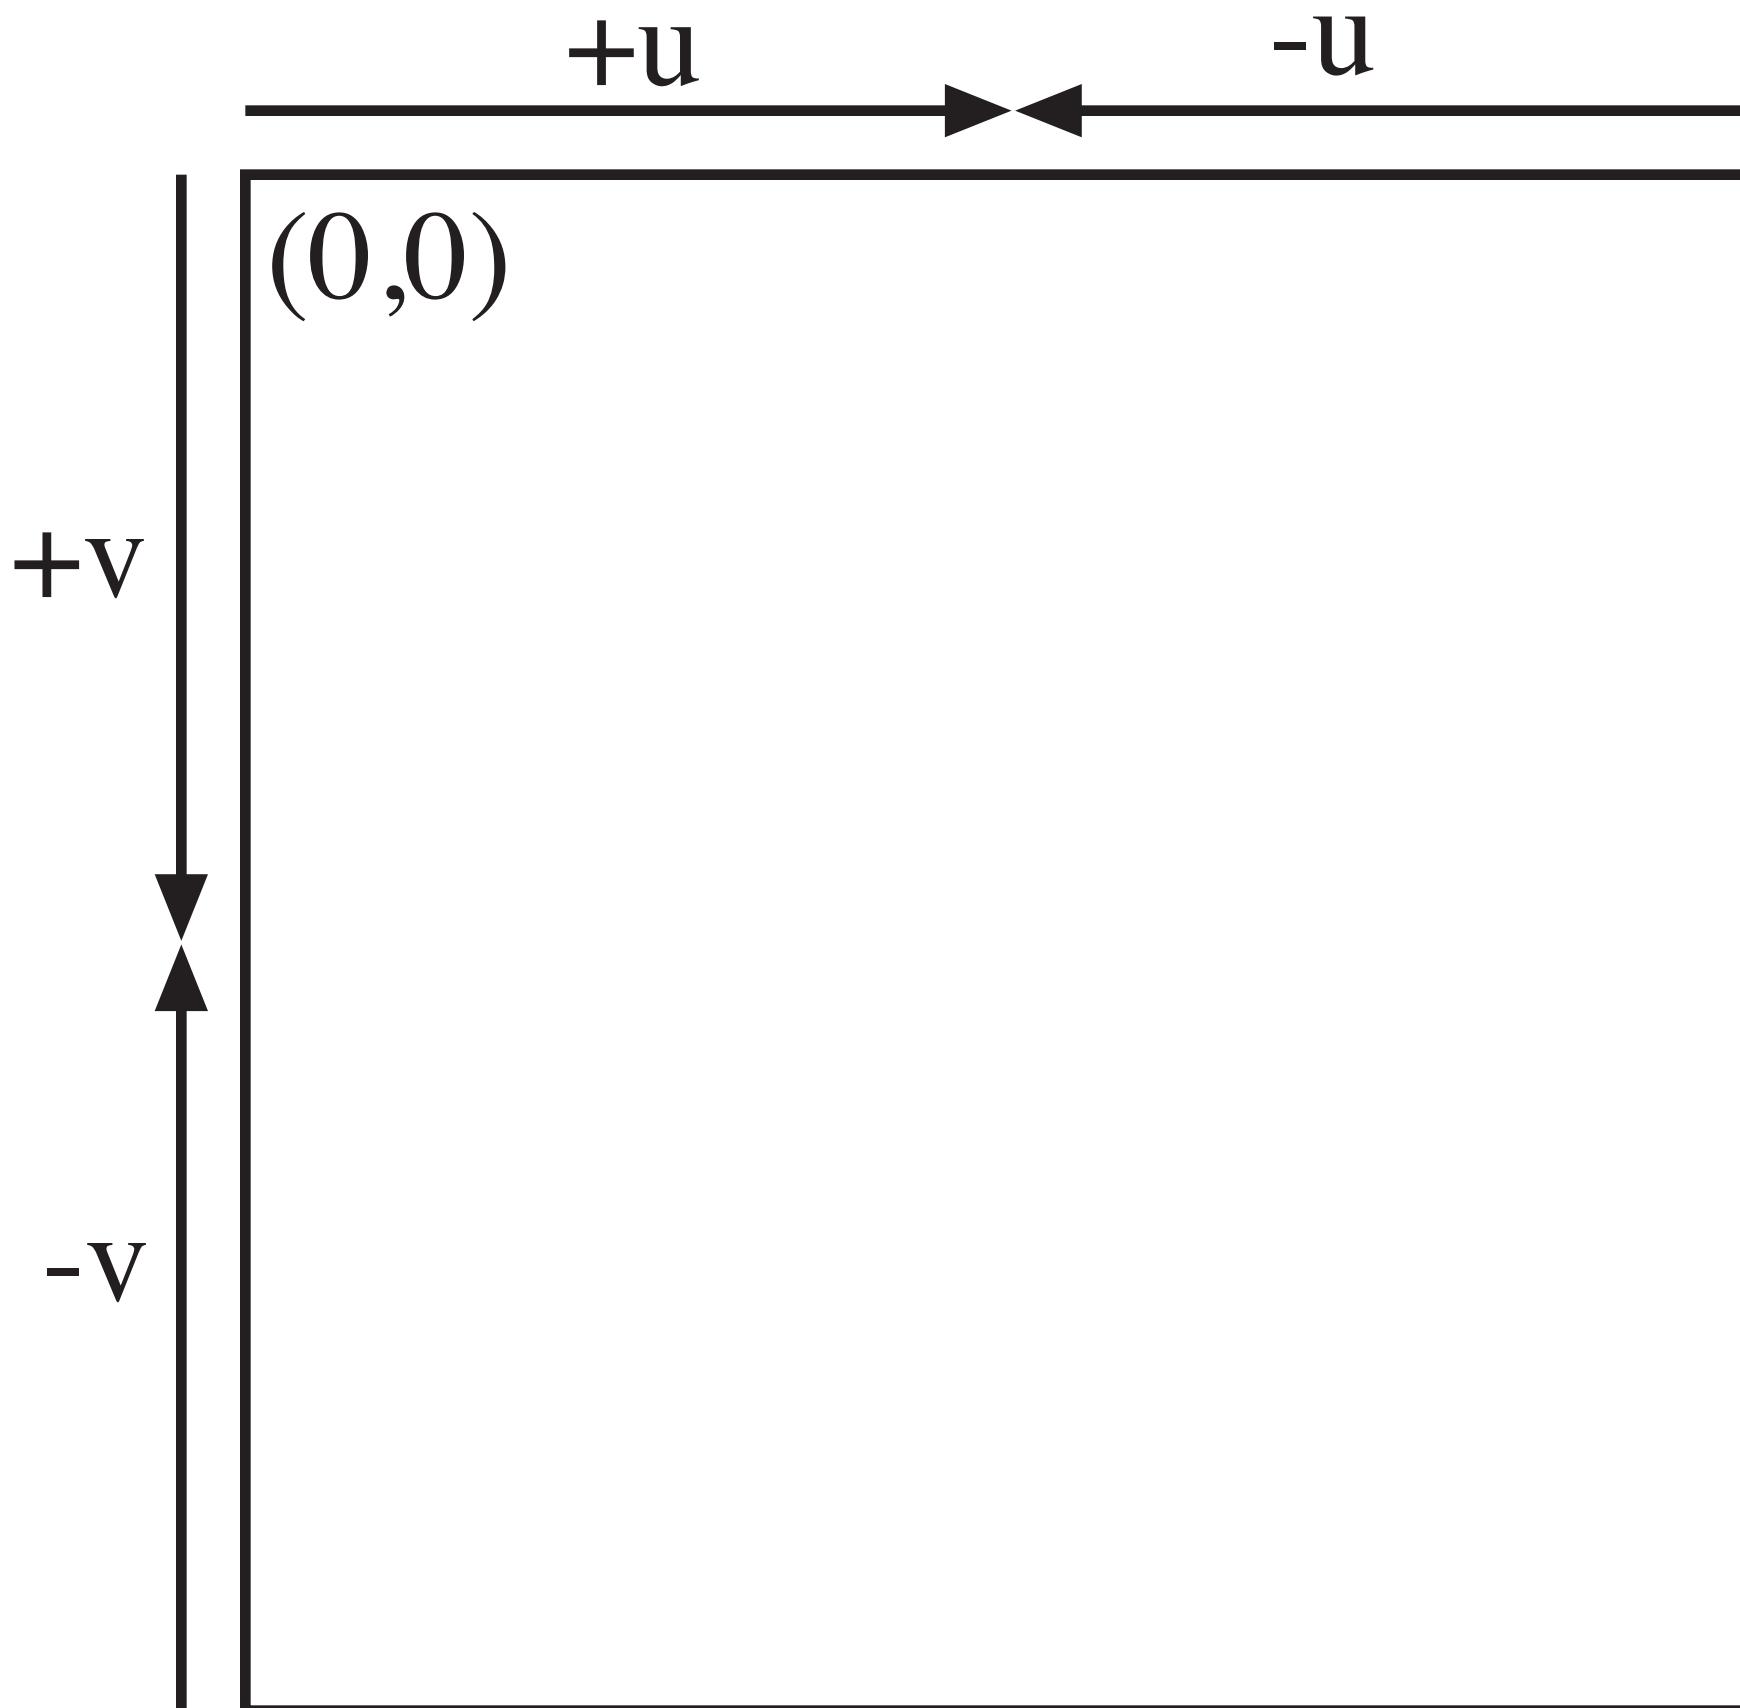
What's does the code look like?

What's the computational complexity?

The 2D Fast Fourier Transform

	DFT	FFT
1-D	$O(N^2)$	$O(N \log N)$
2-D	$O(N^4)$	$O(N^2 \log N)$

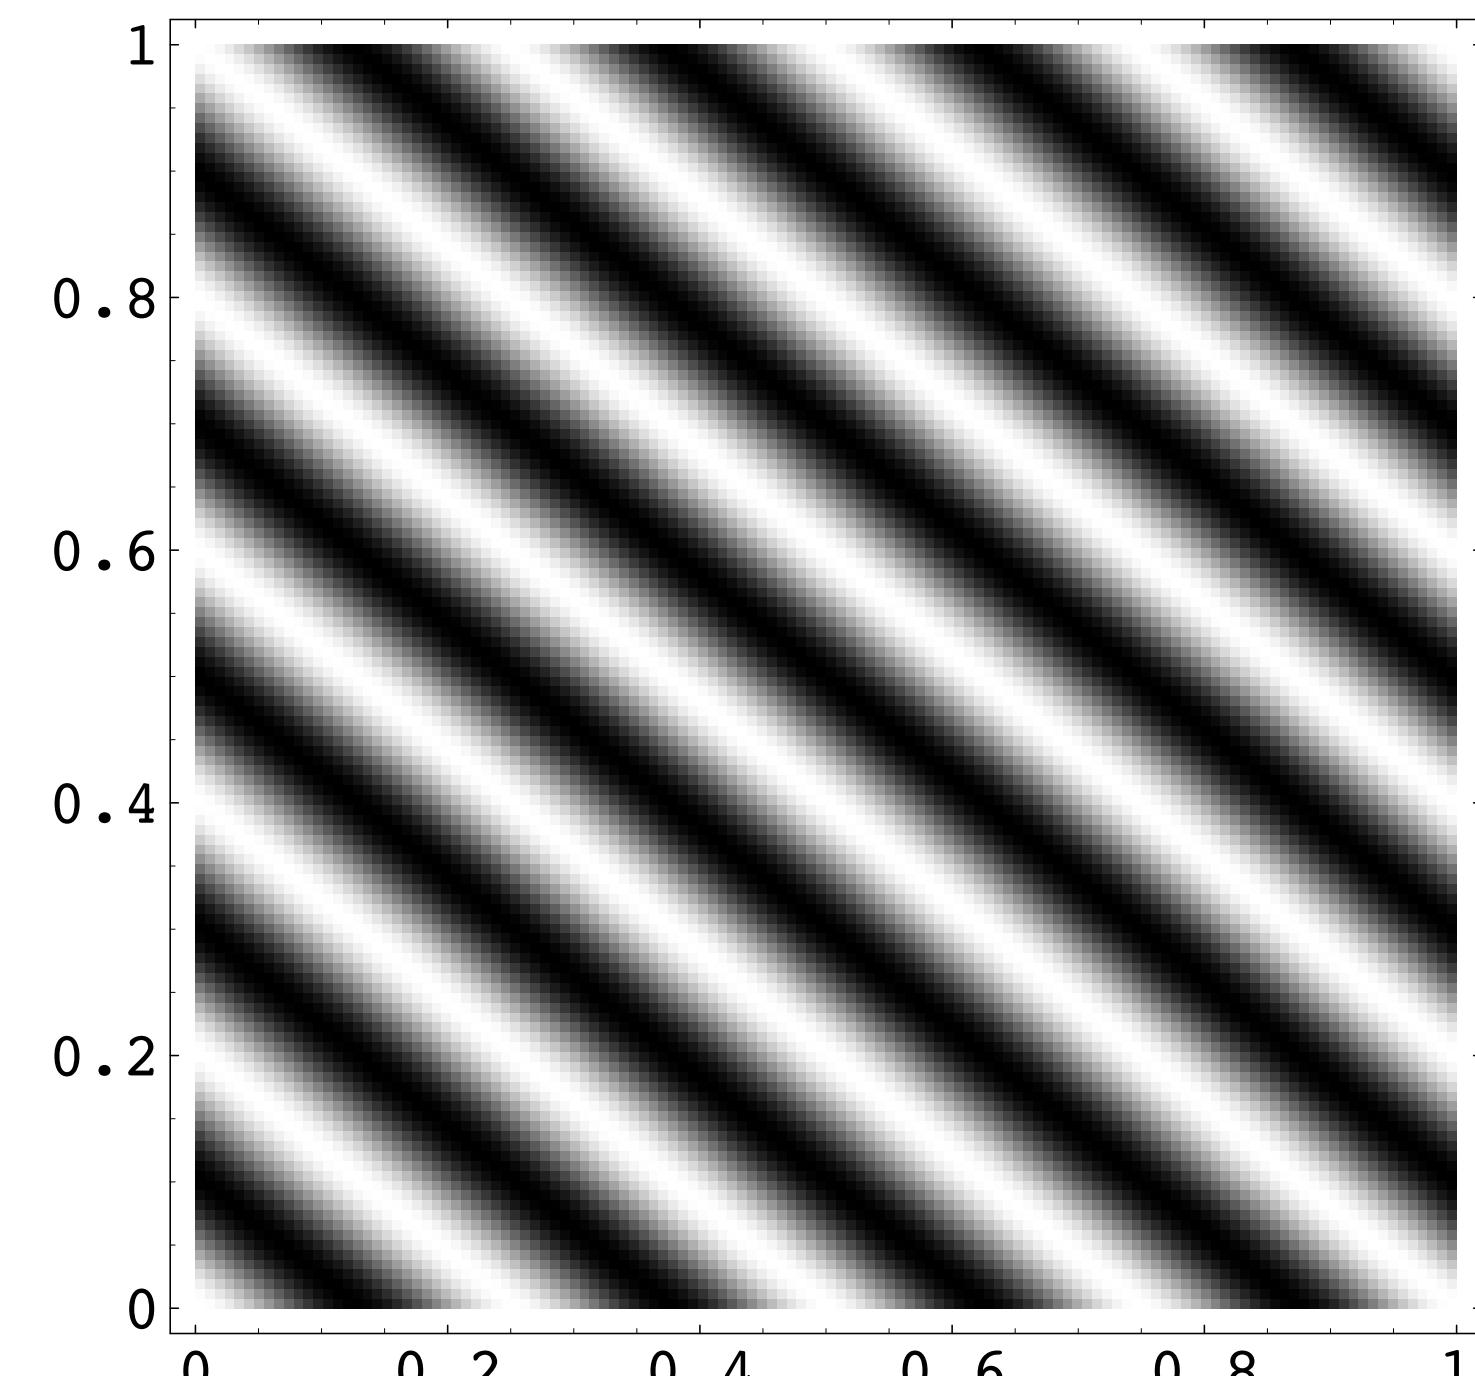
2D DFT Storage



(like 1-D, it's often useful to shift it to better visualize)

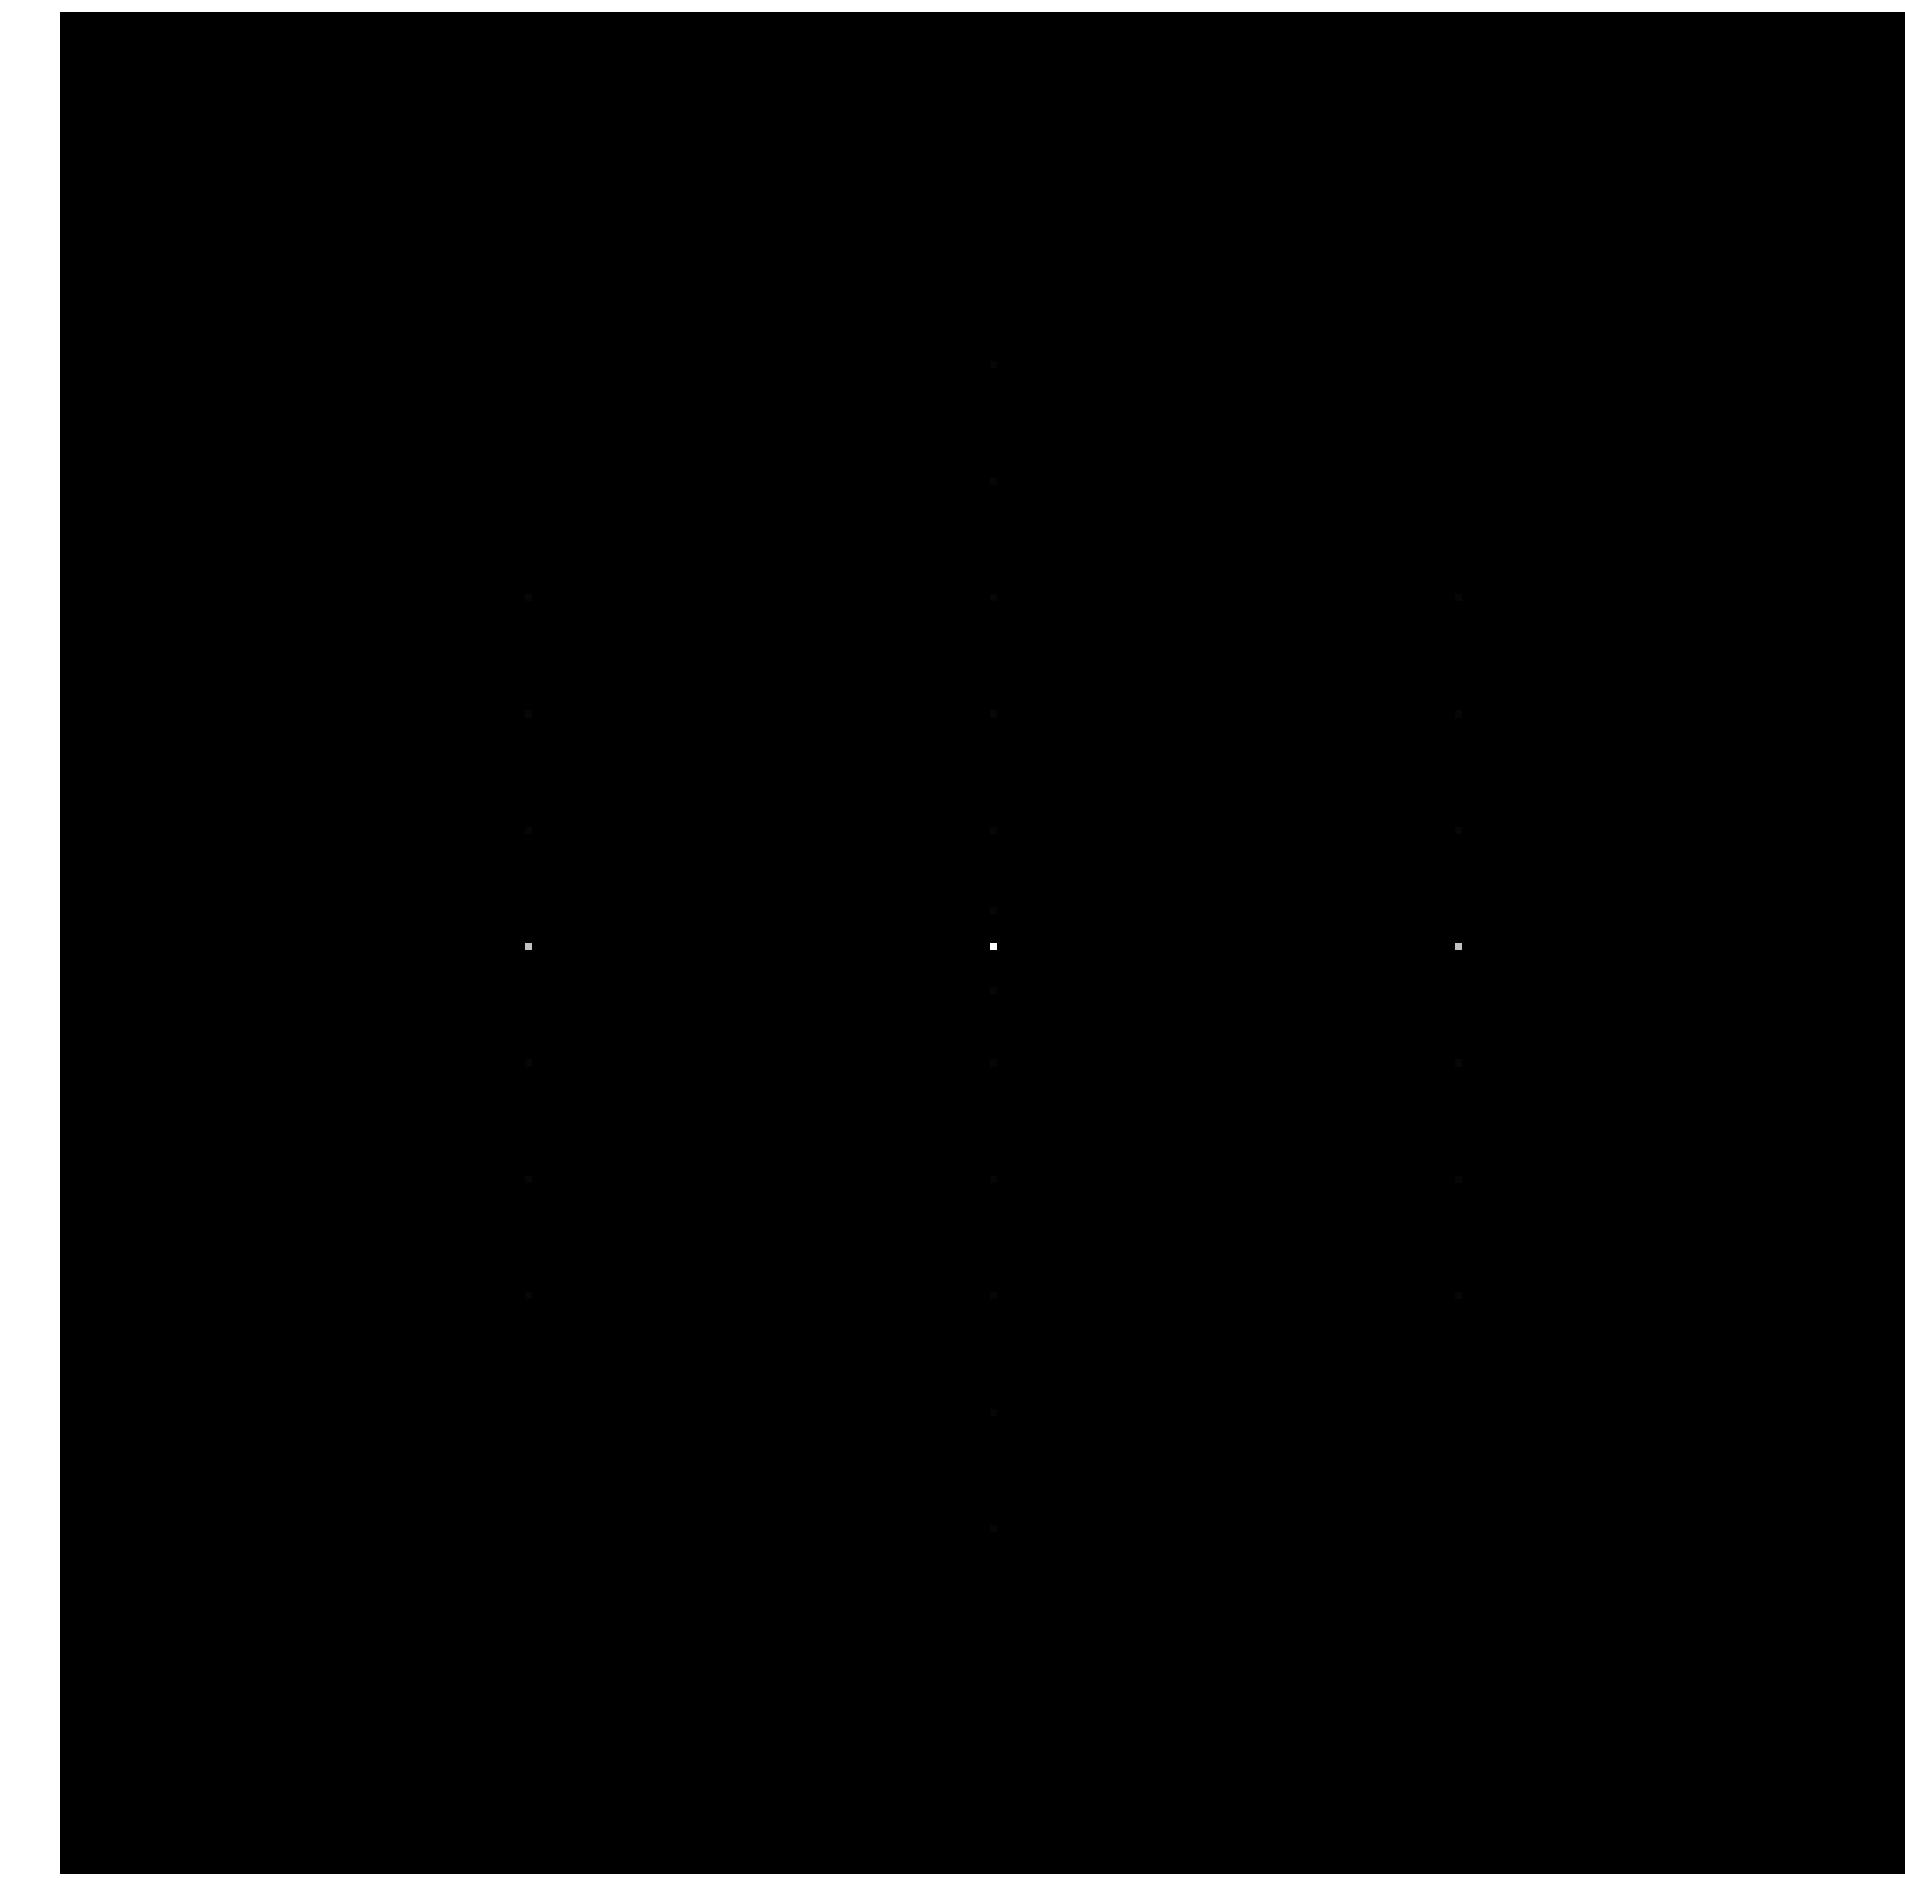
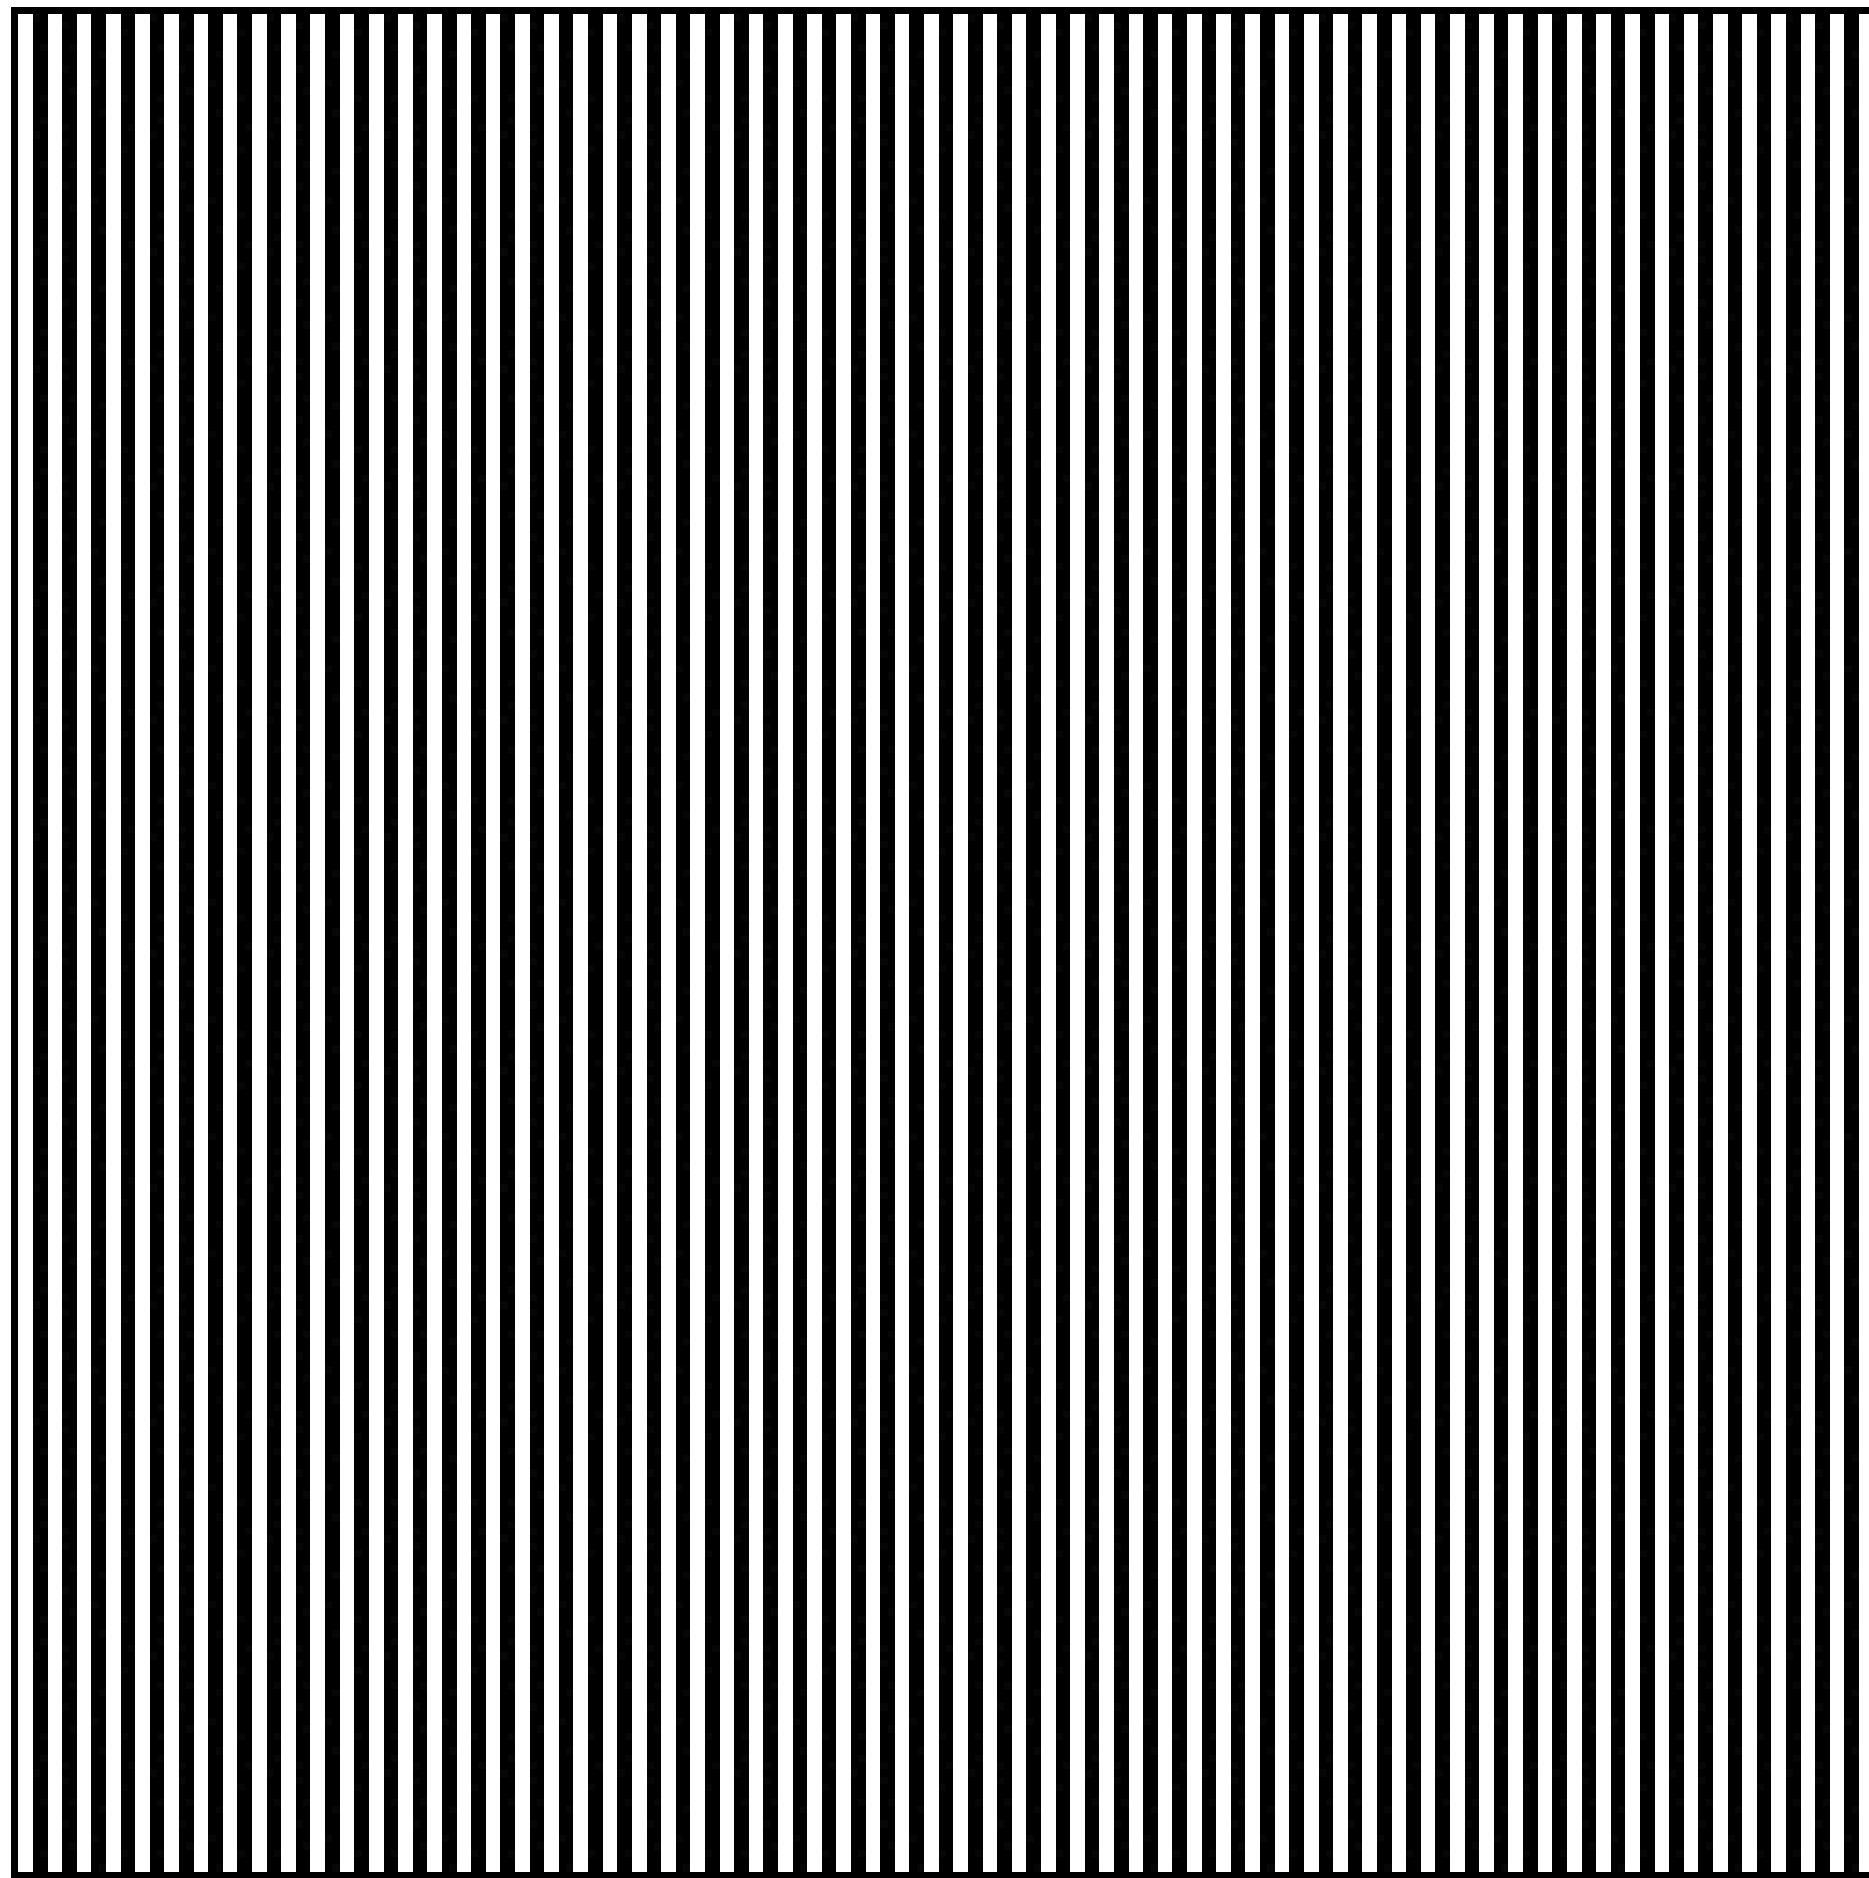
Interpreting the 2D DFT

- Frequencies in x and y directions
 - u = frequency in the x direction
 - v = frequency in the y direction
- A sinusoidal pattern in a single direction
- Useful property: rotating the image rotates its Fourier Transform

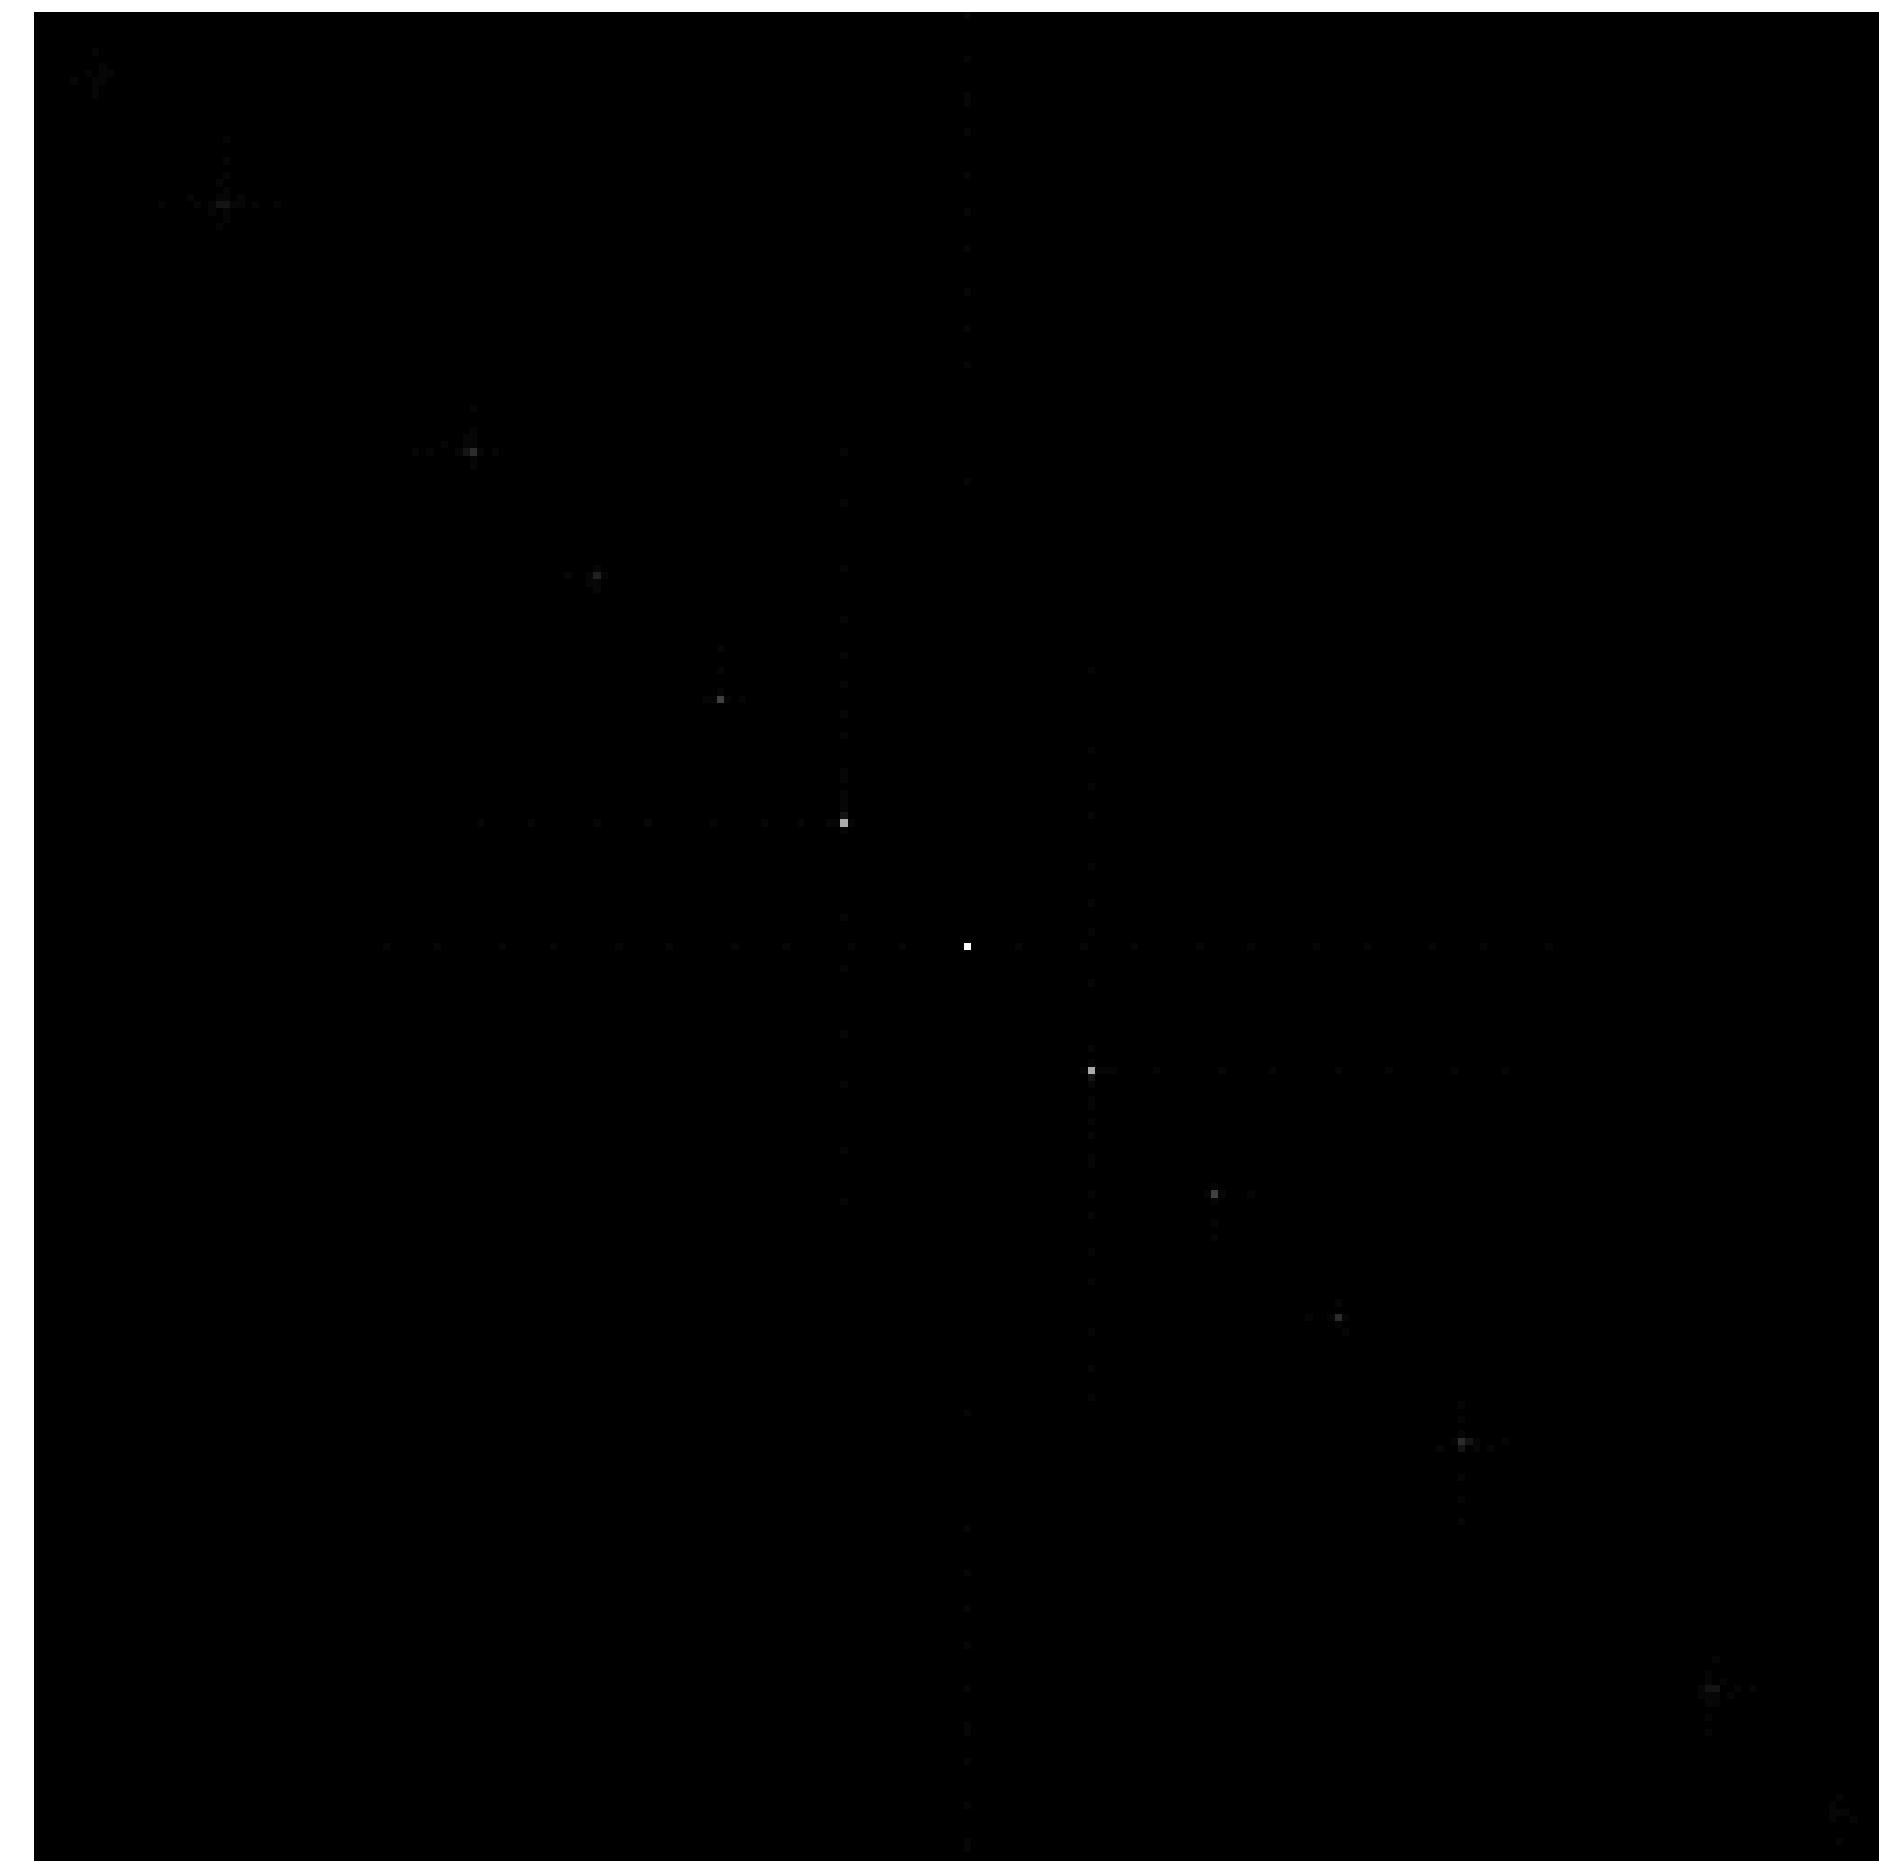
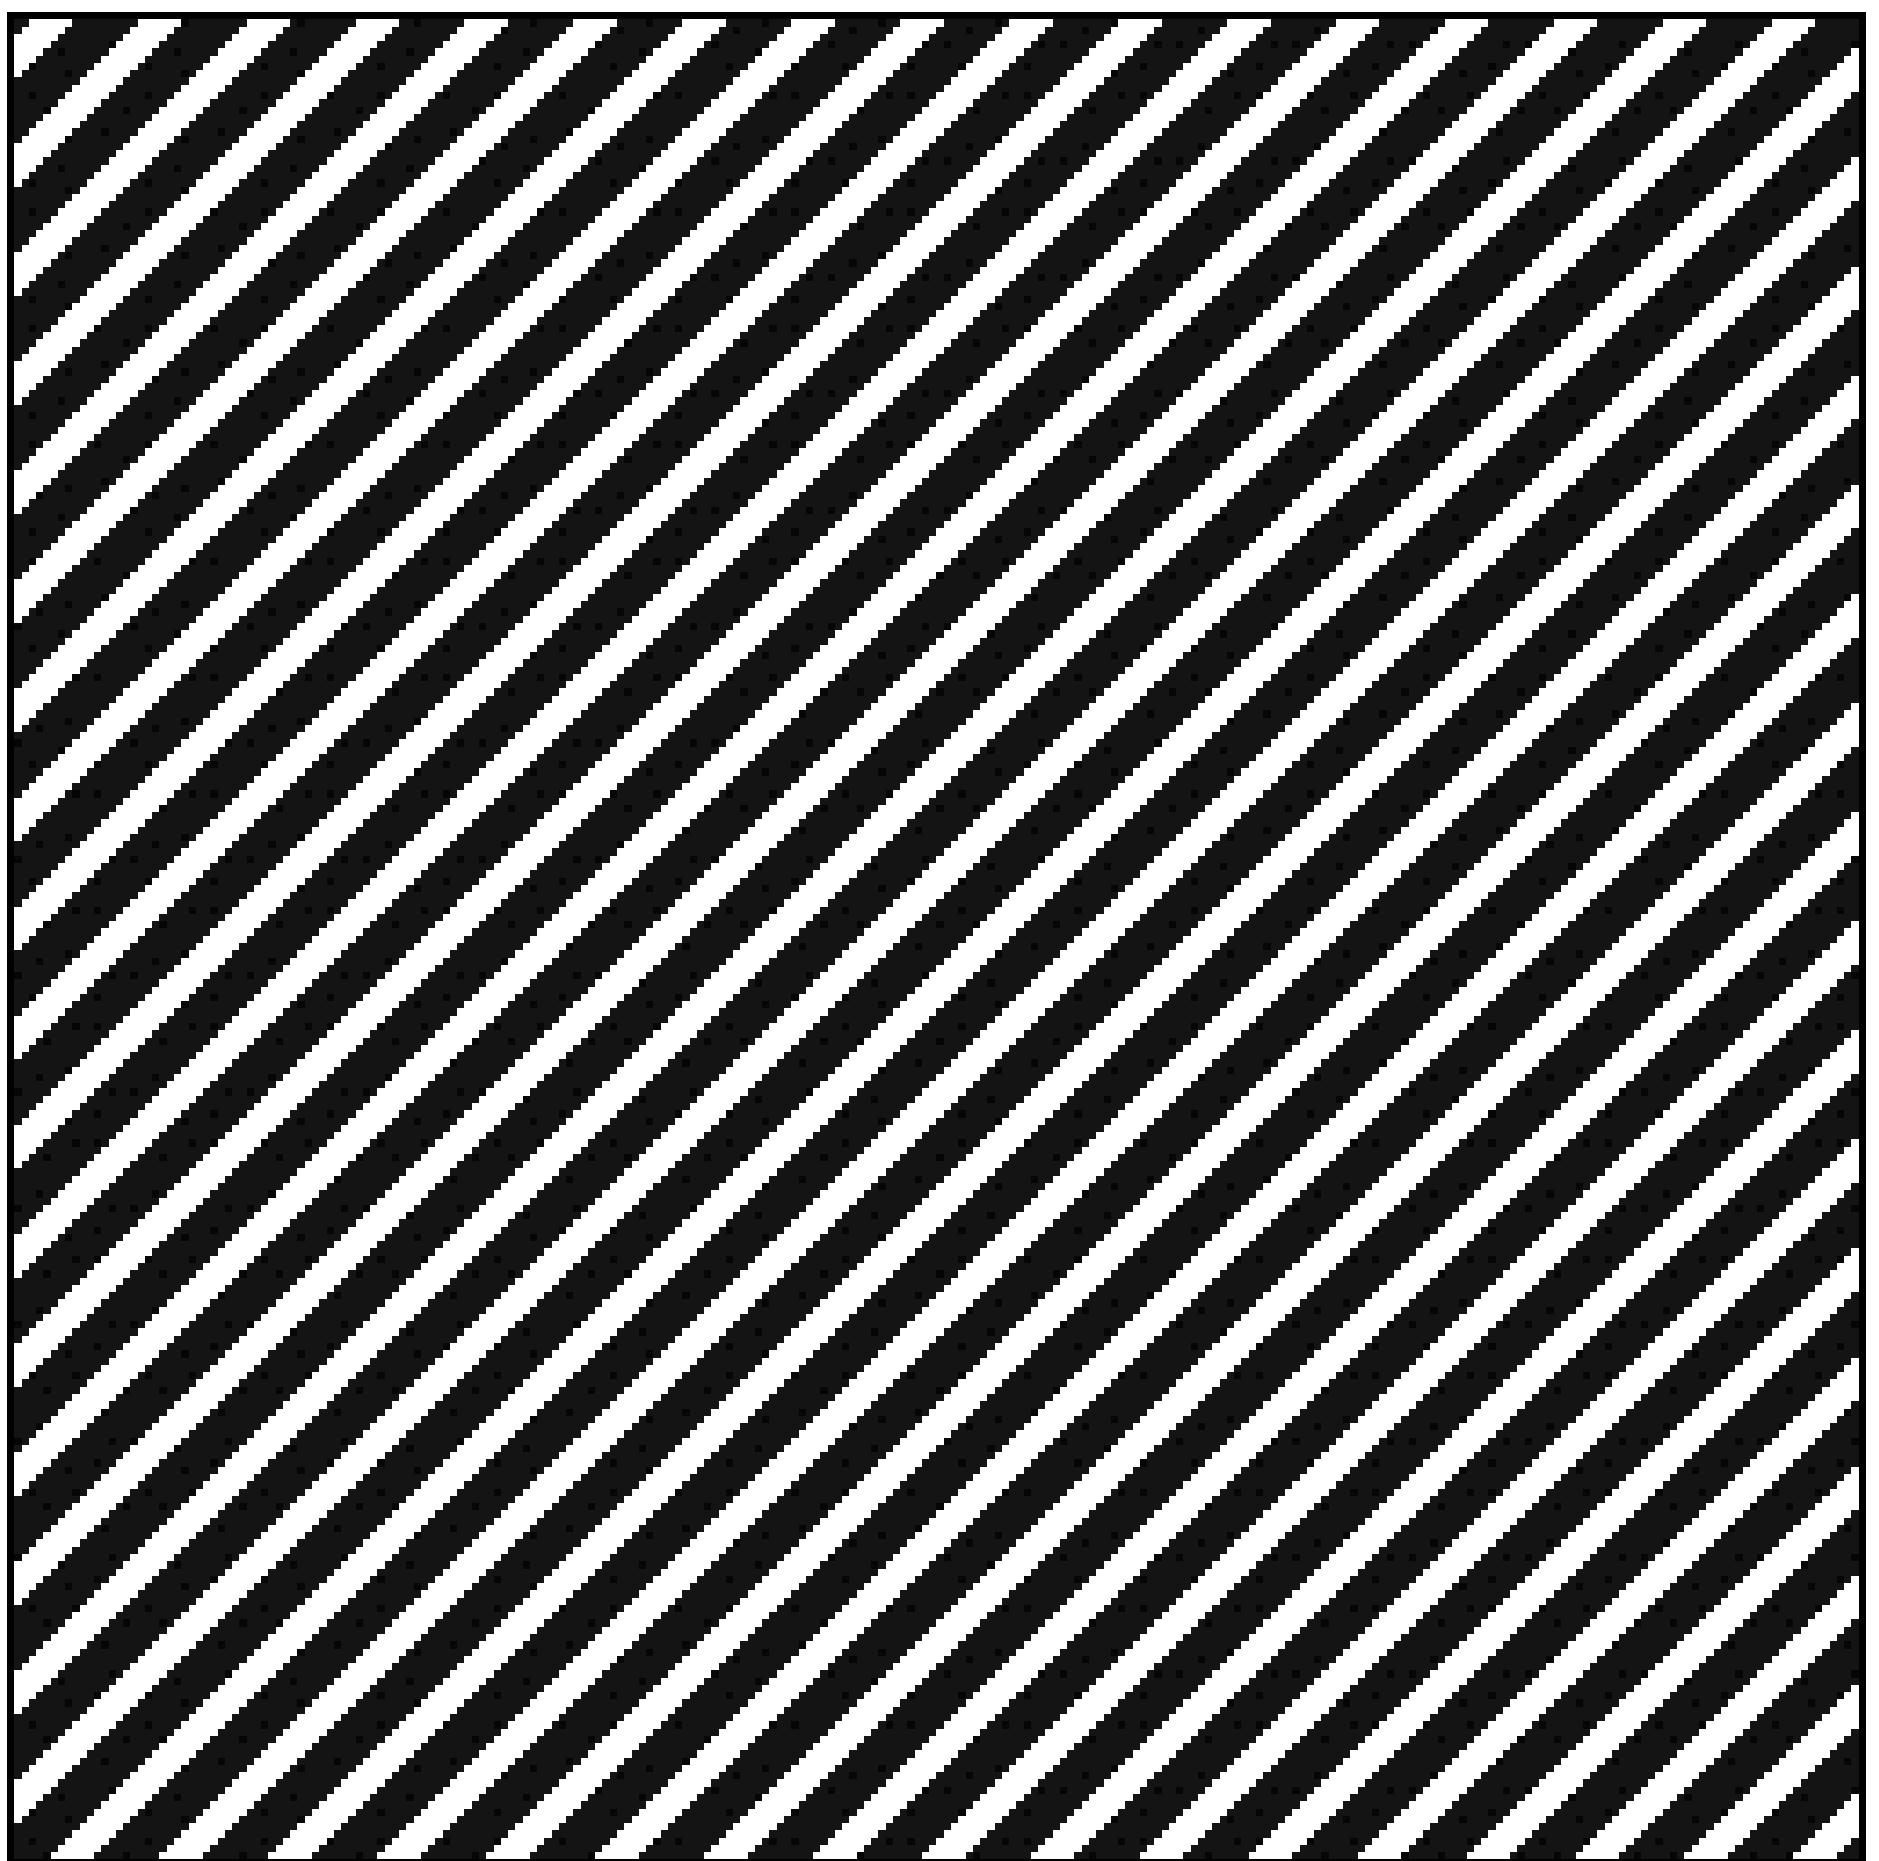


$u = 4, v = 5$

Examples



Examples



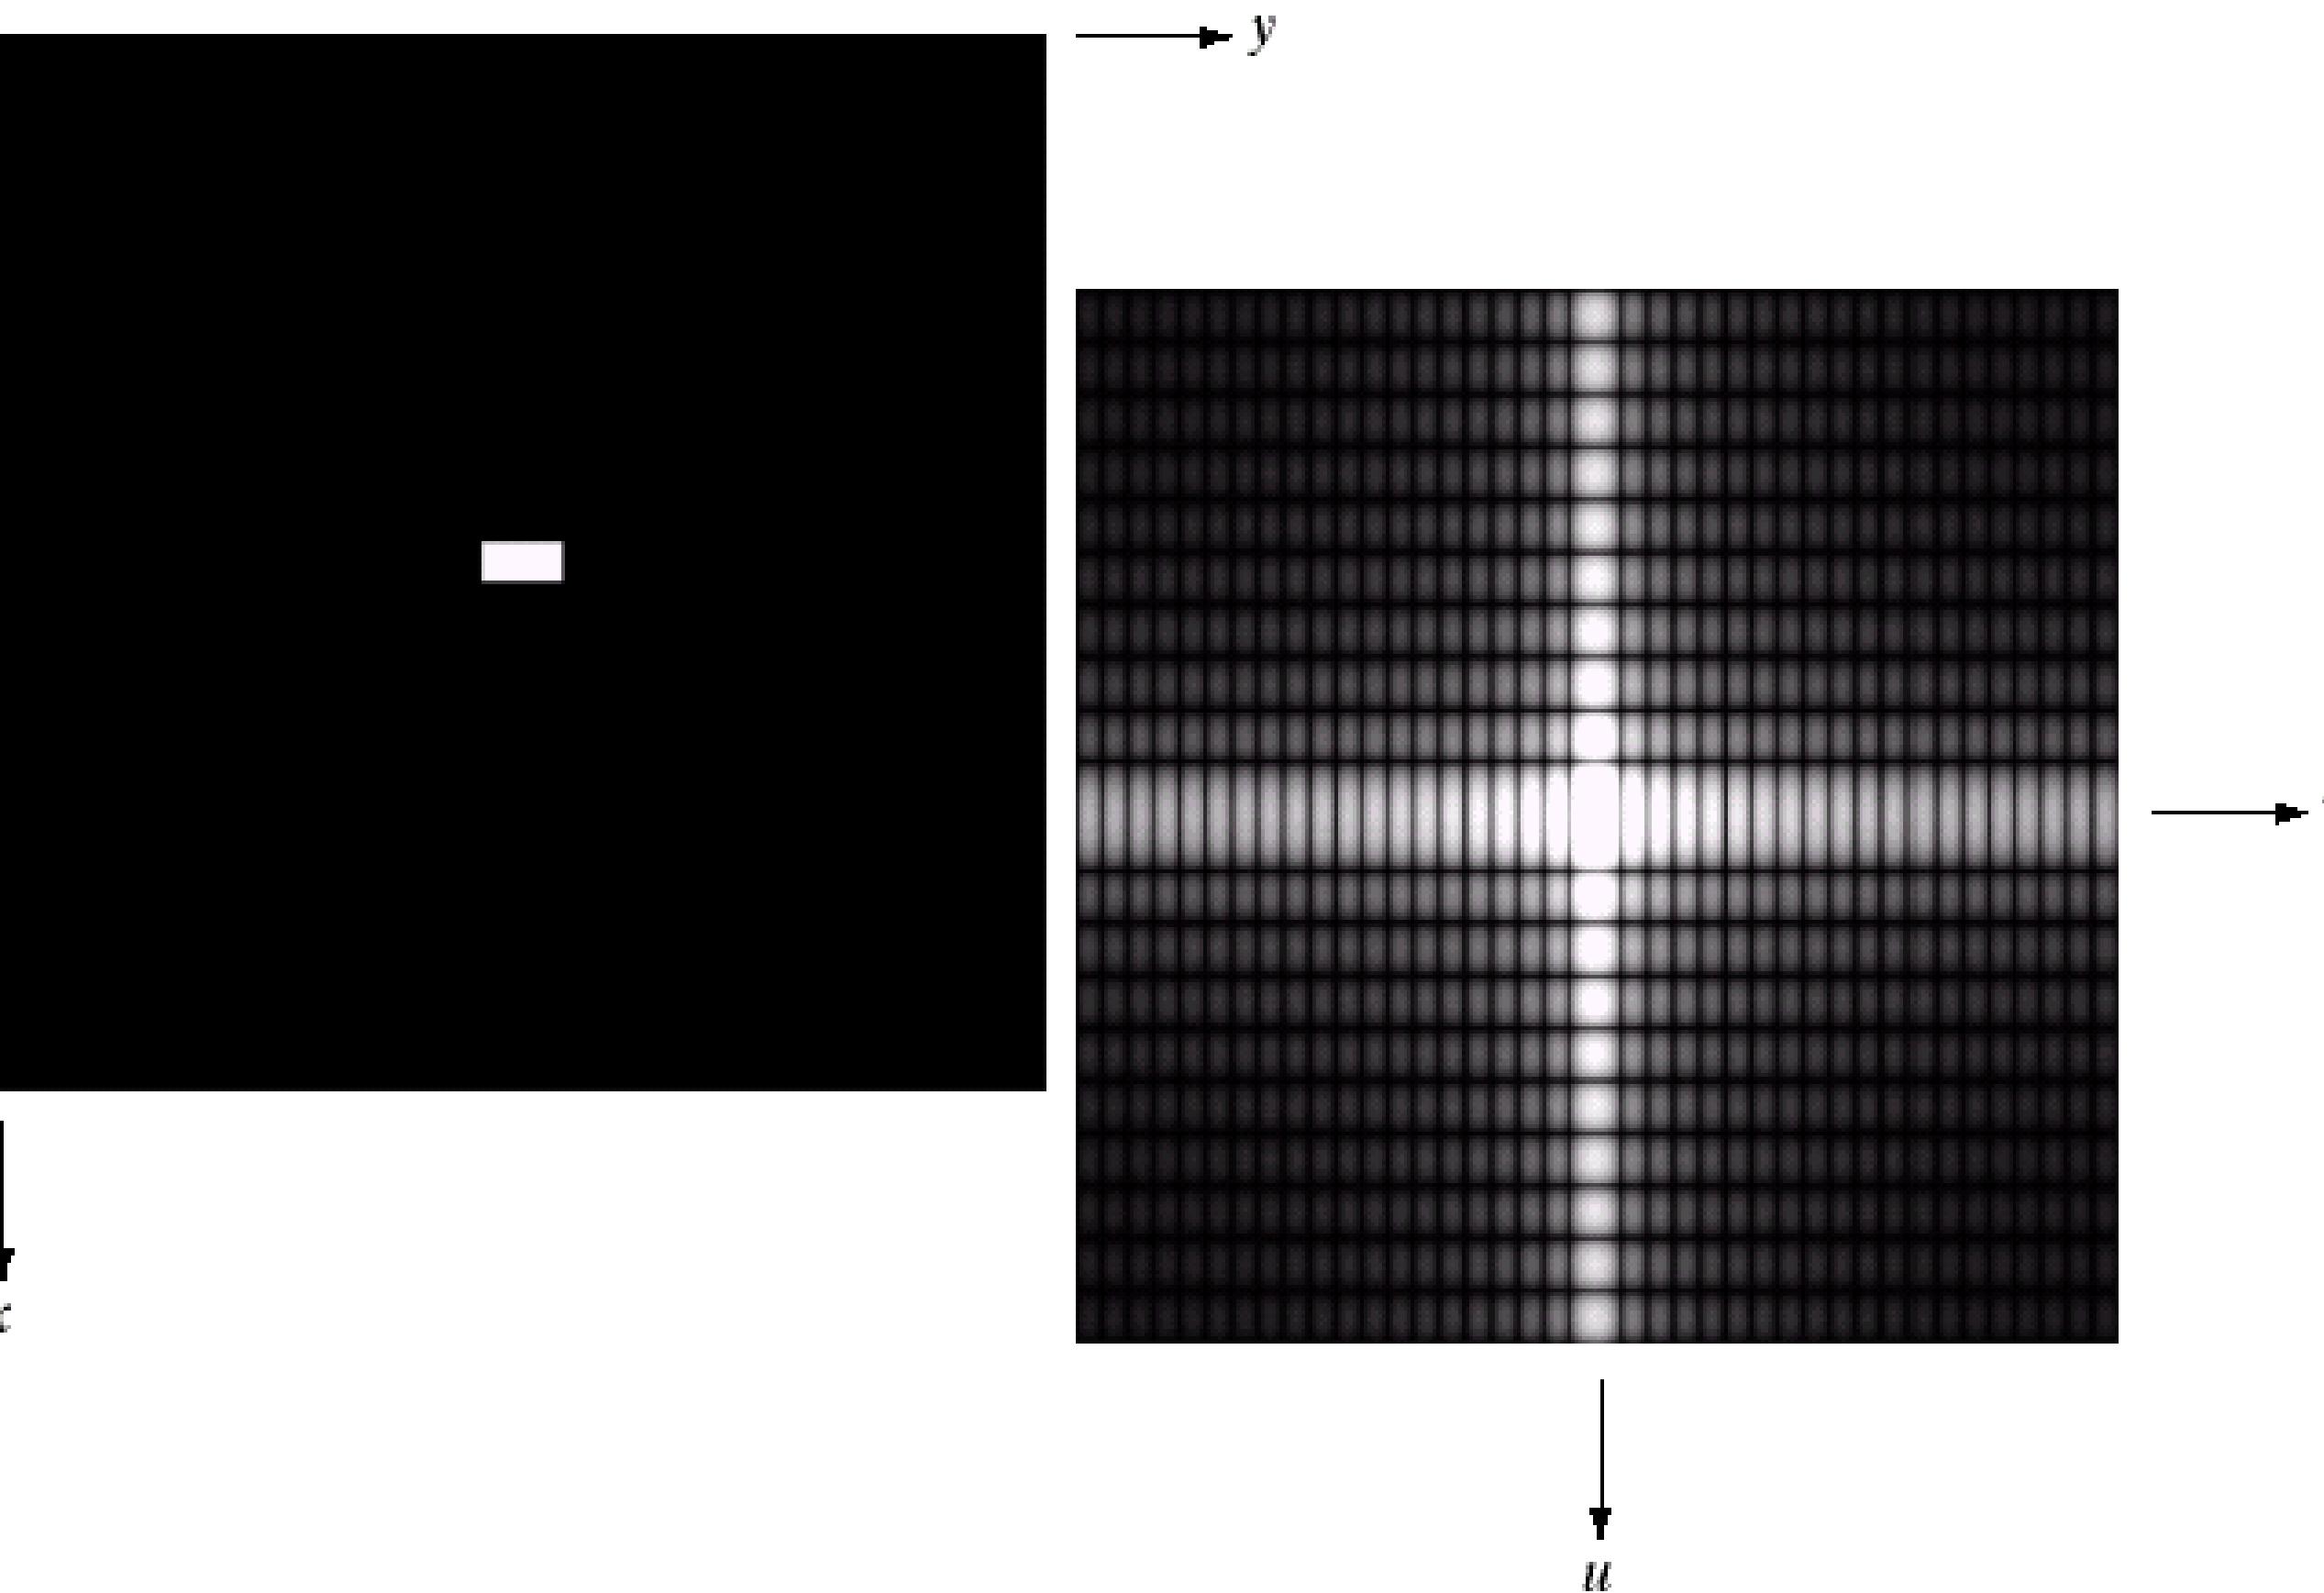
Examples

a b

FIGURE 4.3

(a) Image of a 20×40 white rectangle on a black background of size 512×512 pixels.

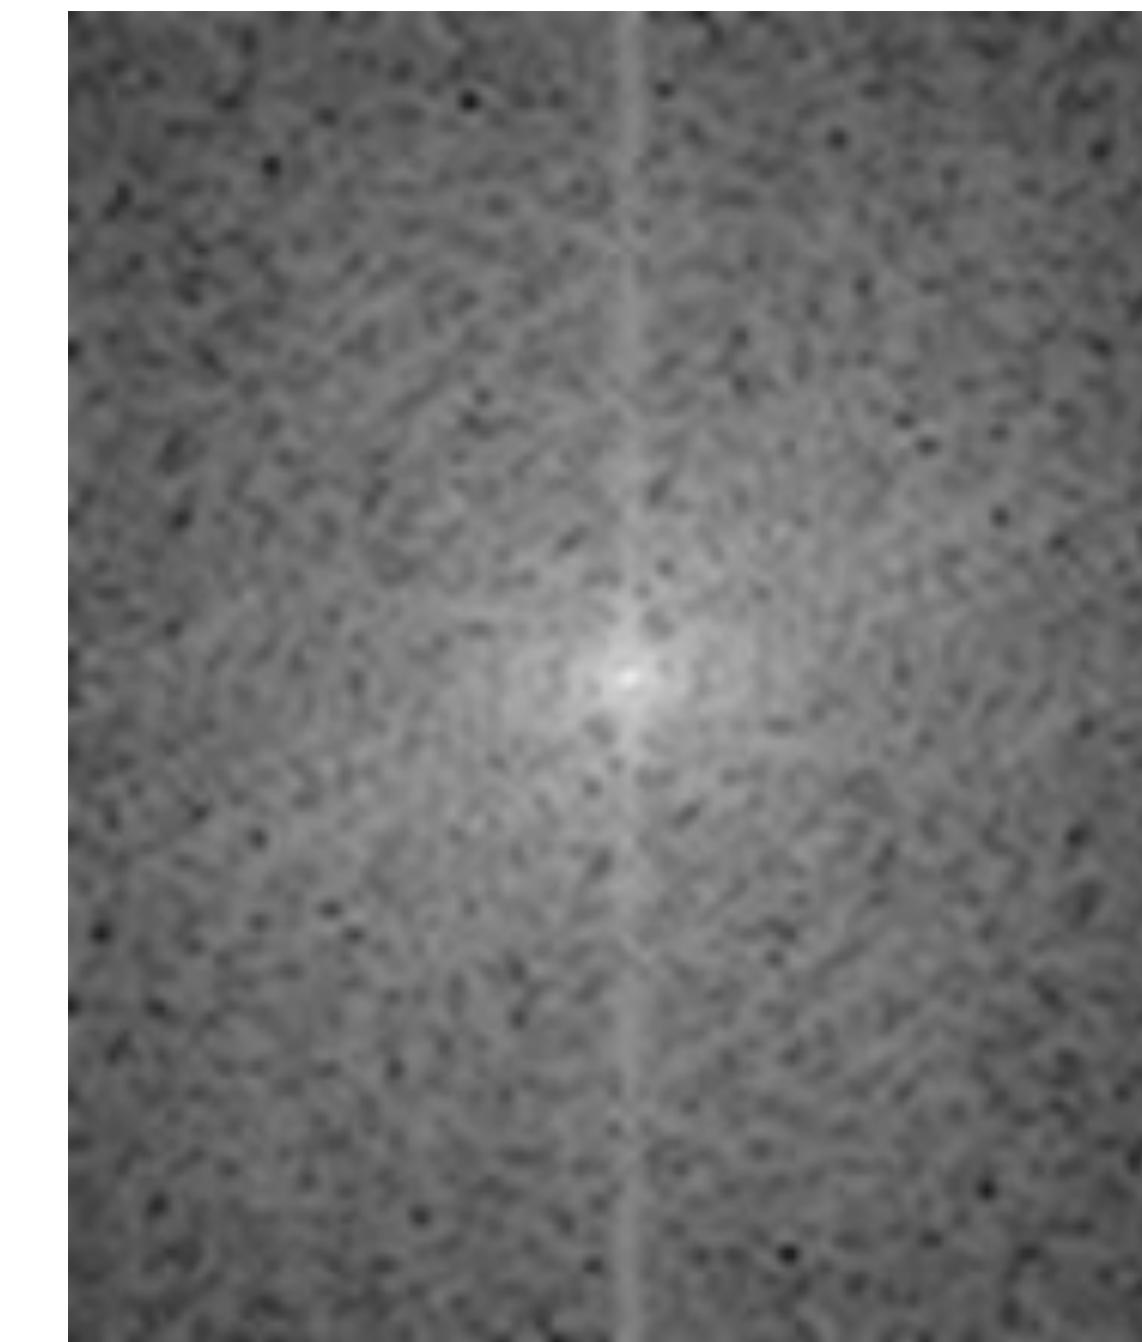
(b) Centered Fourier spectrum shown after application of the log transformation given in Eq. (3.2-2). Compare with Fig. 4.2.



Another Example



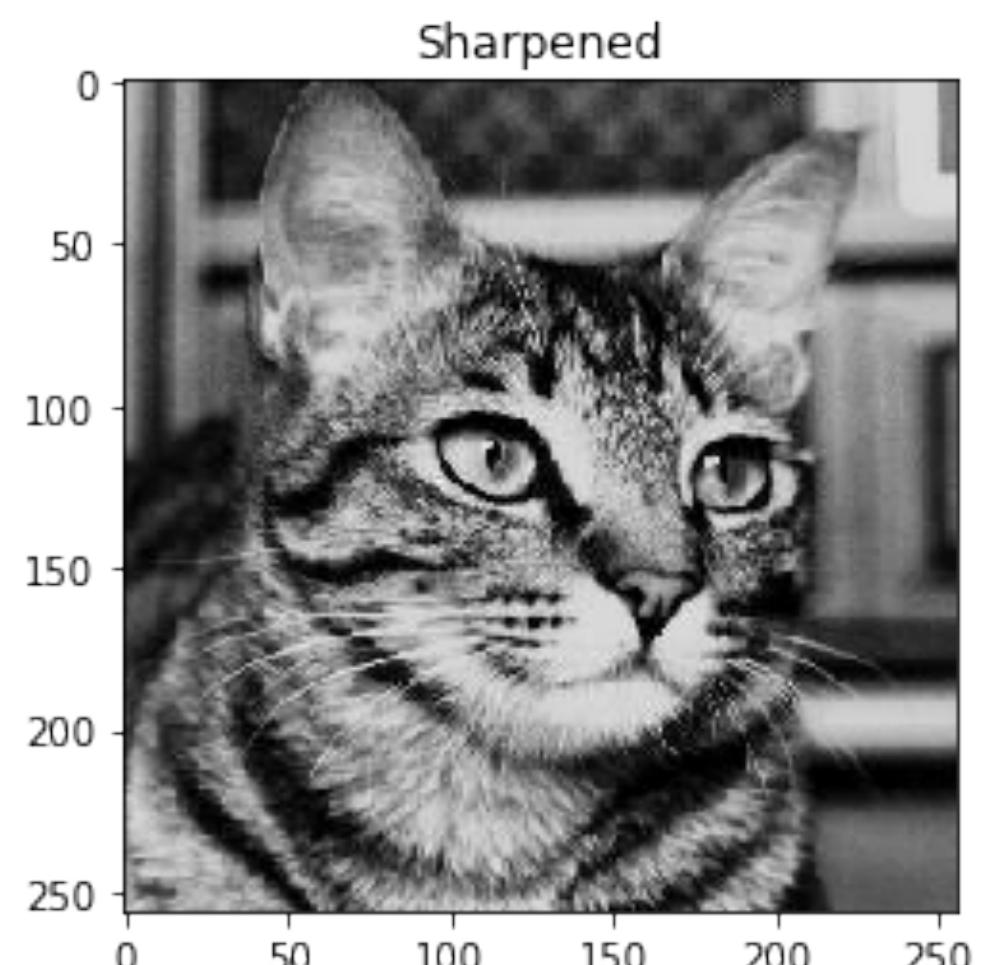
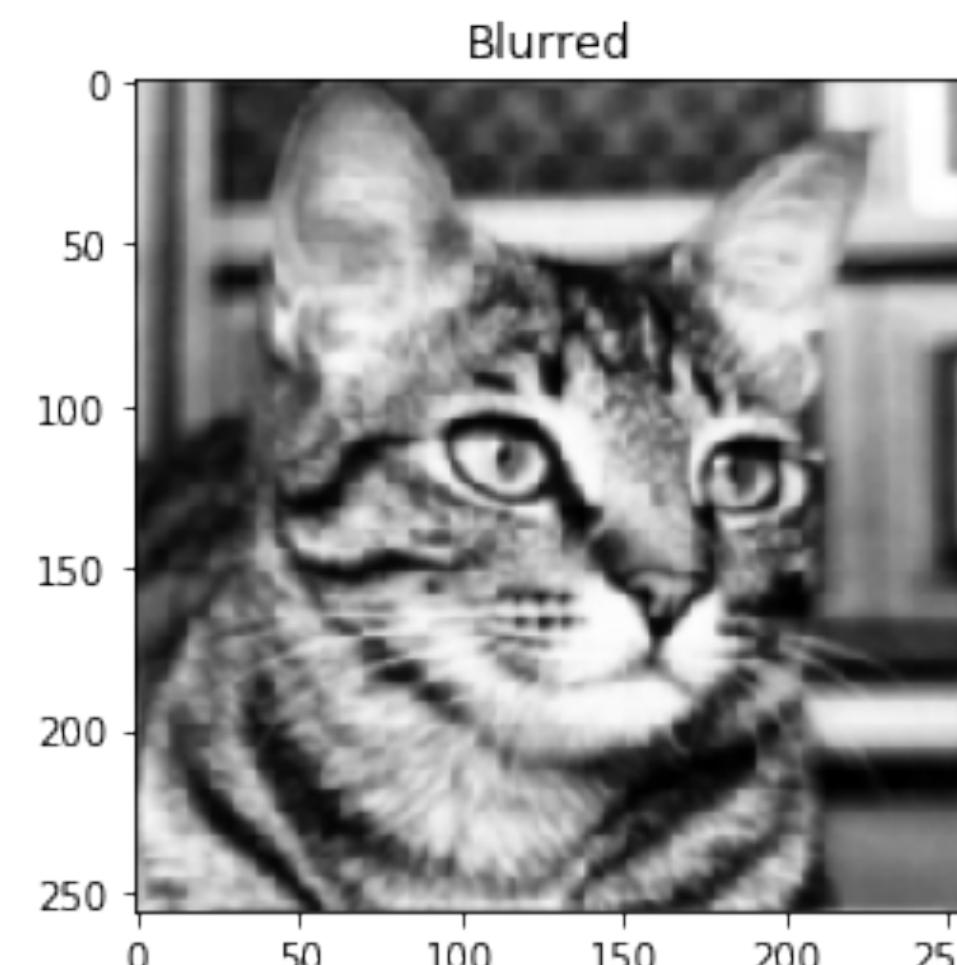
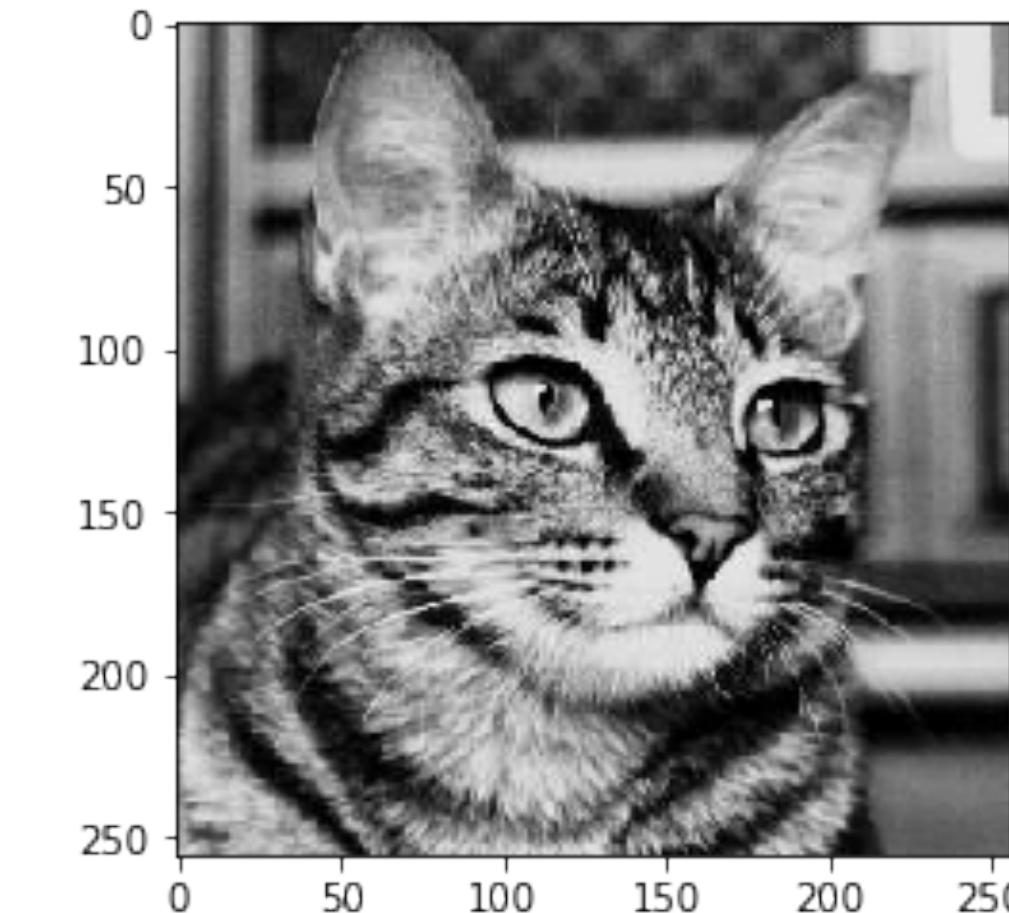
Jean Baptiste
Joseph Fourier



His Fourier Transform

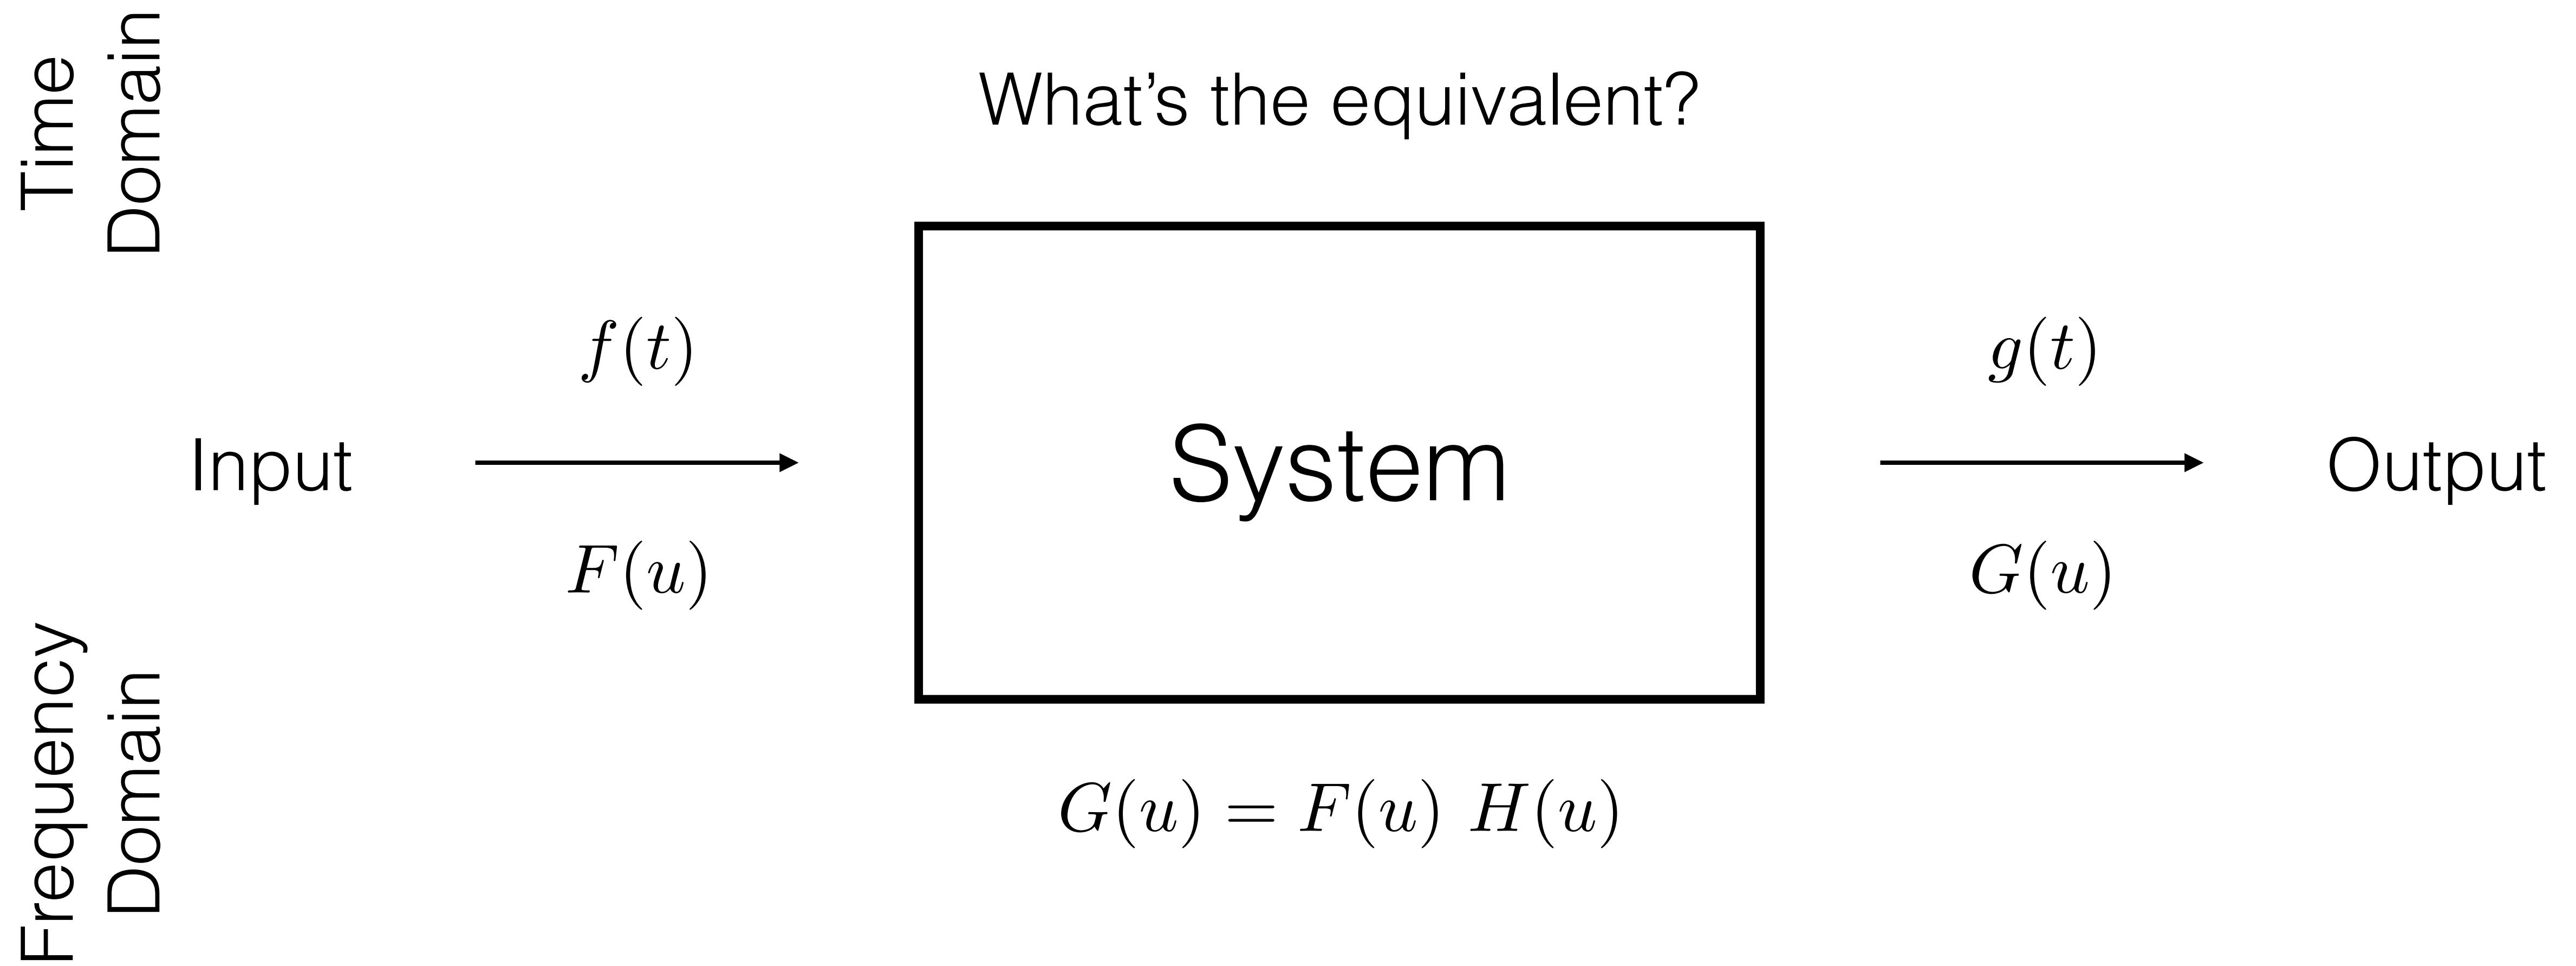
Filtering Images in the Frequency Domain

- Just like we can filter 1-D signals, we can also filter 2-D images
 - Fourier transform the image
 - Multiply by the filter
 - Inverse Fourier Transform
- Low-pass filtering = blurring
- High-pass filtering = edge detection
- High-boost filtering = sharpening



A really important detour...

Time Domain vs. Frequency Domain



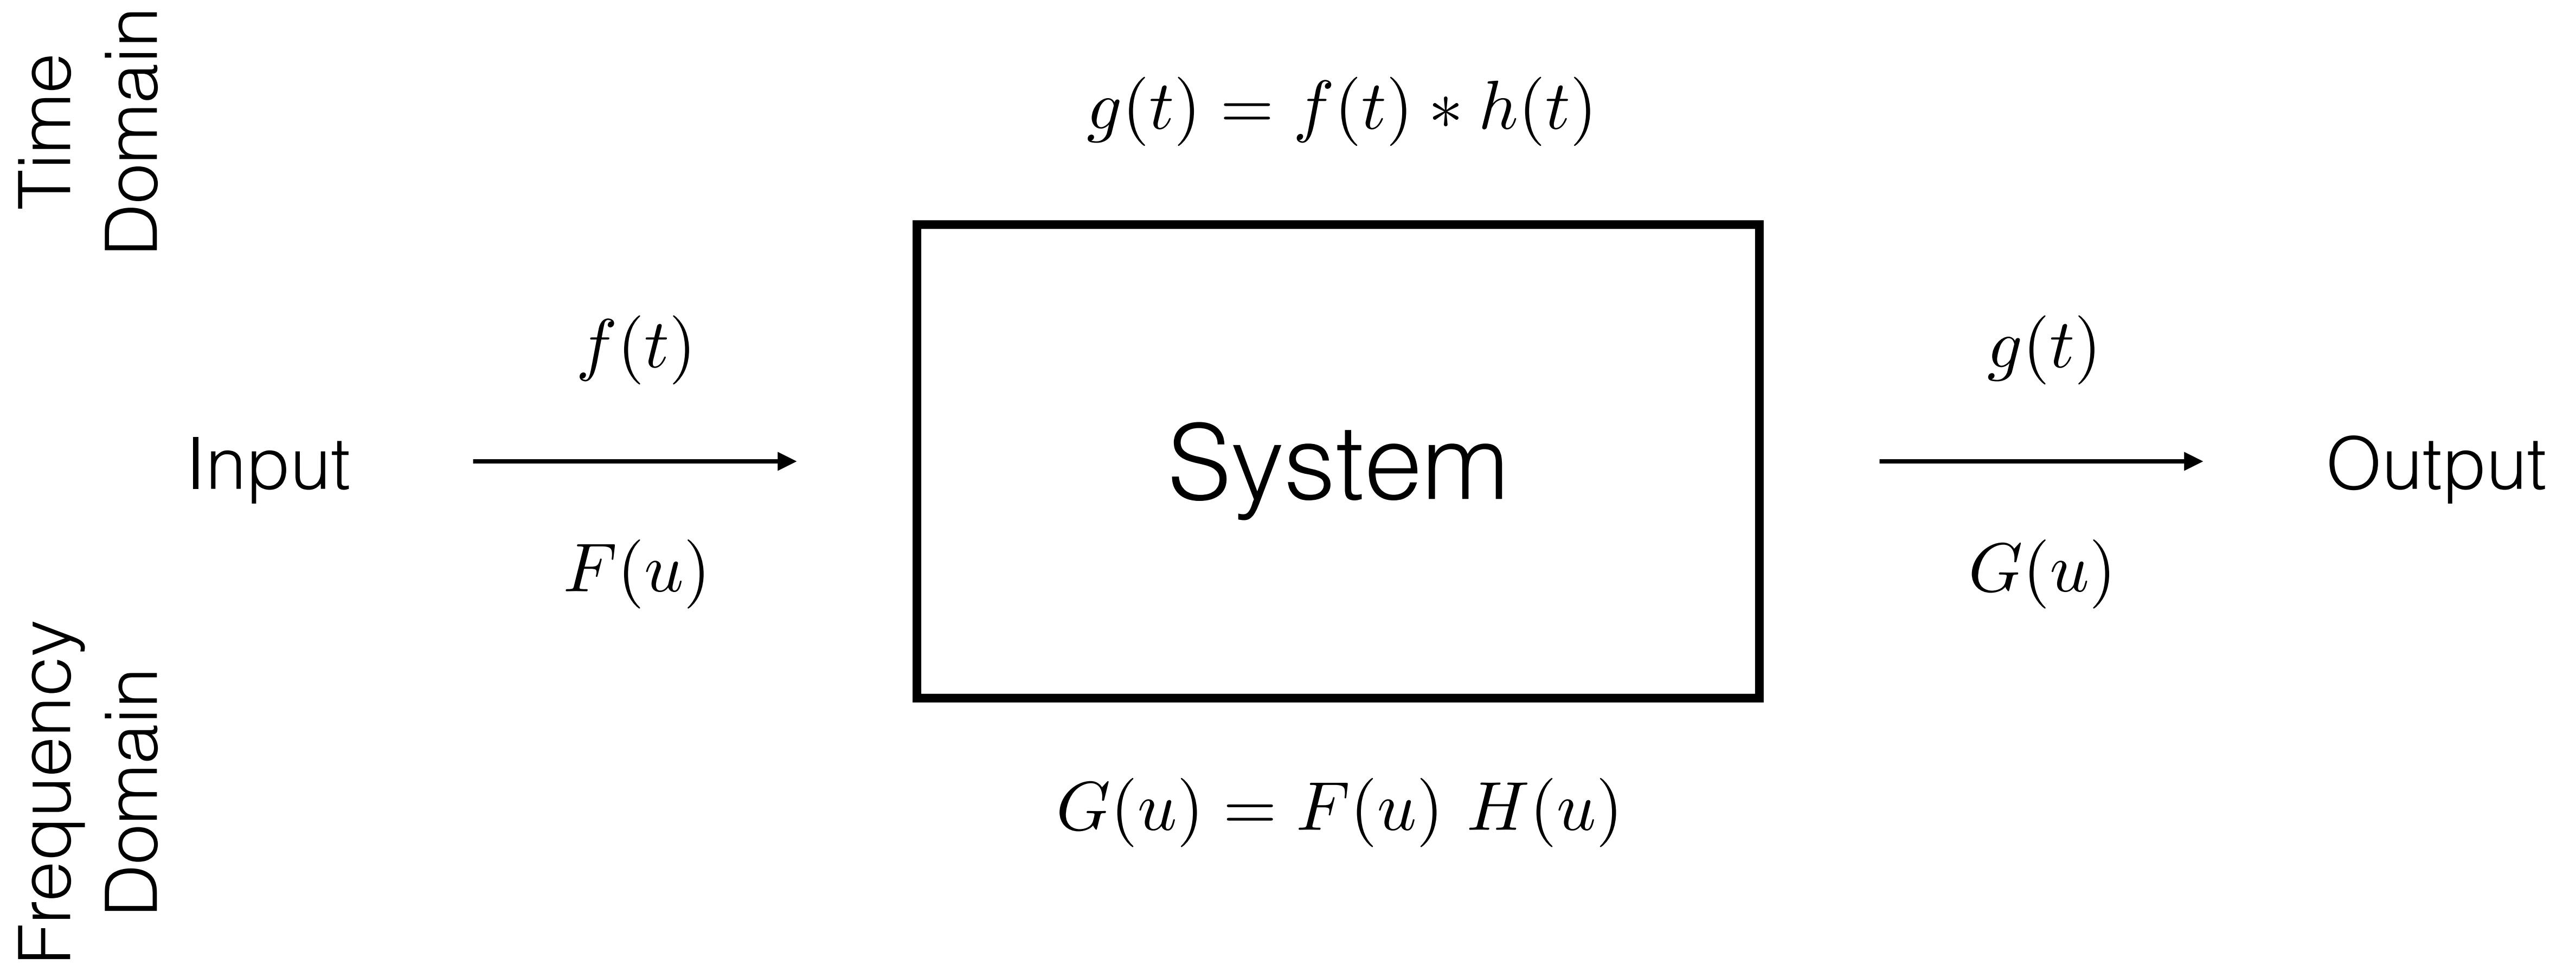
The Convolution Theorem

- **Convolution in one domain is multiplication in the other**
and vice versa
 - Convolution in the time domain and multiplication in the frequency domain are equivalent
 - Multiplying two signals in the time domain convolves their FTs (useful for lots of things)
- Relationship between convolution kernels and transfer functions?

$$f(t) * g(t) = \mathcal{F}(f(t)) \mathcal{F}(g(t))$$

$$f(t) g(t) = \mathcal{F}(f(t)) * \mathcal{F}(g(t))$$

Time Domain vs. Frequency Domain

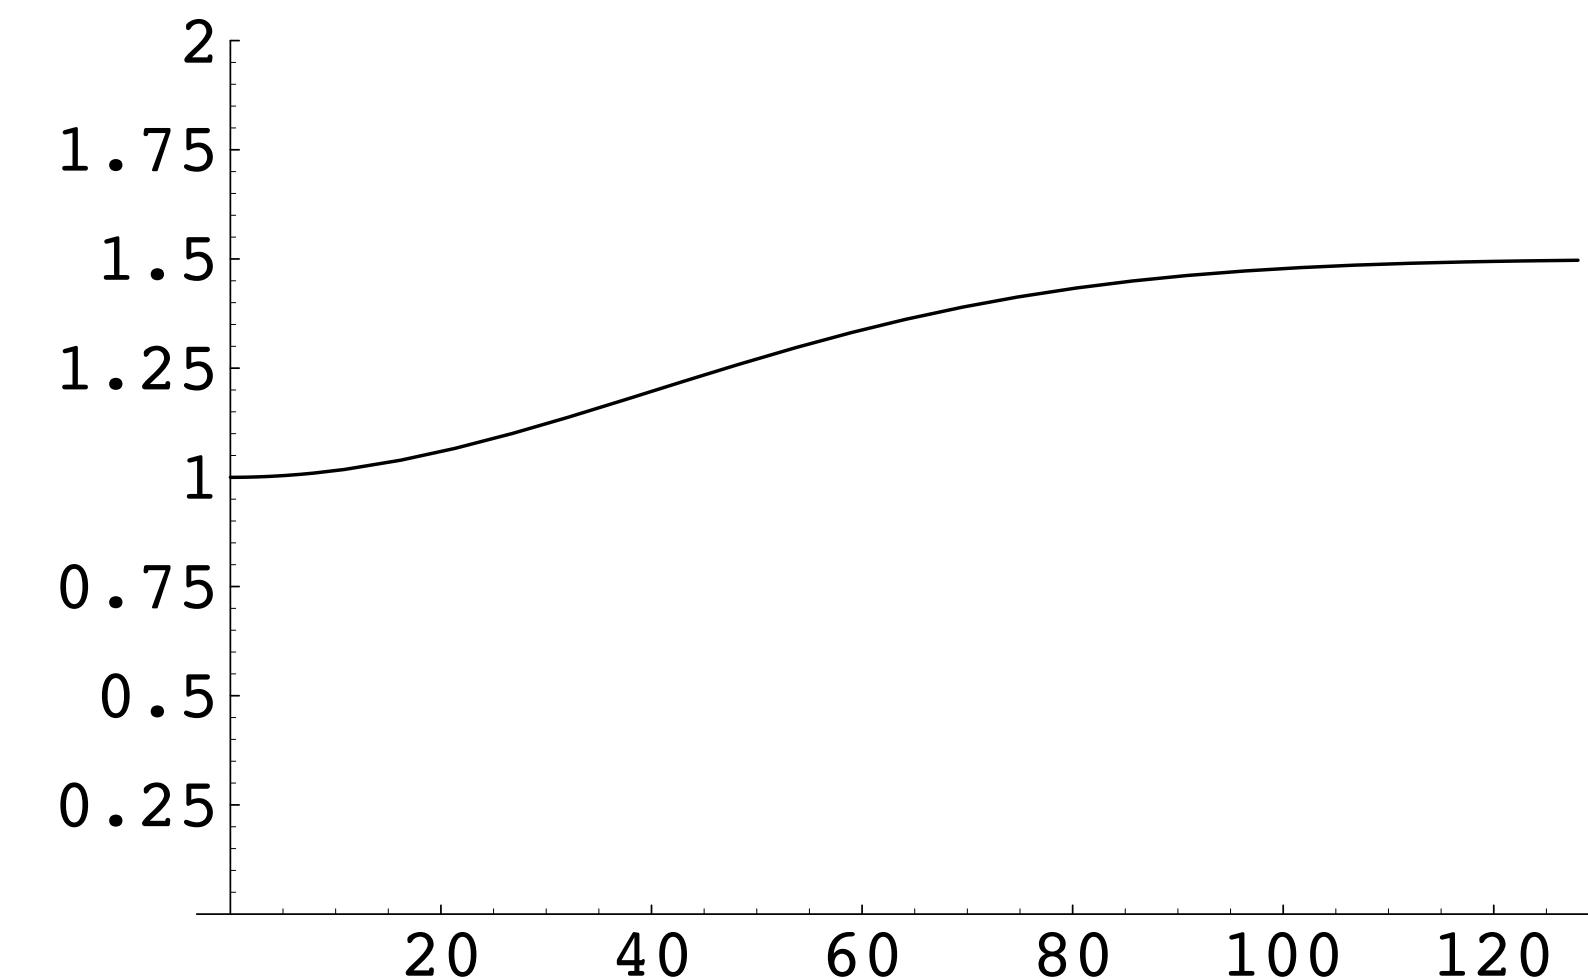


Example: High Boost Filters

- Can construct a high-boost filter by
 - Low-pass filtering the signal
 - Subtracting it from the original
 - Adding a fraction of that back to the original

$$H(u) = 1 + \alpha(1 - e^{-\frac{1}{2}u^2/u_c^2})$$

- Sound familiar?



The Need For Speed...

- Convolution of an $N \times N$ image with a $K \times K$ kernel:

$$O(N^2 K^2)$$

- FFT the signal, FFT the kernel, inverse FFT:

$$O(N^2 \log N)$$

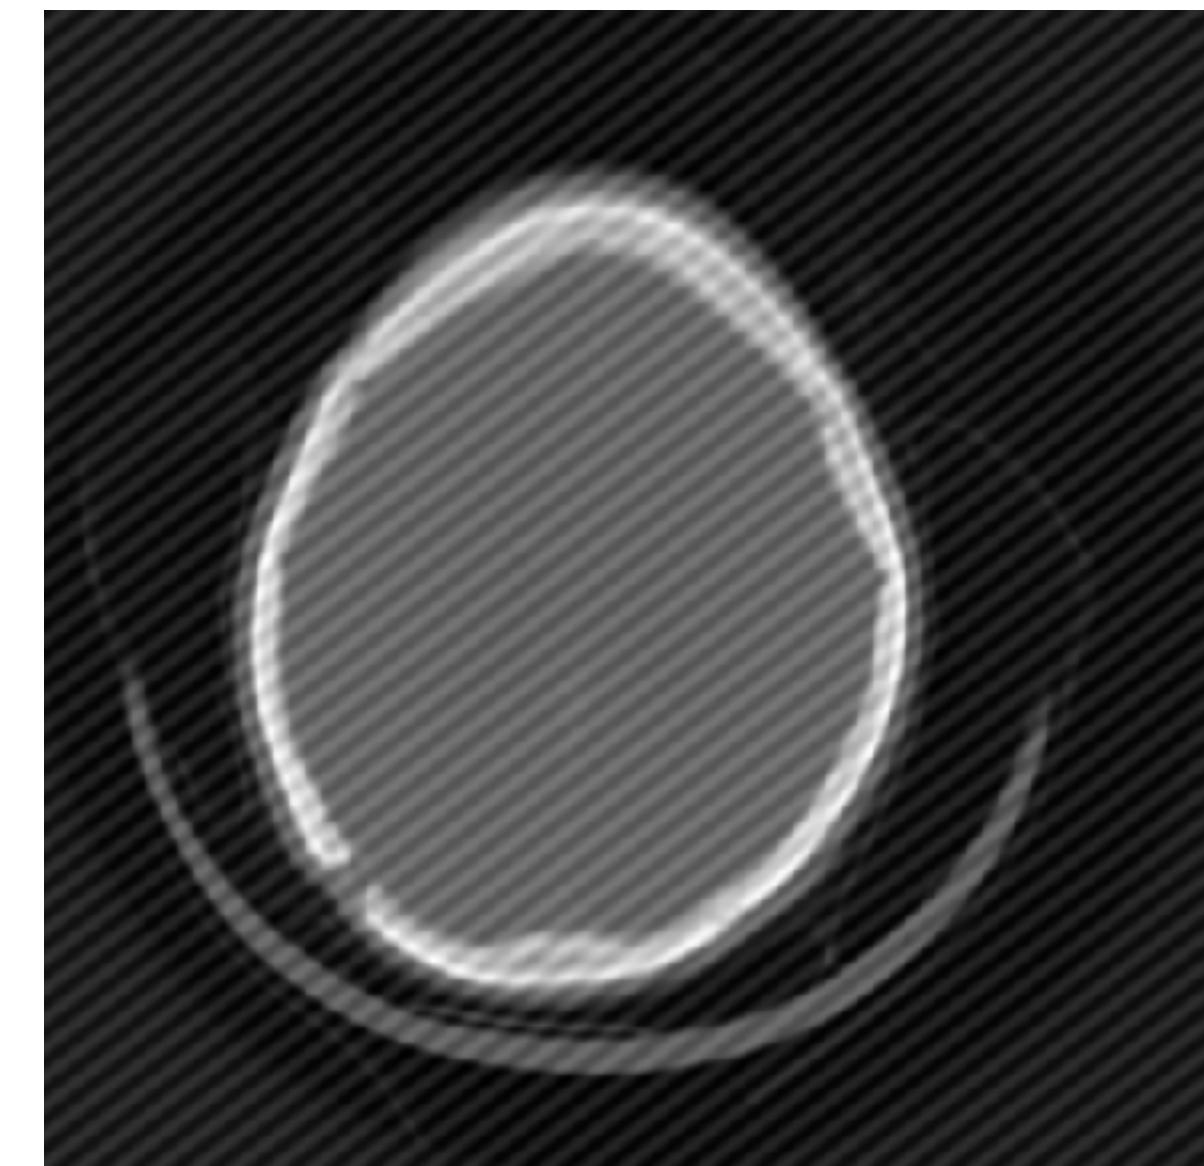
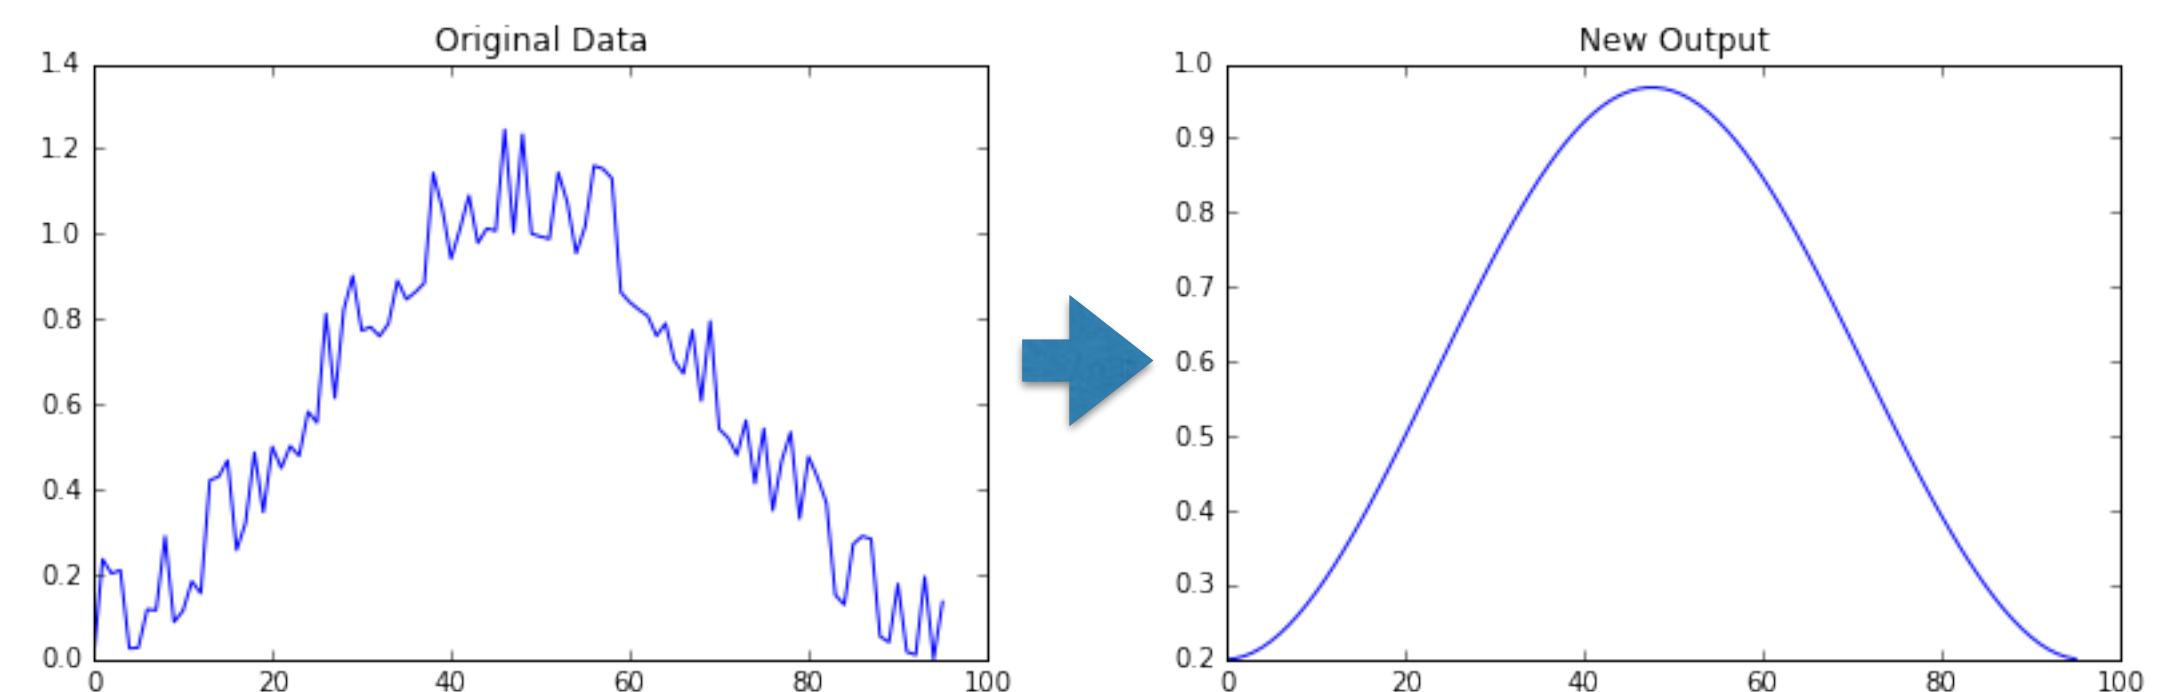
- For large kernels, doing it in the frequency domain can be faster!
- General rule of thumb, the break-even point is 7×7 kernels

Design vs. Implementation

- We can design in either domain
 - Spatial filters in (masks or kernels) in the time/spatial domain
 - Frequency filters in the frequency domain
- And we can implement them in either domain
 - Convolution in the time/spatial domain
 - Multiplication in the frequency domain

Lab 10

- Filtering 1-D signals
 - Synthetic example (noisy function)
 - Audio example
- Filtering 2-D images
 - Smoothing/sharpening
 - Isolating and removing a single-frequency interference pattern



Coming up...

- Other uses of these ideas in Computer Science
- Applications in Computer Graphics (and CS 455)
- Applications in Computer Vision (and CS 450)



Other Applications in CS

CS 355: Introduction to Graphics and Image Processing

Math Applications in CS

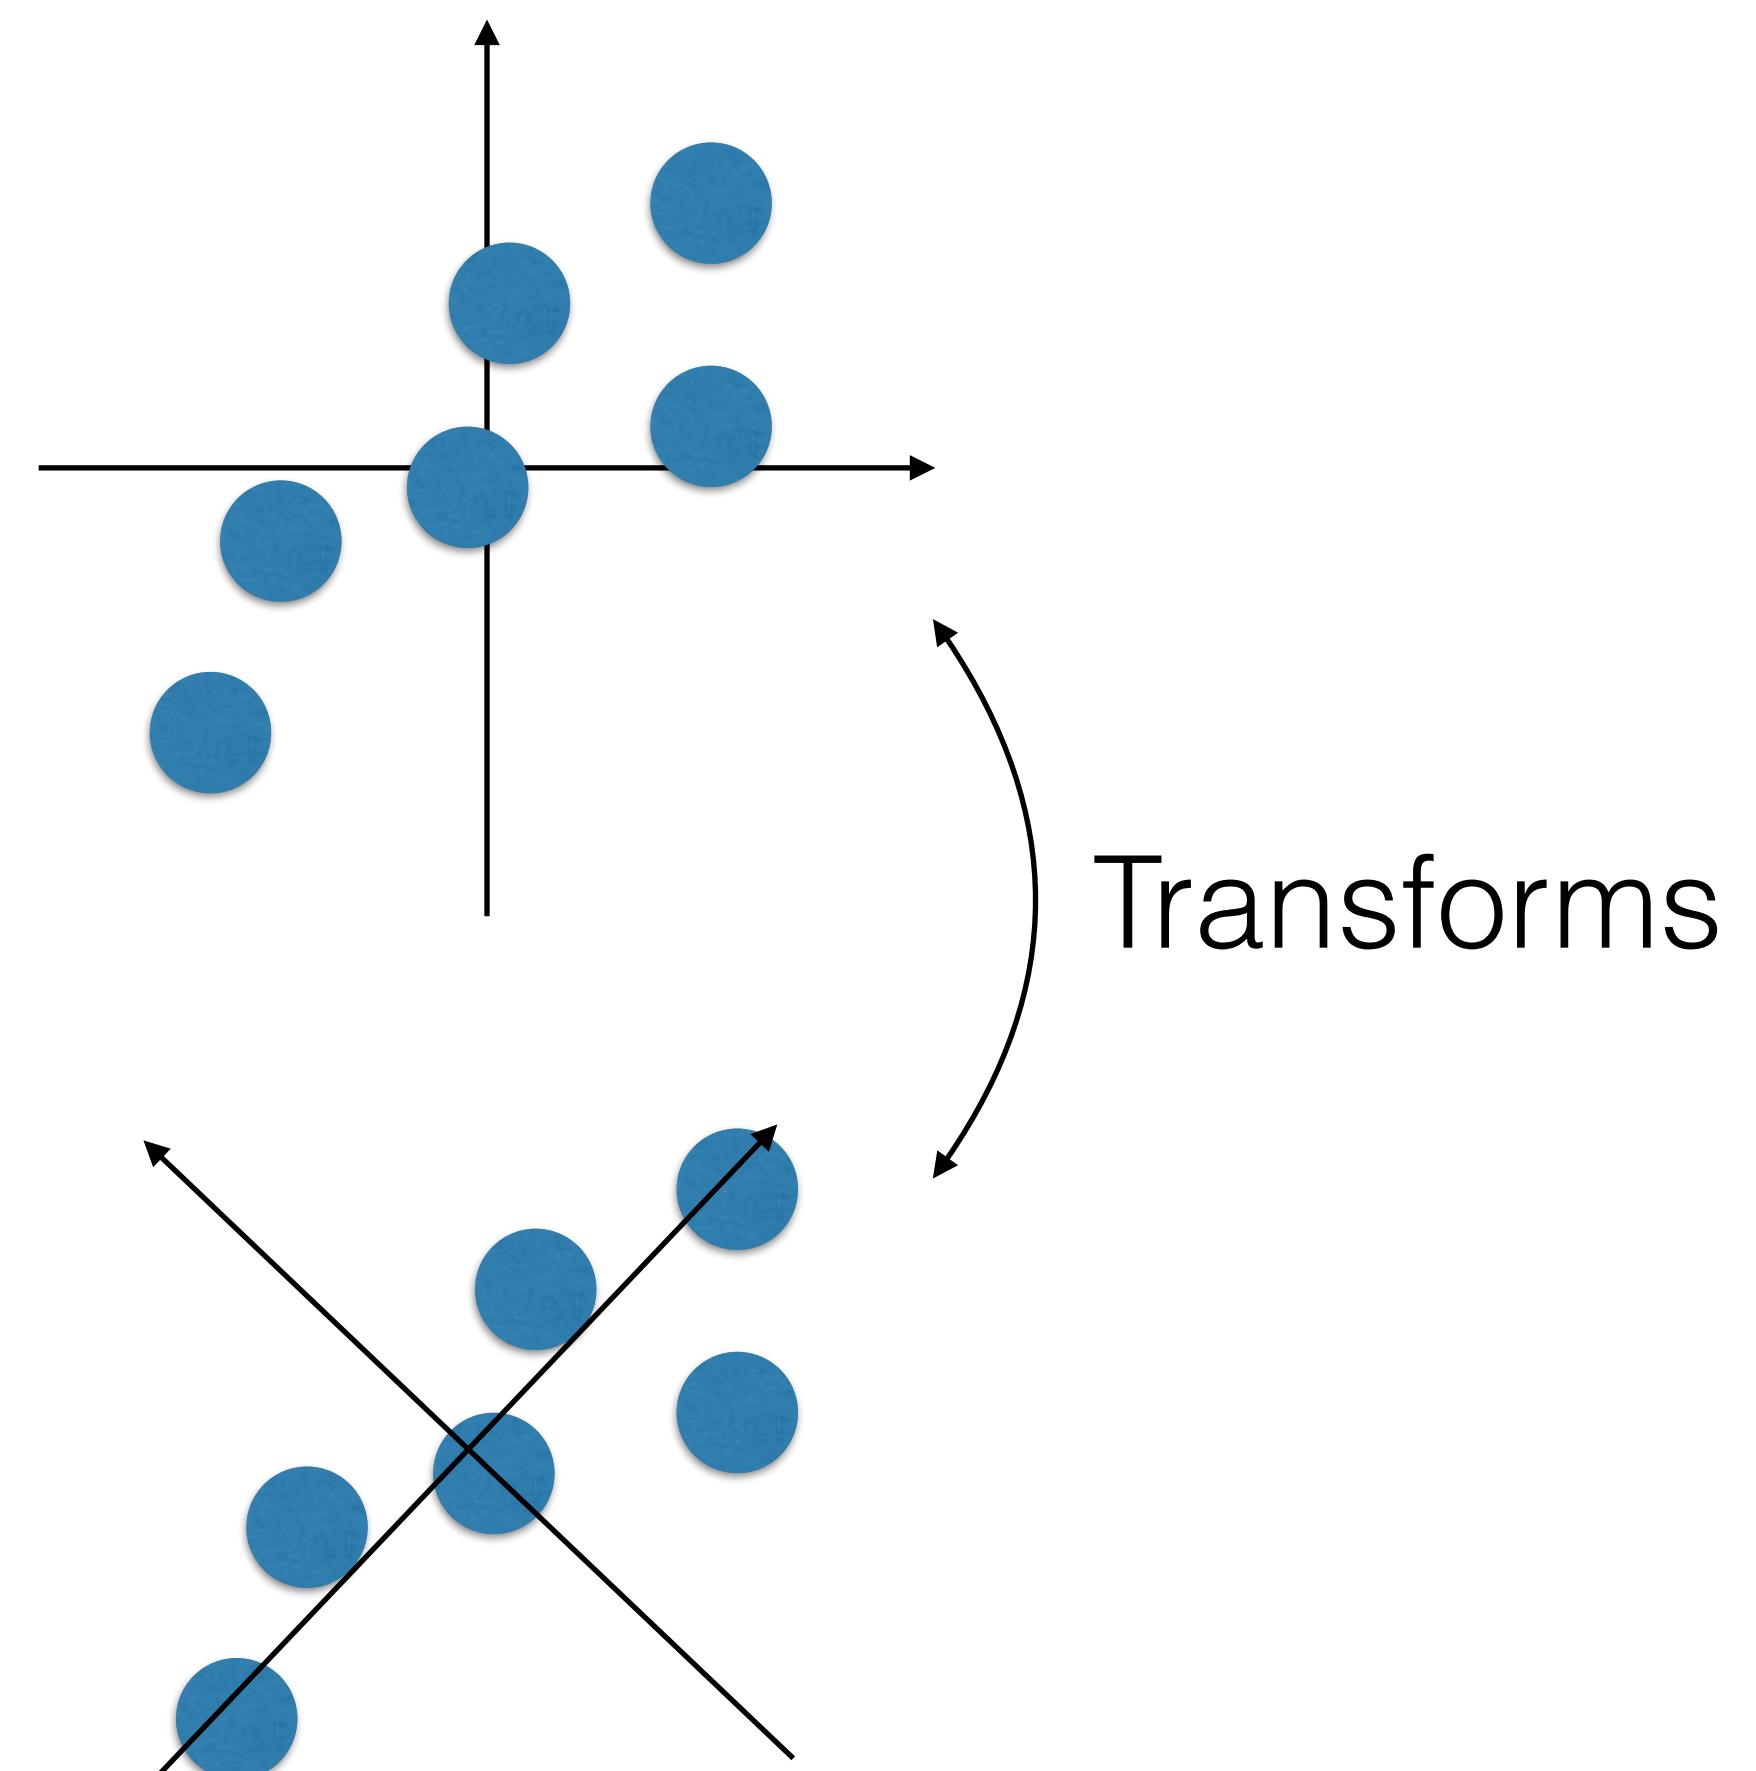
- All the “discrete math” from CS 235, 252, etc.
 - Sets, relations, functions, ...
- Linear algebra
 - Geometric transformations
 - Data transforms
 - Systems of equations
 - Eigensystems

Data as Vectors

Lots of things can be thought of as points/vectors in some space

Data Transforms

- Common pattern in data analysis:
 - Represent data as vectors
 - Convert to a different coordinate system
 - Analyze (or change!) while in that coordinate system
 - Convert back (if needed)



Data Transforms

- Same form as any other change of coordinates
 - Transform using dot products
 - Convert back using weighted sum

$$\begin{array}{c} \text{transformed} \\ \text{data} \\ \downarrow \\ g[i] = \sum_k f[k] e_i[k] \end{array} \quad \begin{array}{c} \text{original} \\ \text{data} \\ \downarrow \\ f[k] = \sum_i g[i] e_i[k] \end{array} \quad \begin{array}{c} \text{basis} \\ \text{set} \\ \downarrow \end{array}$$

Fourier Analysis

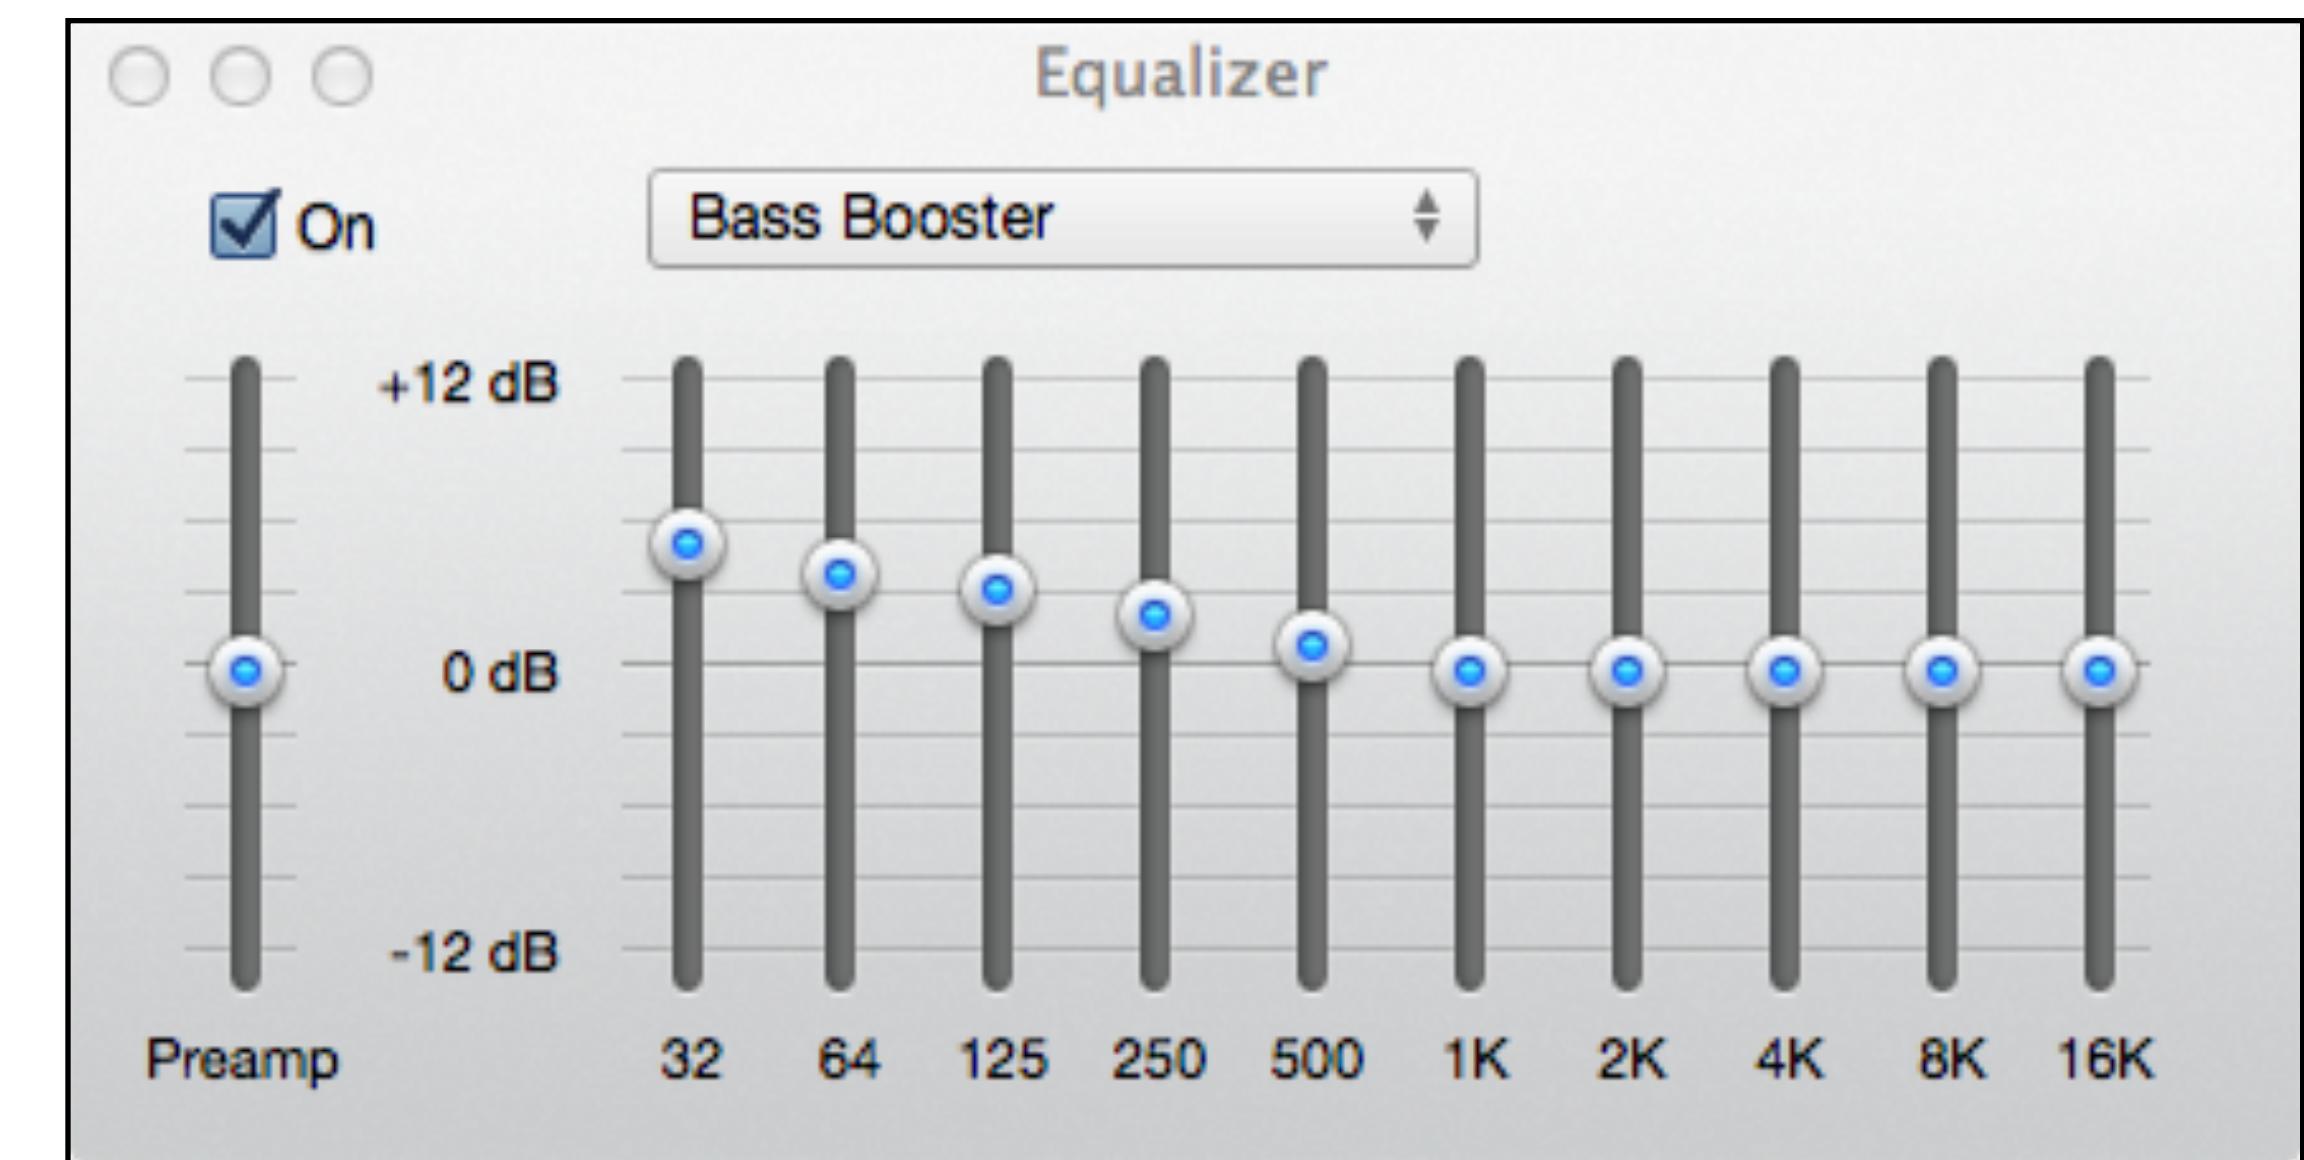
- Sampled sines and cosines of different frequencies form an orthonormal basis set
- Can decompose any waveform into a weighted sum of sines and cosines of different frequencies
- Great for analysis, manipulation, etc.

$$c[u] = \frac{1}{N} \sum_{u=0}^{N-1} f[t] \cos(2\pi ut/N)$$

$$s[u] = \frac{1}{N} \sum_{u=0}^{N-1} f[t] \sin(2\pi ut/N)$$

Frequency Manipulation

- Can use frequency-based representation to do manipulation
 - Boost bass/treble
 - Boost typical range of human speaking (hearing aides do this)
 - Suppress unwanted sounds
 - Smoothing / sharpening
 - And lots more...



Audio Compression

- Fact: your ear doesn't hear all frequencies equally well
(and it's different for everybody)
- Idea: don't not spend as many bits of precision on the ones we don't as hear well anyway

Audio Compression

- Compression:
 - Convert to a frequency-based representation
 - Use more bits to store the coefficients for the frequencies we hear better; fewer for the ones we don't hear well
 - Store in this form
- Decompression:
 - Use lossy coefficients and convert back to a PCM representation
 - Play!

This is how MP3 compression works!

Image Compression

- Can we do something similar with images?
- Fact: your eye is less accurate in sensing very rapid changes in brightness across an image (fine texture)
- Idea: use the same approach in the 2D frequency domain for images
- This is the basis of JPEG
(uses Discrete Cosine Transform instead of Fourier)

Classification

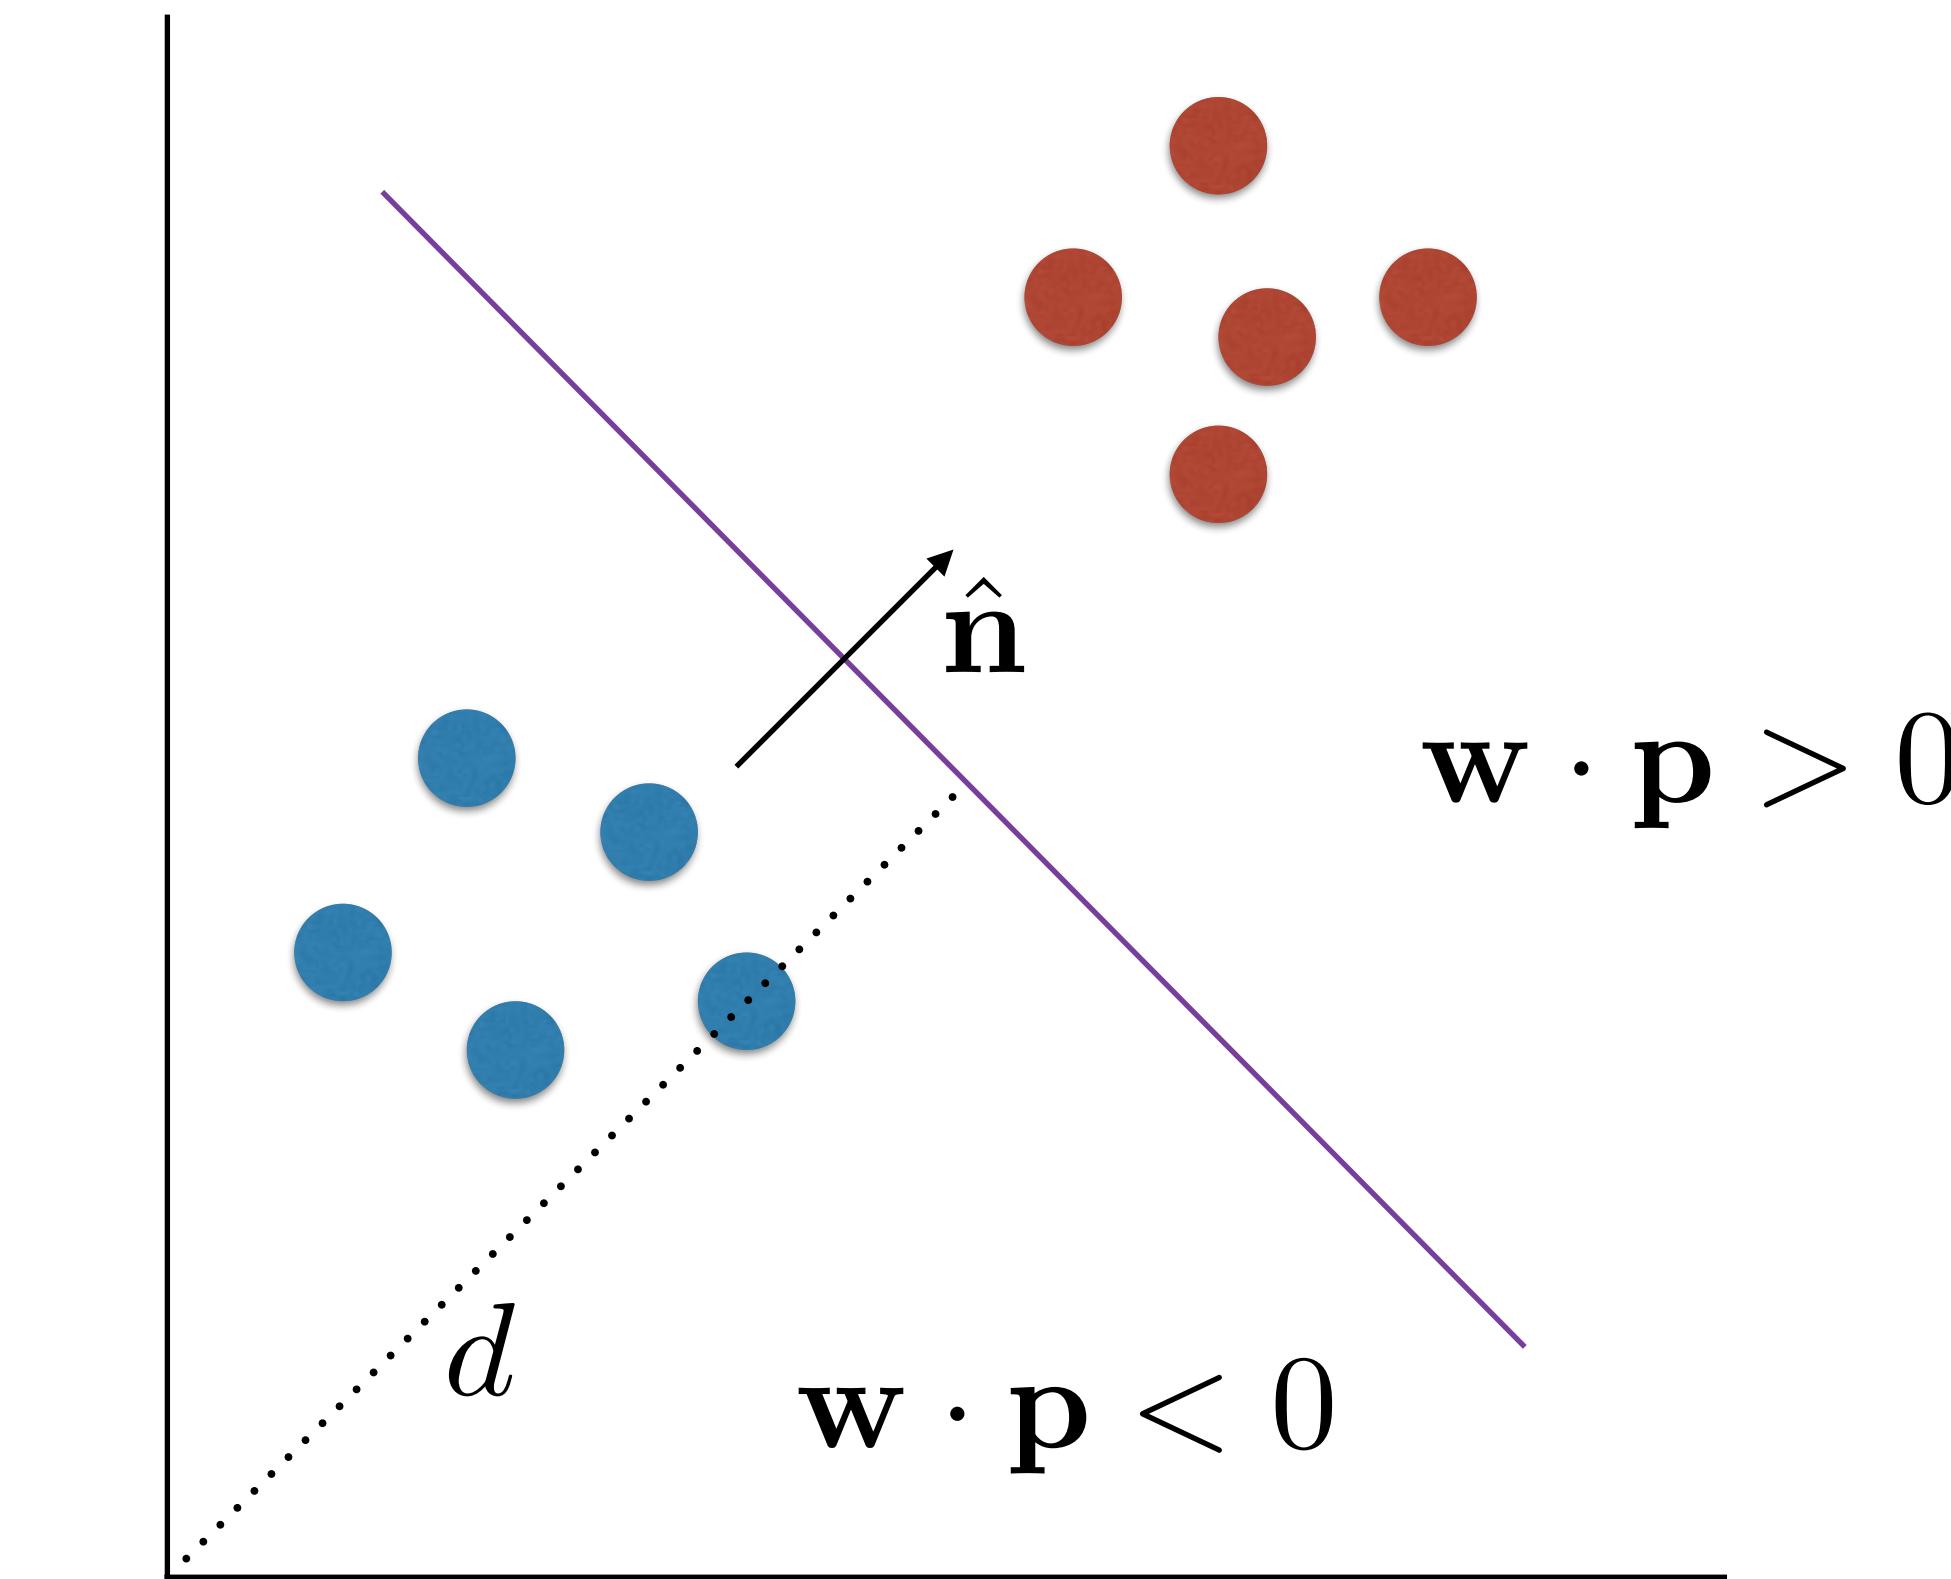
- Simple problem:
 - Two classes of things, with lots of examples of each
 - New thing — what kind is it?
- Approach:
 - Measure “features” of the things
 - Put features together in a vector
 - Look at the problem geometrically
 - Changing the coordinate system can make a huge difference!

(basis for pattern recognition, machine learning, other AI, ...)

Classification

$$\mathbf{p} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \\ 1 \end{bmatrix}$$

$$\mathbf{w} = \begin{bmatrix} n_1 \\ n_2 \\ \vdots \\ n_k \\ -d \end{bmatrix}$$



Classification

- But what if it's complicated, and you can't easily solve for \mathbf{w} ?
- Can you iteratively tweak the values in \mathbf{w} until you "get it right"?
- The entries of \mathbf{w} act as "weights" in a weighted combination of features, so it's called a *weight vector*

$$\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_k \\ w_{k+1} \end{bmatrix} \quad \mathbf{p} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \\ 1 \end{bmatrix}$$

$$\mathbf{w} \cdot \mathbf{p} = \sum_{i=1}^k w_i x_i + w_{k+1}$$

This is basically the heart of what neural networks do!

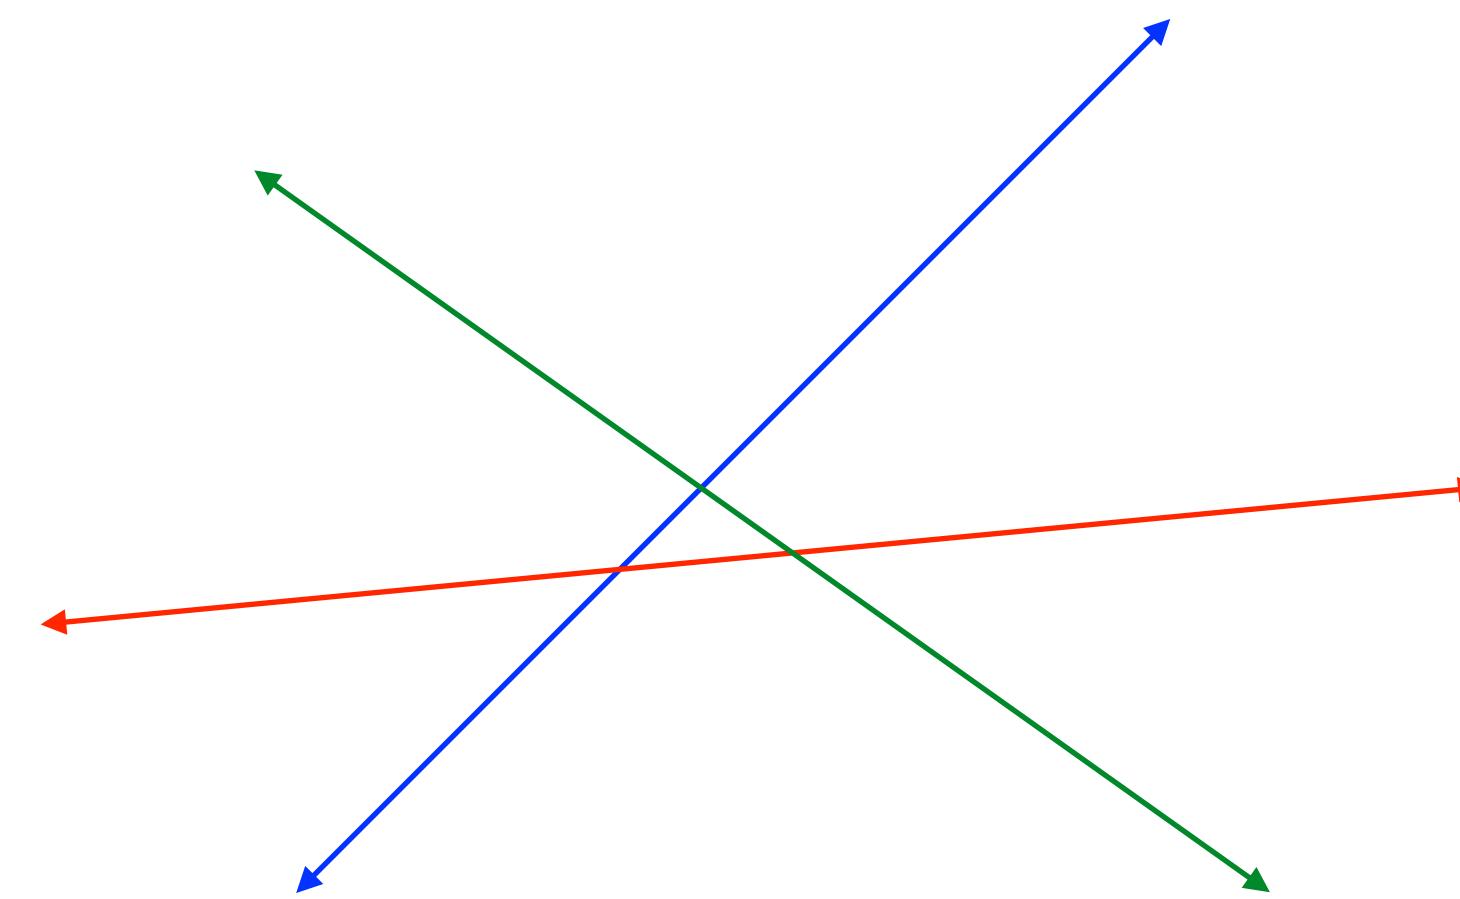
Systems of Equations

- If you need to solve for n unknowns, you need n equations, right?
- What if the data is noisy?
- Idea: get more data and let the noise “average out”
- But now there are too many equations!

$$\mathbf{Ax} = \mathbf{b}$$

Overconstrained Systems

- When you have too many equations, there may not be a solution
- Idea: get as close as possible to fitting all of the equations
- If you use a squared-error metric, this leads to a *least-squares solution*



$$\text{minimize } \|\mathbf{Ax} - \mathbf{b}\|^2$$

Eigensystems

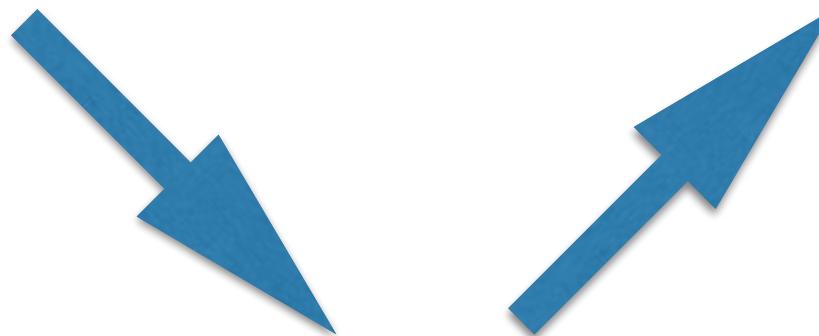
- The actions of some matrices are sometimes described more easily when converted to another coordinate system
- Some matrices are *pure scaling* along certain key directions
- A useful tool for analyzing these are *eigensystems*
 - Eigenvectors - directions of scaling
 - Eigenvalues - amount of scaling in each direction

Eigensystems

What does this matrix do?

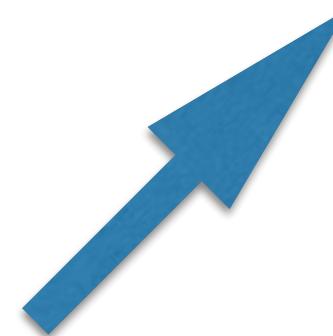
$$\begin{bmatrix} 3 & 1 \\ 1 & 3 \end{bmatrix}$$

Scales by 4 in the direction [1,1]
and by 2 in the direction [-1,1]



Eigenvectors:

$$\begin{bmatrix} 1 \\ 1 \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$



Eigenvalues:

$$4, 2$$

Coming up...

- Applications in Computer Vision (and CS 450) — Dr. Farrell
- Applications in Computer Graphics (and CS 455) — Dr. Egbert